

Introduction to Deep Learning

Pierre Colombo

pierre.colombo@centralesupelec.fr

MICS
CentraleSupélec Université Paris-Saclay



CentraleSupélec

université
PARIS-SACLAY



I. Theoretical Class

1. Introduction

2. Deep Learning Architecture

3. Performing the learning phase

4. From theory to practice

II. Deep Learning in action

Curriculum Vitae

Education & Diploma

- 2018.** Diplôme D'Ingénieur, Supelec, *France*
- 2018.** MSc in Computer Science, Ecole Polytechnique Fédérale de Lausanne, *Switzerland*
- 2021.** PhD in Computer Science, Telecom Paris, Institut Polytechnique de Paris, *France*
Title: [Learning to represent and generate text using information measures](#)
- 2022 .** *Post-Doctoral Researcher*, CentraleSupelec, Laboratoire des Signaux et Systèmes, CNRS, Université Paris-Saclay, CentraleSupelec
- 2022 - Present.** *Invited Member (external) of Comète Inria Team*, Laboratoire d'informatique de l'École polytechnique, Ecole polytechnique

Present Occupation

- 2022 - Present.** Assistant Prof au MICS (CentraleSupelec)

Past Research Experiences



2016 (3 months). Research Intern, Procter & Gamble, **Kronberg, Germany**



2016 - 2017 (6 months). Lab Assistant, EPFL Laboratory for Quantum Gases, **Lausanne, Switzerland**

Developing various sensors algorithms (Python/C++) for a quantum experiment on fermi gas.



2017 (6 months). Research Intern, Swisscom Digital Lab, **Lausanne, Switzerland**

Research on NLP: Information Retrieval for a product oriented chatbot



2017 - 2018 (6 months). Research Intern, IBM Research, **Zürich, Switzerland**

Research on deep learning for cancer treatment



2018 (6 months). Lab Associate, Disney Research, **Los Angeles, USA**

Research on Natural Language Generation, 1 workshop paper, 1 paper accepted at NAACL, 1 US Patent



2018 - 2021 (36 months). PhD Student (CIFRE), Telecom Paris (S2A) & IBM France, **Paris, France**

11 first author papers published in A conferences (A*), 4 papers accepted as oral presentation (top 5%), 1 best student paper award

What Is Deep Learning?

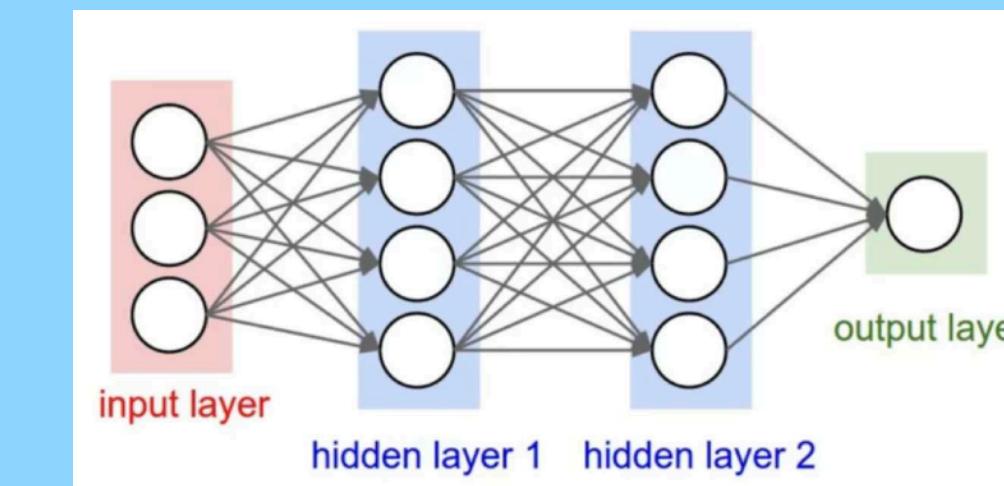
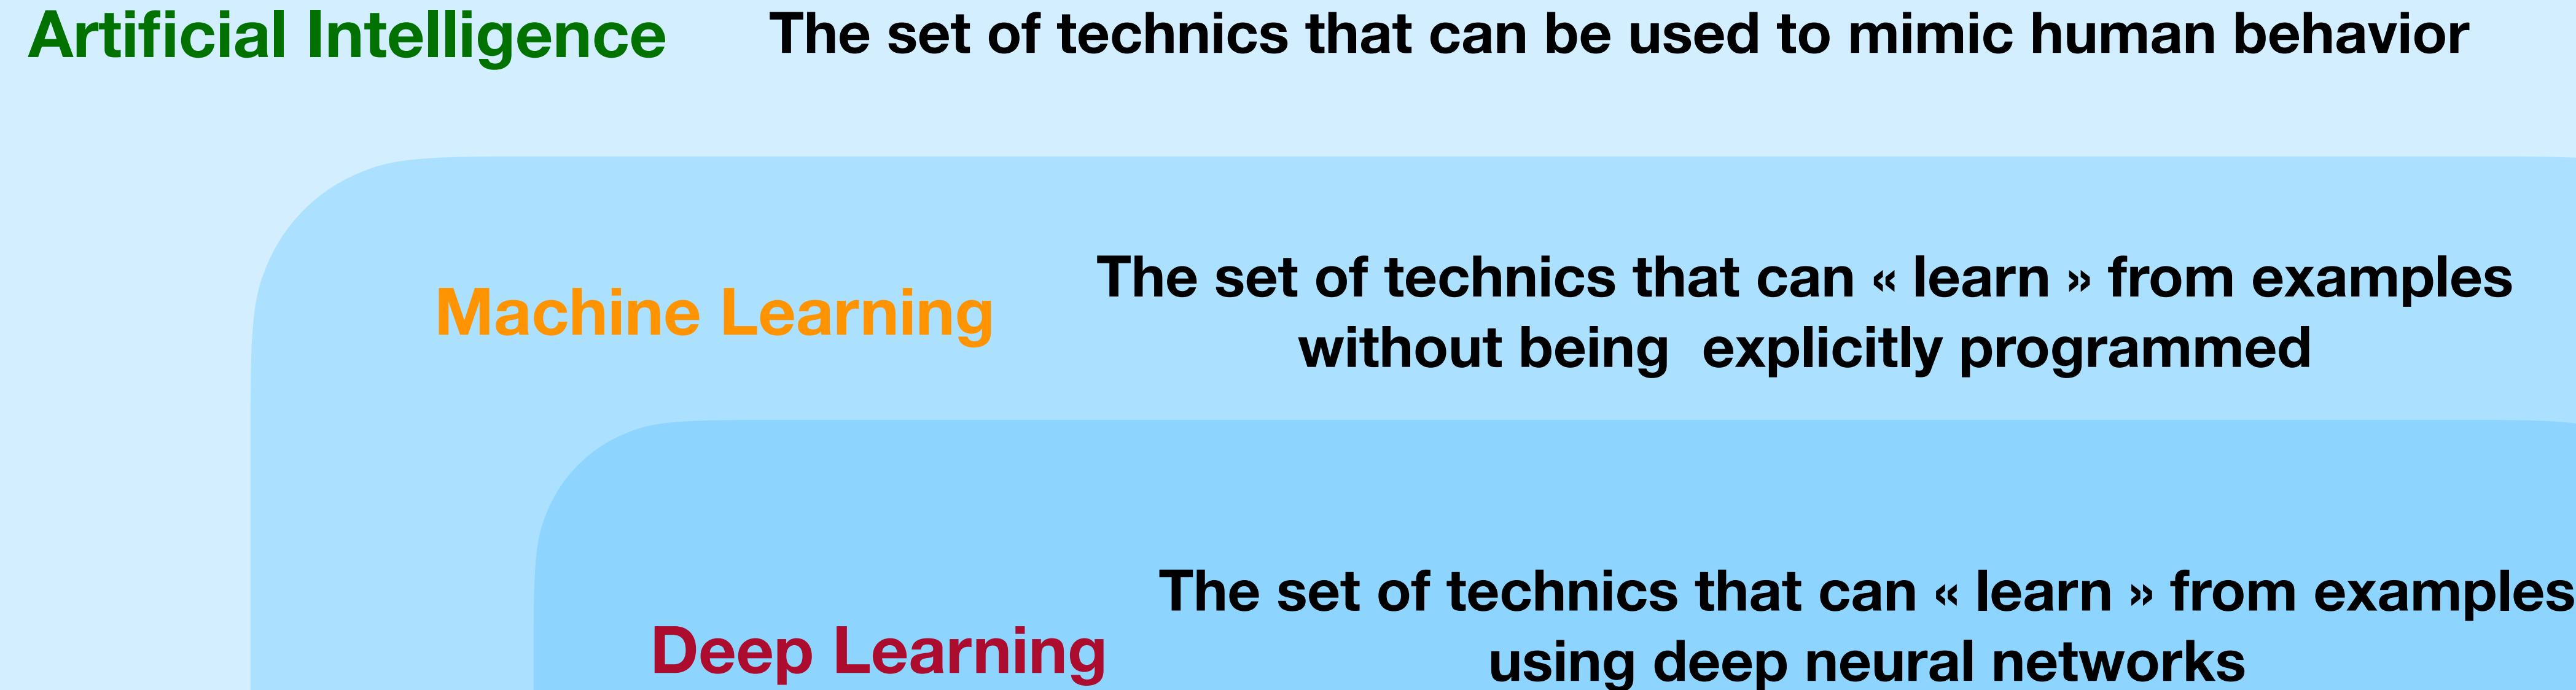
What Is Deep Learning?

Artificial Intelligence **The set of technics that can be used to mimic human behavior**

What Is Deep Learning?

Artificial Intelligence	The set of technics that can be used to mimic human behavior
Machine Learning	The set of technics that can « learn » from examples without being explicitly programmed

What Is Deep Learning?



What are the goals of this class?

What are the goals of this class?

What to expect?

What are the goals of this class?

What to expect?

- 1. Have a **high level view** of how neural networks works**
- 2. Getting **familiar with general concepts** such as overfitting, regularization,....**
- 3. **Code and train** your first neural network**

What are the goals of this class?

What to expect?

1. Have a **high level view** of how neural networks works
2. Getting **familiar with general concepts** such as overfitting, regularization,....
3. **Code and train** your first neural network

What not to expect?

What are the goals of this class?

What to expect?

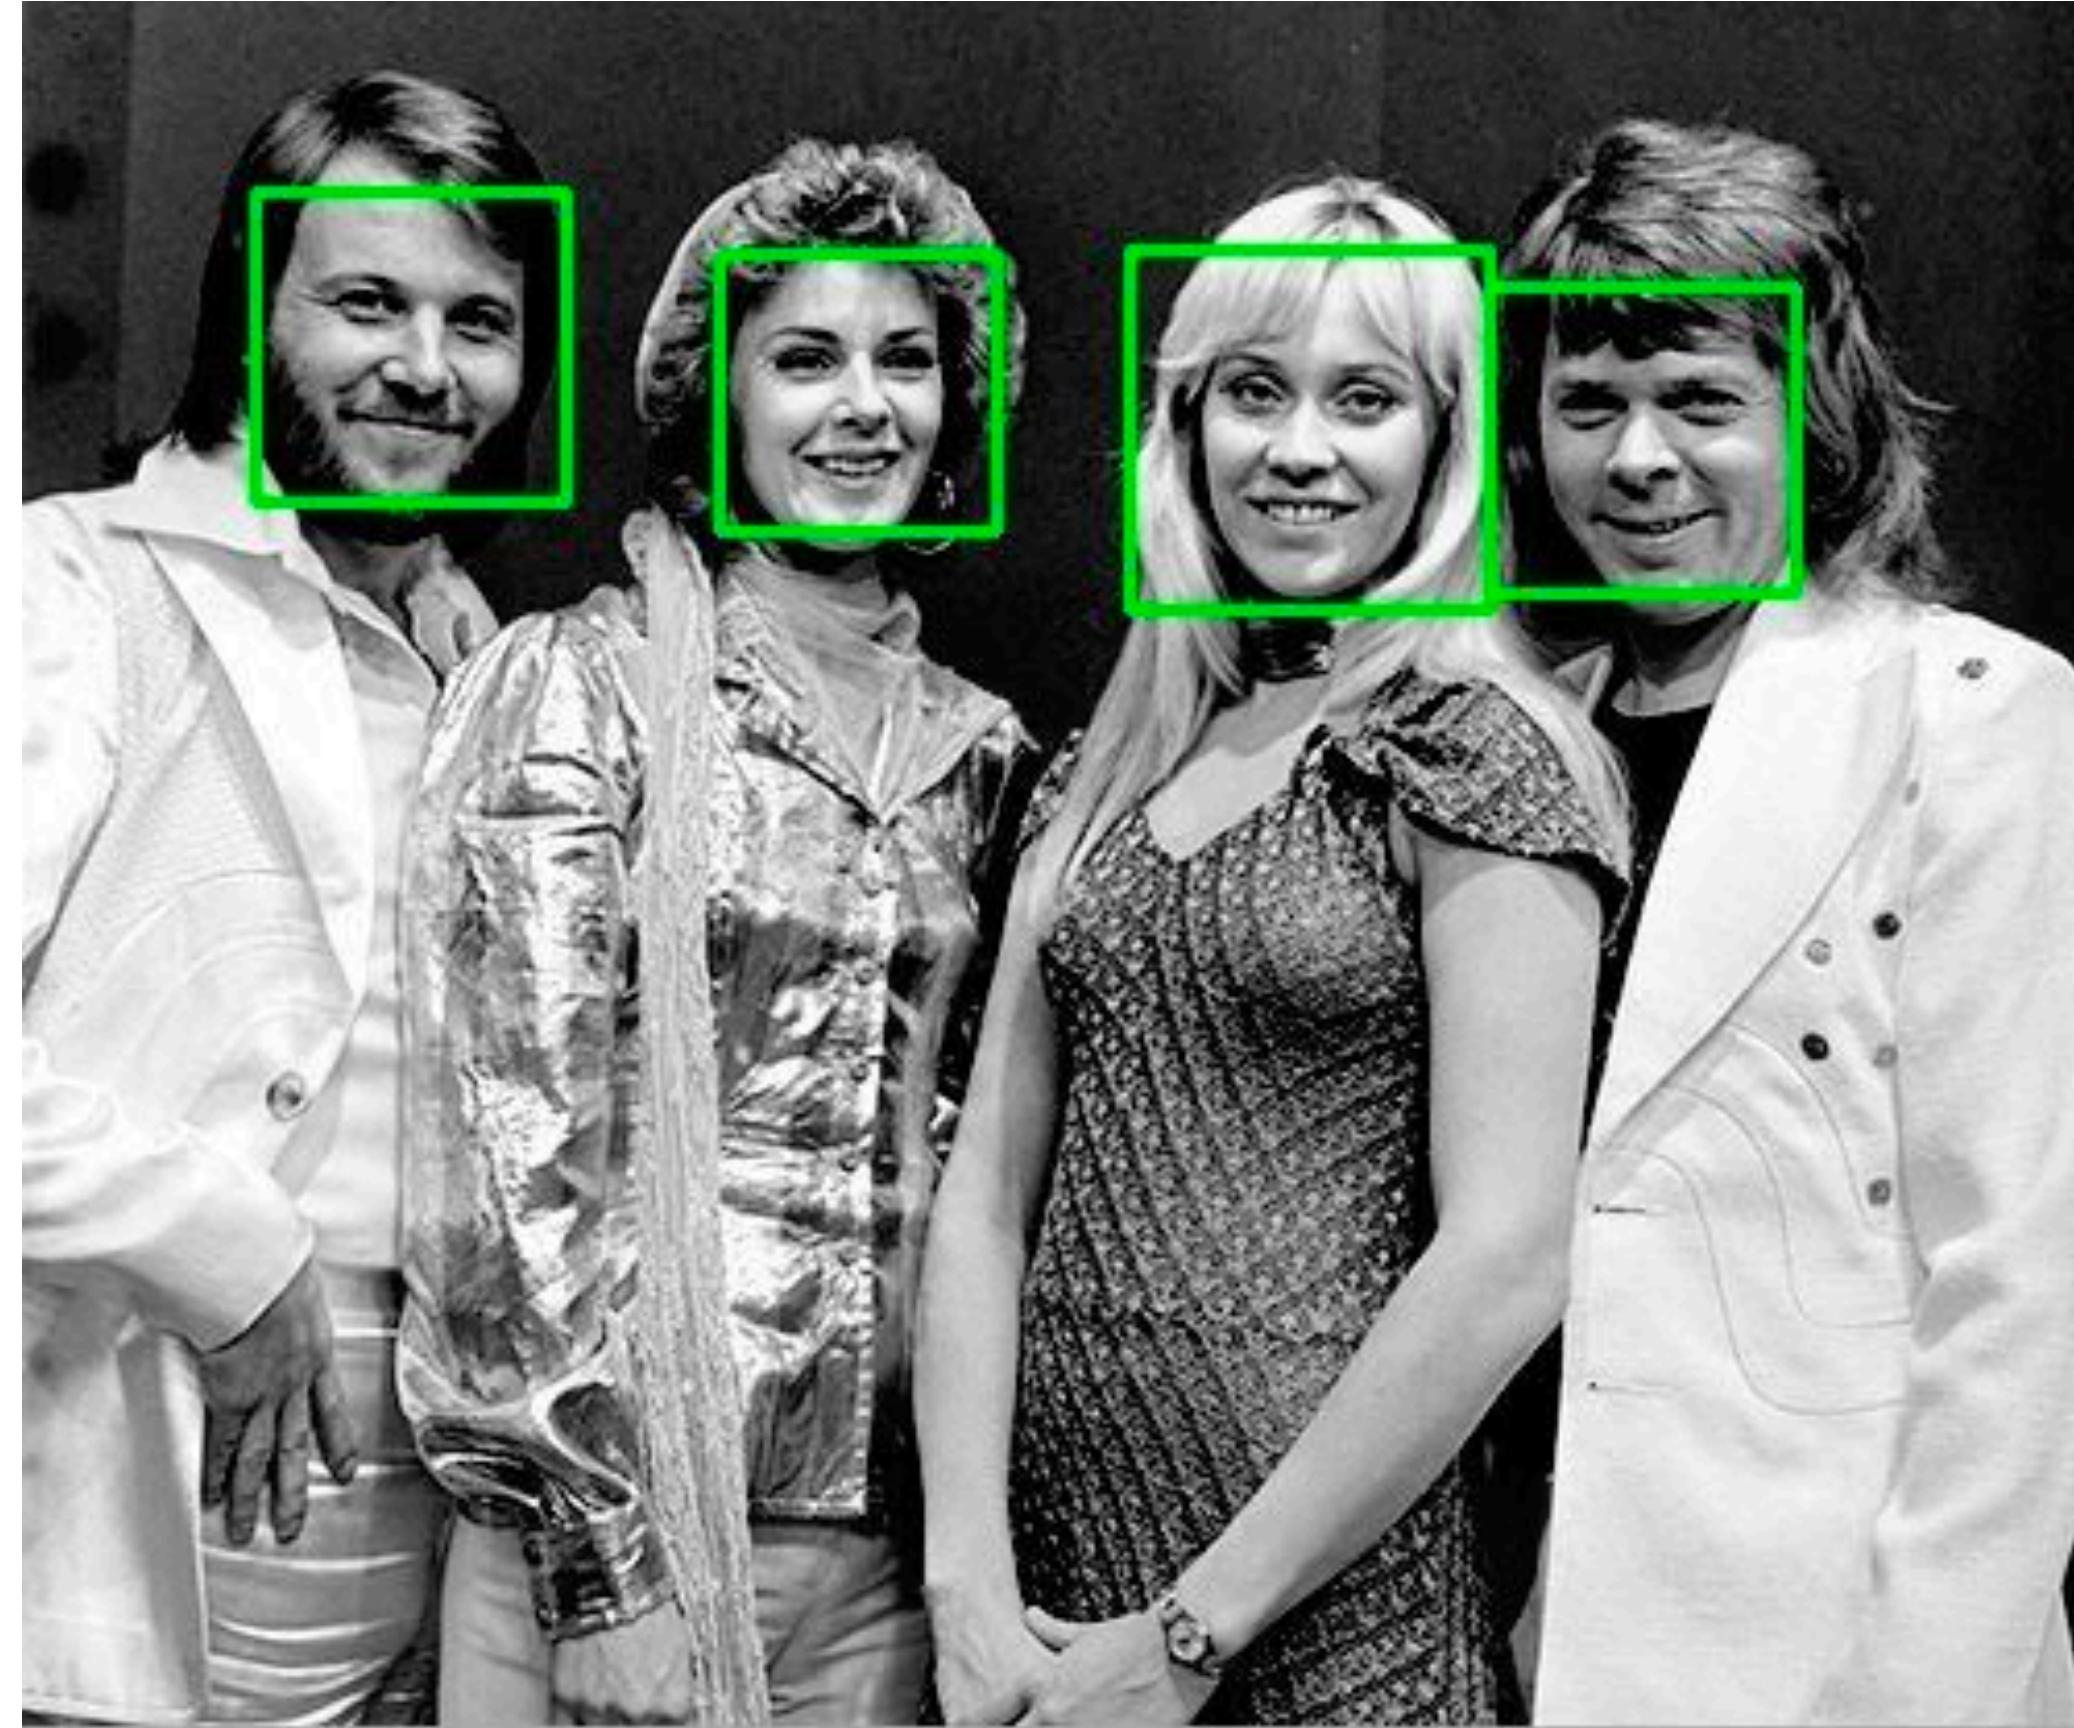
1. Have a **high level view** of how neural networks works
2. Getting **familiar with general concepts** such as overfitting, regularization,....
3. **Code and train** your first neural network

What not to expect?

1. You will not become an expert in deep learning
2. You (most likely) won't be able to develop your own deep learning model
3. You won't know the state-of-the-art technics to deploy neural networks

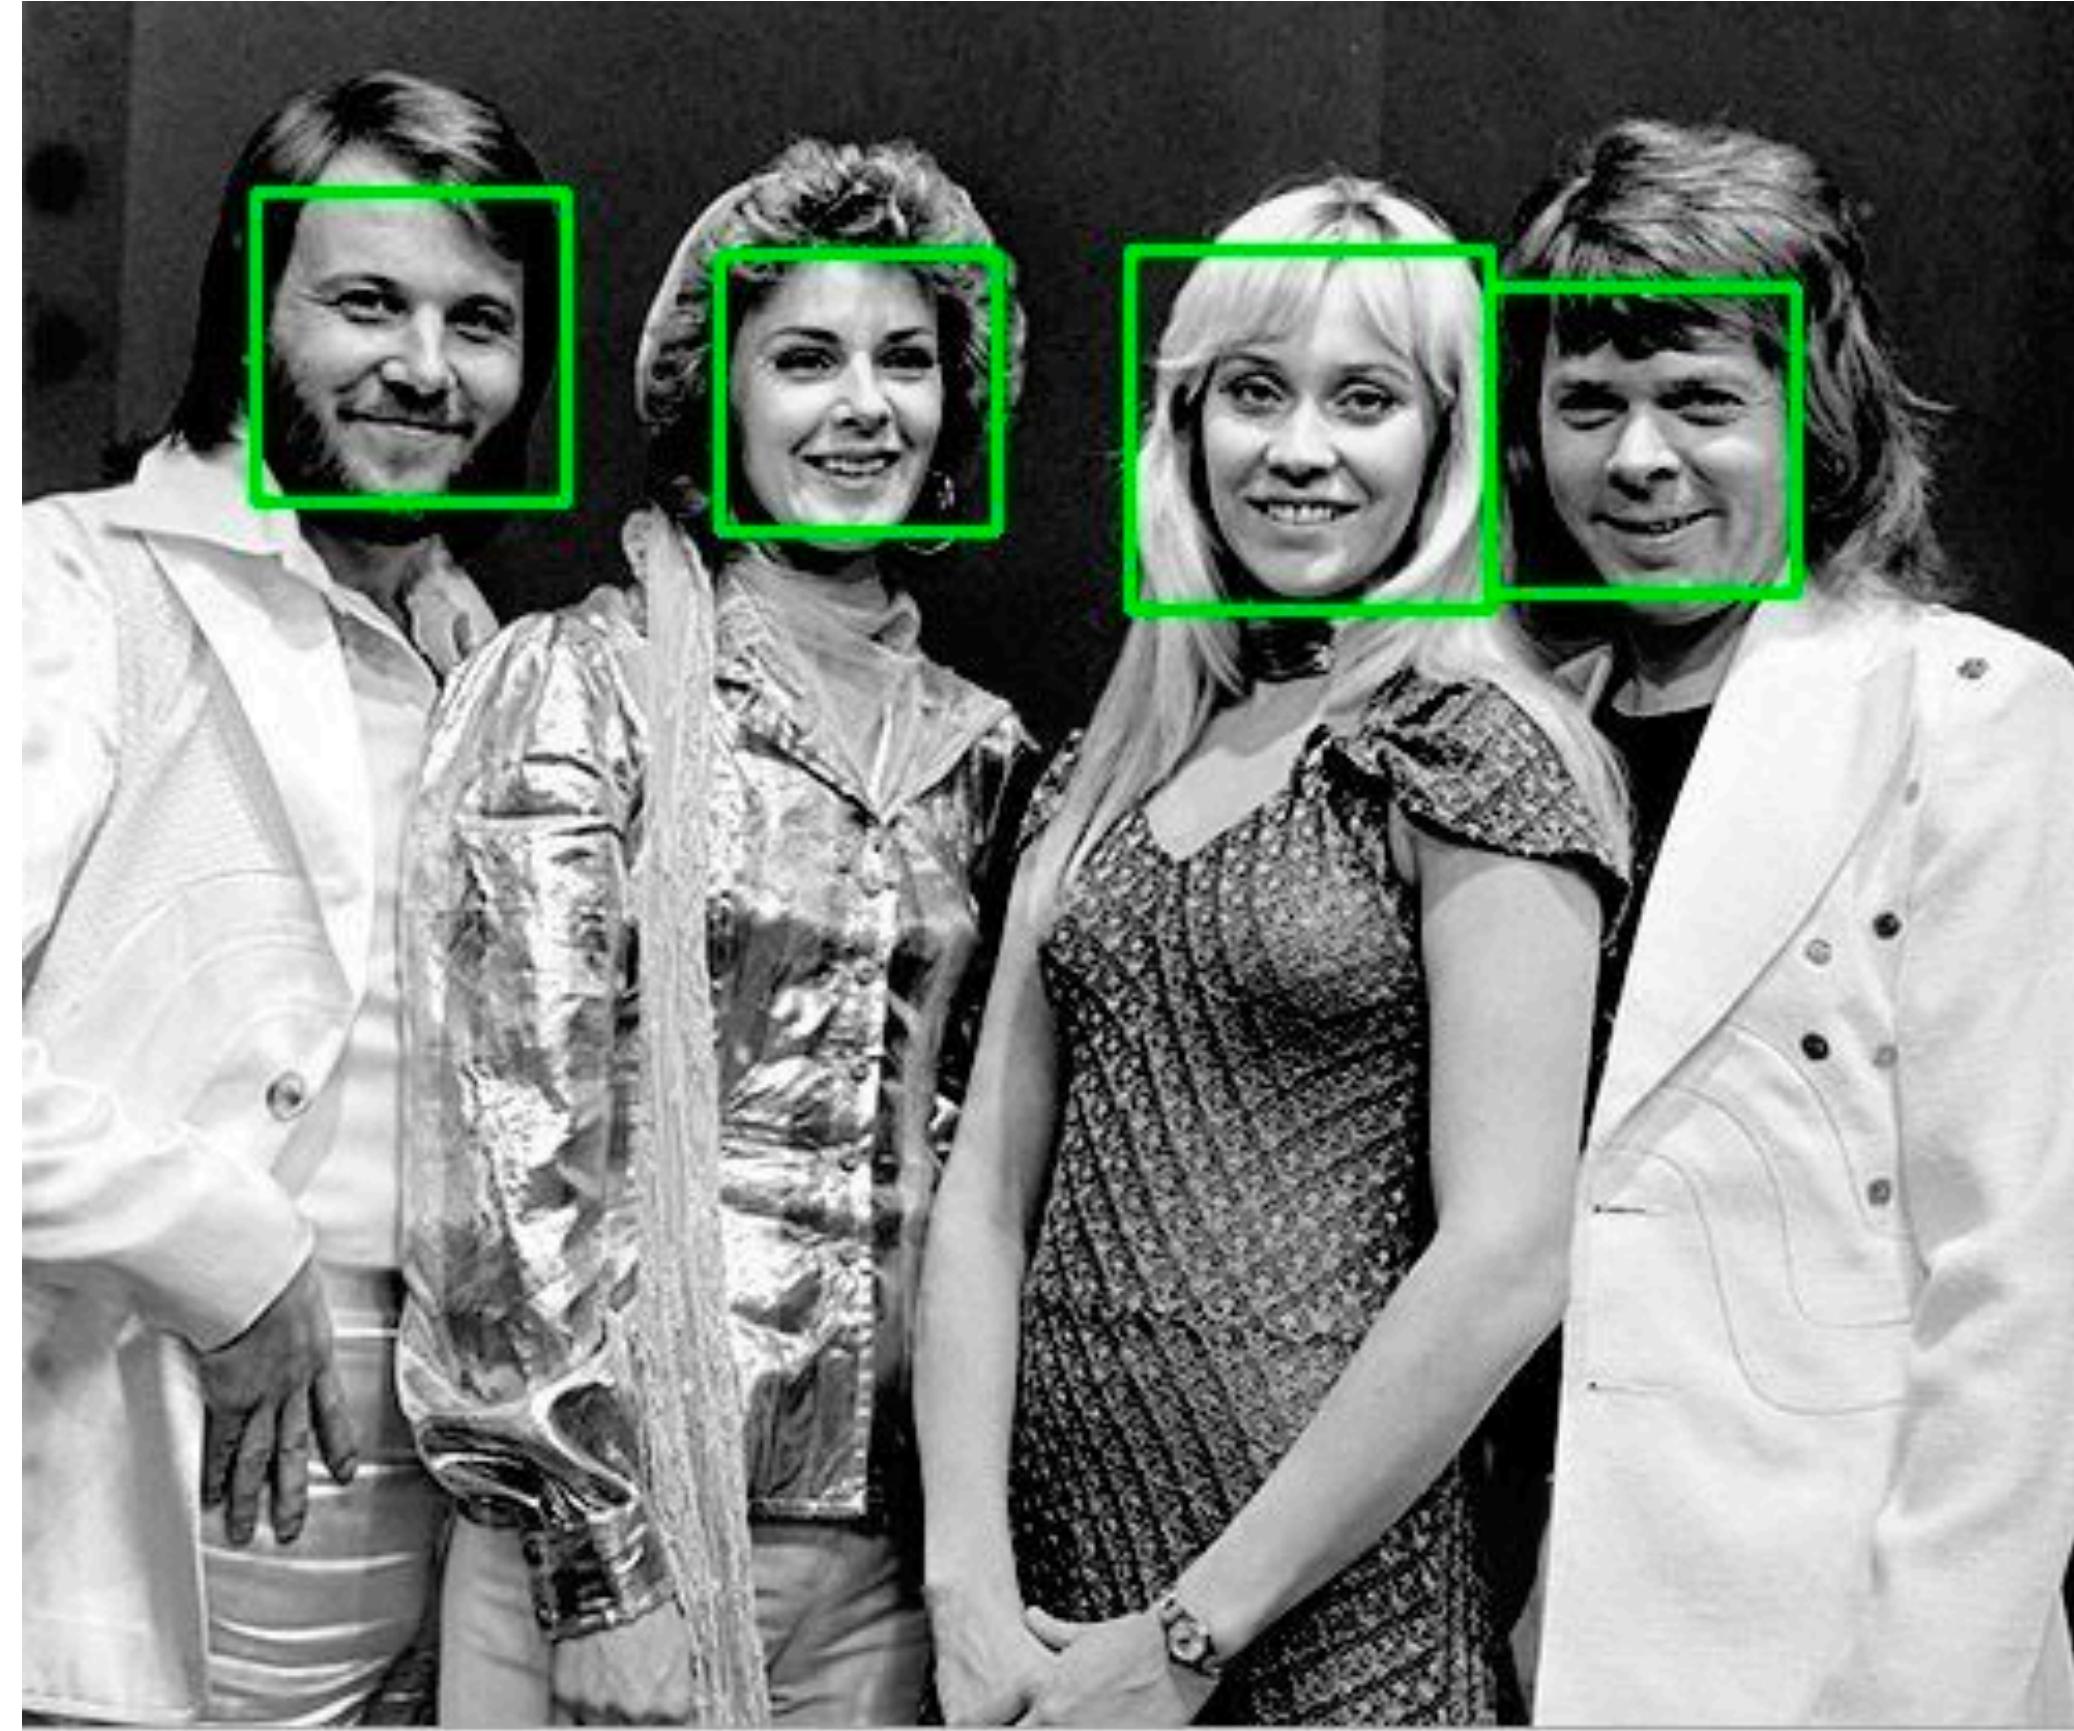
Facial Detection / Recognition

Facial Detection / Recognition



Face detection

Facial Detection / Recognition



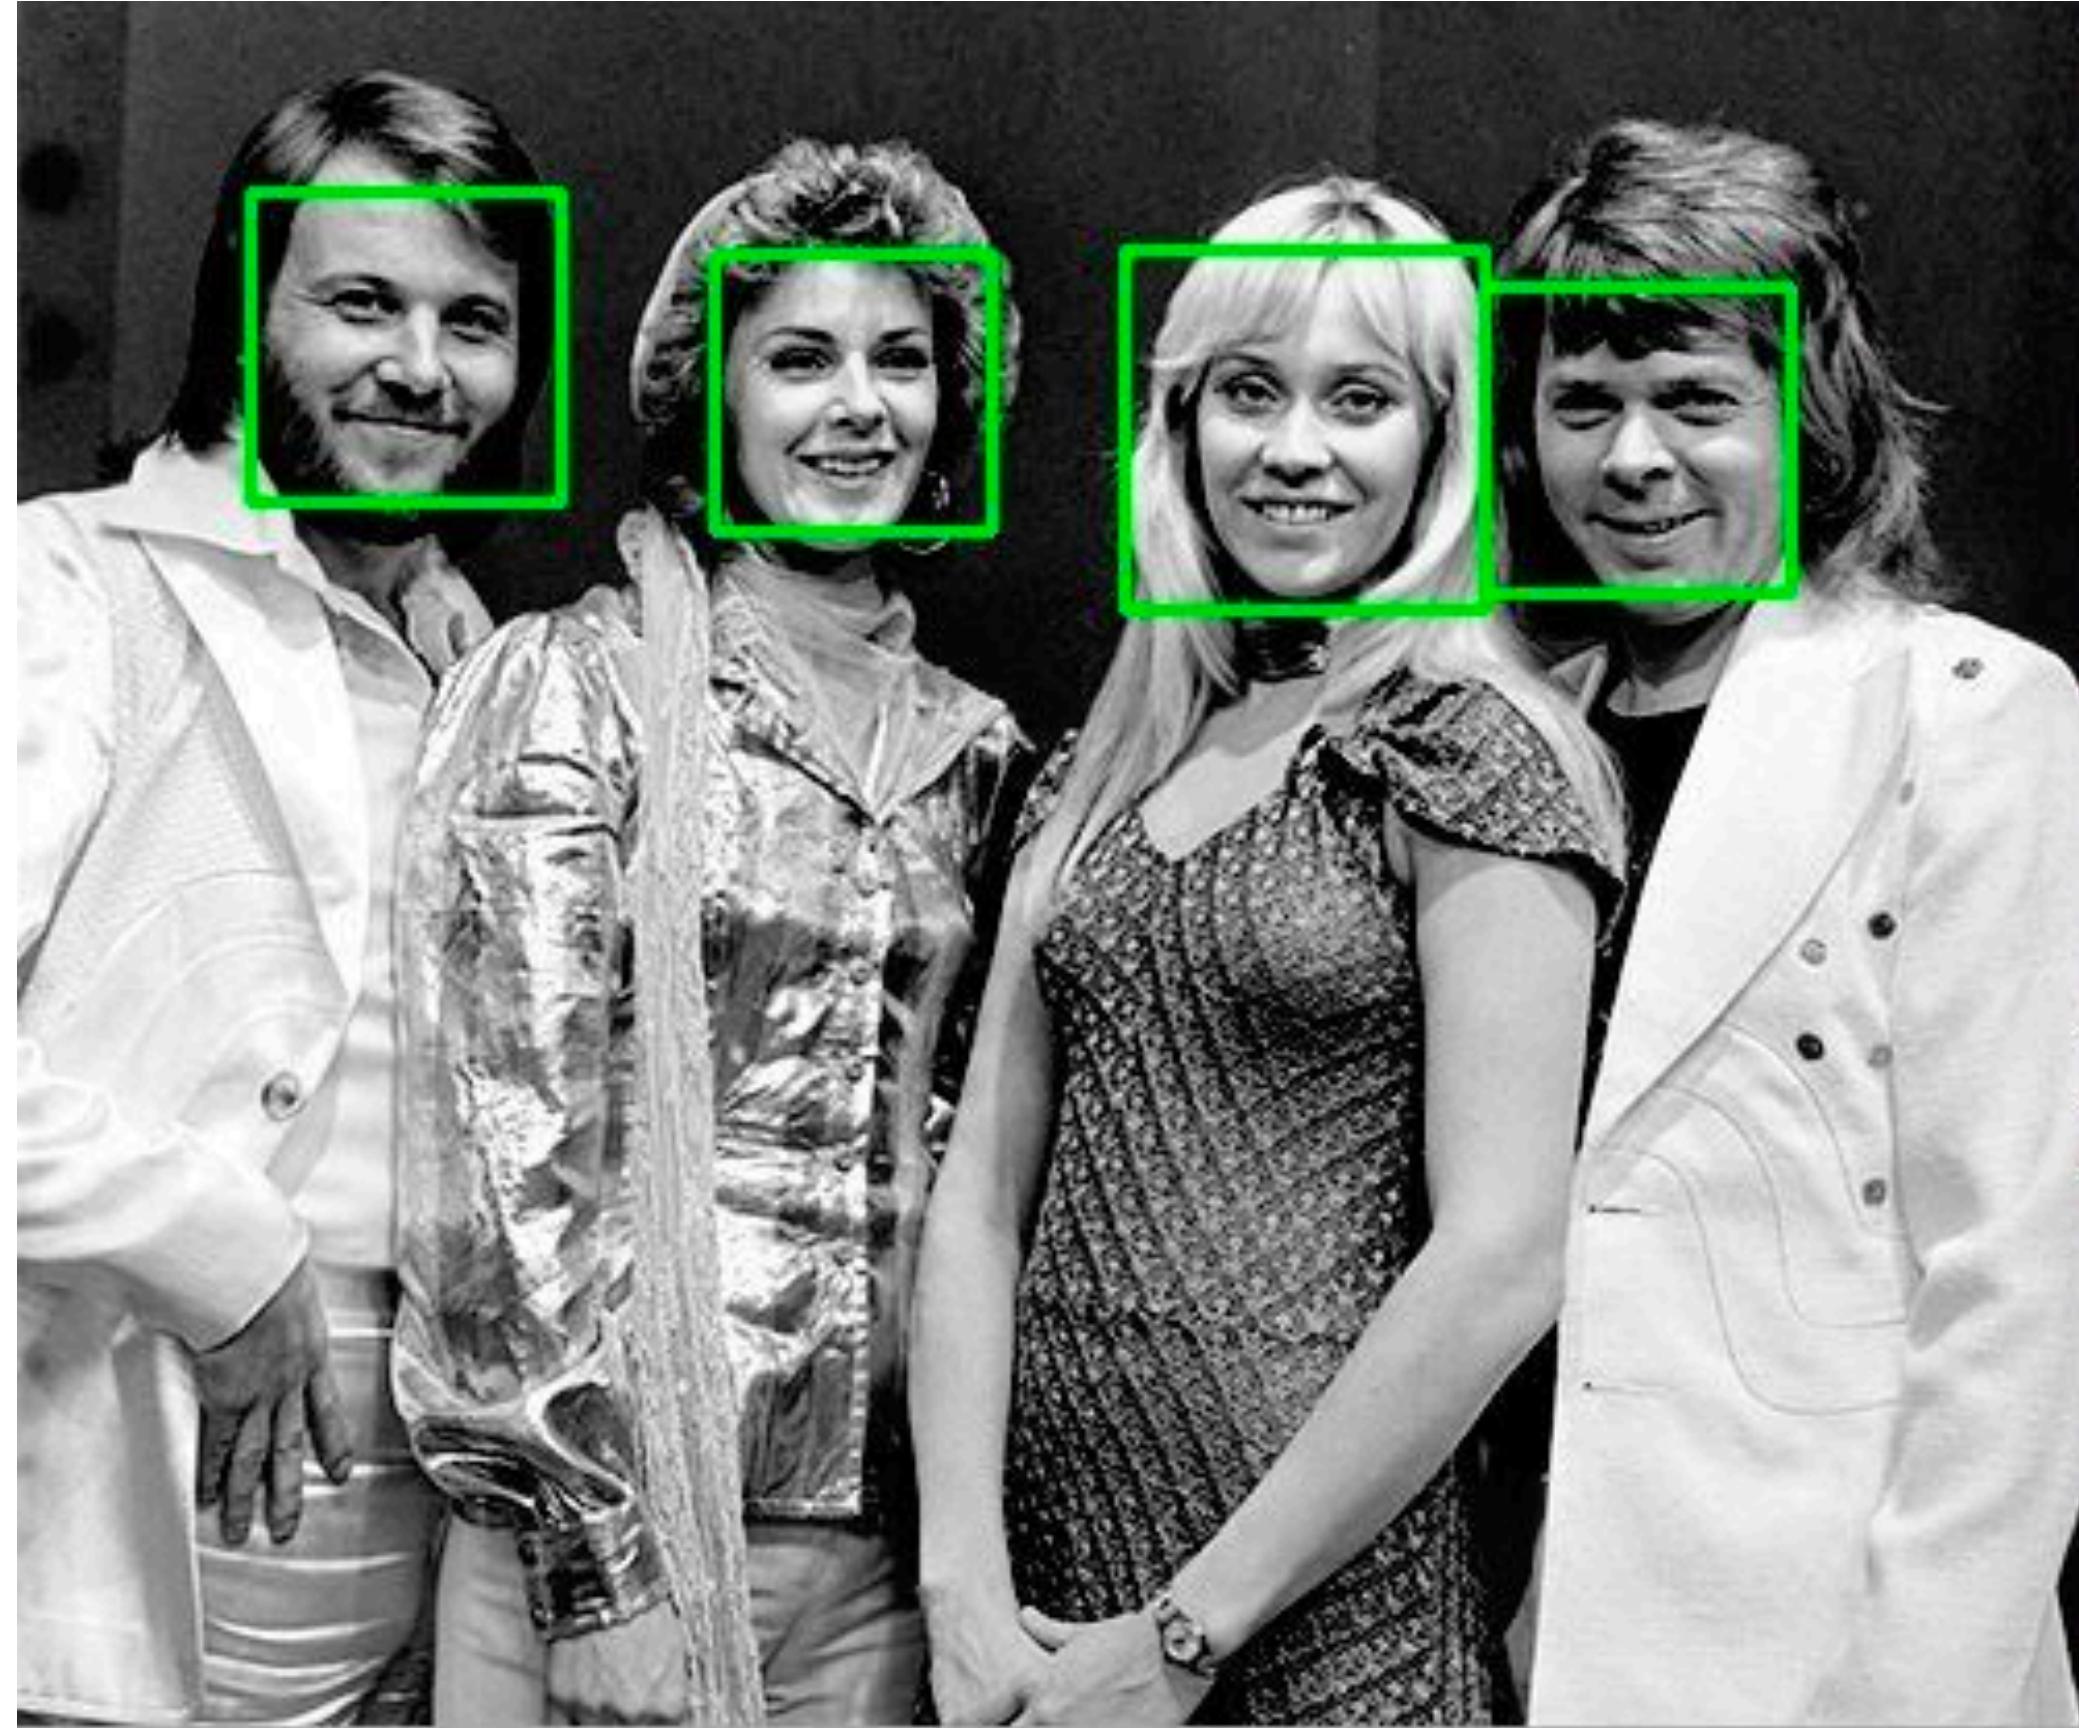
Face detection



Unlock your smarphone

Take pictures

Facial Detection / Recognition



Face detection

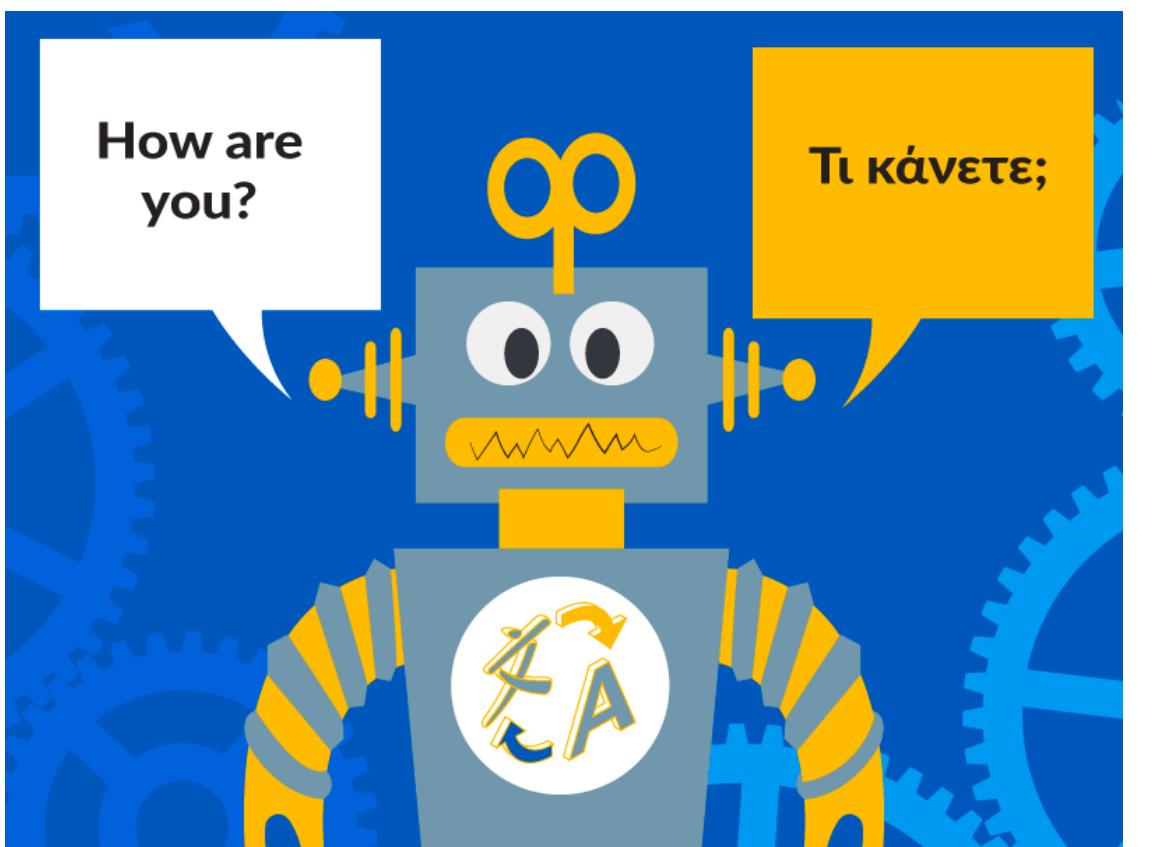


Unlock your smartphone

Take pictures

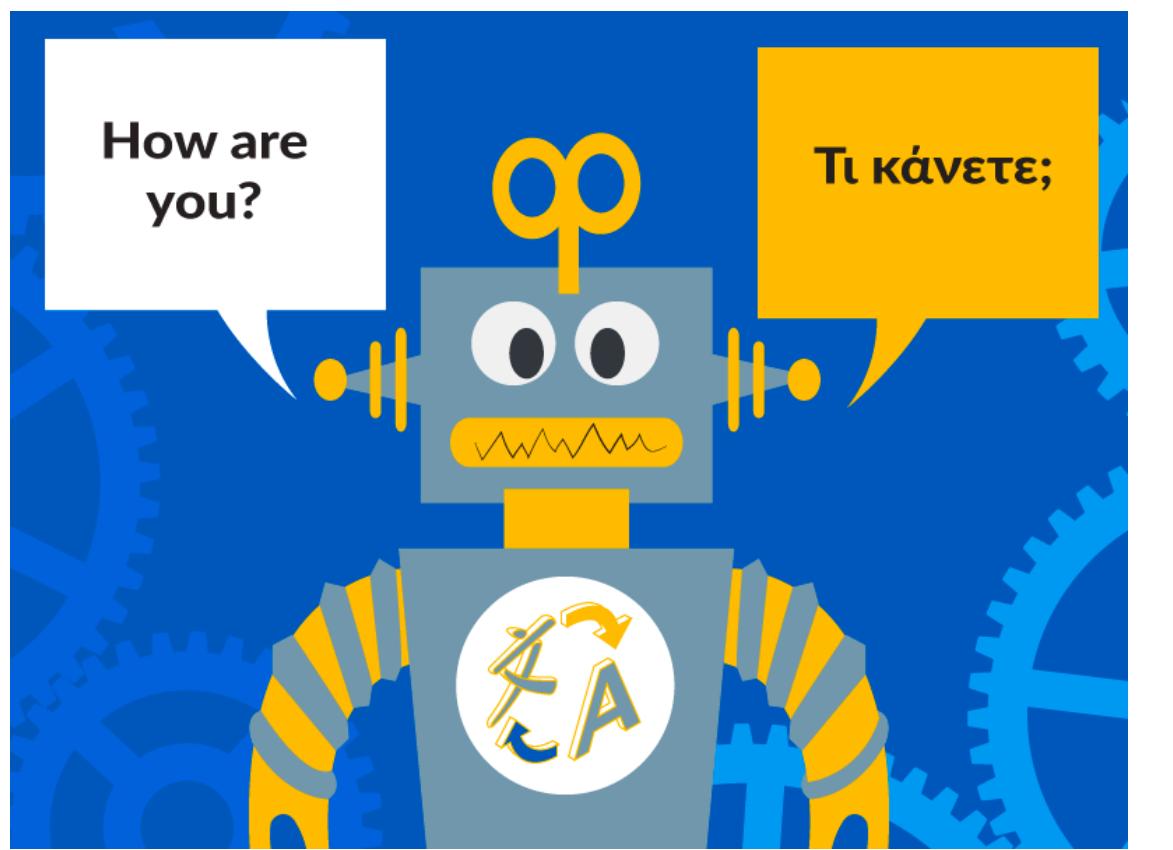


Airport



Translation

Translation / Information retrieval

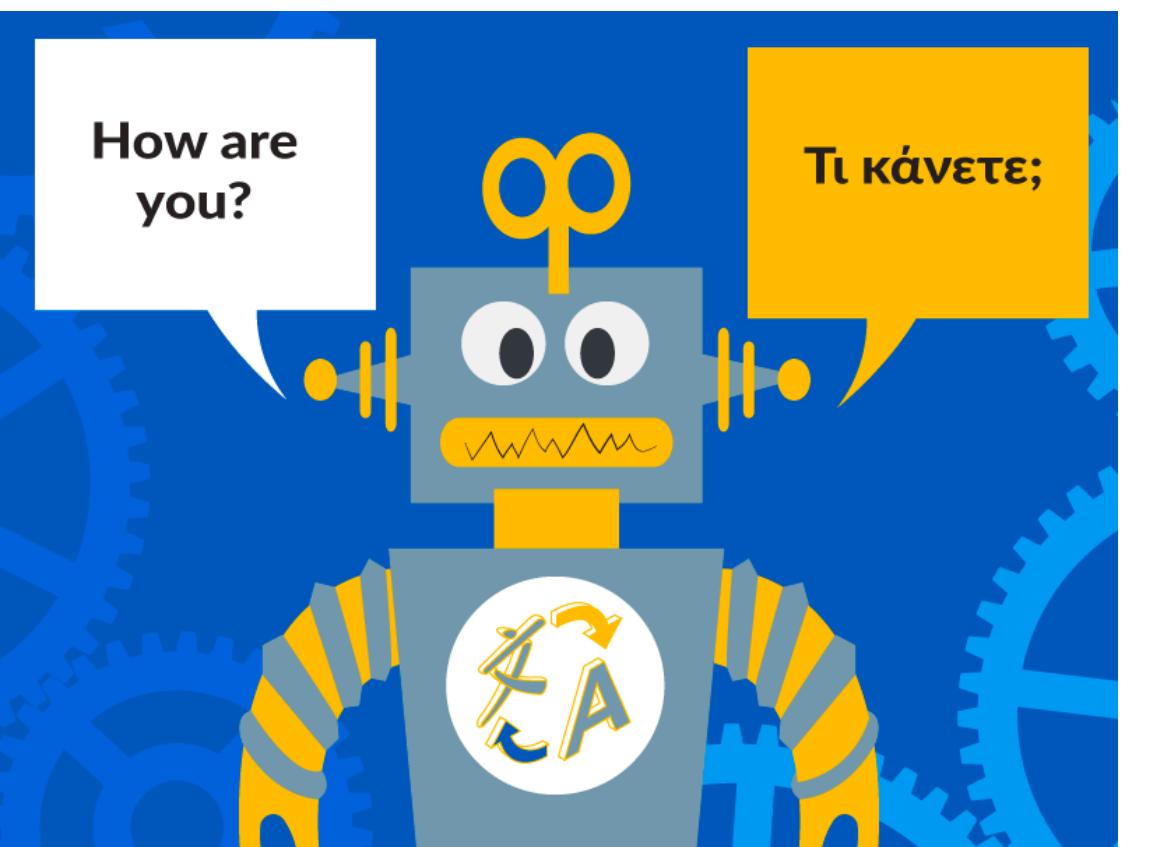


Translation



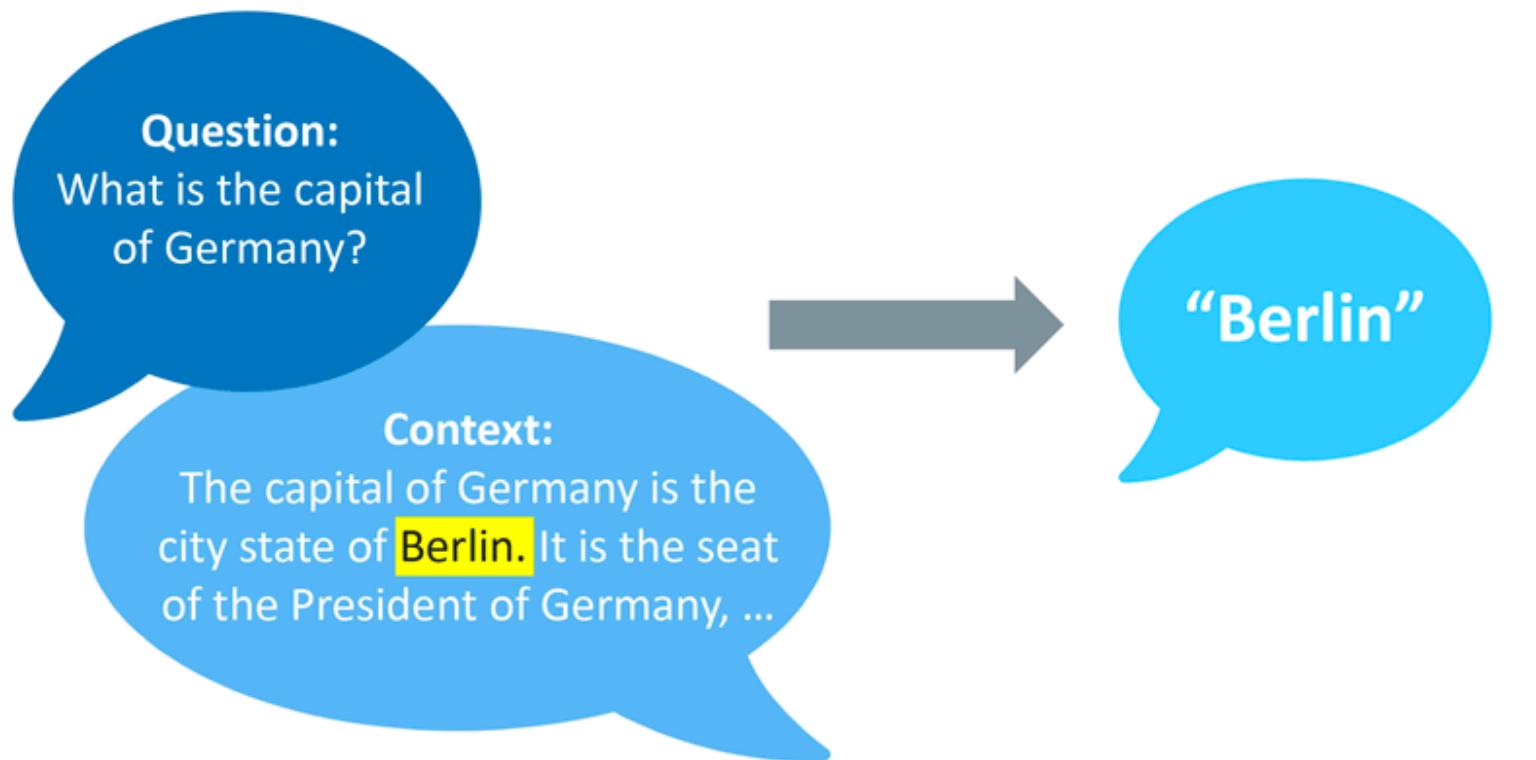
Information Retrieval

Translation / Information retrieval



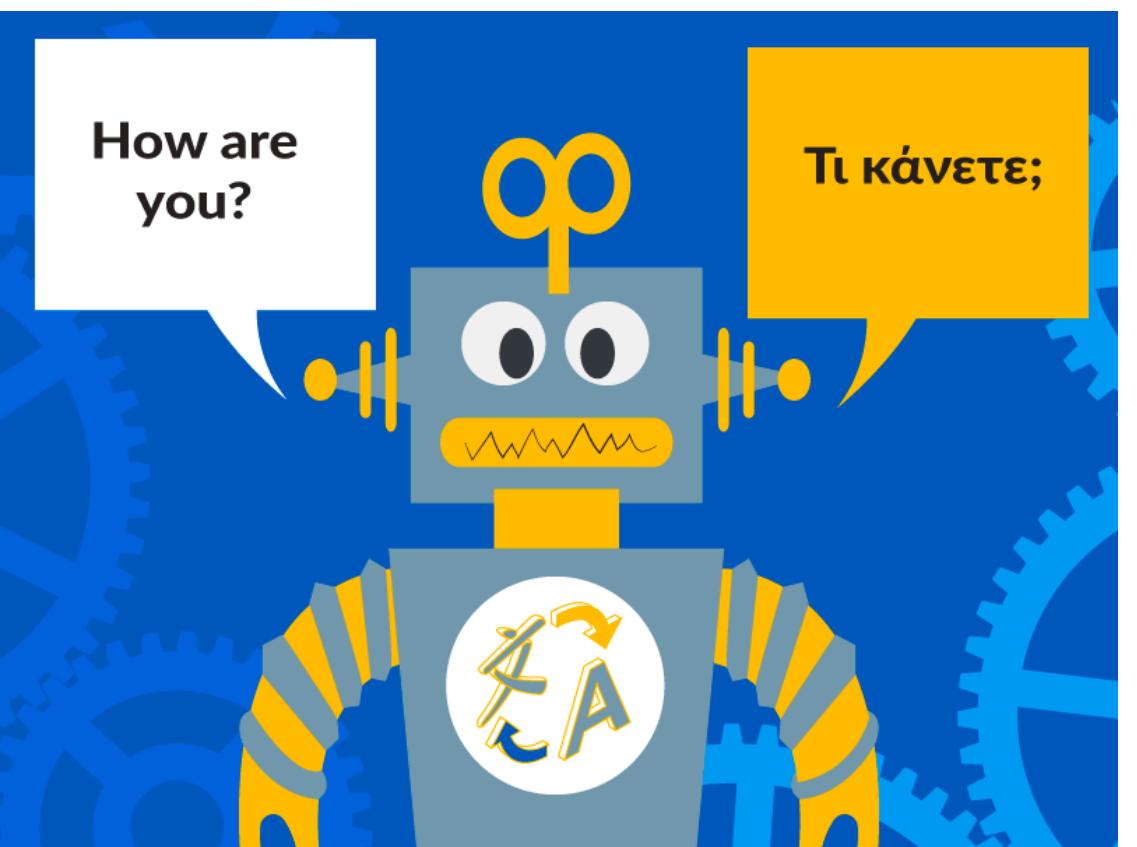
Information Retrieval

Translation

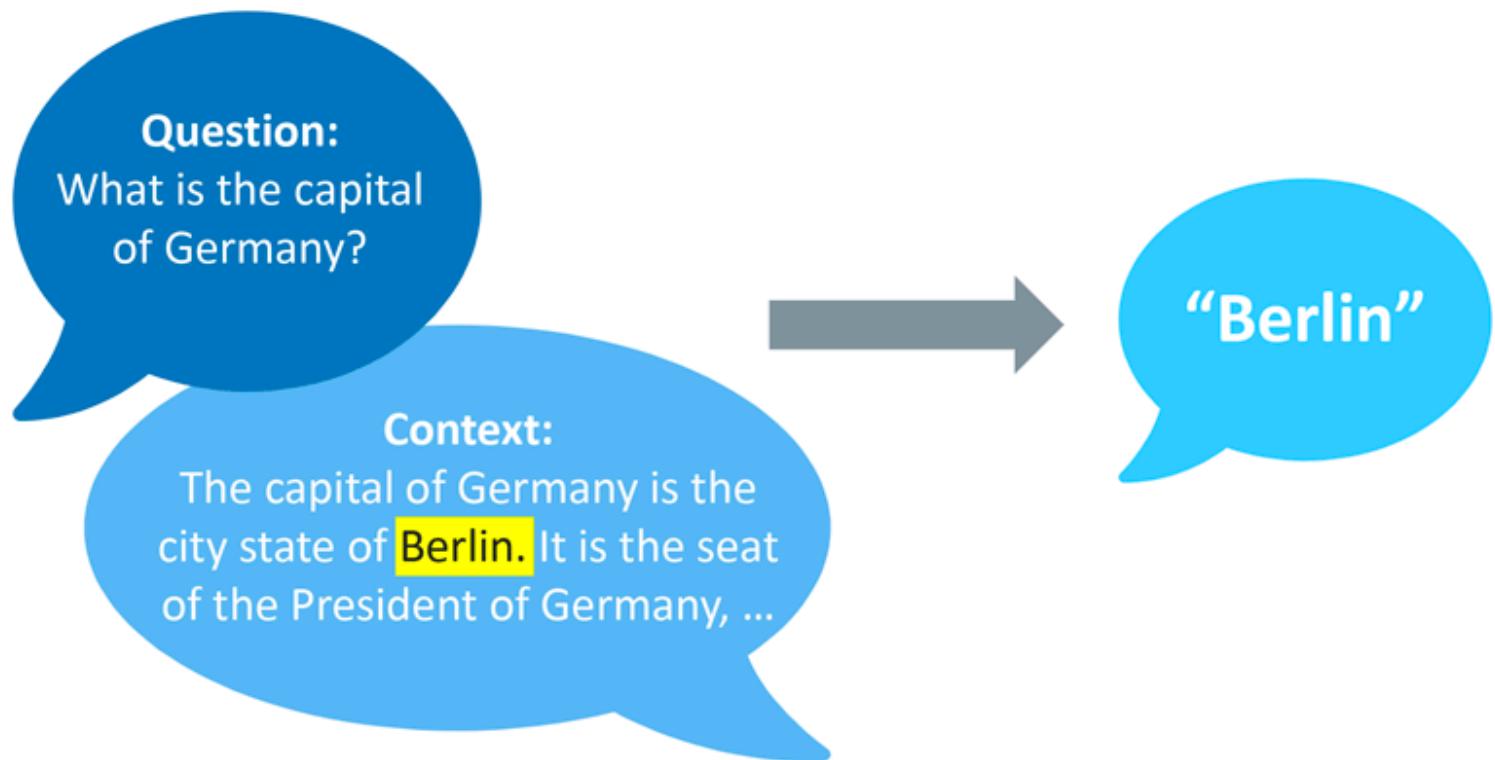


Question Answering

Translation / Information retrieval



Translation



Question Answering



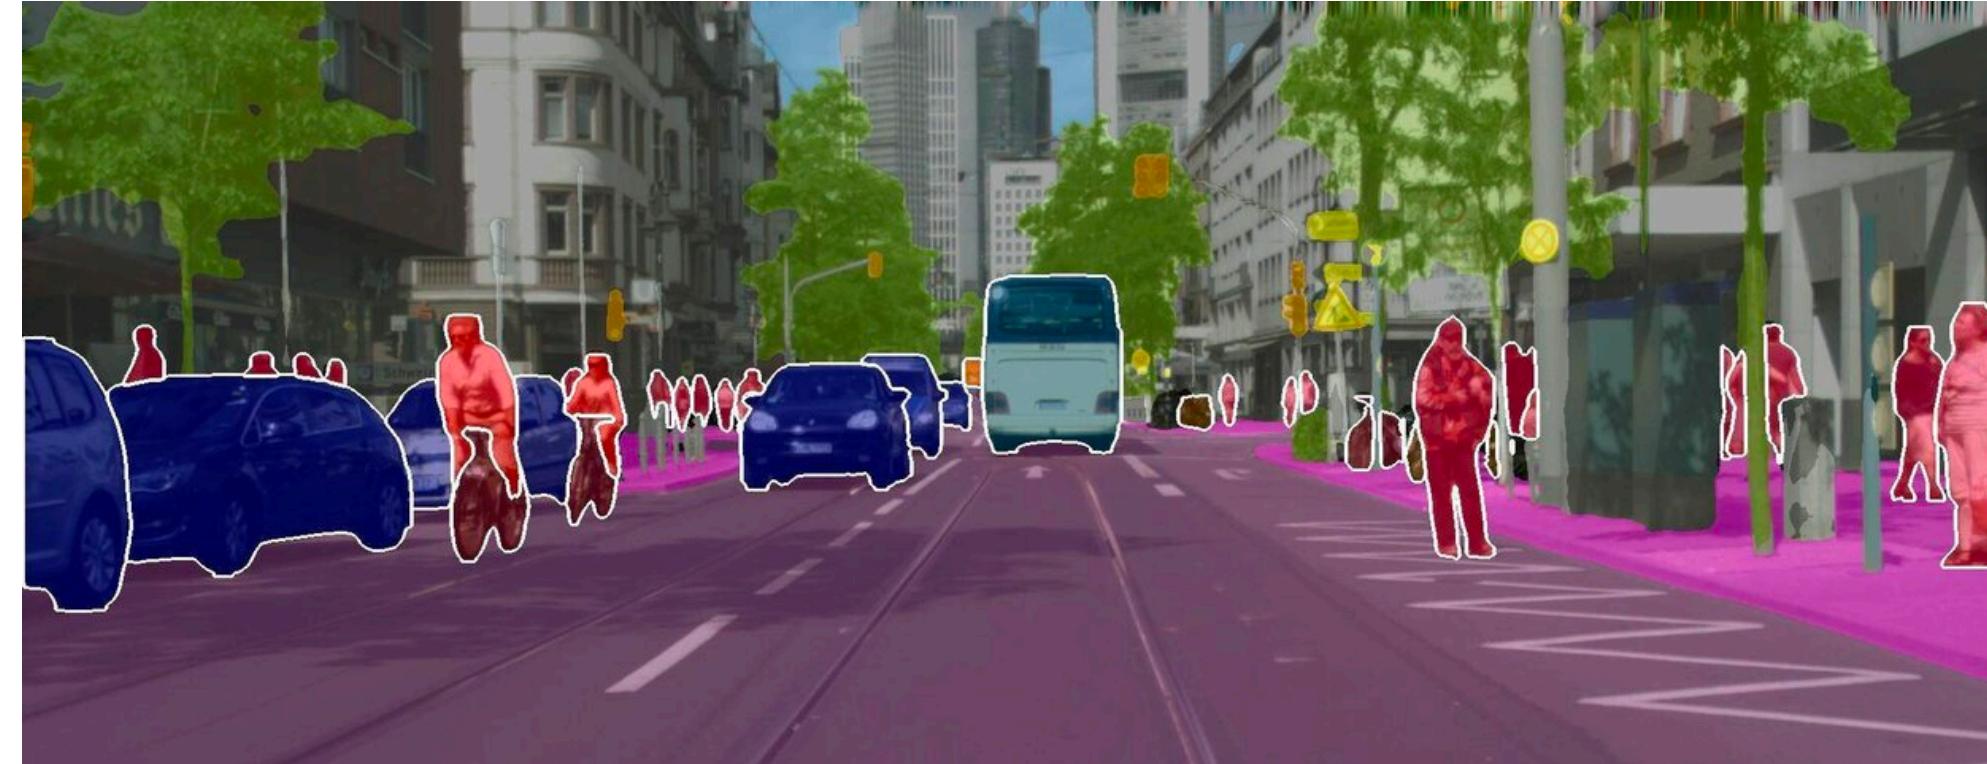
Information Retrieval



Social Media Analysis

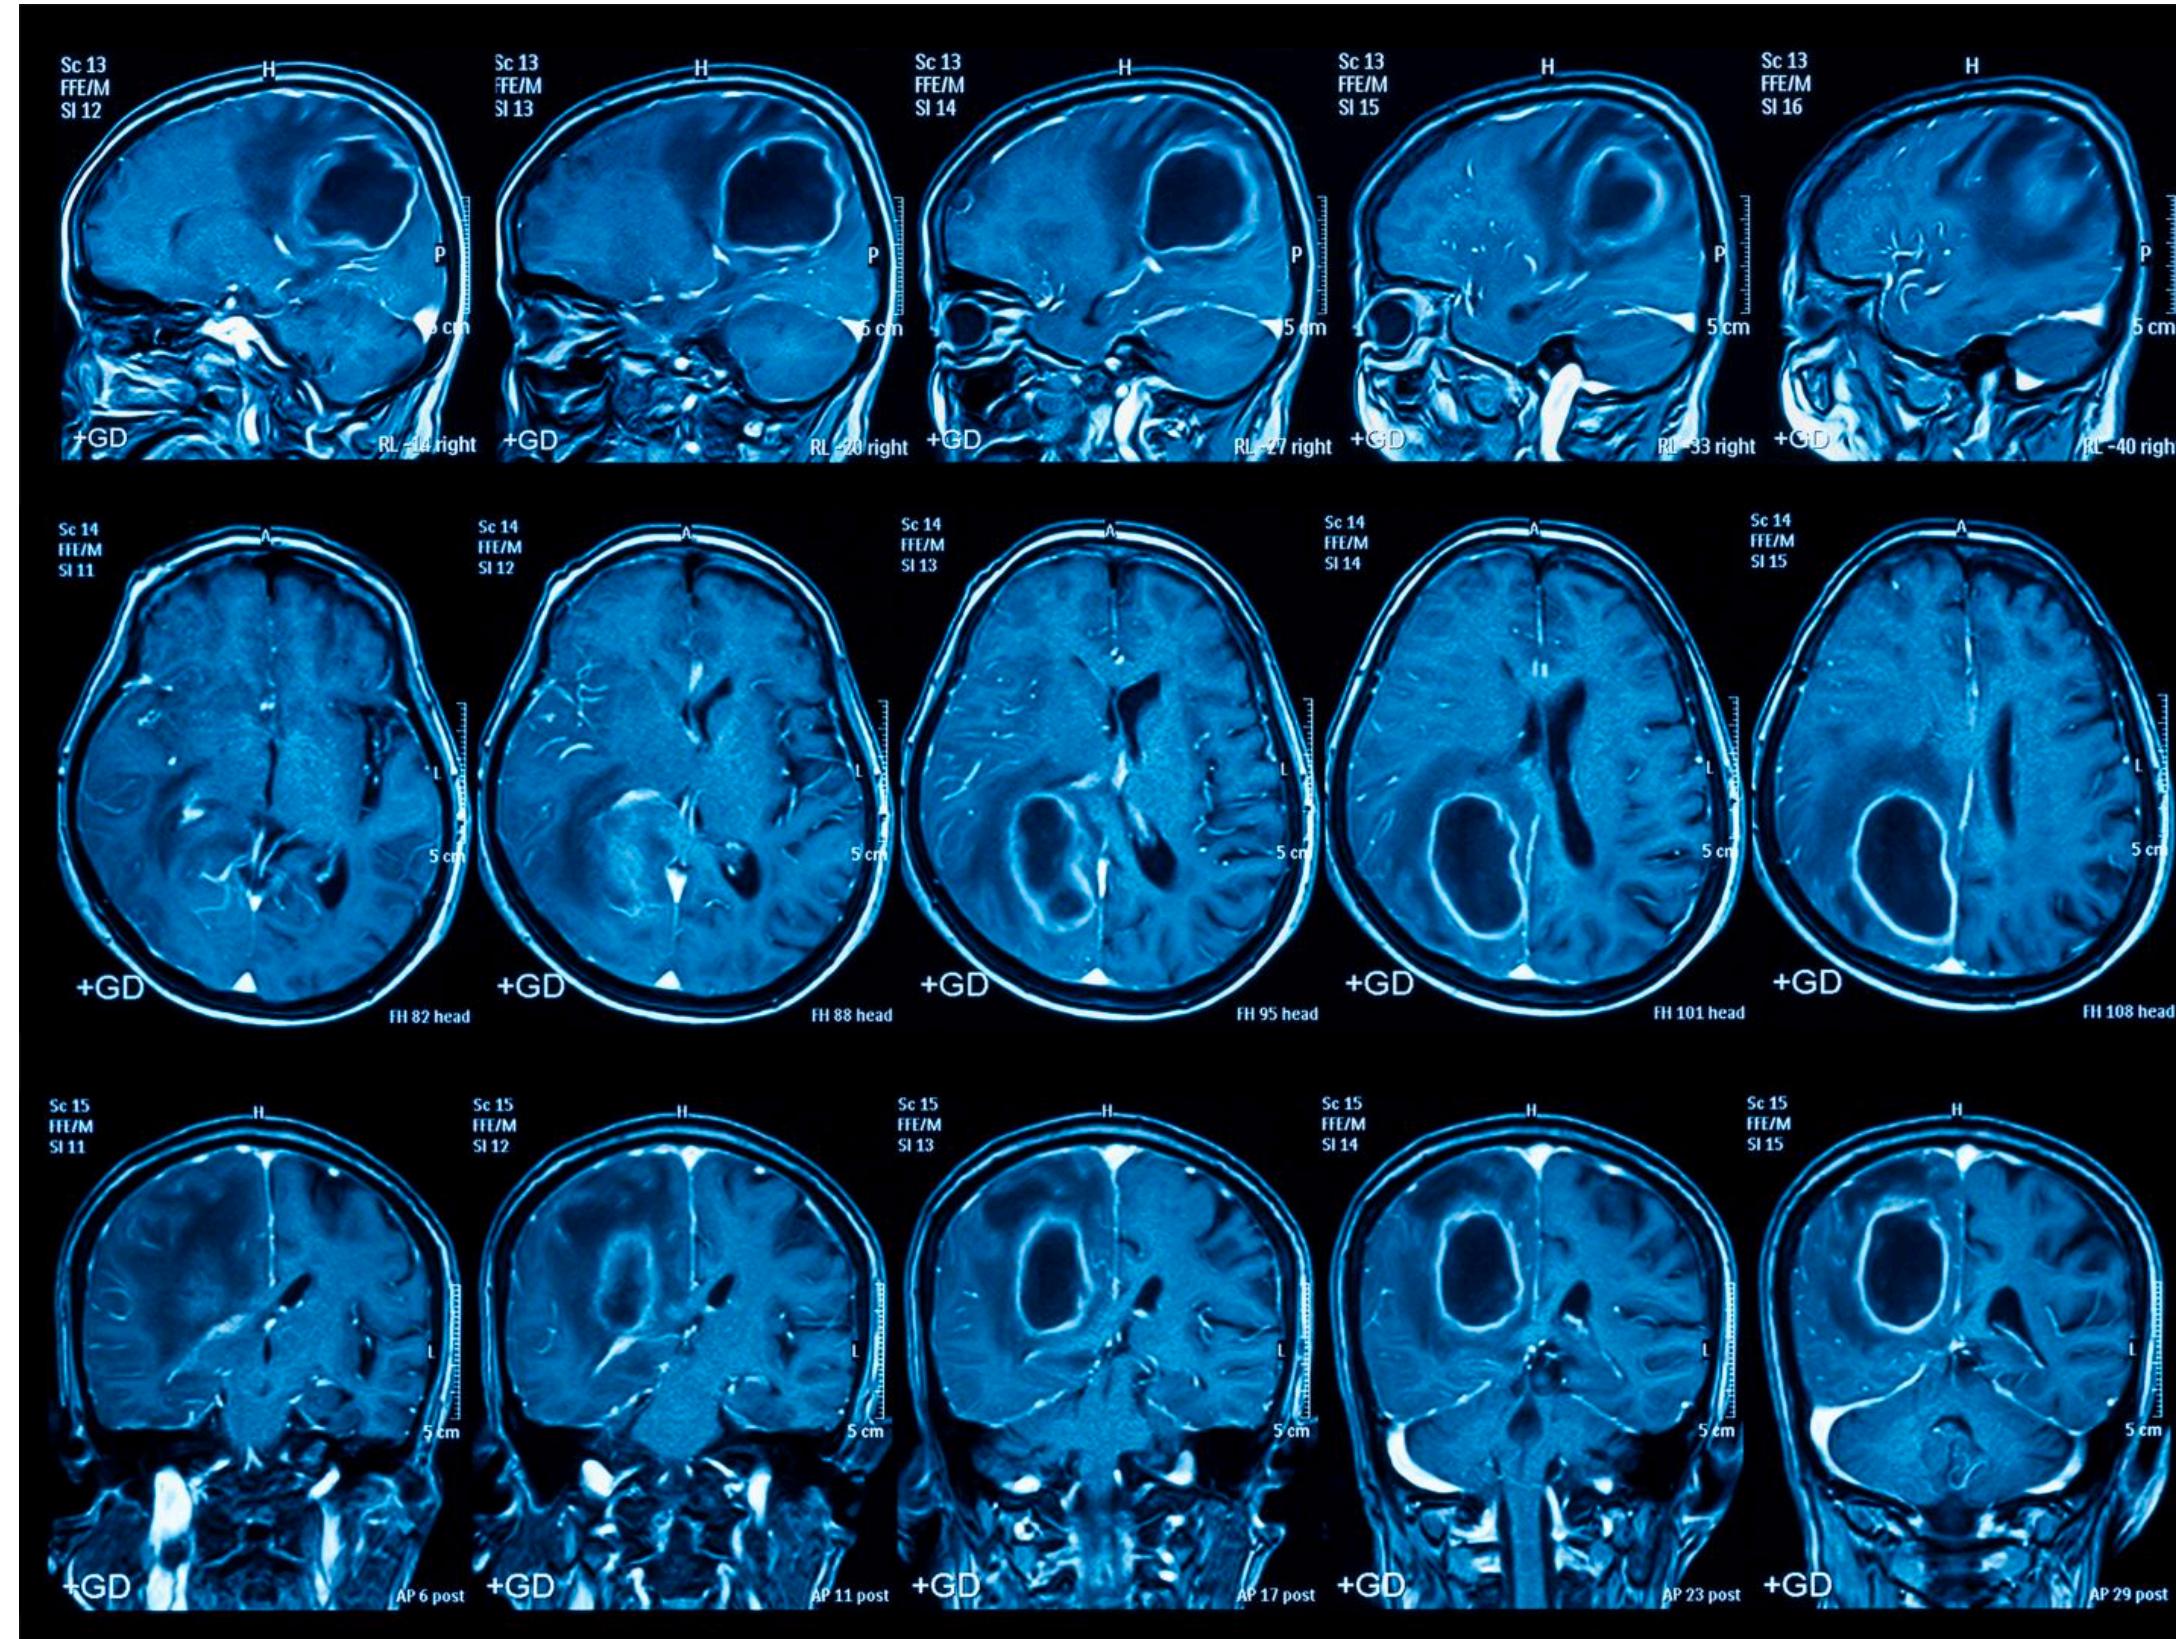
Autonomous Driving

Autonomous Driving



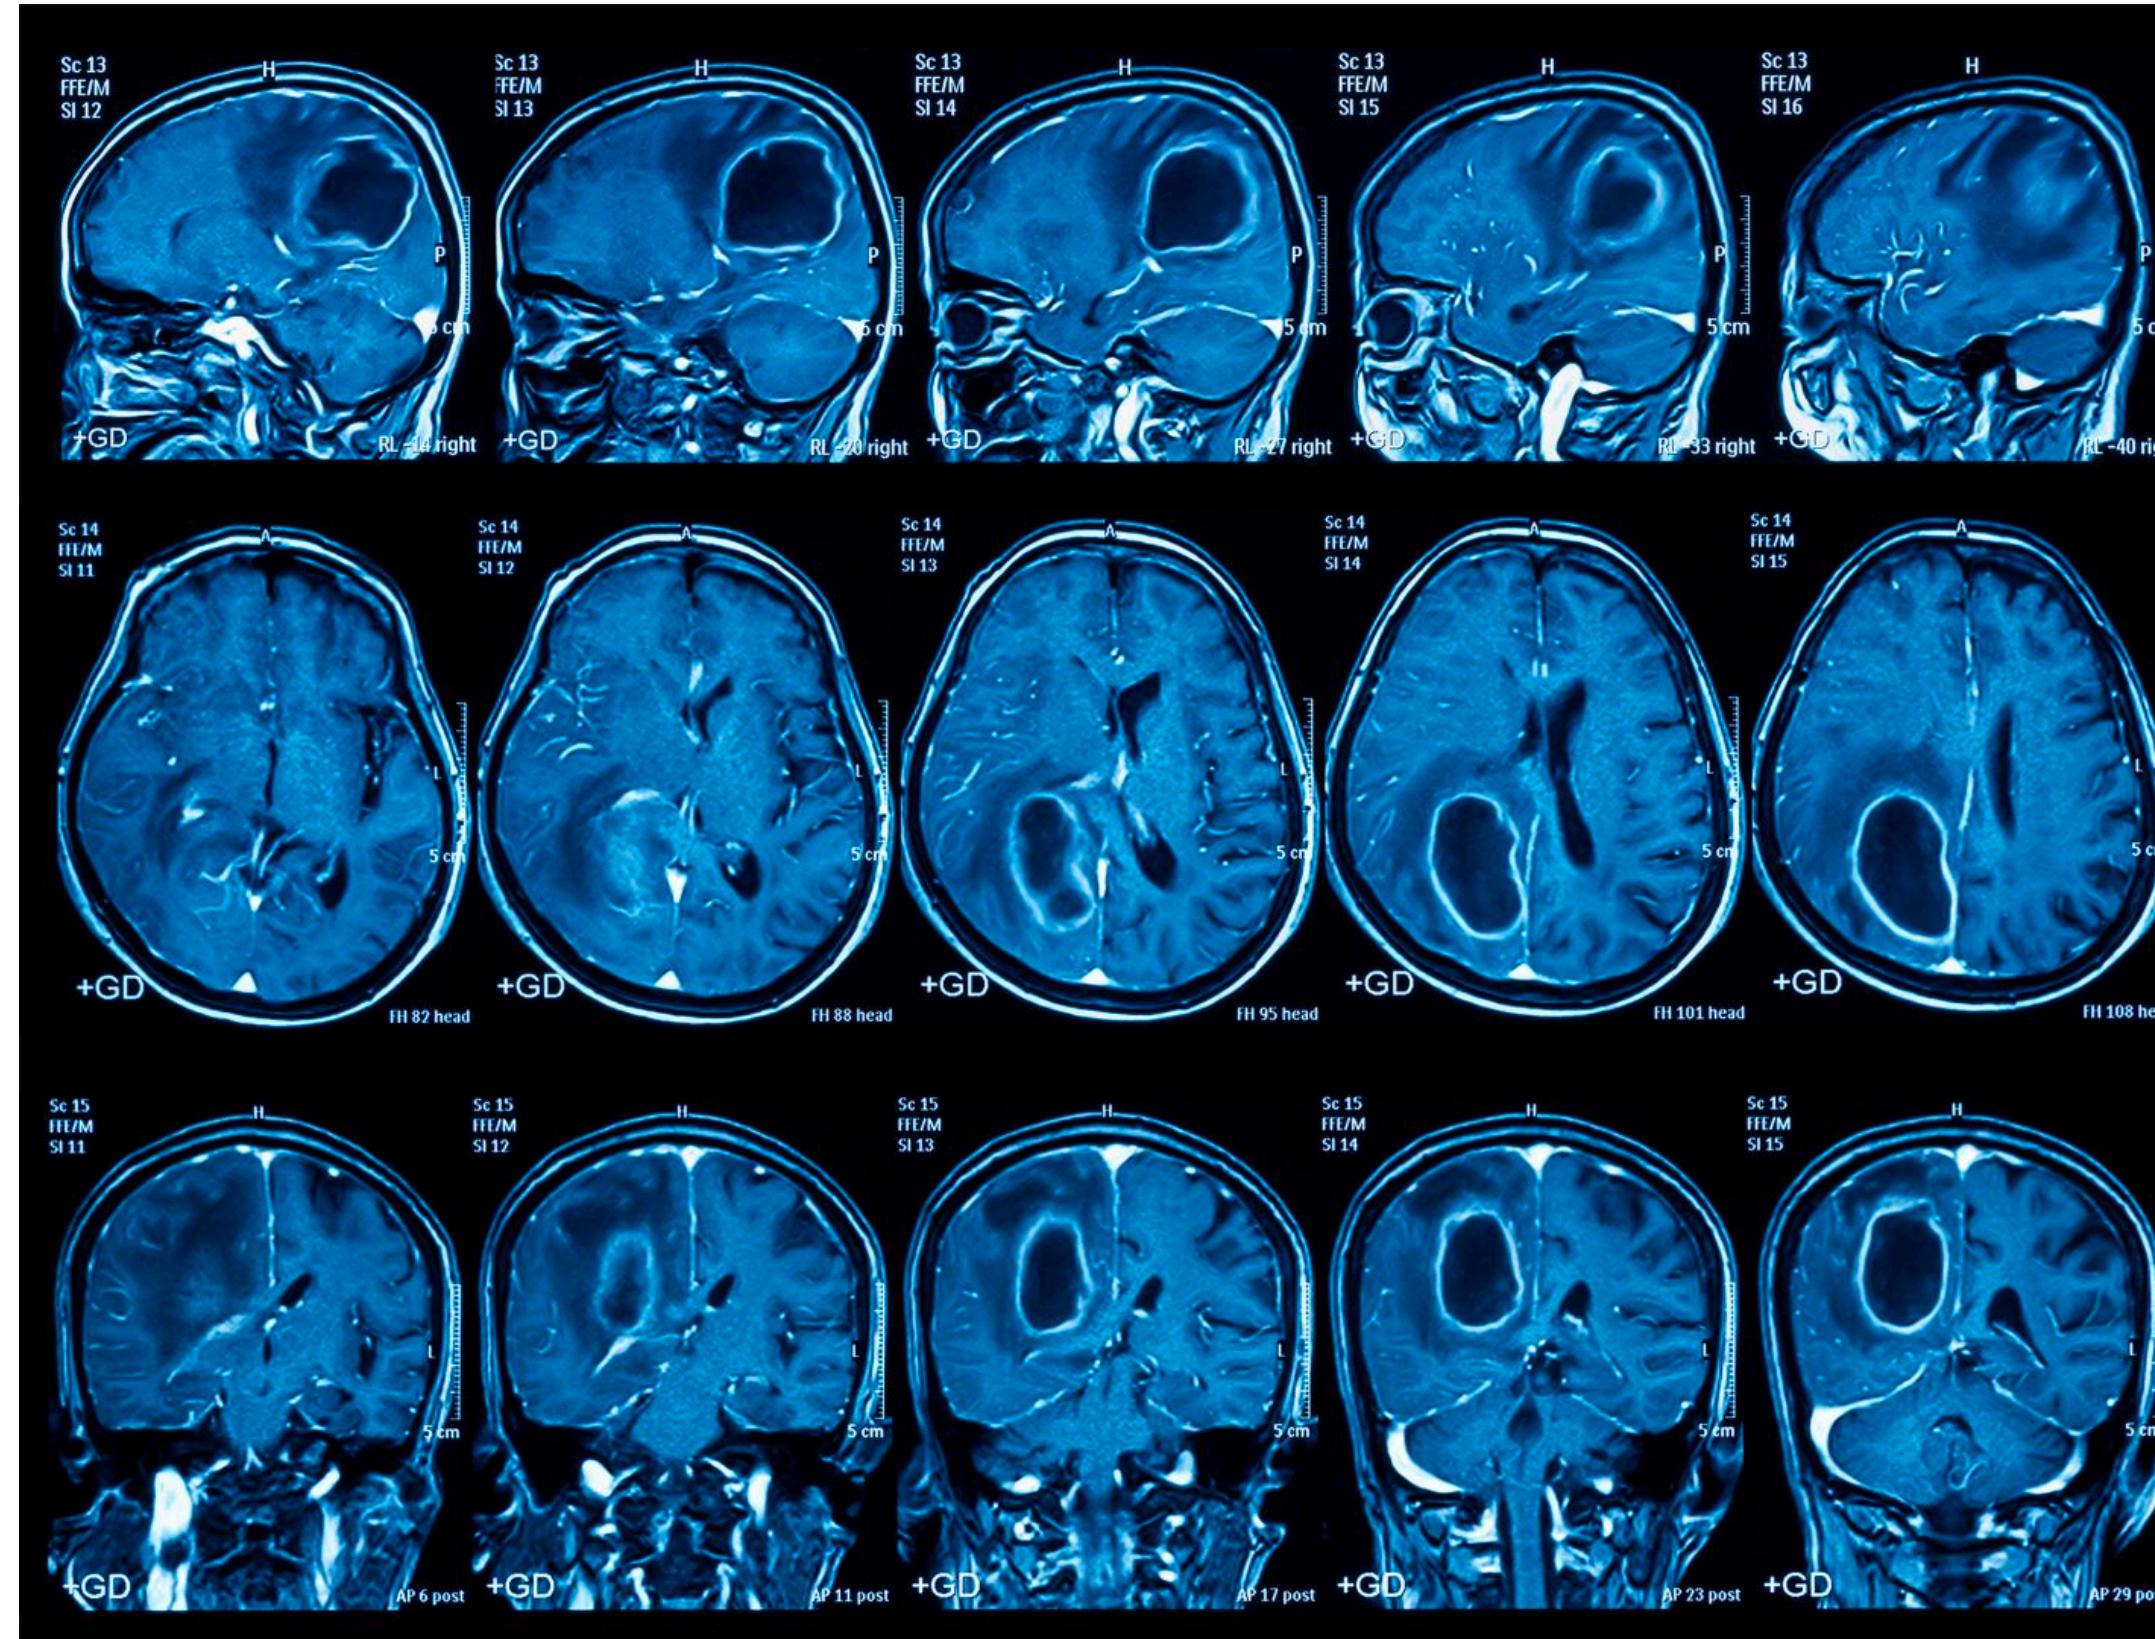
Medical Application

Medical Application



Medical Imaging

Medical Application



Medical Imaging



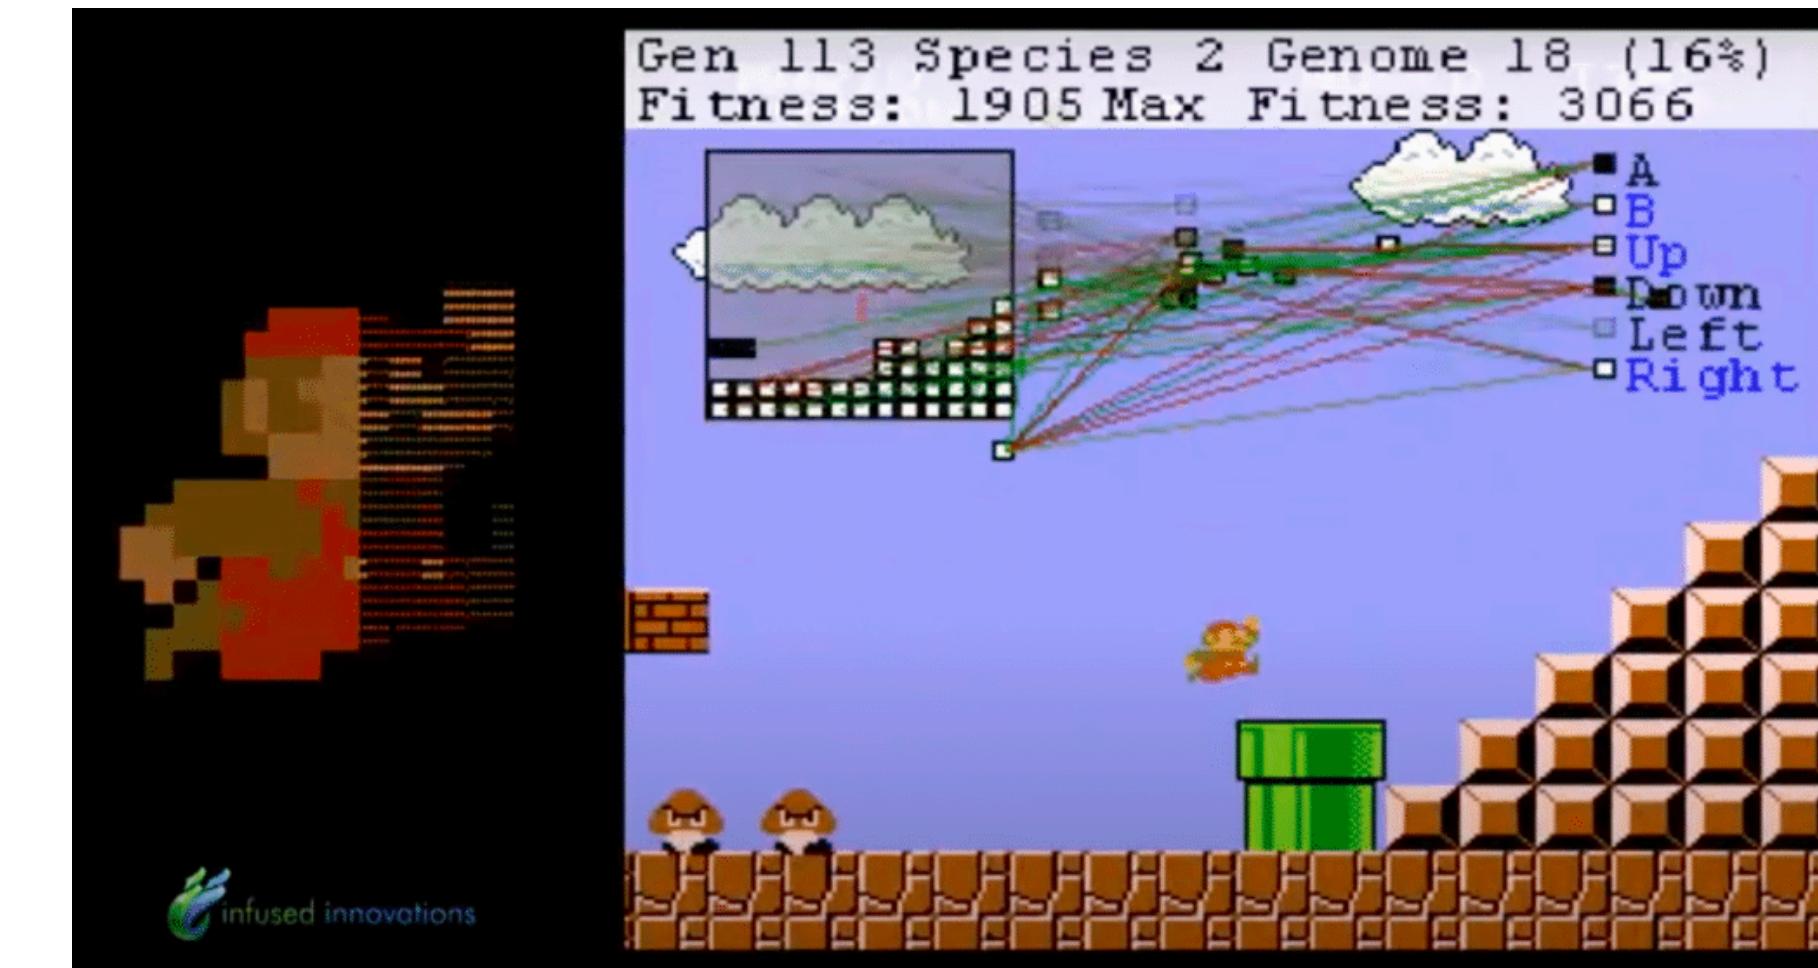
Drug Discovery



Sport Analysis



Sport Analysis



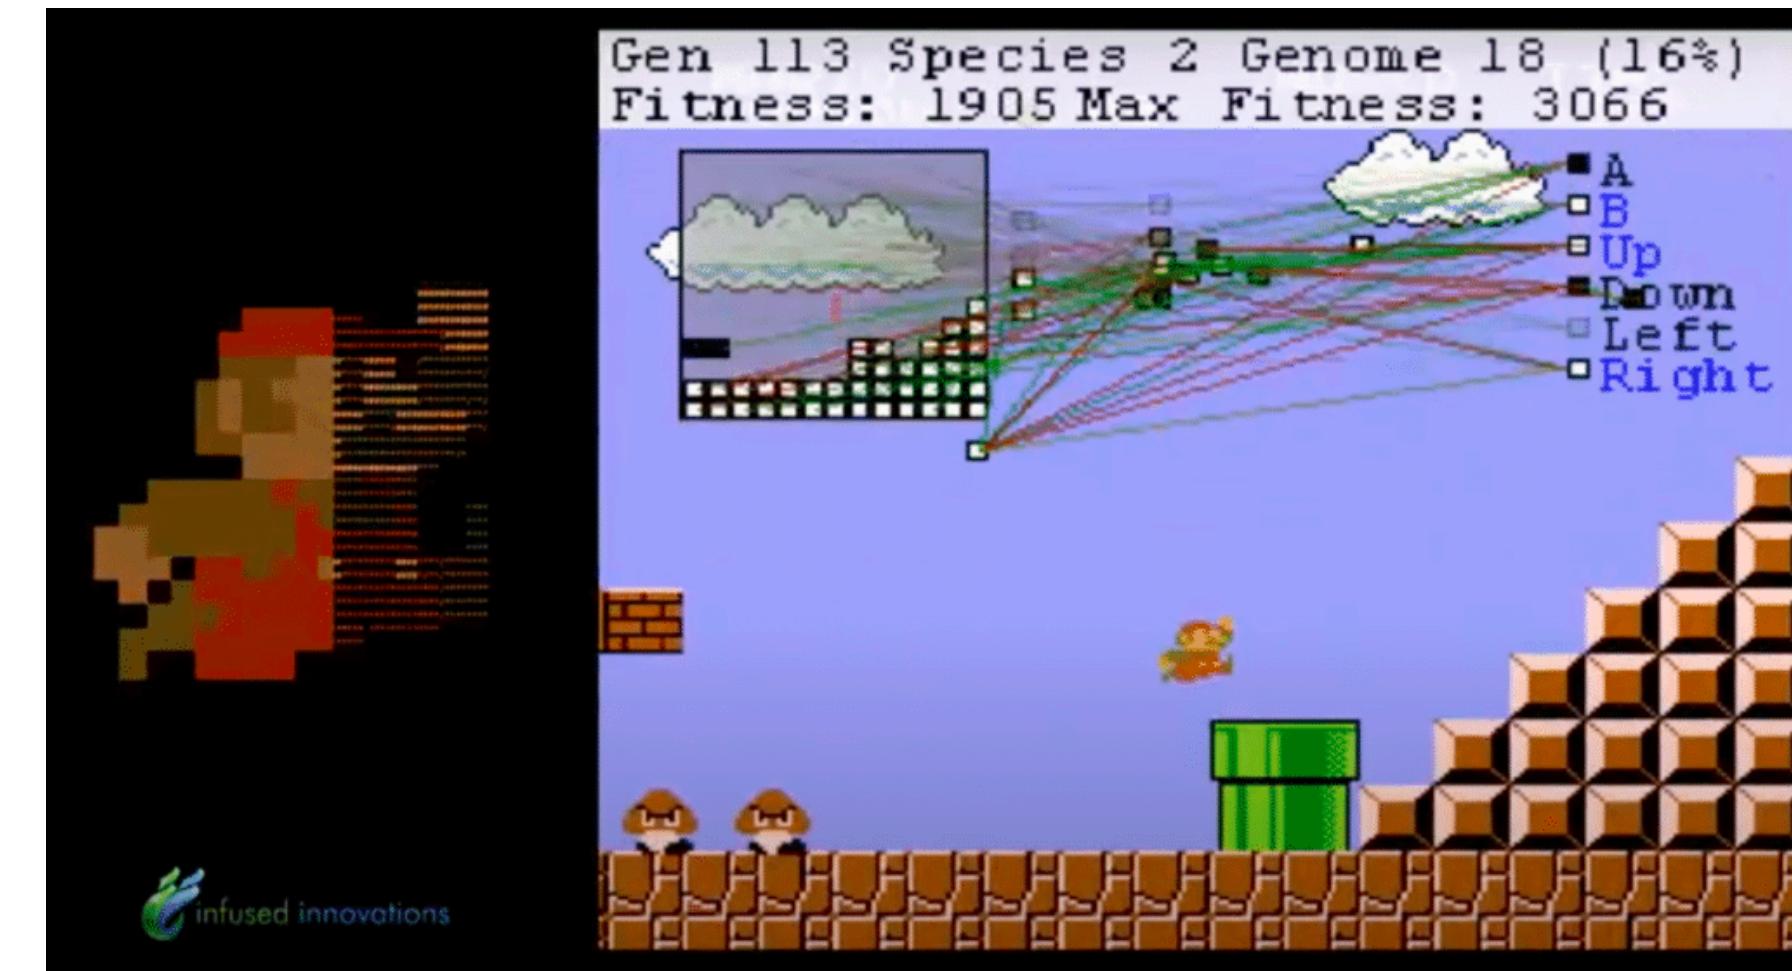
Video Games



Sport Analysis

$$\begin{cases} \frac{dx}{dt} = \sigma [(y(t) - x(t)] \\ \frac{dy}{dt} = \rho x(t) - y(t) - x(t) z(t) \\ \frac{dz}{dt} = x(t) y(t) - \beta z(t) \end{cases}$$

ODE



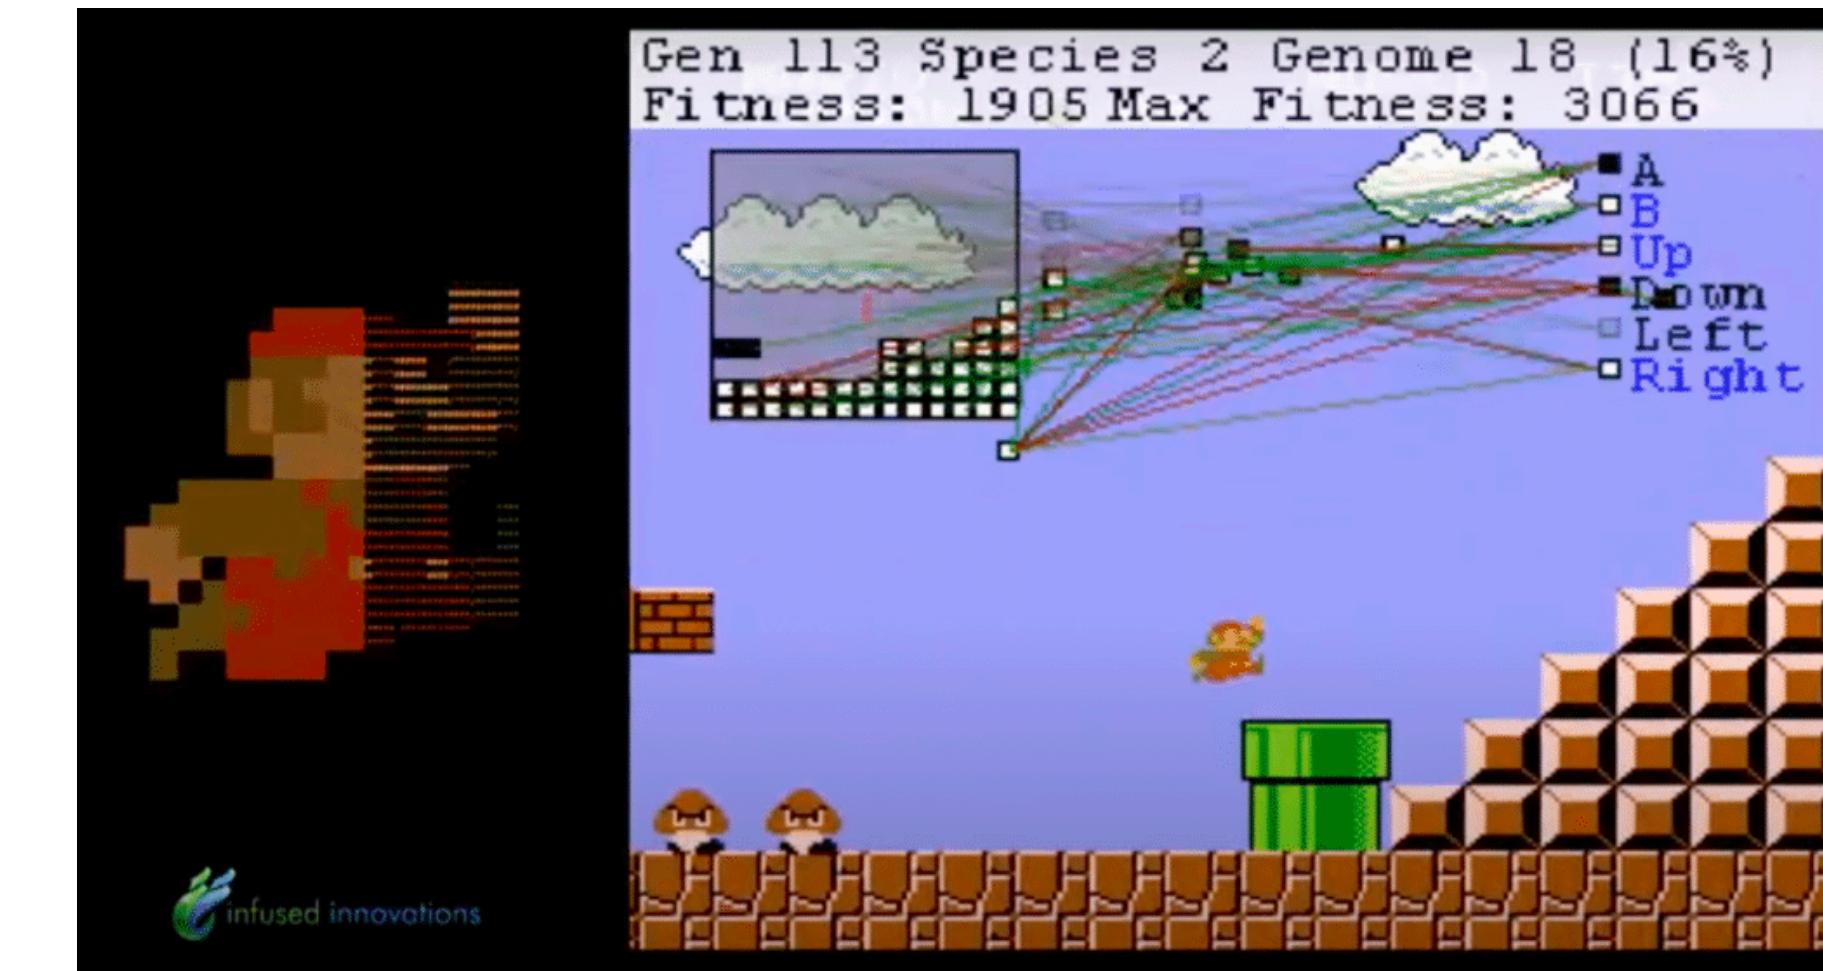
Video Games



Sport Analysis

$$\begin{cases} \frac{dx}{dt} = \sigma [(y(t) - x(t))] \\ \frac{dy}{dt} = \rho x(t) - y(t) - x(t) z(t) \\ \frac{dz}{dt} = x(t) y(t) - \beta z(t) \end{cases}$$

ODE



Video Games



Chess

Why and when to use deep learning in practice ?

Why and when to use deep learning in practice ?

Deep learning rely on the concept of learning from examples.

Why and when to use deep learning in practice ?

Deep learning rely on the concept of learning from examples.

You need « a lot of examples »

Why and when to use deep learning in practice ?

Deep learning rely on the concept of learning from examples.

You need « a lot of examples »

How much in practice?

Why and when to use deep learning in practice ?

Deep learning rely on the concept of learning from examples.

You need « a lot of examples »

How much in practice?

Open research question

Why and when to use deep learning in practice ?

Deep learning rely on the concept of learning from examples.

You need « a lot of examples »

How much in practice?

Open research question

Why should you use deep learning?

Hand-crafted features are time consuming, not scalable and in practice sub-optimal

Deep Learning aims at automatically learning features directly from data!

Why and when to use deep learning in practice ?

Deep learning rely on the concept of learning from examples.

You need « a lot of examples »

How much in practice?

Open research question

Why should you use deep learning?

Hand-crafted features are **time consuming, not scalable and in practice sub-optimal**

Deep Learning aims at automatically learning **features** directly from data!



Pattern,
discriminative characteristics

Why and when to use deep learning in practice ?

Deep learning rely on the concept of **learning from examples**.

You need « a lot of examples »

How much in practice?

Open research question

Why should you use deep learning?

Hand-crafted features are **time consuming, not scalable and in practice sub-optimal**

Deep Learning aims at automatically learning **features** directly from data!

MNIST

Pattern,
discriminative characteristics



Why and when to use deep learning in practice ?

Deep learning rely on the concept of **learning from examples**.

You need « a lot of examples »

How much in practice?

Open research question

Why should you use deep learning?

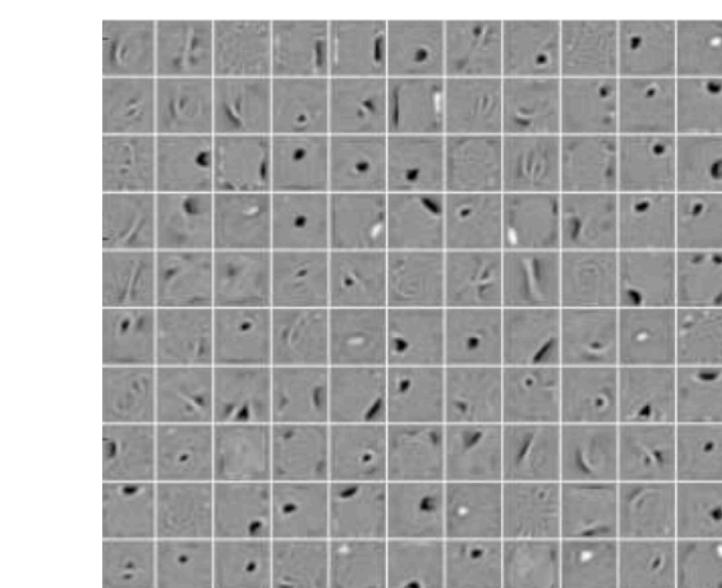
Hand-crafted features are **time consuming, not scalable and in practice sub-optimal**

Deep Learning aims at automatically learning **features** directly from data!

MNIST



Features Learnt by the NN



Pattern,
discriminative characteristics

Why and when to use deep learning in practice ?

Deep learning rely on the concept of **learning from examples**.

You need « a lot of examples »

How much in practice?

Open research question

Why should you use deep learning?

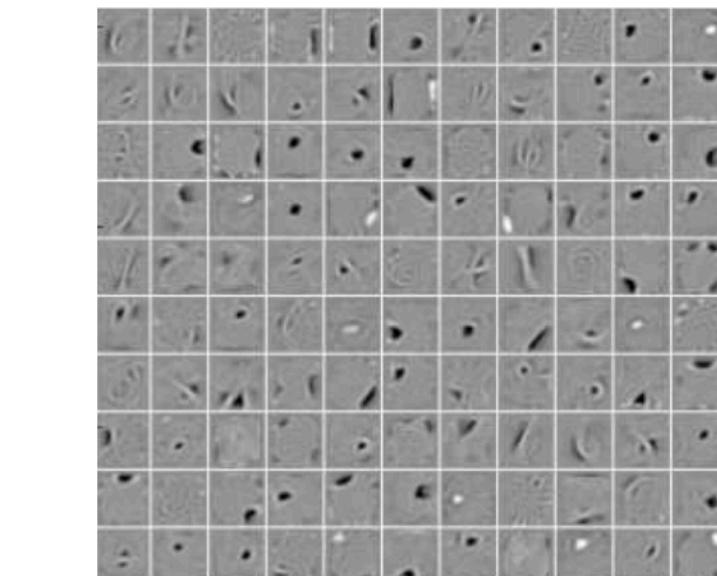
Hand-crafted features are **time consuming, not scalable and in practice sub-optimal**

Deep Learning aims at automatically learning **features** directly from data!

MNIST



Features Learnt by the NN



Pattern,
discriminative characteristics

These features
are hard to guess !

A brief history of deep learning



A brief history of deep learning



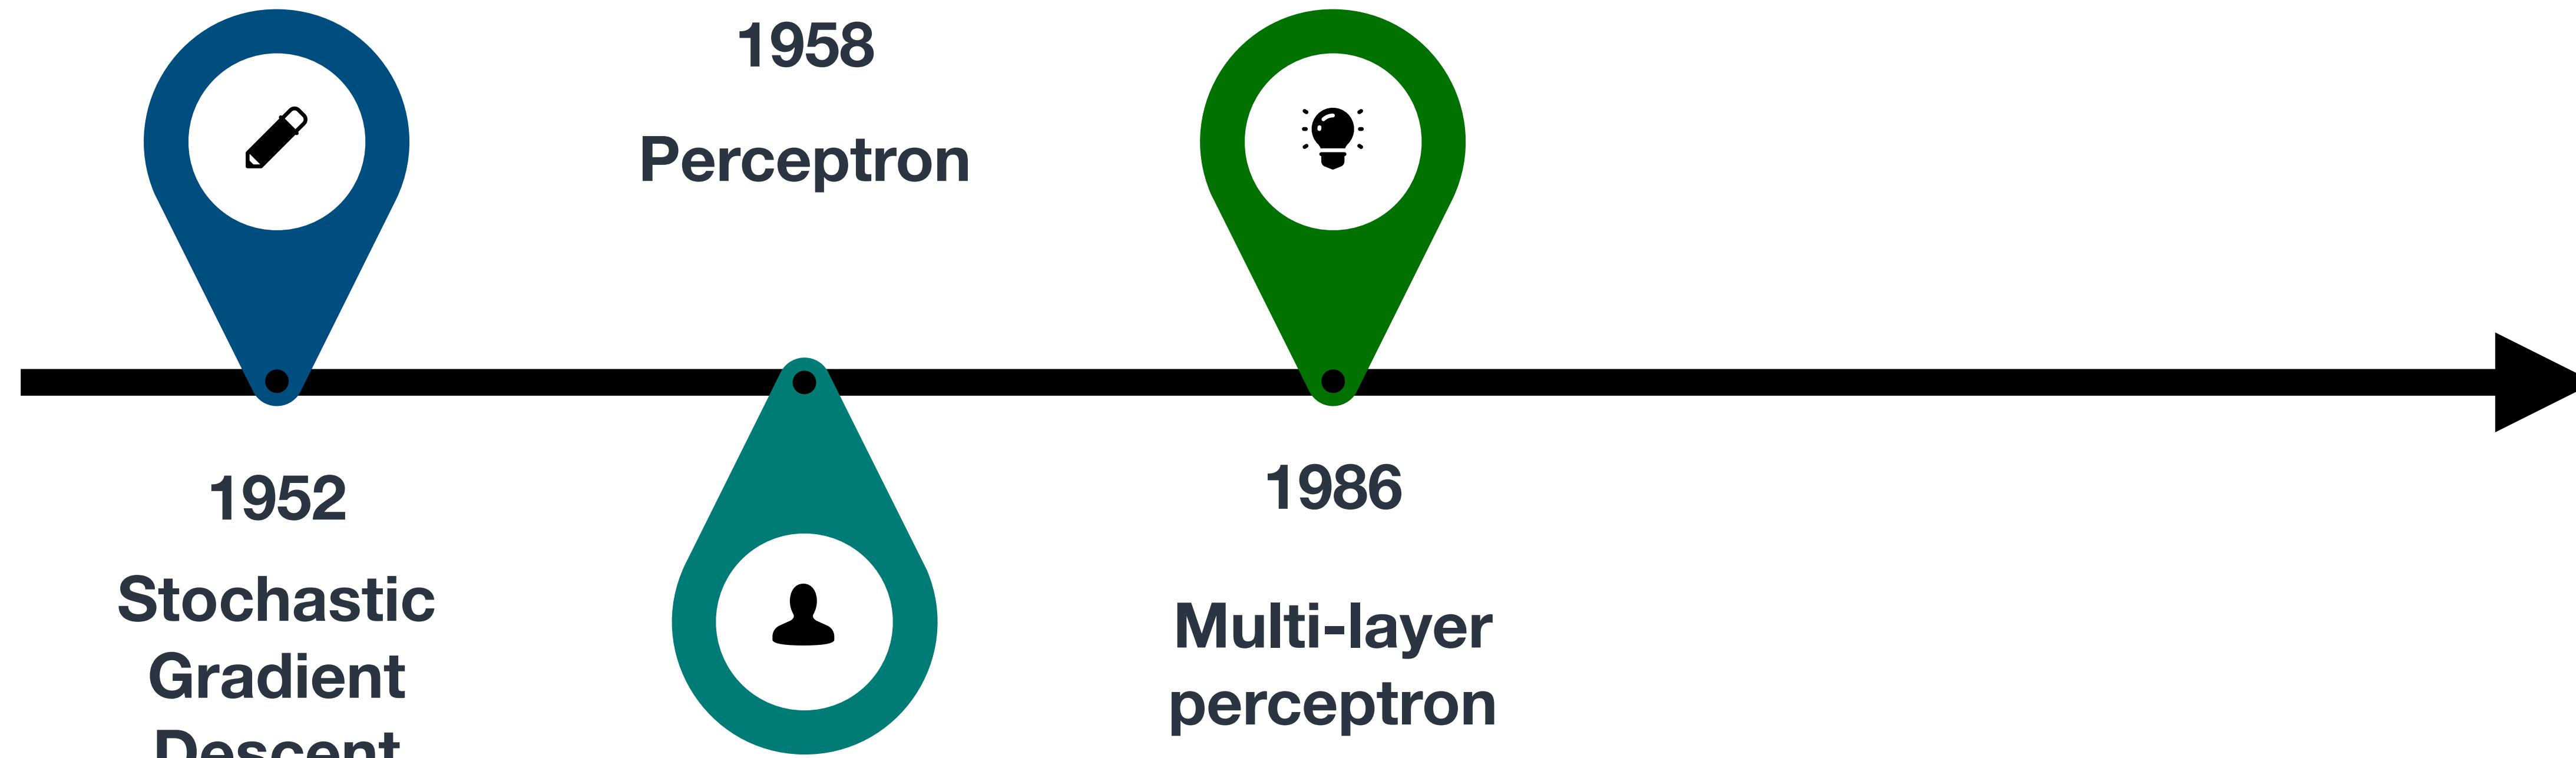
1952

**Stochastic
Gradient
Descent**

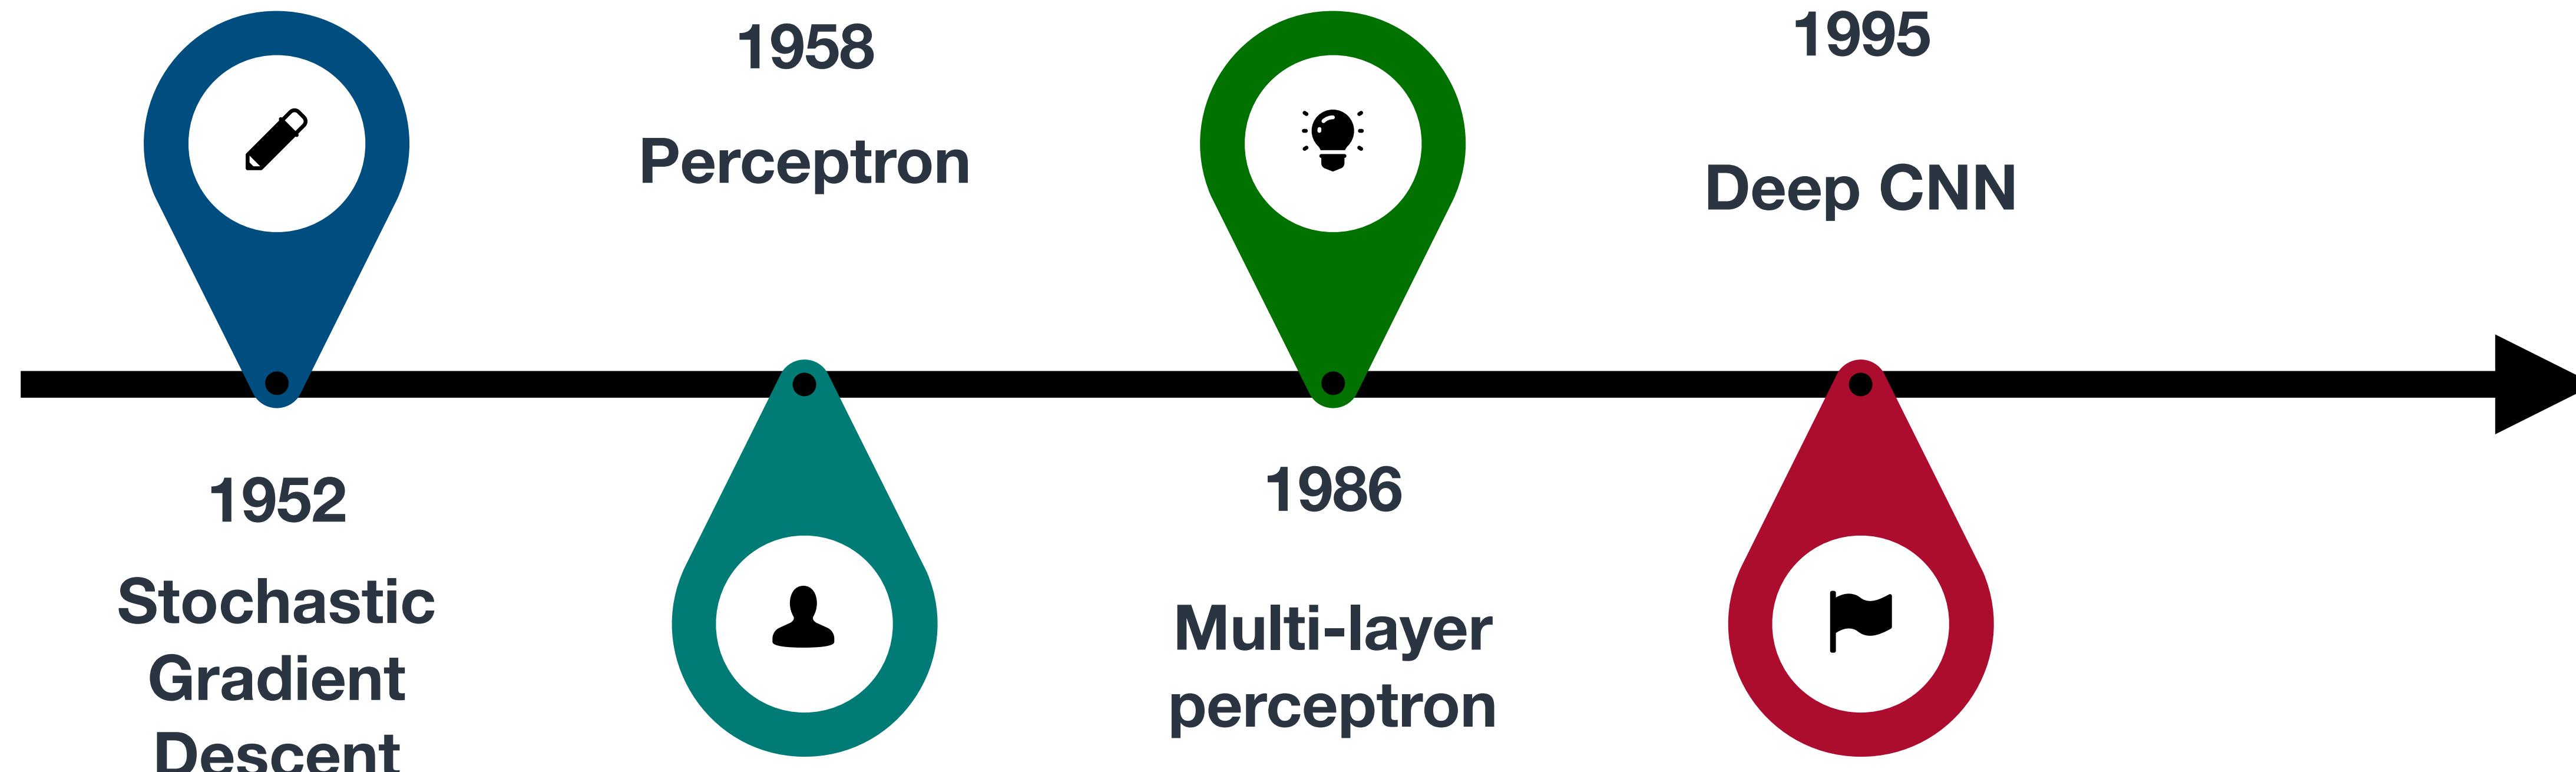
A brief history of deep learning



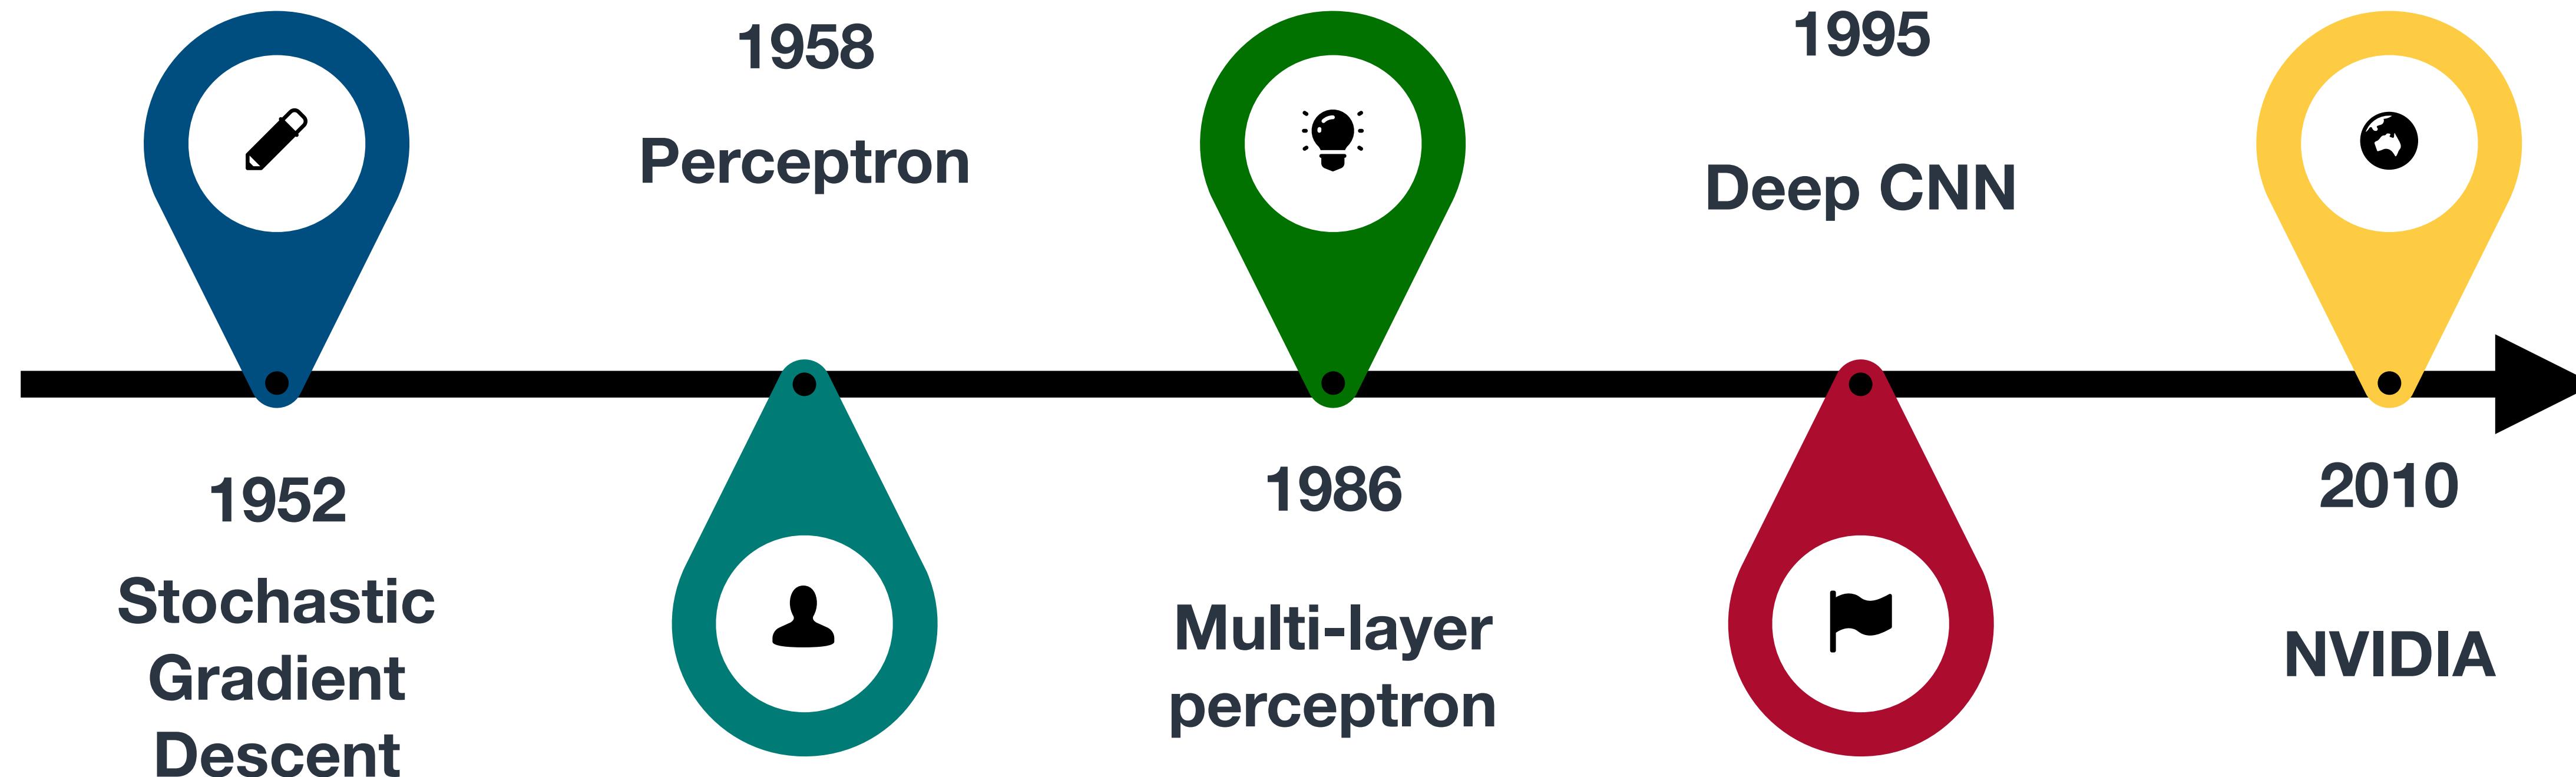
A brief history of deep learning



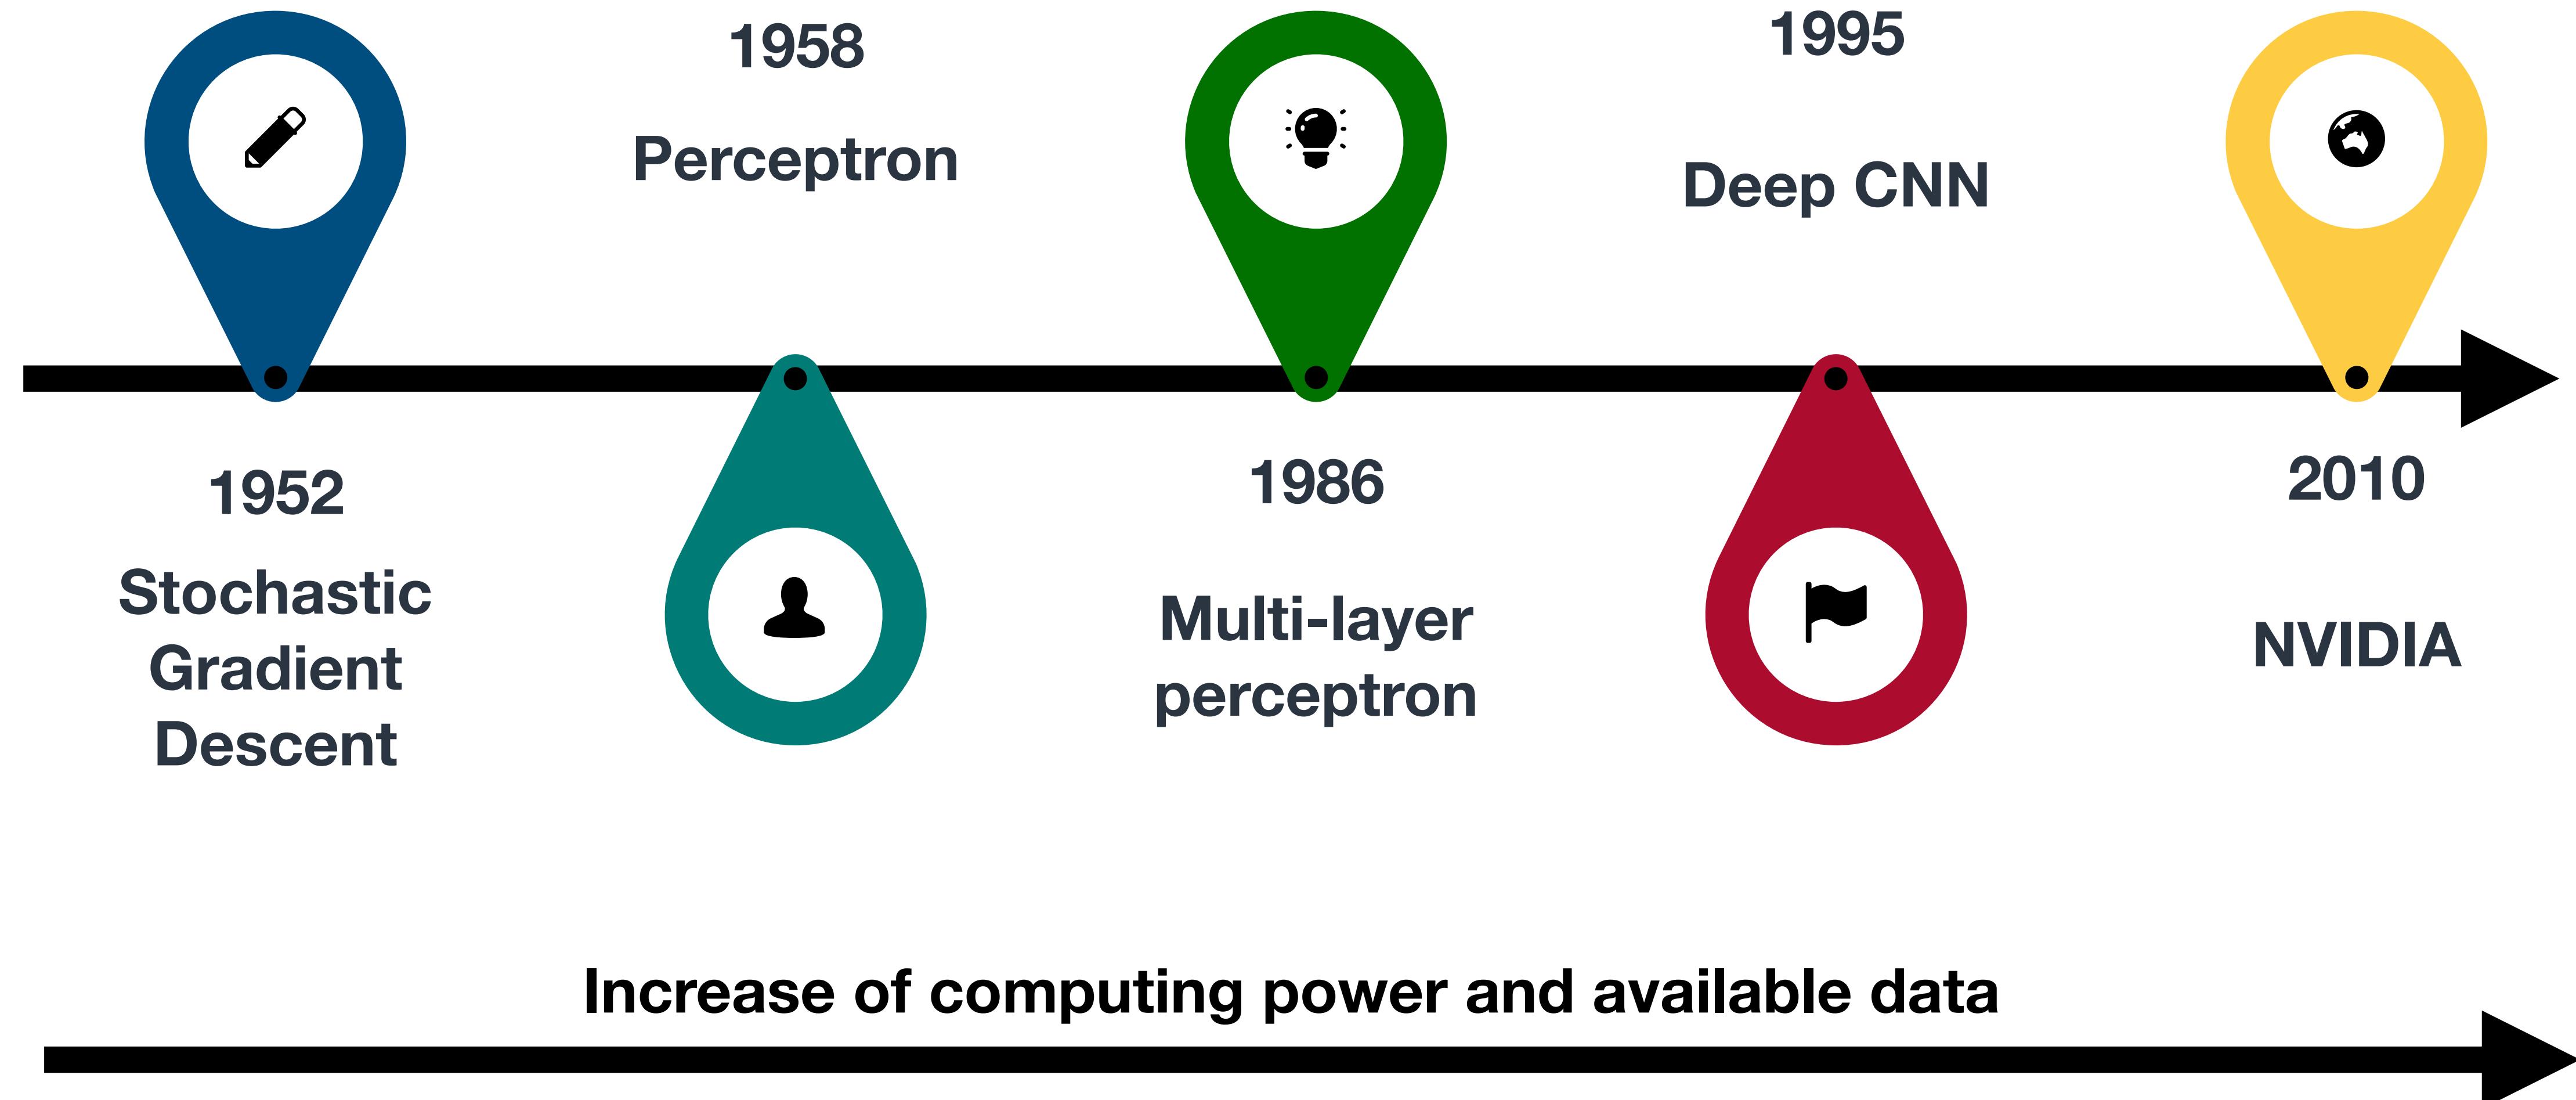
A brief history of deep learning



A brief history of deep learning



A brief history of deep learning



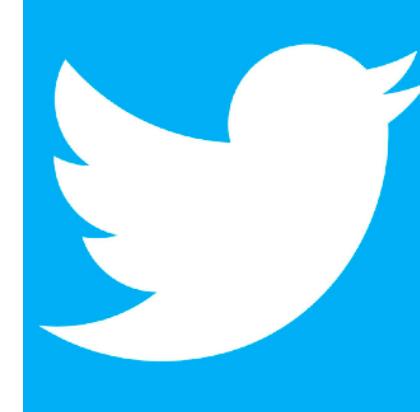
Why Now ?

Deep learning date back decades, so why now?

With internet we have access to massive amount of data



WIKIPÉDIA
L'encyclopédie libre



With GPUs we have access to massive amount of GPUs

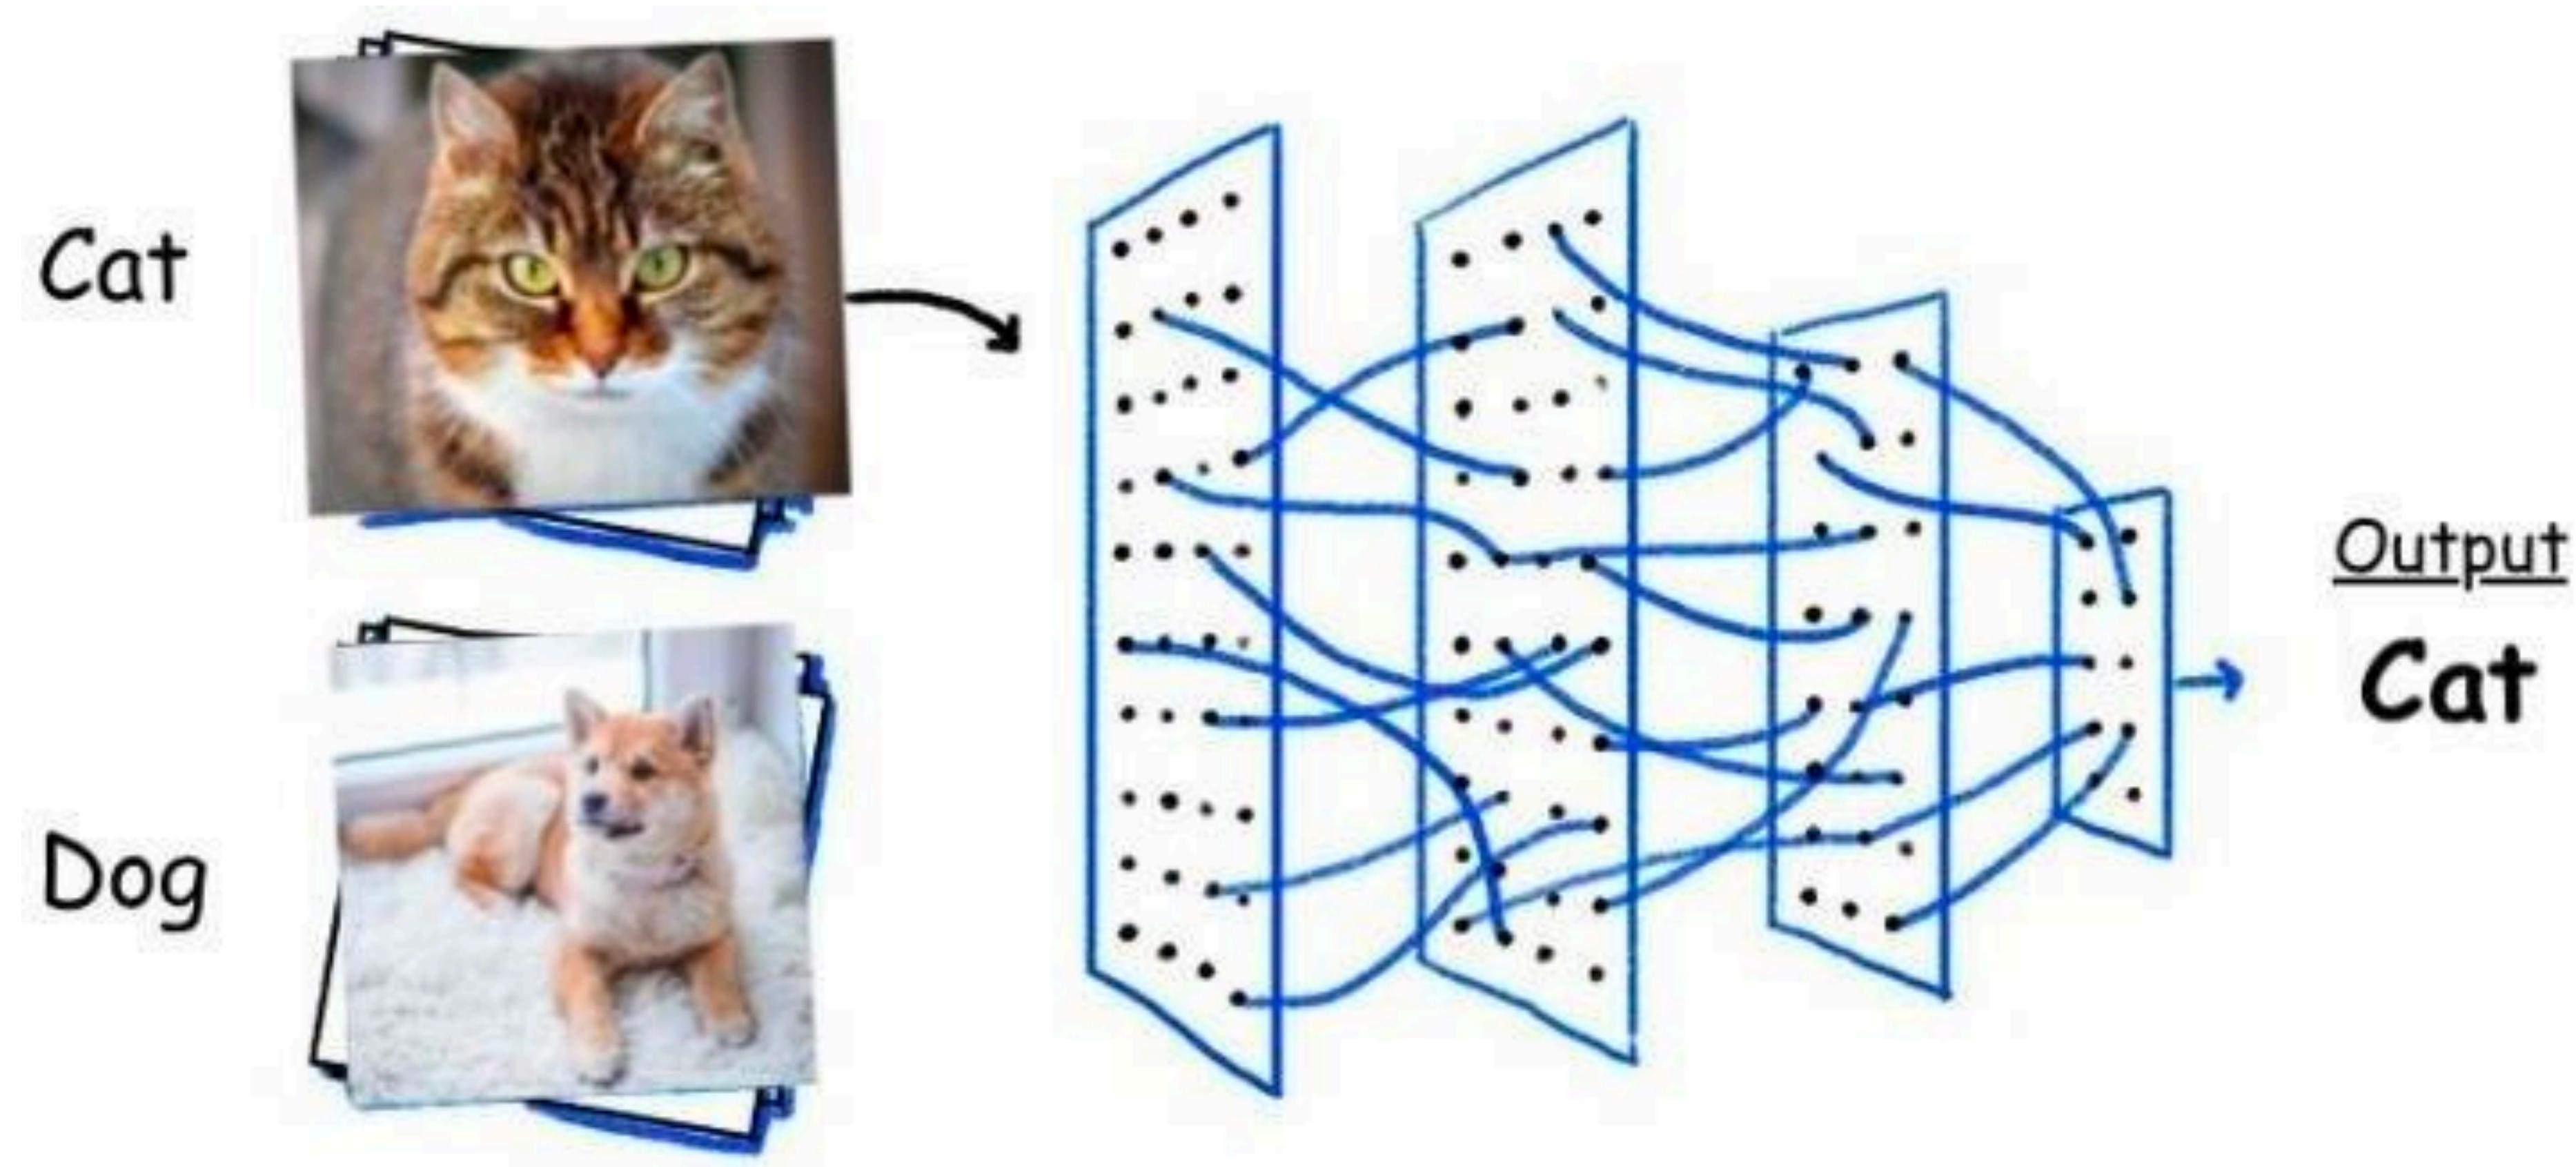
Jeanzay etc..

With the huge effort of Google, Facebook and others we have access to high level Software

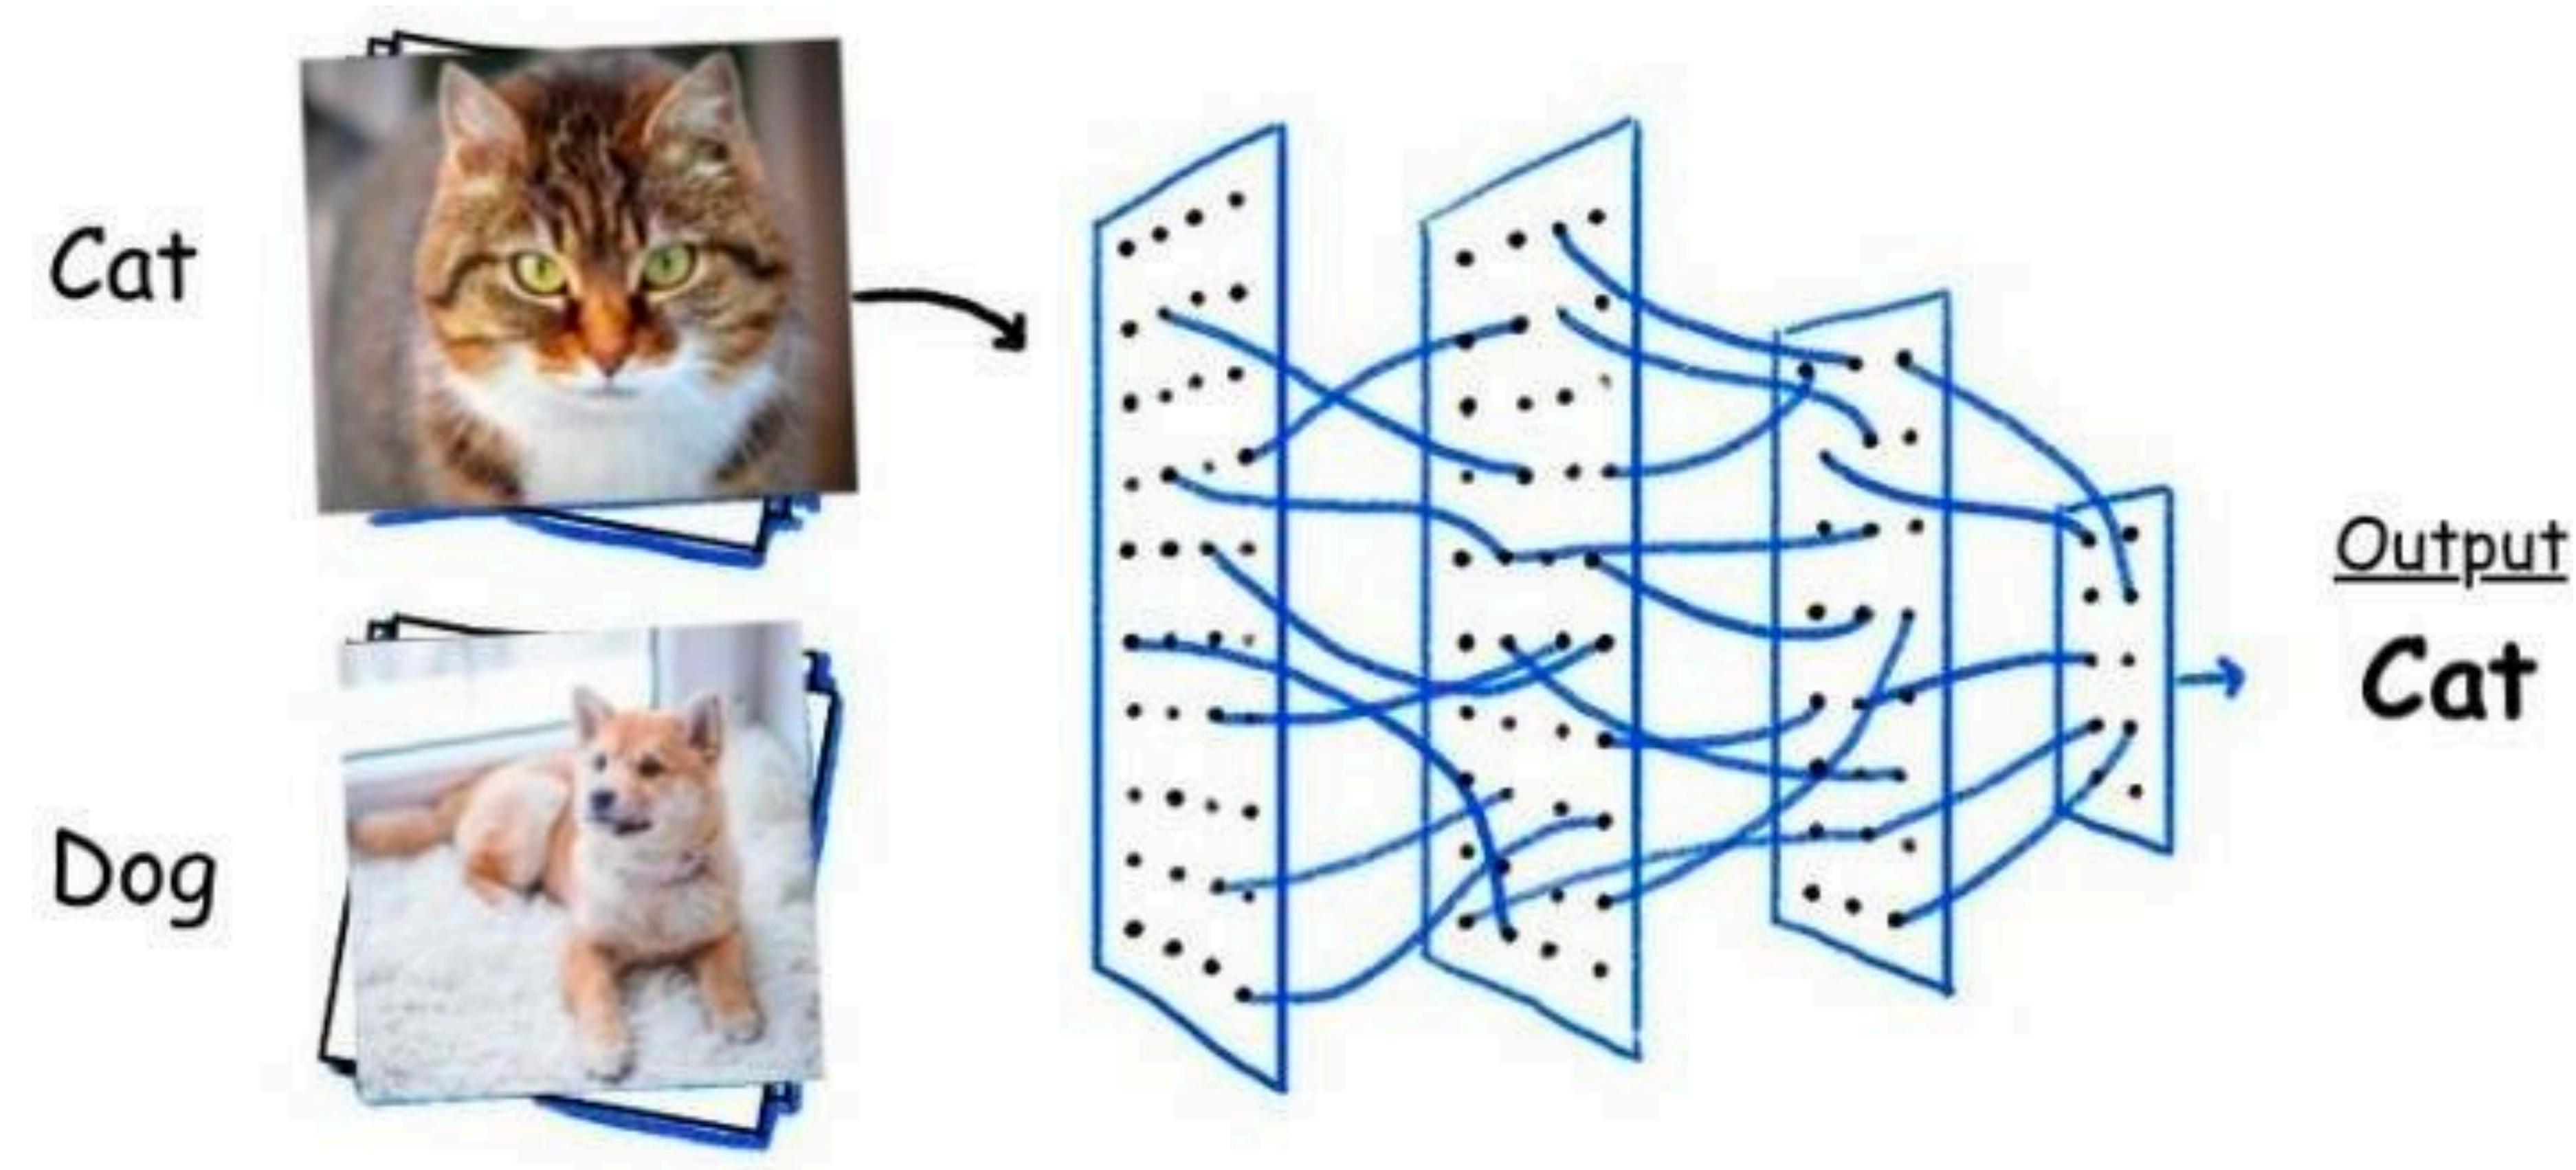


Example of Task to solve with a neural network (I)

Example of Task to solve with a neural network (I)



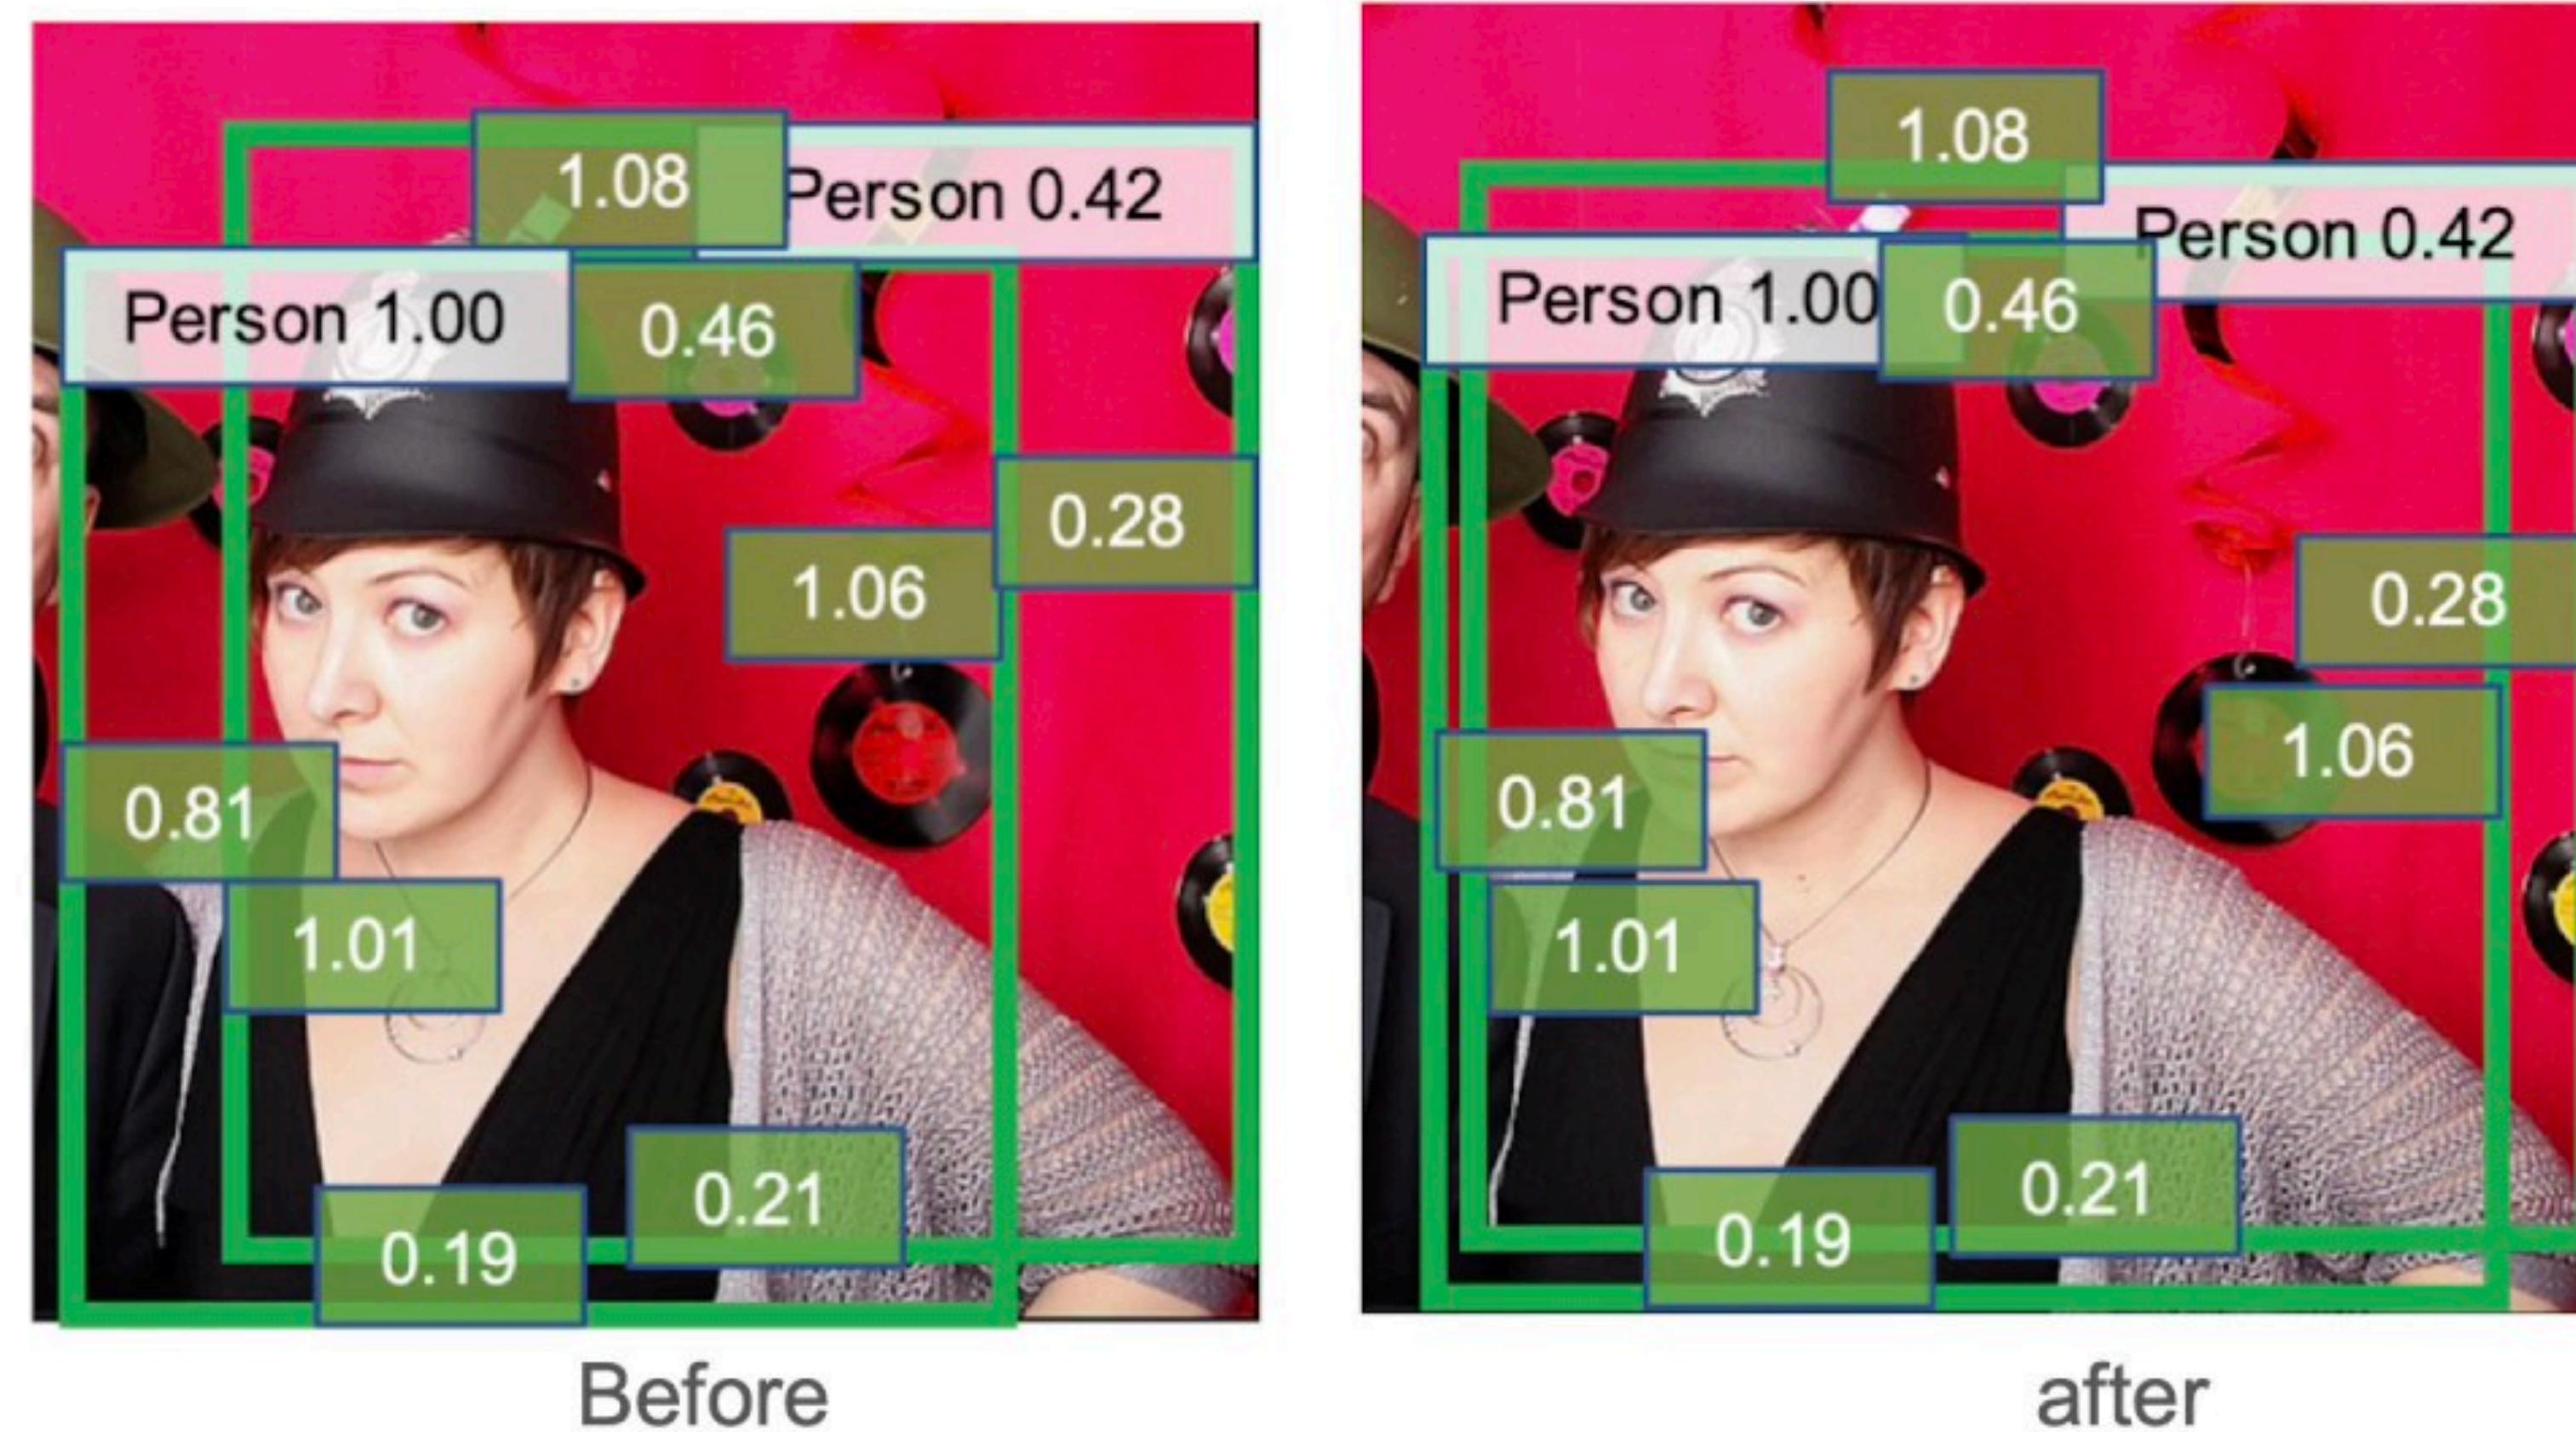
Example of Task to solve with a neural network (I)



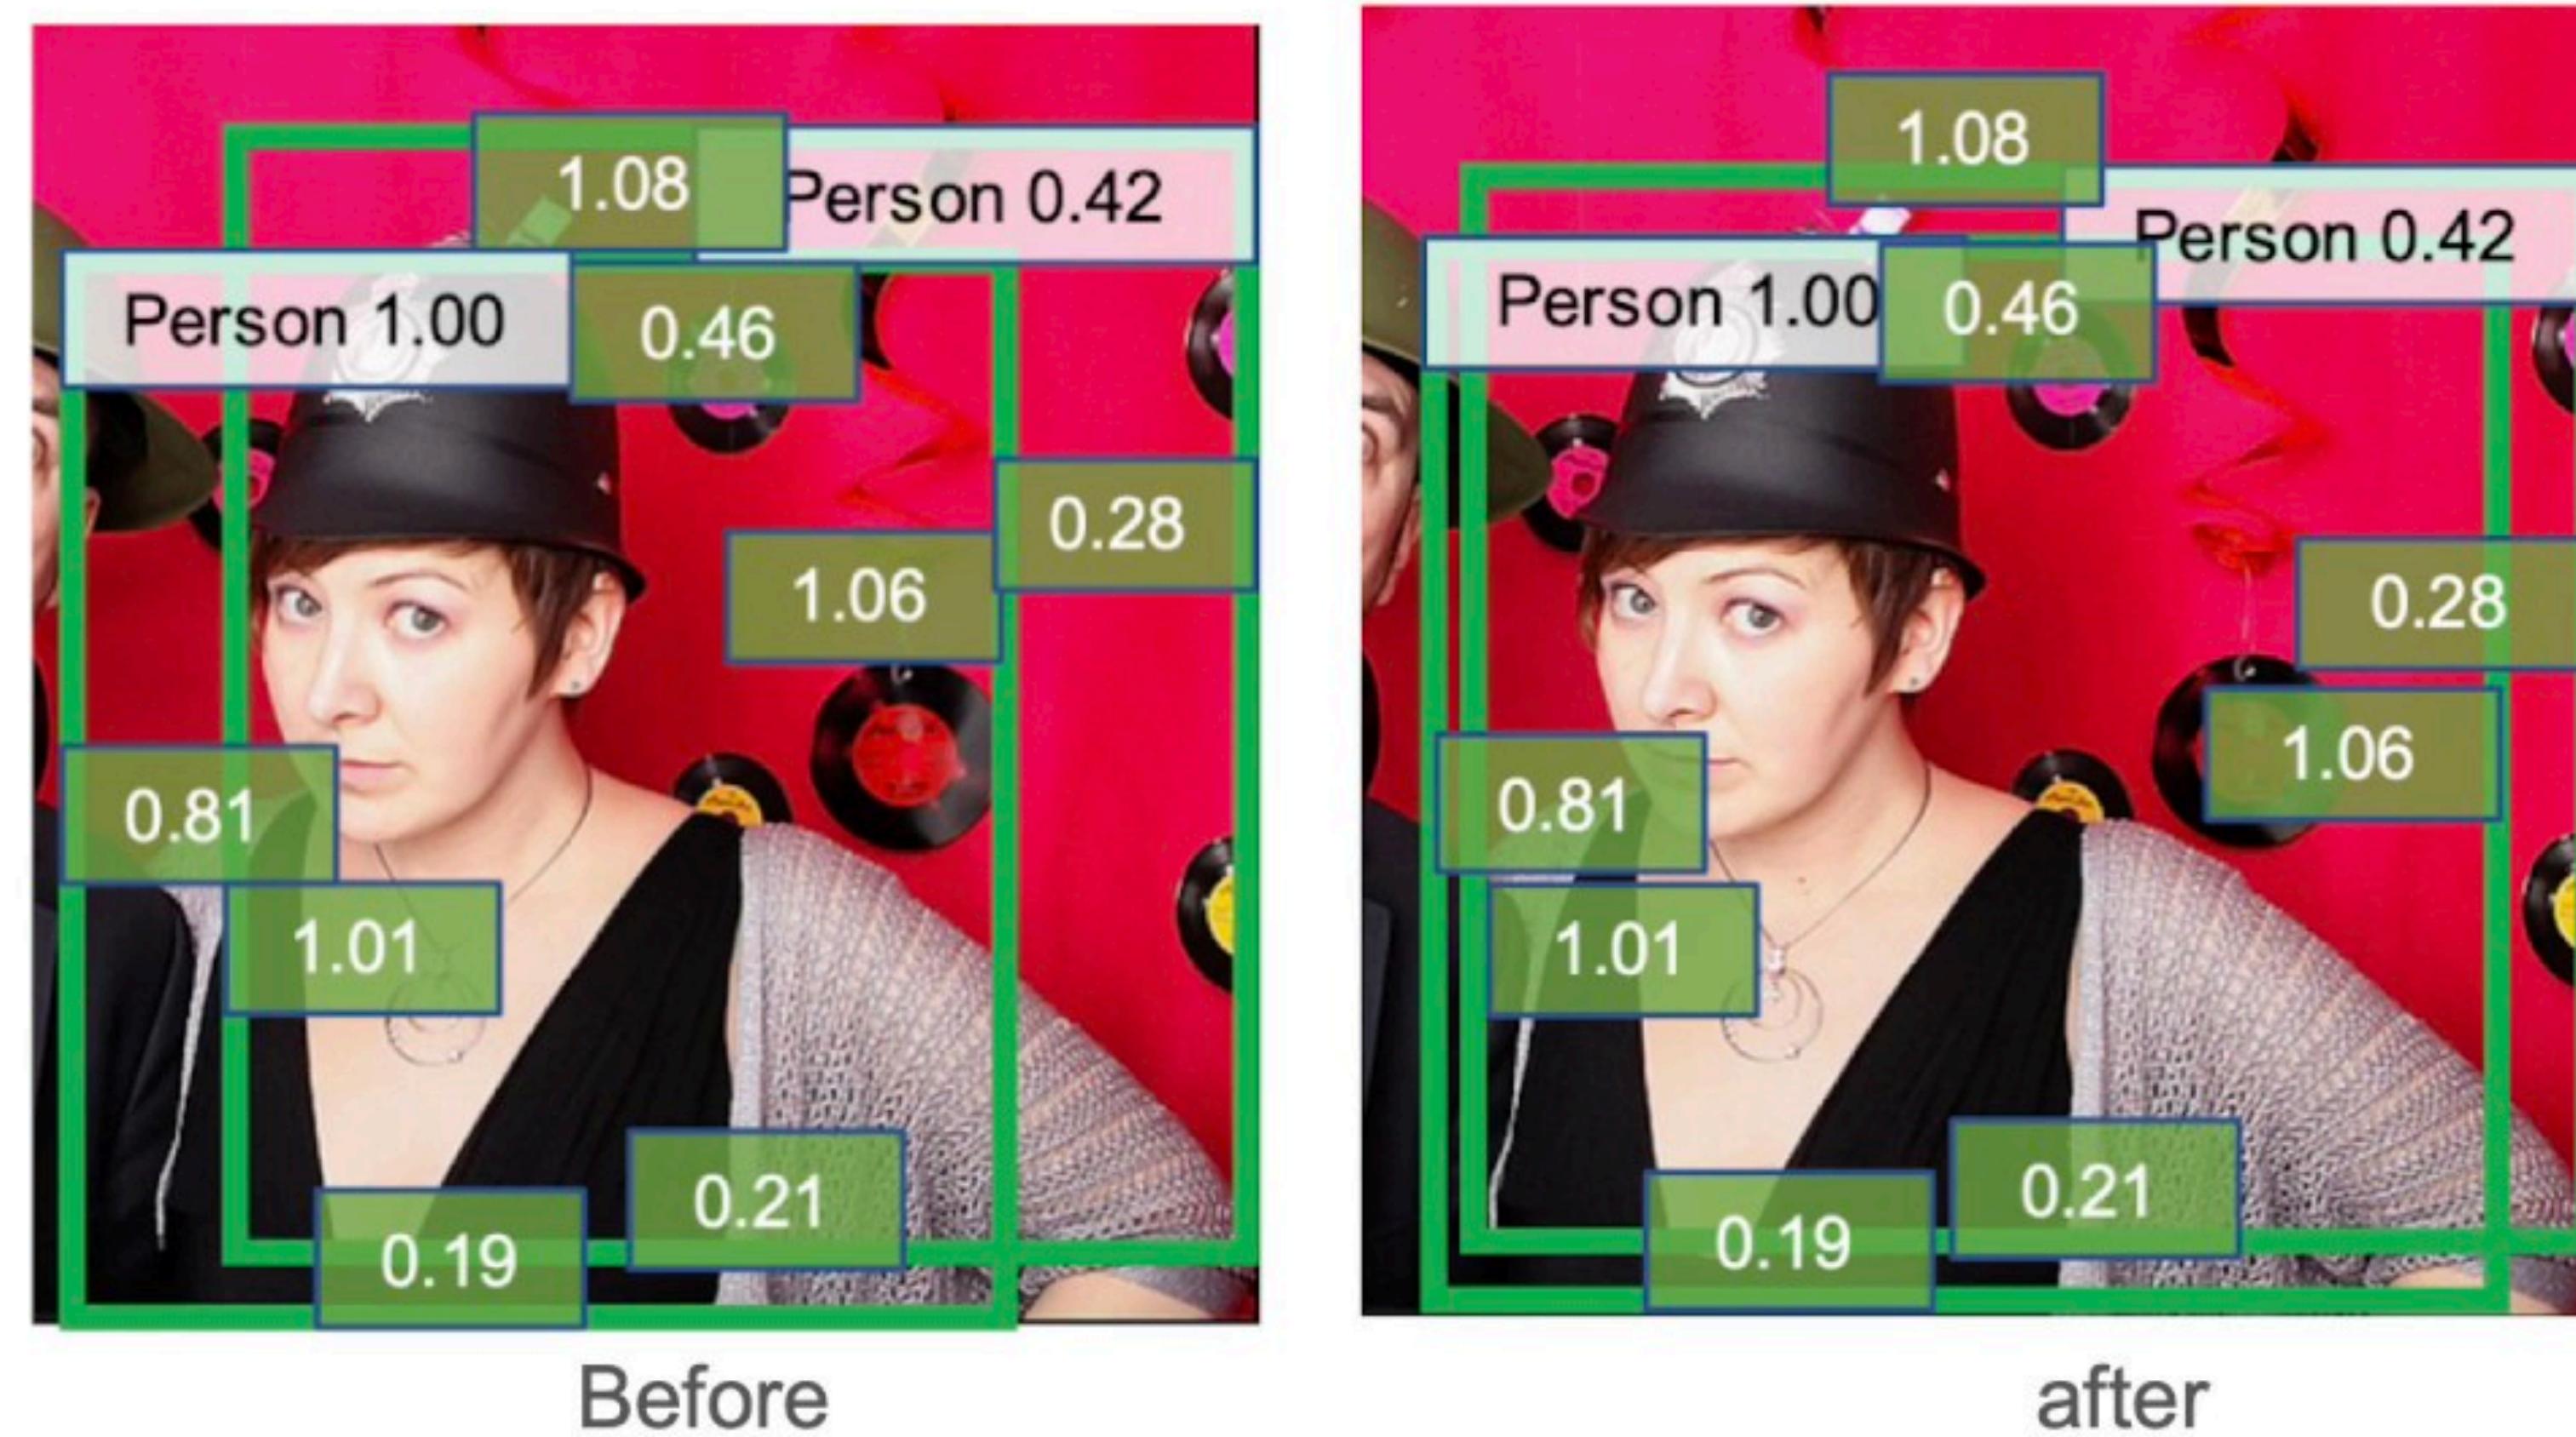
NN can do that without any image processing background

Example of Task to solve with a neural network (II)

Example of Task to solve with a neural network (II)



Example of Task to solve with a neural network (II)



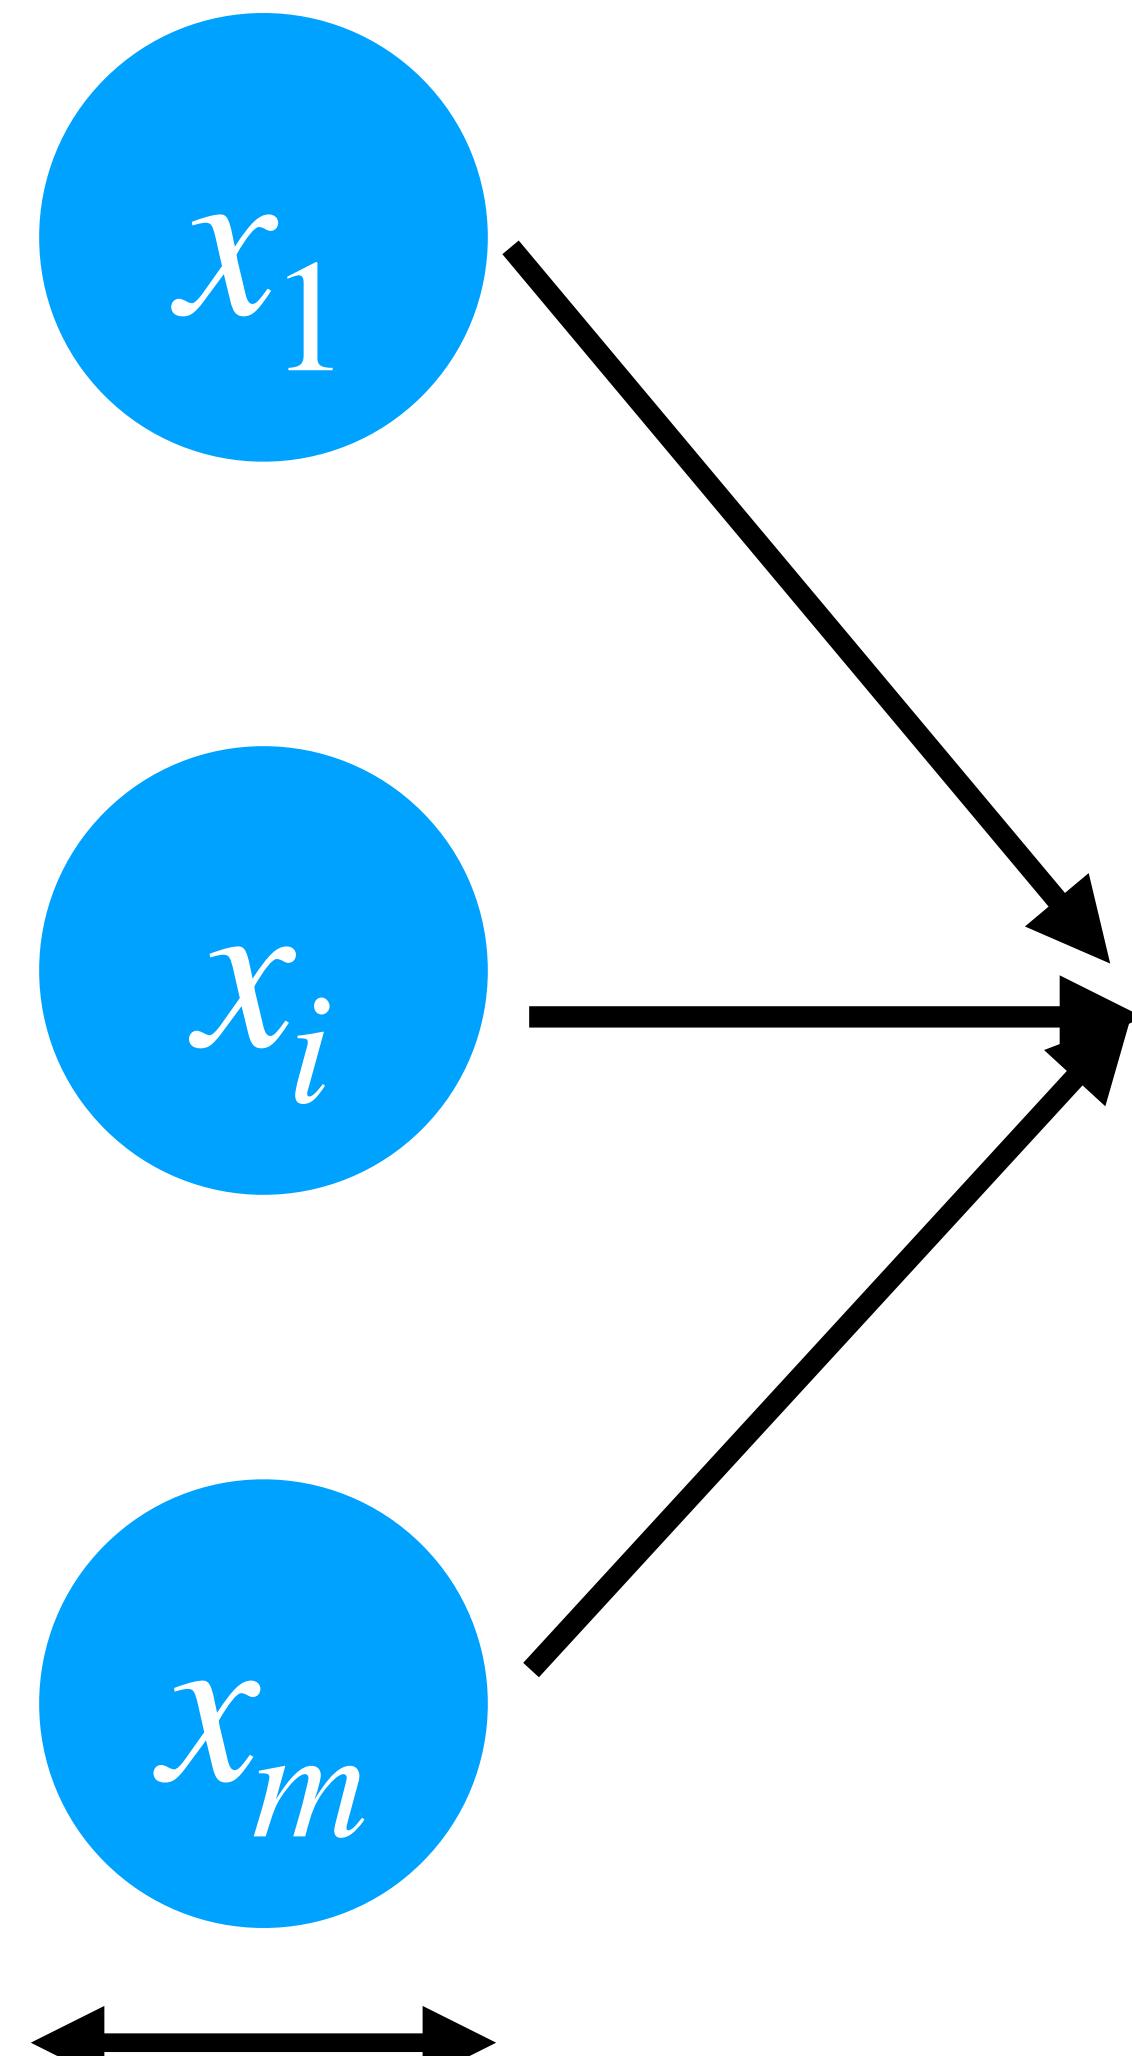
NN can do that without any image processing background

**Design the architecture
of the neural network**



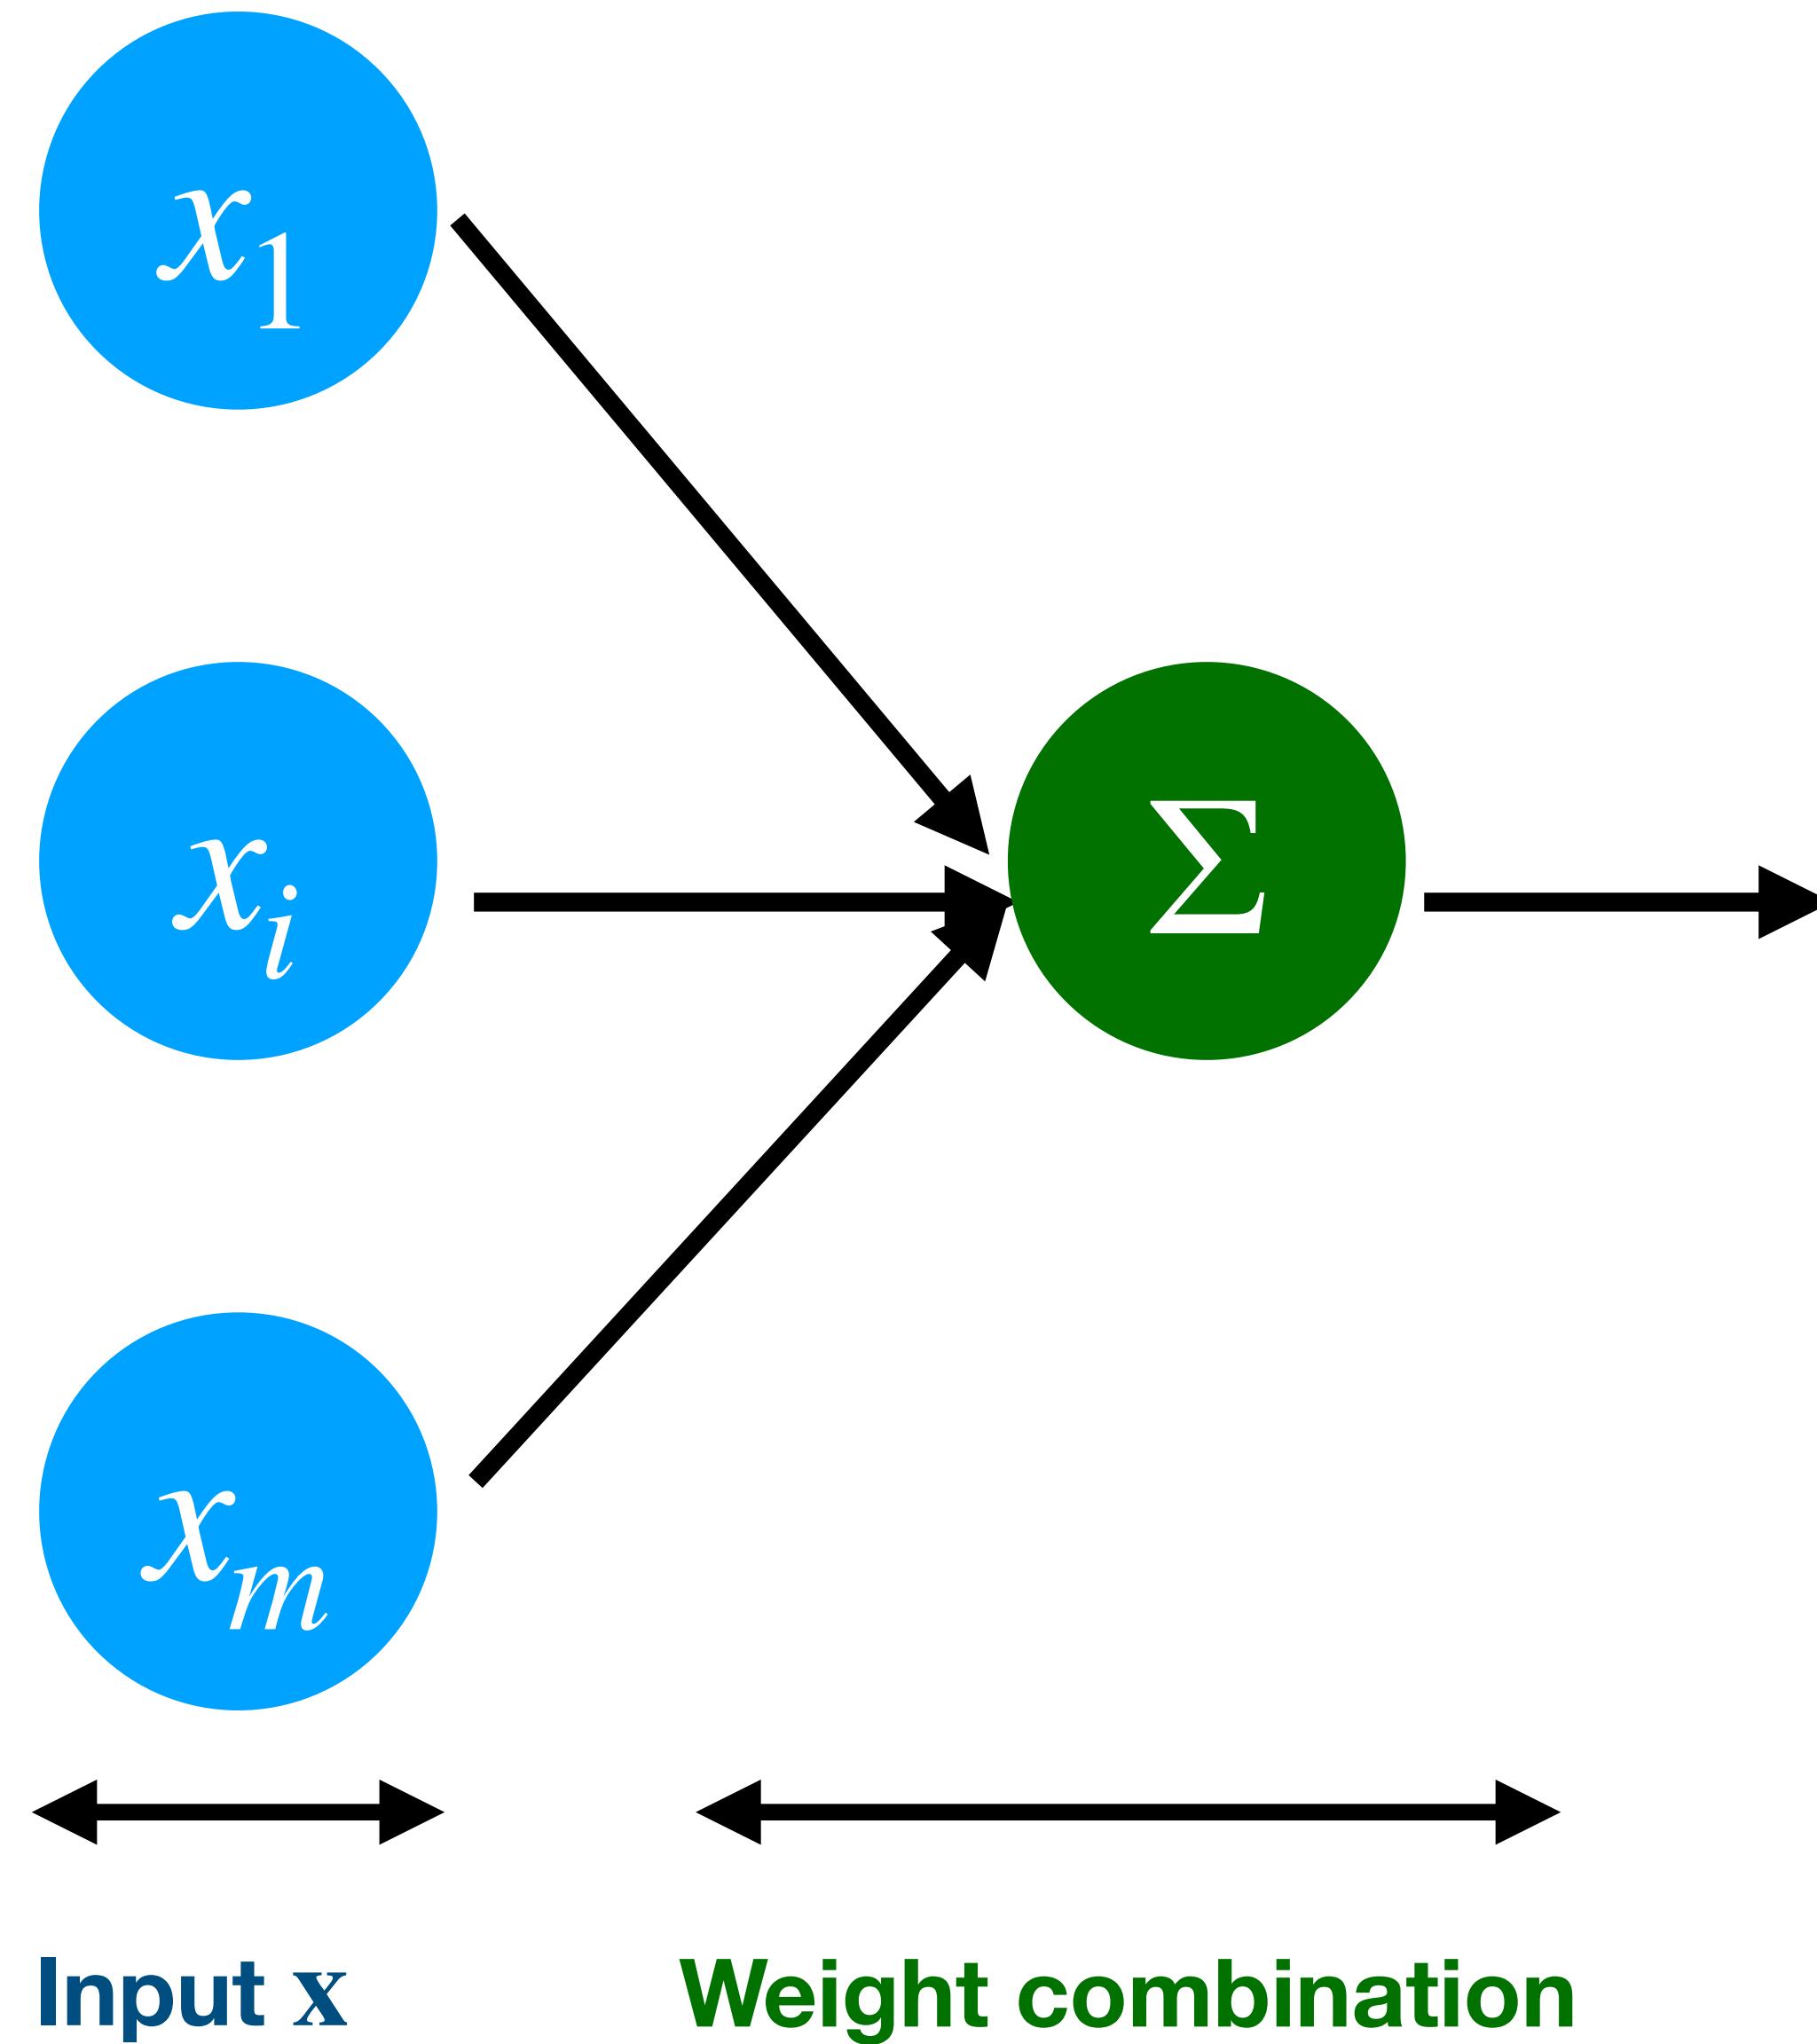
The Perceptron: the most basic neural network

The Perceptron: the most basic neural network

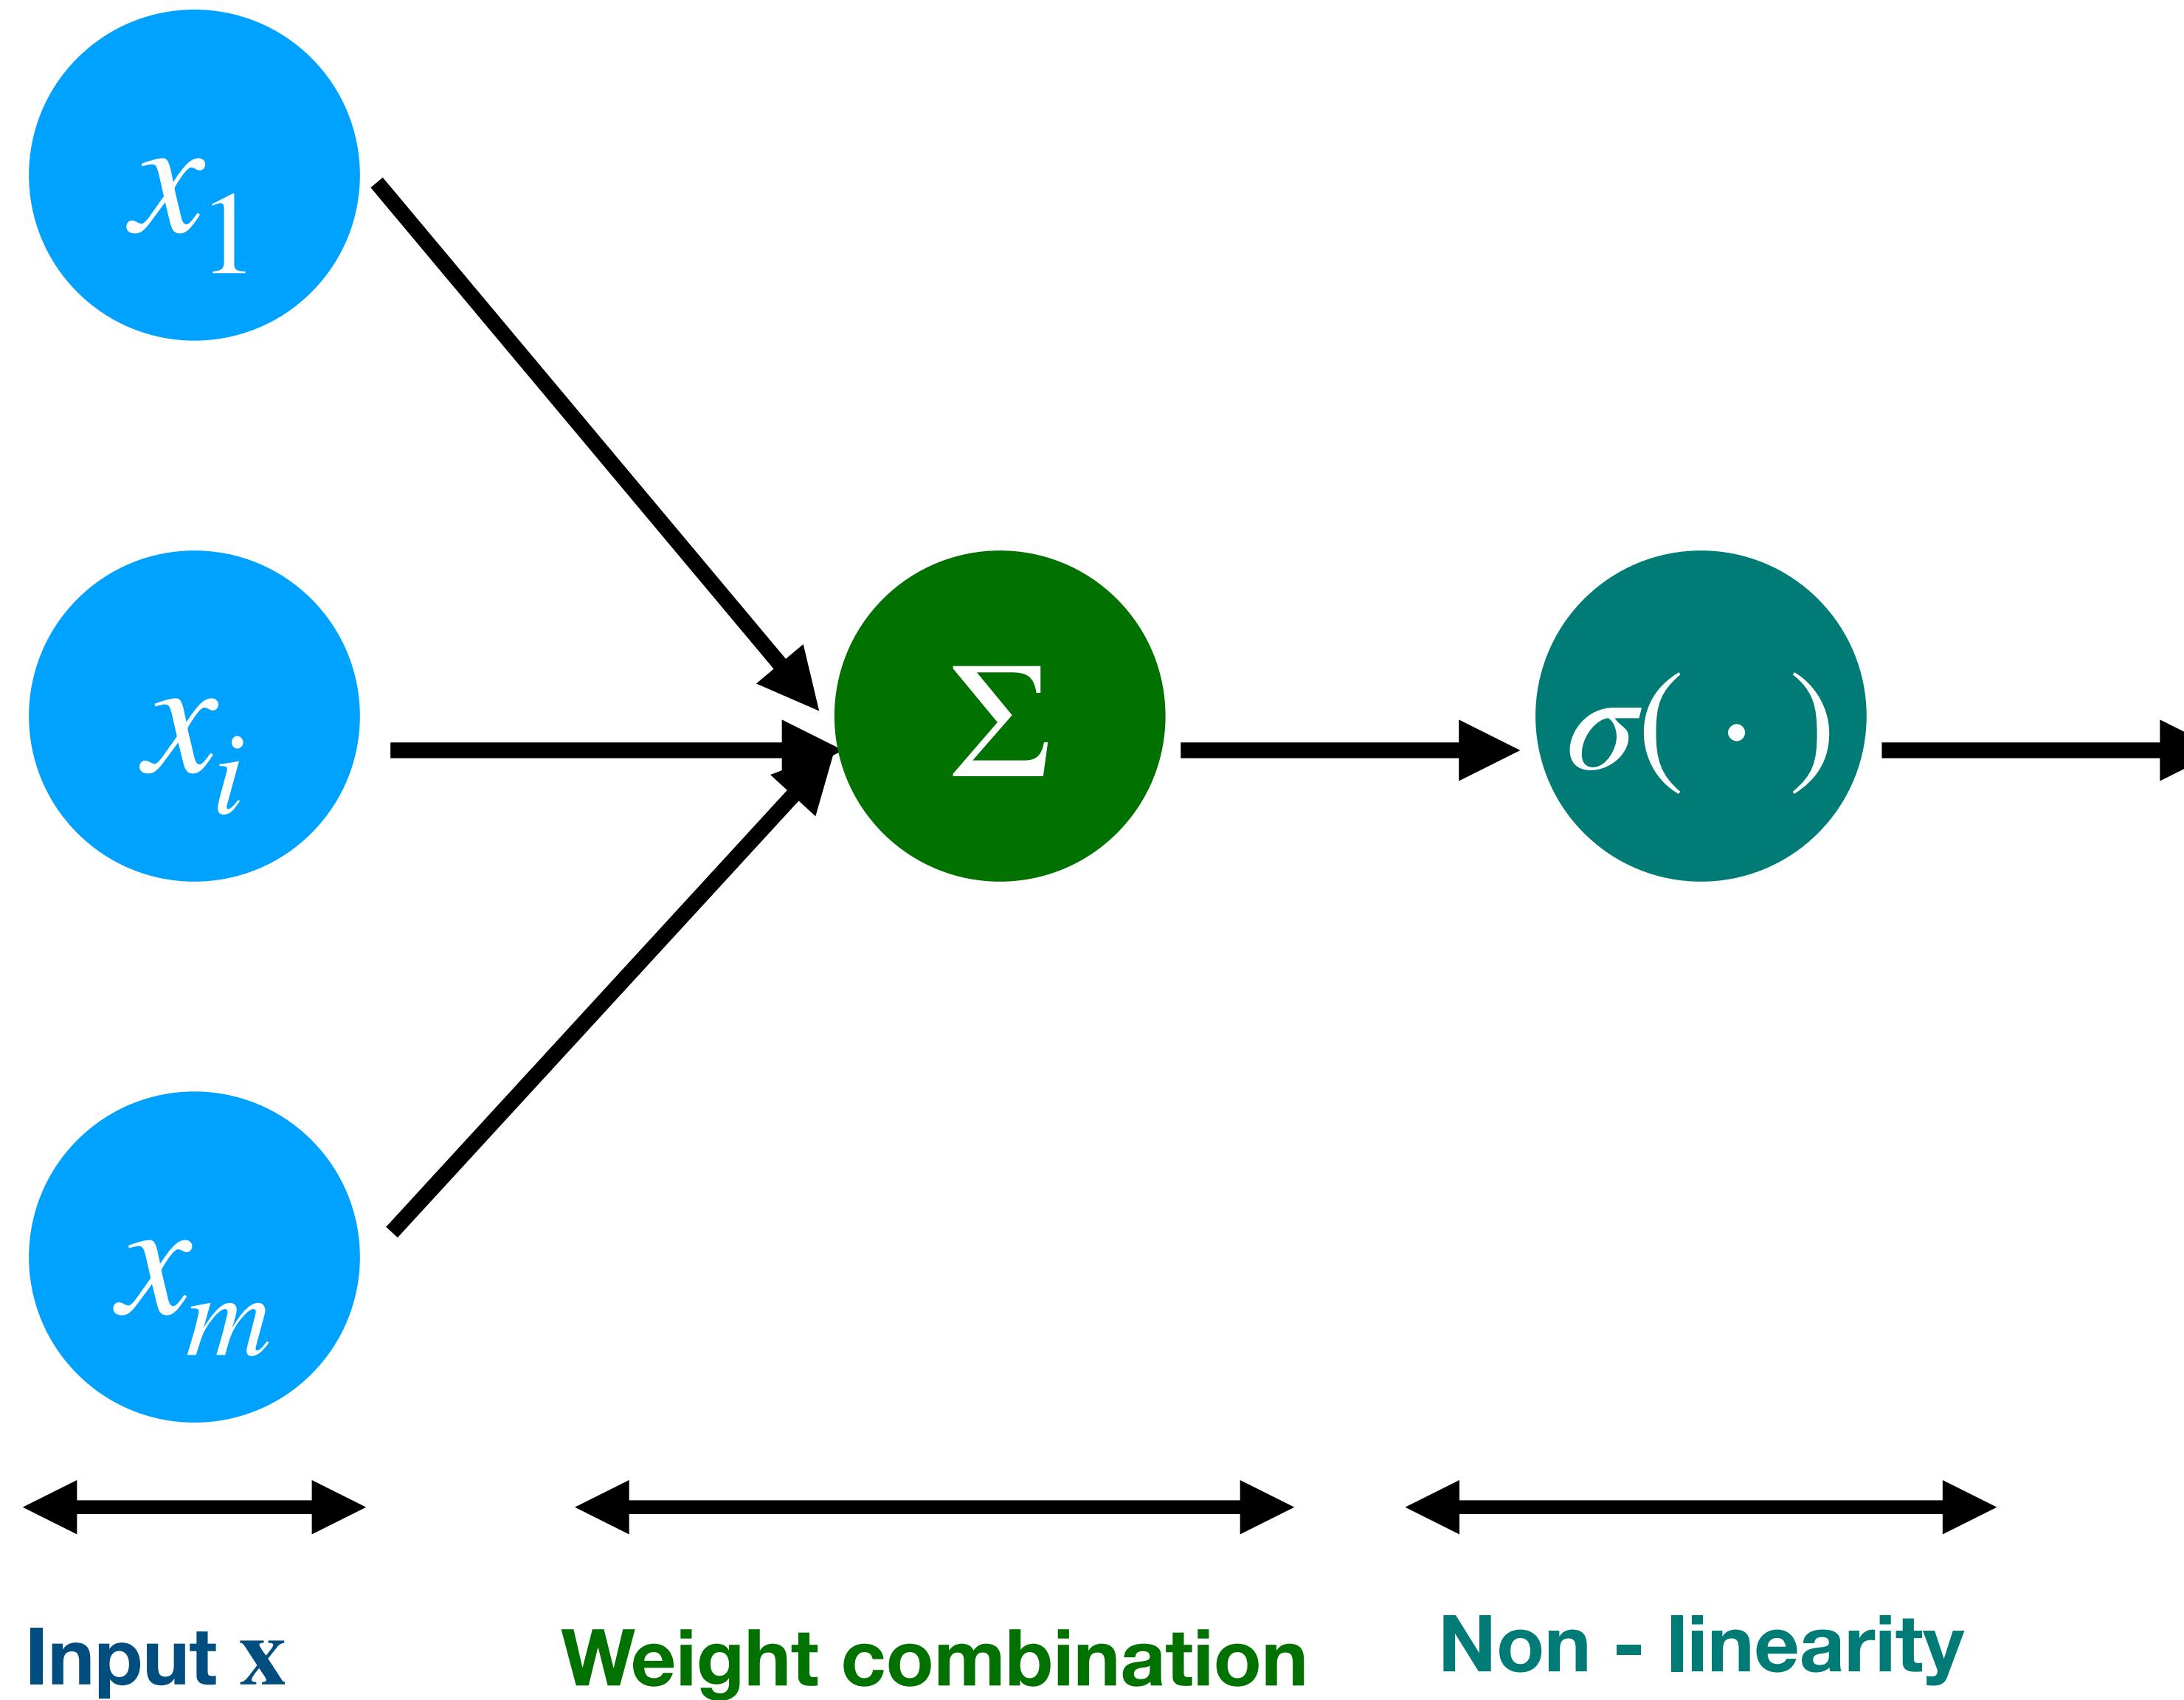


Input \mathbf{x}

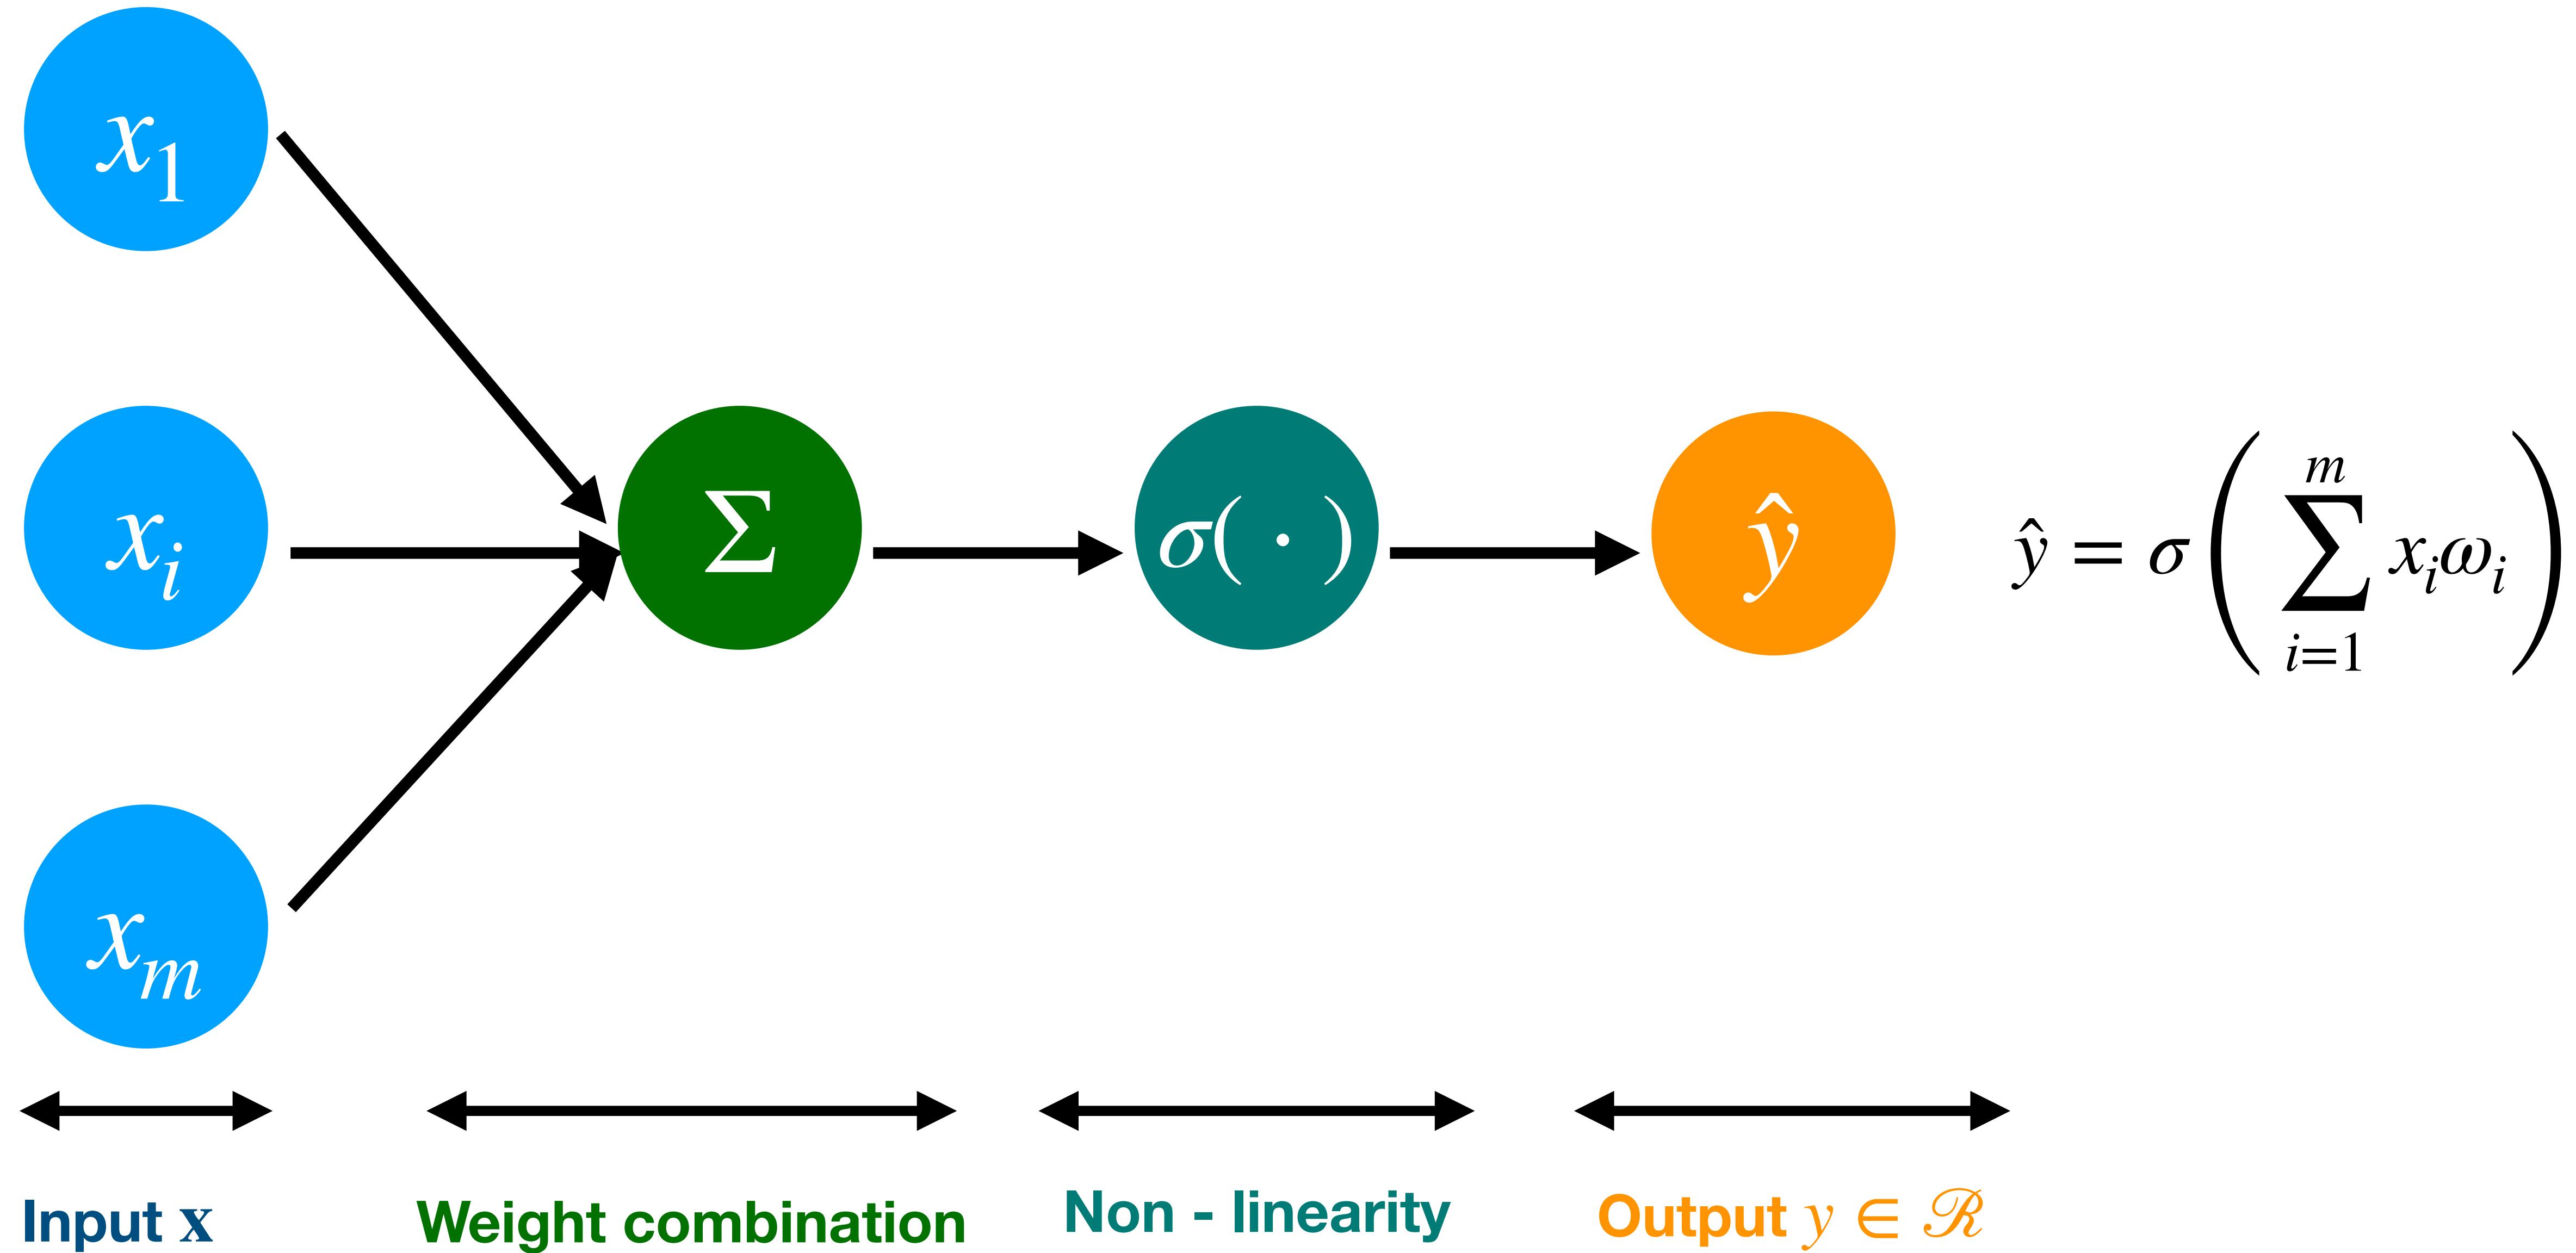
The Perceptron: the most basic neural network



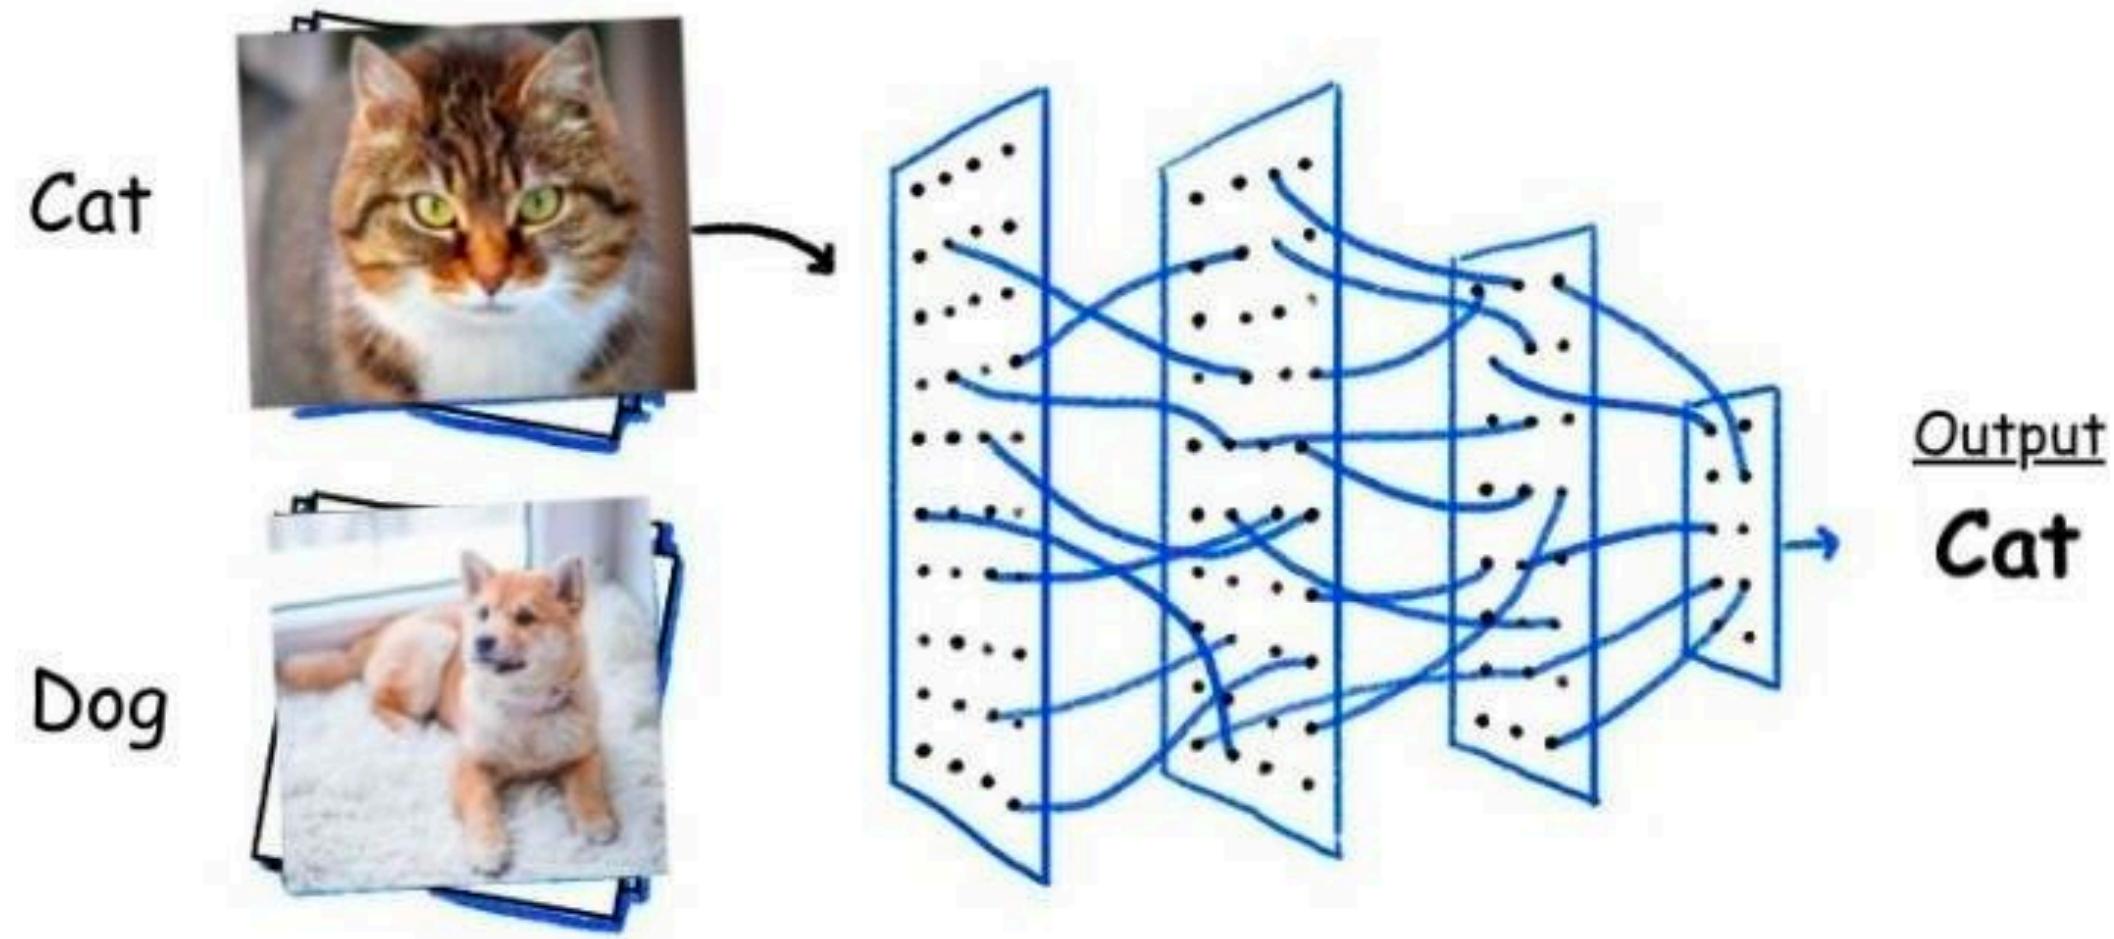
The Perceptron: the most basic neural network



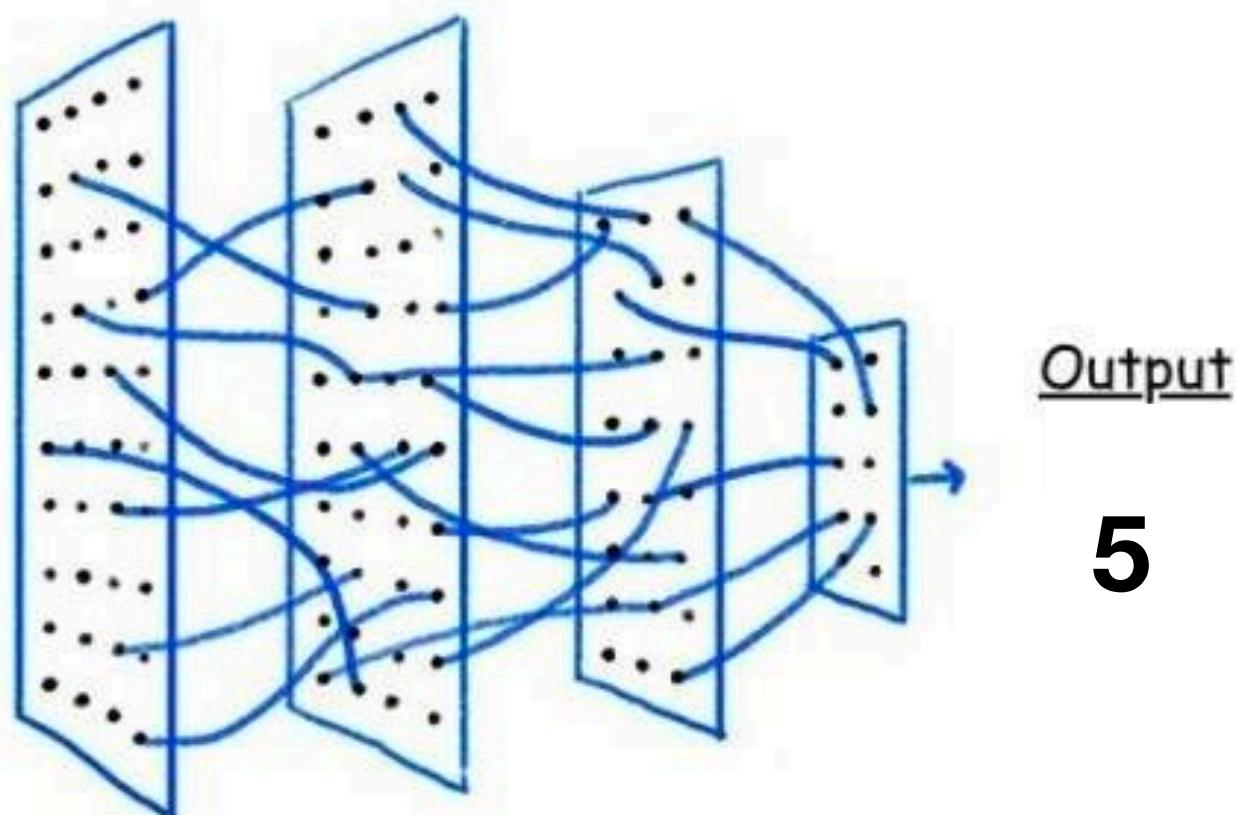
The Perceptron: the most basic neural network



Example of Inputs



I really love
my professor



x_1, \dots, x_m corresponds to the **pixels of the image**.

\hat{y} corresponds to the **any scalar > 0 if the Input is an cat**

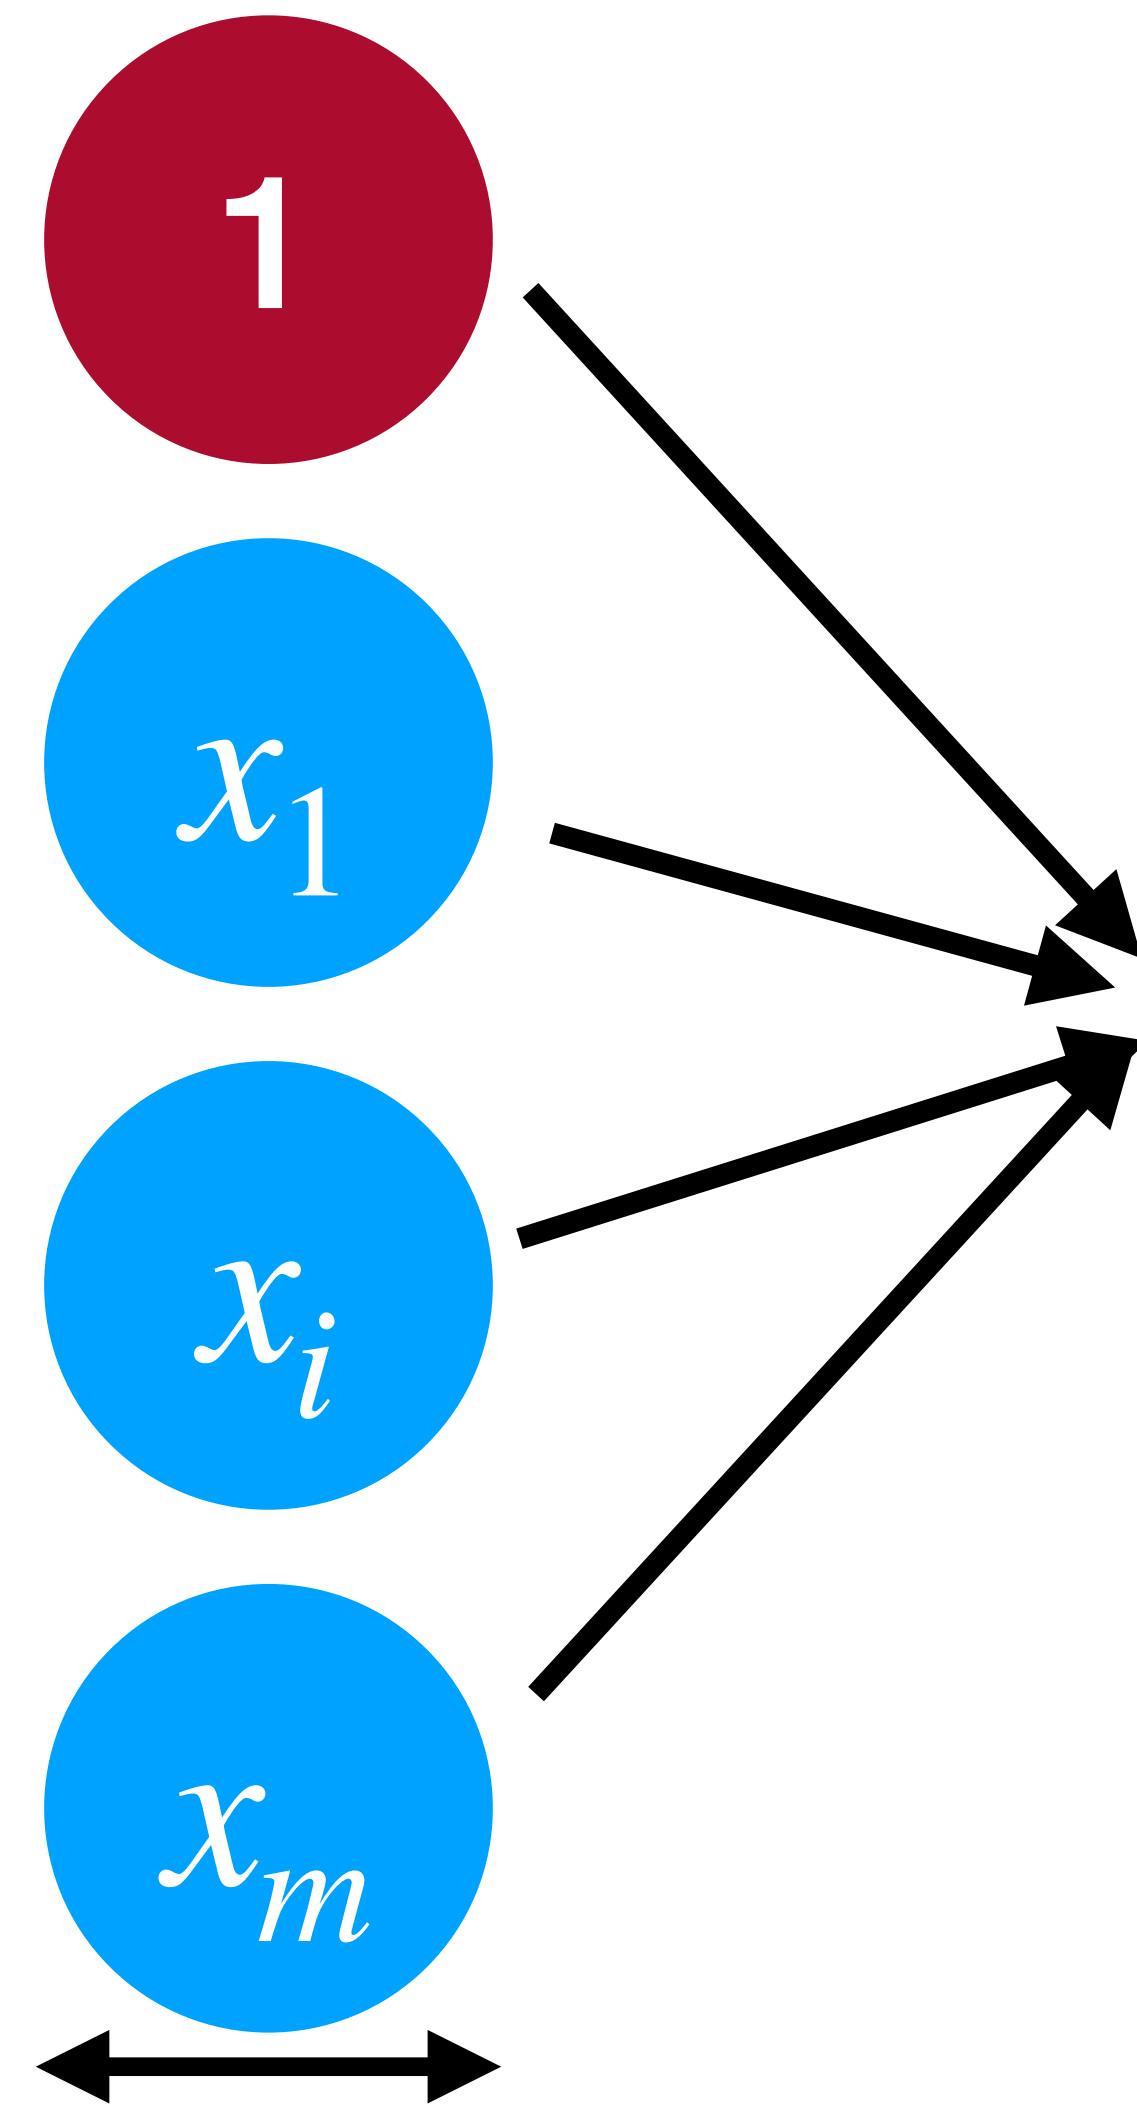
\hat{y} corresponds to the **any scalar < 0 if the Input is an dog**

x_1, \dots, x_m corresponds to the **index of the word in text**

\hat{y} corresponds to the **any scalar between [0,5] and model the sentiment of the text.**

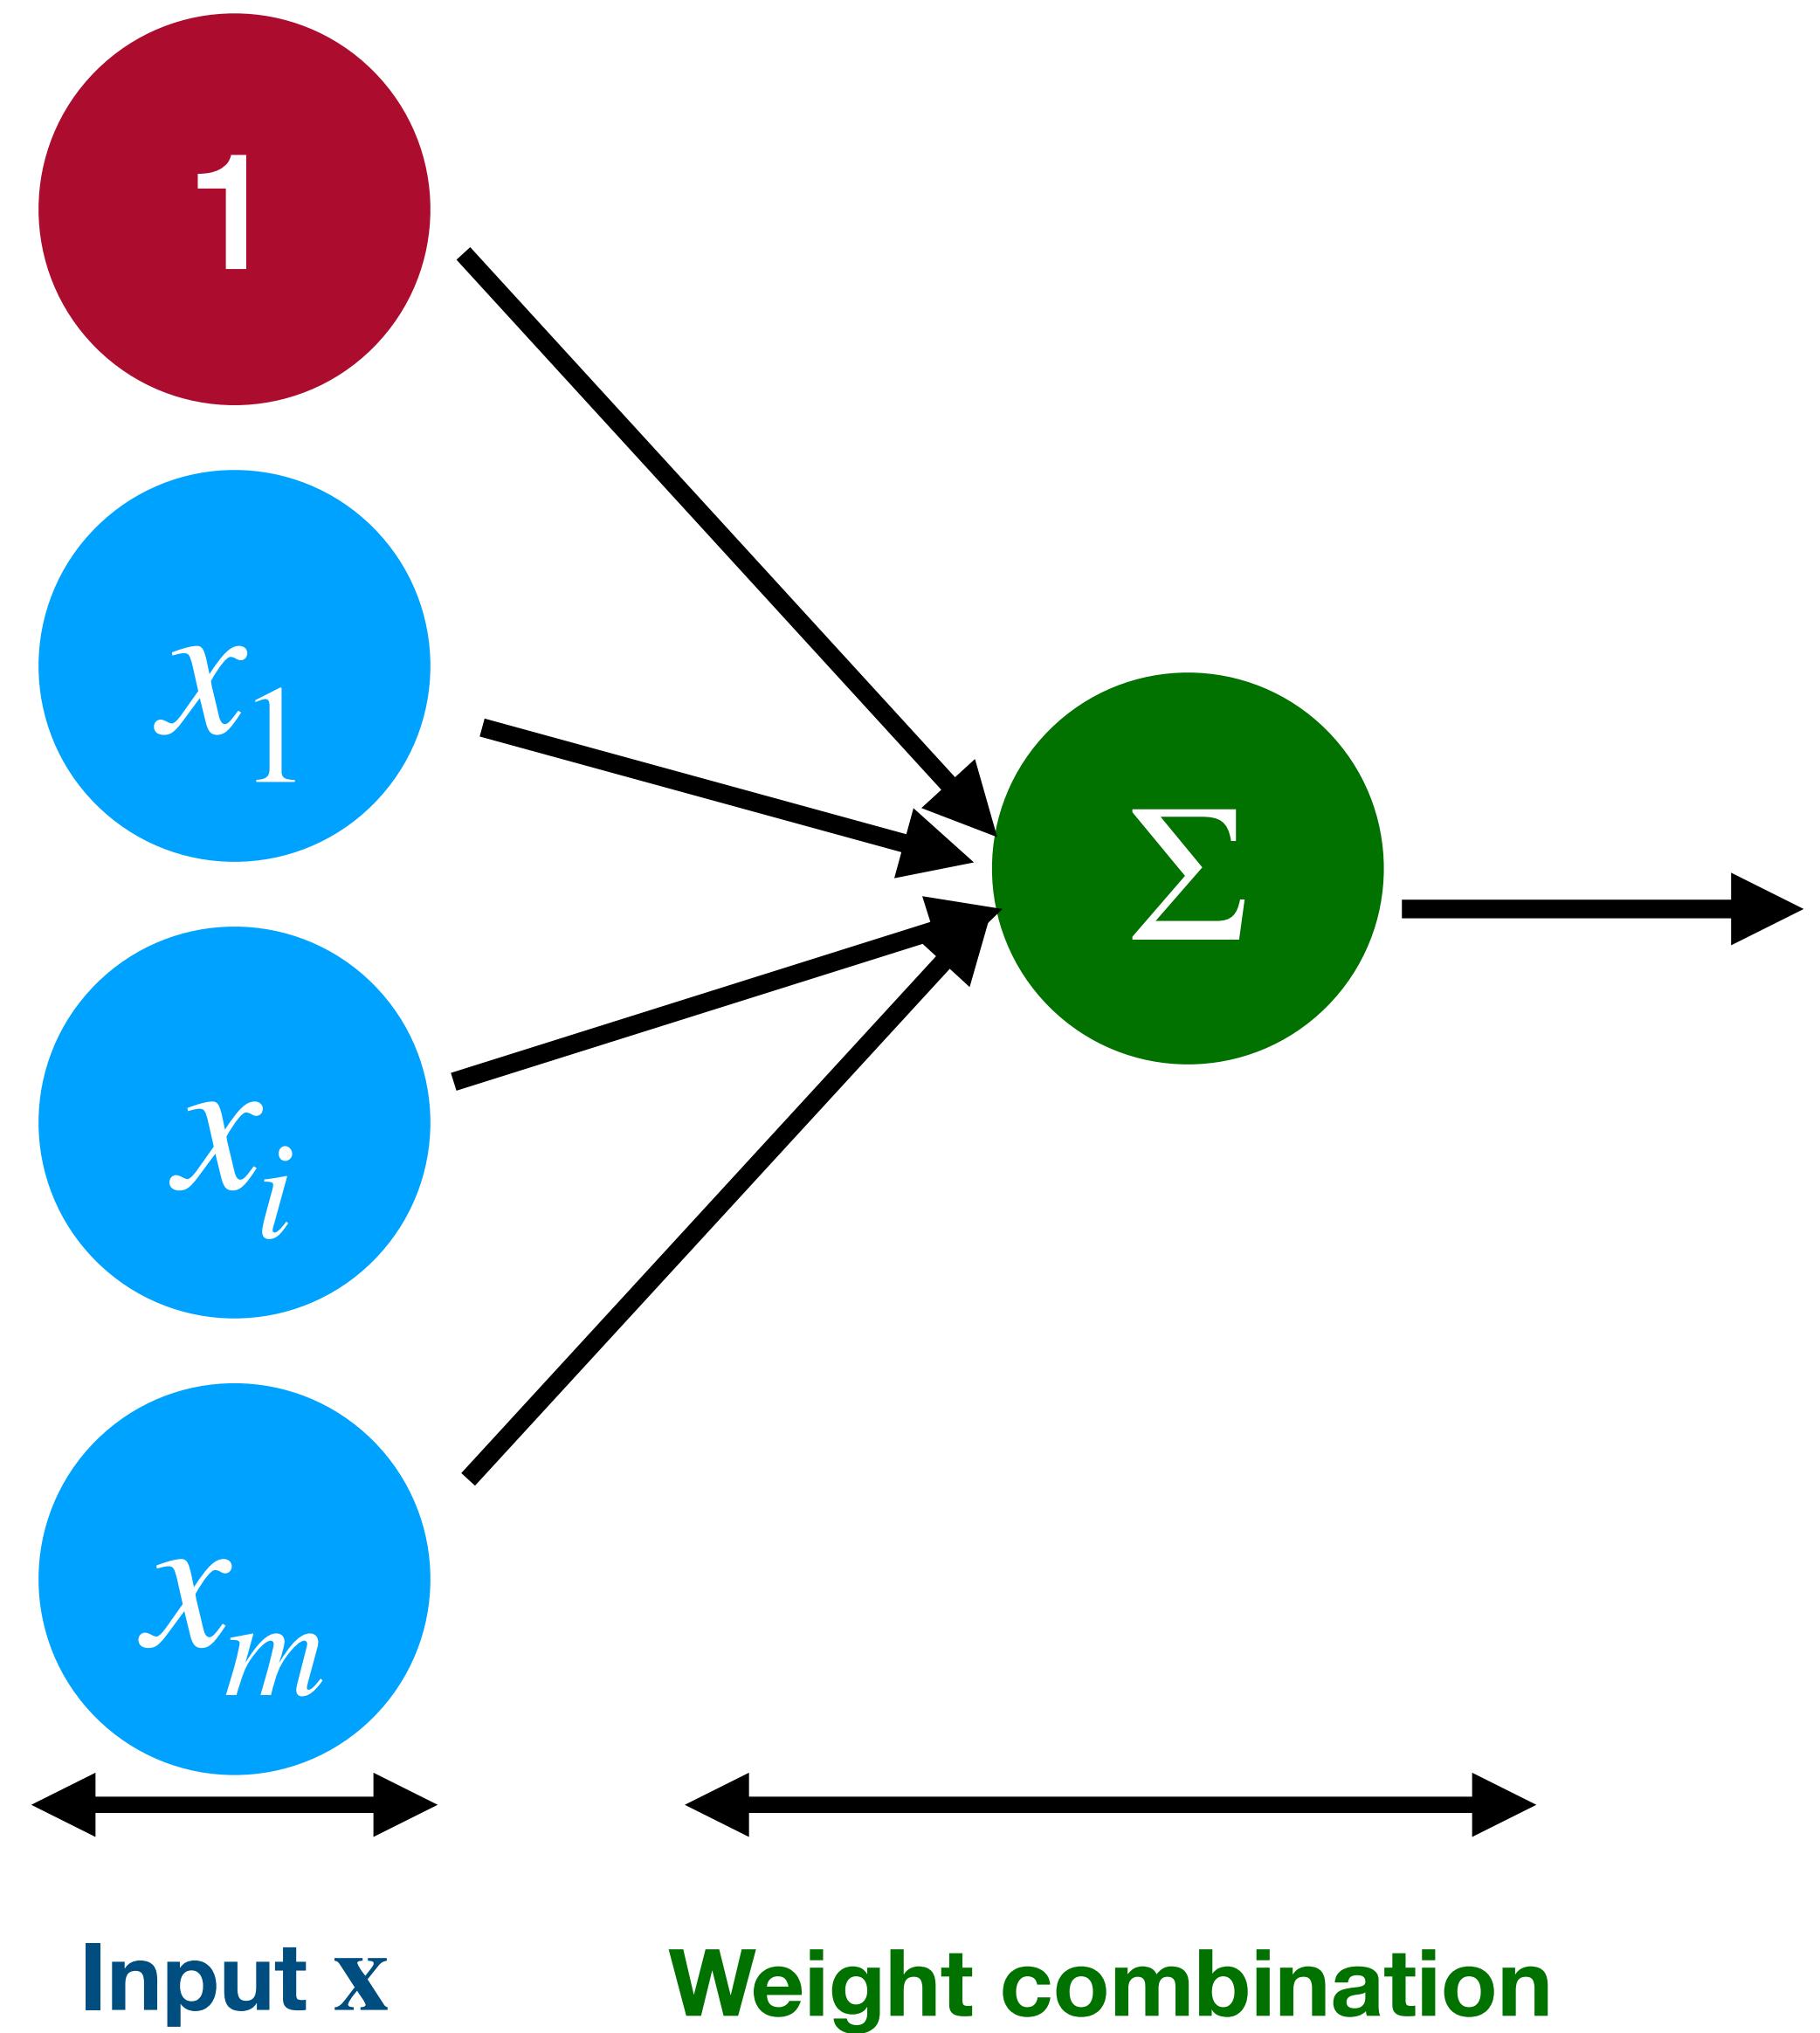
Do not forget the Biases for the forward propagation

Do not forget the Bias for the forward propagation

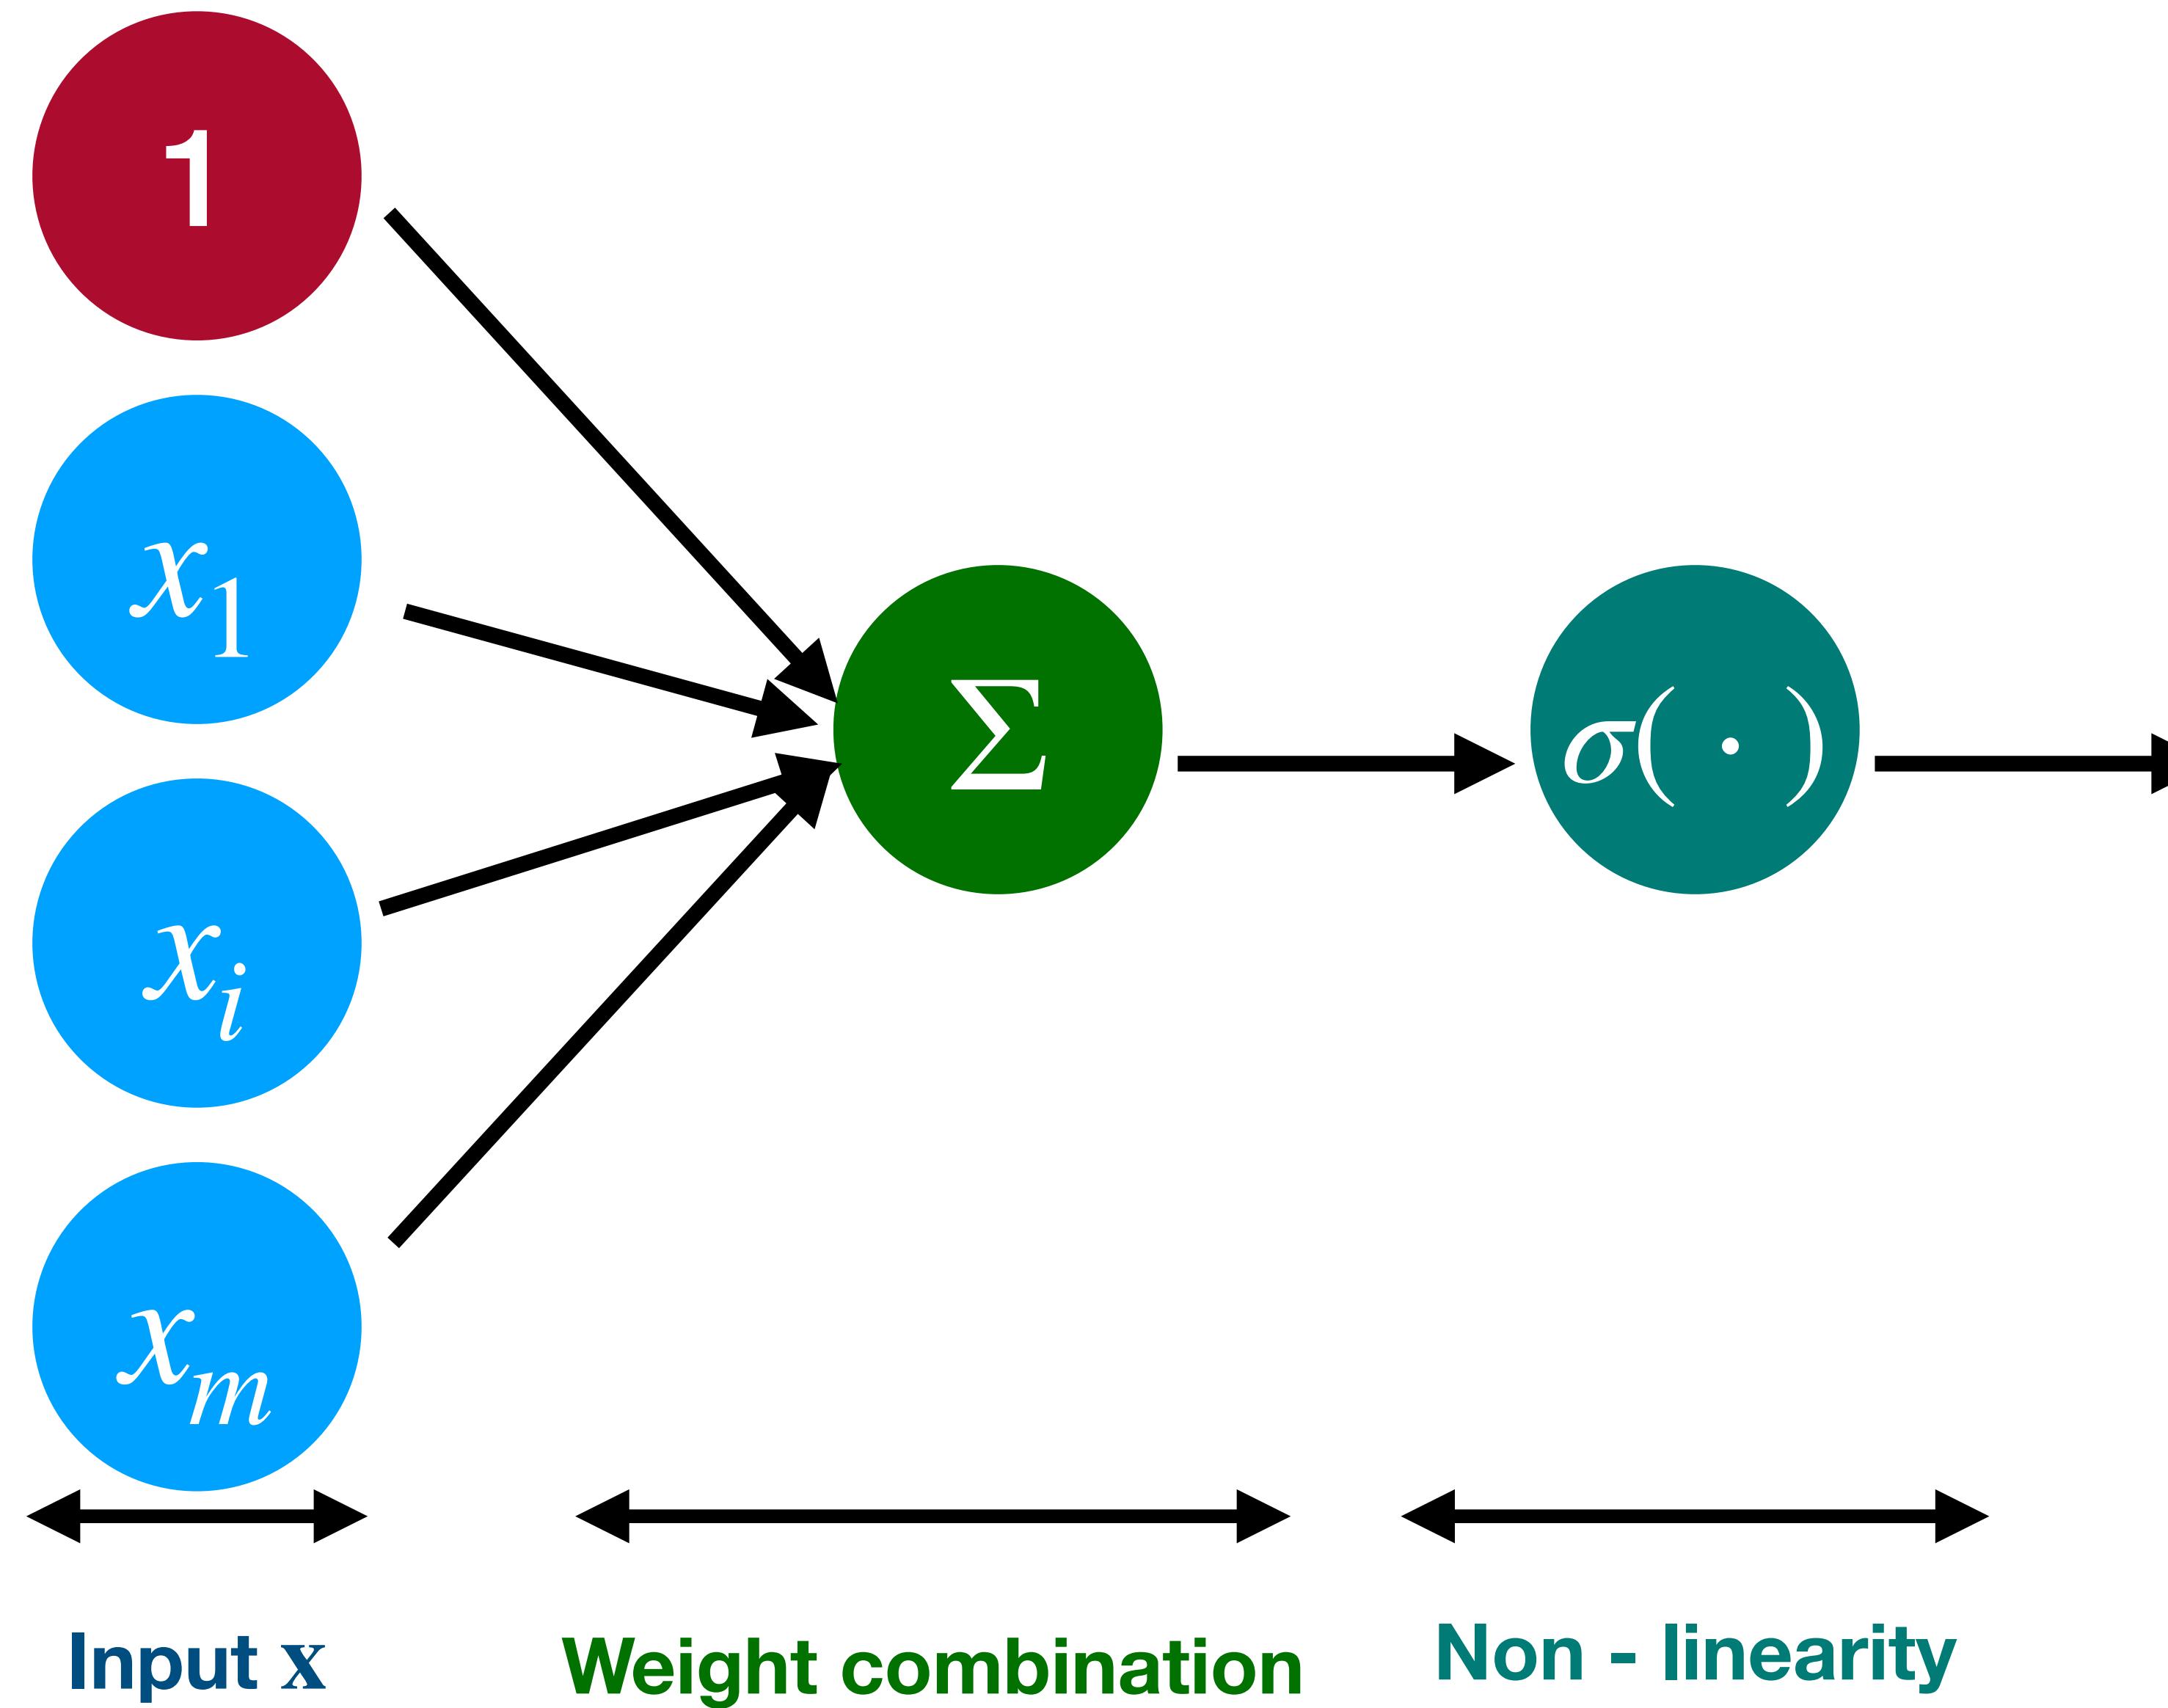


Input \mathbf{x}

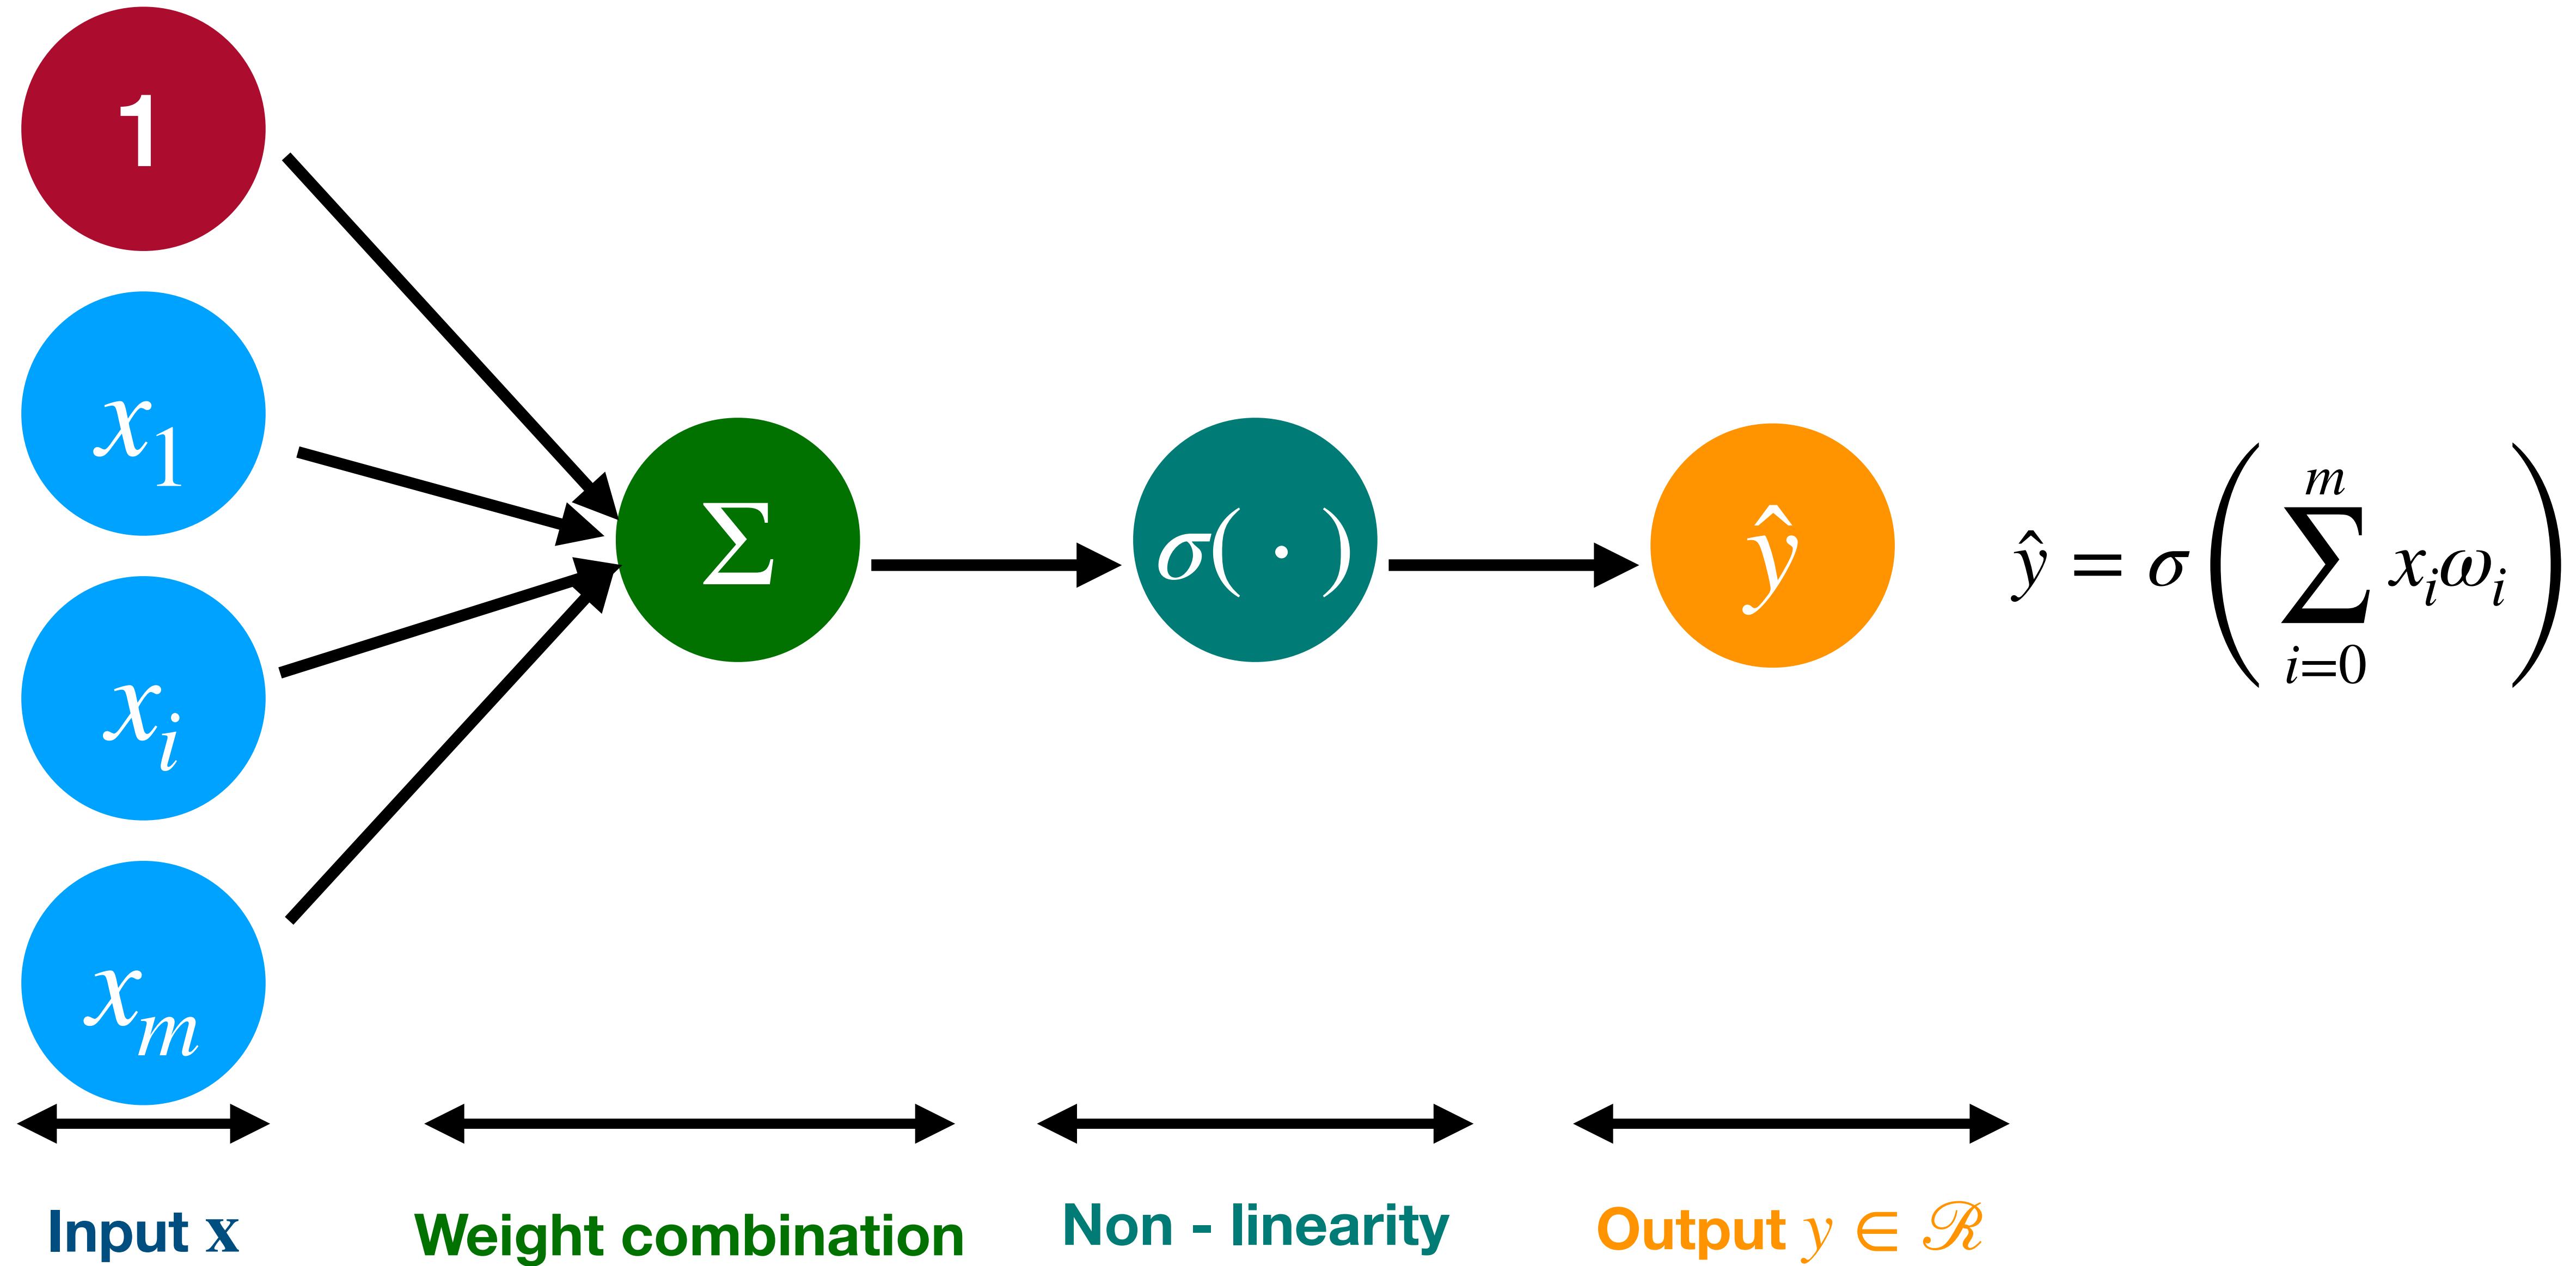
Do not forget the Bias for the forward propagation



Do not forget the Bias for the forward propagation



Do not forget the Bias for the forward propagation



Towards a Matrix Formulation

Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Matrix Formulation

$$\hat{y} = \sigma \left(\mathbf{X}^T \mathbf{W} \right)$$

Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Matrix Formulation

$$\hat{y} = \sigma \left(\mathbf{X}^T \mathbf{W} \right)$$

Where

$$\mathbf{X} = [x_0, \dots, x_m]$$

$$\mathbf{W} = [\omega_0, \dots, \omega_m]$$

Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Matrix Formulation

$$\hat{y} = \sigma \left(\mathbf{X}^T \mathbf{W} \right) \rightarrow \text{Simple Dot Product !}$$

Where

$$\mathbf{X} = [x_0, \dots, x_m]$$

$$\mathbf{W} = [\omega_0, \dots, \omega_m]$$

Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Matrix Formulation

$$\hat{y} = \sigma \left(\mathbf{X}^T \mathbf{W} \right) \rightarrow$$

Where

$$\mathbf{X} = [x_0, \dots, x_m]$$

Simple Dot Product !

Neural Network Weights

$$\mathbf{W} = [\omega_0, \dots, \omega_m]$$



Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Matrix Formulation

$$\hat{y} = \sigma \left(\mathbf{X}^T \mathbf{W} \right) \rightarrow$$

Where

$$\mathbf{X} = [x_0, \dots, x_m]$$

Simple Dot Product !

Neural Network Weights

$$\mathbf{W} = [\omega_0, \dots, \omega_m]$$



For an image of size H,W we have $\mathbf{X} \in \mathcal{R}^{H \times W}$ but $\hat{y} \in \mathcal{R}$

Towards a Matrix Formulation

Before

$$\hat{y} = \sigma \left(\sum_{i=0}^m x_i \omega_i \right)$$

Matrix Formulation

$$\hat{y} = \sigma \left(\mathbf{X}^T \mathbf{W} \right) \rightarrow$$

Where

$$\mathbf{X} = [x_0, \dots, x_m]$$

Simple Dot Product !

Neural Network Weights

$$\mathbf{W} = [\omega_0, \dots, \omega_m]$$

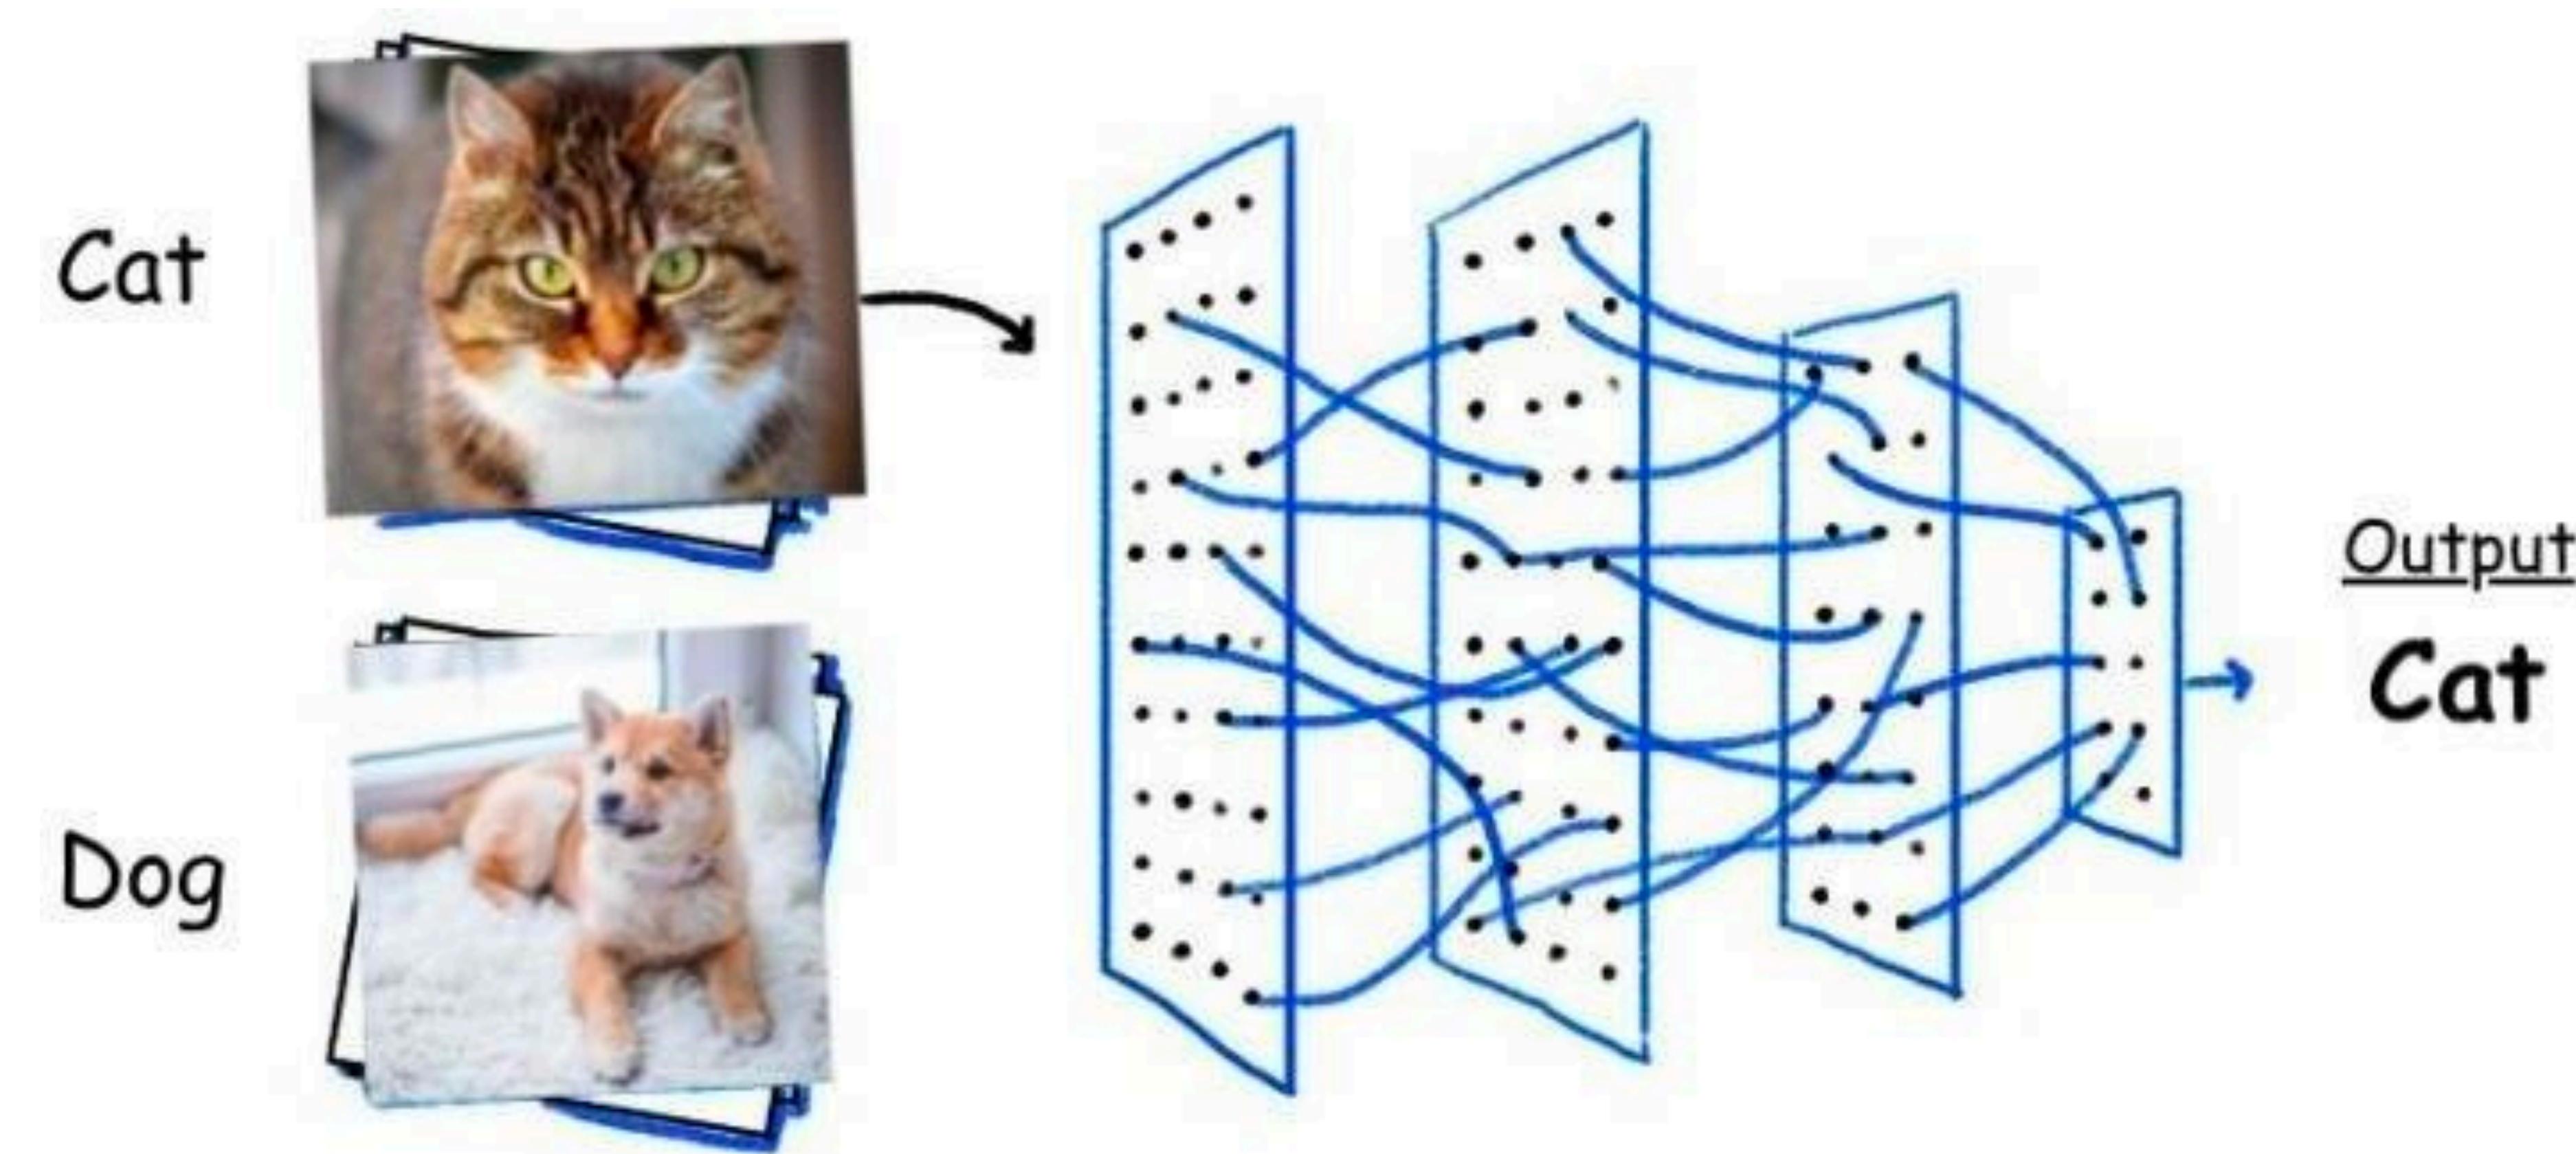


For an image of size H, W we have $\mathbf{X} \in \mathcal{R}^{H \times W}$ but $\hat{y} \in \mathcal{R}$

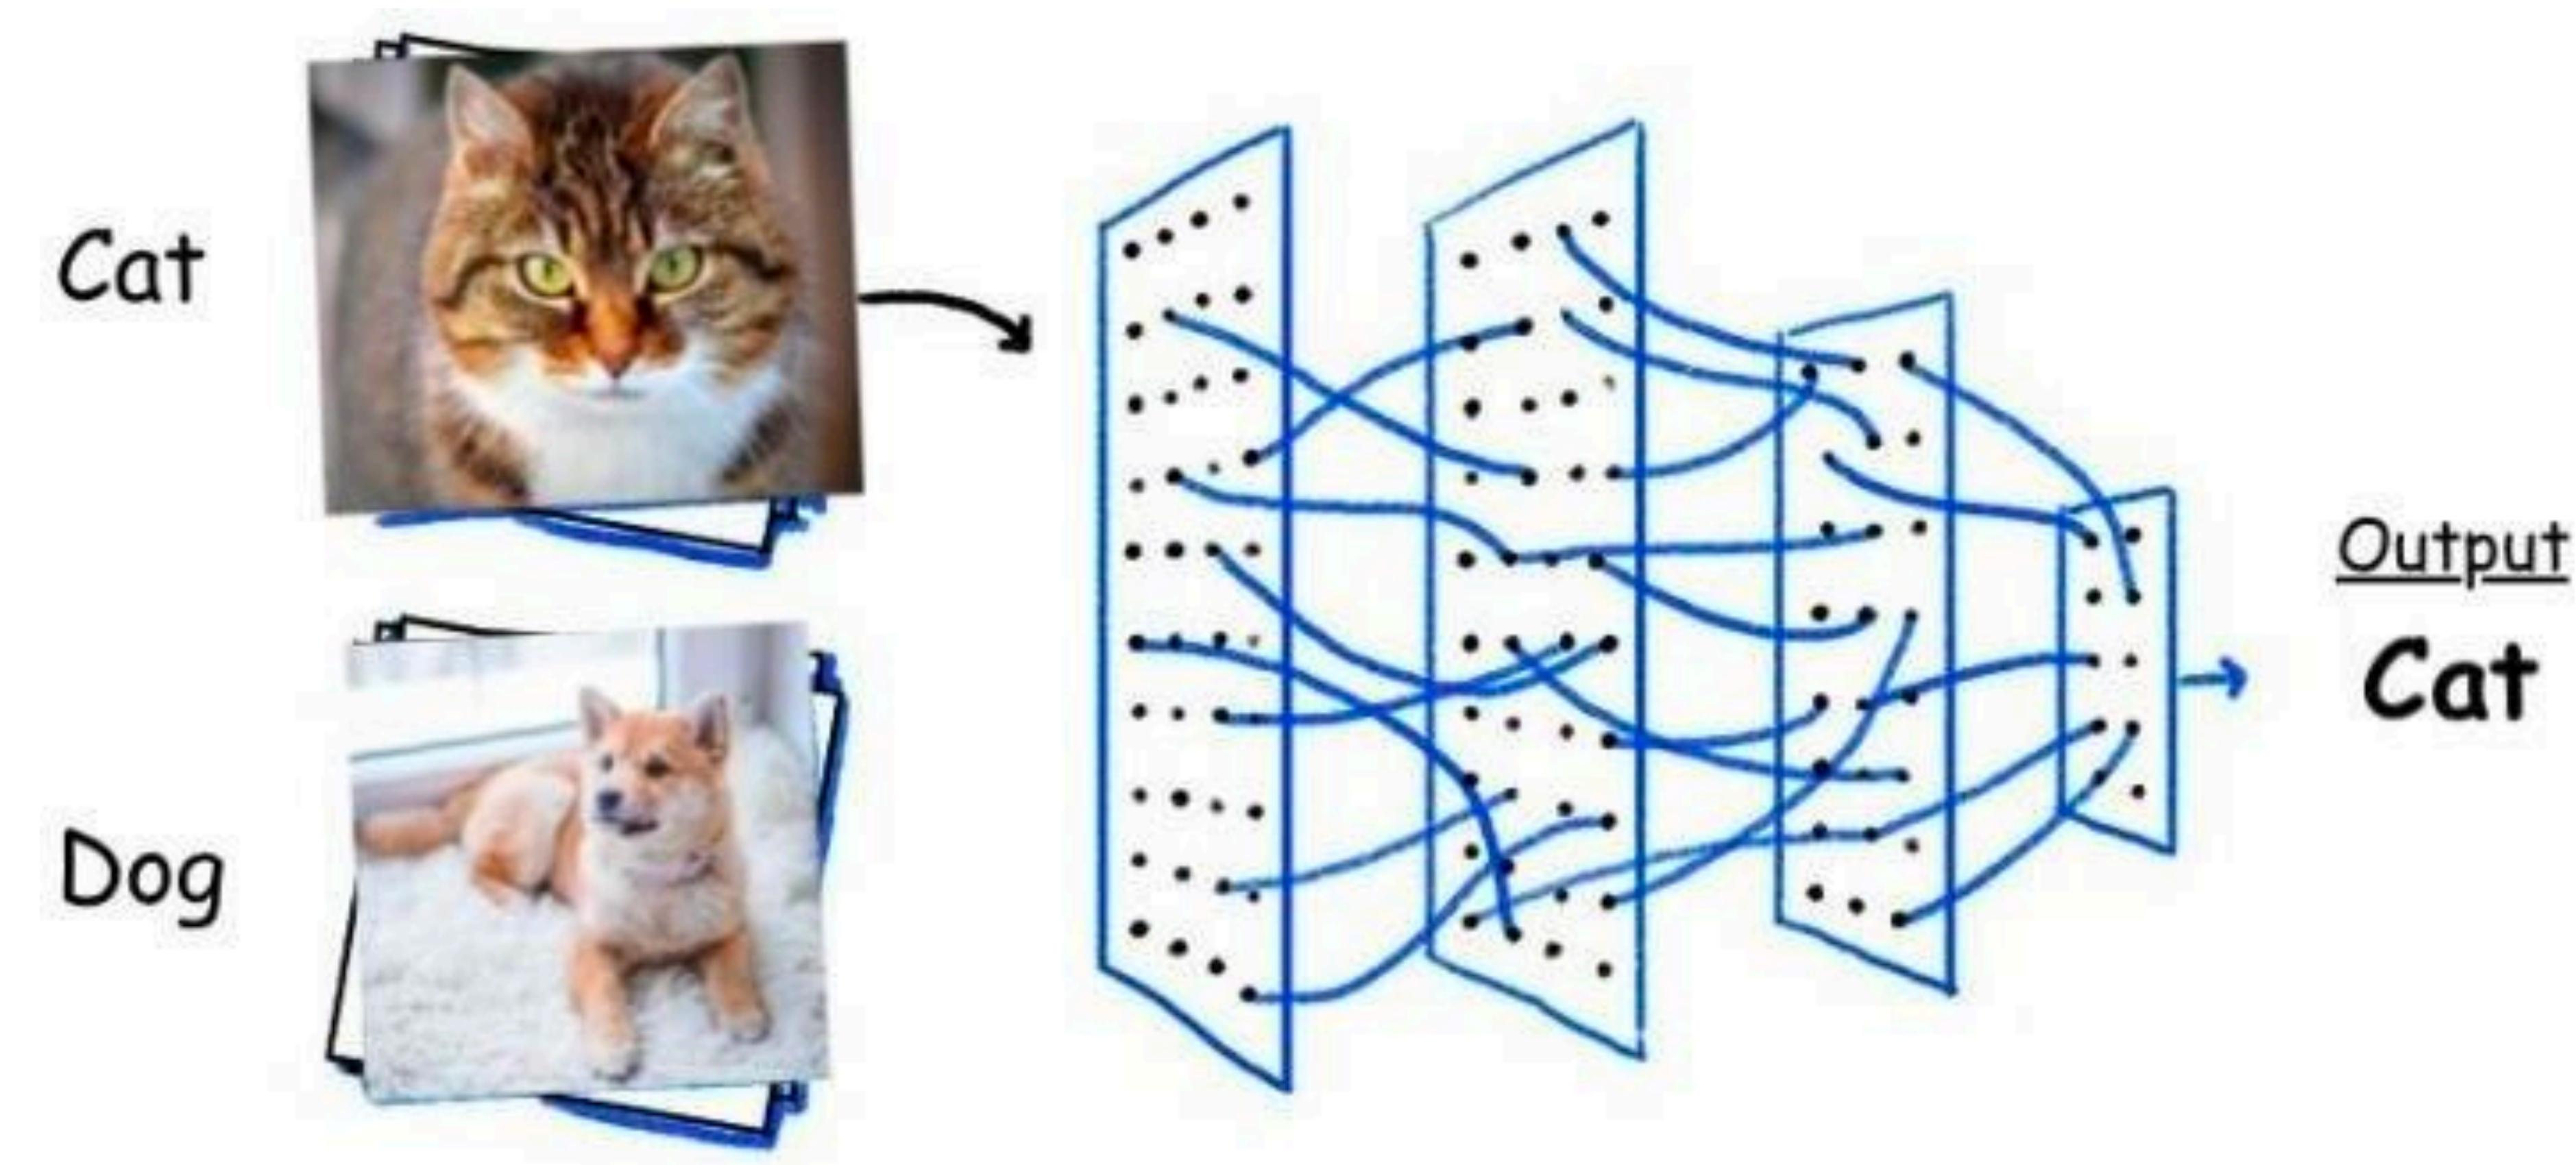
Be aware !!!!! $\mathbf{X} \in \mathcal{R}^d$ but $\hat{y} \in \mathcal{R}$

Let's go back to cat and dog classification

Let's go back to cat and dog classification



Let's go back to cat and dog classification



For an image of size H, W we have $X \in \mathcal{R}^{H \times W}$ but $\hat{y} \in [0, 1]$

Let's sum up so far

Let's sum up so far

1. The basic bloc to build a Neural Network

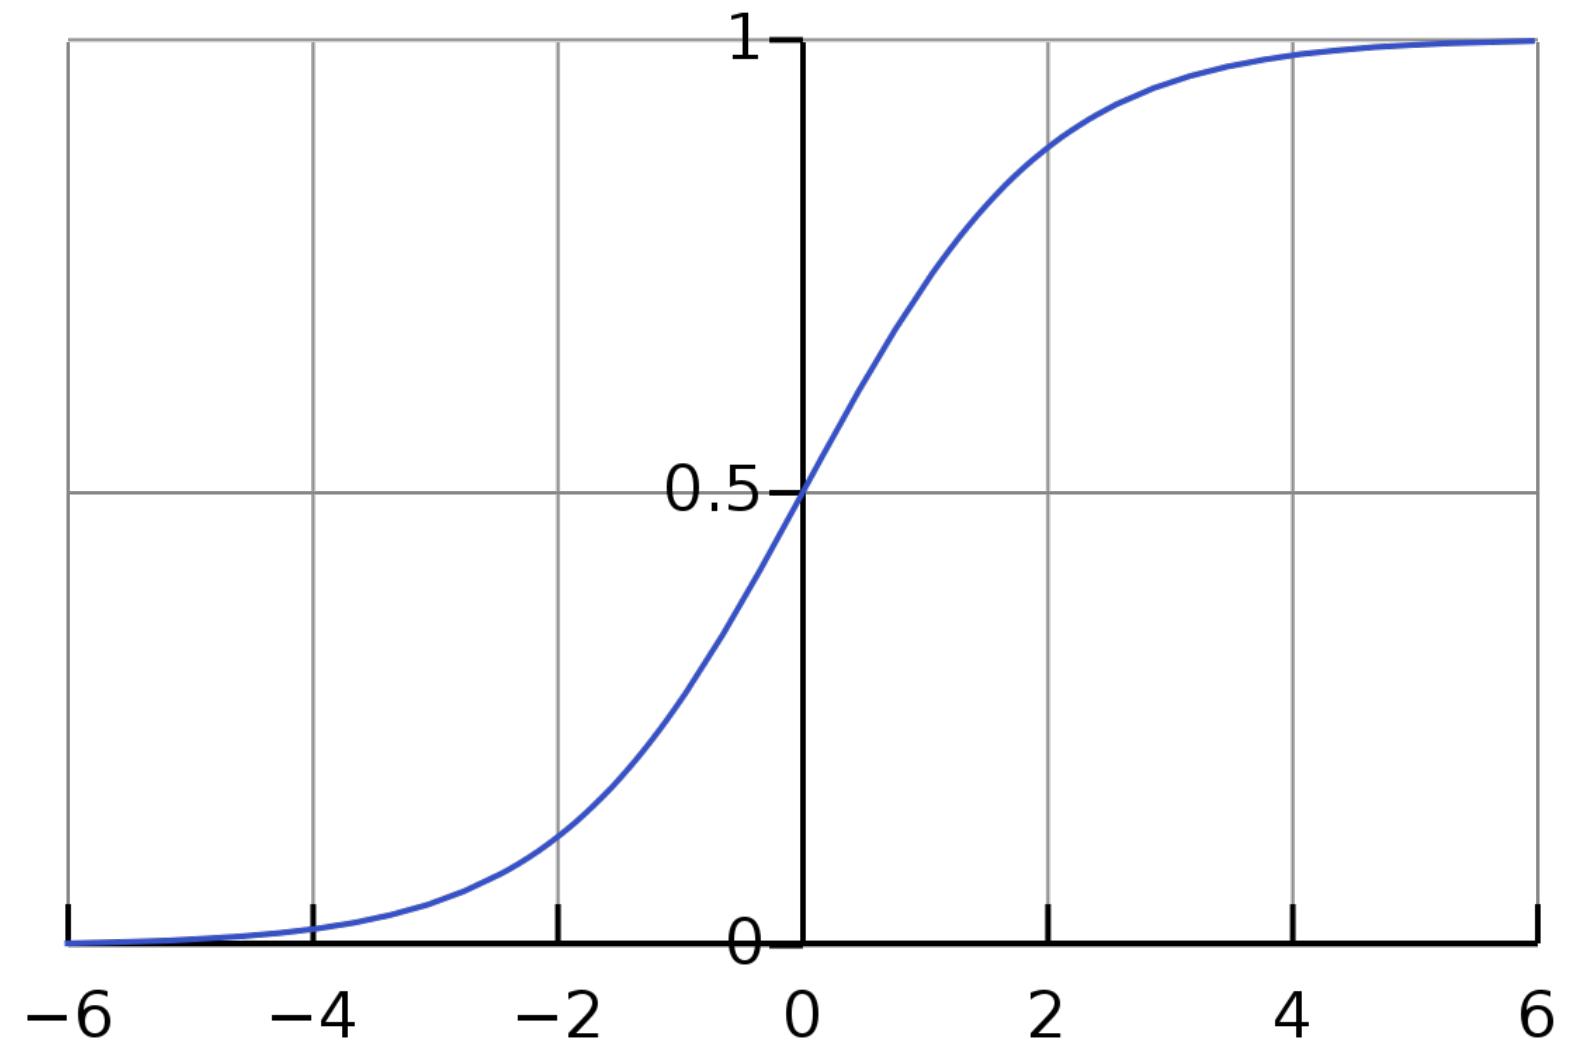
We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

What are the most common activation functions?

What are the most common activation functions?

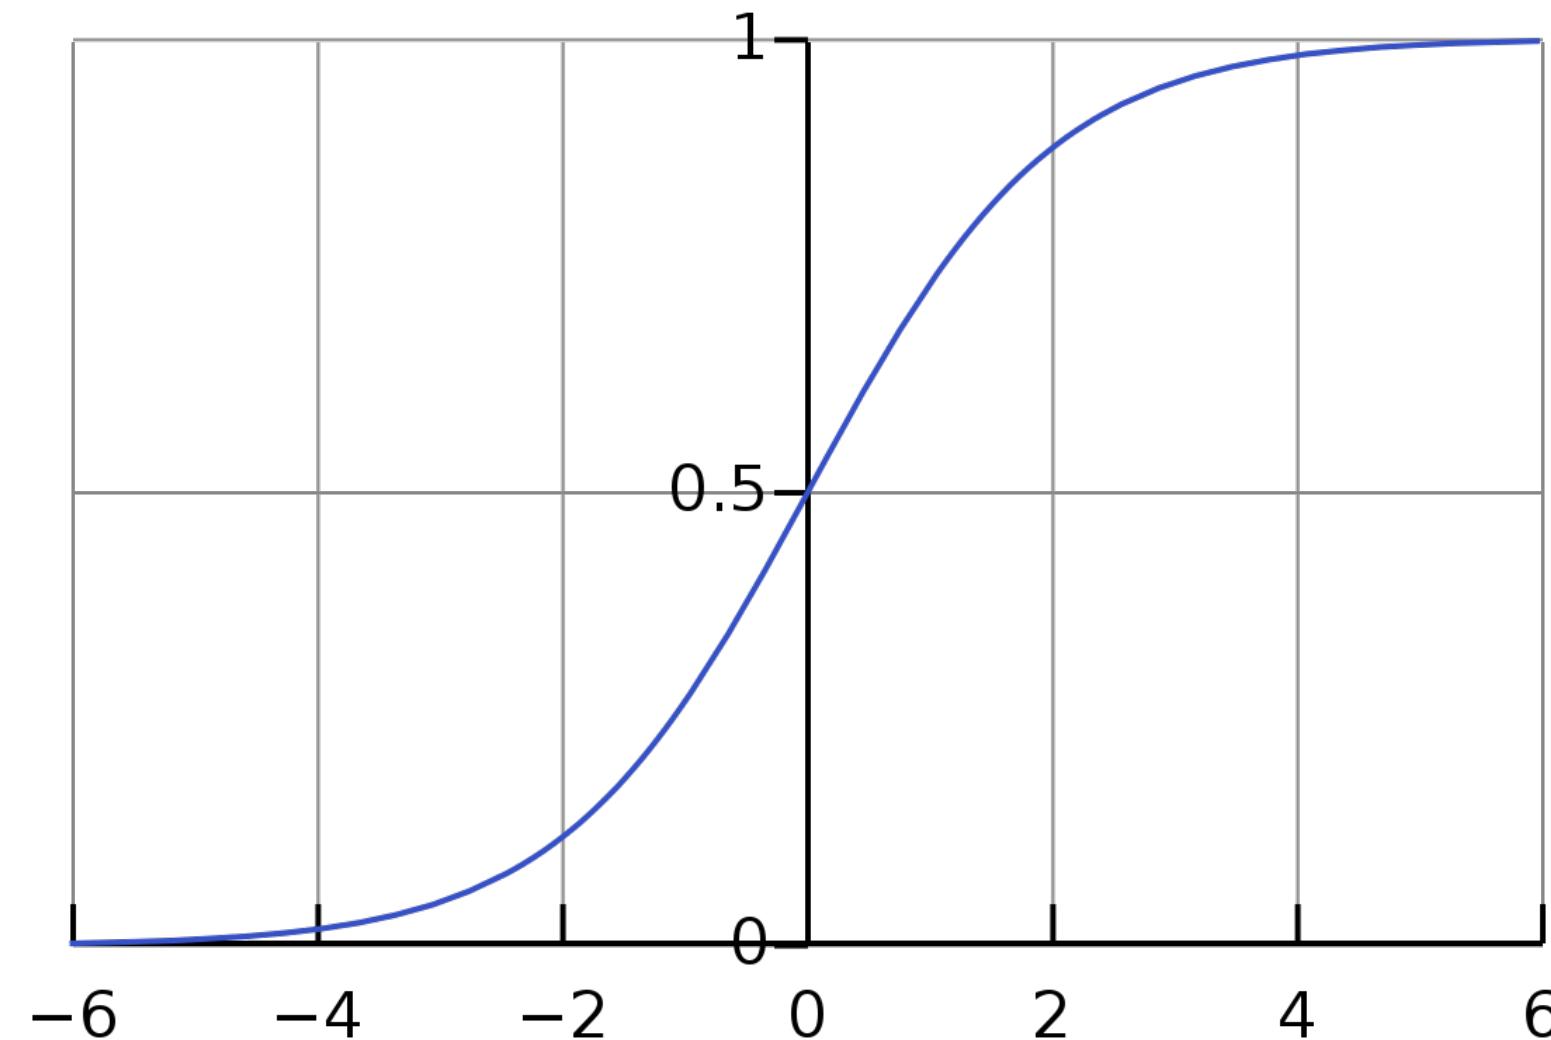
Sigmoid



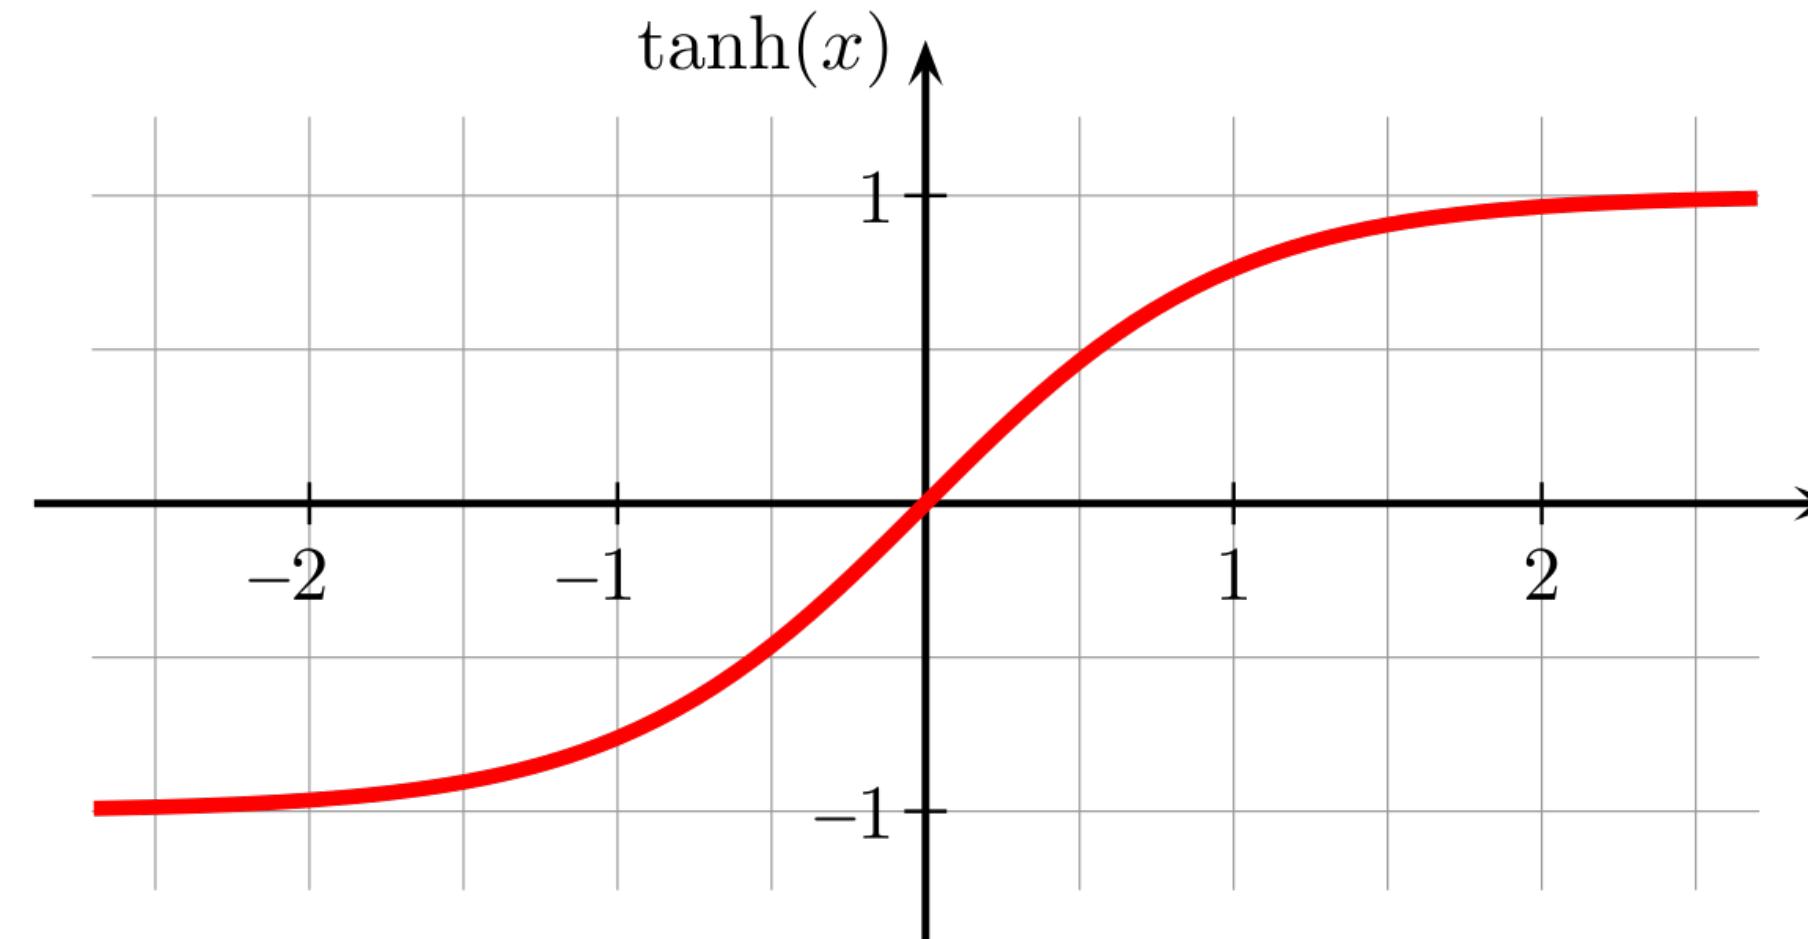
$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

What are the most common activation functions?

Sigmoid



Tanh

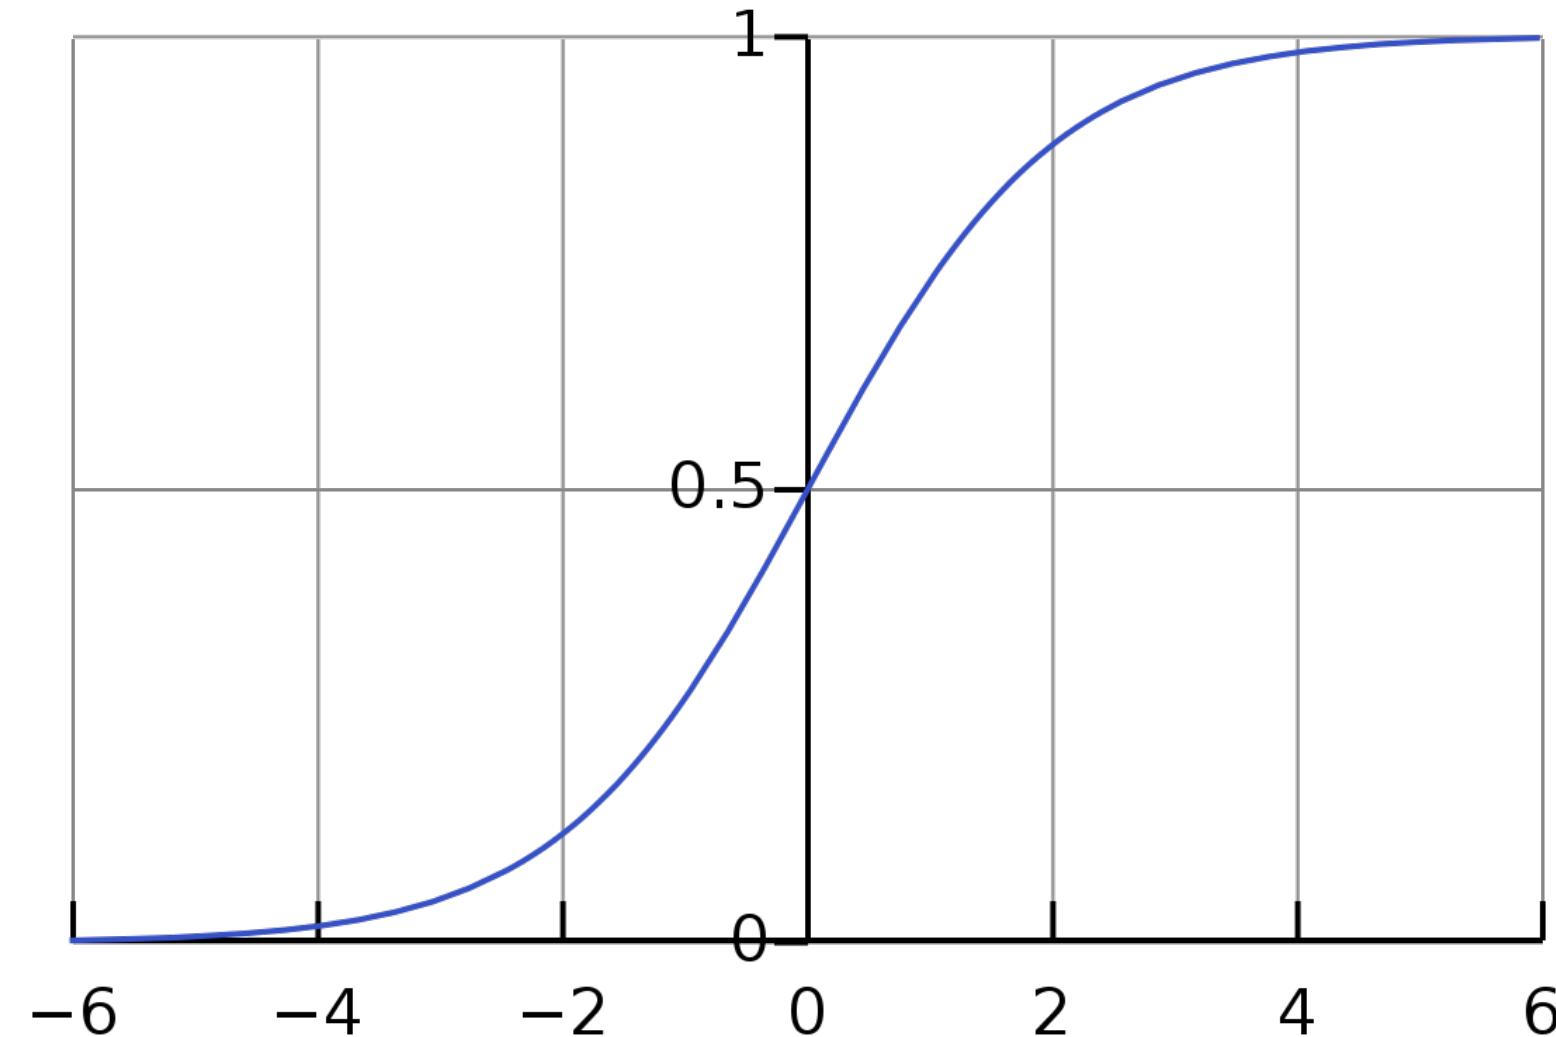


$$\sigma(x) = \frac{1}{1 + e^x}$$

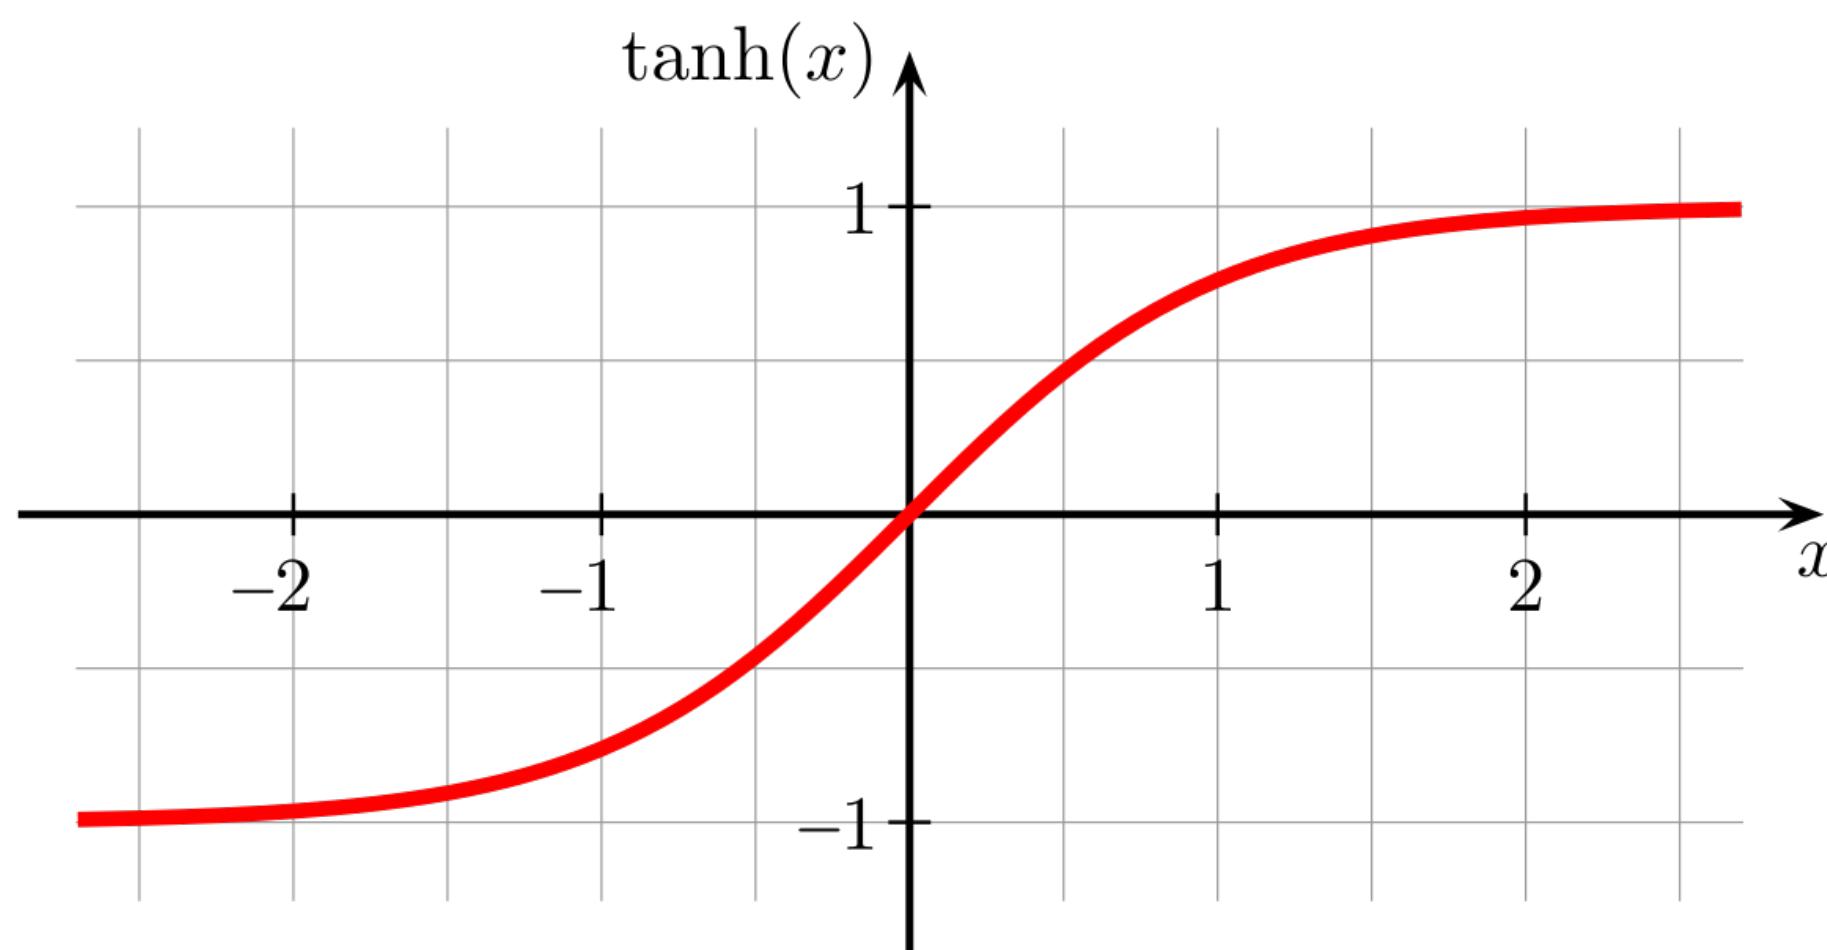
$$\sigma(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

What are the most common activation functions?

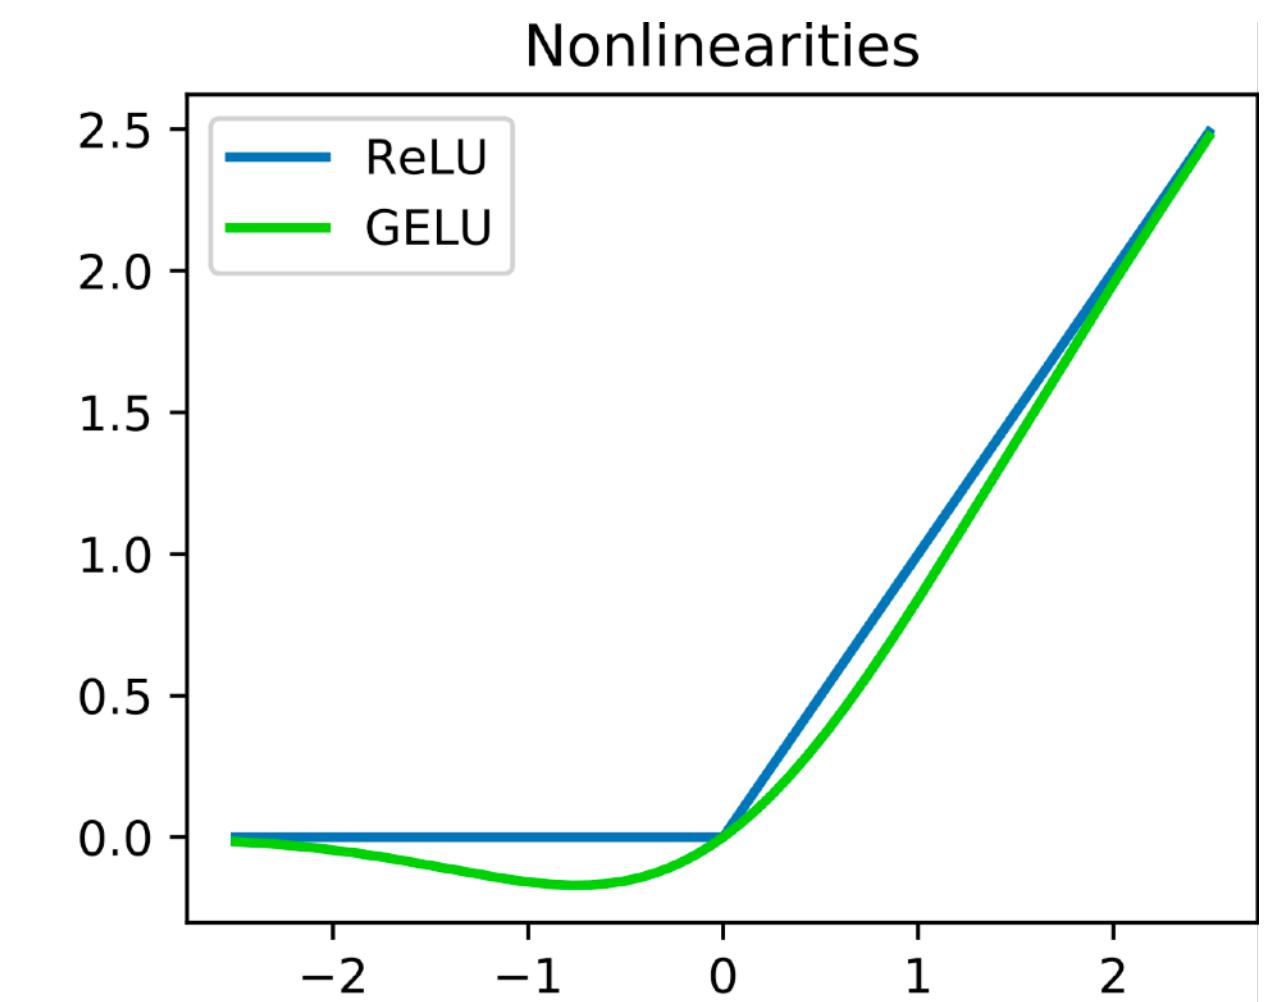
Sigmoid



Tanh



ReLU/GeLU



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

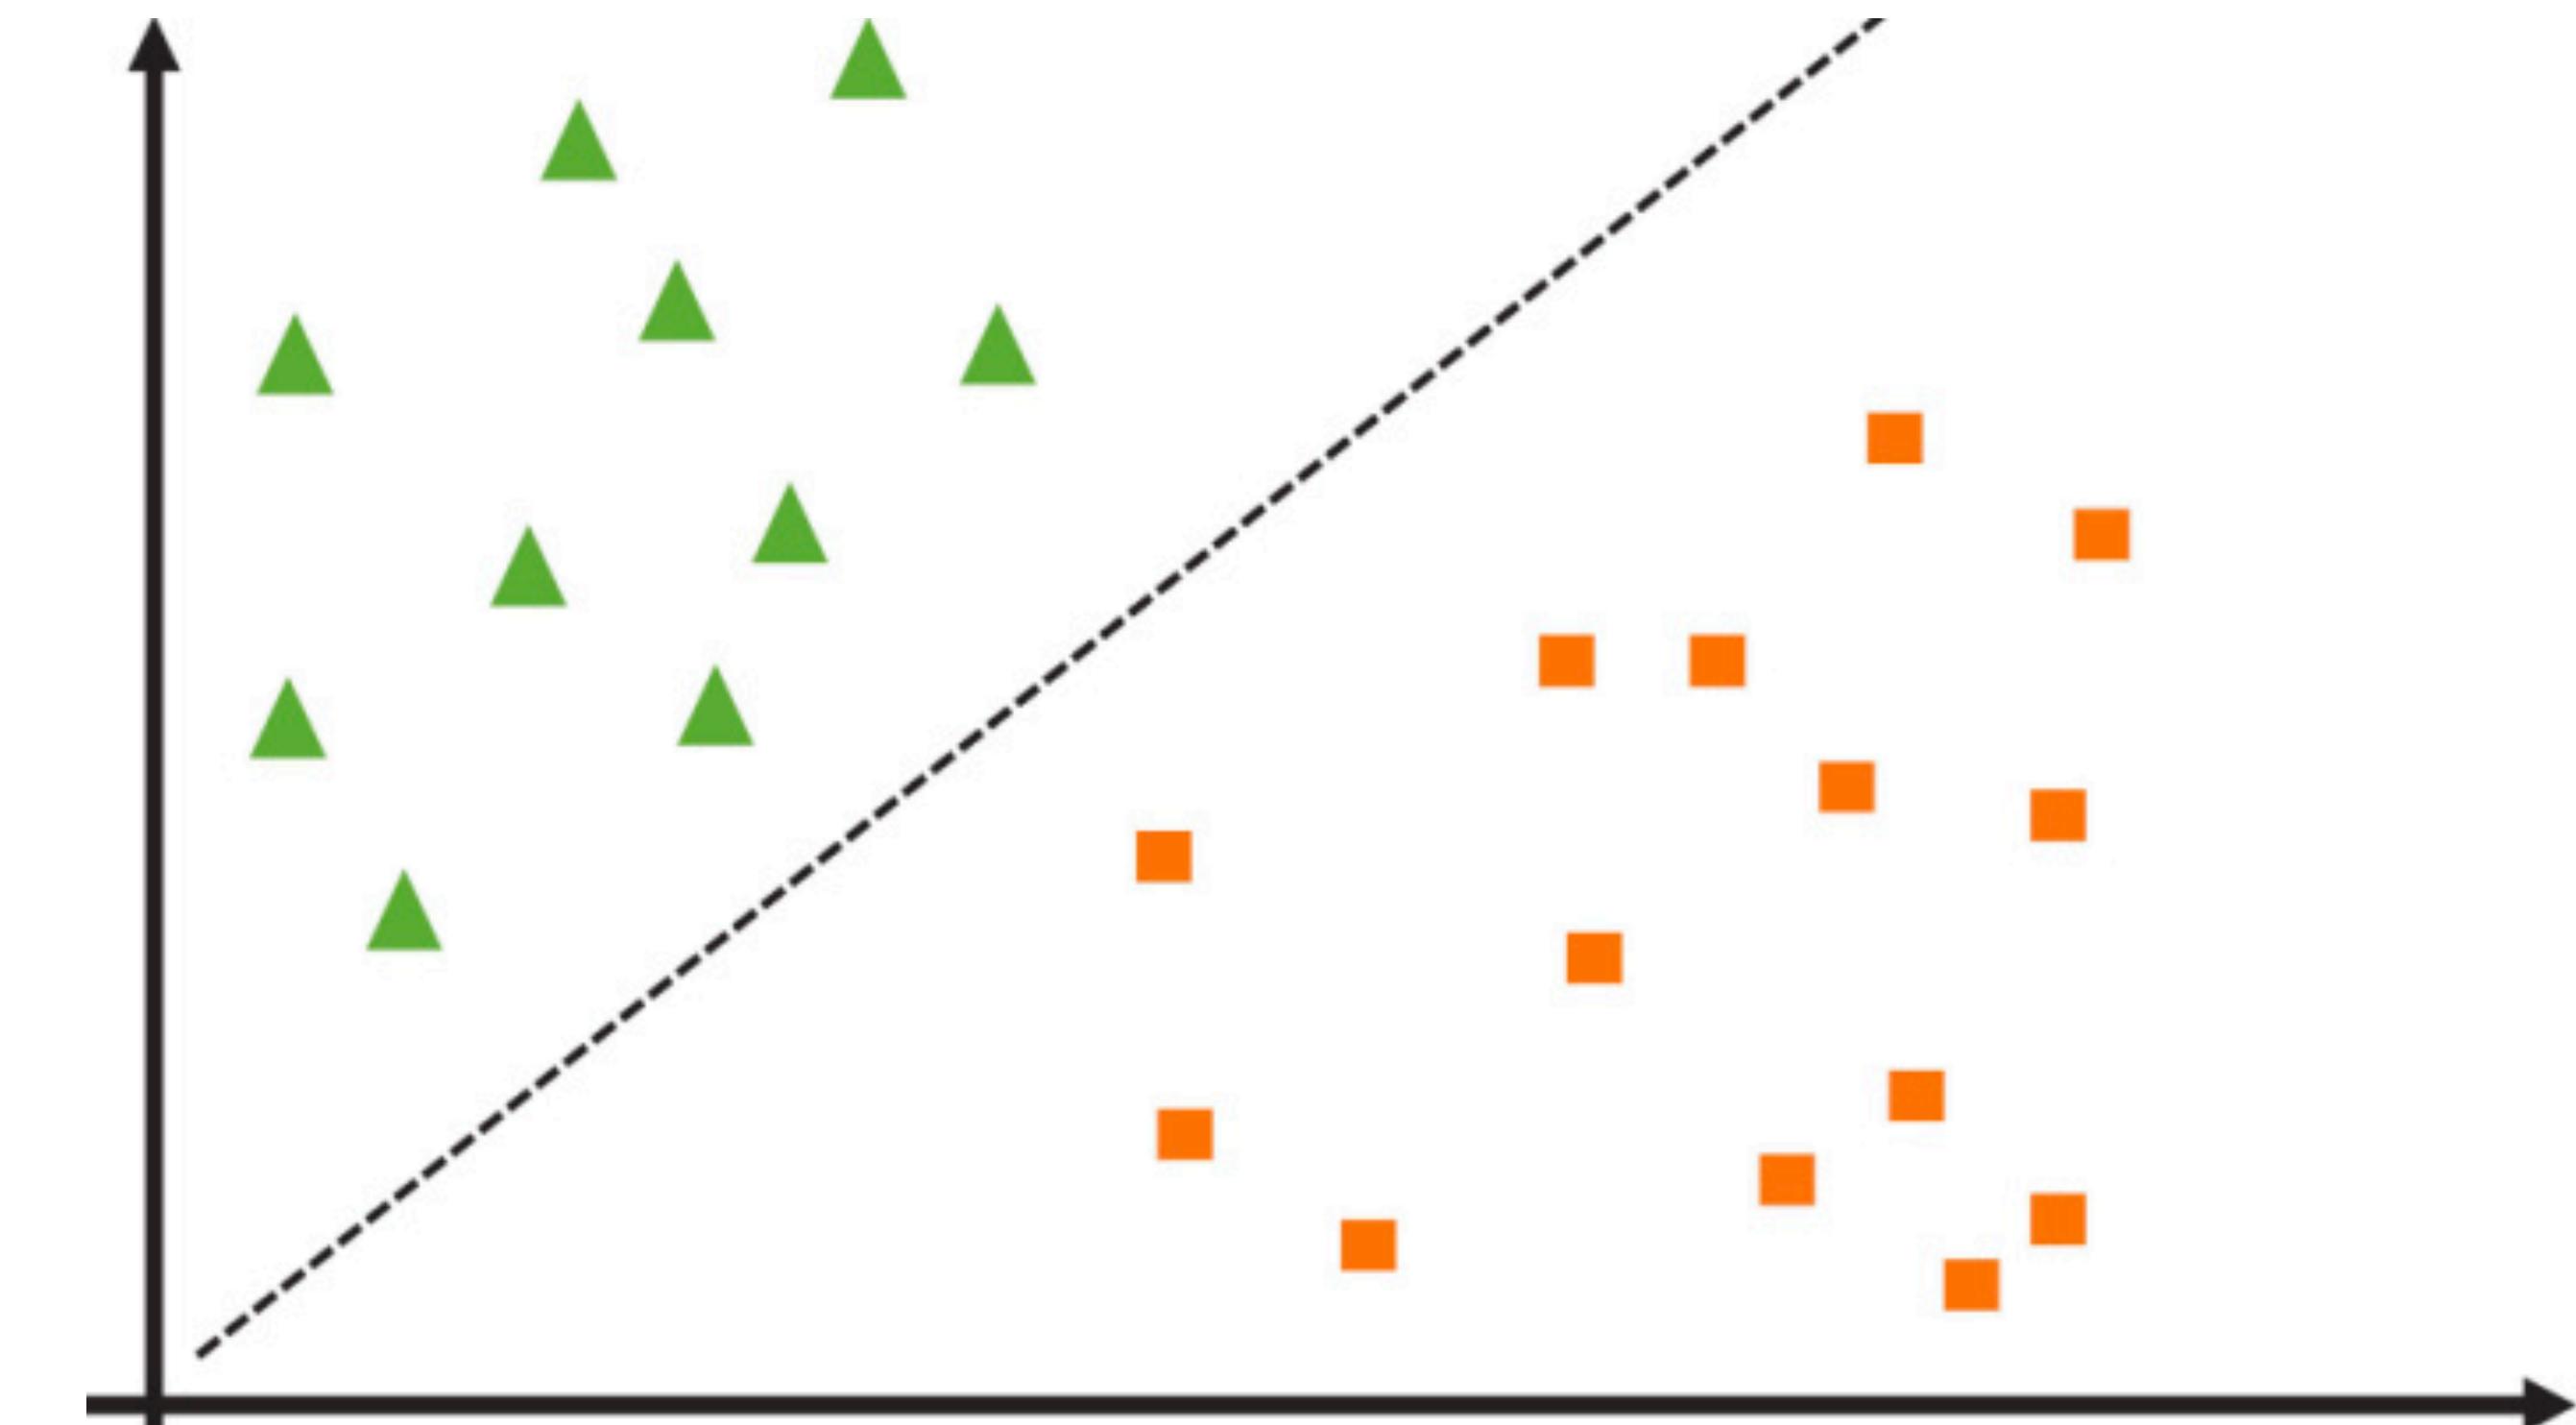
$$\sigma(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\sigma(x) = \max(0, x)$$

Why should I use activations functions?

Why should I use activations functions?

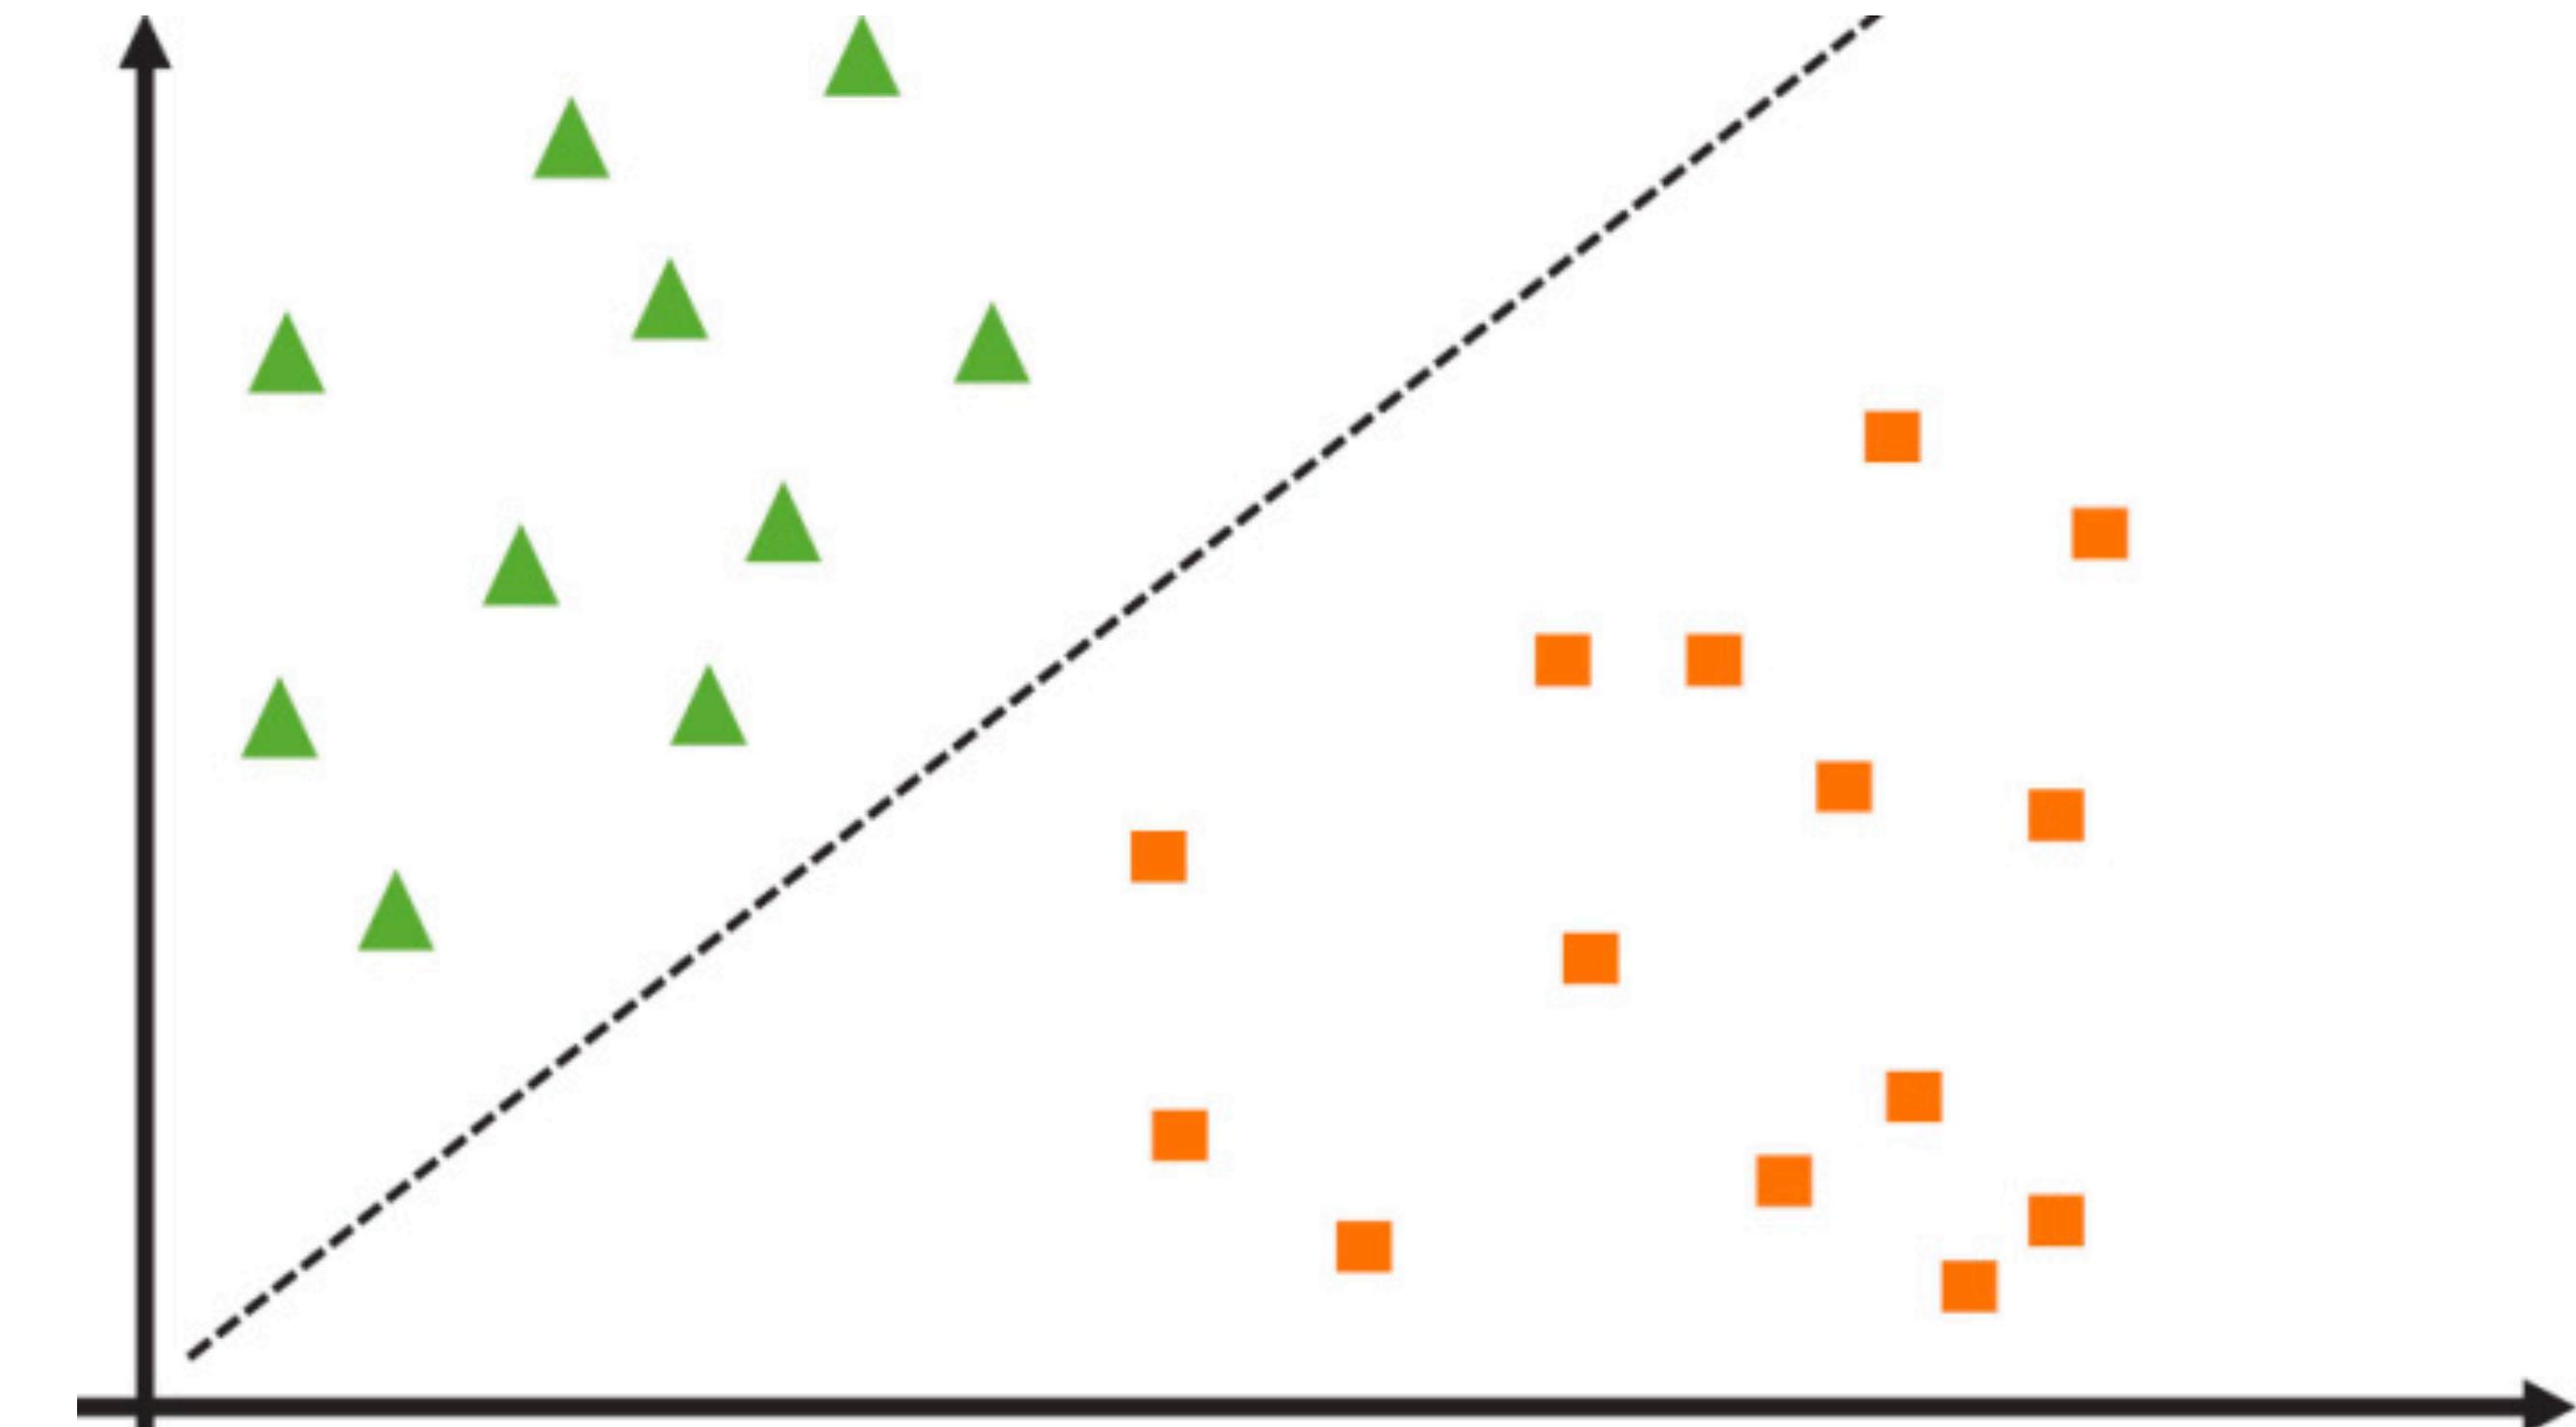
**Suppose your goal is to separate/
classify, the green triangles from
the orange squares.....**



Why should I use activations functions?

**Suppose your goal is to separate/
classify, the green triangles from
the orange squares.....**

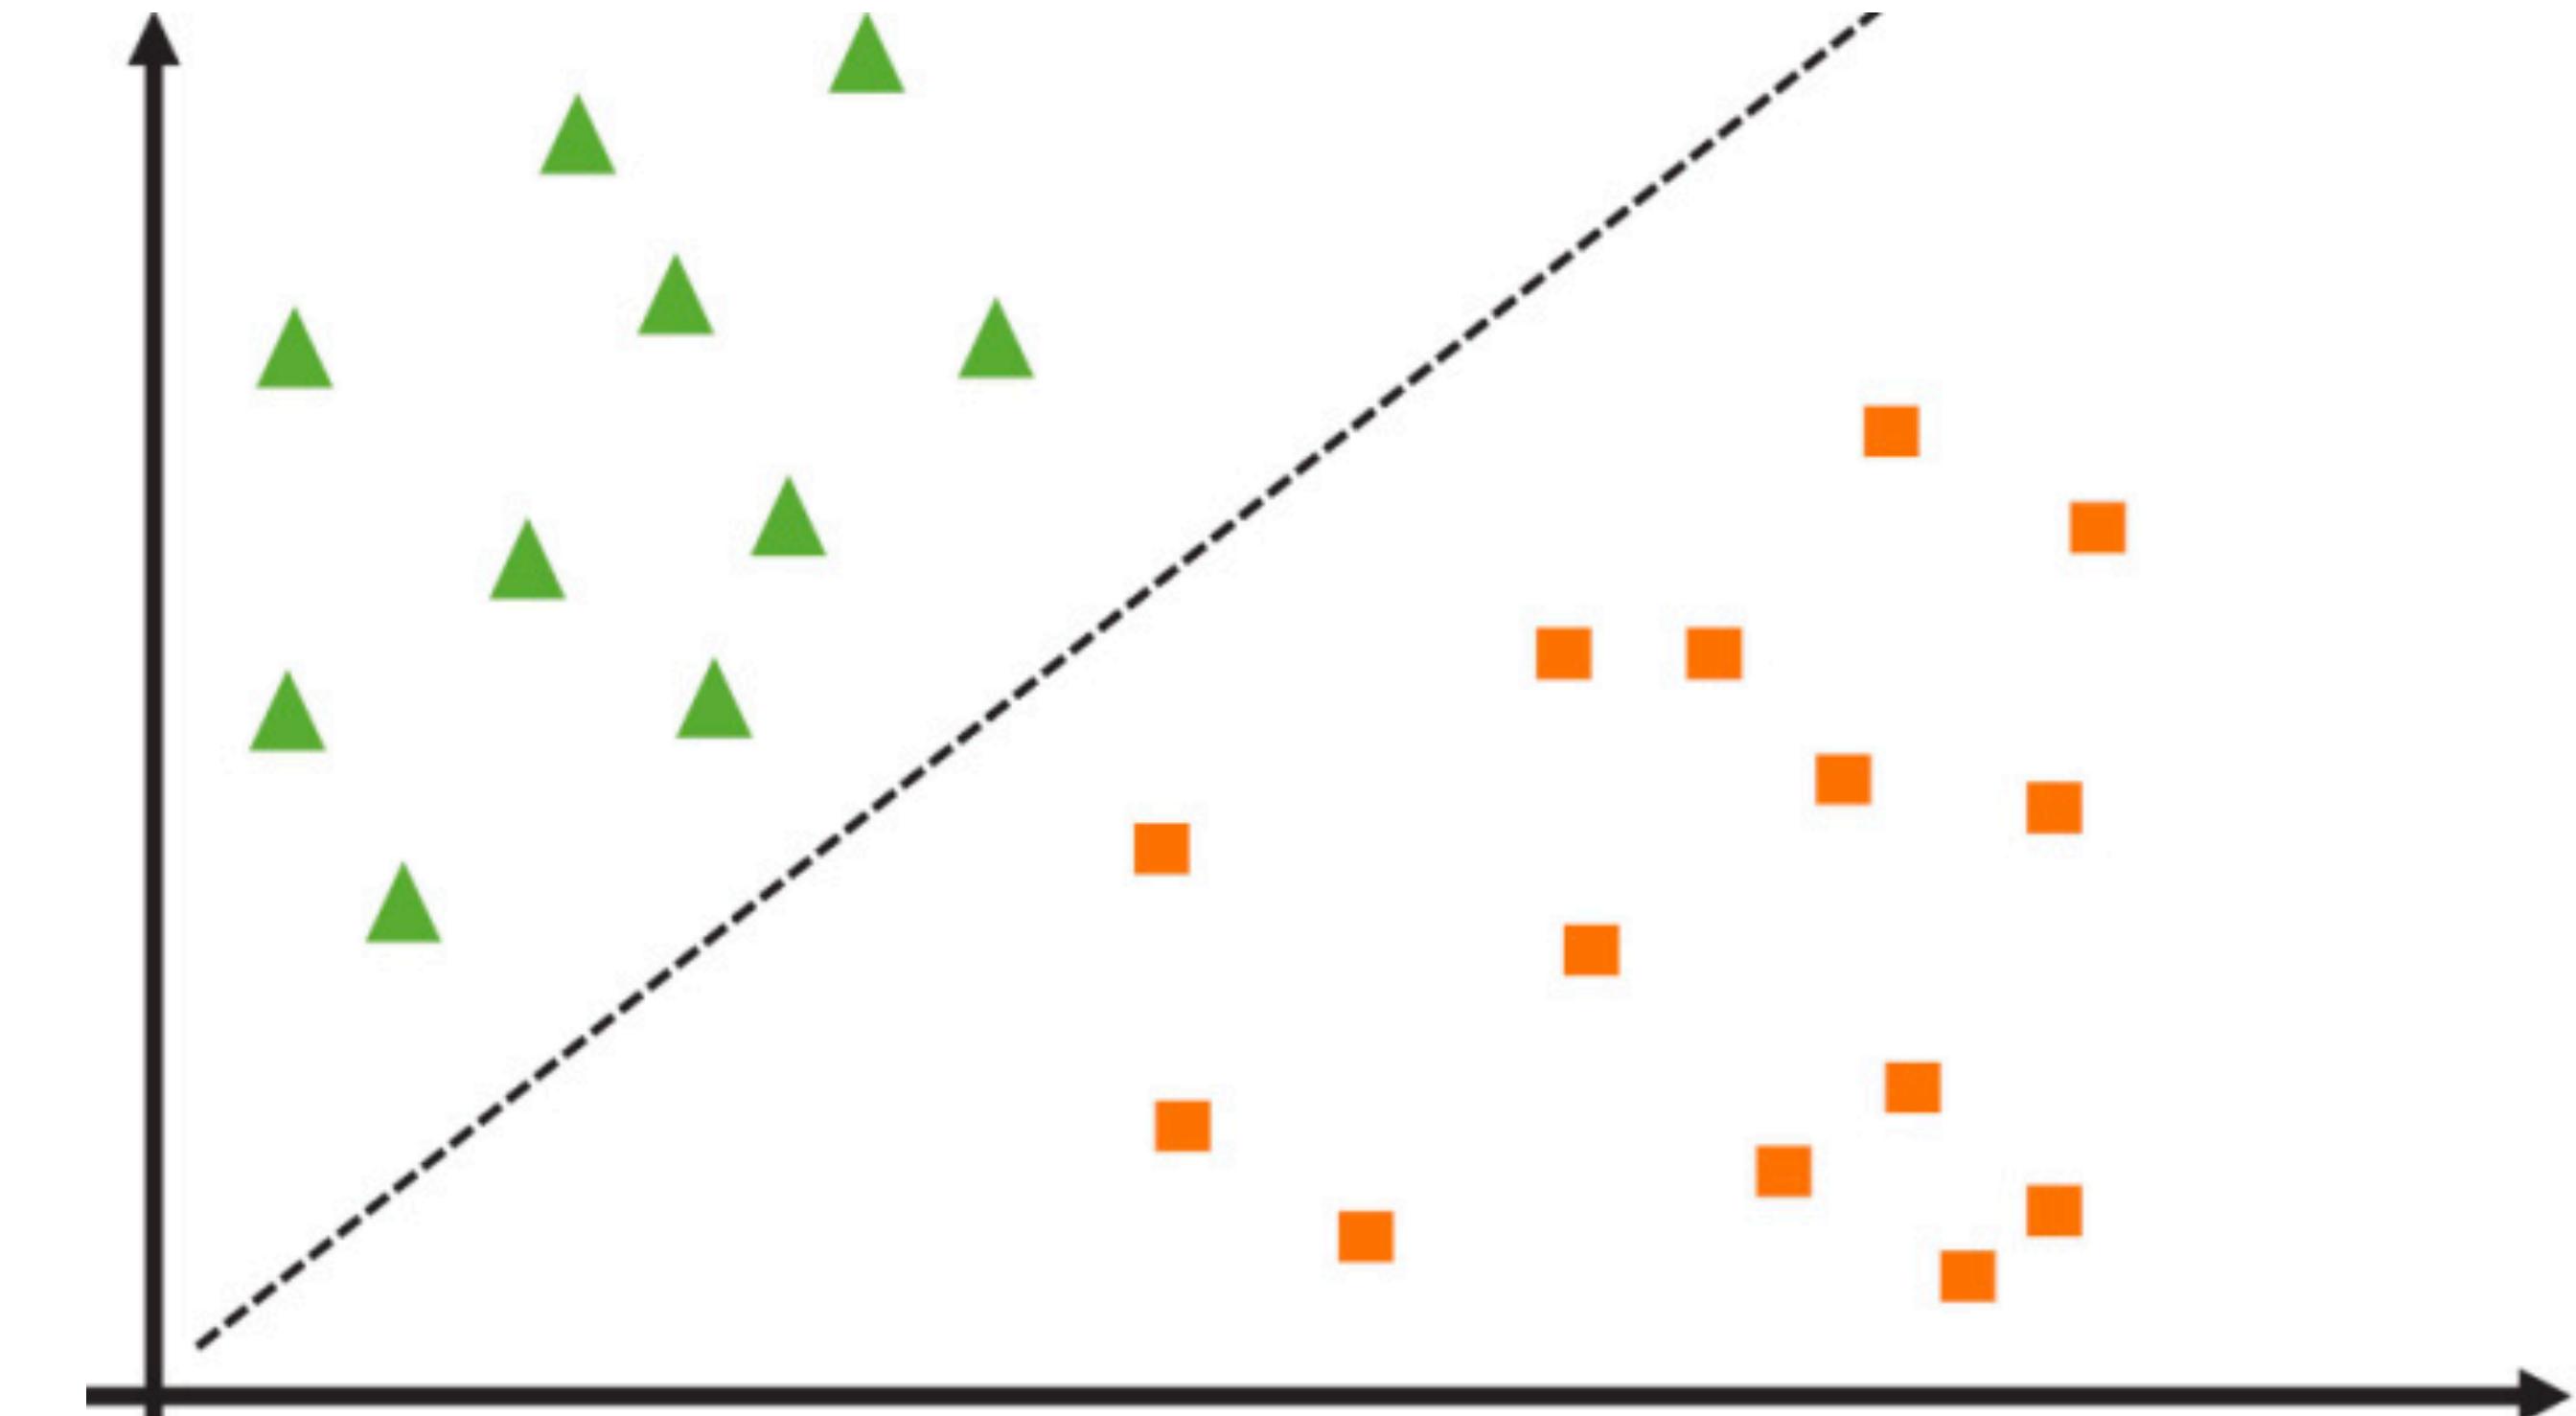
**If you do not use linearity
you decision
boundary will **an** hyperplan**



Why should I use activations functions?

Suppose your goal is to **separate/classify**, the green triangles from the orange squares.....

If you do not use linearity
you decision
boundary will **an hyperplan**

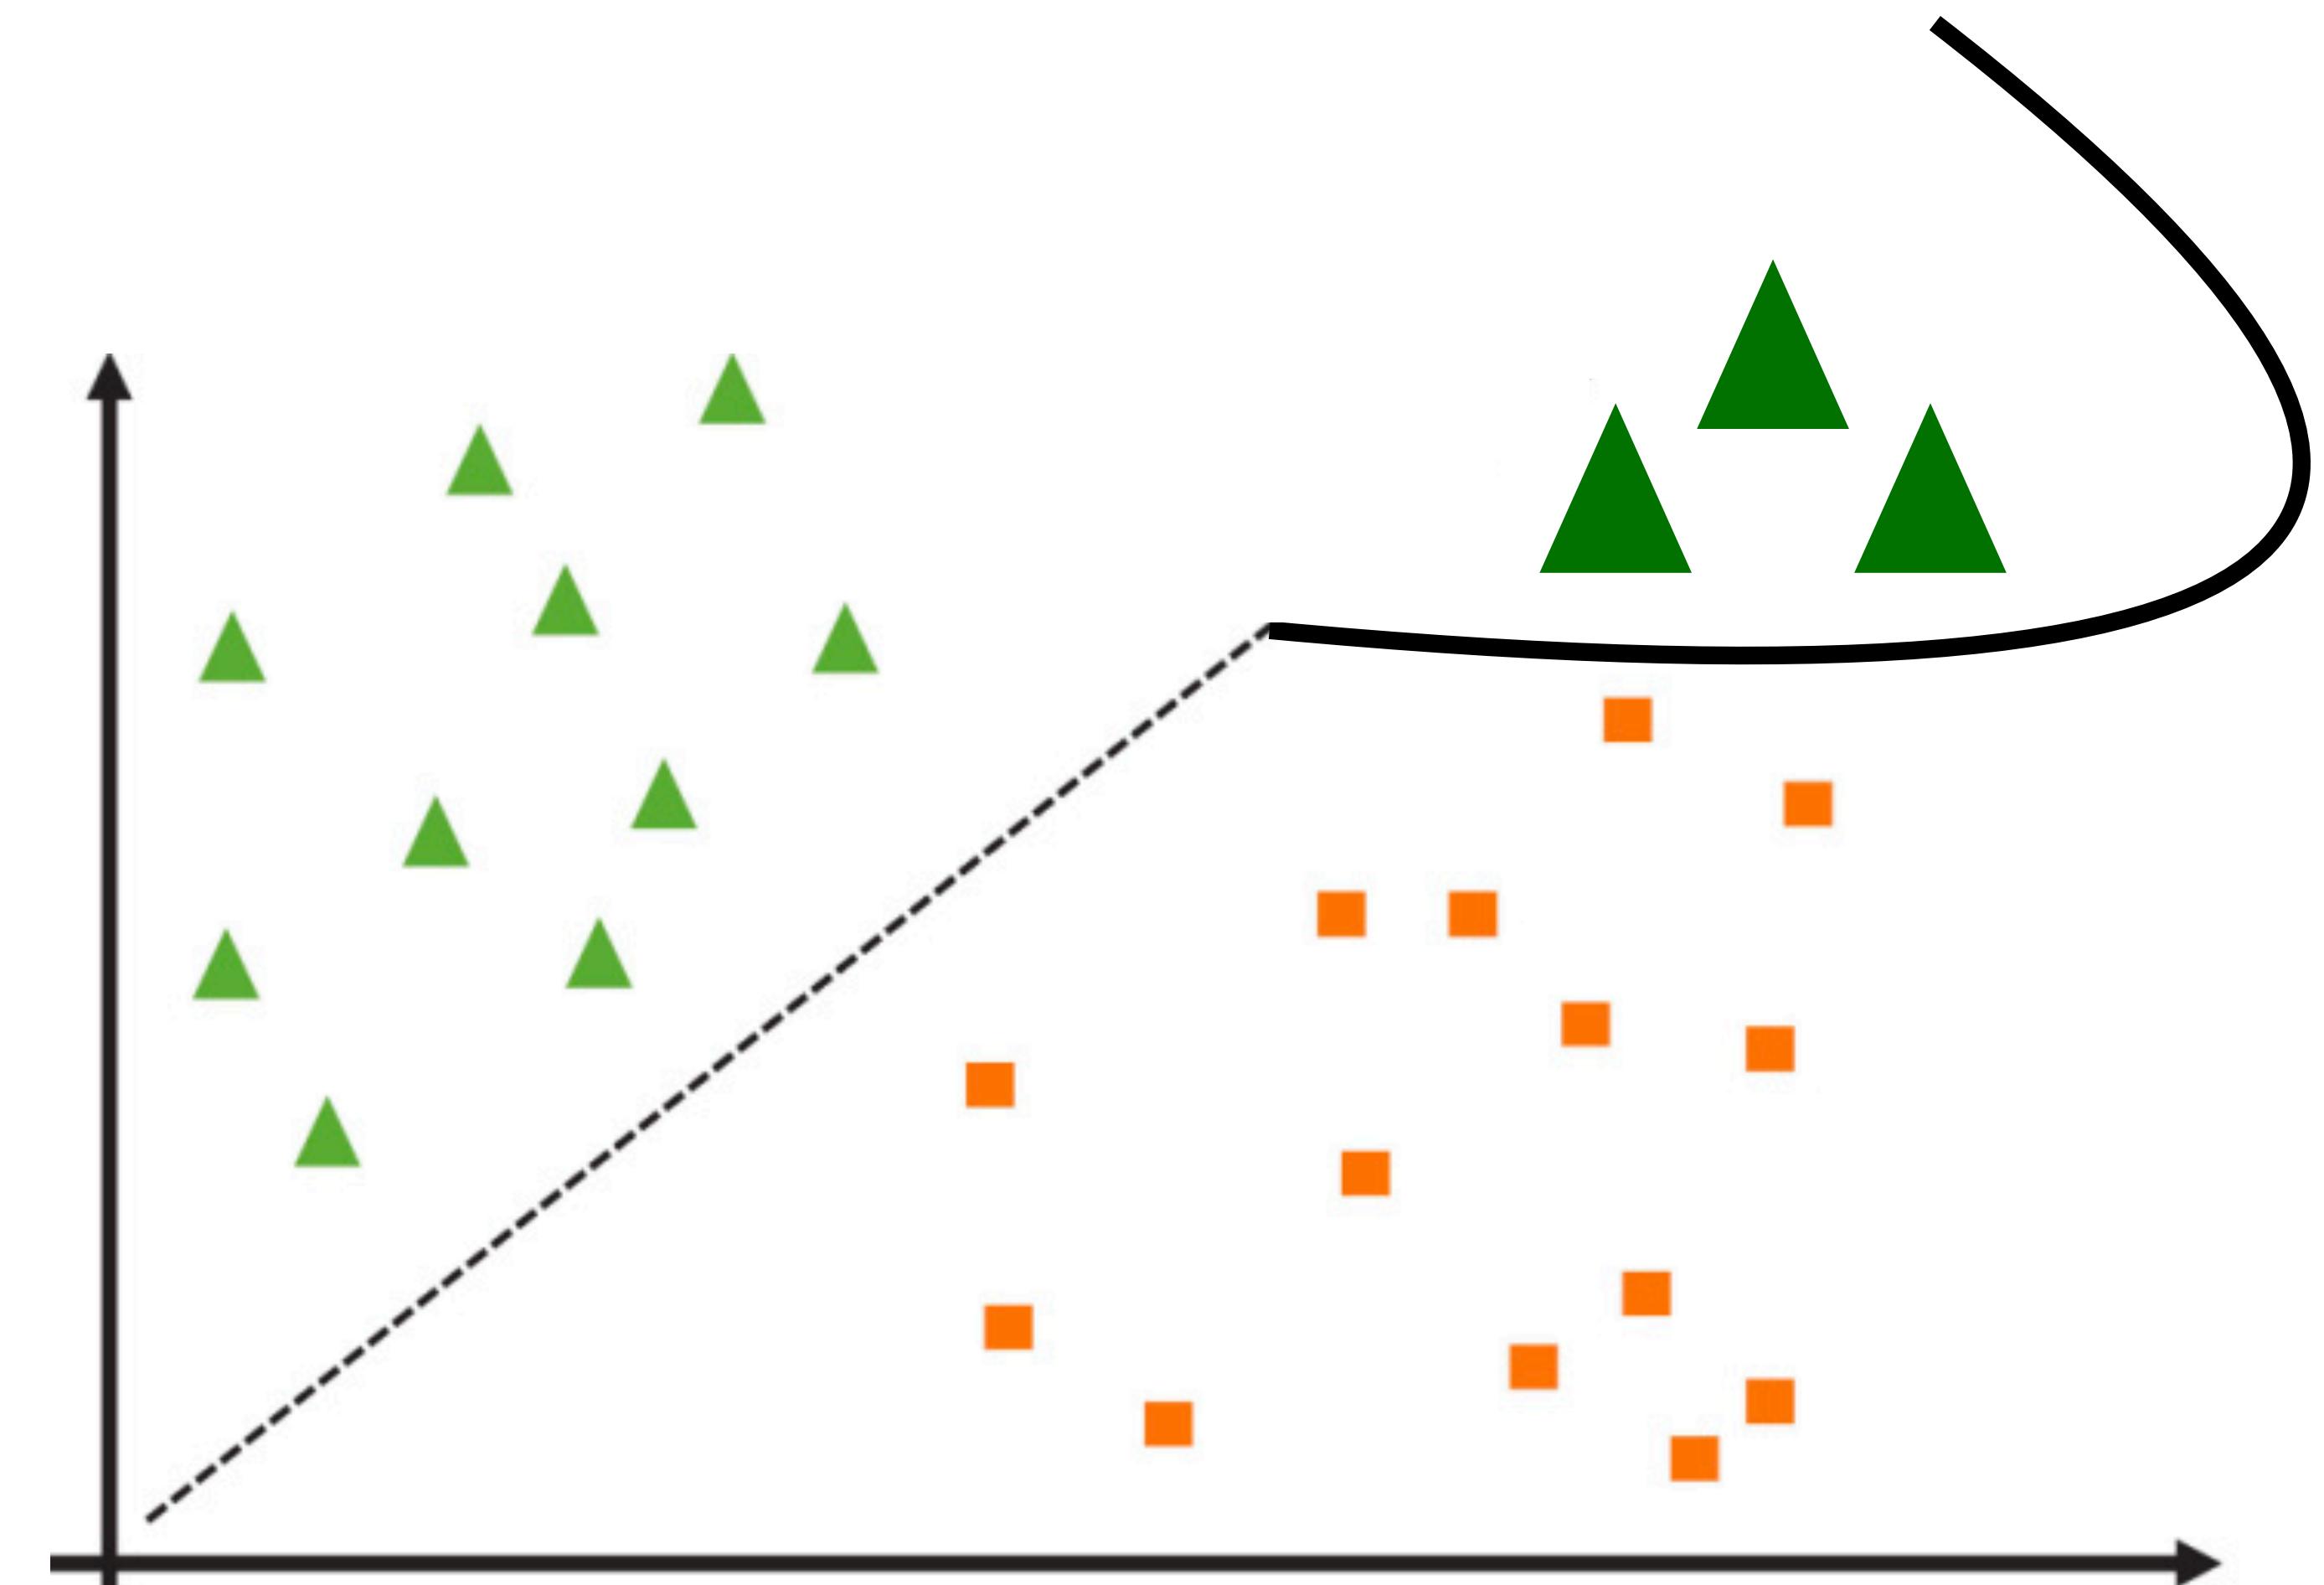


In this example 2D points represent the input (e.g. images/text)

Why should I use activations functions?

Why should I use activations functions?

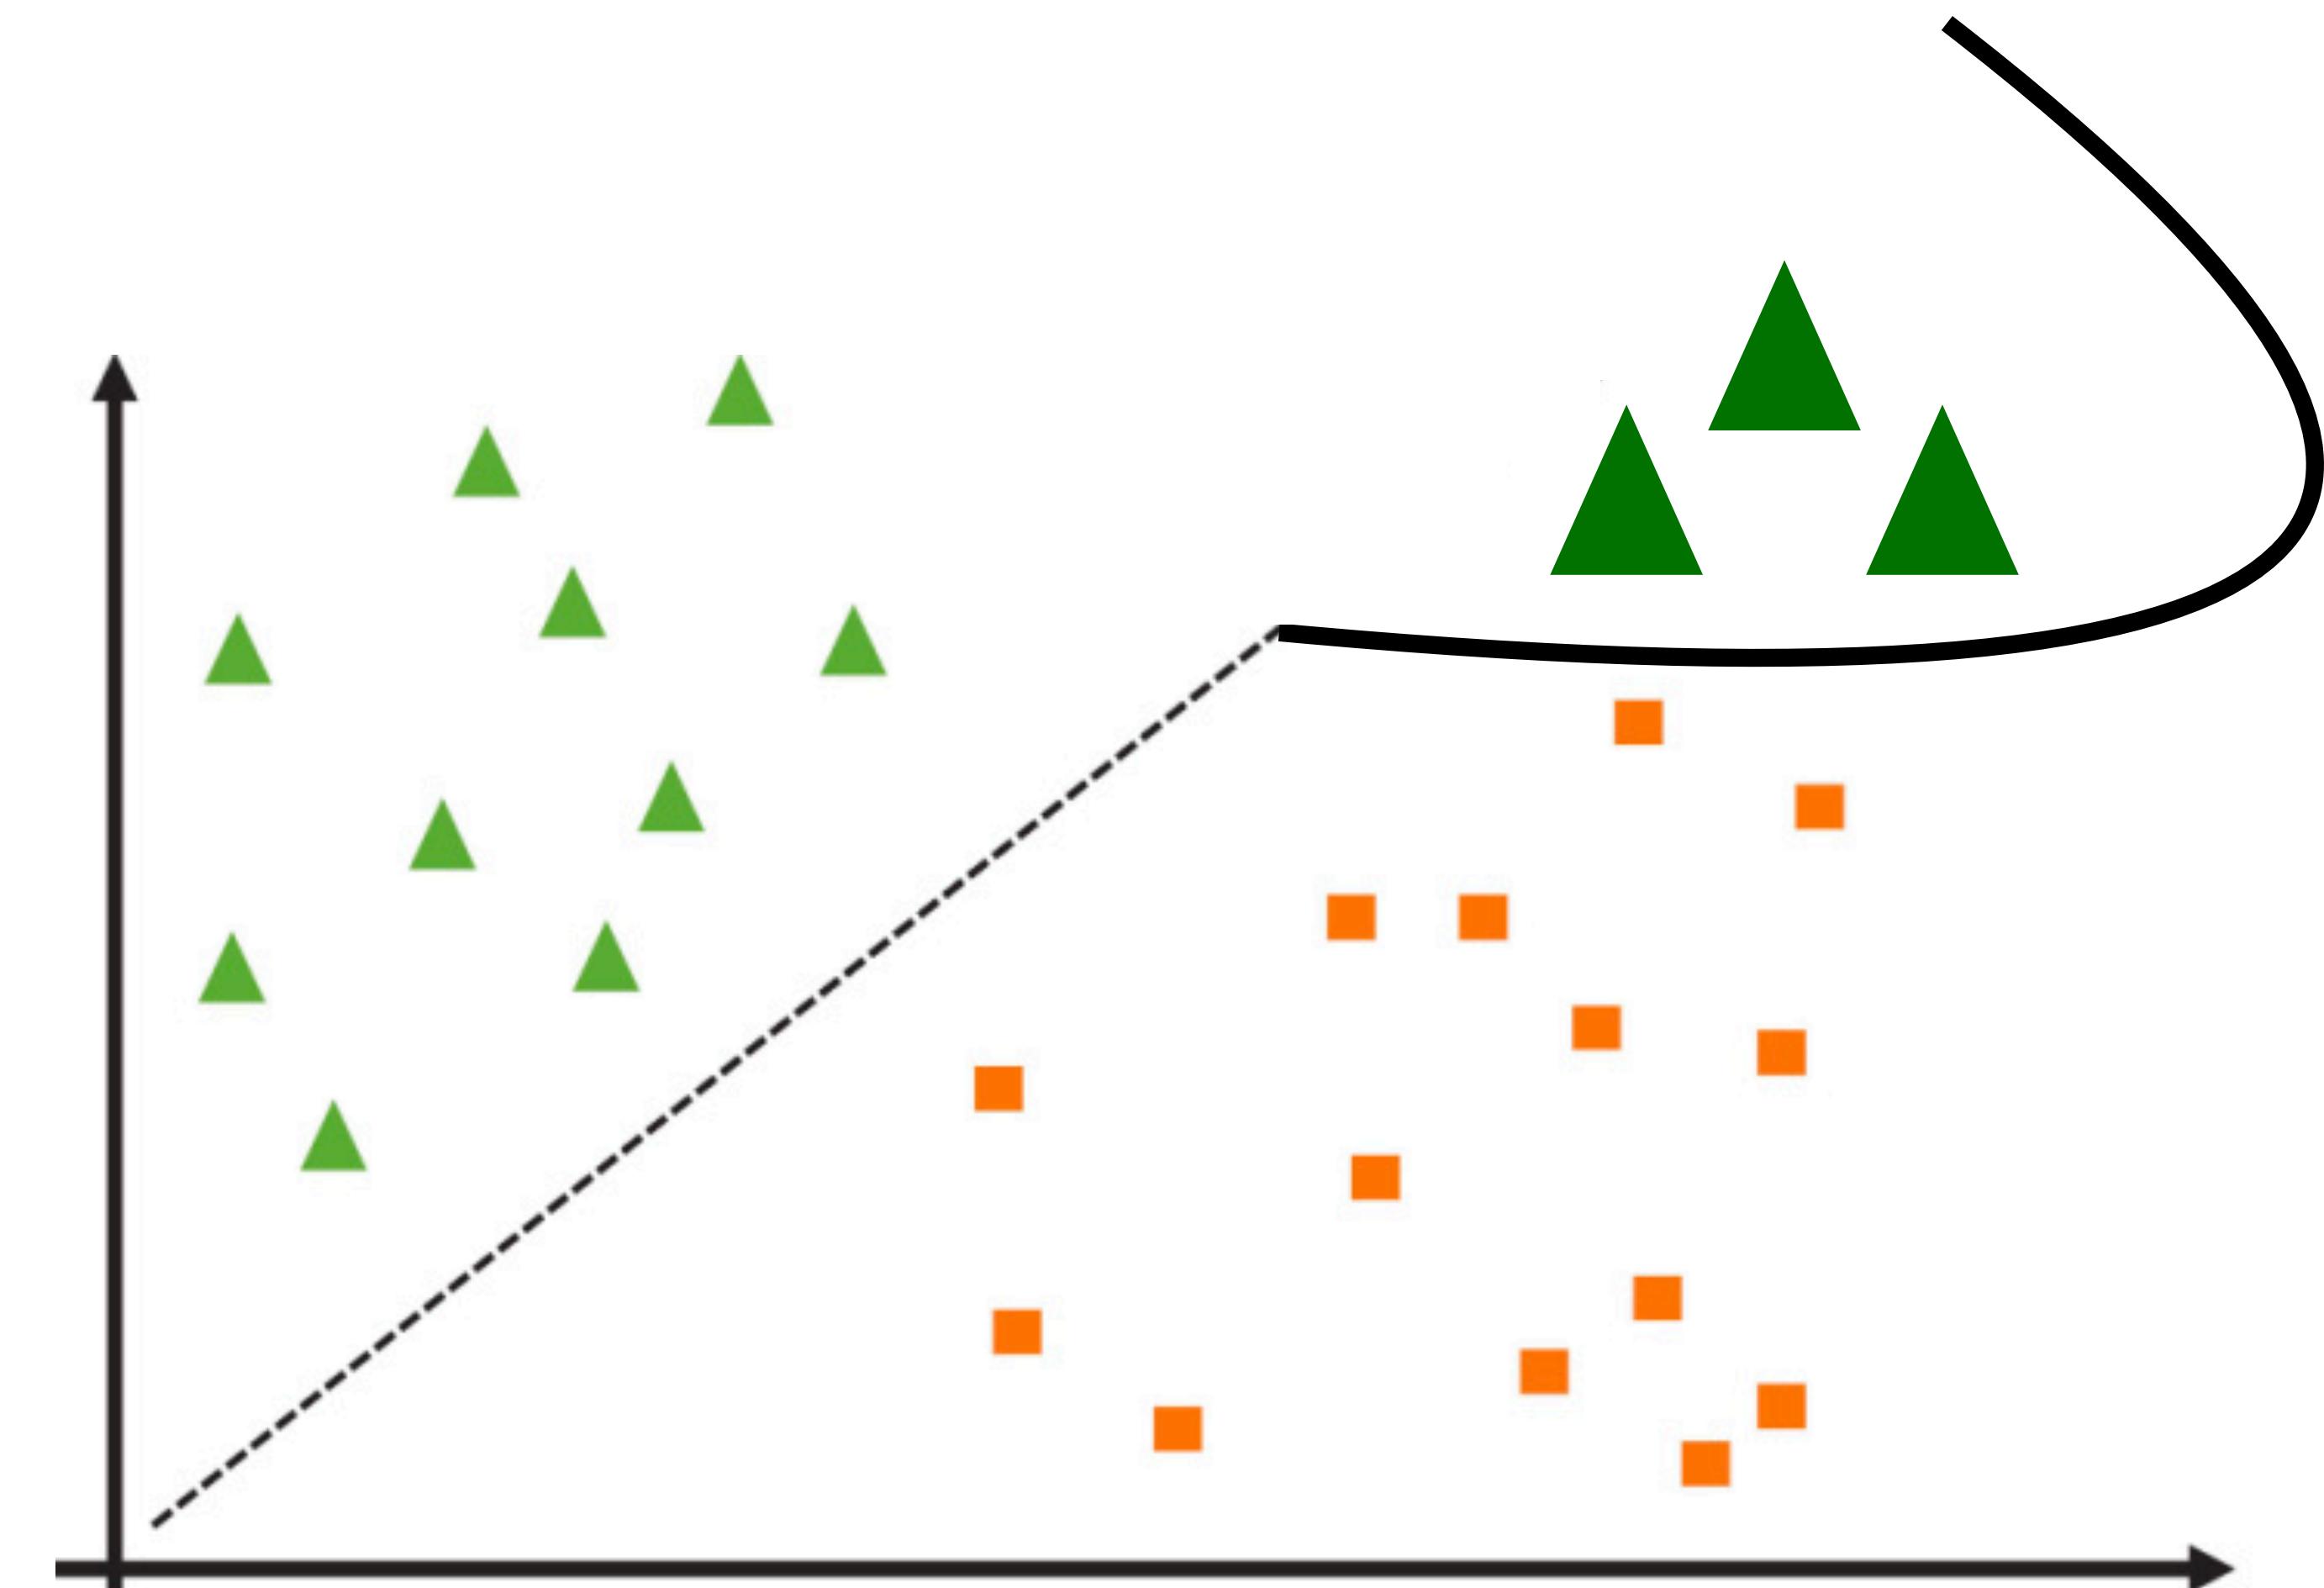
**Suppose your goal is to separate,
the green triangles from the orange
squares.....**



Why should I use activations functions?

**Suppose your goal is to separate,
the green triangles from the orange
squares.....**

**Thanks to the linearity you
can have more « complex
decision functions ».**

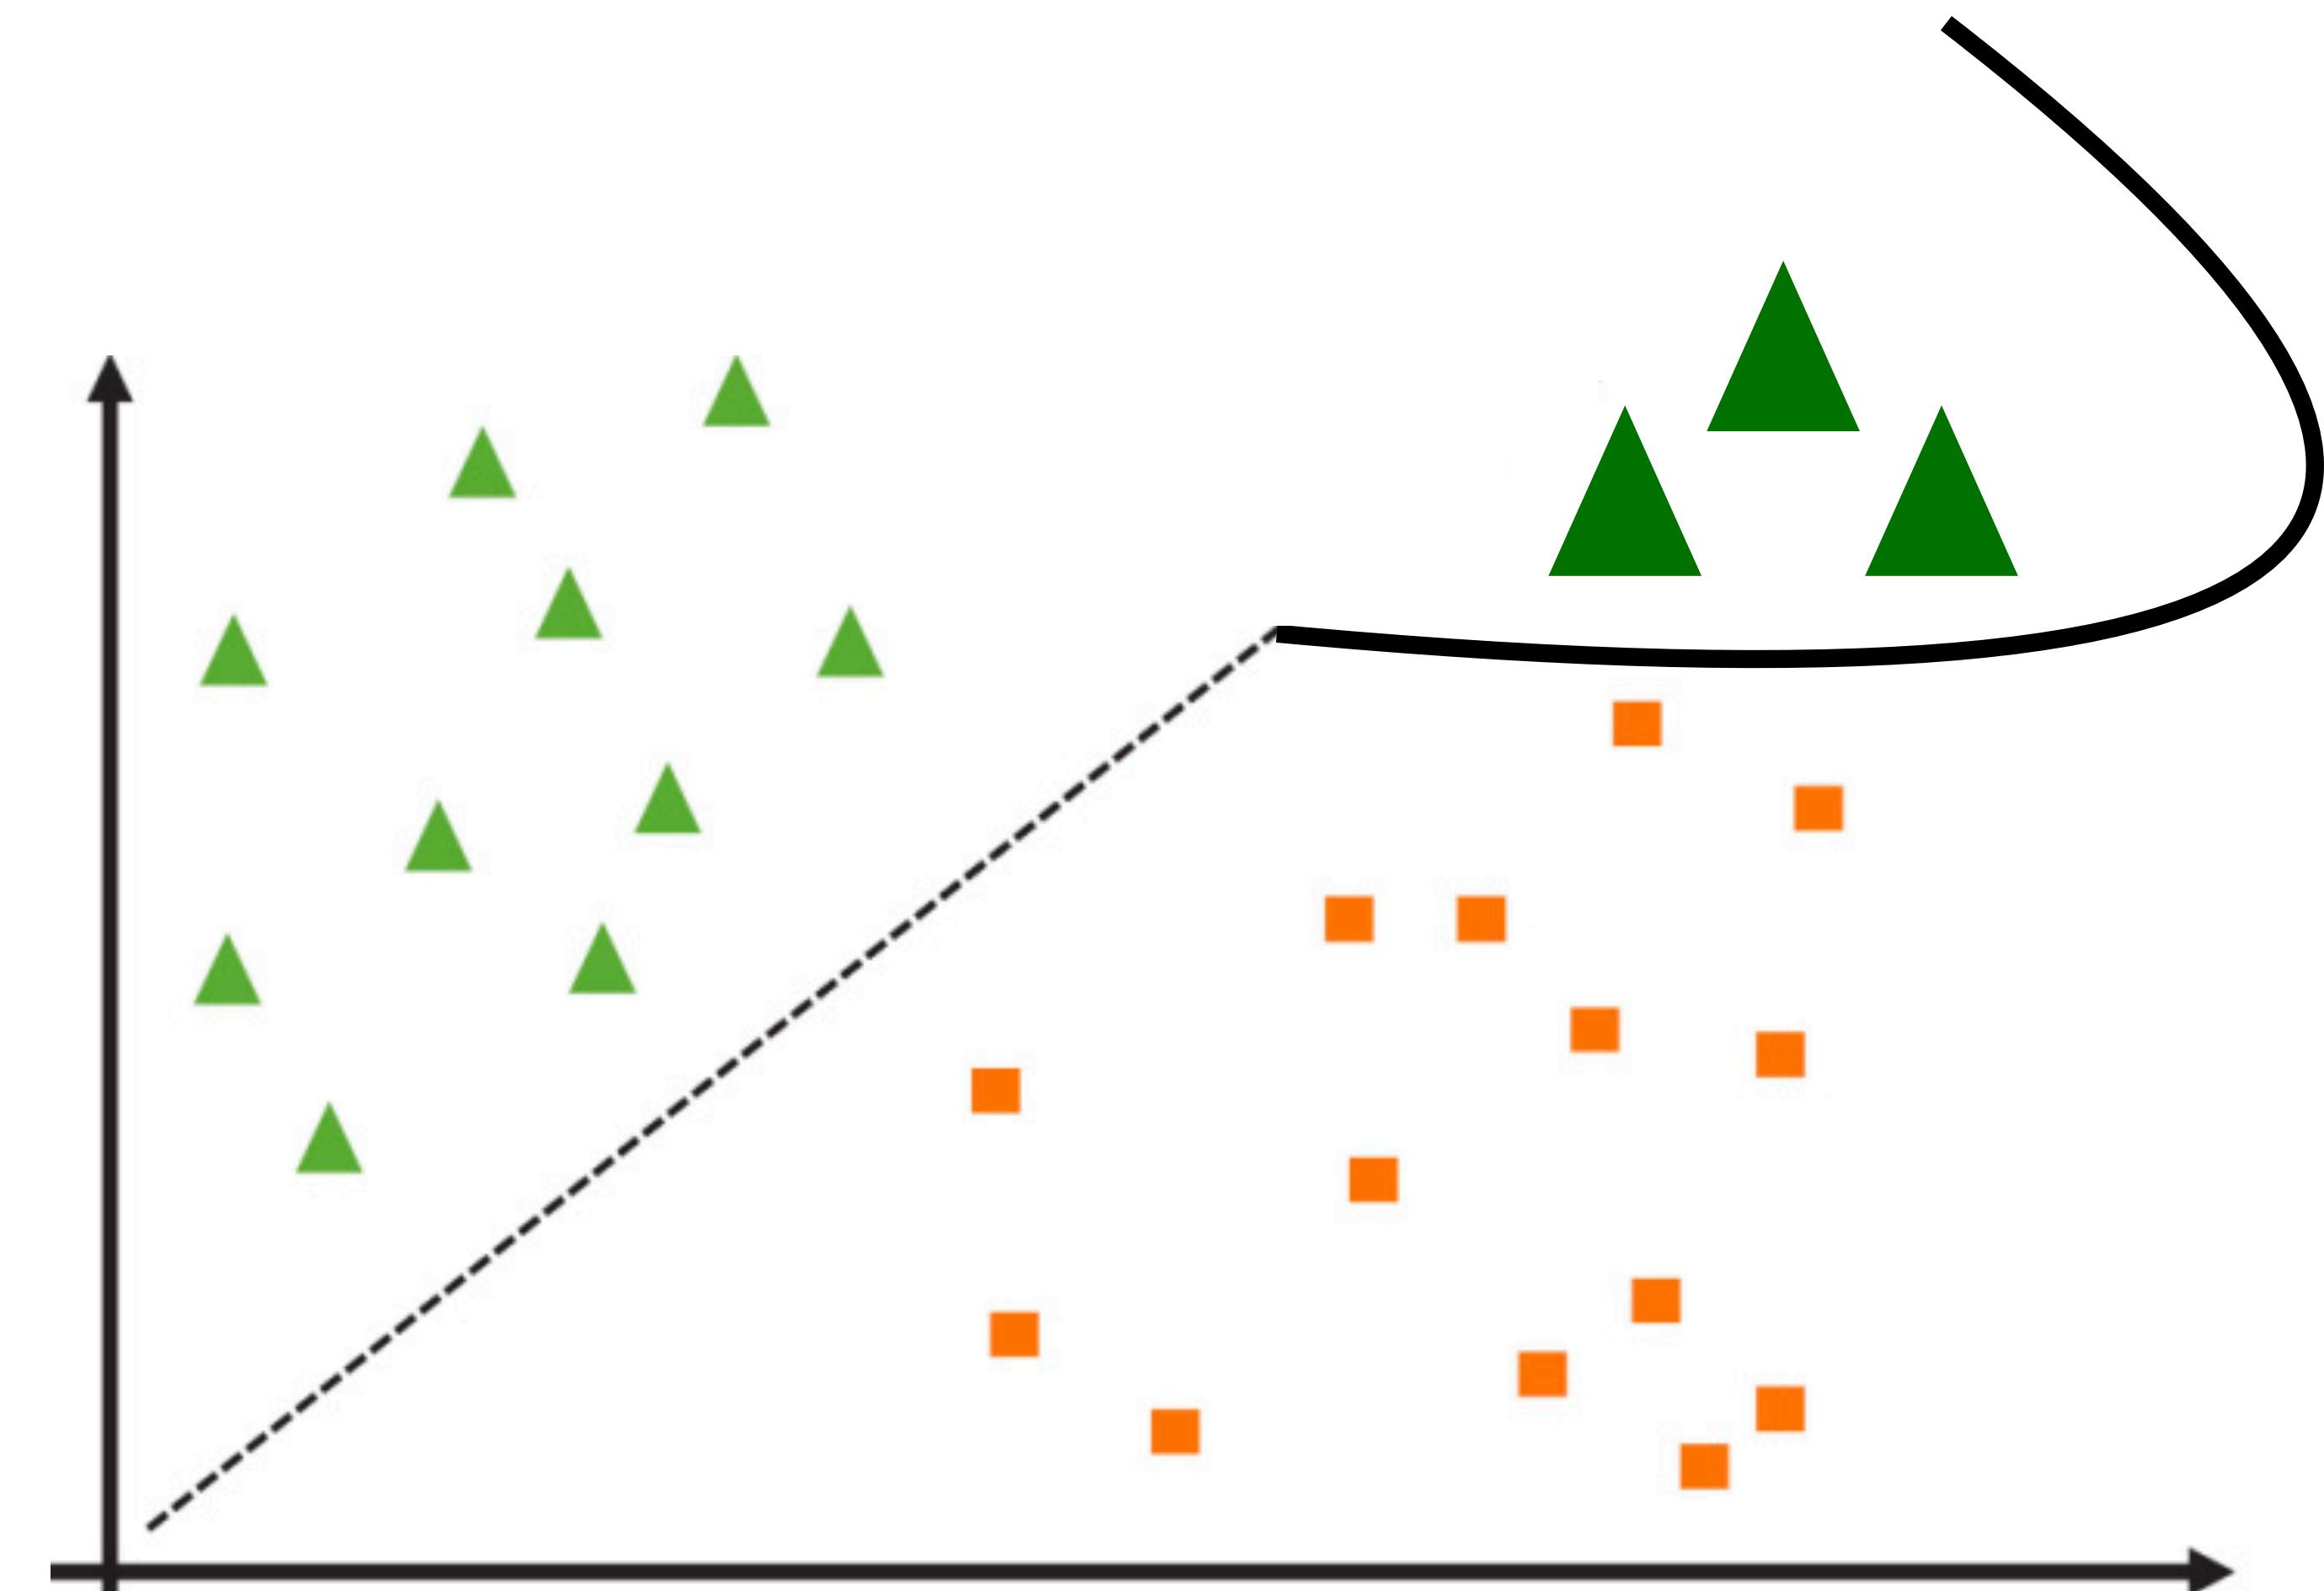


Why should I use activations functions?

**Suppose your goal is to separate,
the green triangles from the orange
squares.....**

**Thanks to the linearity you
can have more « complex
decision functions ».**

**See universal
approximation Theorem.**



Let's sum up so far

1. The basic bloc to build a Neural Network

We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

1. The basic bloc to build a Neural Network

We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

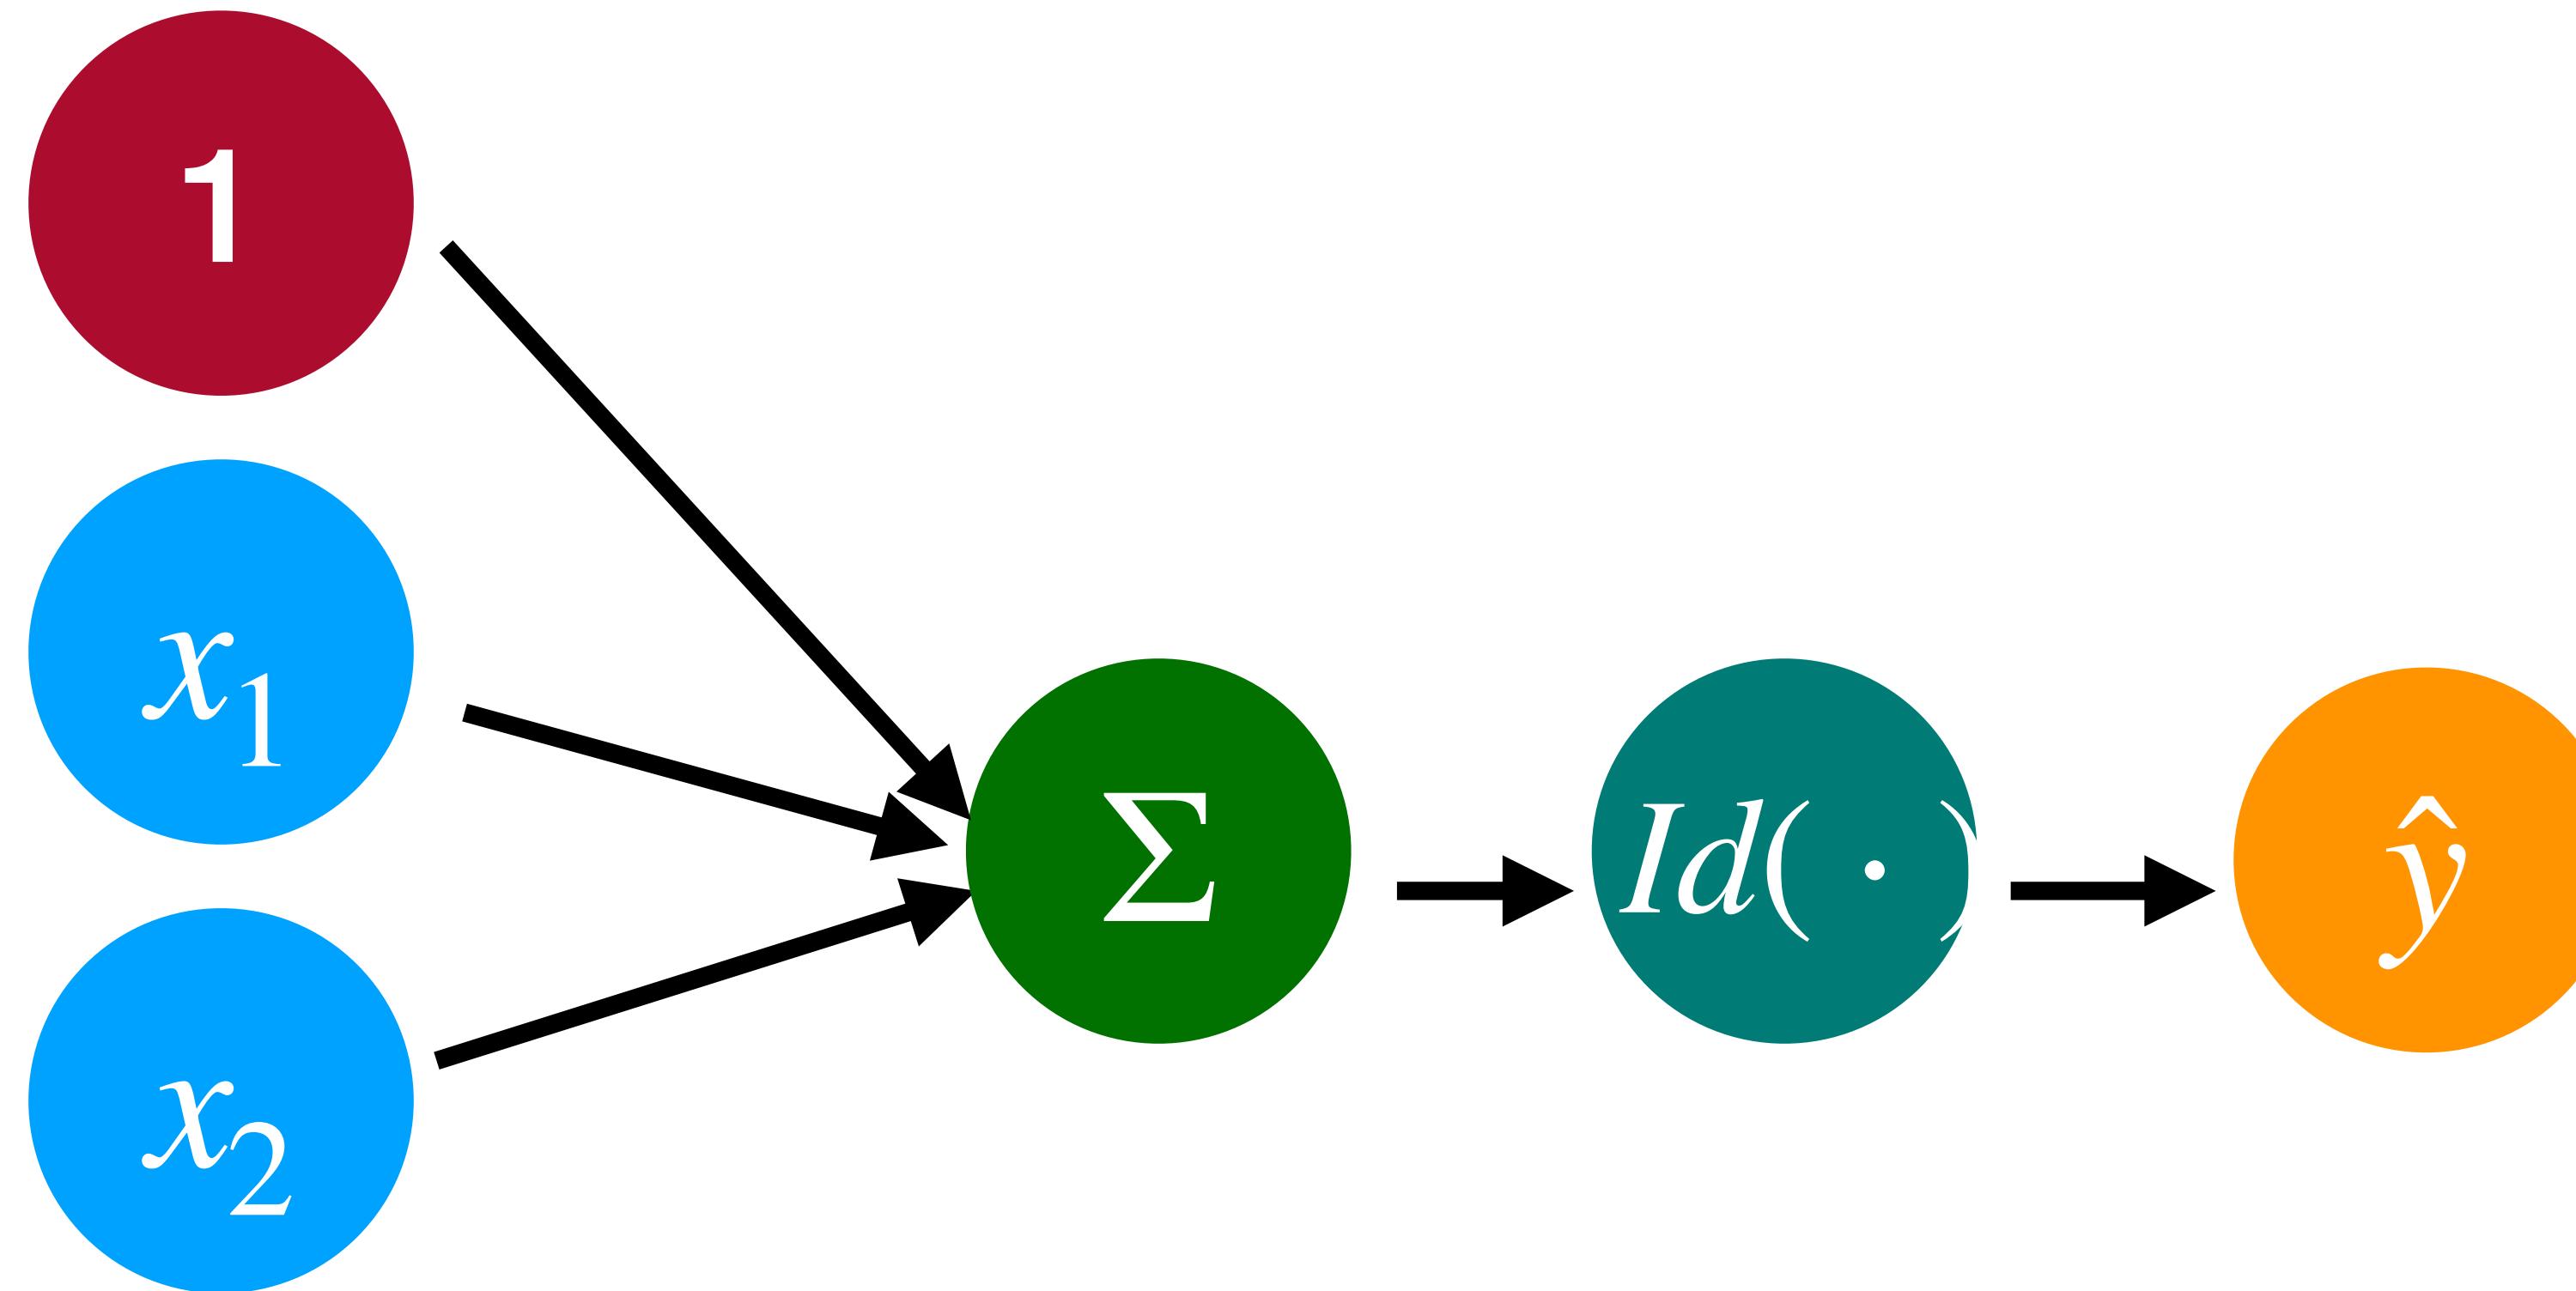
2. The role of activation functions

They complexity the neural network and allow to fit more complex data

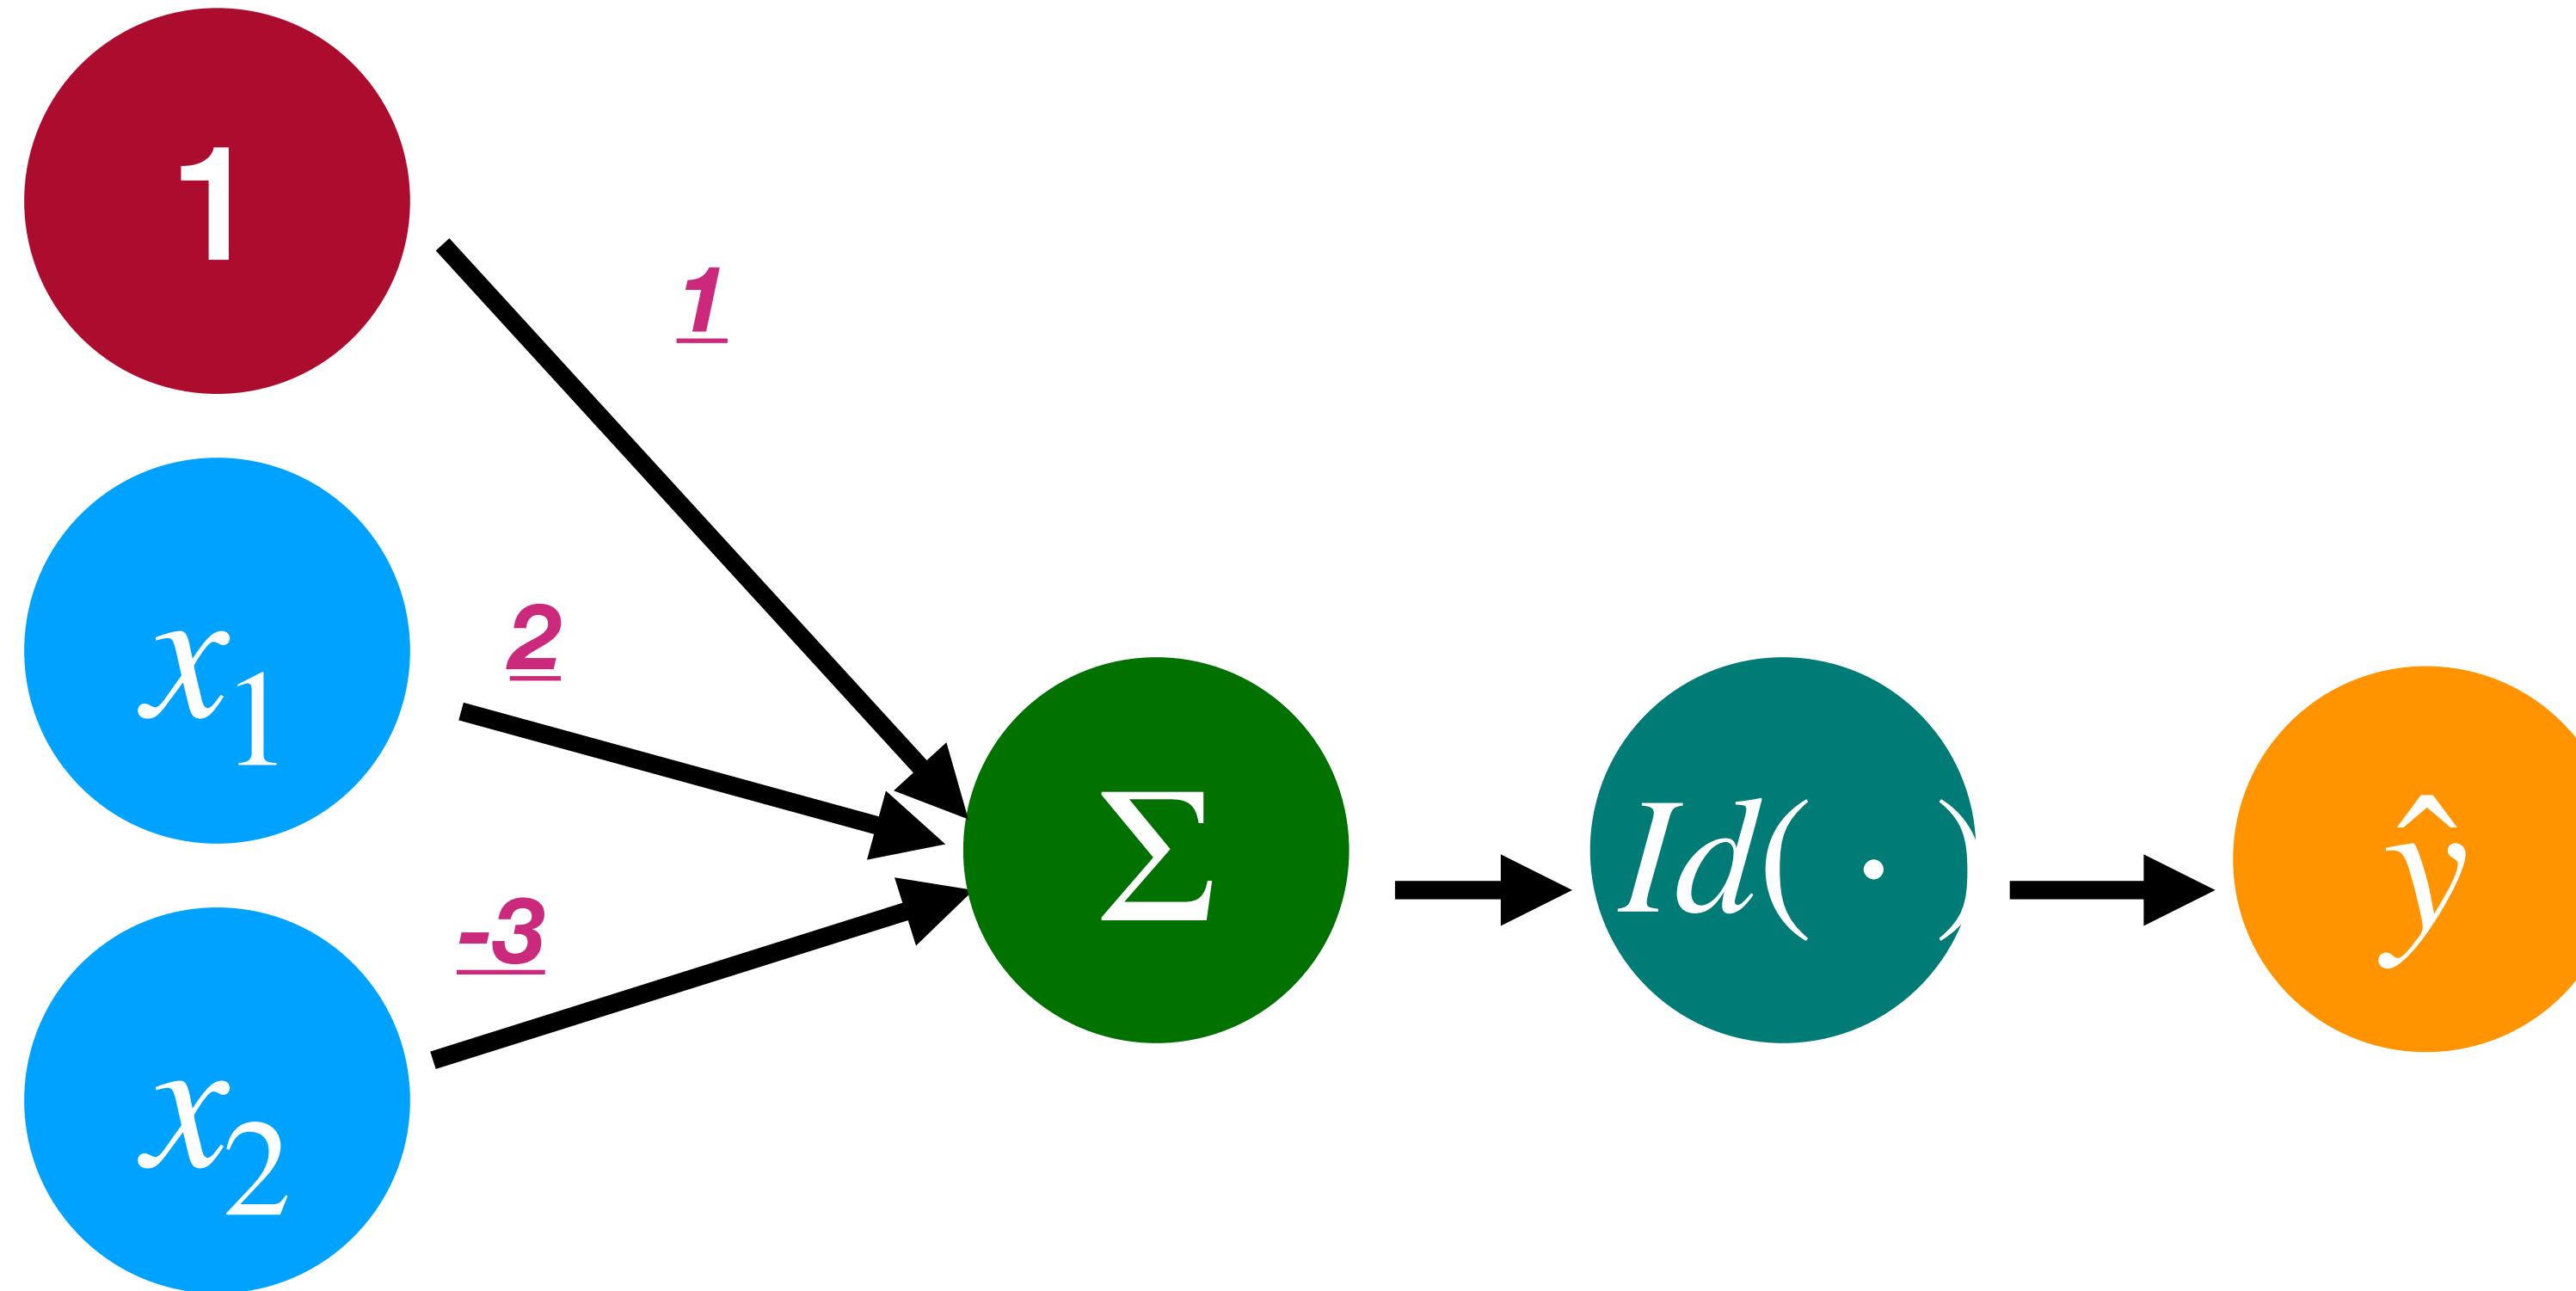
They are added at the end of each layer and there is a large variety of them

Our Perceptron to distinguish triangles and square

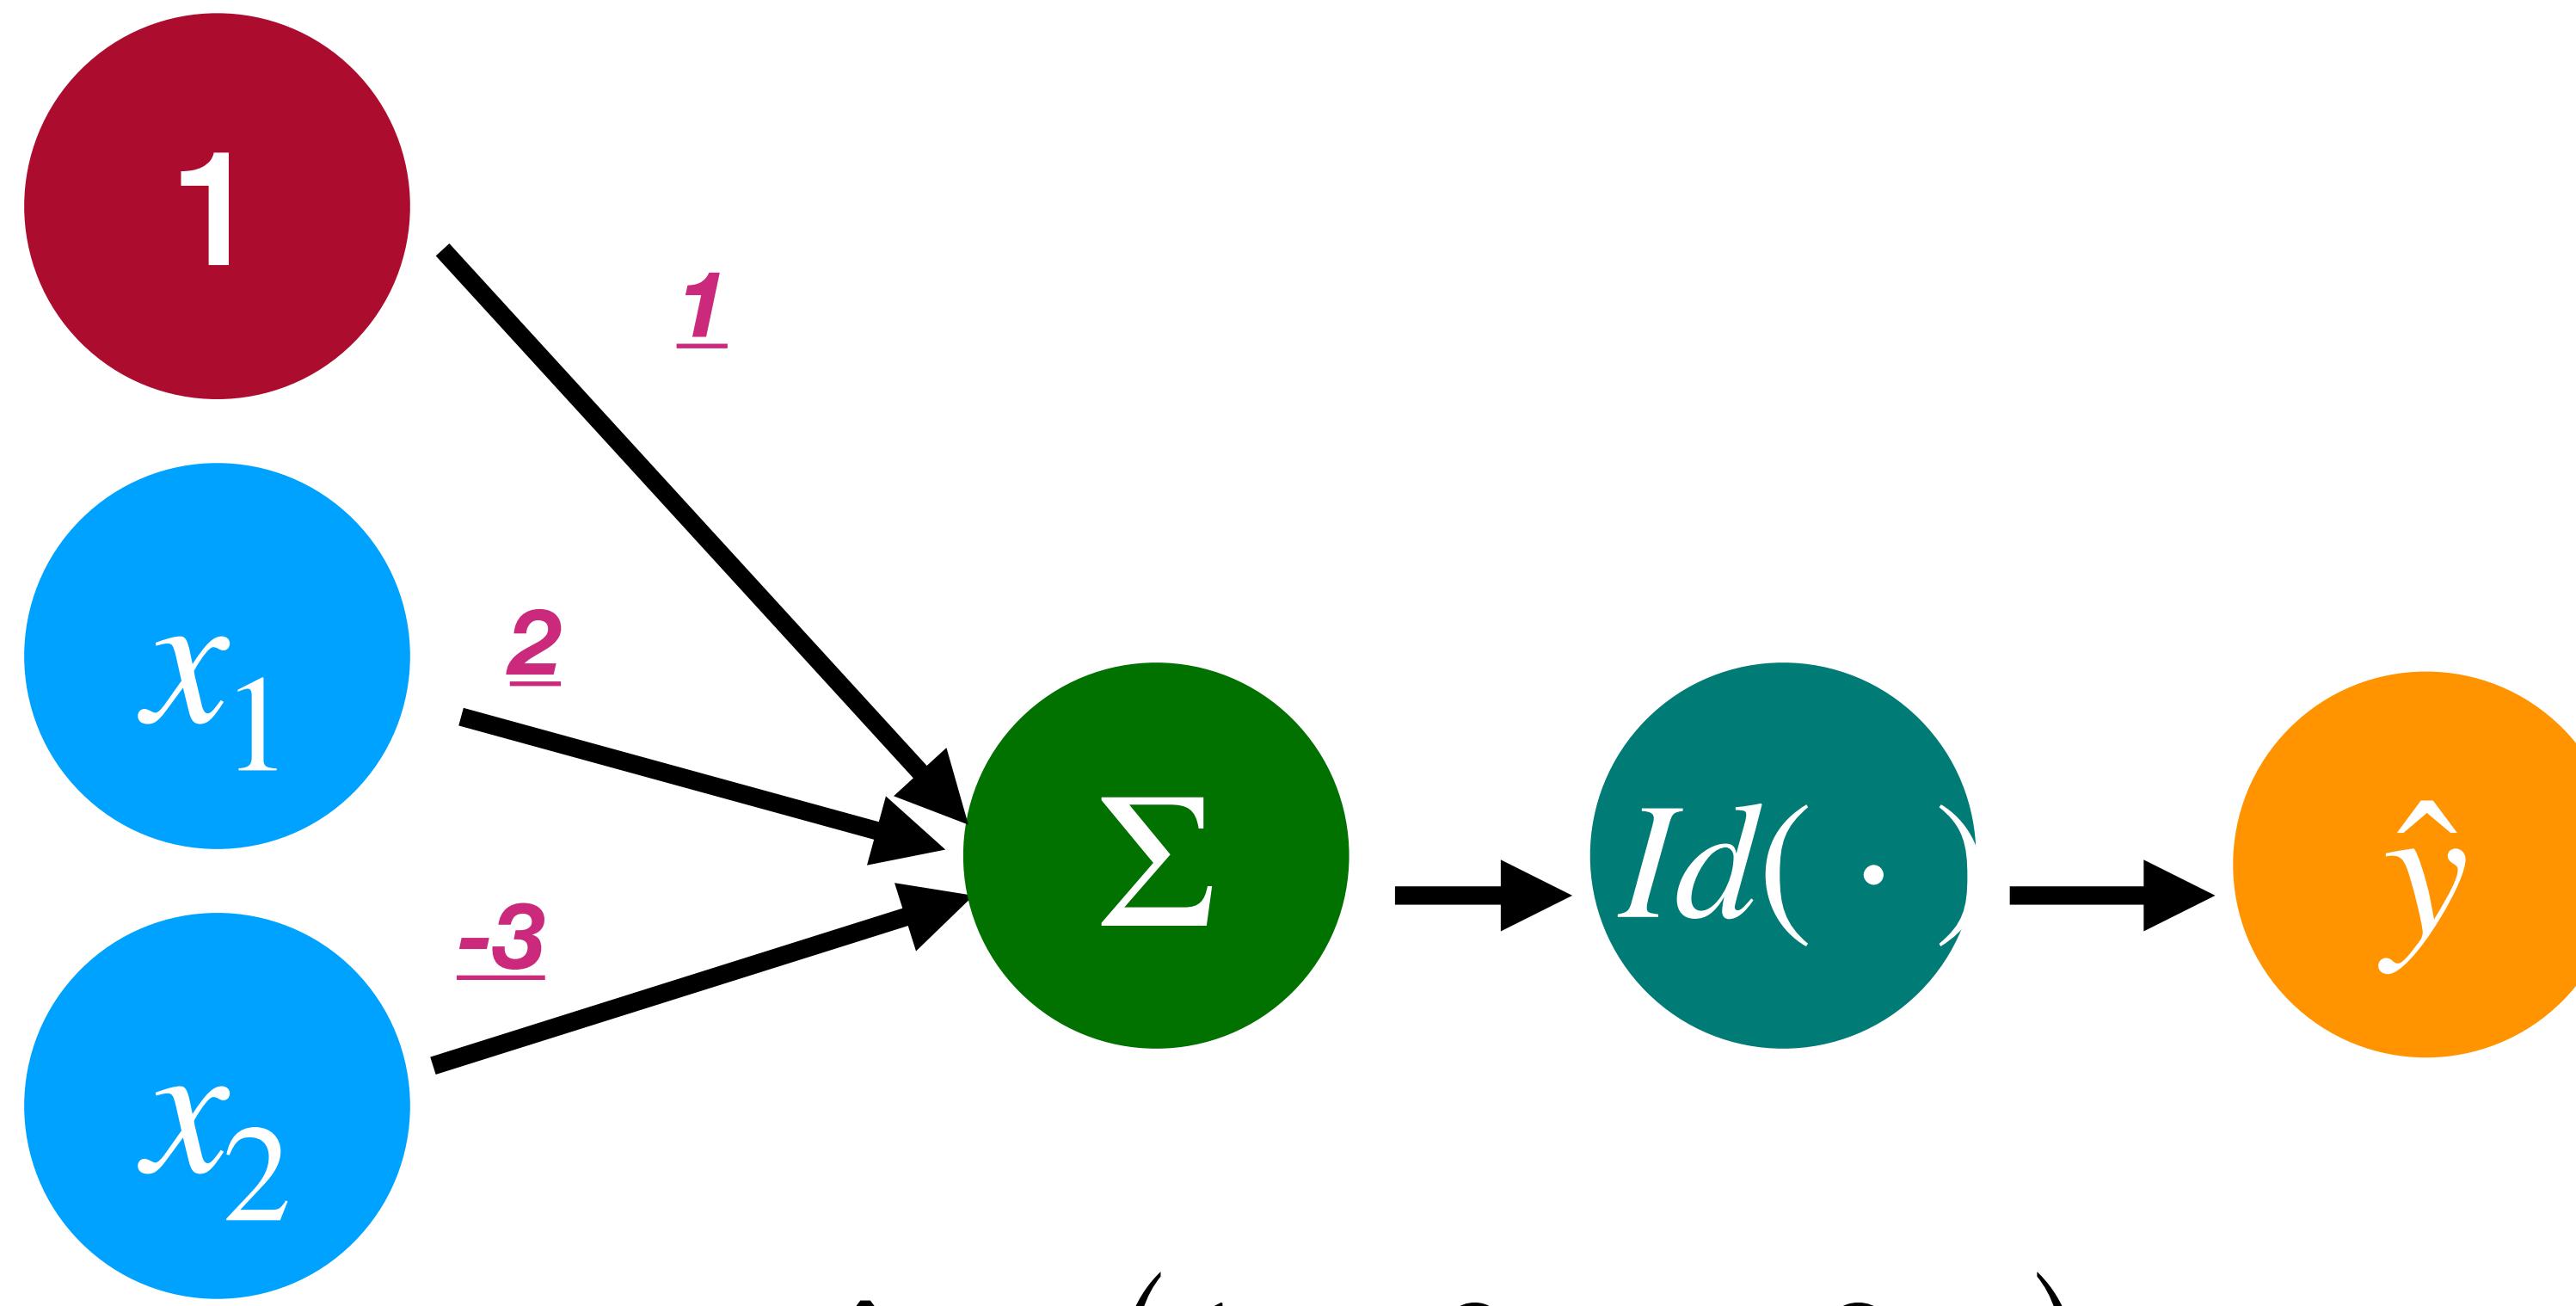
Our Perceptron to distinguish triangles and square



Our Perceptron to distinguish triangles and square

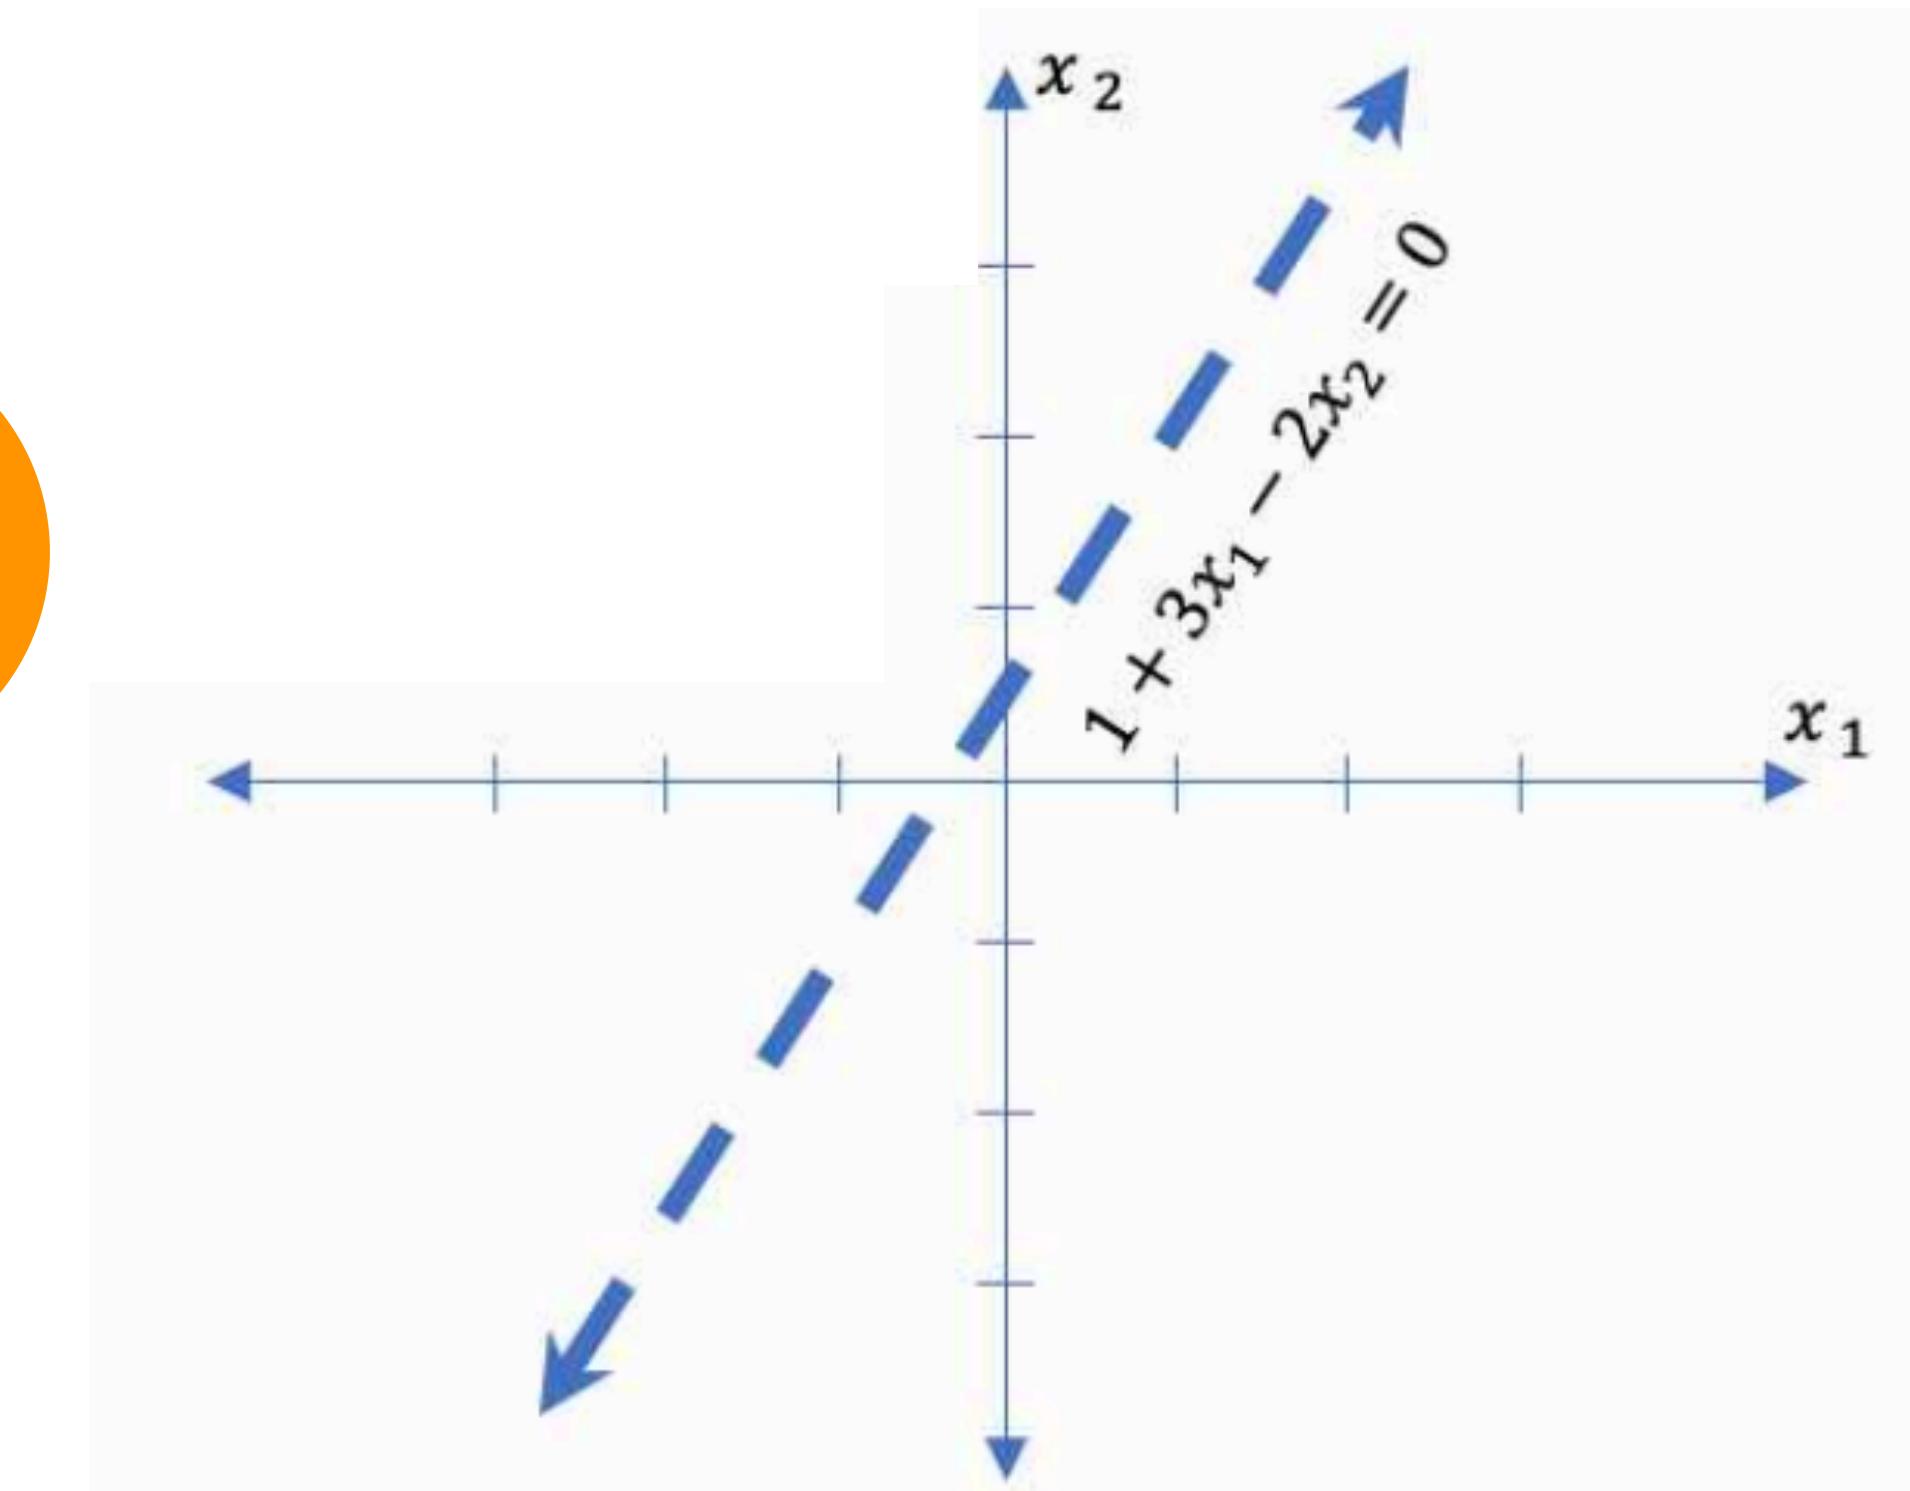
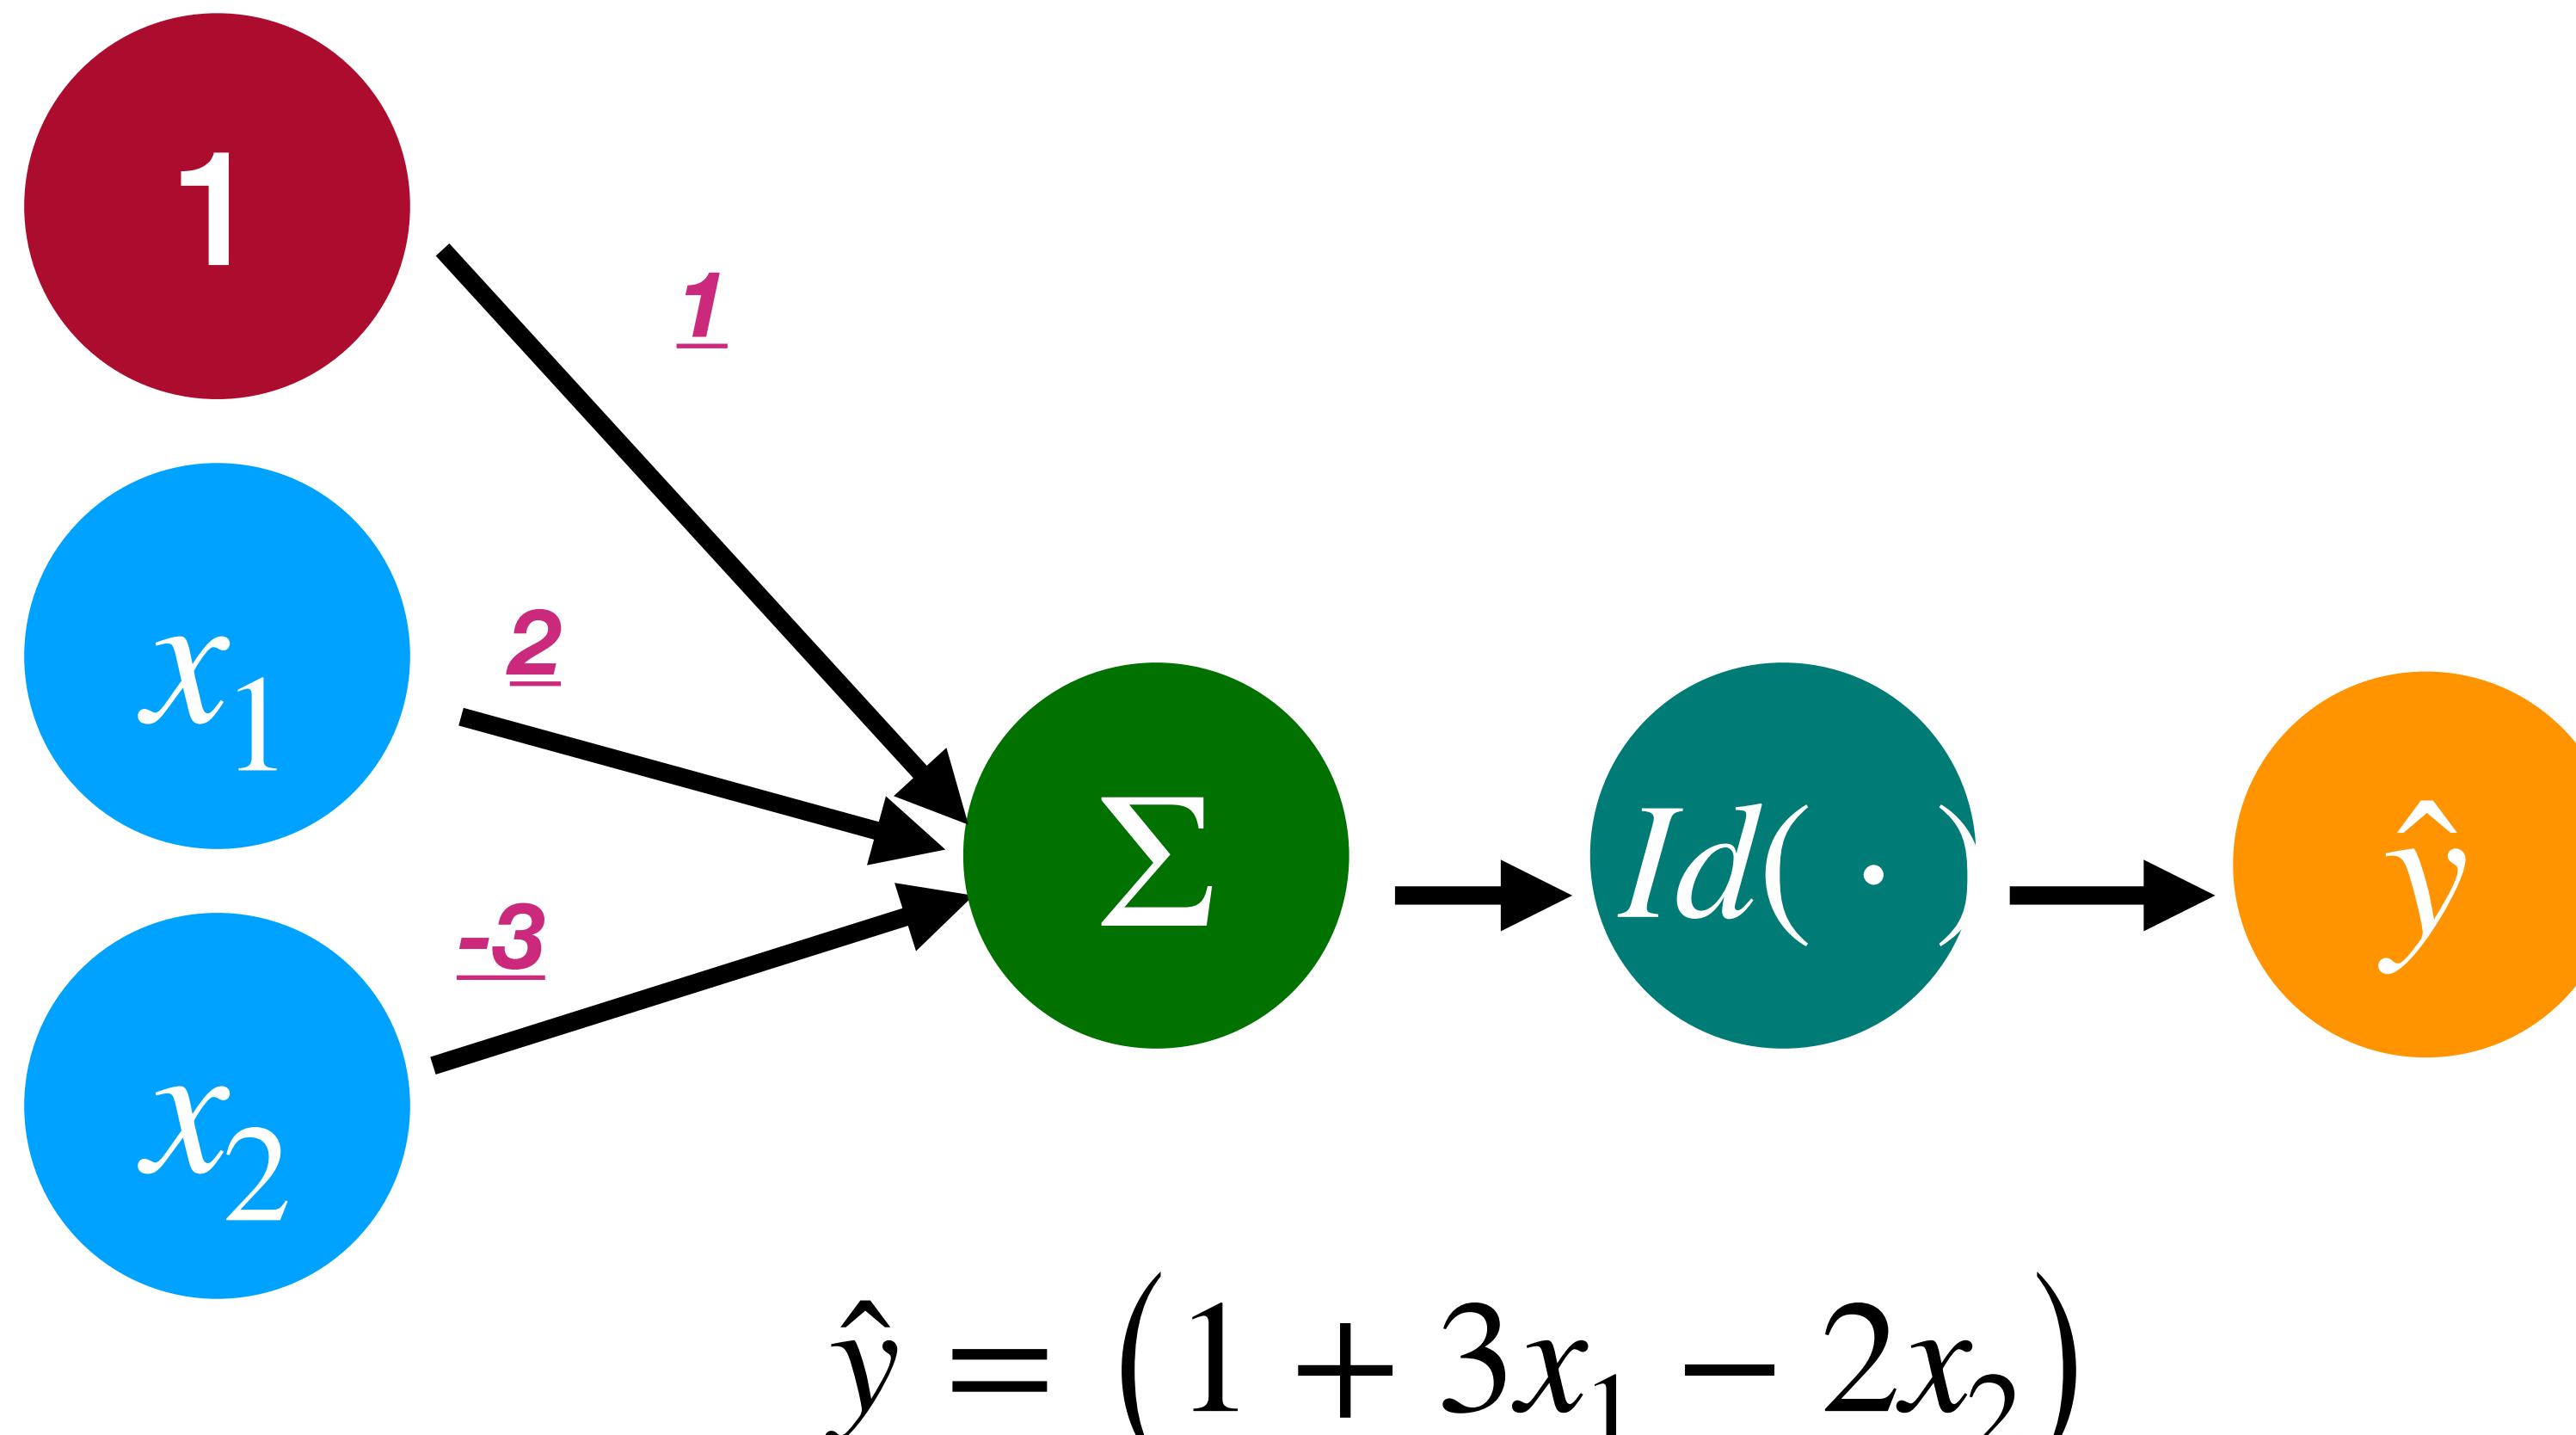


Our Perceptron to distinguish triangles and square

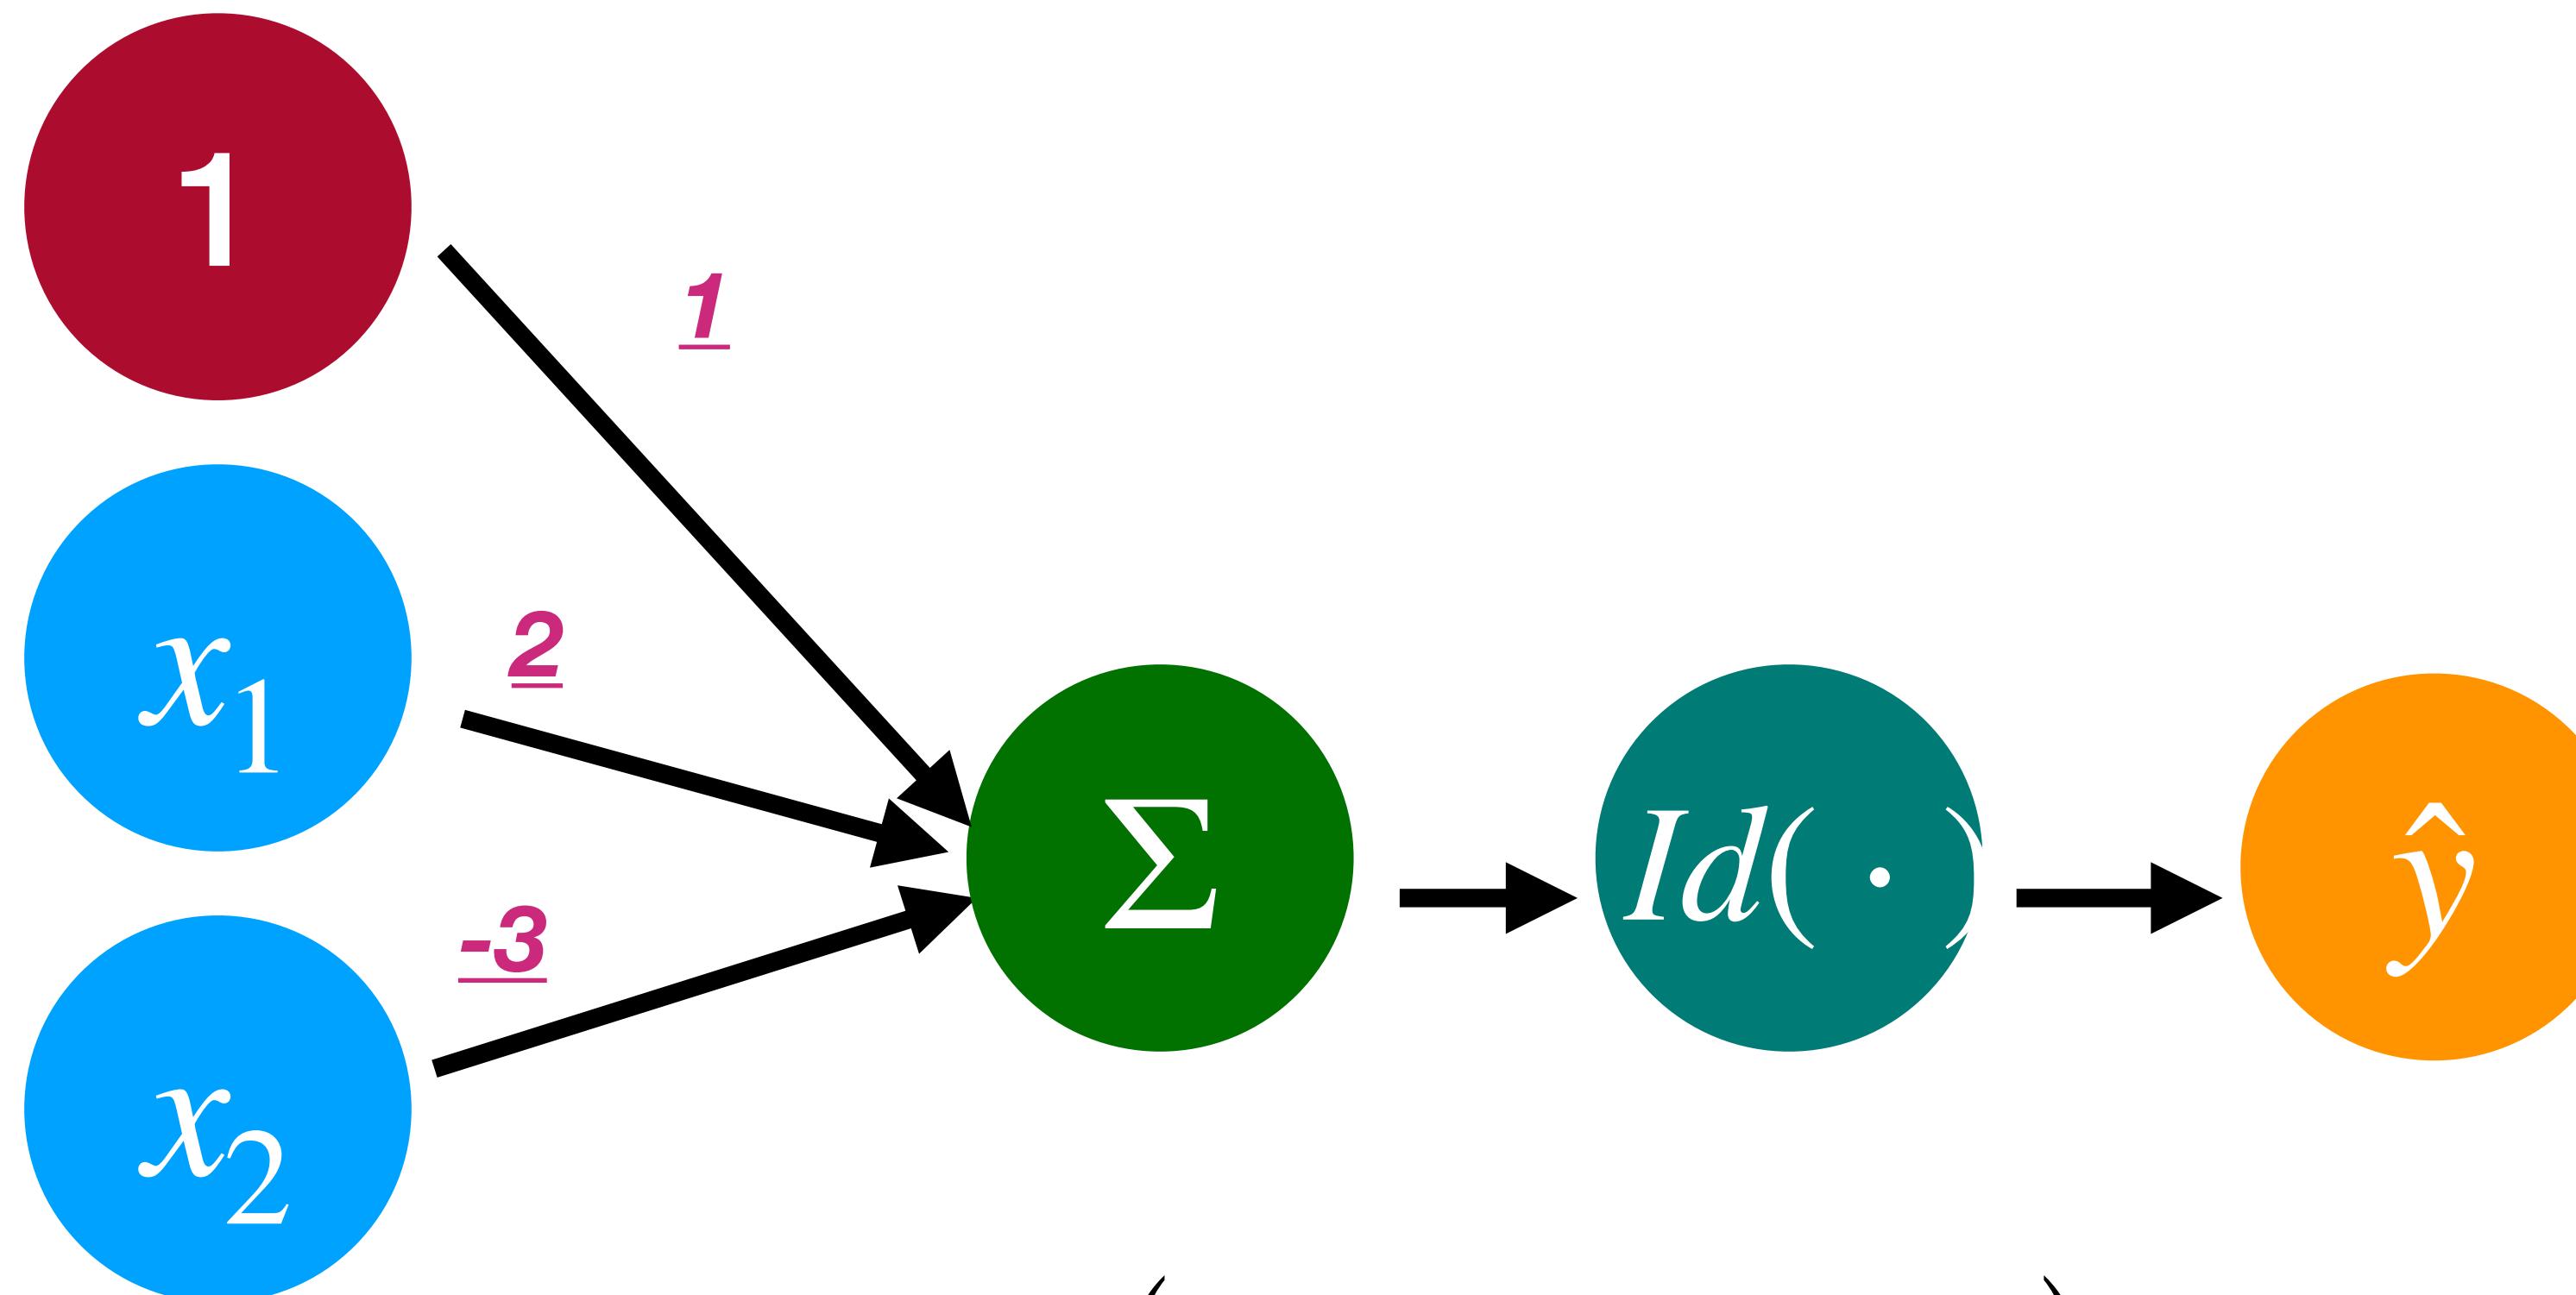


$$\hat{y} = (1 + 3x_1 - 2x_2)$$

Our Perceptron to distinguish triangles and square

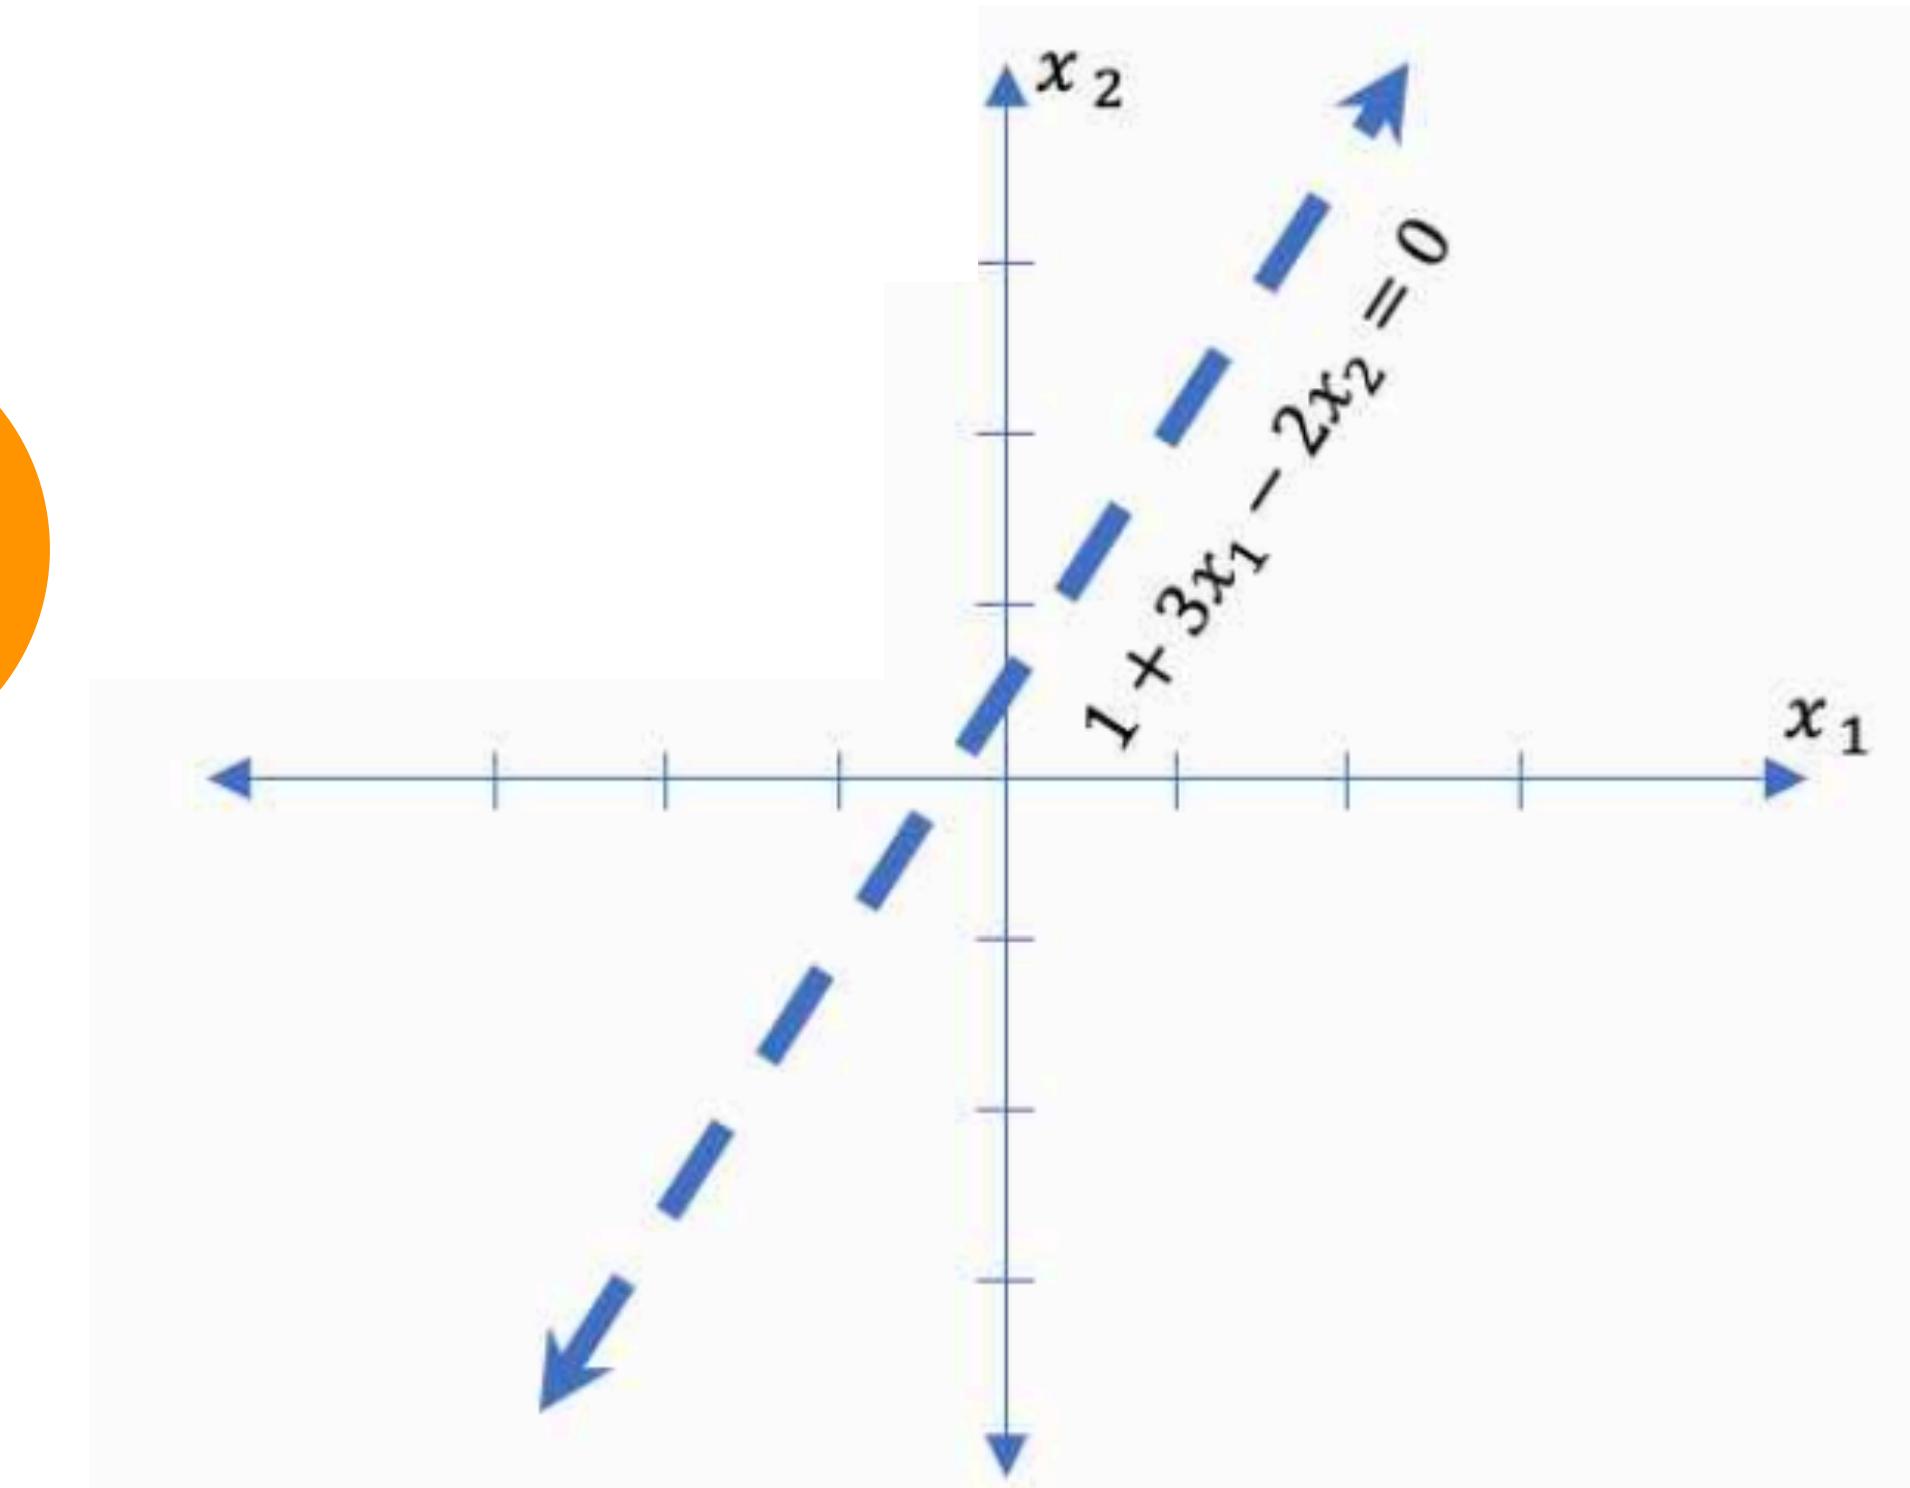


Our Perceptron to distinguish triangles and square

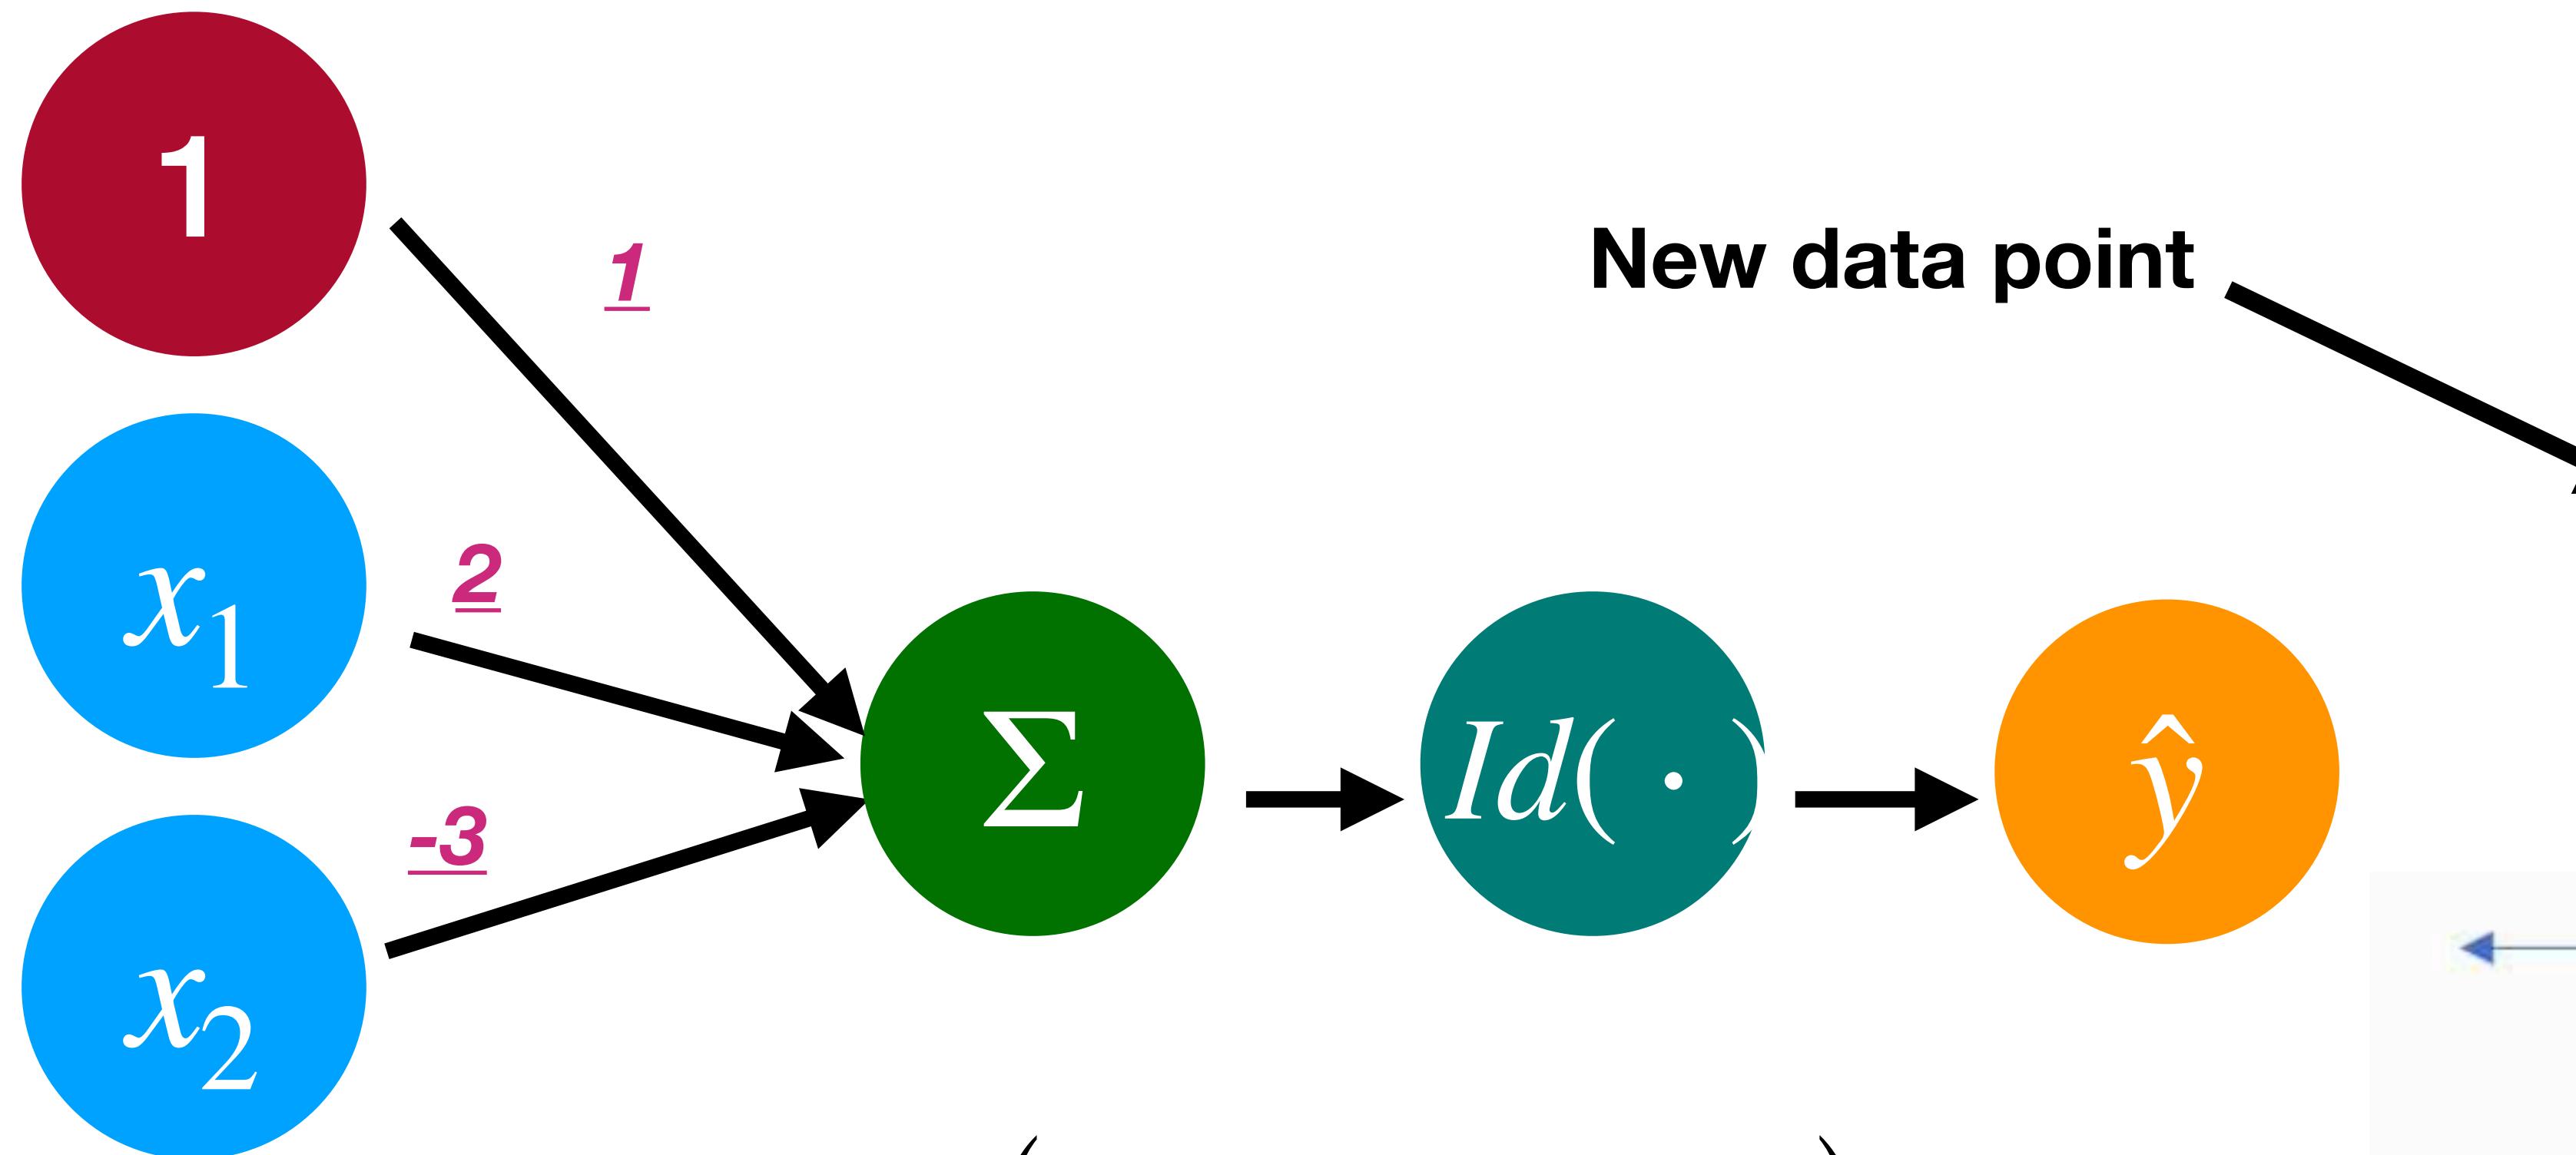


corresponds to the any scalar > 0 if the
Input is an triangle

corresponds to the any scalar < 0 if the
Input is an square

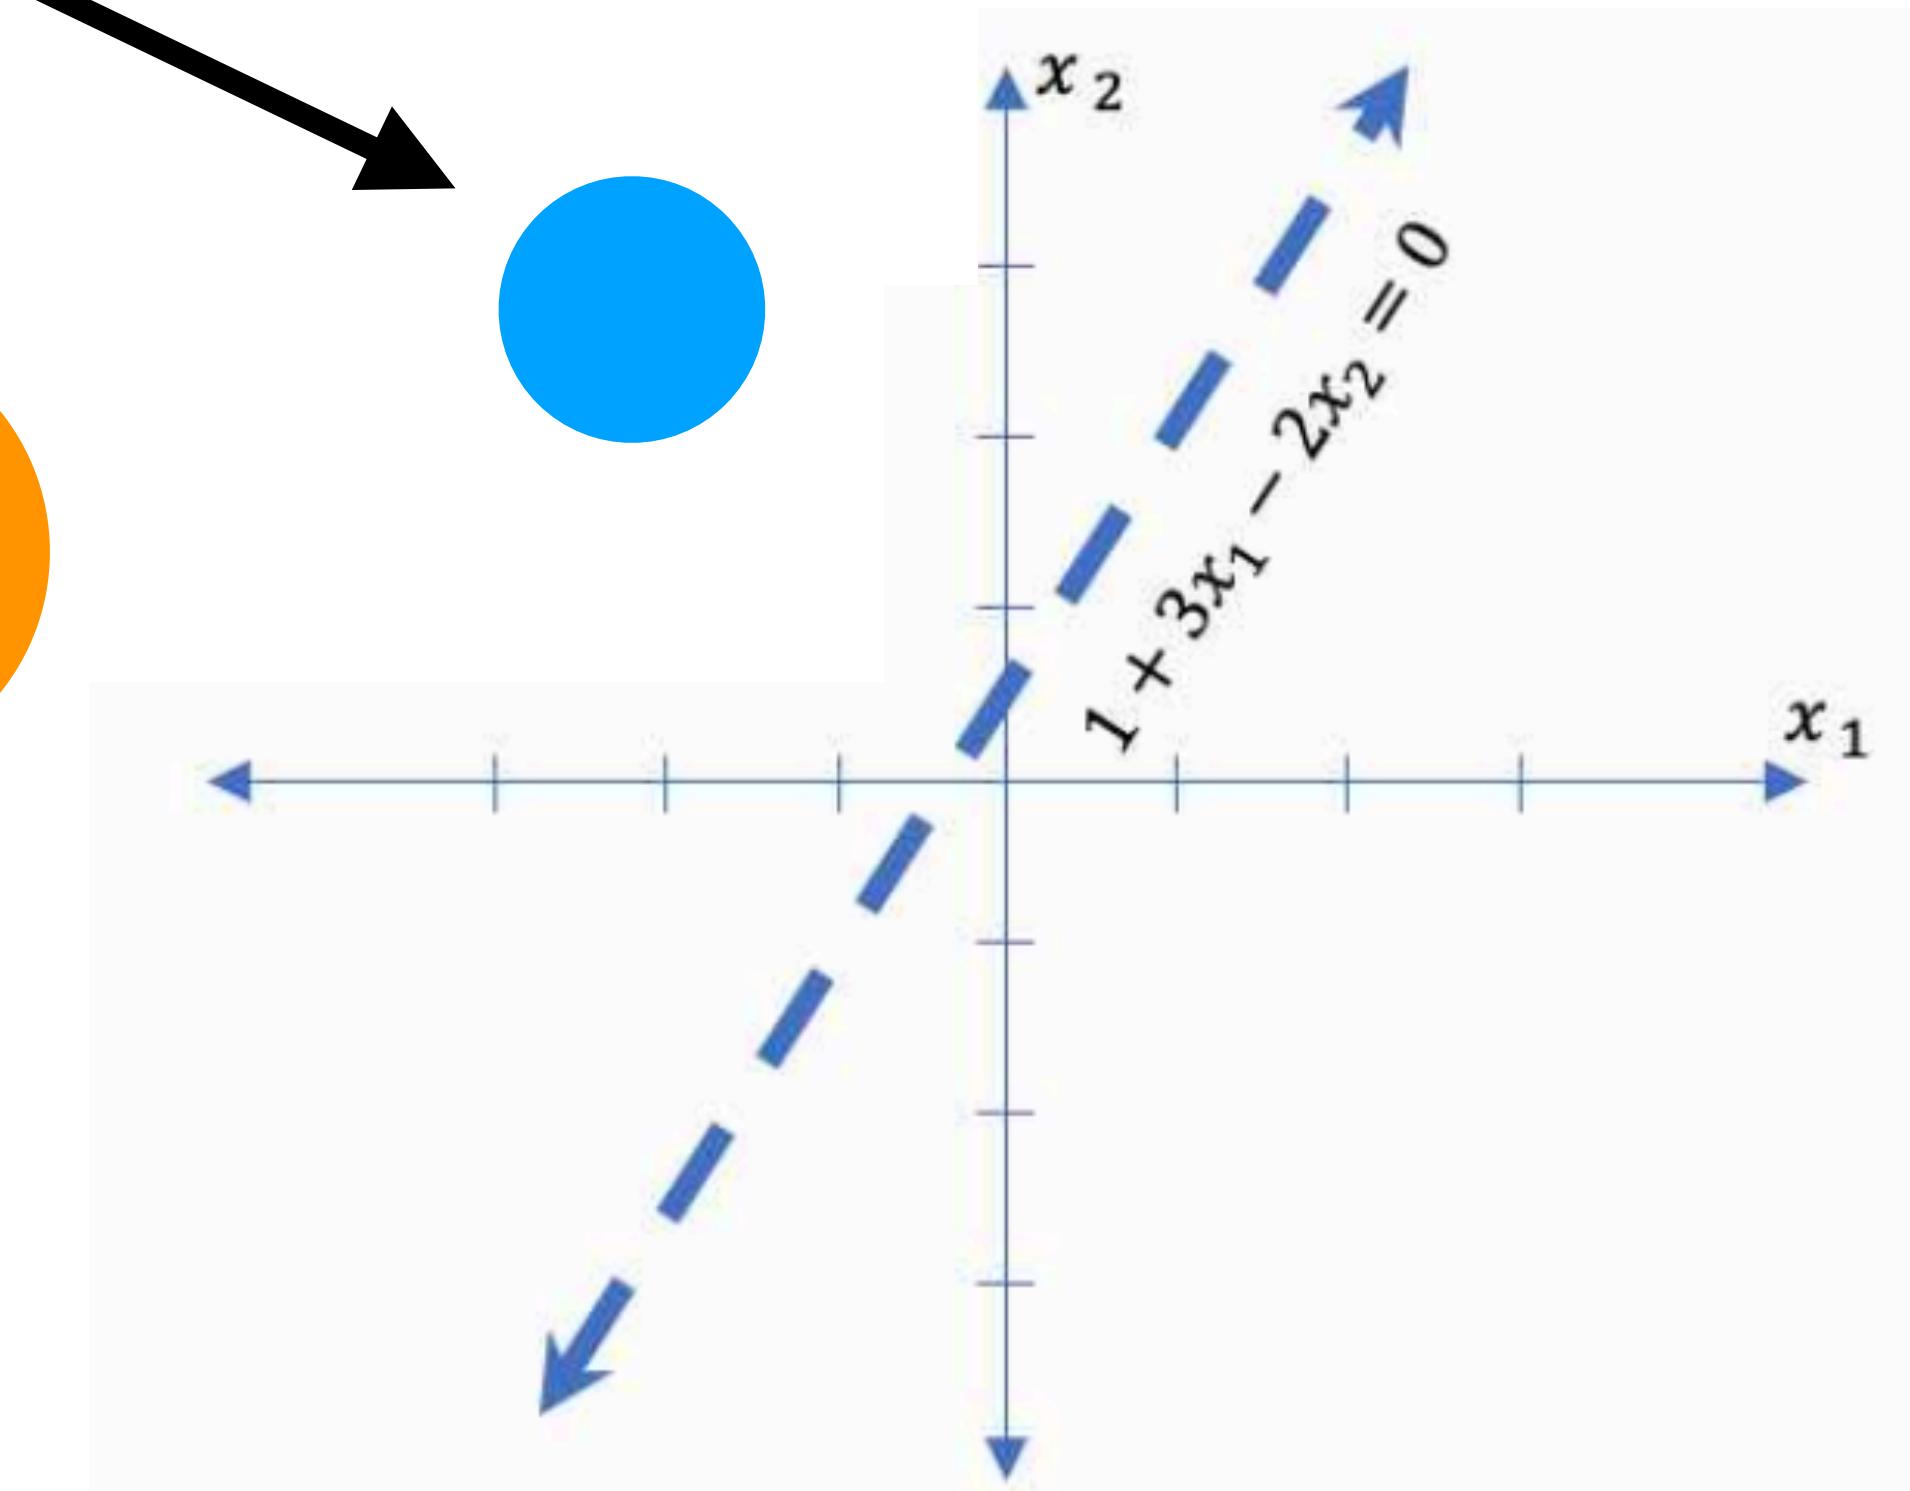


Our Perceptron to distinguish triangles and square

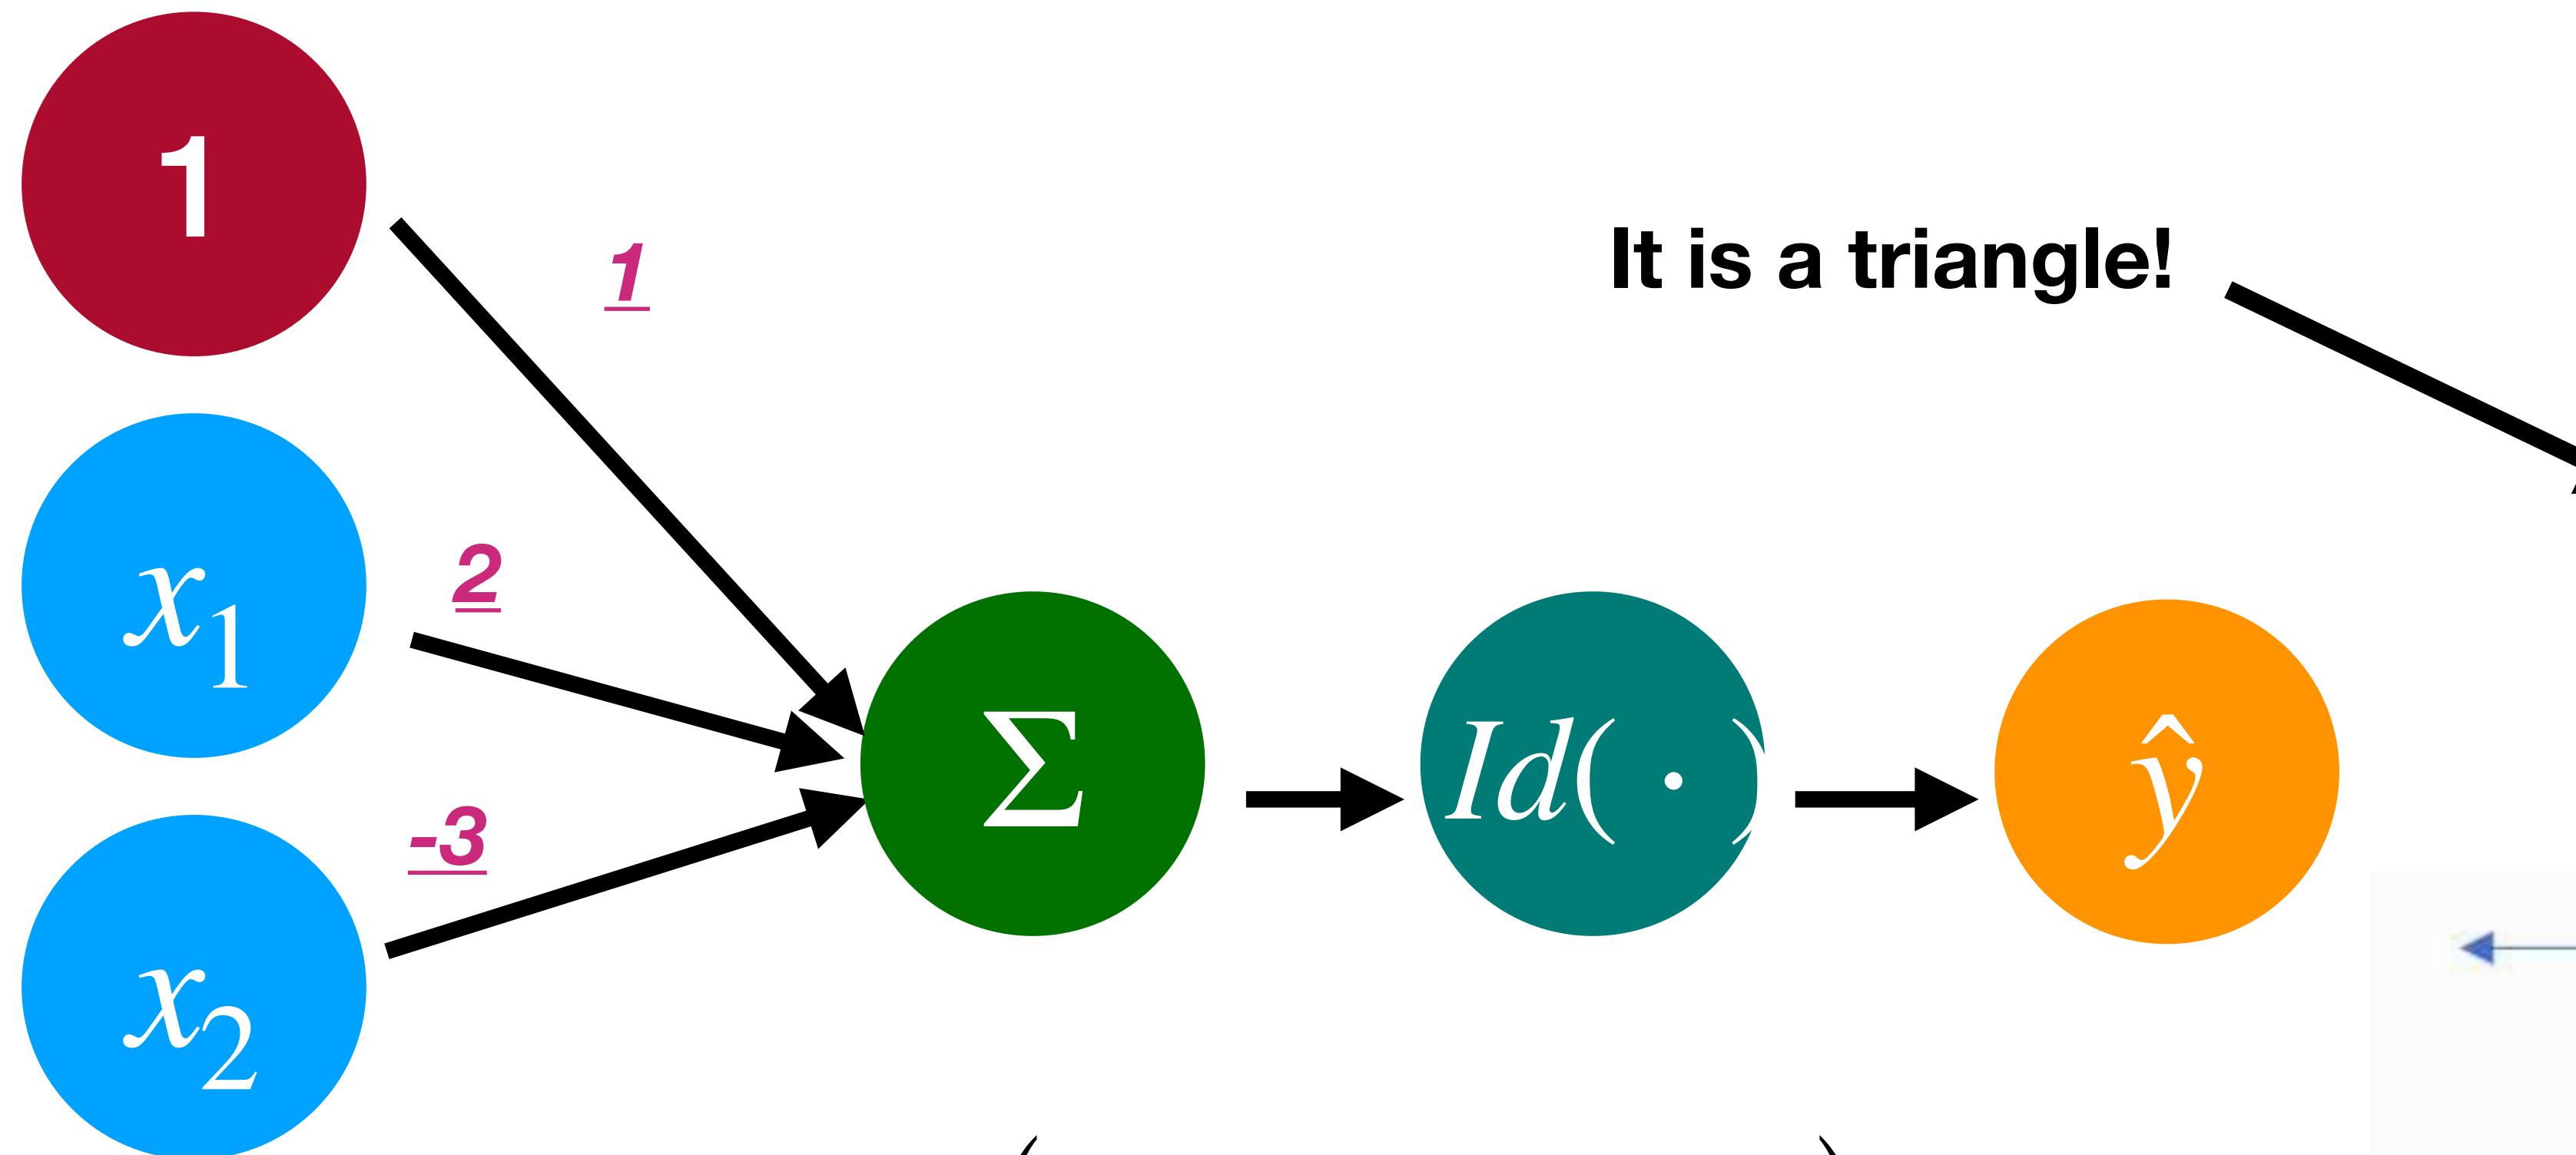


corresponds to the any scalar > 0 if the
Input is an triangle

corresponds to the any scalar < 0 if the
Input is an square

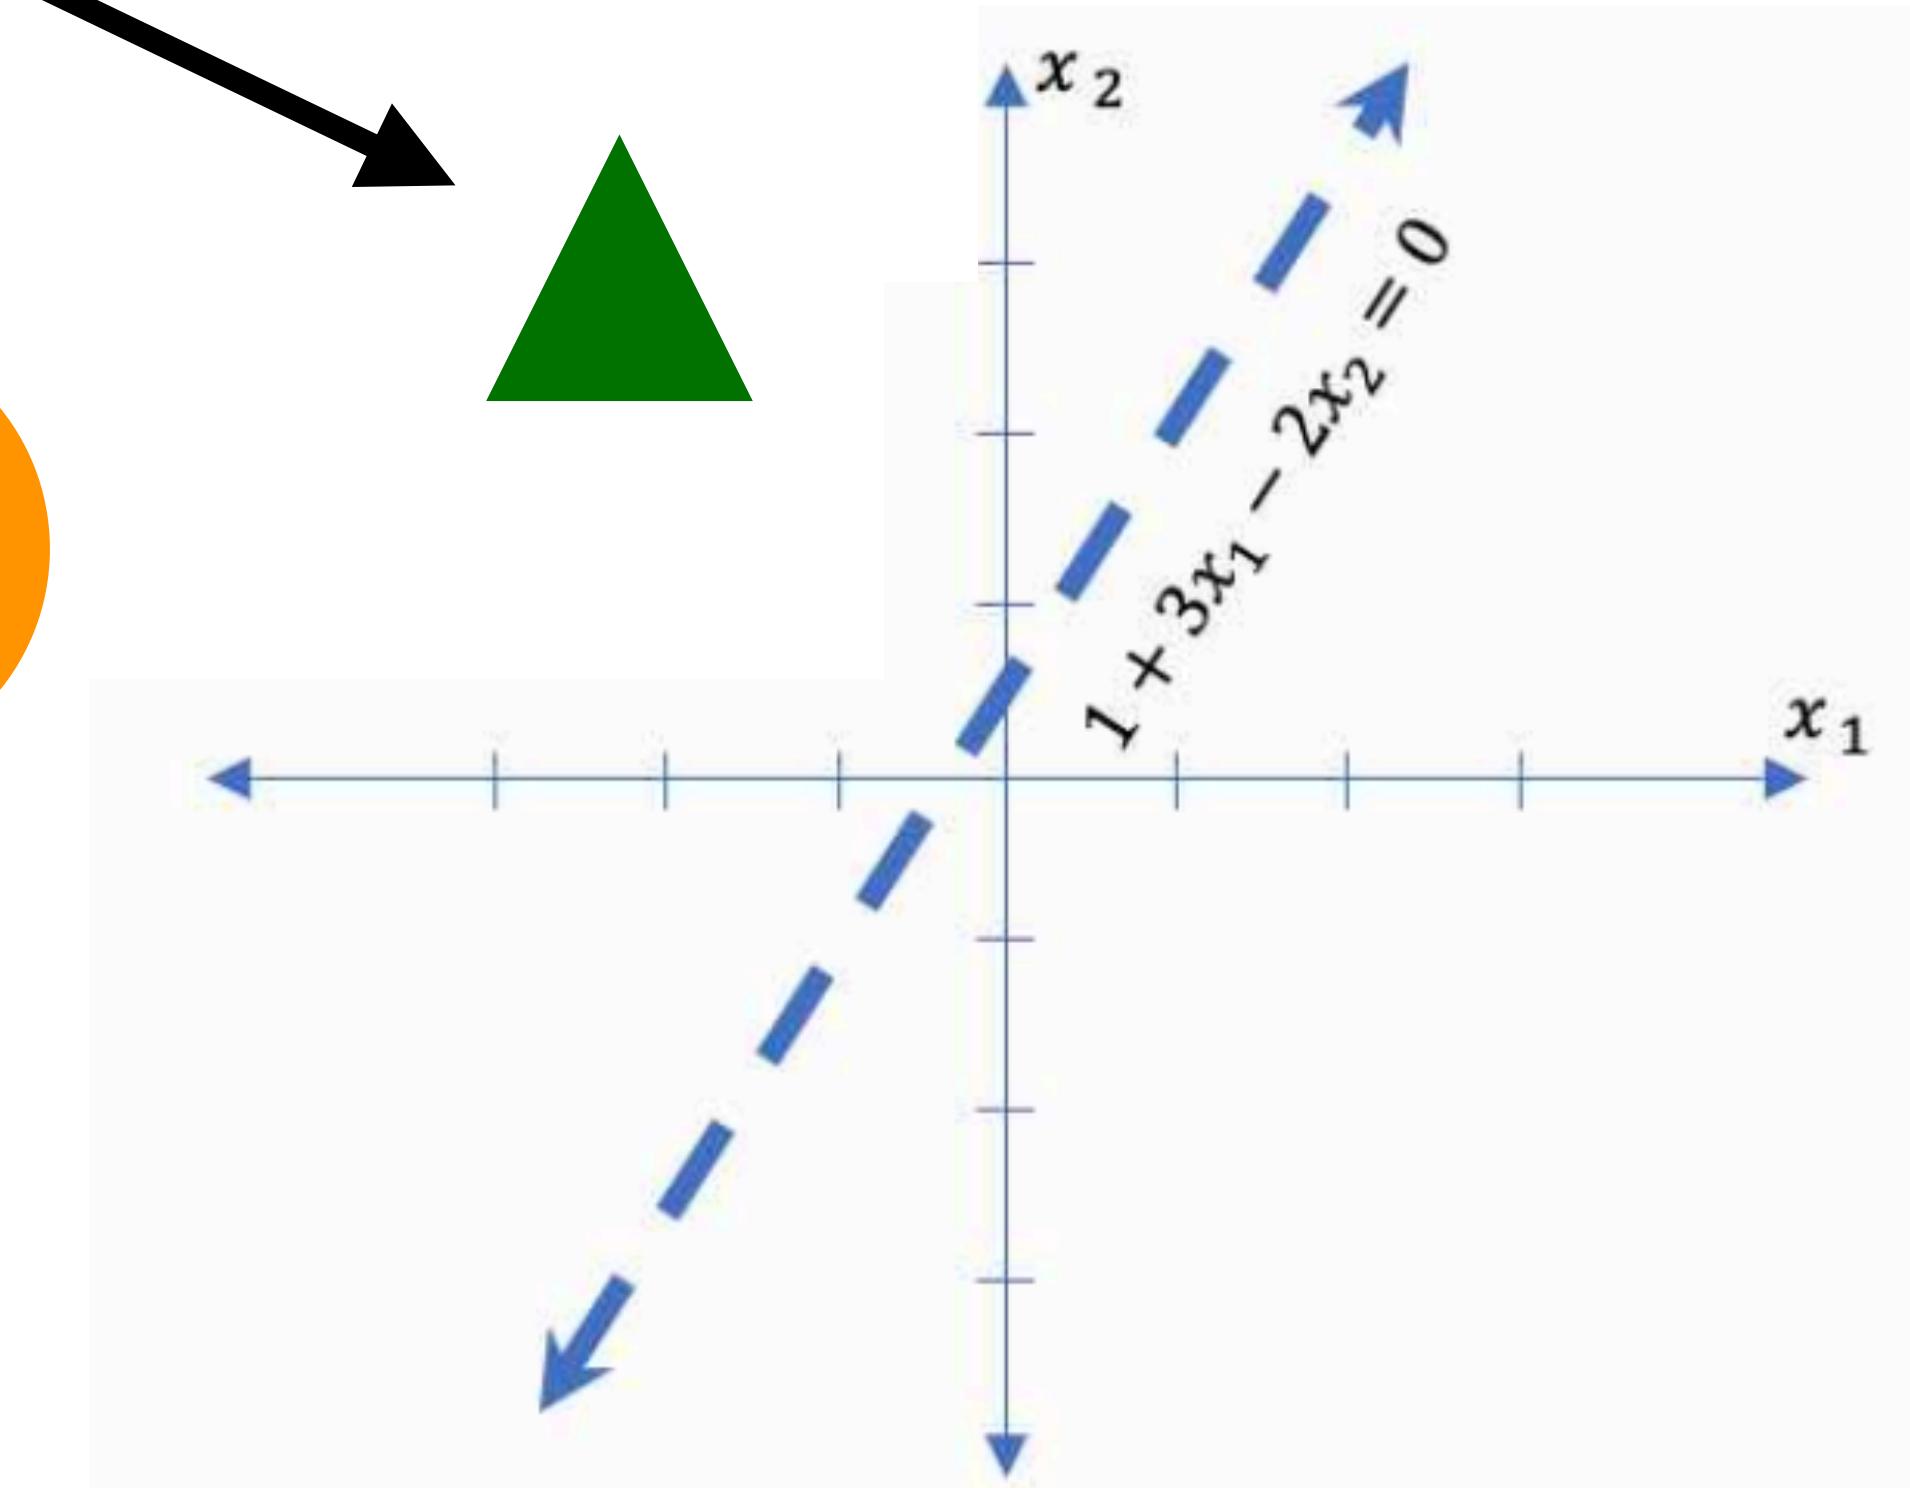


Our Perceptron to distinguish triangles and square

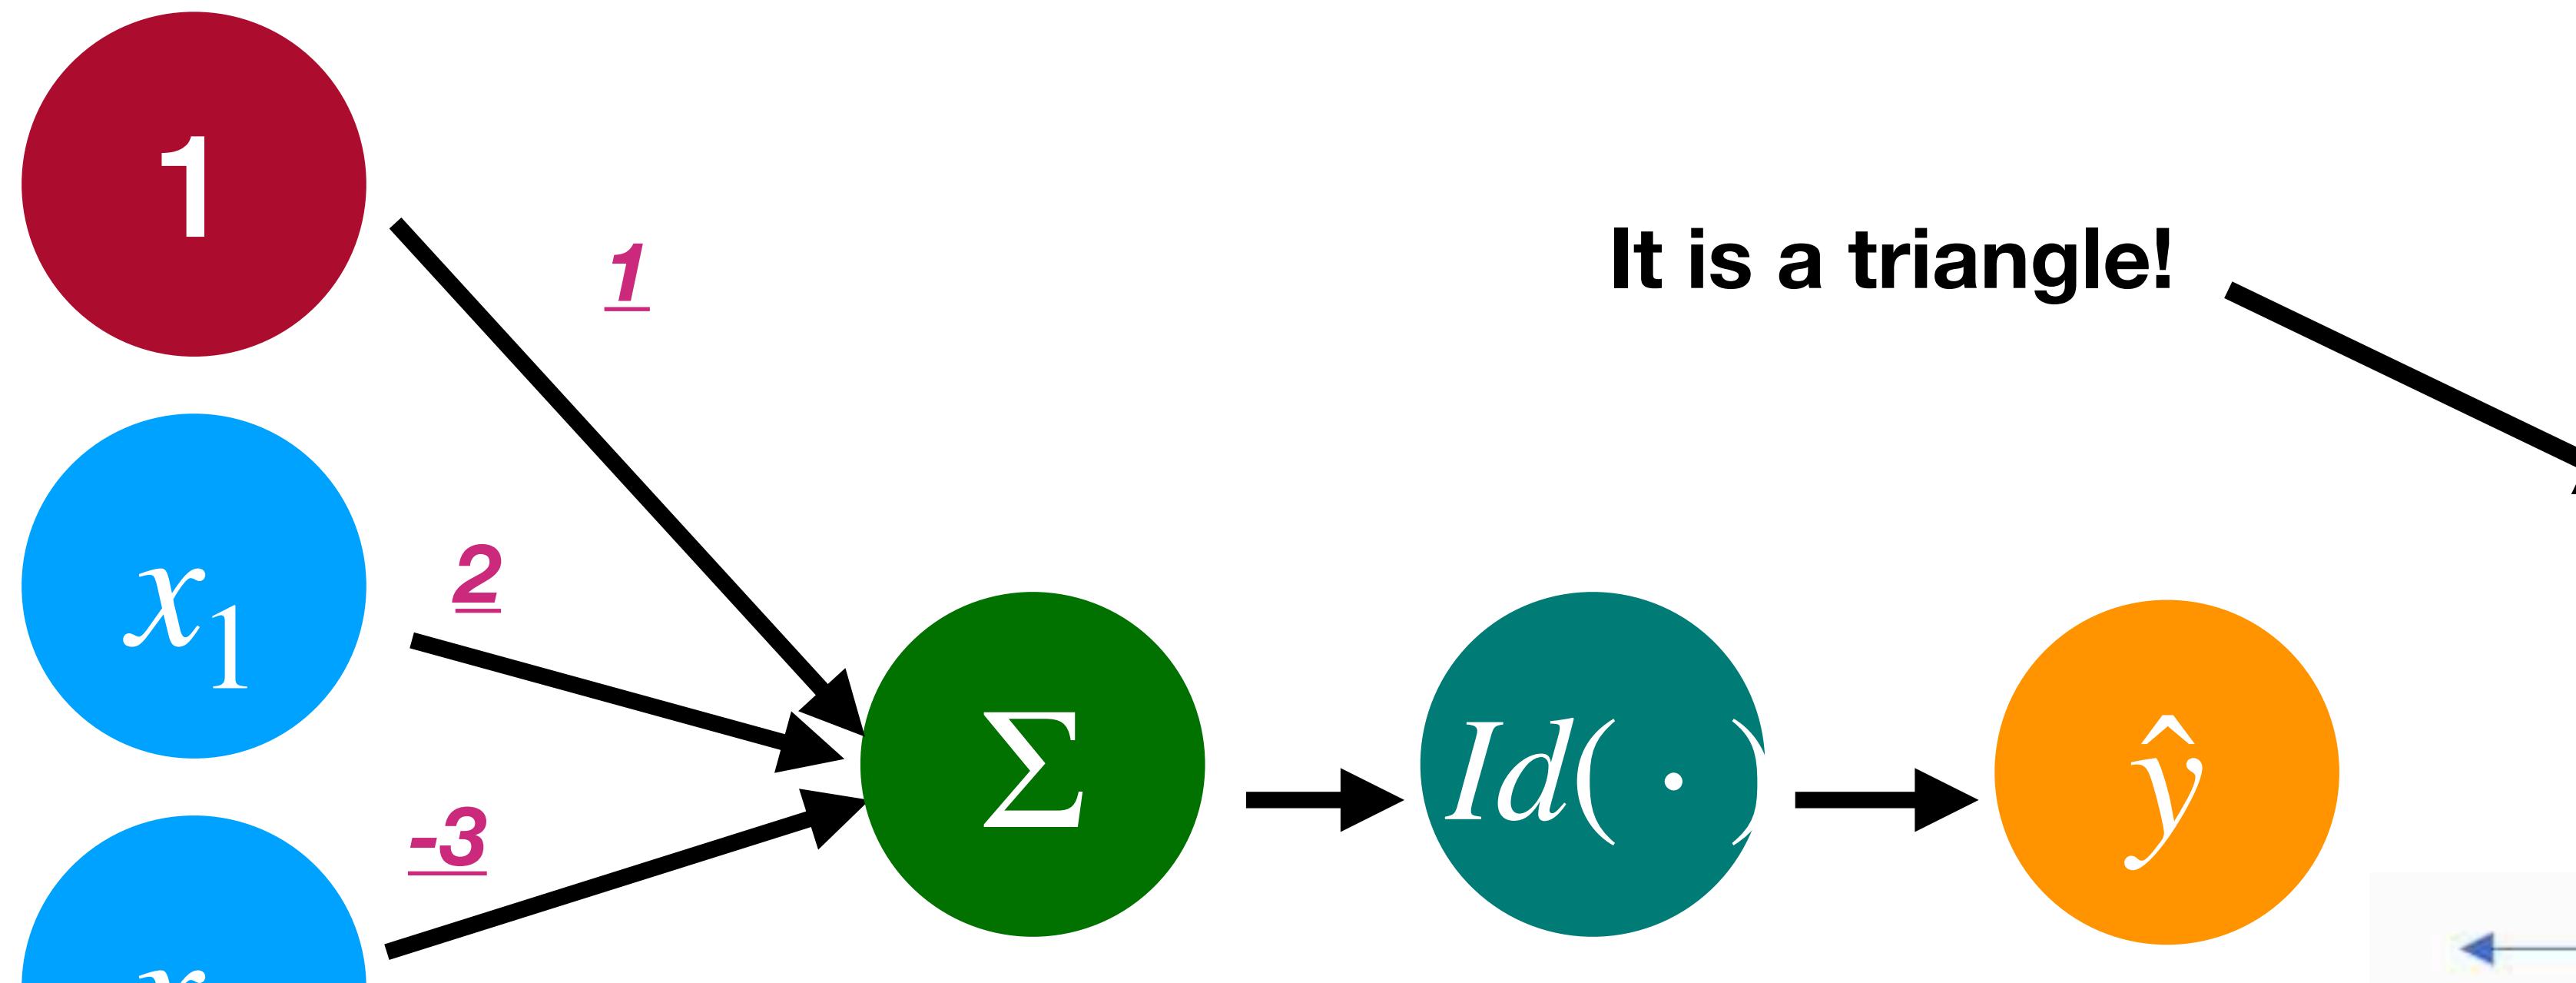


corresponds to the any scalar > 0 if the
Input is an triangle

corresponds to the any scalar < 0 if the
Input is an square



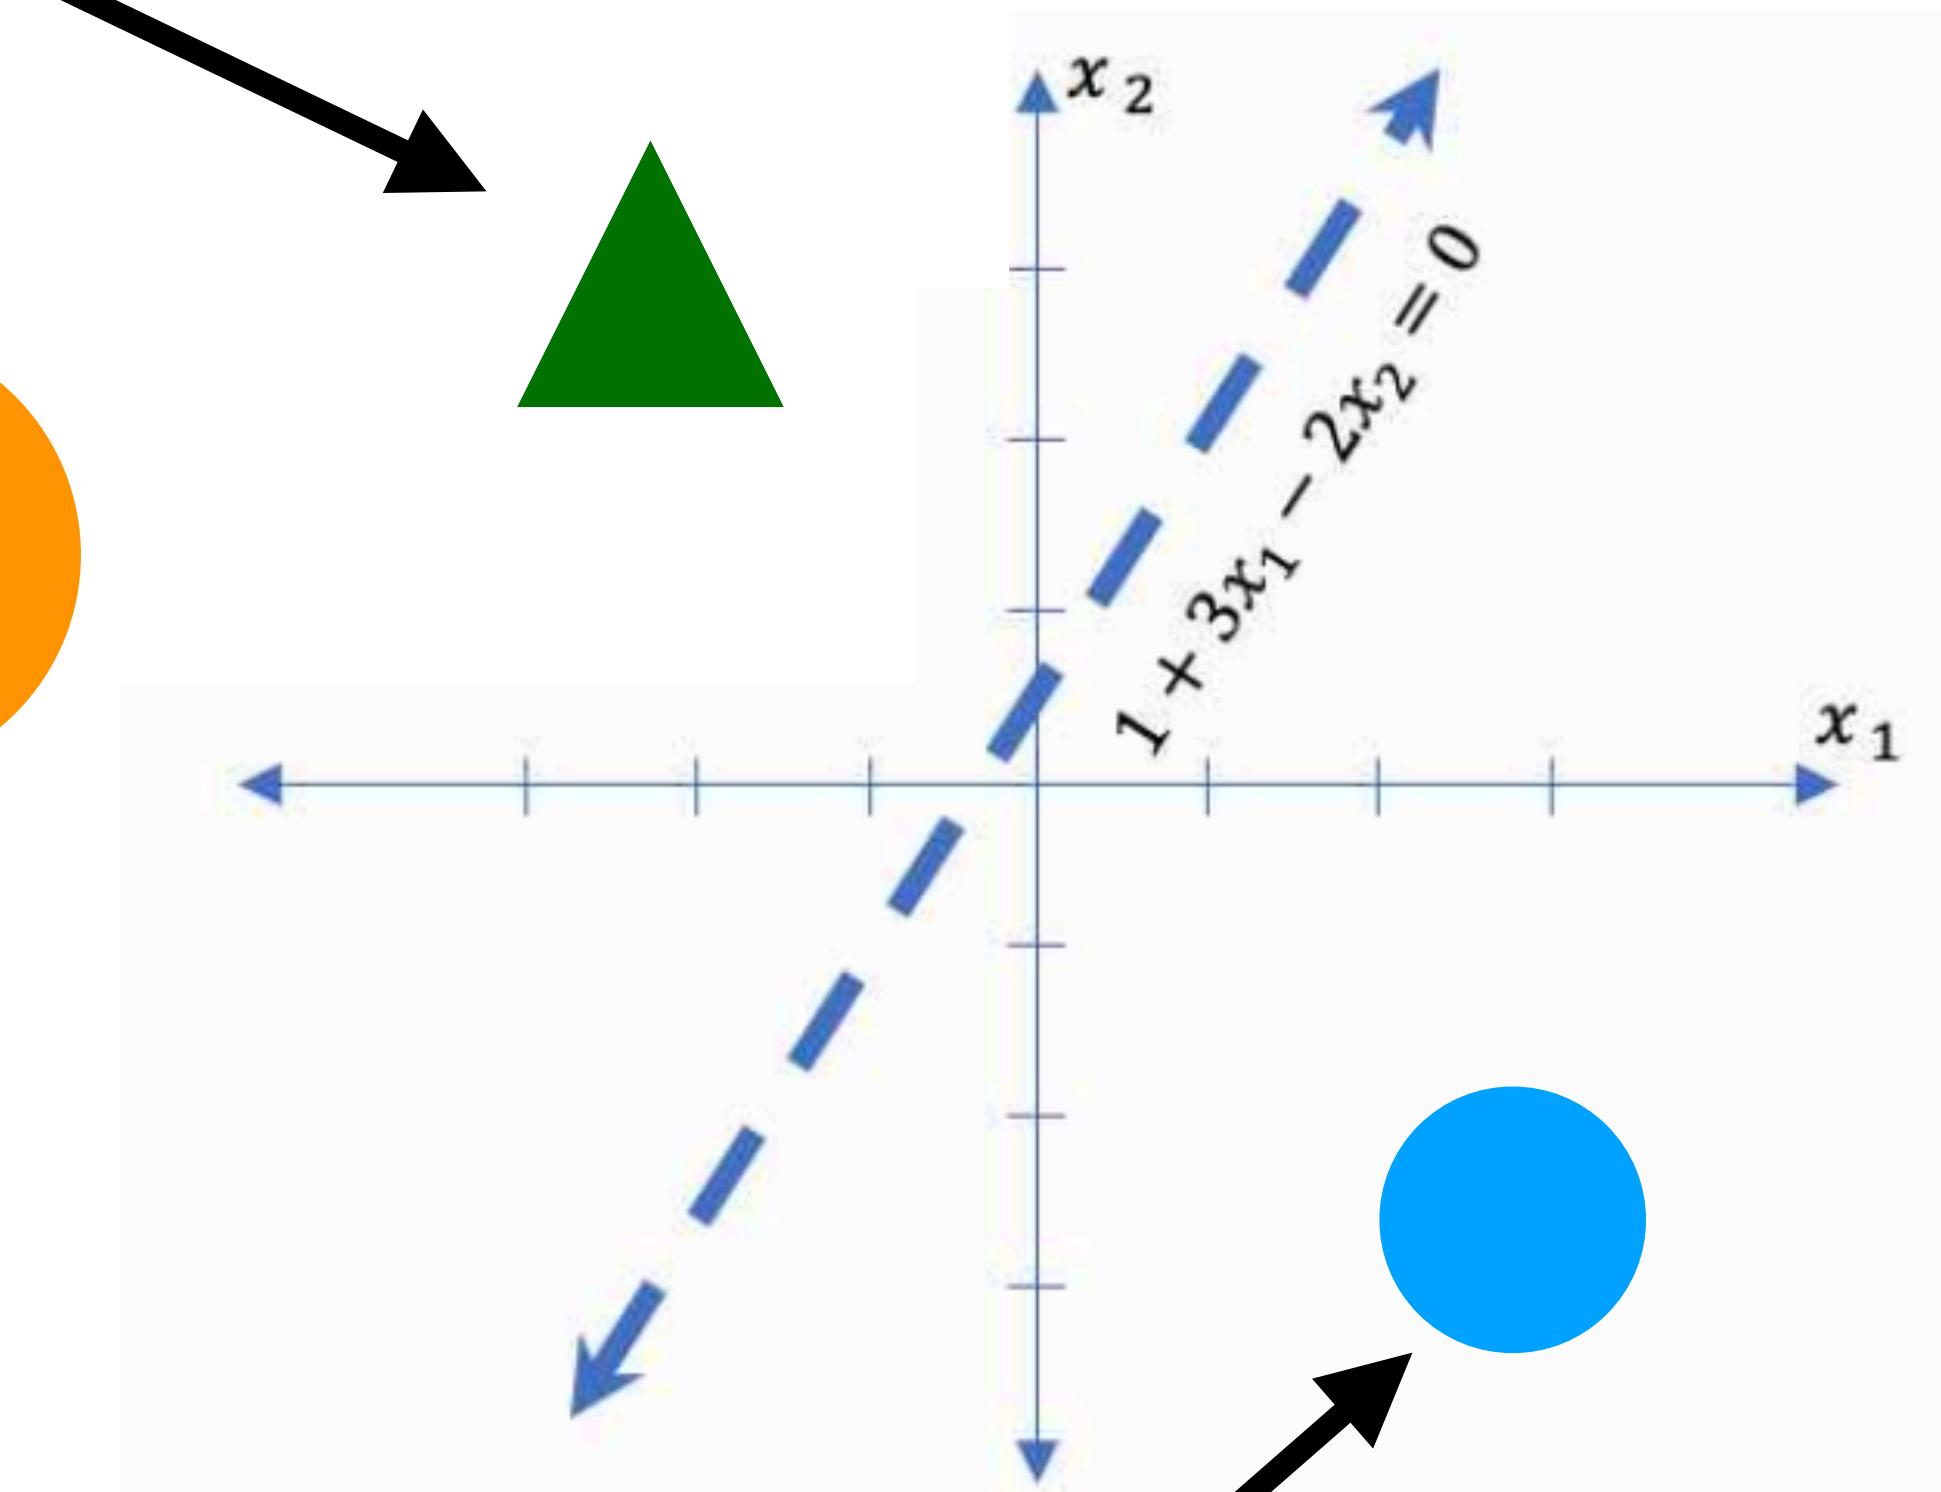
Our Perceptron to distinguish triangles and square



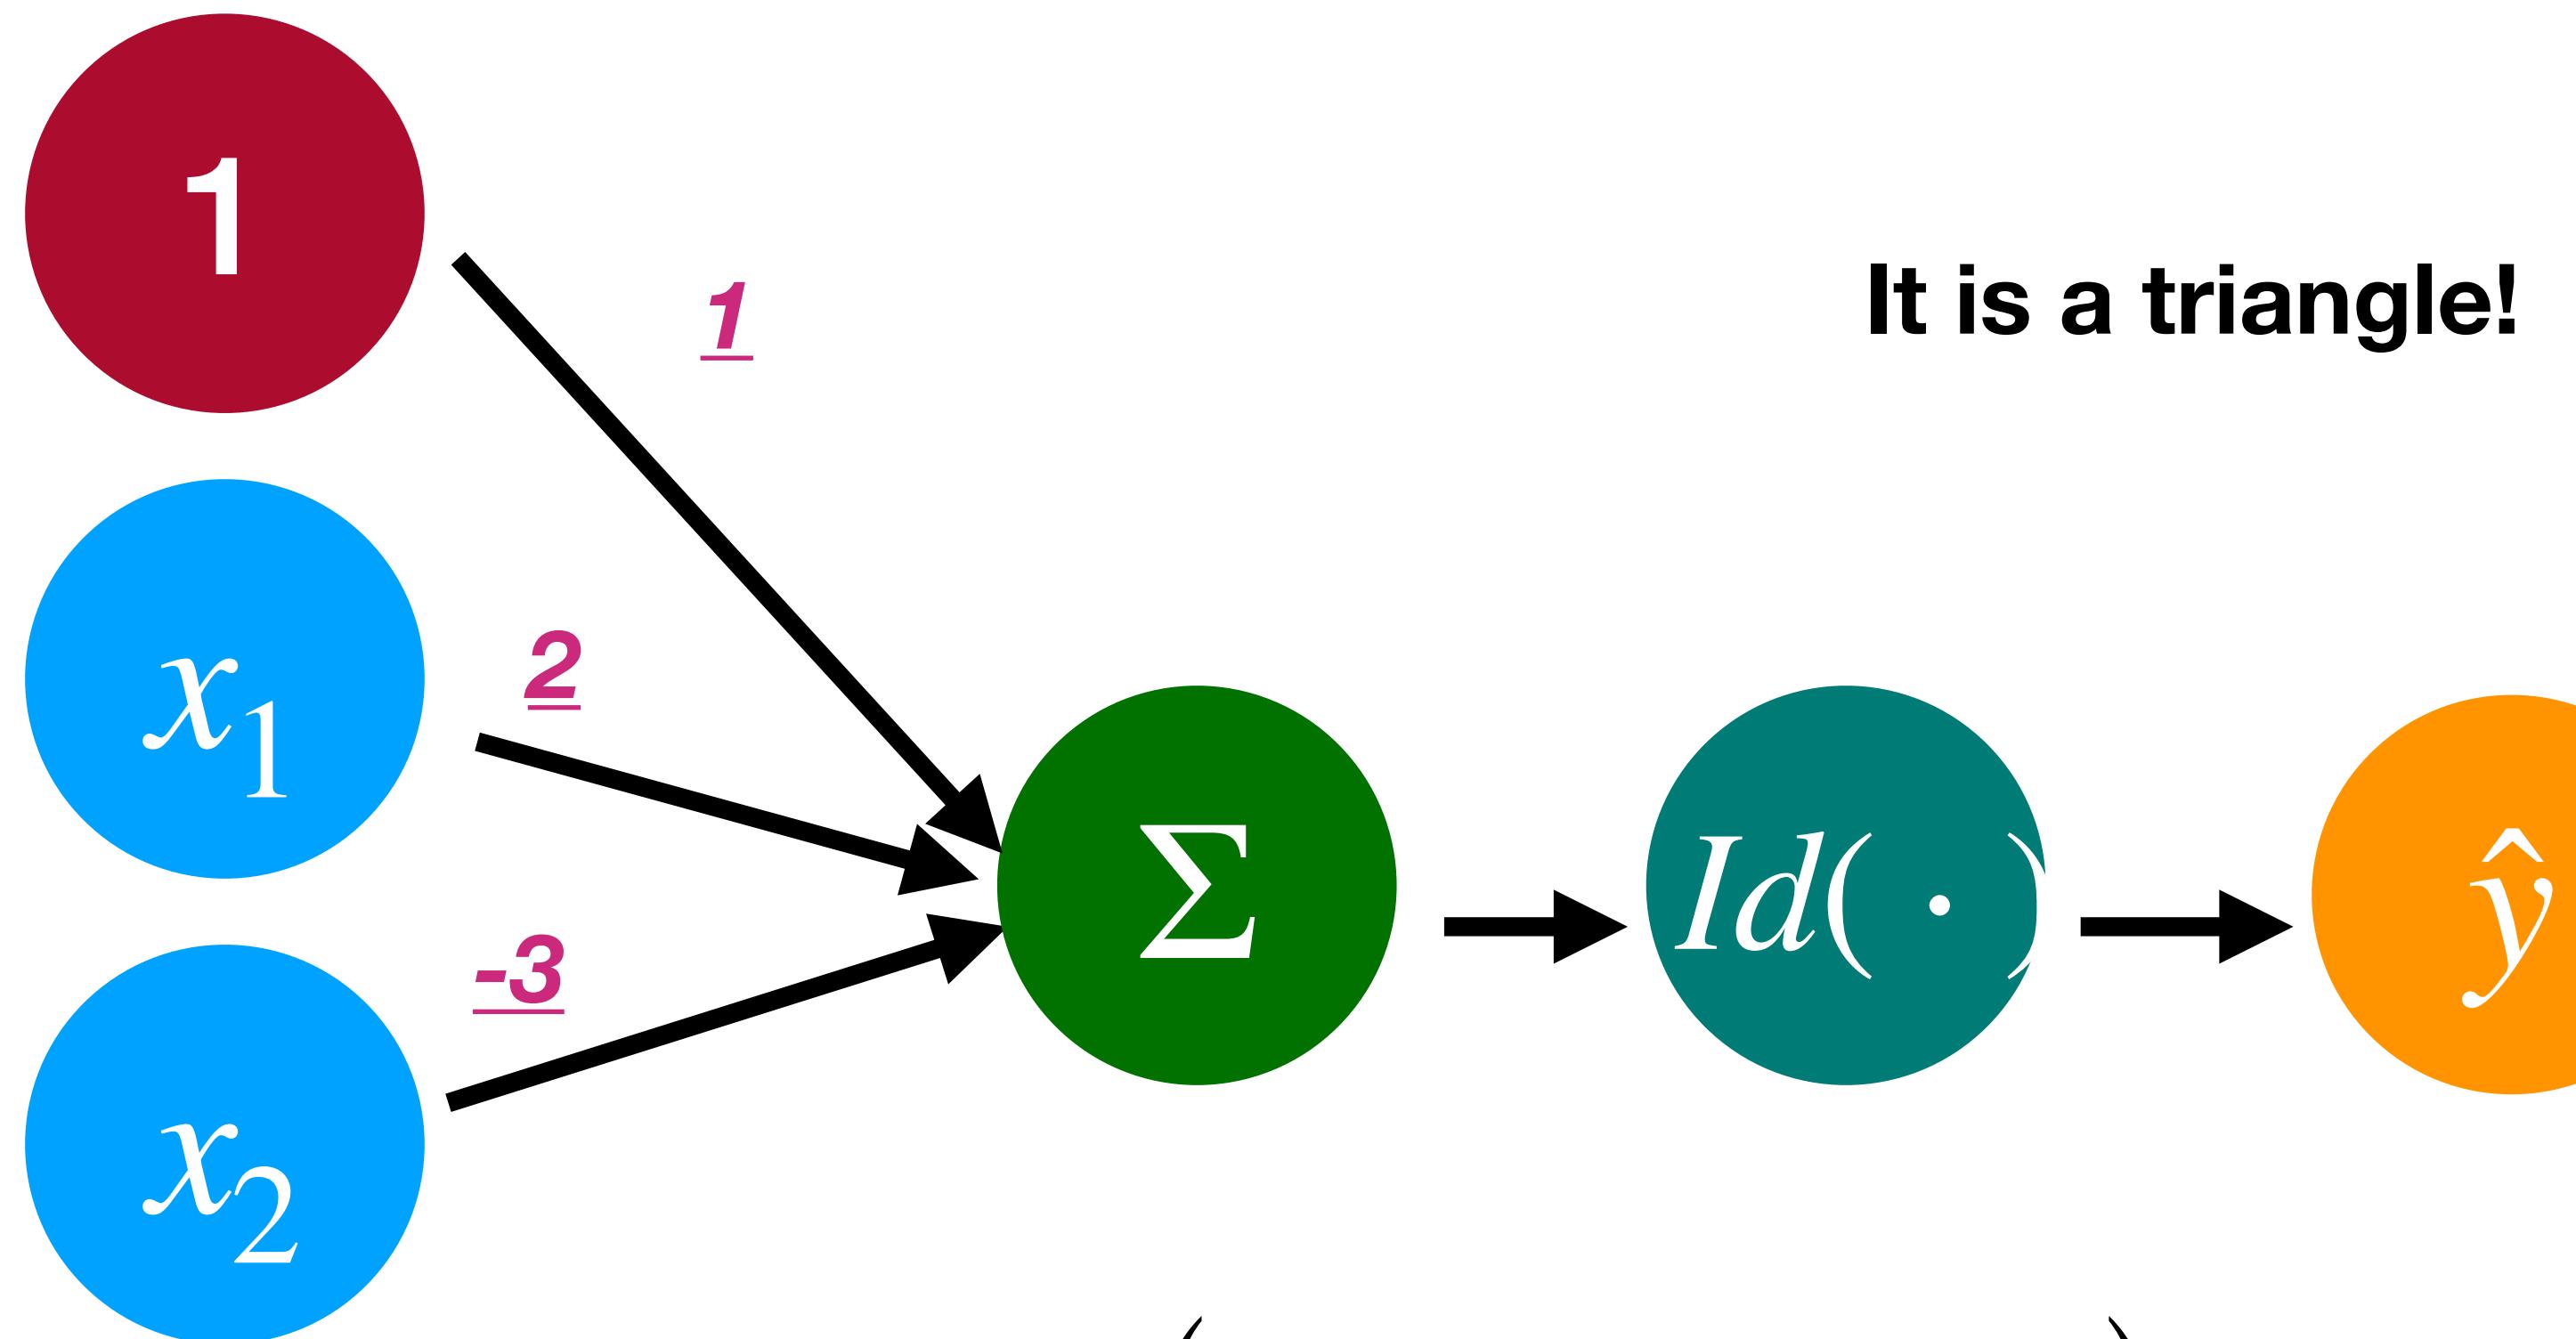
$$\hat{y} = (1 + 3x_1 - 2x_2)$$

corresponds to the any scalar > 0 if the
Input is an triangle

corresponds to the any scalar < 0 if the
Input is an square

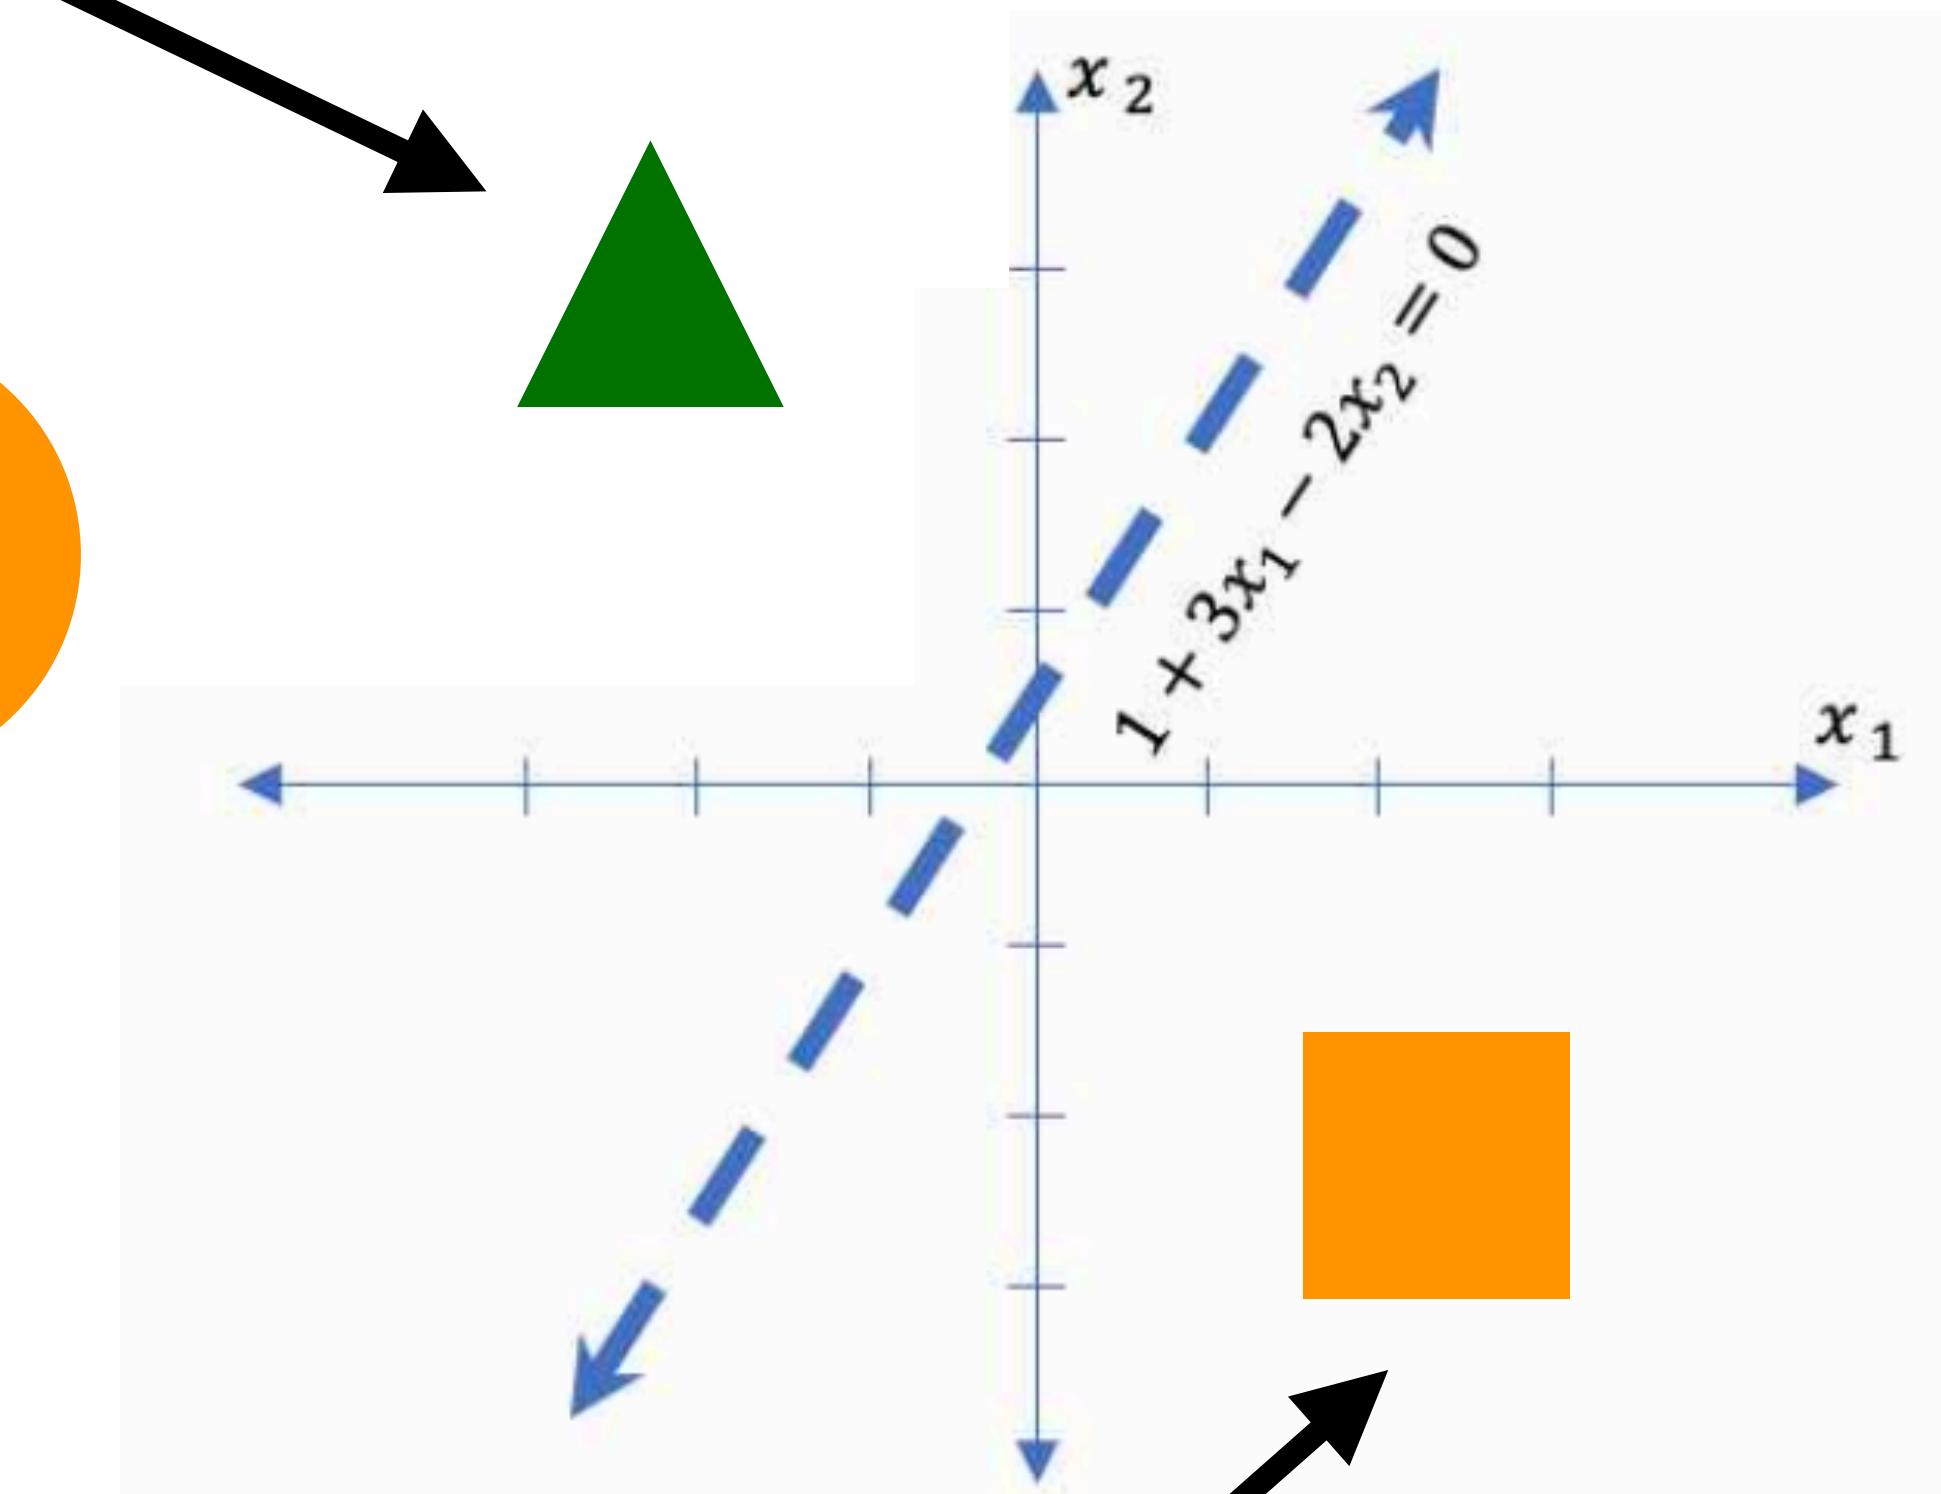


Our Perceptron to distinguish triangles and square



corresponds to the any scalar > 0 if the
Input is an triangle

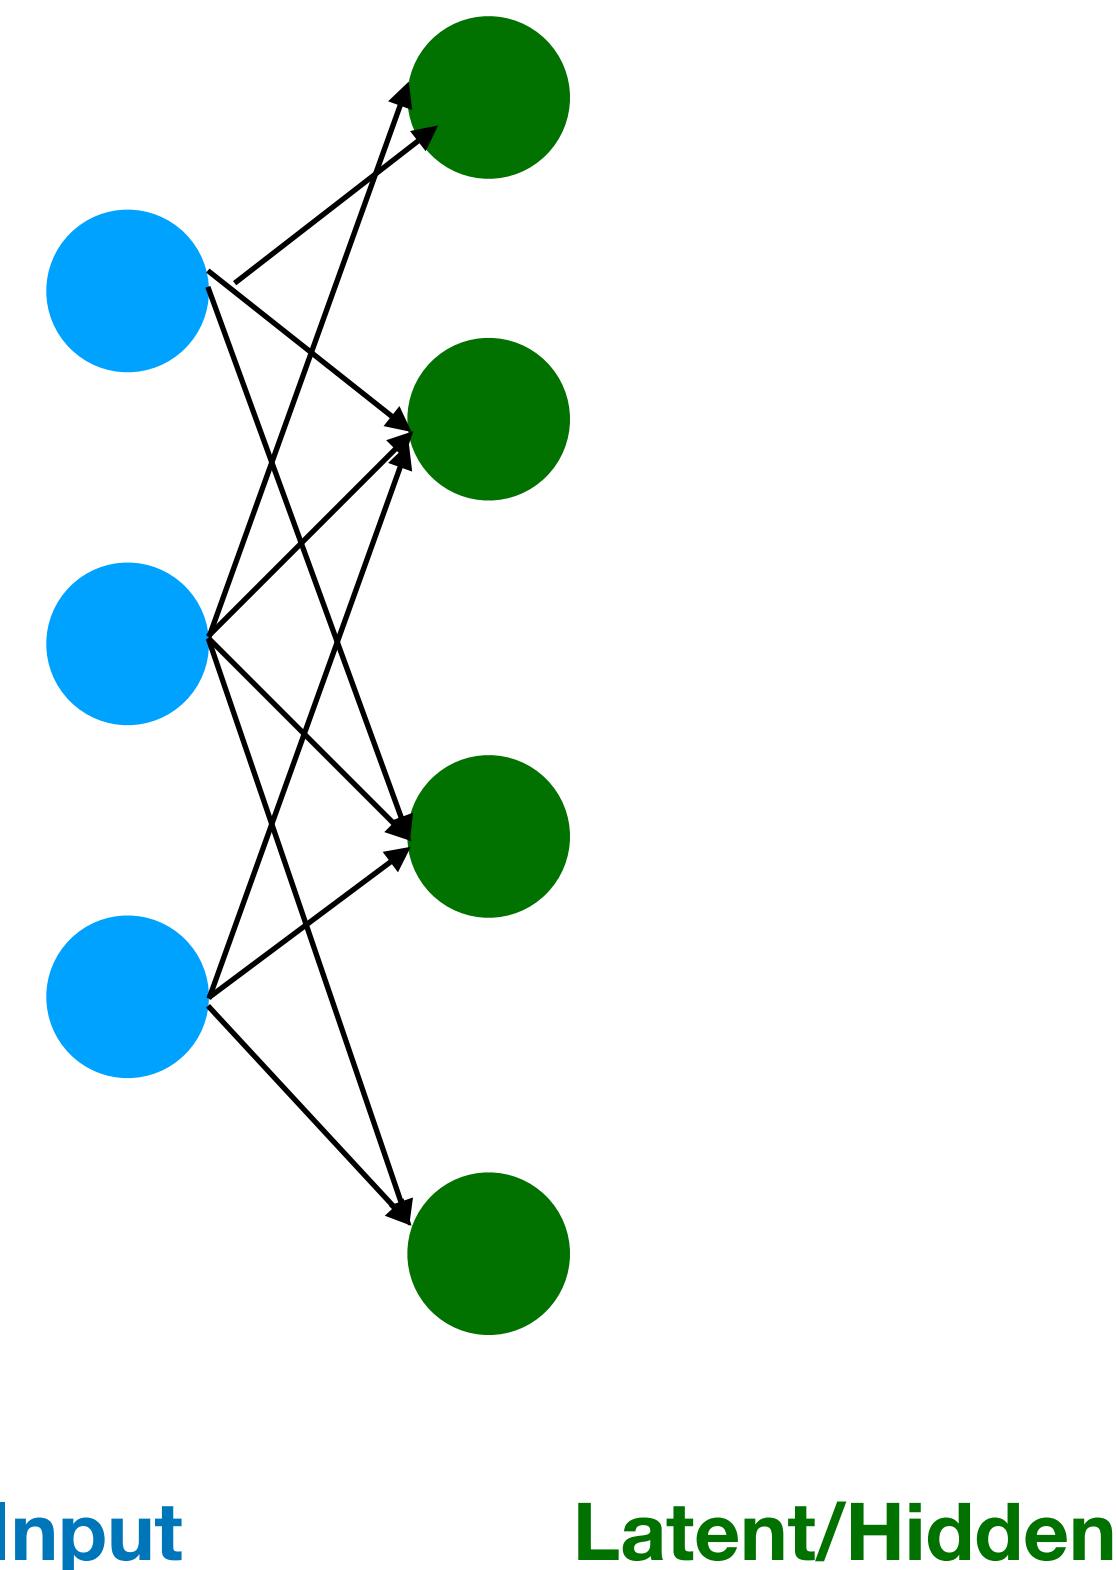
corresponds to the any scalar < 0 if the
Input is an square



It is a square

Towards more complex architectures

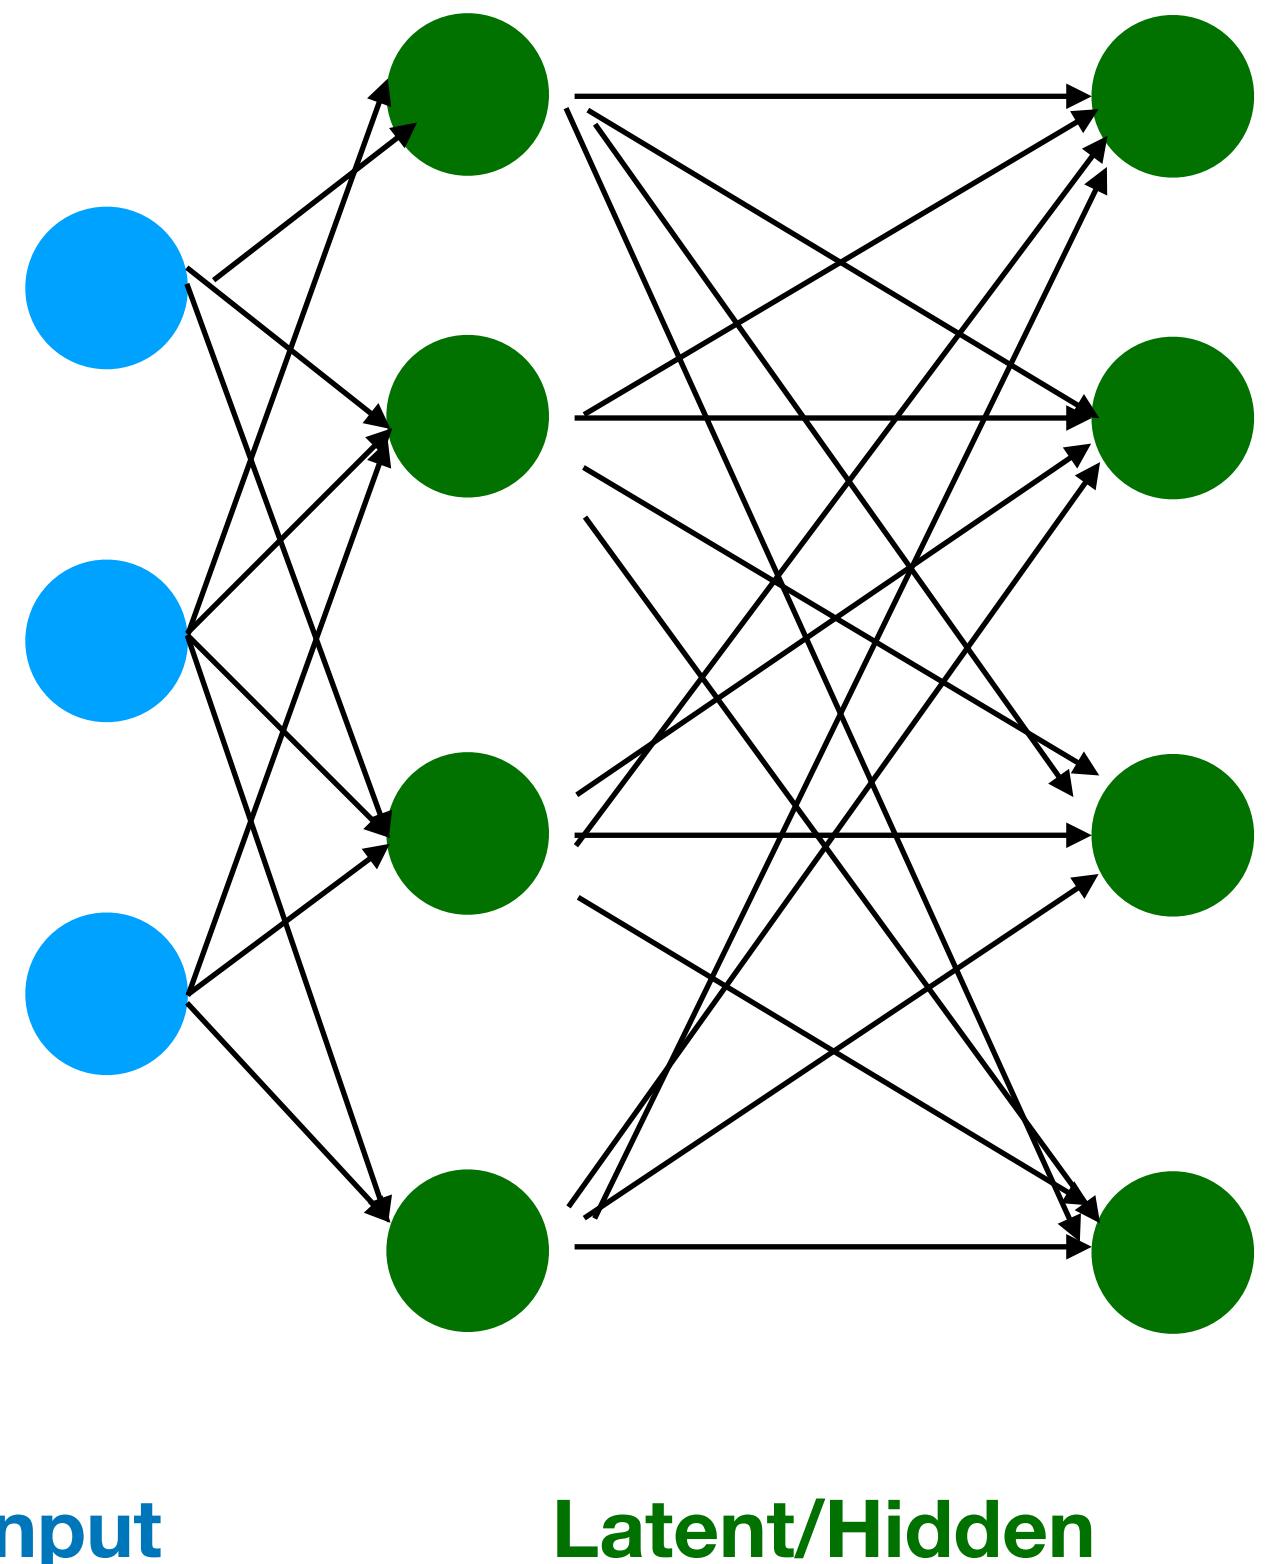
Towards more complex architectures



Input

Latent/Hidden

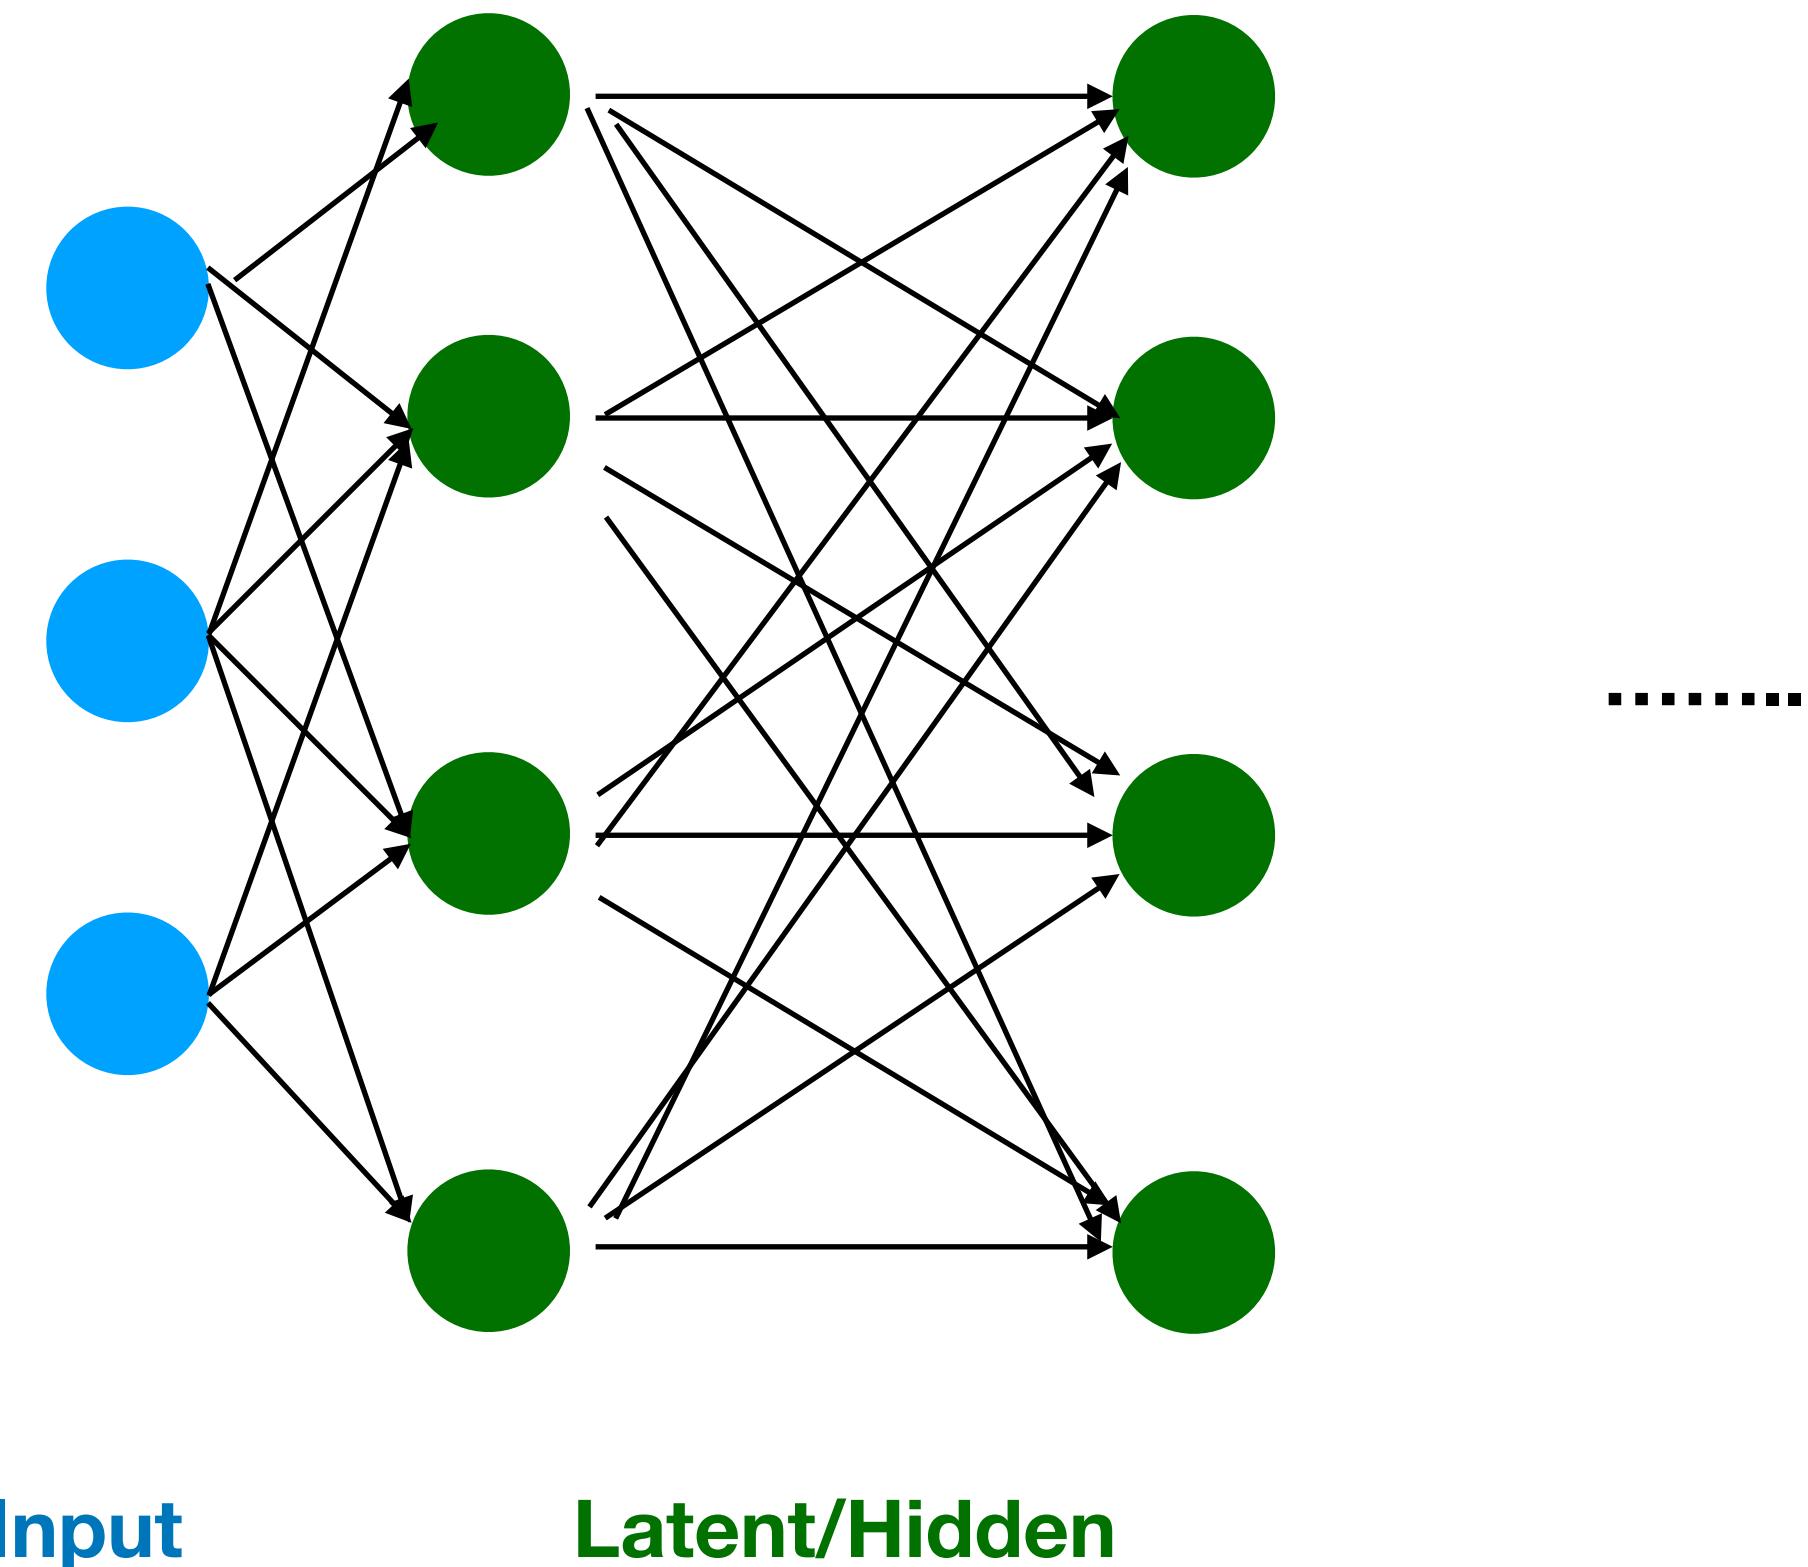
Towards more complex architectures



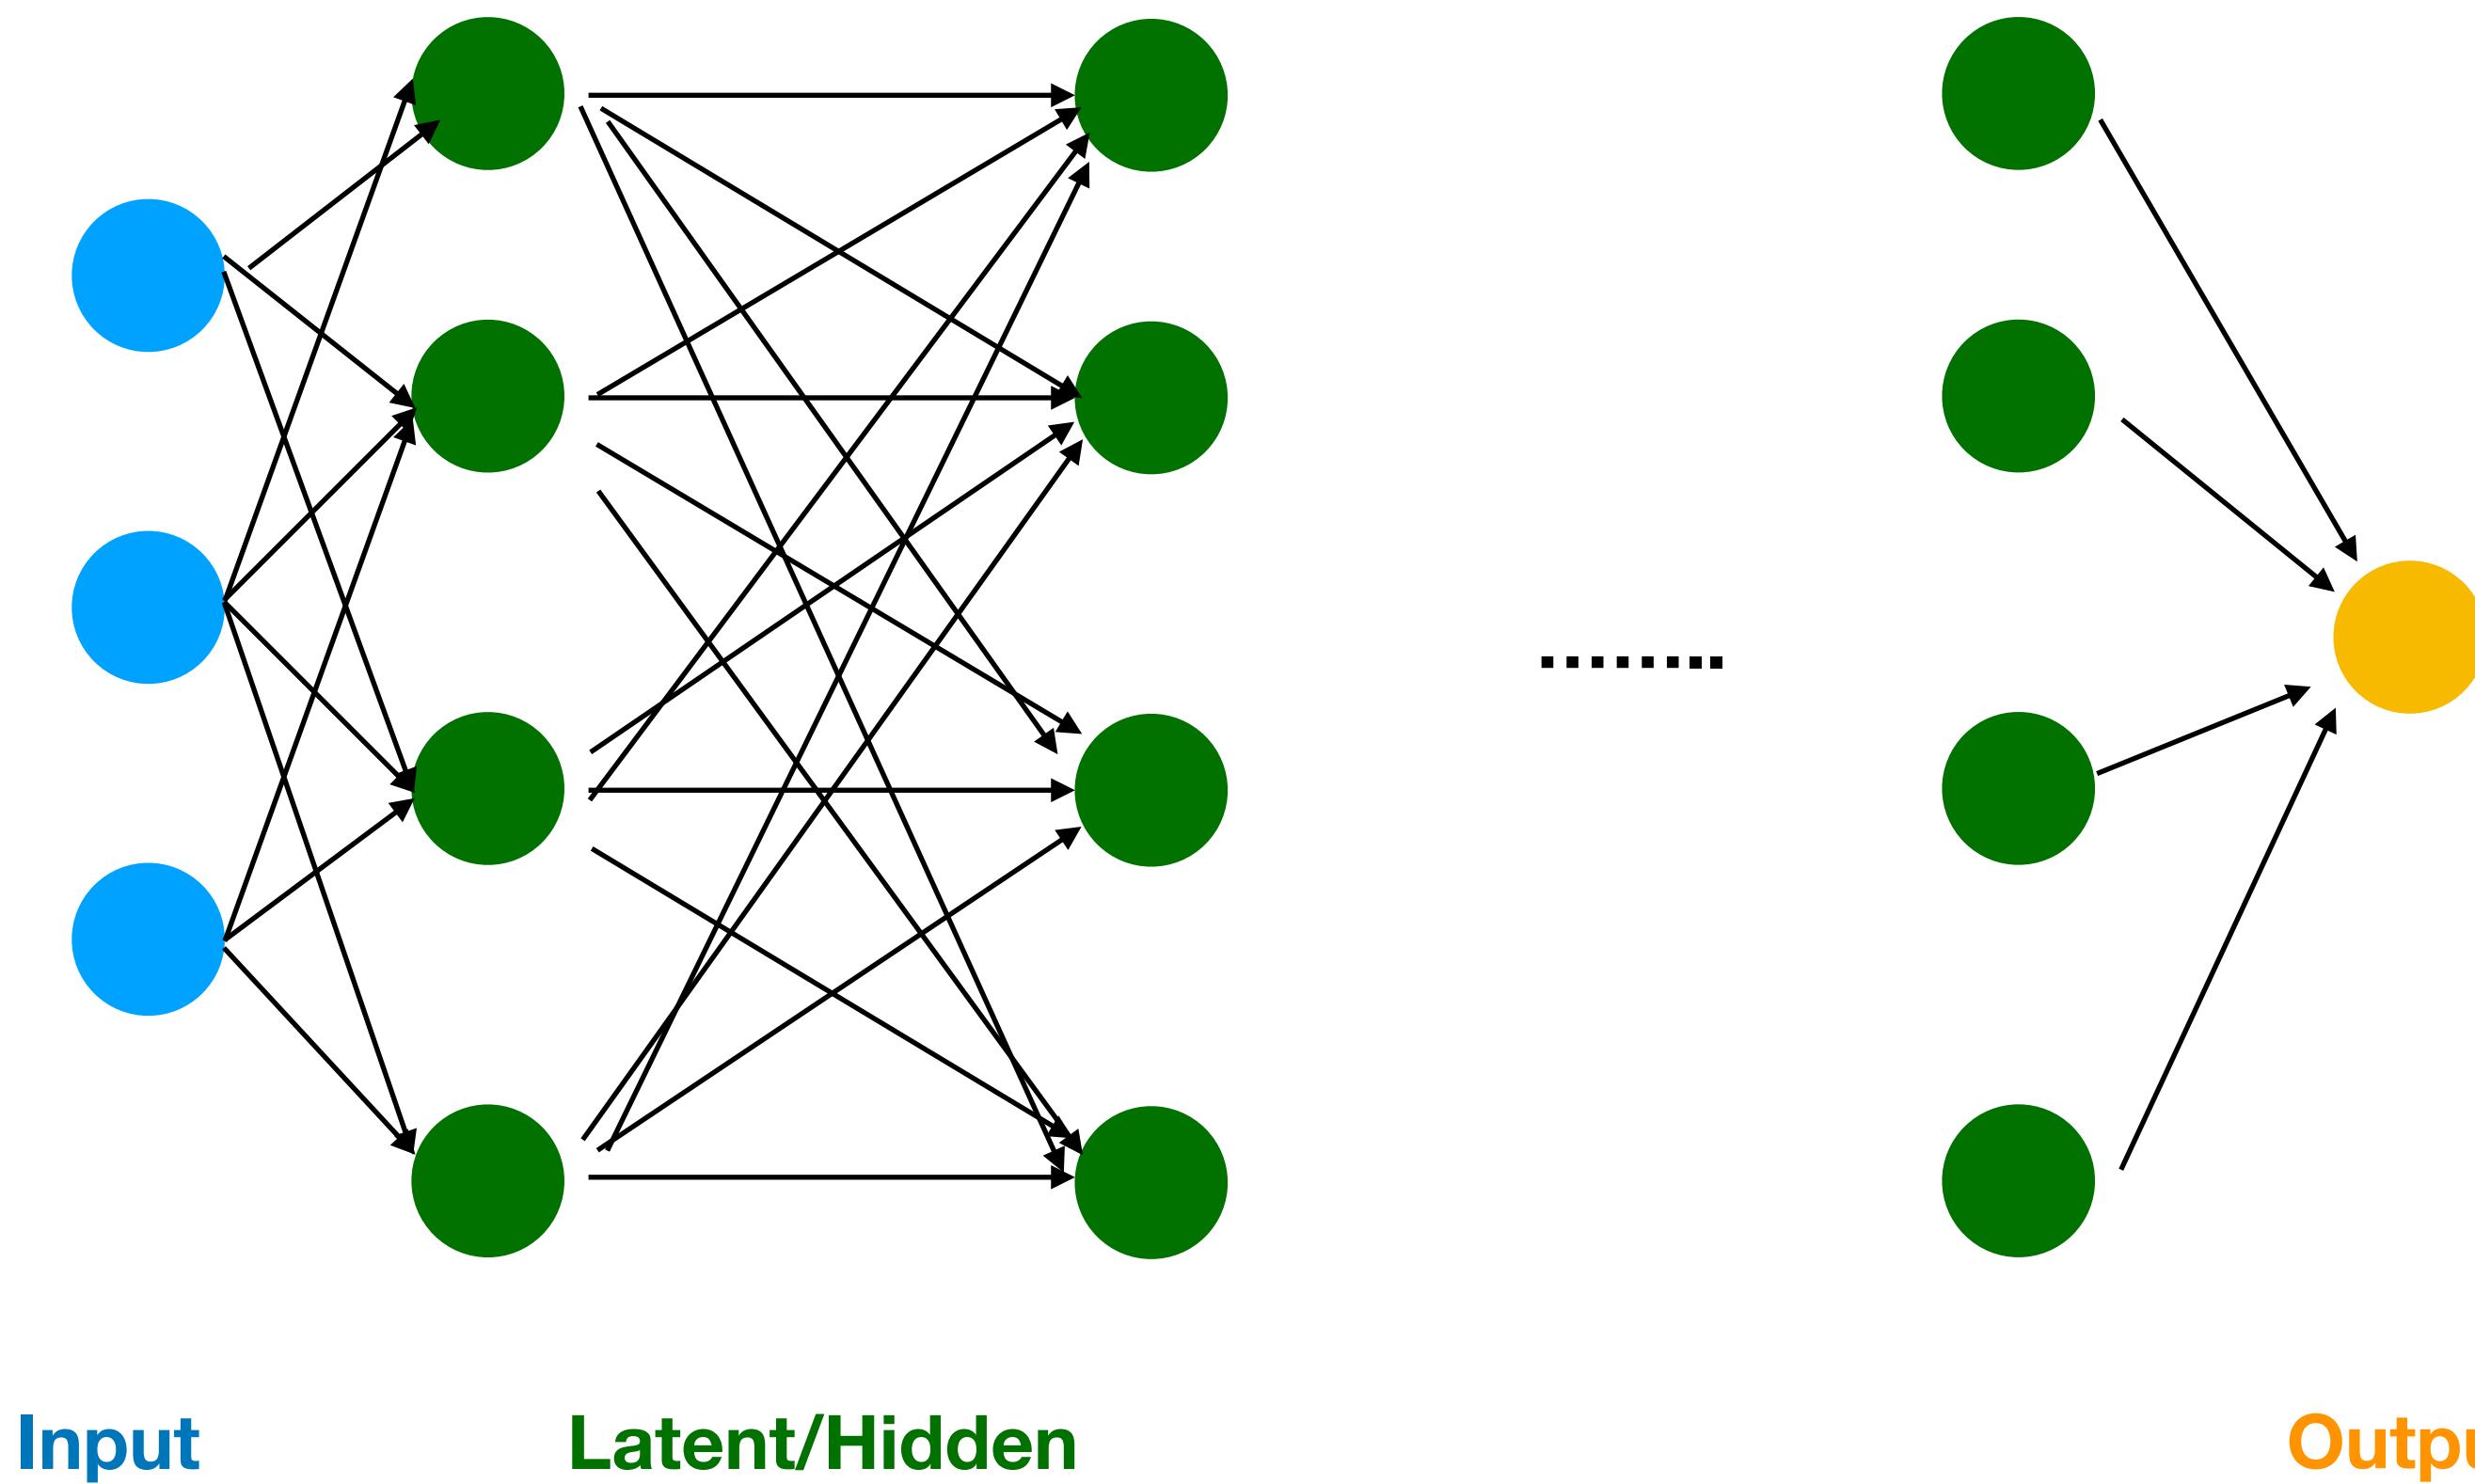
Input

Latent/Hidden

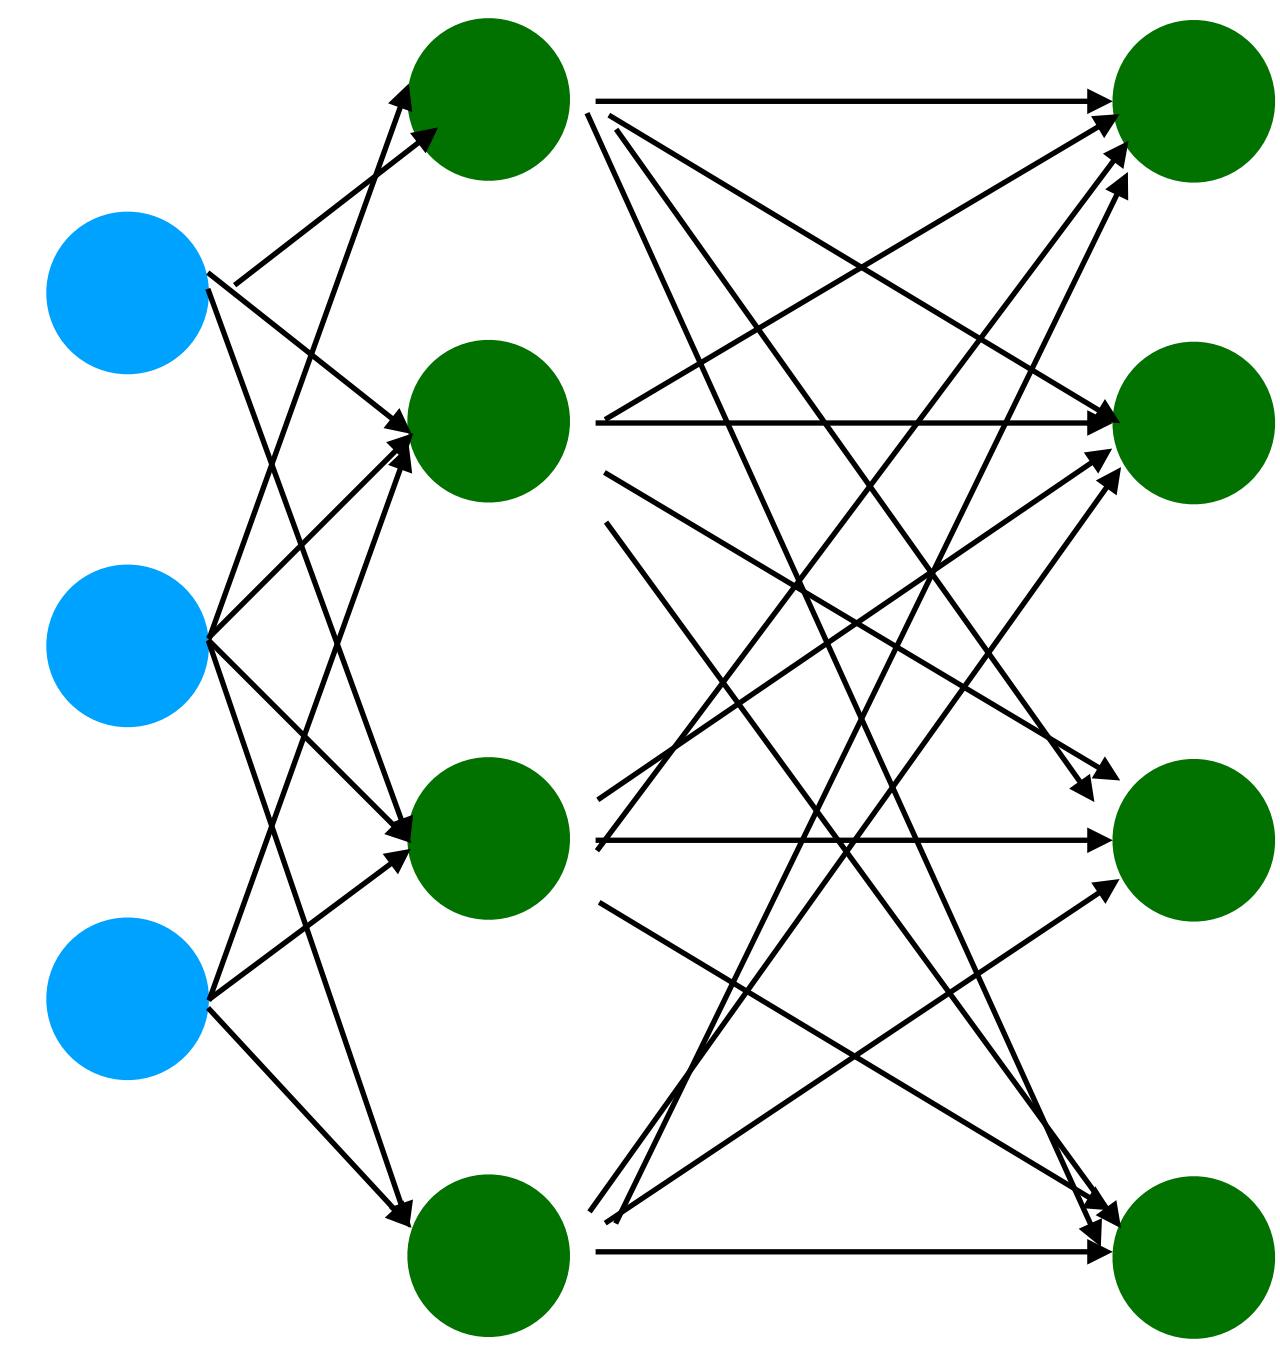
Towards more complex architectures



Towards more complex architectures



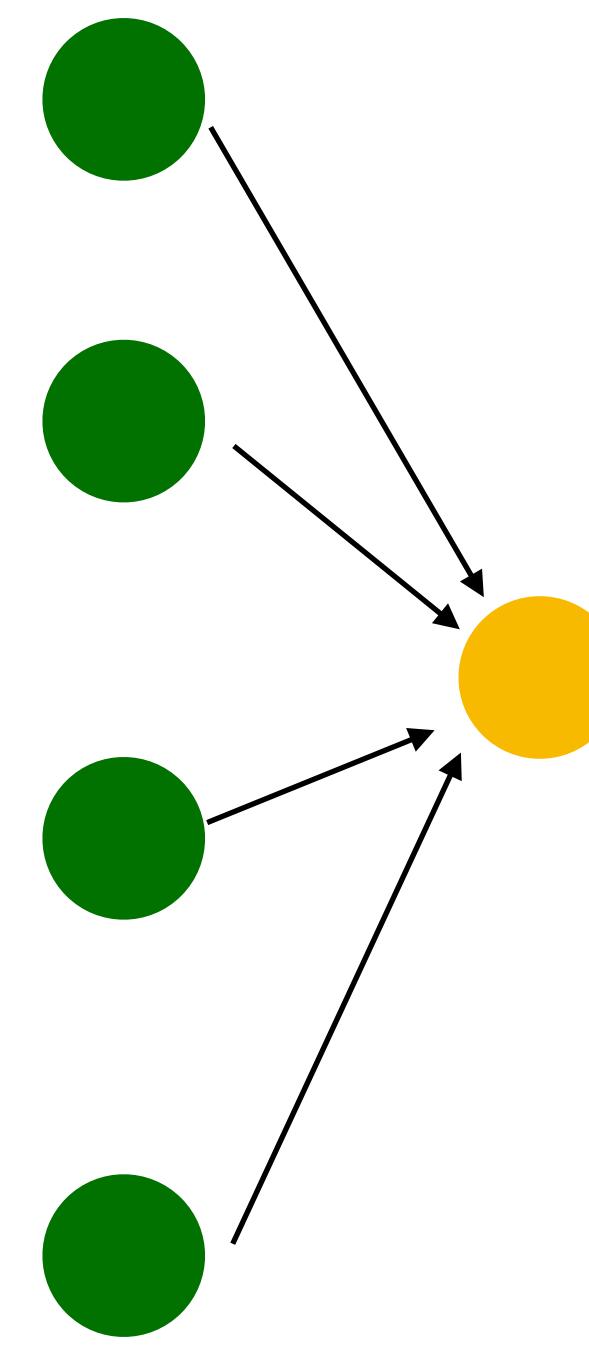
Towards more complex architectures



Input

Latent/Hidden

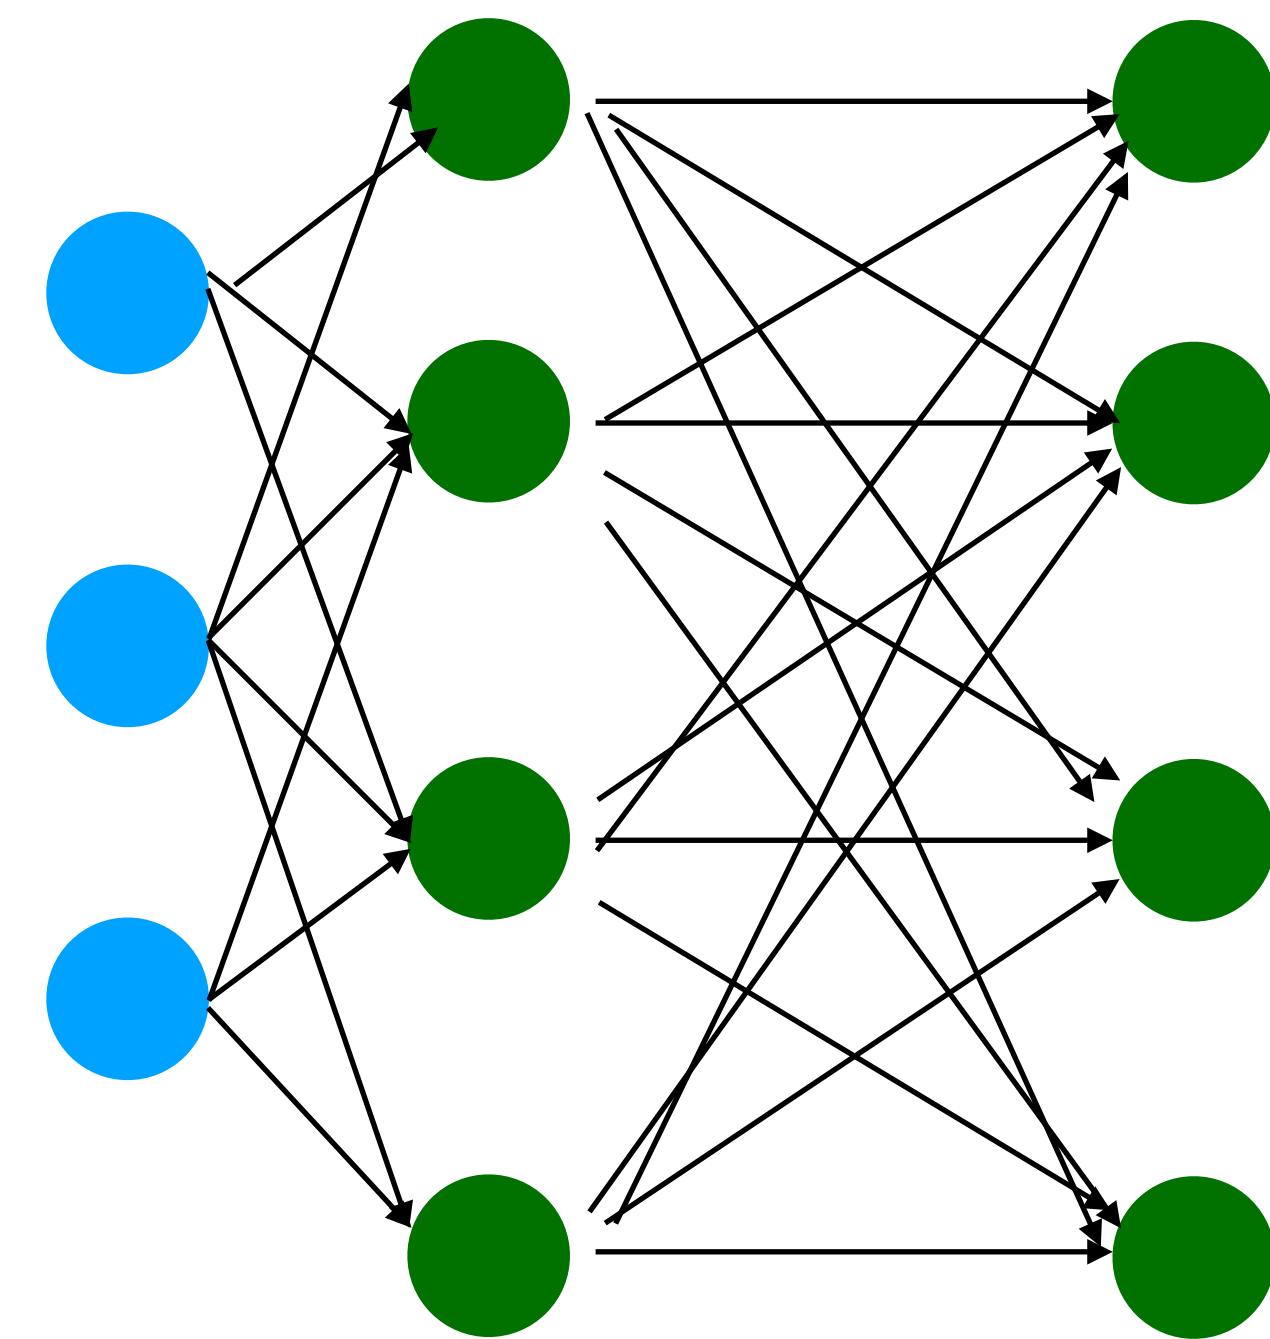
.....



Output

$$\hat{y} = \sigma \left(\dots \sigma \left(W_1 \sigma \left((X^T W_0)^T \right) \right) \right)$$

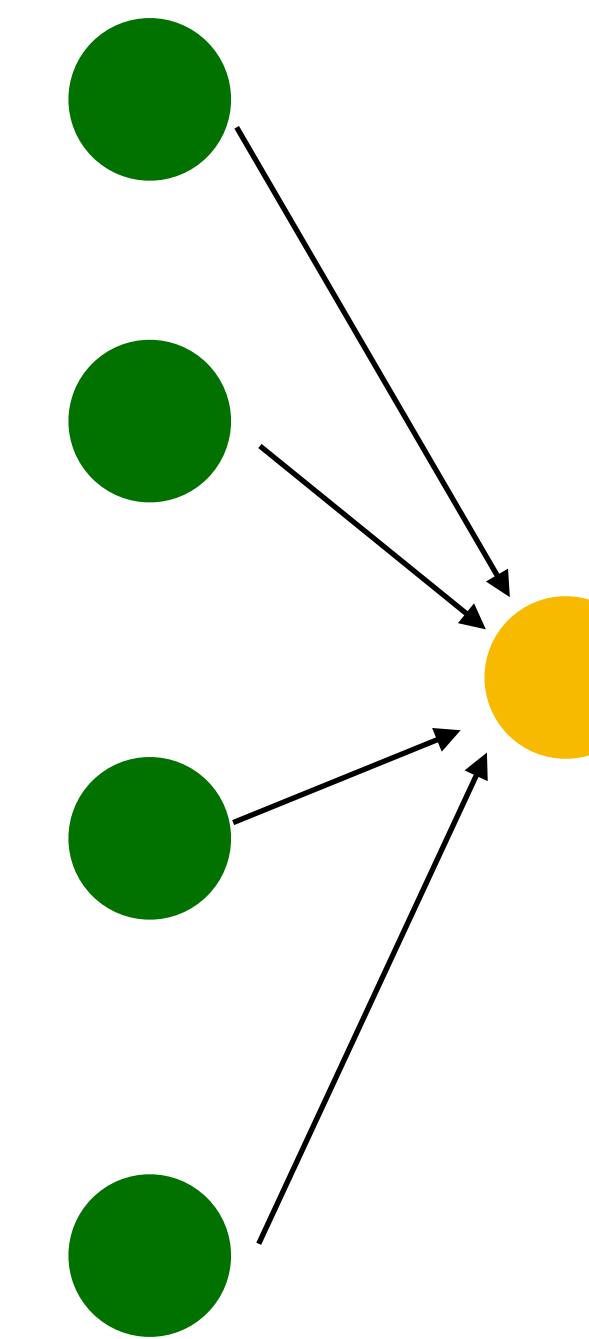
Towards more complex architectures



Input

Latent/Hidden

.....

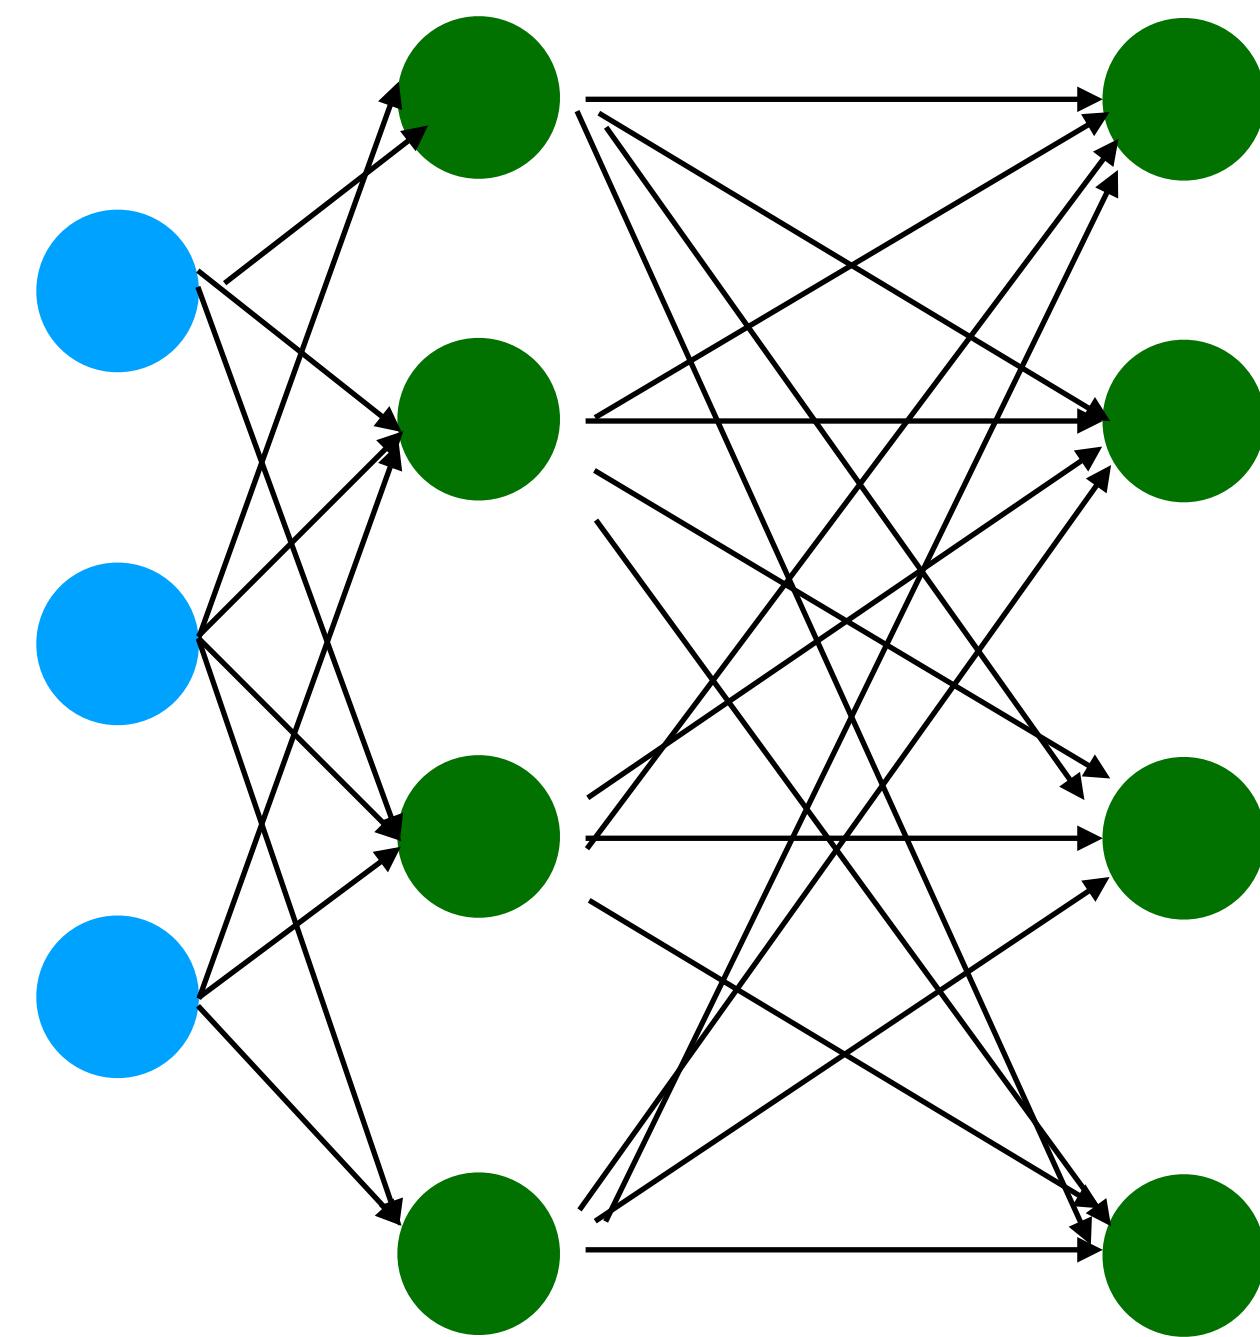


Output

$$\hat{y} = \sigma \left(\dots \sigma \left(W_1 \sigma \left((X^T W_0)^T \right) \right) \right)$$

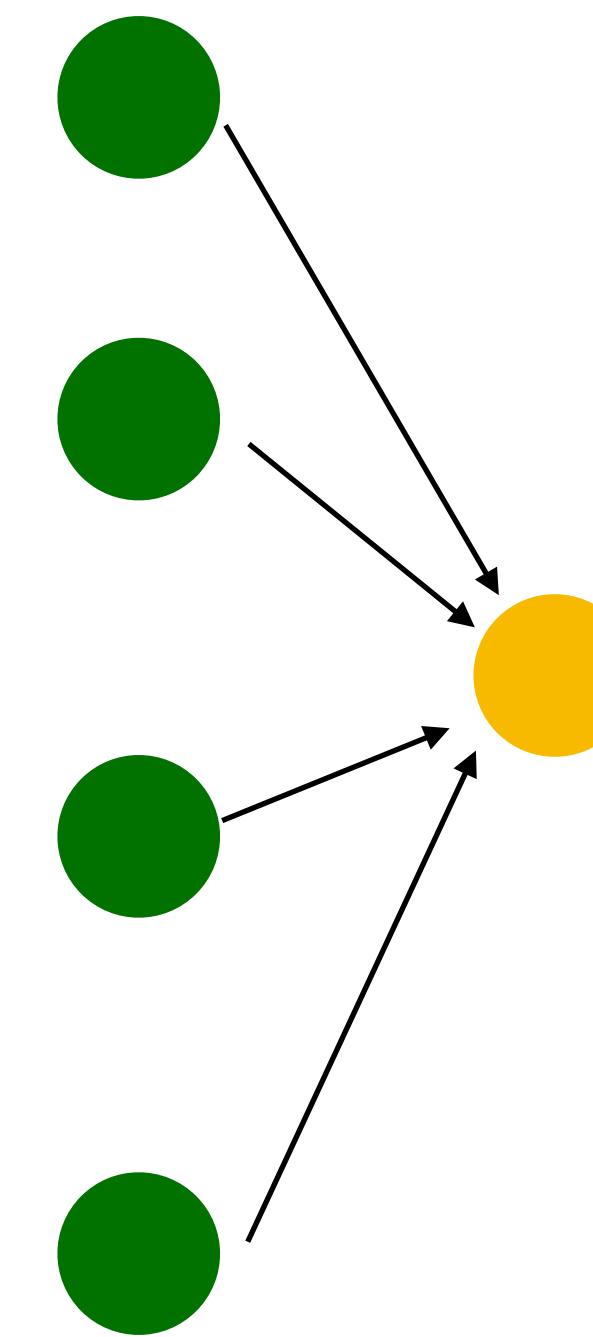
Adding layers is equivalent to adding new matrix and activation functions !

Towards more complex architectures



Input

Latent/Hidden



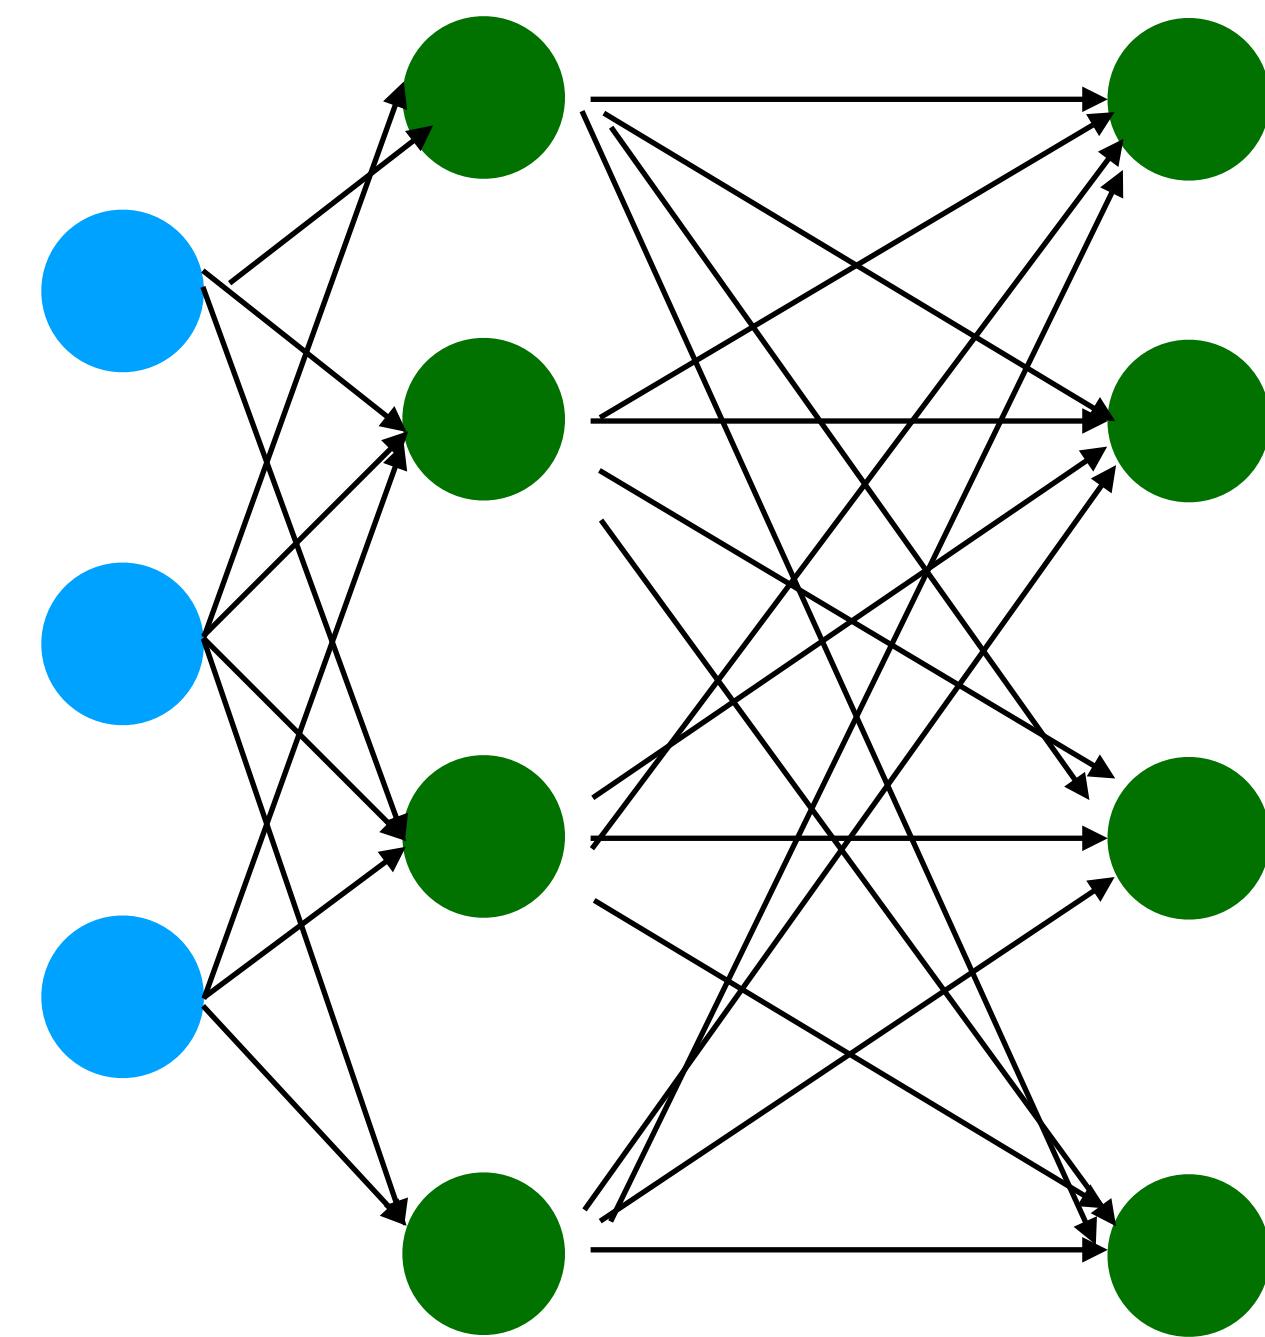
Output

$$\hat{y} = \sigma \left(\dots \sigma \left(W_1 \sigma \left((X^T W_0)^T \right) \right) \right)$$

Adding layers is equivalent to adding new matrix and activation functions !

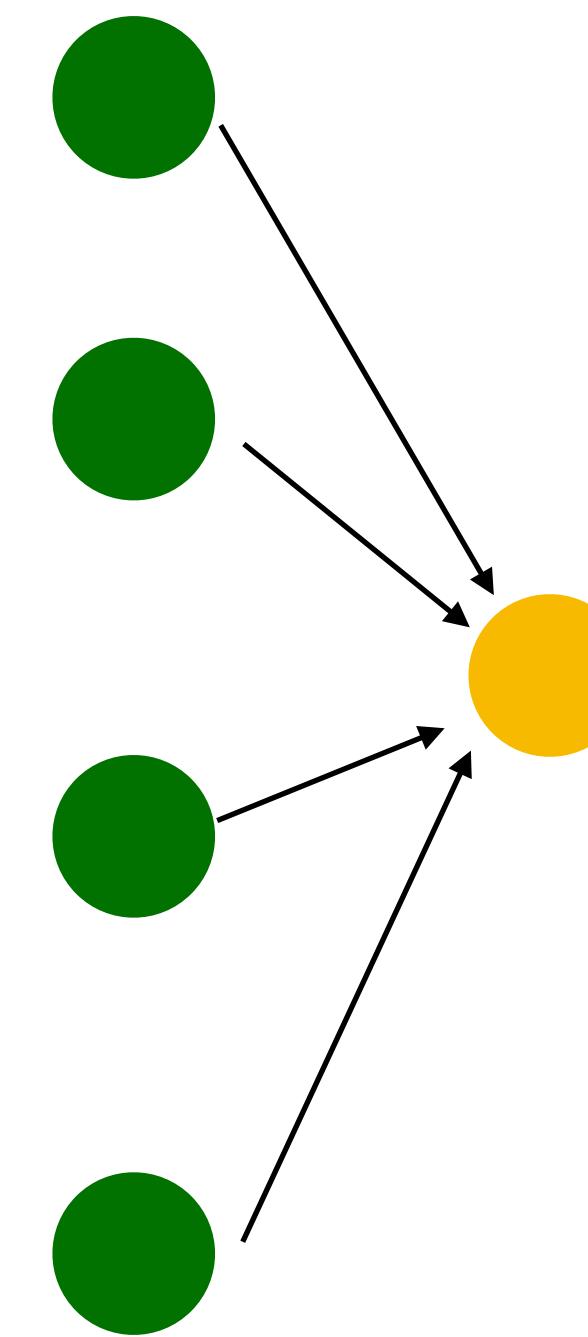
How do I choose the network architecture?

Towards more complex architectures



Input

Latent/Hidden

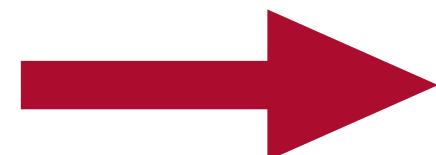


Output

$$\hat{y} = \sigma \left(\dots \sigma \left(W_1 \sigma \left((X^T W_0)^T \right) \right) \right)$$

Adding layers is equivalent to adding new matrix and activation functions !

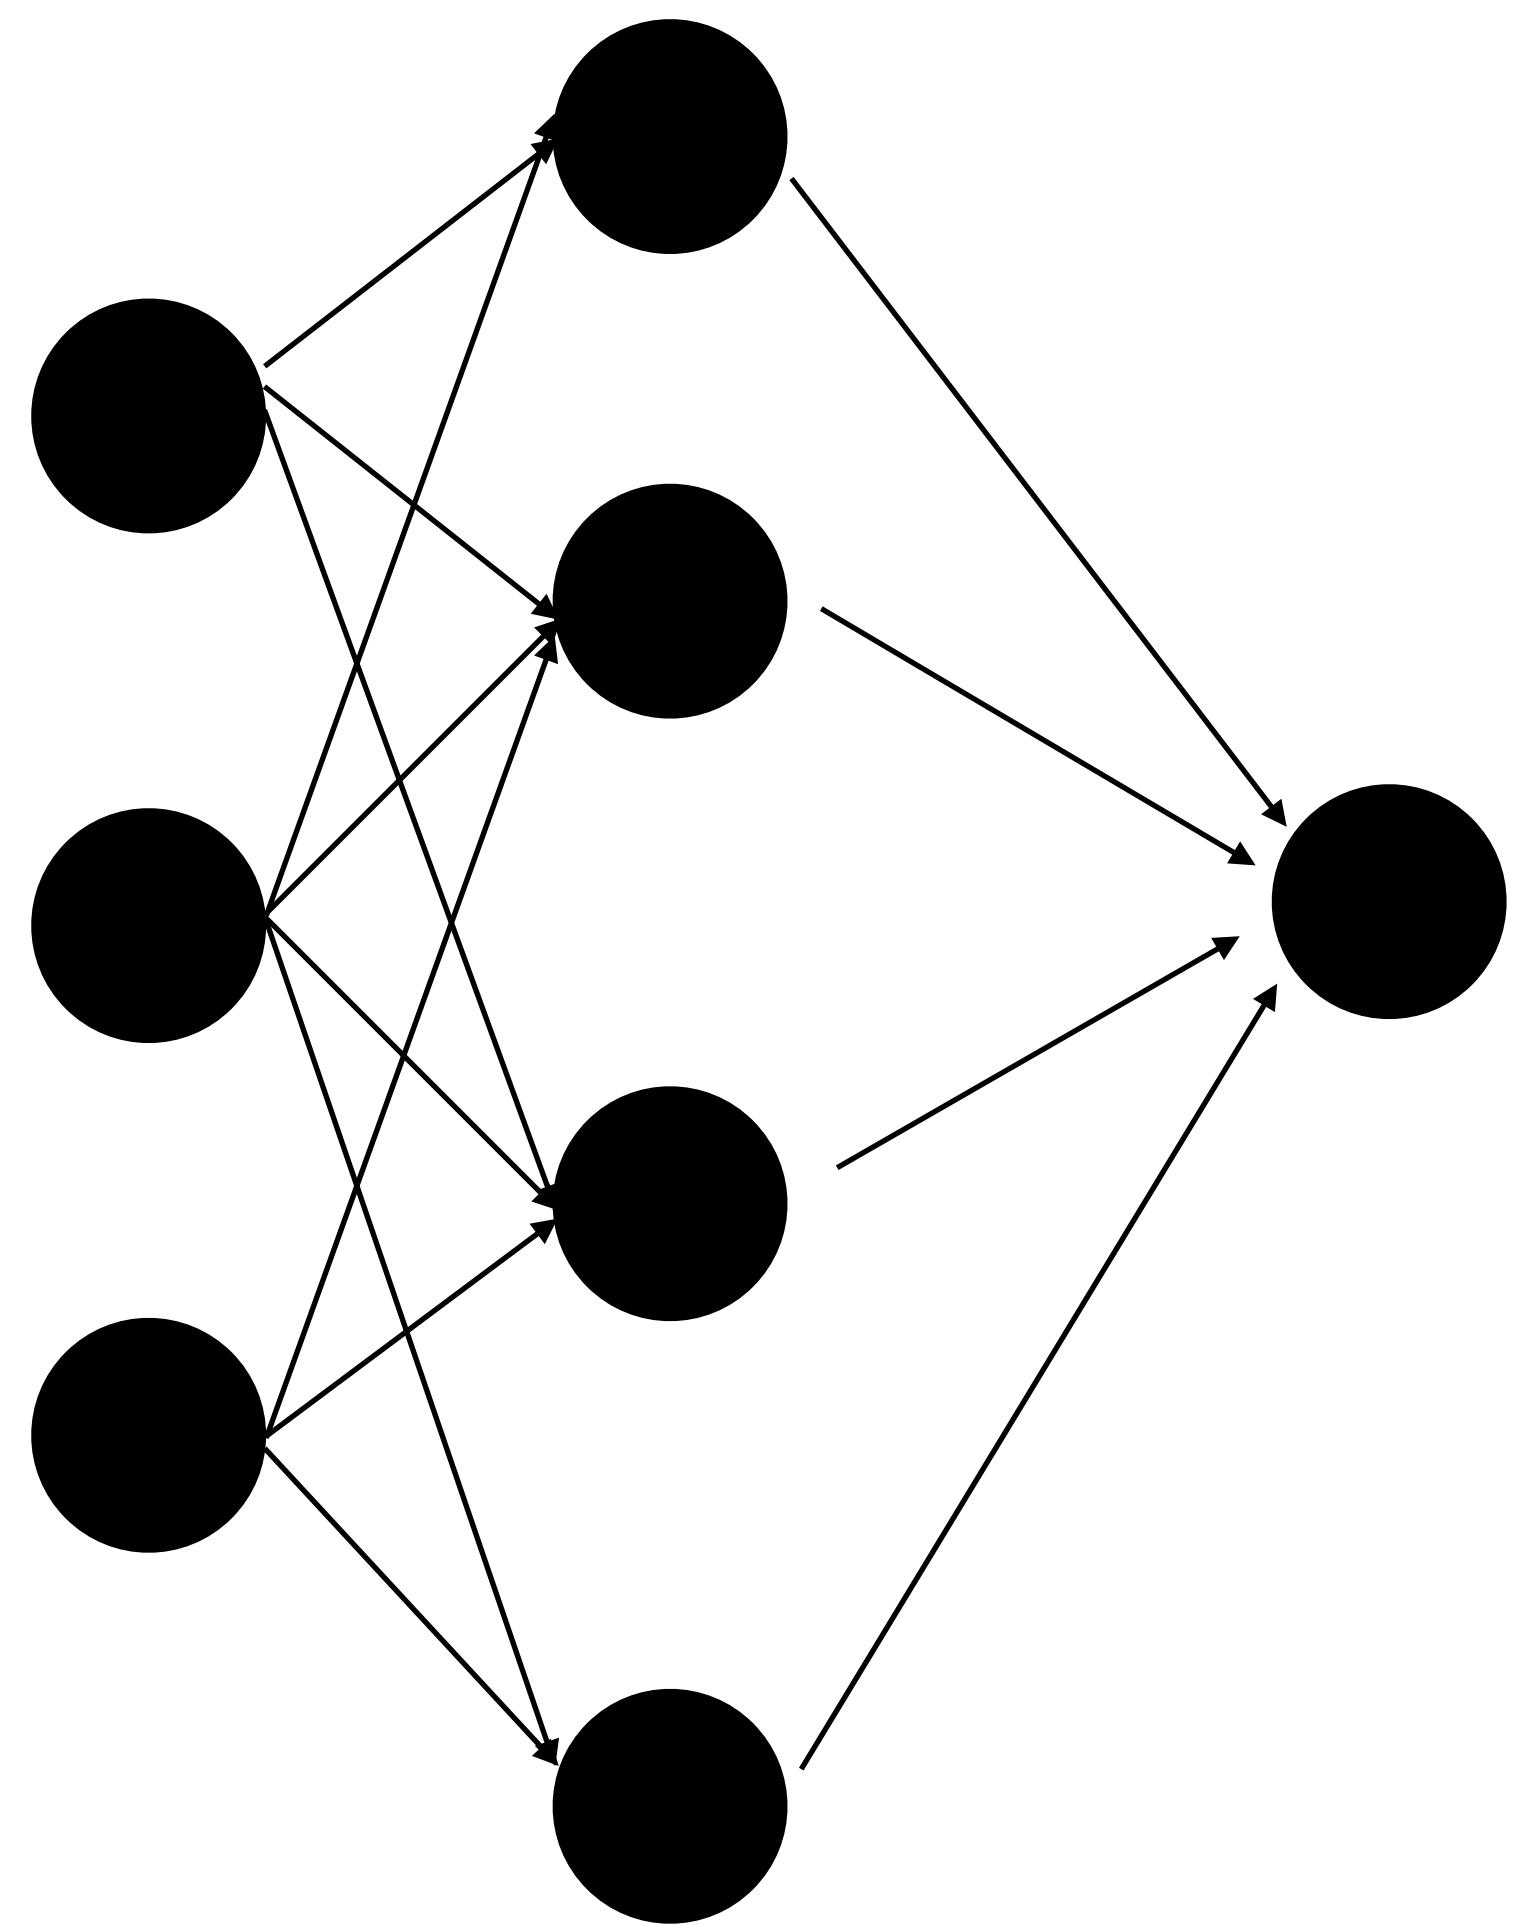
How do I choose the network architecture?



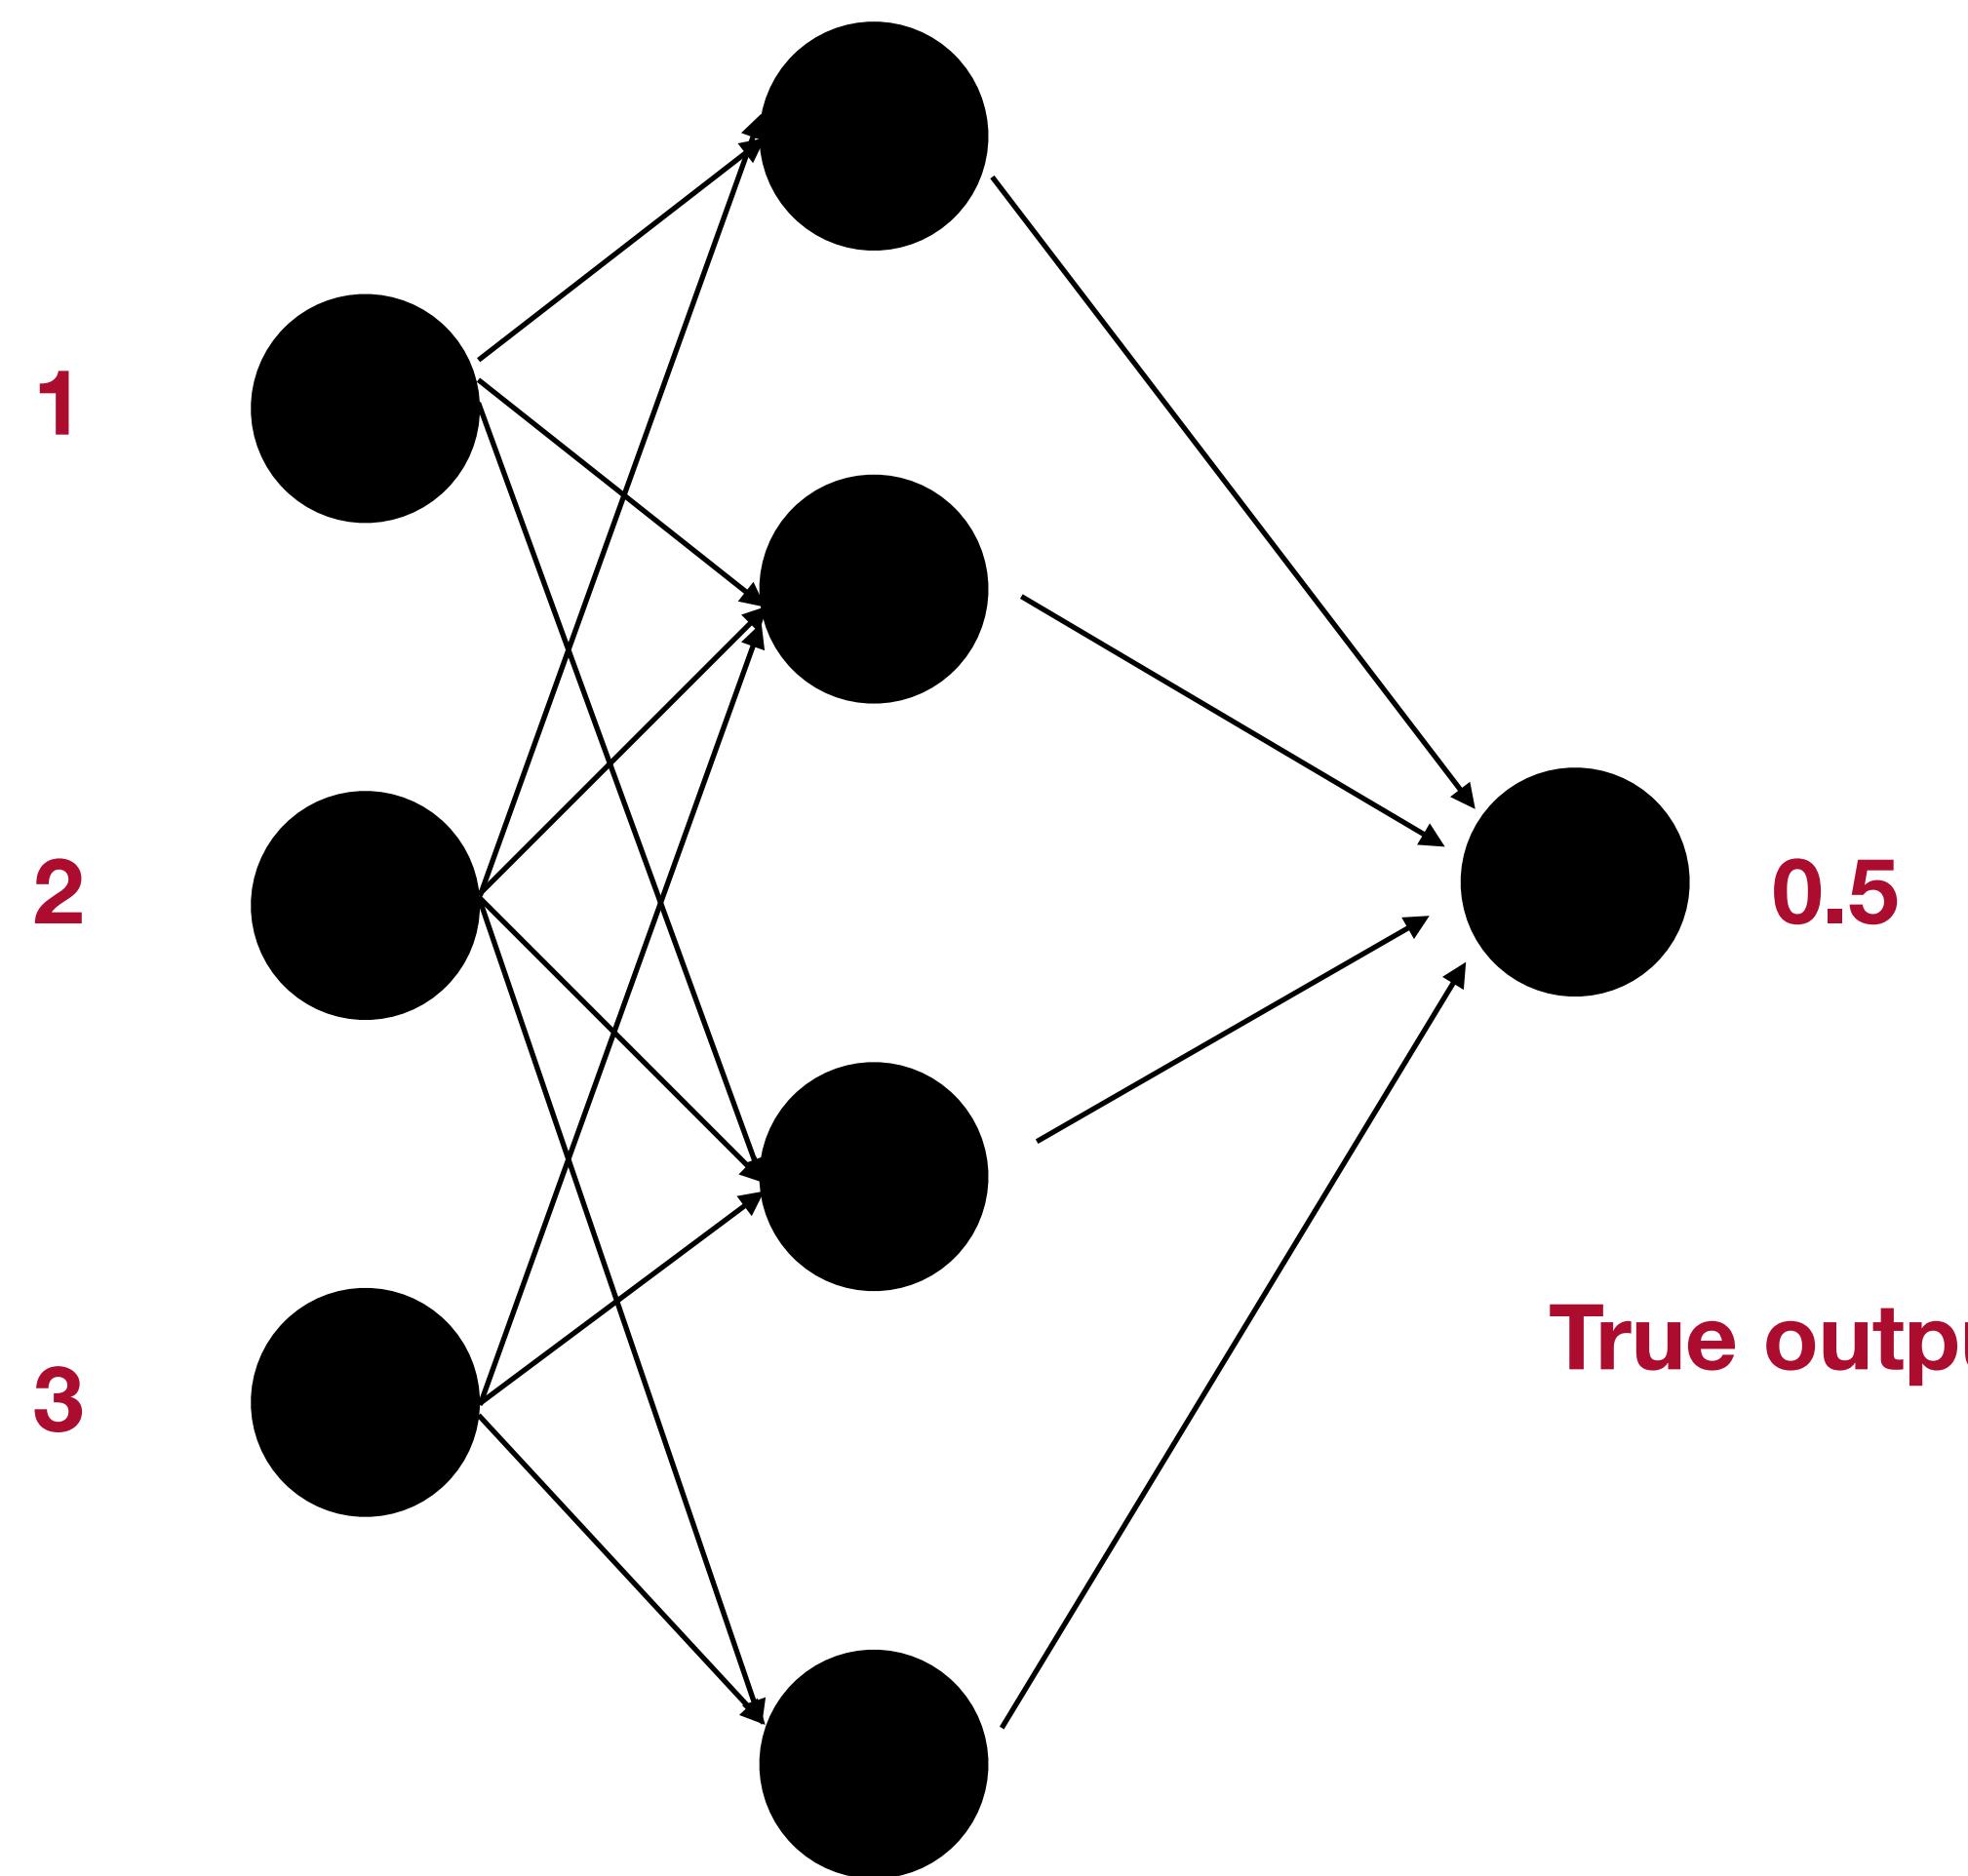
Research papers!

Another example with a one layer NN

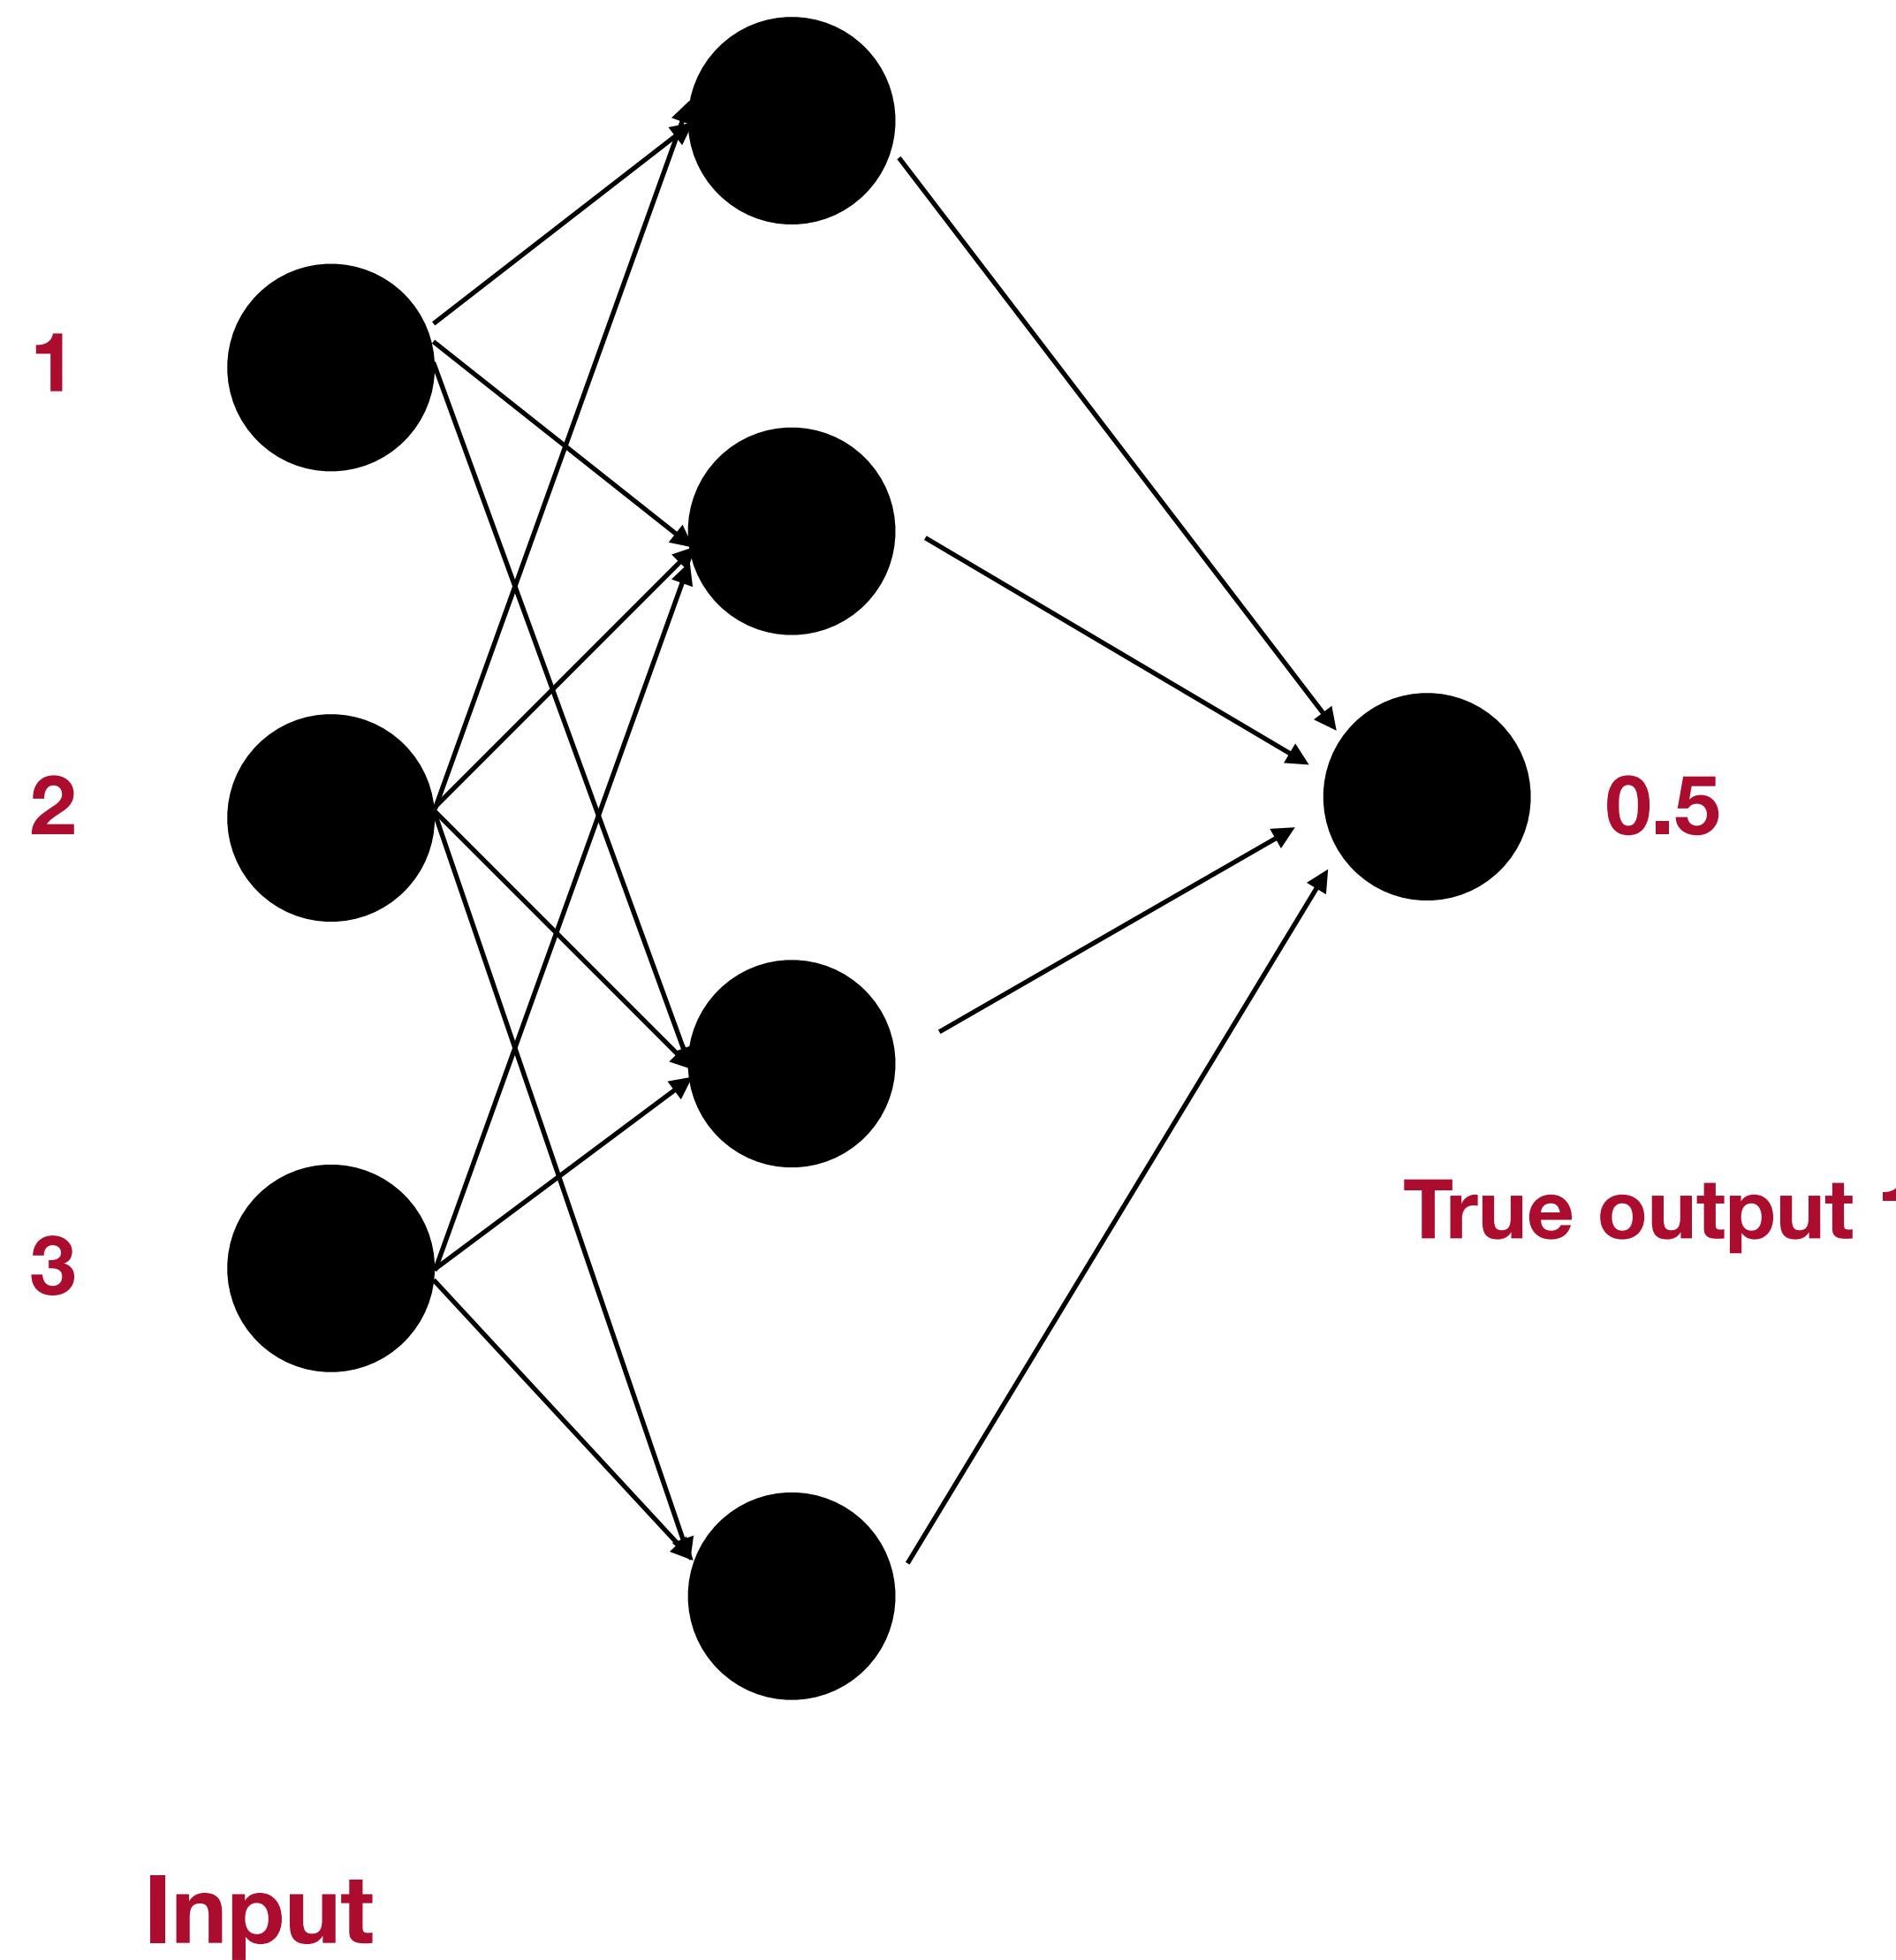
Another example with a one layer NN



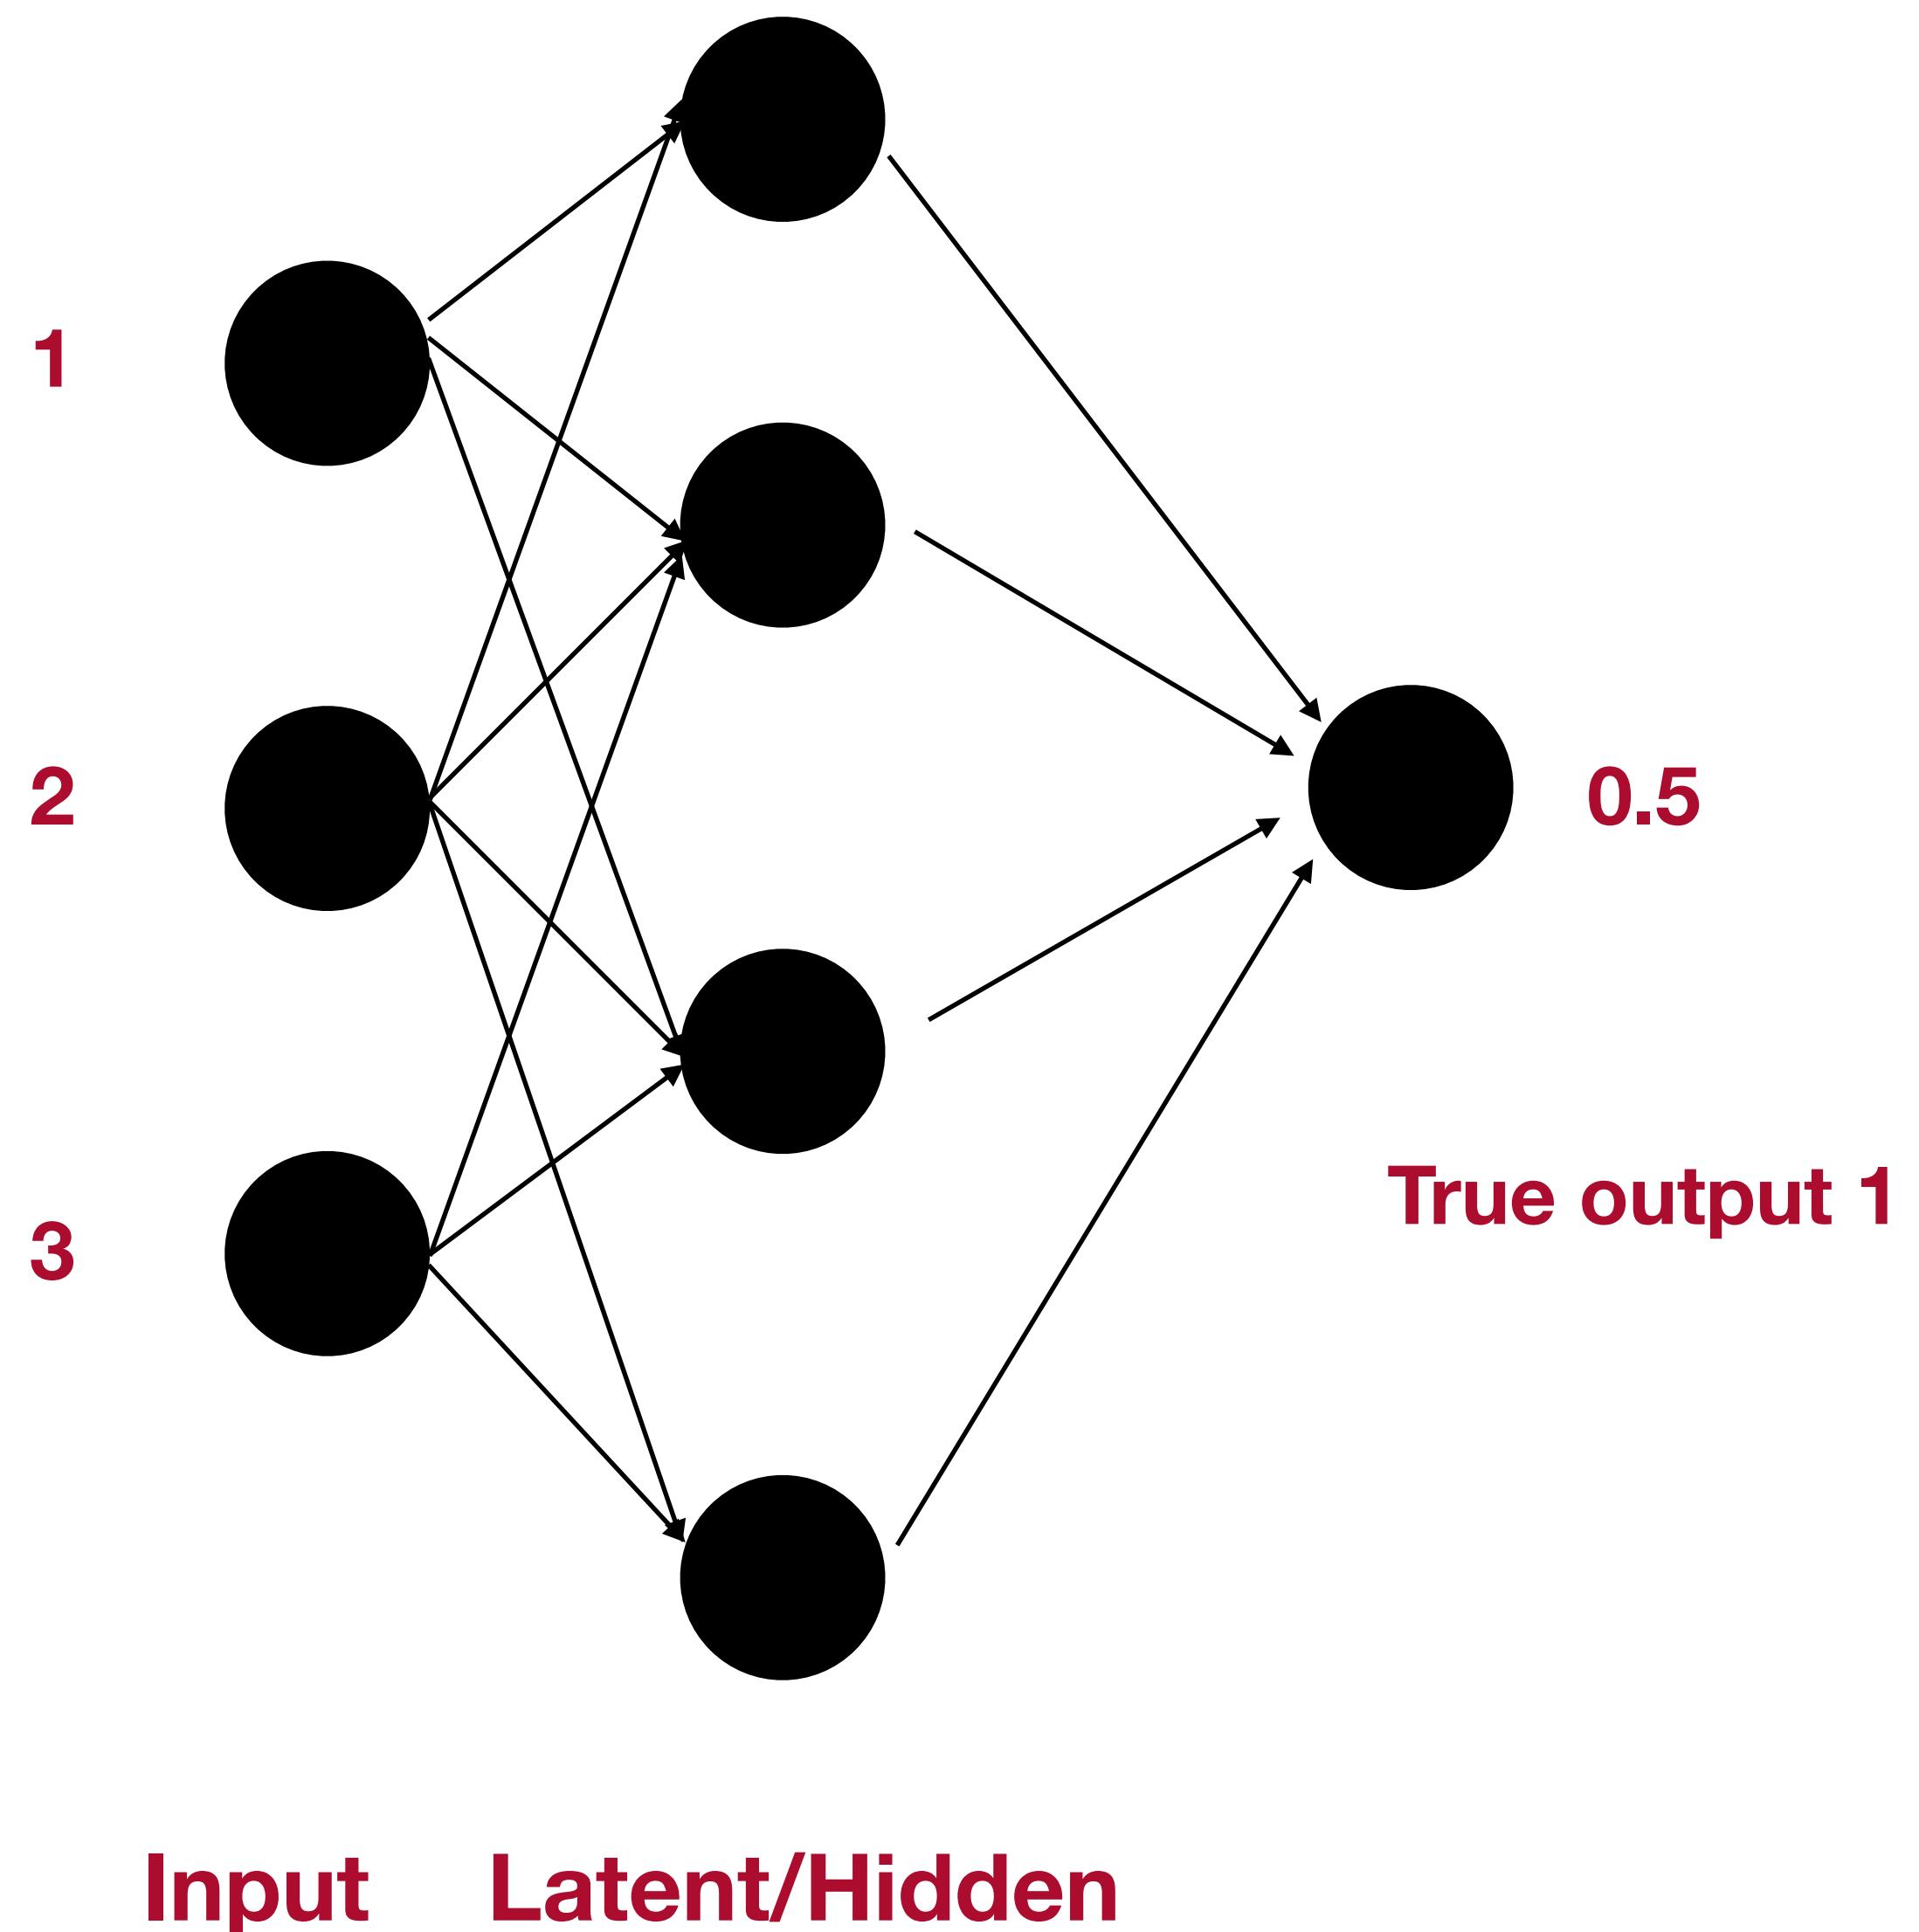
Another example with a one layer NN



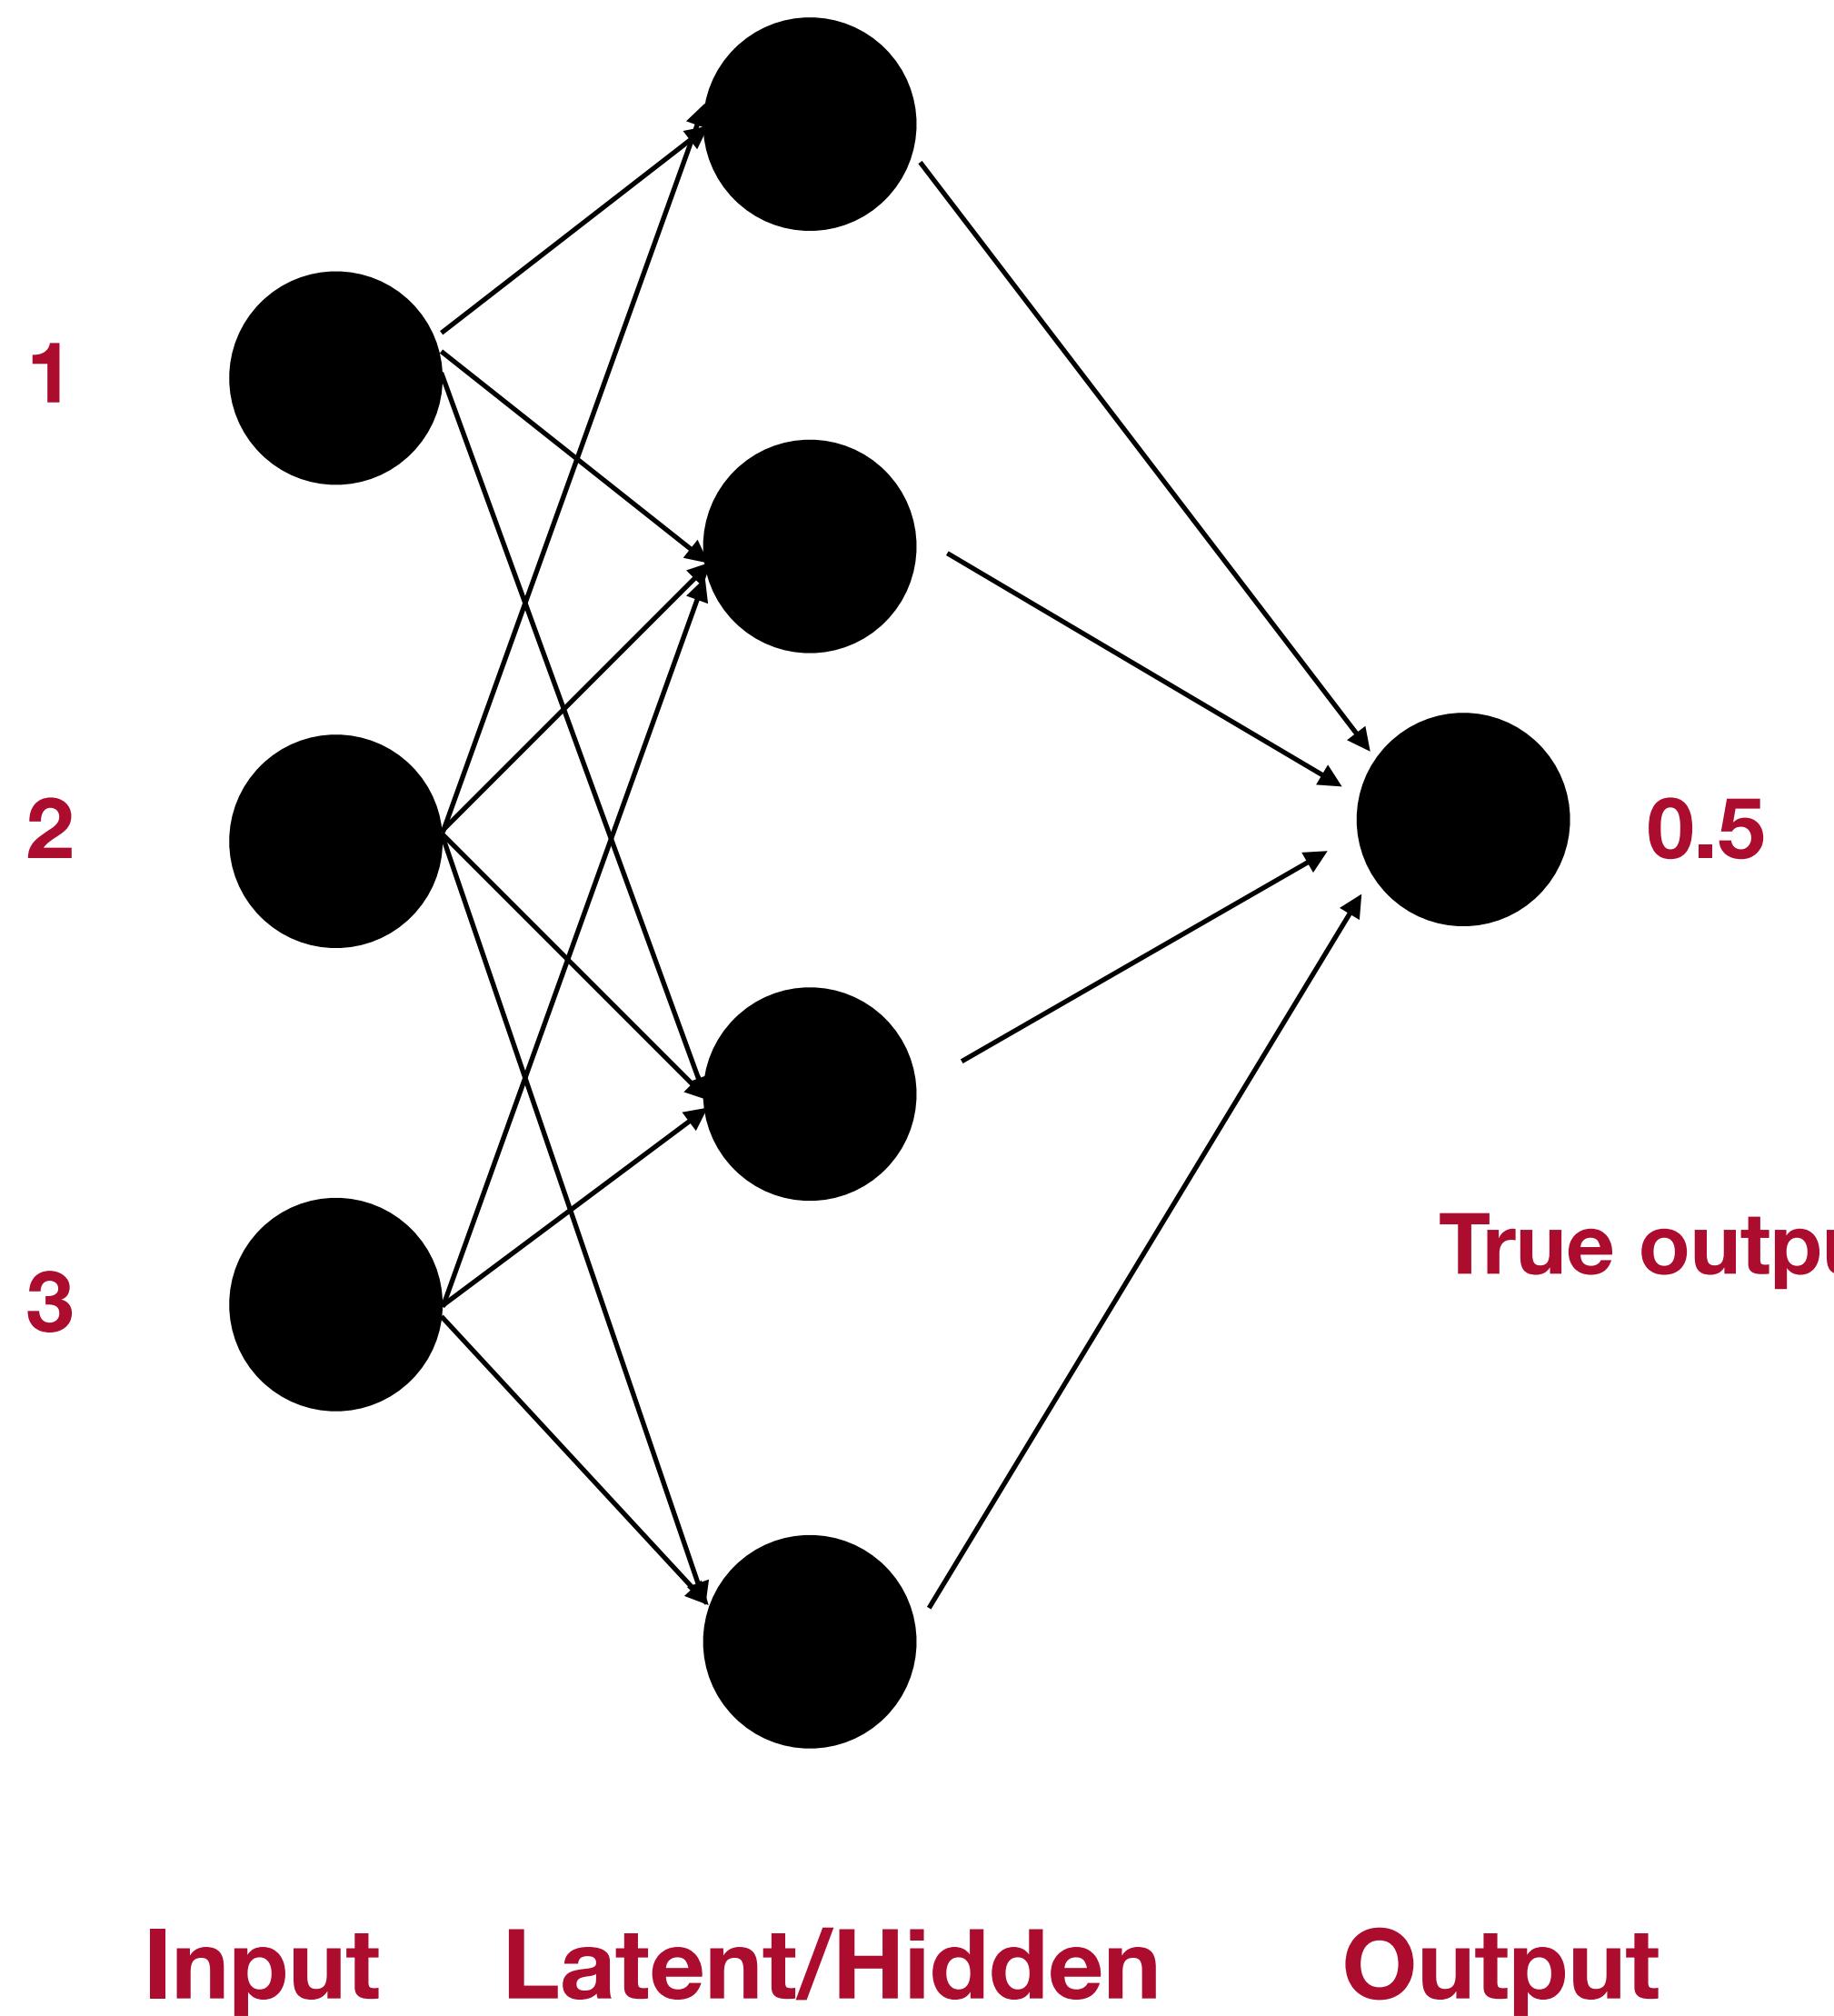
Another example with a one layer NN



Another example with a one layer NN

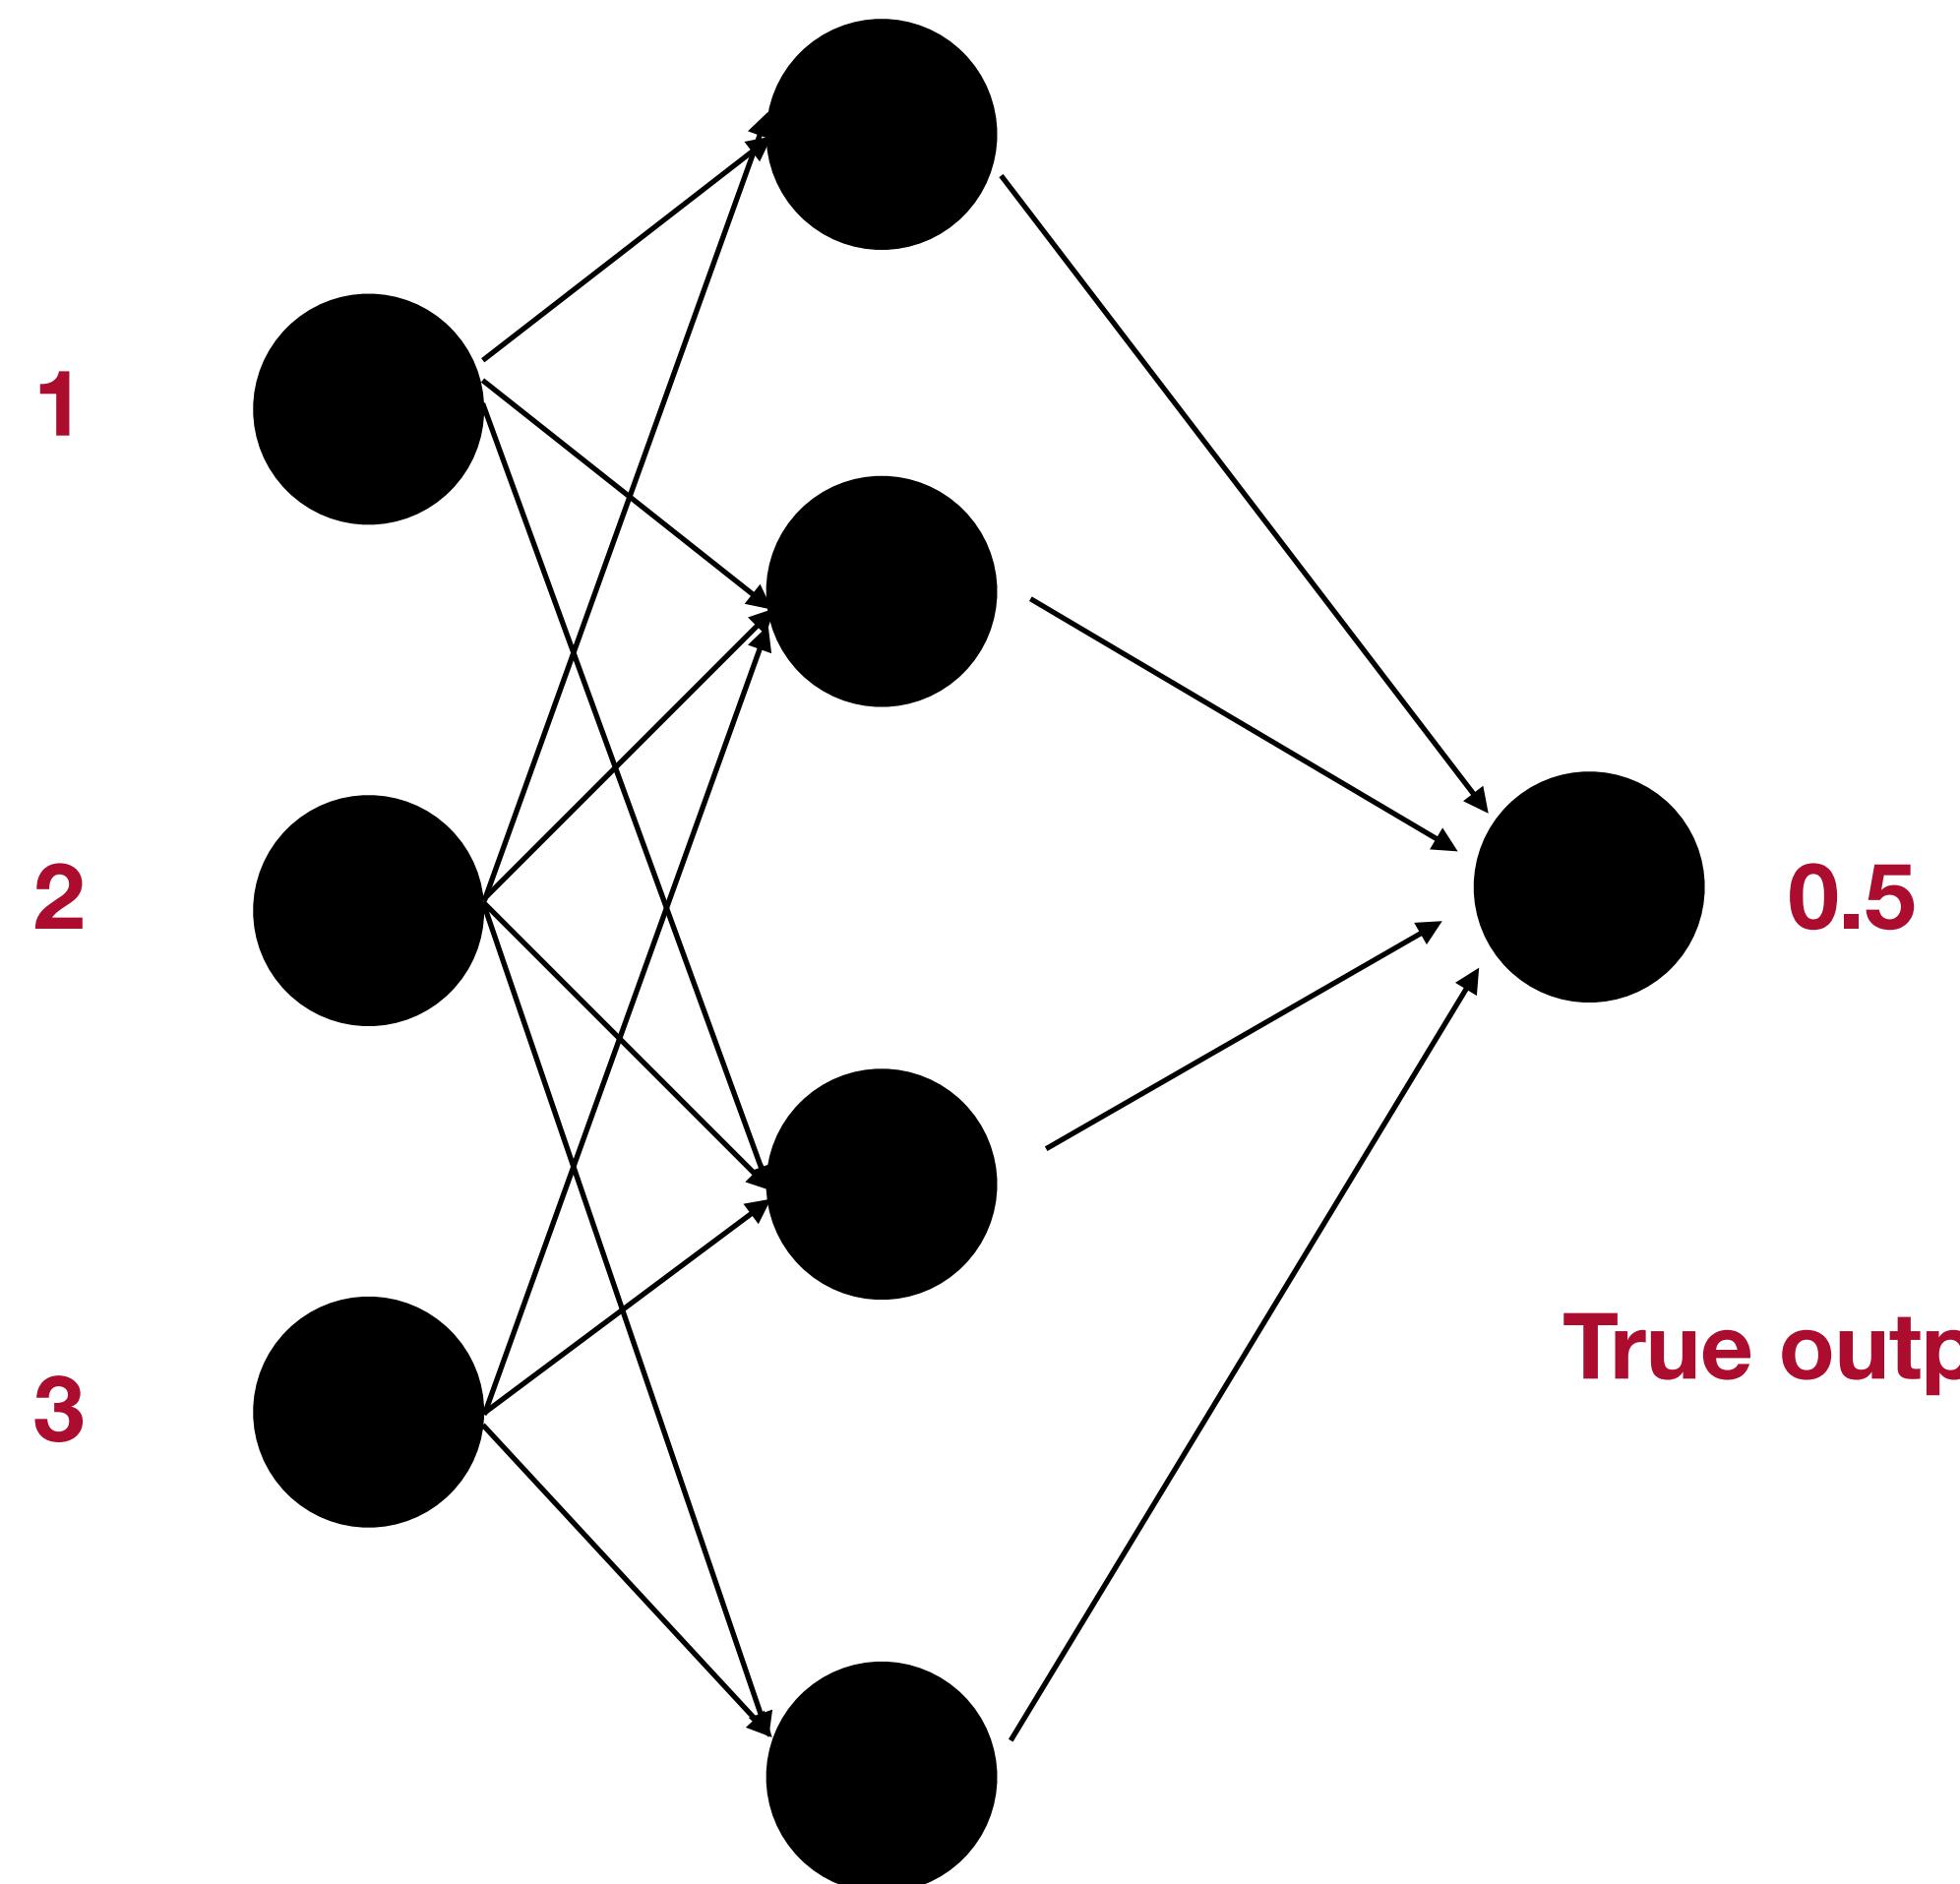


Another example with a one layer NN



Suppose that I give as input $\mathbf{X} = [1, 2, 3]$
and obtain as output $\hat{y} = 0.5$

Another example with a one layer NN

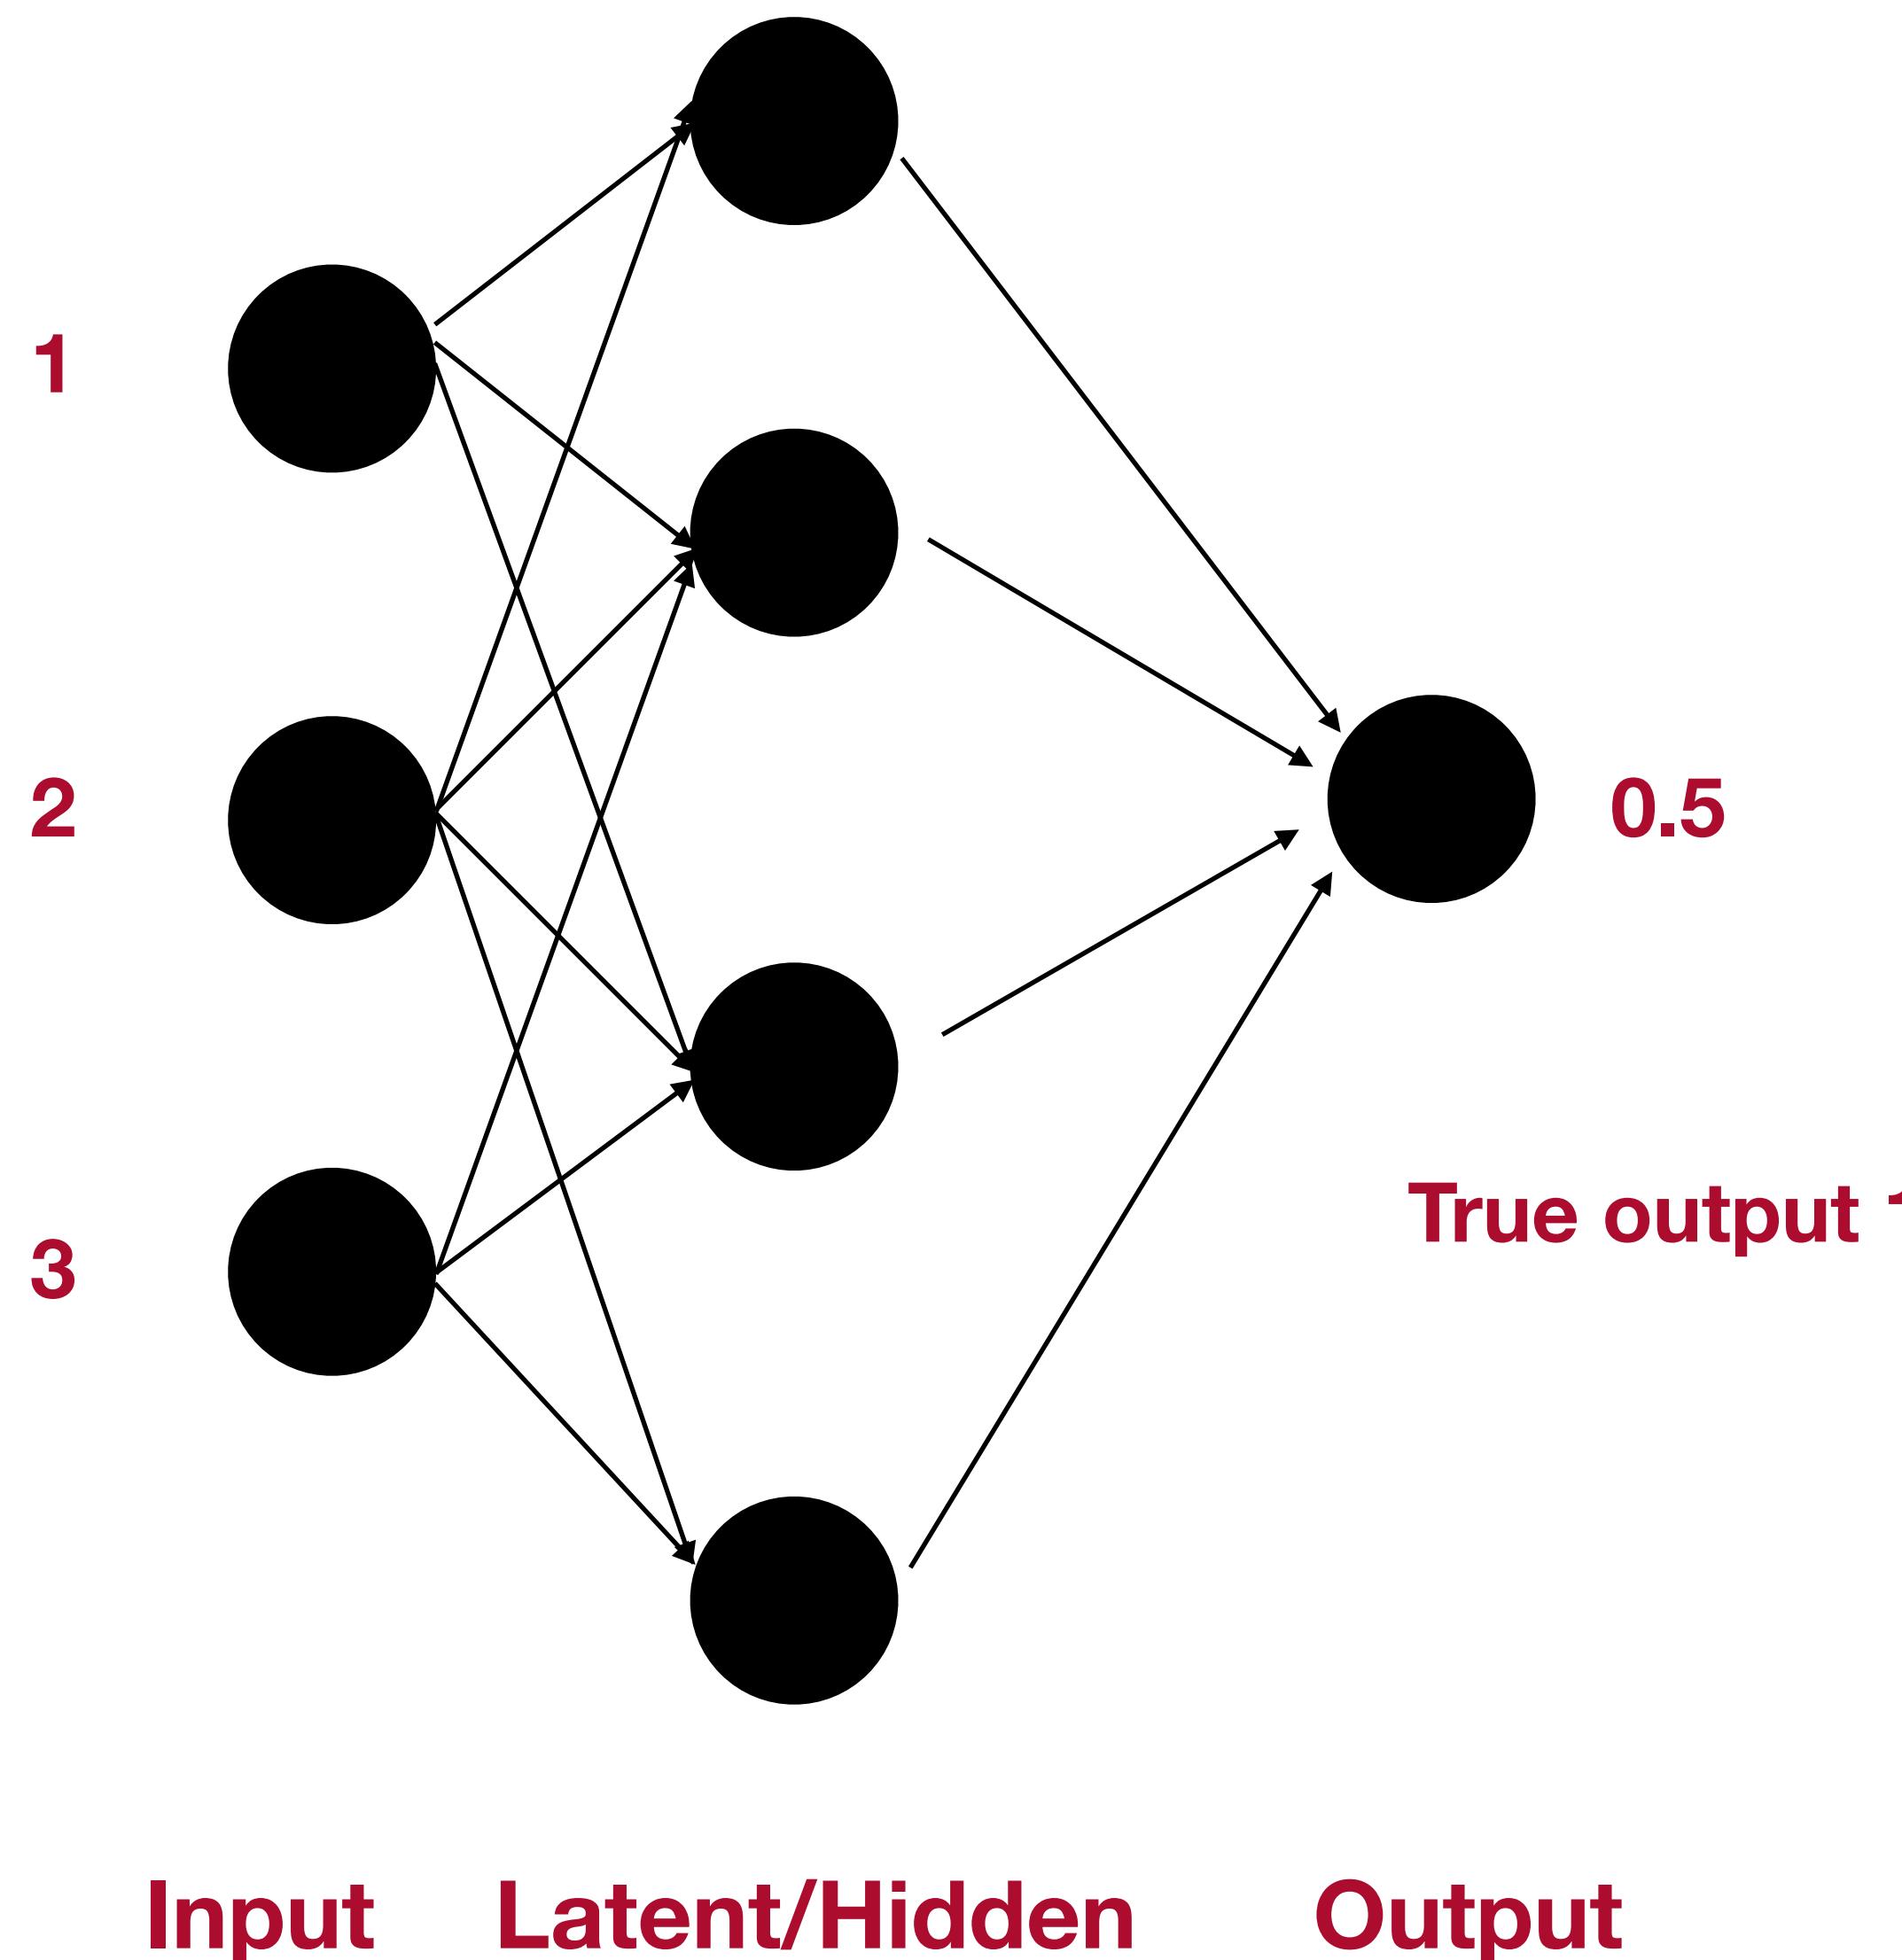


Suppose that I give as input $\mathbf{X} = [1, 2, 3]$
and obtain as output $\hat{y} = 0.5$

How to measure the quality of the prediction?

Input Latent/Hidden Output

Another example with a one layer NN

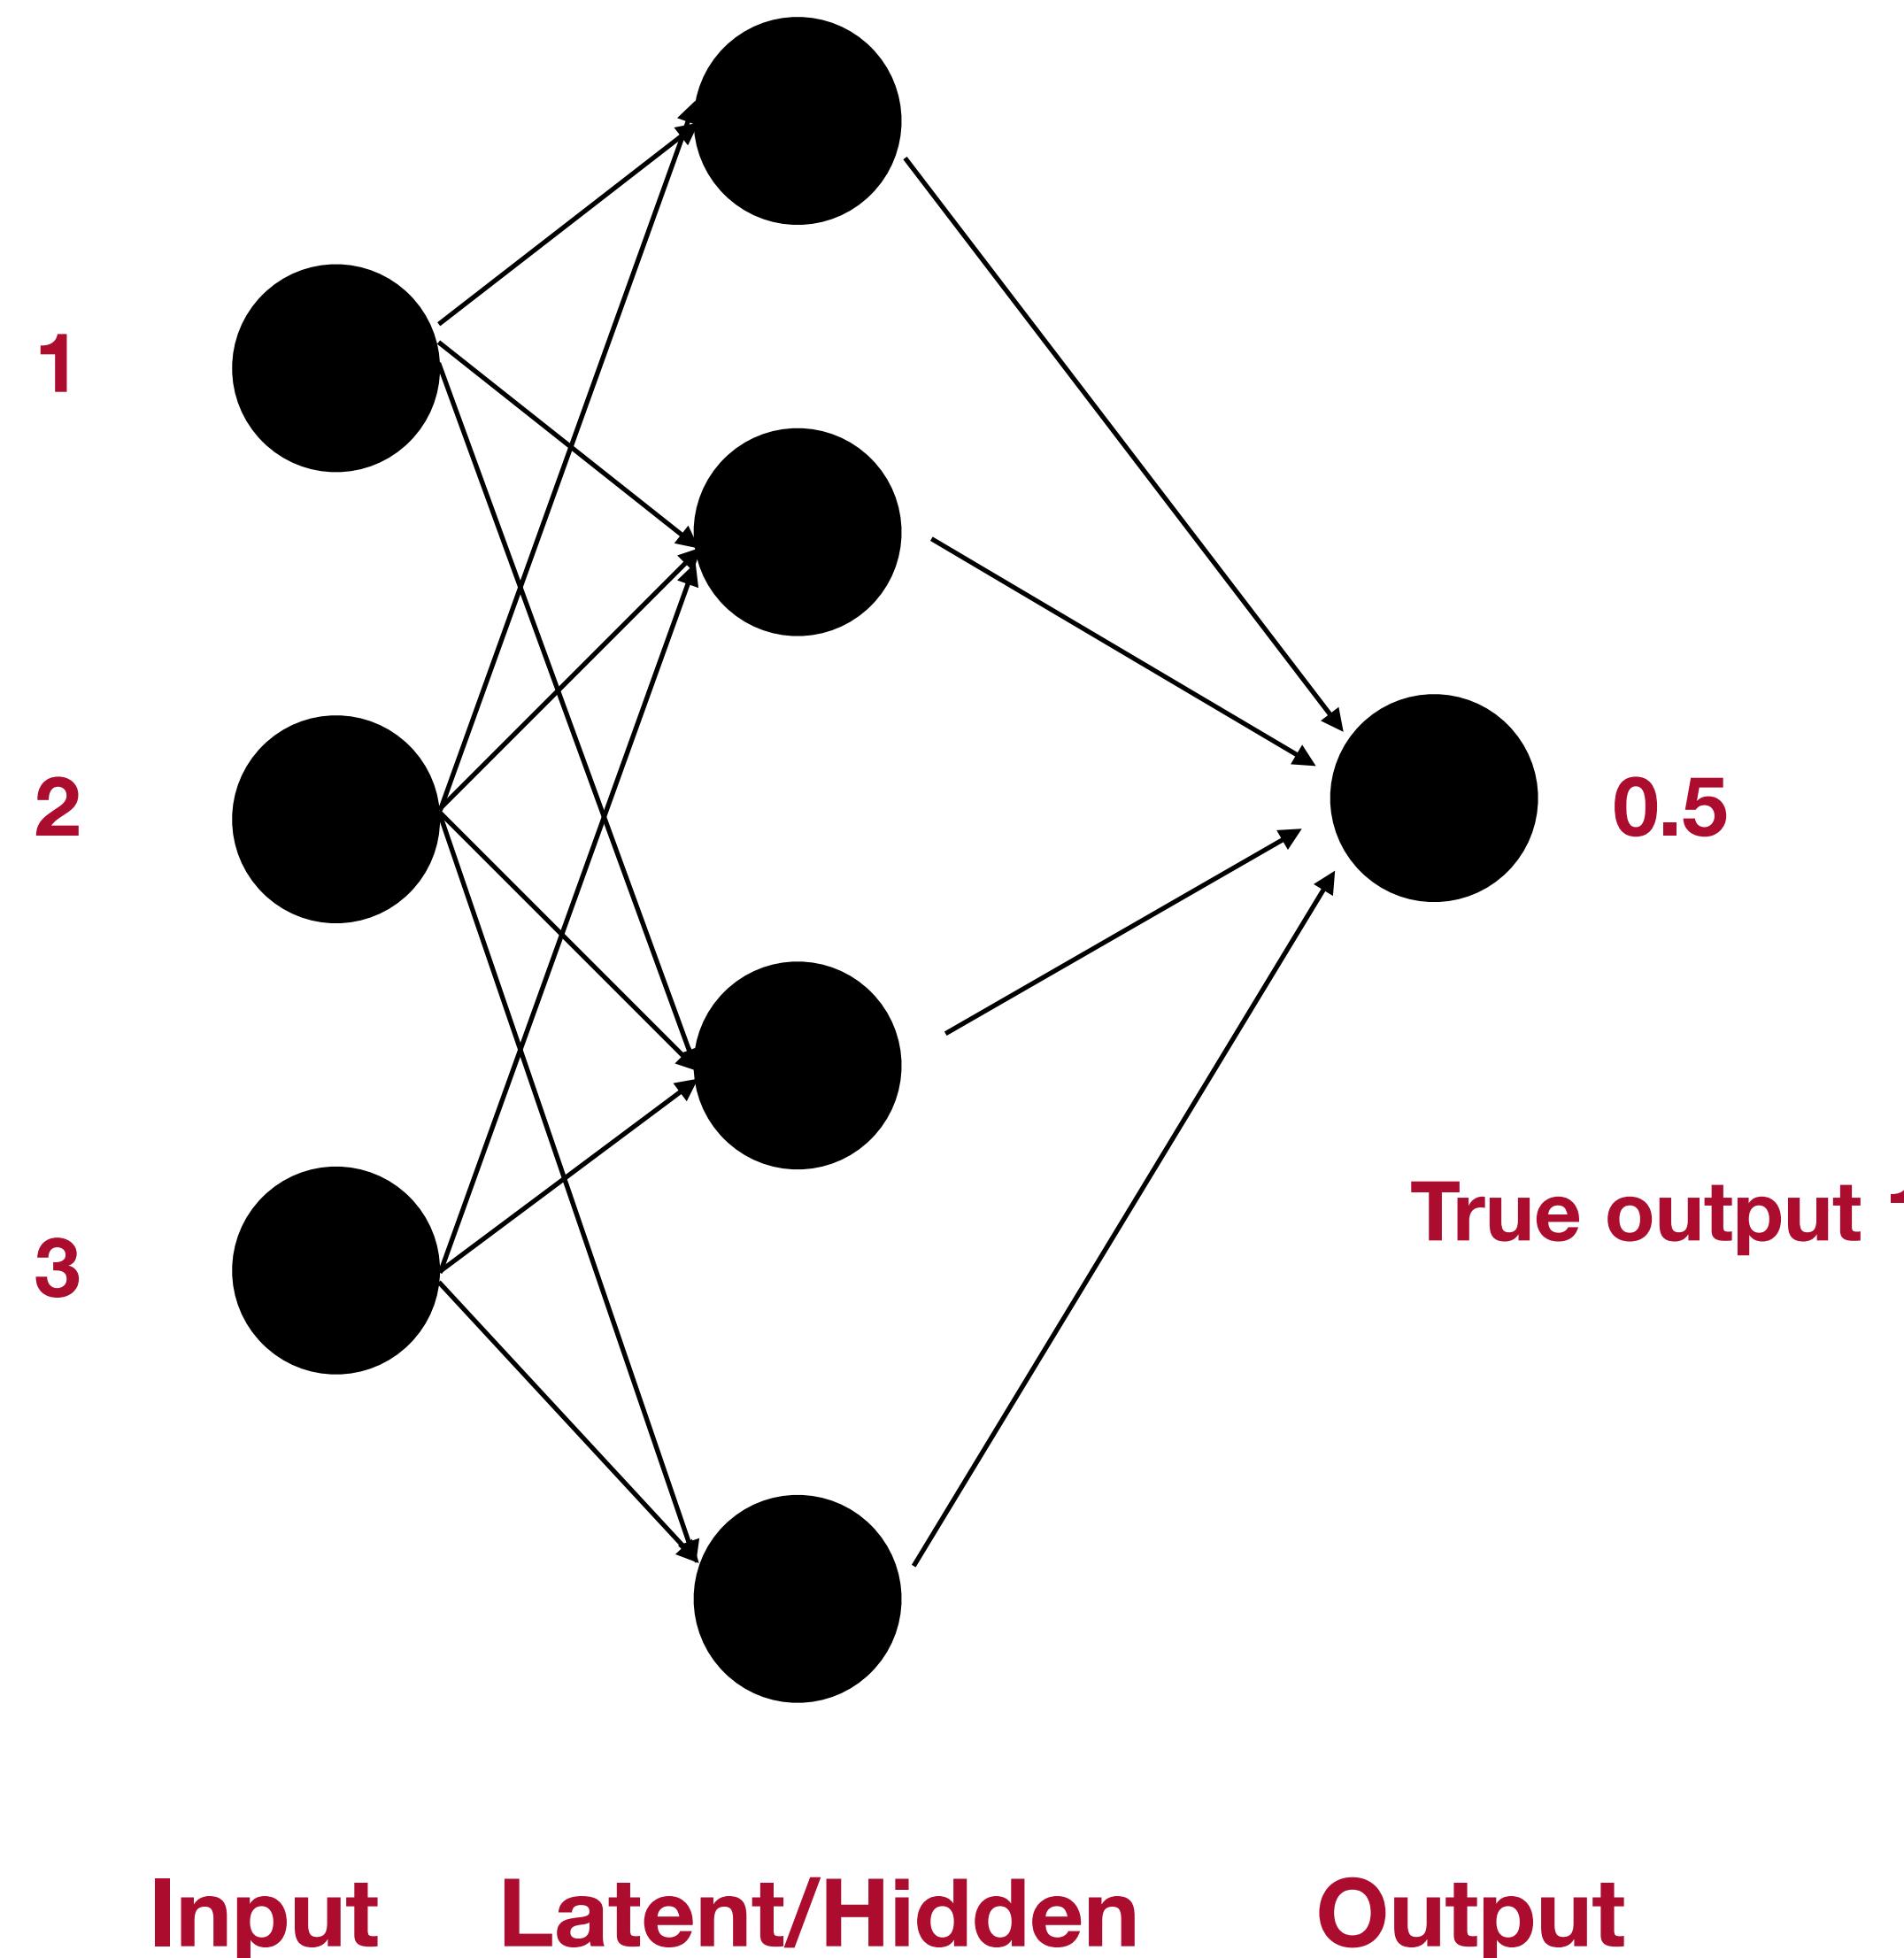


Suppose that I give as input $\mathbf{X} = [1, 2, 3]$
and obtain as output $\hat{y} = 0.5$

How to measure the quality of the prediction?

The loss of our network
will measure the cost of Incorrect predictions

Another example with a one layer NN



Suppose that I give as input $\mathbf{X} = [1, 2, 3]$
and obtain as output $\hat{y} = 0.5$

How to measure the quality of the prediction?

The loss of our network
will measure the cost of Incorrect predictions

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}, \mathbf{W}, \sigma)$$

A glance at existing loss functions !

A glance at existing loss functions !

In what follows the Neural Network will be denoted as f_θ

A glance at existing loss functions !

In what follows the Neural Network will be denoted as f_θ

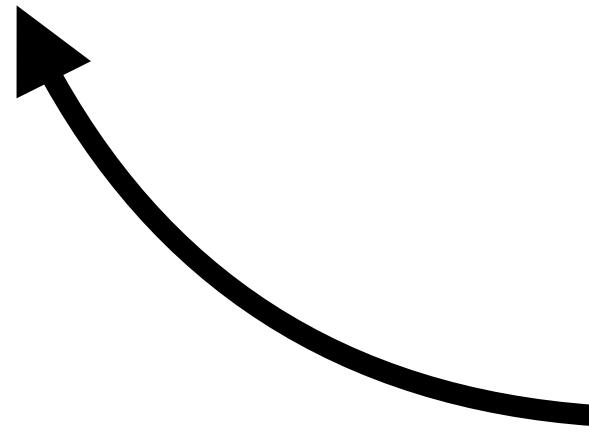
$\theta = [W_1, \dots, W_k]$ the set of weight matrices

A glance at existing loss functions !

In what follows the Neural Network will be denoted as f_θ

$$\theta = [W_1, \dots, W_k]$$

the set of weight matrices



each W_i corresponds to the weights of the layer i

A glance at existing loss functions !

In what follows the Neural Network will be denoted as f_θ

$$\theta = [W_1, \dots, W_k]$$

the set of weight matrices

each W_i corresponds to the weights of the layer i

f_θ is obtained via composition
of W_i and the activation functions

Empirical Losses

Neural Networks learn from data but so far we assumed no data

Neural Networks learn from data but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

Neural Networks learn from data but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

Neural Networks learn from data but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

each \mathbf{X}_i corresponds to
an input sample (image)

Neural Networks learn from data but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

each \mathbf{X}_i corresponds to
an input sample (image)

The empirical loss $J(\theta)$ measures the total loss over our entire dataset

Neural Networks learn from data but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

each \mathbf{X}_i corresponds to
an input sample (image)

The empirical loss $J(\theta)$ measures the total loss over our entire dataset

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_\theta(\mathbf{X}_i), y_i\right)$$

Neural Networks learn from data but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

each \mathbf{X}_i corresponds to
an input sample (image)

The empirical loss $J(\theta)$ measures the total loss over our entire dataset

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L} \left(f_\theta \left(\mathbf{X}_i \right), y_i \right)$$

Prediction \hat{y}_i

Neural Networks **learn from data** but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

each \mathbf{X}_i corresponds to
an input sample (image)

The empirical loss $J(\theta)$ measures the **total loss over our entire dataset**

Cost function
Empirical Risk
Objective function

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_\theta(\mathbf{X}_i), y_i\right)$$

Prediction \hat{y}_i

Neural Networks **learn from data** but so far we assumed no data

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

each y_i
corresponds to
label (dog/cat)

each \mathbf{X}_i corresponds to
an input sample (image)

The empirical loss $J(\theta)$ measures the **total loss over our entire dataset**

Cost function
Empirical Risk
Objective function

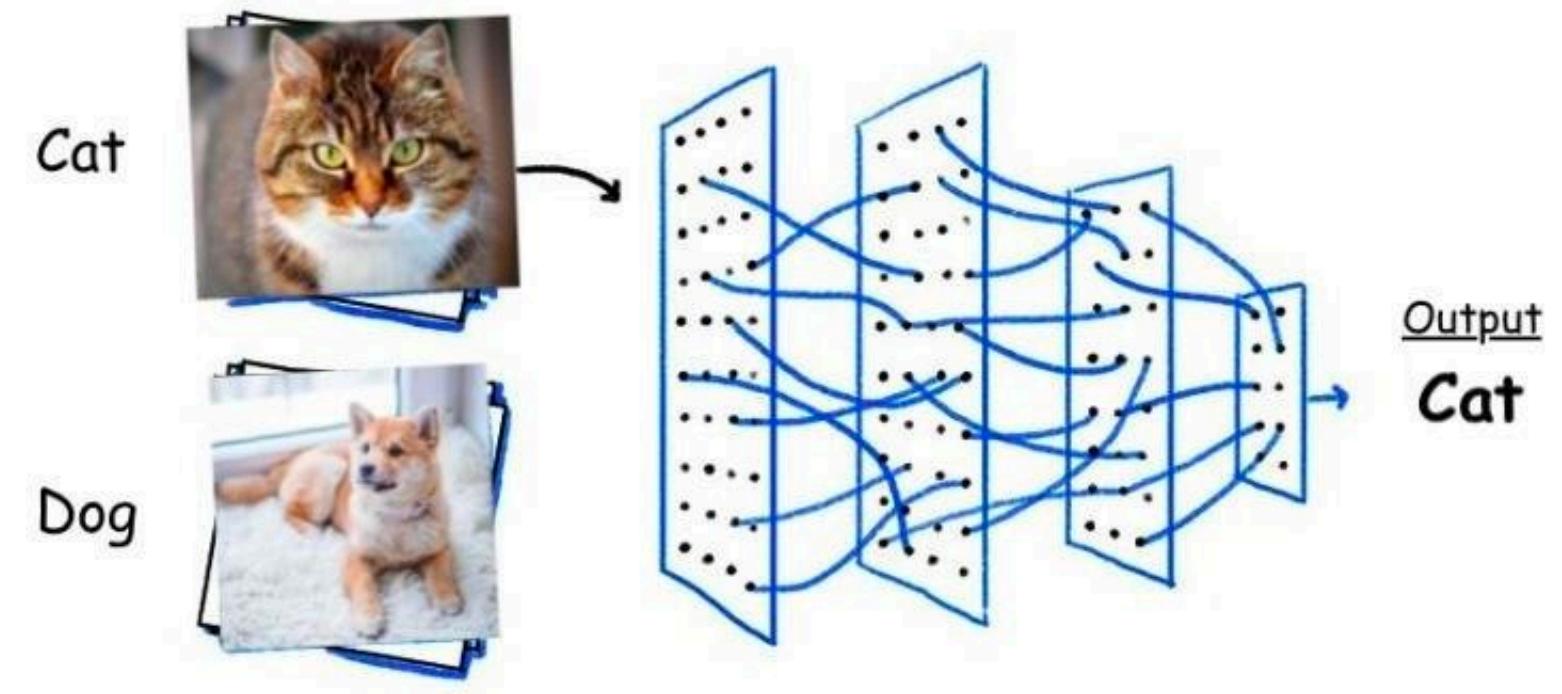
$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L} \left(f_\theta (\mathbf{X}_i), y_i \right)$$

Prediction \hat{y}_i

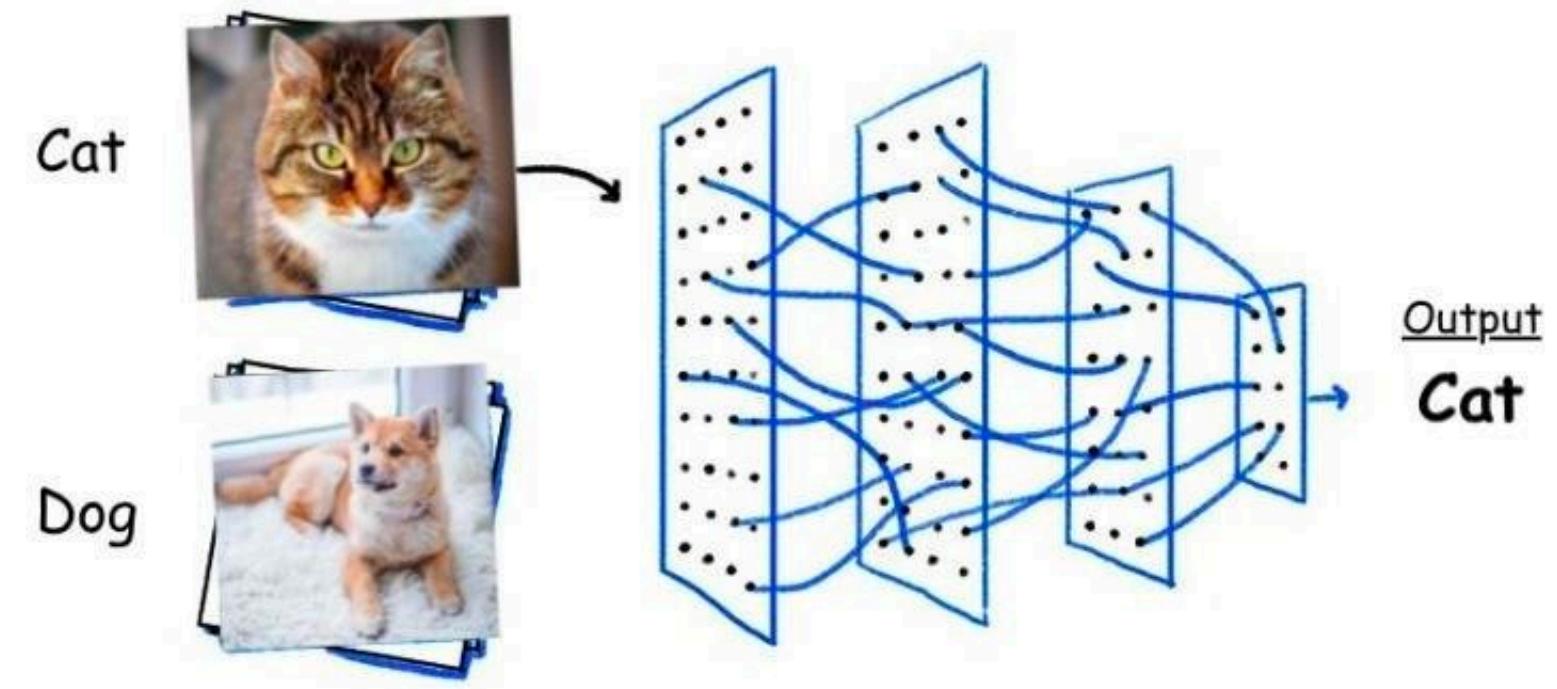
Each type of problem induces a **different loss function** (e.g. Classification, Regression)

Classification via Cross Entropy Losses

Classification via Cross Entropy Losses

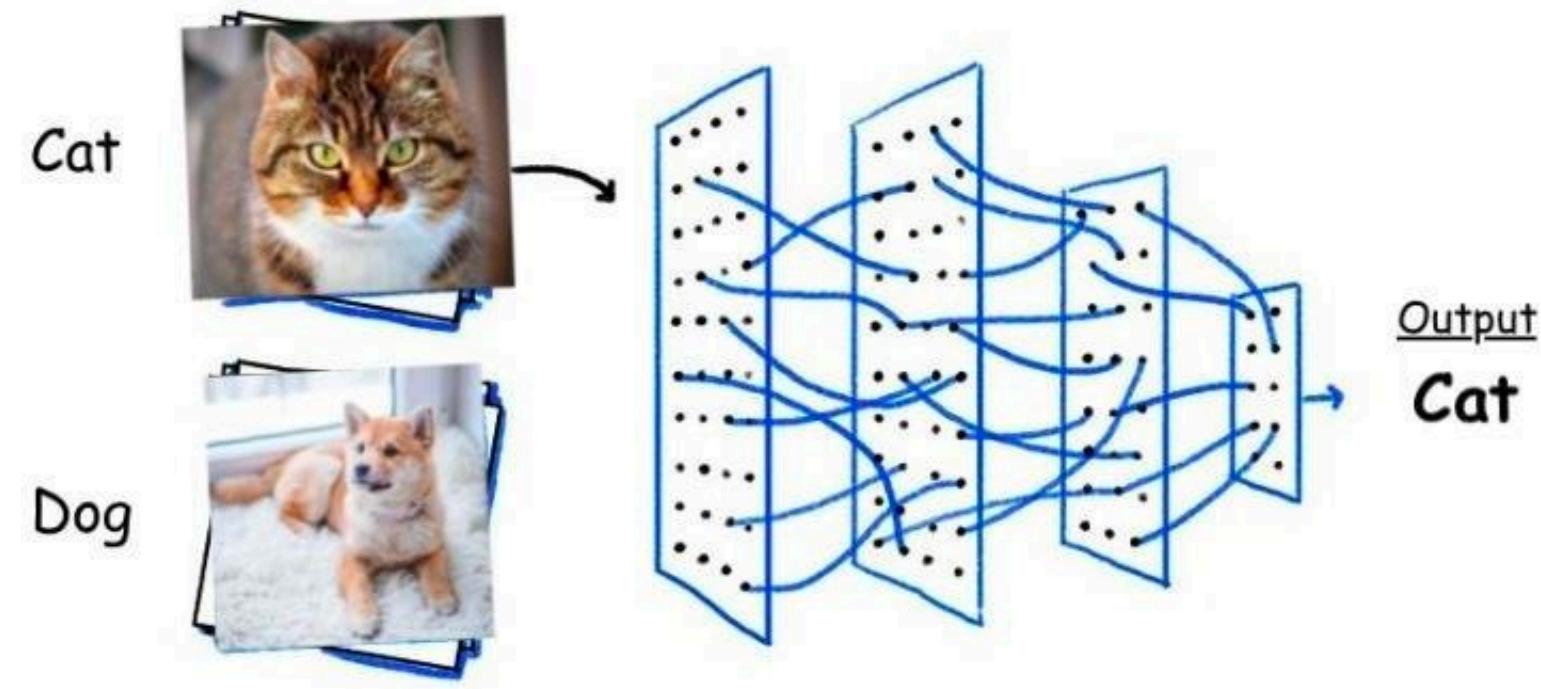


Classification via Cross Entropy Losses



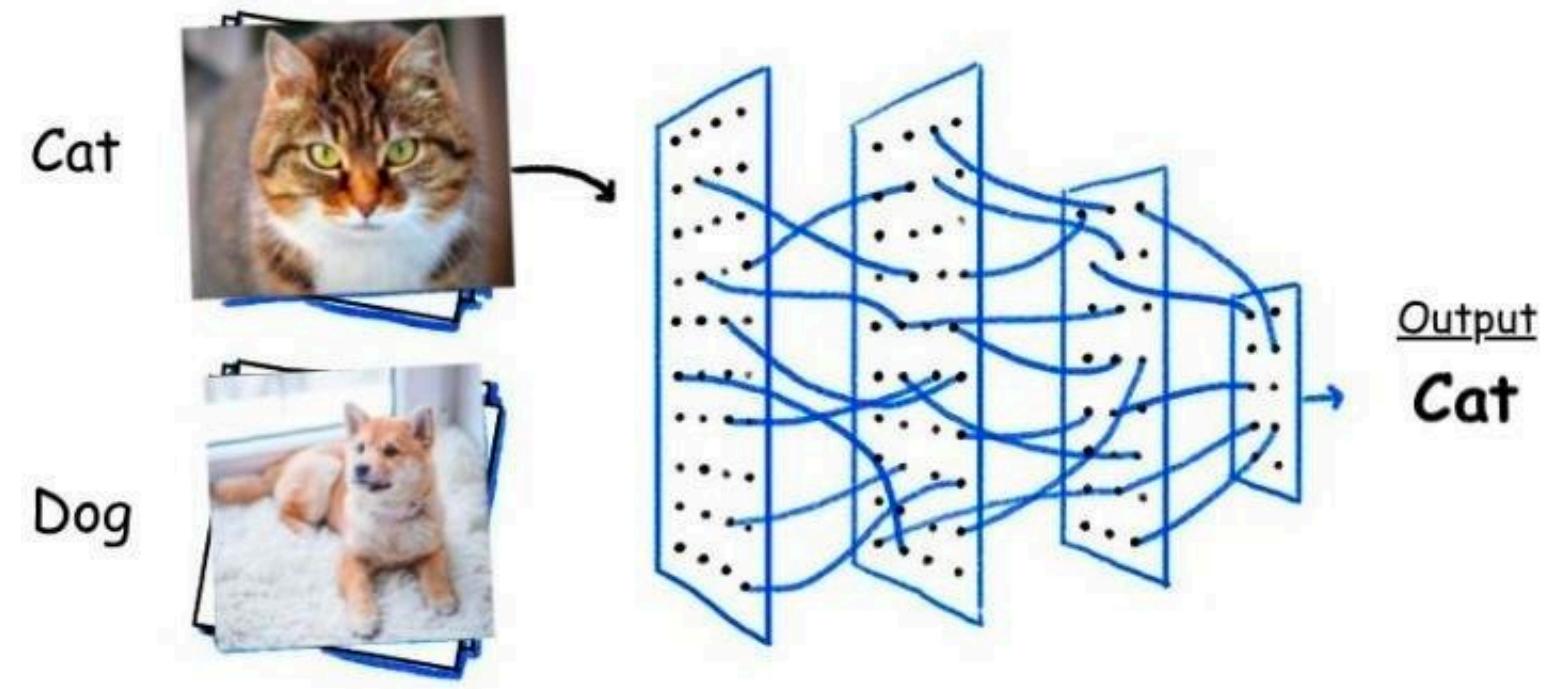
We want to classify an input picture as a cat or a dog

Classification via Cross Entropy Losses



We want to **classify** an input picture as a cat or a dog

Classification via Cross Entropy Losses

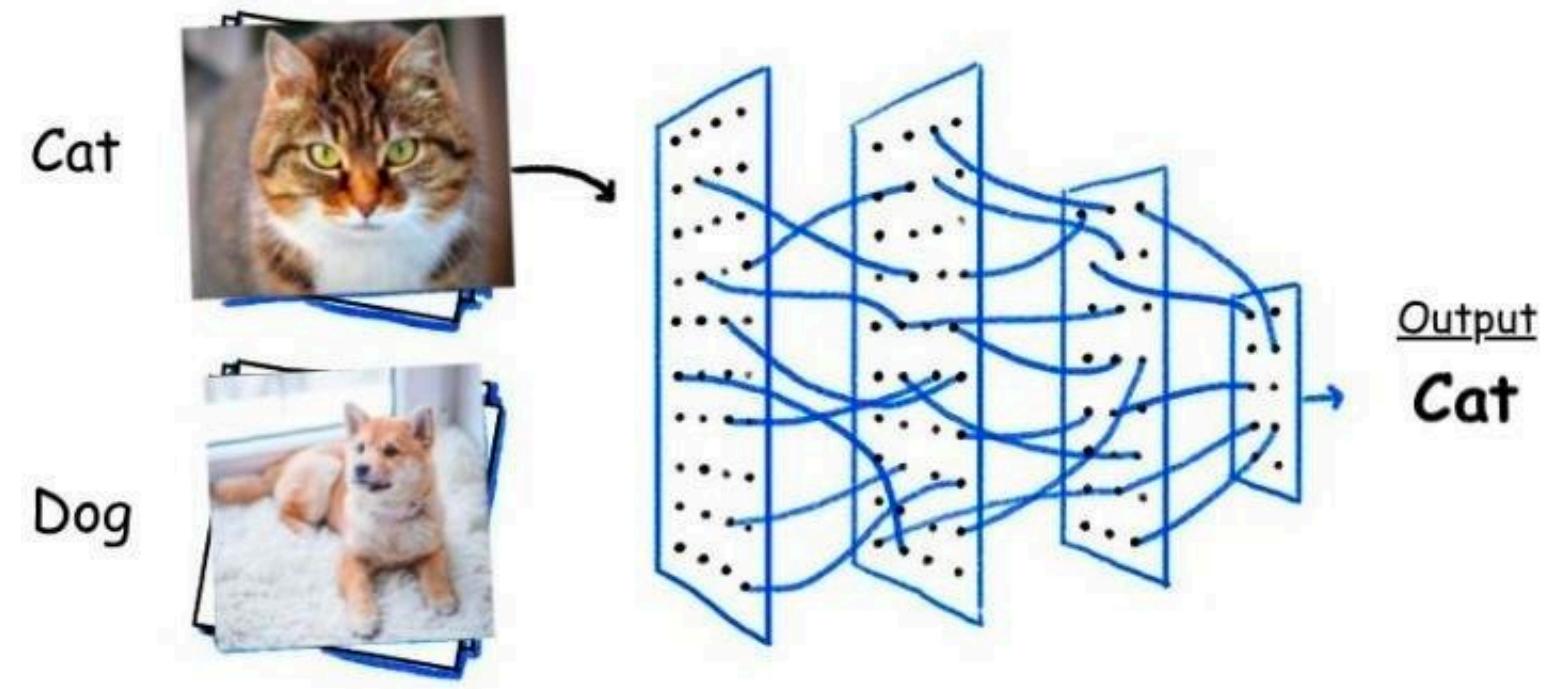


We want to **classify** an input picture as a cat or a dog

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}\left(\mathbf{X}^i\right), y^i\right)$$

—————>

Classification via Cross Entropy Losses



$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}\left(\mathbf{X}^i\right), y^i\right)$$

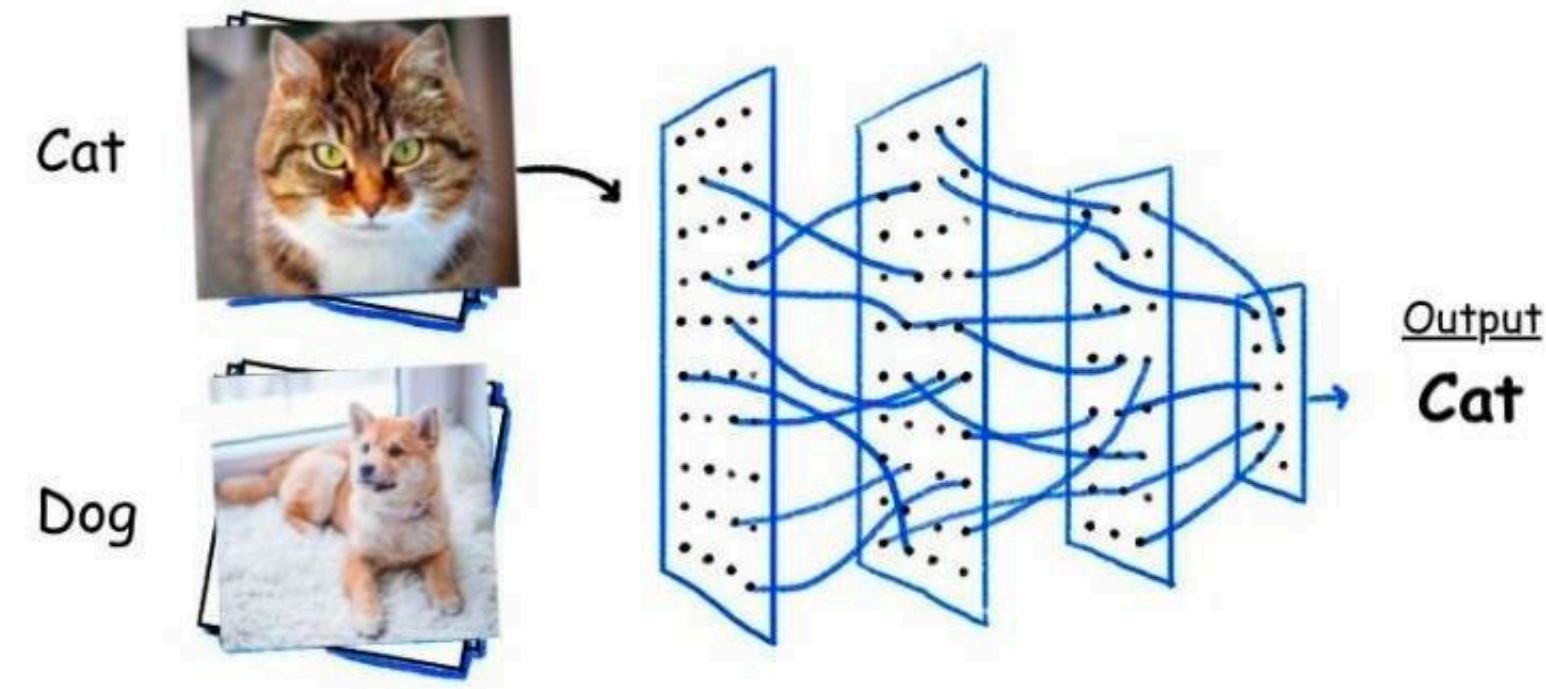


We want to **classify** an input picture as a cat or a dog

We will rely on the **binary cross entropy**

The **cross-entropy loss** can be used
with a classification model

Classification via Cross Entropy Losses



We want to **classify** an input picture as a cat or a dog

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}\left(\mathbf{X}^i\right), y^i\right)$$

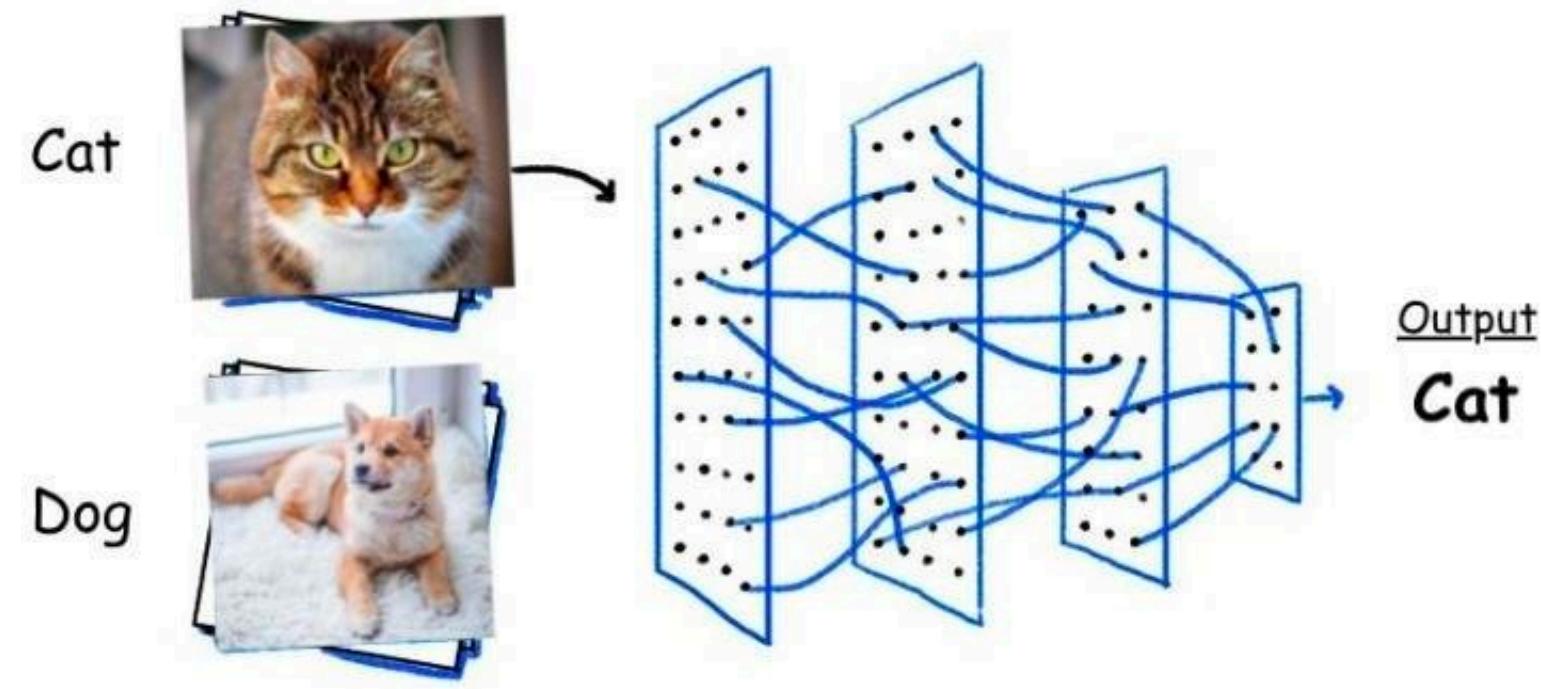


We will rely on the **binary cross entropy**

The **cross-entropy loss** can be used
with a classification model

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_{\theta}\left(\mathbf{X}^i\right) \right) + (1 - y_i) \times \log \left(1 - \left(f_{\theta}\left(\mathbf{X}^i\right) \right) \right)$$

Classification via Cross Entropy Losses



We want to **classify** an input picture as a cat or a dog

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}\left(\mathbf{X}^i\right), y^i\right)$$



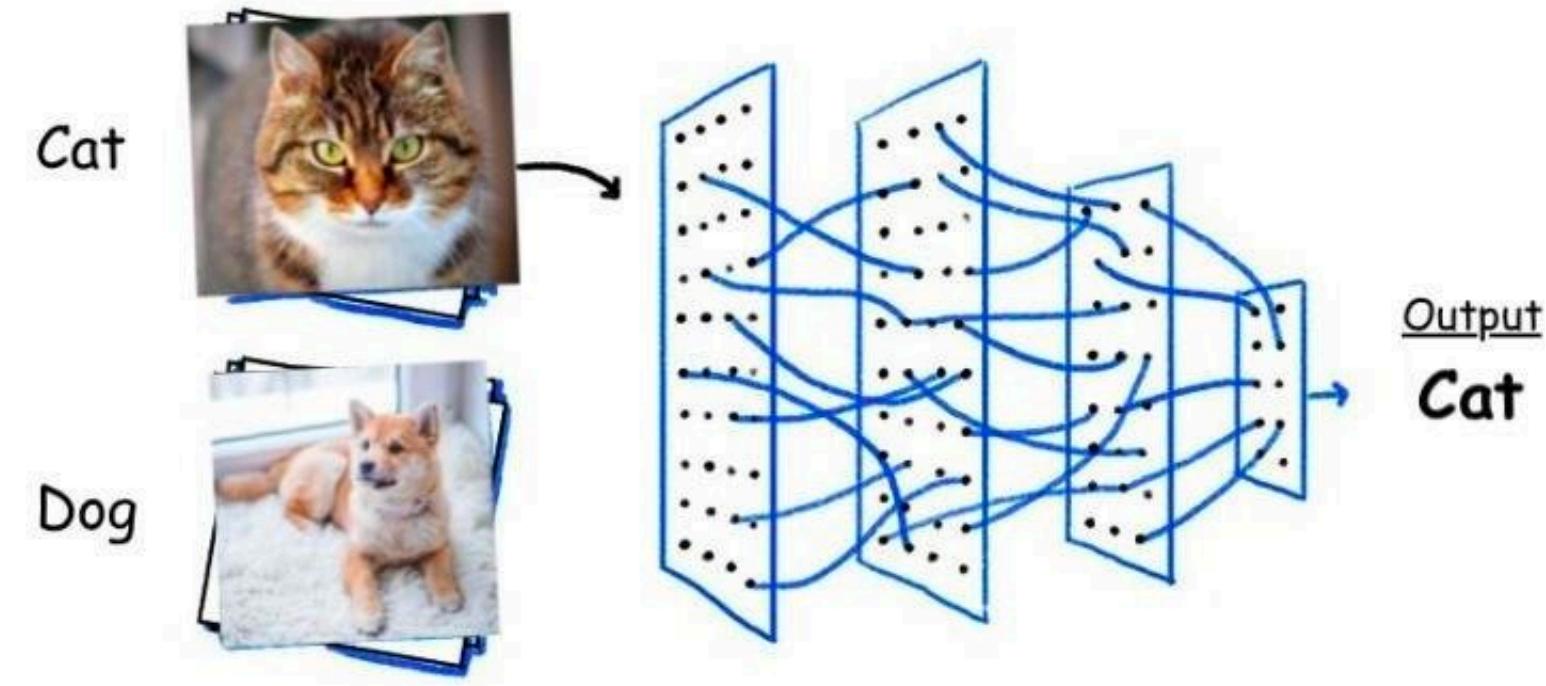
We will rely on the **binary cross entropy**

The **cross-entropy loss** can be used
with a classification model

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_{\theta}\left(\mathbf{X}^i\right) \right) + (1 - y_i) \times \log \left(1 - \left(f_{\theta}\left(\mathbf{X}^i\right) \right) \right)$$

True label

Classification via Cross Entropy Losses



We want to **classify** an input picture as a cat or a dog

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}\left(\mathbf{X}^i\right), y^i\right)$$



We will rely on the **binary cross entropy**

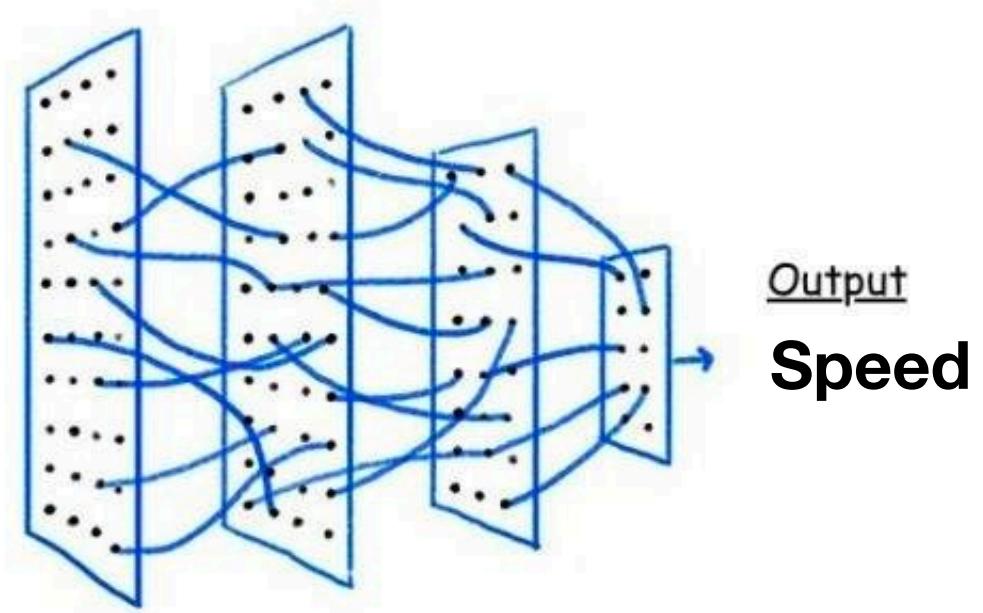
The **cross-entropy loss** can be used
with a classification model

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_{\theta}\left(\mathbf{X}^i\right) \right) + (1 - y_i) \times \log \left(1 - \left(f_{\theta}\left(\mathbf{X}^i\right) \right) \right)$$

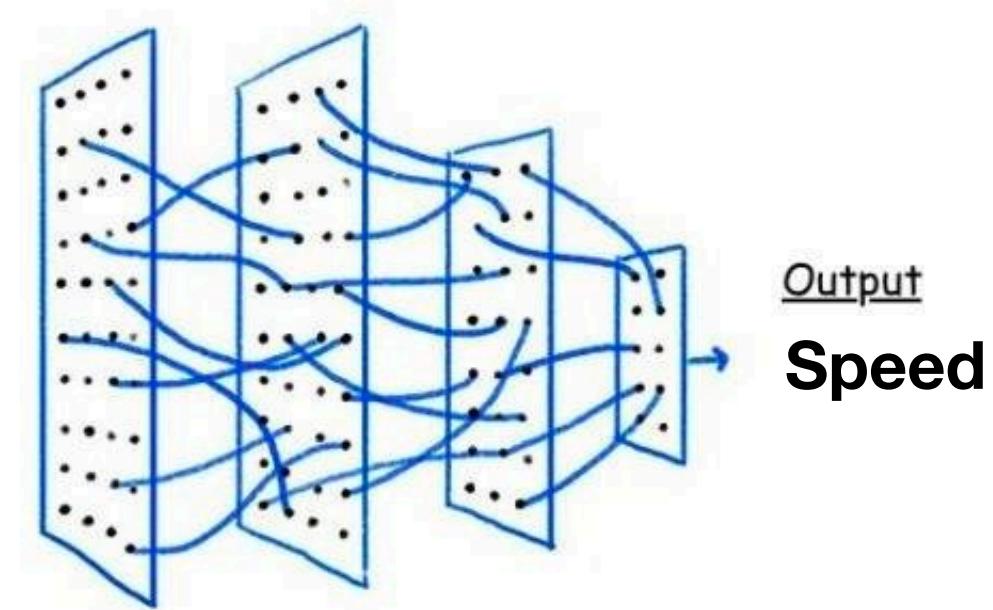
Prediction \hat{y}_i True label

Regression via Mean Square Error or Mean Absolute Error

Regression via Mean Square Error or Mean Absolute Error

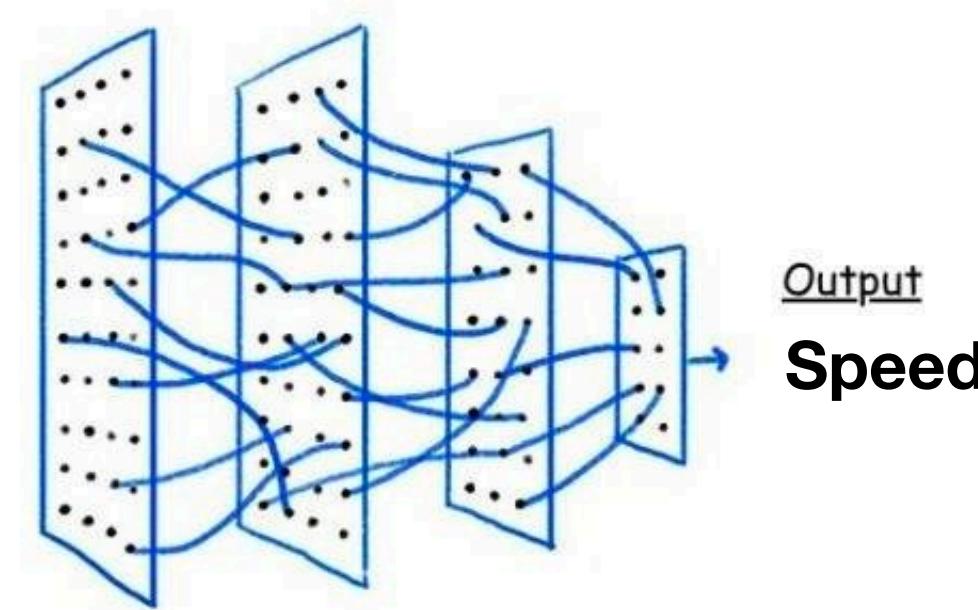


Regression via Mean Square Error or Mean Absolute Error



We want to predict the maximum speed of the car

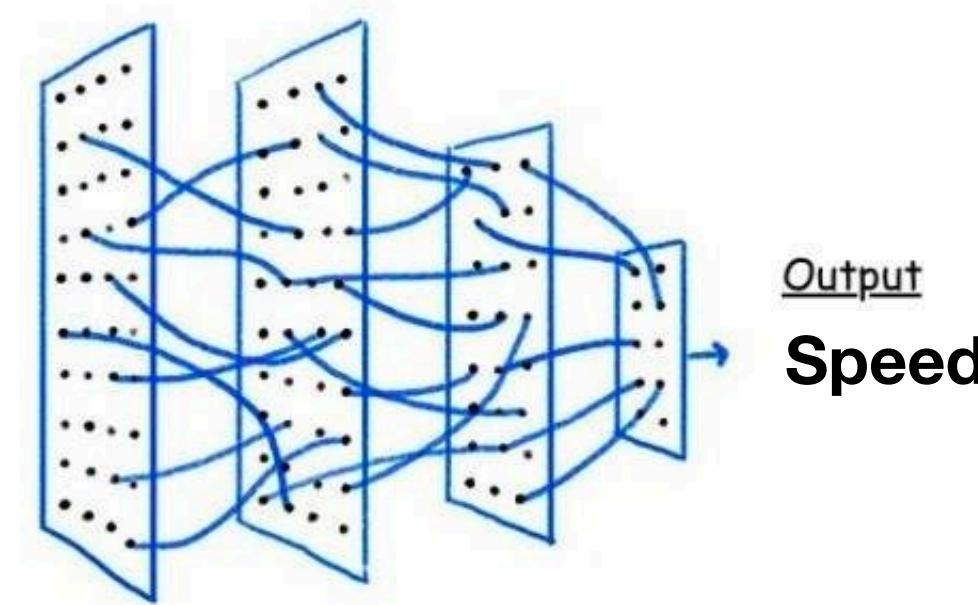
Regression via Mean Square Error or Mean Absolute Error



Real value!

We want to predict the **maximum speed of the car**

Regression via Mean Square Error or Mean Absolute Error

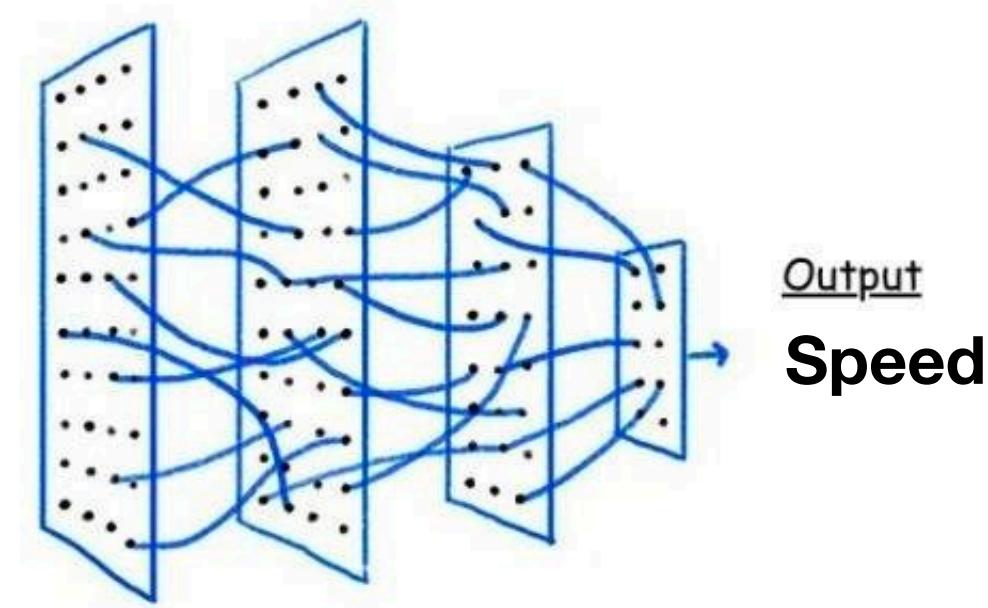


Real value!

We want to predict the **maximum speed of the car**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L} \left(f_{\theta} (\mathbf{X}^i), y^i \right)$$

Regression via Mean Square Error or Mean Absolute Error



$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}\left(\mathbf{X}^i\right), y^i\right)$$



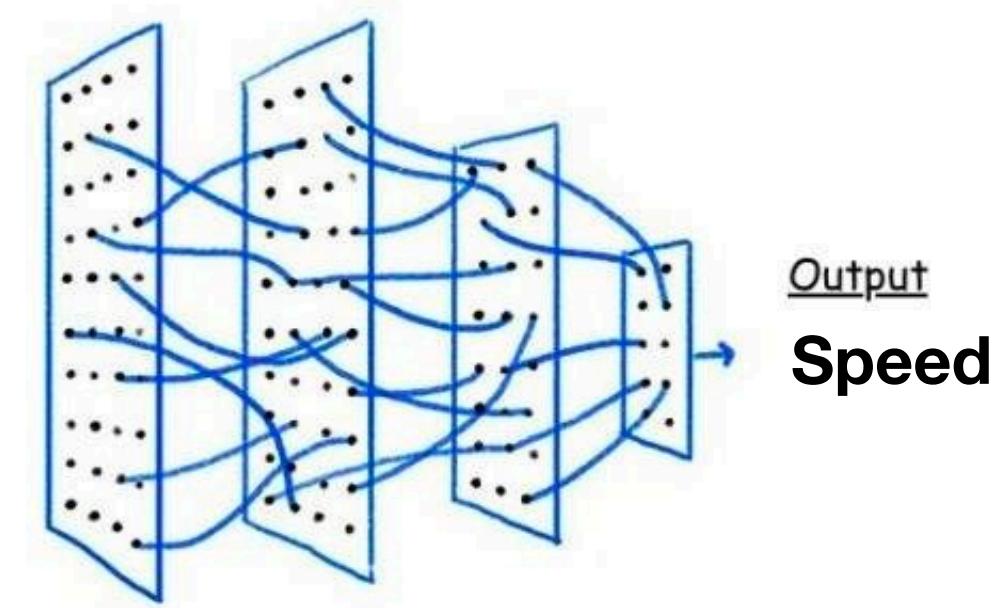
Real value!

We want to predict the **maximum speed of the car**

We will rely on the mean square error

The **MAE or MSE loss** can be used
with a regression model

Regression via Mean Square Error or Mean Absolute Error



Real value!

We want to predict the **maximum speed of the car**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}(\mathbf{X}^i), y^i\right)$$



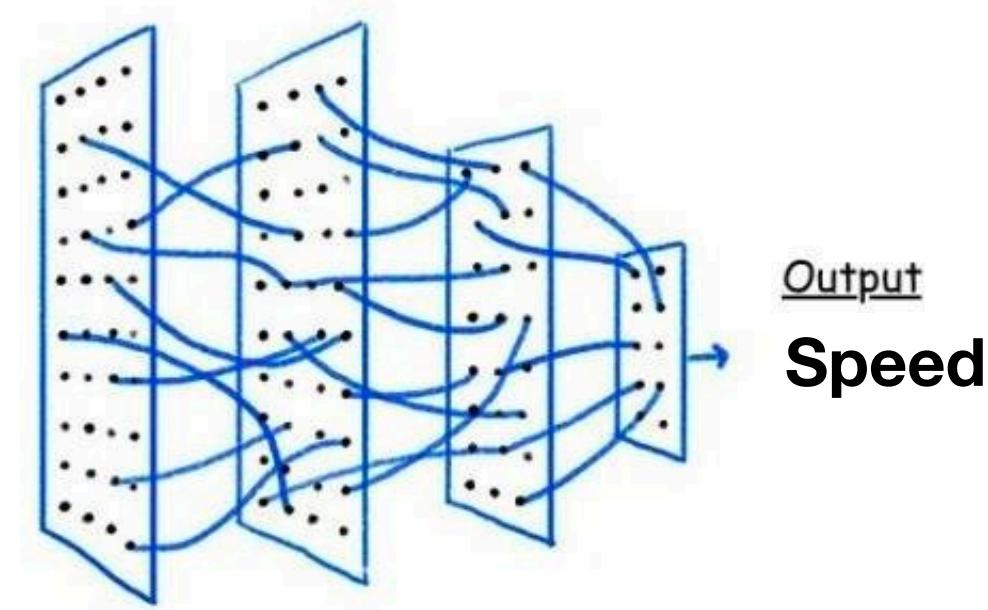
We will rely on the mean square error

The **MAE or MSE loss** can be used
with a regression model

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_{\theta}(\mathbf{X}^i) \right) \right)^2$$

Mean Square Error

Regression via Mean Square Error or Mean Absolute Error



Real value!

We want to predict the **maximum speed of the car**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}\left(f_{\theta}(\mathbf{X}^i), y^i\right)$$



We will rely on the mean square error

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_{\theta}(\mathbf{X}^i) \right) \right)^2$$

Mean Square Error

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_{\theta}(\mathbf{X}^i) \right) \right|$$

Mean Absolute Error

Let's sum up so far

1. The basic bloc to build a Neural Network

We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

2. The role of activation functions

They complexity the neural network and allow to fit more complex data

They are added at the end of each layer and there is a large variety of them

Let's sum up so far

1. The basic bloc to build a Neural Network

We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

2. The role of activation functions

They complexity the neural network and allow to fit more complex data

They are added at the end of each layer and there is a large variety of them

3. The cost function

It allows to measure how « far » is the neural network from the reality

How to perform the « learning phase »?

How to optimize the weights (learning algorithm)

How to optimize the weights (learning algorithm)

We want to find the network's weight that achieve the lowest cost/loss

How to optimize the weights (learning algorithm)

We want to find the network's weight that achieve the lowest cost/loss

$$\theta^* = \operatorname{argmin}_{\theta} J(\theta)$$

How to optimize the weights (learning algorithm)

We want to find the network's weight that achieve the lowest cost/loss

$$\theta^* = \operatorname{argmin}_{\theta} J(\theta)$$

Remember: our loss is a function of the network weights $\theta = [W_1, \dots, W_k]$

How to optimize the weights (learning algorithm)

We want to find the network's weight that achieve the lowest cost/loss

$$\theta^* = \operatorname{argmin}_{\theta} J(\theta)$$

Remember: our loss is a function of the network weights $\theta = [W_1, \dots, W_k]$

We will use gradient descent to find the best weights

Loss Optimization via Gradient Descent

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: **Find the best neural network weights**

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: **Find the best neural network weights**

1. Initialize θ

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: **Find the best neural network weights**

1. Initialize θ

Usually sampled from a Gaussian distribution

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: **Find the best neural network weights**

1. Initialize θ

Usually sampled from a Gaussian distribution

2. While it has not converged

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: **Find the best neural network weights**

1. Initialize θ

Usually sampled from a Gaussian distribution

2. While it has not converged

Compute the gradient $\nabla_\theta J(\theta)$

This operation is called back-propagation

Loss Optimization via Gradient Descent

Requirements: **Gradient Descent can be done if we can compute the derivative of $J(\theta)$**

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: **Find the best neural network weights**

1. Initialize θ

Usually sampled from a Gaussian distribution

2. While it has not converged

Compute the gradient $\nabla_\theta J(\theta)$

This operation is called back-propagation

Update the weights $\theta = \theta - \lambda \nabla_\theta J(\theta)$

This operation is called gradient update

Loss Optimization via Gradient Descent

Requirements: Gradient Descent can be done if we can compute the derivative of $J(\theta)$

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: Find the best neural network weights

1. Initialize θ

Usually sampled from a Gaussian distribution

2. While it has not converged

Compute the gradient $\nabla_\theta J(\theta)$

This operation is called back-propagation

λ is the learning rate

Update the weights $\theta = \theta - \lambda \nabla_\theta J(\theta)$

This operation is called gradient update

Loss Optimization via Gradient Descent

Requirements: Gradient Descent can be done if we can compute the derivative of $J(\theta)$

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i - \left(f_\theta(\mathbf{X}^i) \right) \right)^2 \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N \left| y_i - \left(f_\theta(\mathbf{X}^i) \right) \right| \quad J(\theta) = \frac{1}{N} \sum_{i=1}^N y_i \times \log \left(f_\theta(\mathbf{X}^i) \right) + (1 - y_i) \times \log \left(1 - \left(f_\theta(\mathbf{X}^i) \right) \right)$$

Goal: Find the best neural network weights

1. Initialize θ

Usually sampled from a Gaussian distribution

2. While it has not converged

Compute the gradient $\nabla_\theta J(\theta)$

This operation is called back-propagation

λ is the learning rate

Update the weights $\theta = \theta - \lambda \nabla_\theta J(\theta)$

This operation is called gradient update

3. Return the final weights θ

θ should be closed to θ^*

Some remarks

About Backpropagation

This step compute the gradient of the error w.r.t to the weights

This step is compute intensive. This step is usually done on GPU

For most of the Neural Network you won't be able to write the close form solution.

About Backpropagation

This step compute **the gradient of the error w.r.t to the weights**

This step is compute intensive. This step is usually done on **GPU**

For most of the Neural Network **you won't be able to write the close form solution.**

About the **optimization algorithm**

Finding a **good learning rate λ** is crucial in practice.

There exists many alternative to Gradient Descent (e.g. SGD, Adam, RMSProp)

About Backpropagation

This step compute the gradient of the error w.r.t to the weights

This step is compute intensive. This step is usually done on GPU

For most of the Neural Network you won't be able to write the close form solution.

About the optimization algorithm

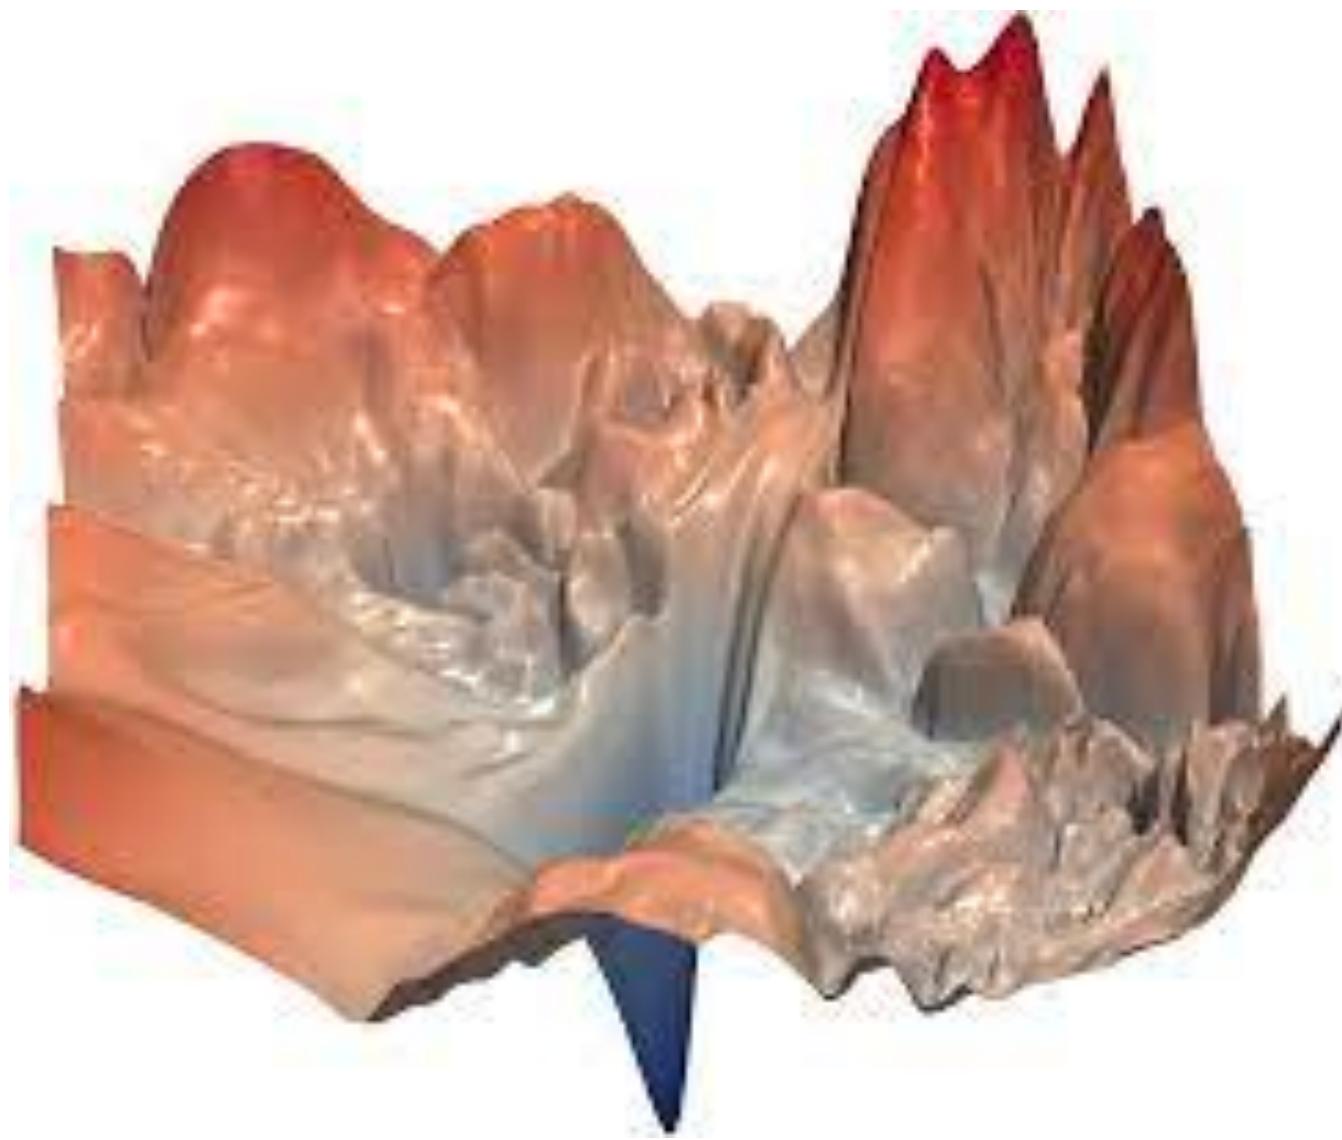
Finding a good learning rate λ is crucial in practice.

There exists many alternative to Gradient Descent (e.g. SGD, Adam, RMSProp)

Speed, speed, speed.....

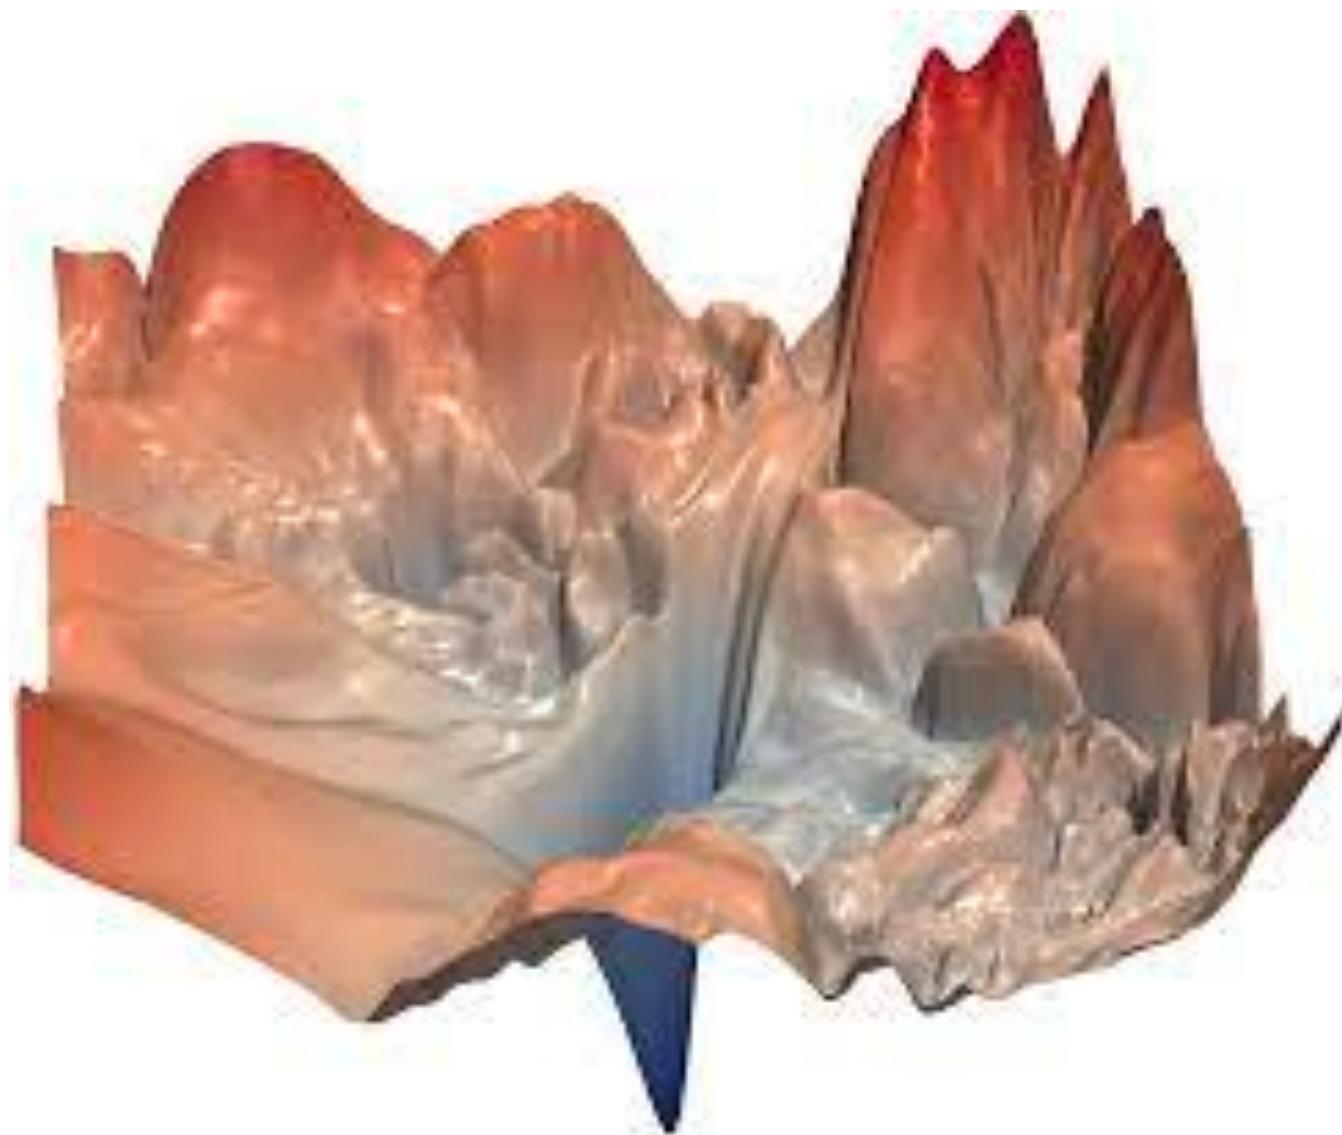
The difficulty to train neural networks

The difficulty to train neural networks



Example of loss landscape of NN

The difficulty to train neural networks

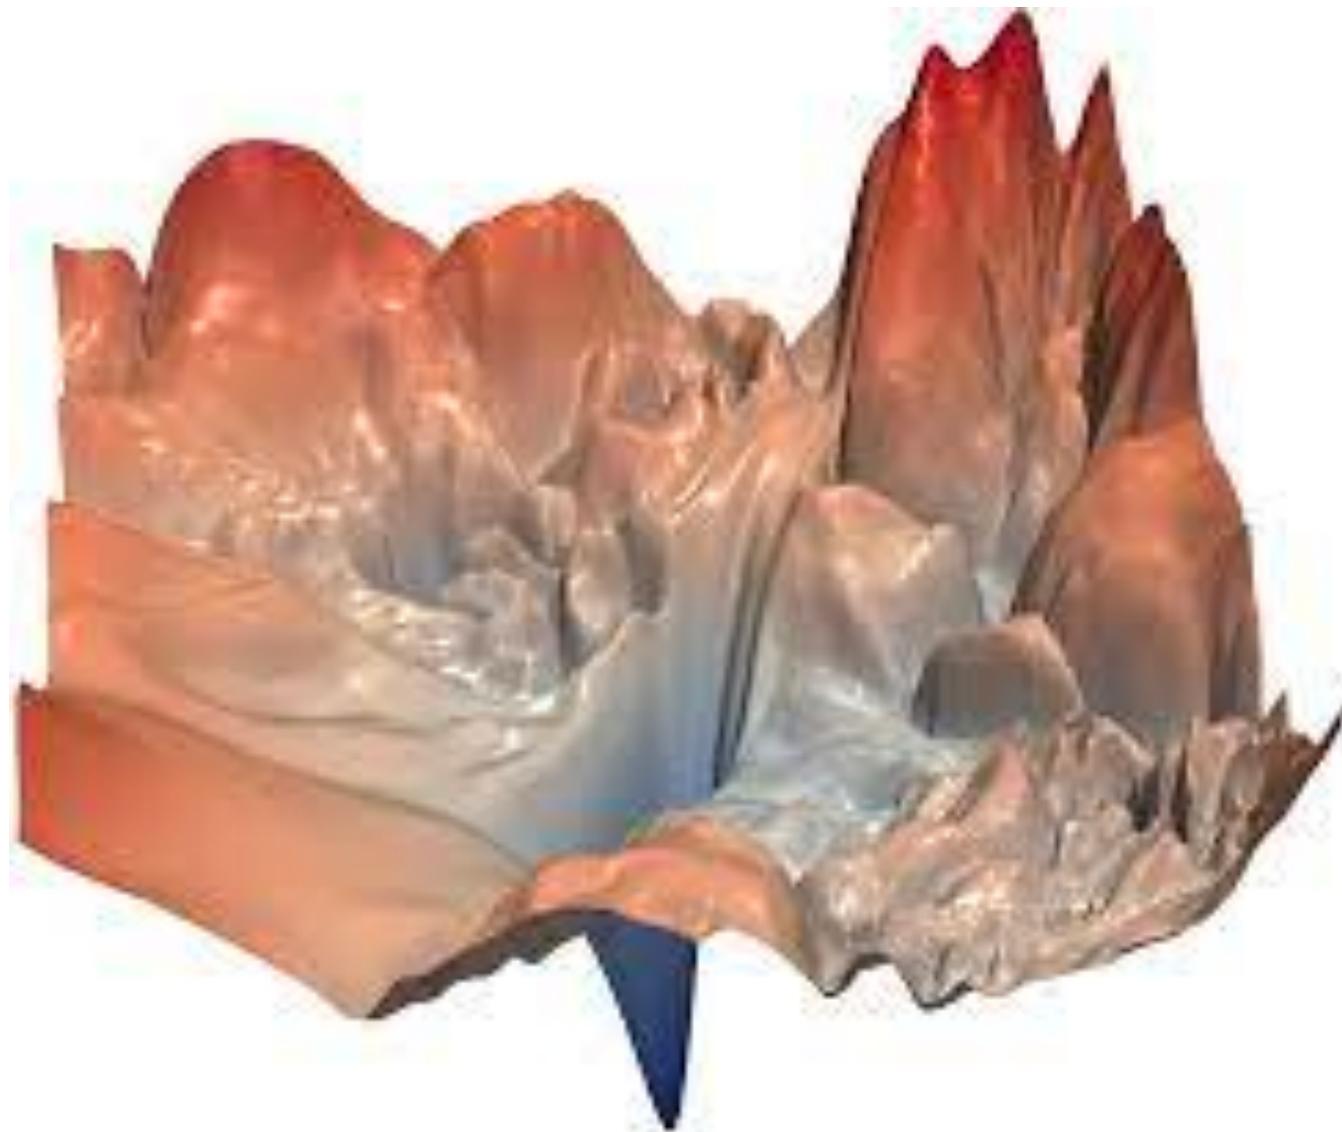


What do you notice?



Example of loss landscape of NN

The difficulty to train neural networks



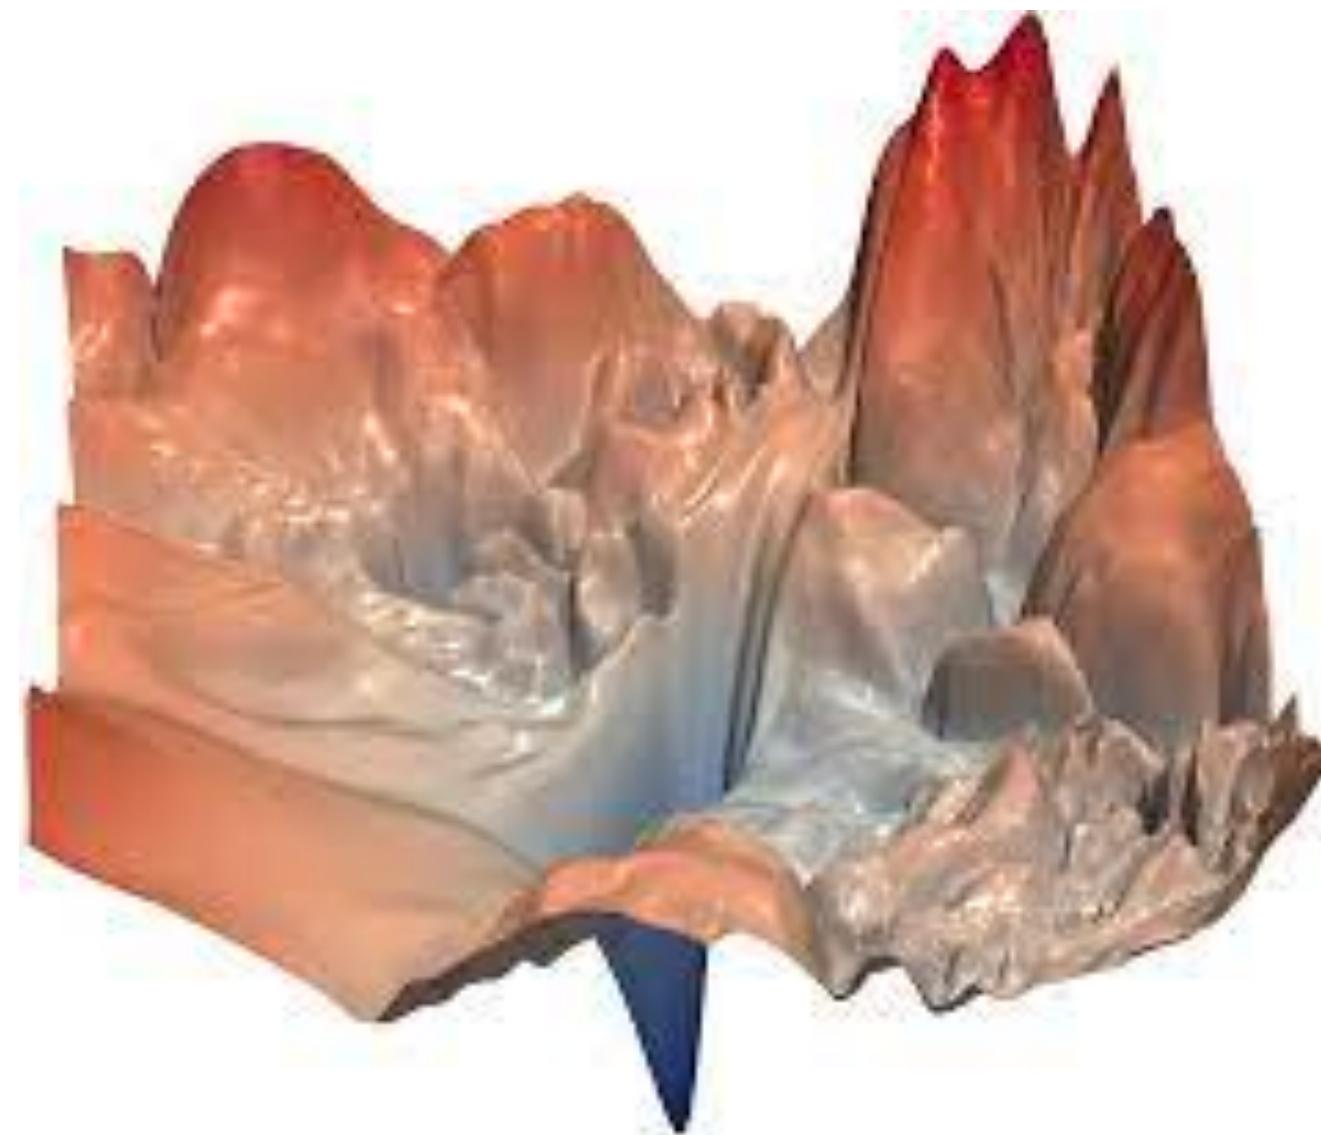
What do you notice?



This is not convex at all !

Example of loss landscape of NN

The difficulty to train neural networks



Example of loss landscape of NN

What do you notice?



This is not convex at all !

Gradient descent can be stuck in local optima

Initialization will matter a lot!

Setting the learning rate will also matter!

How to deal with it? Some recipe

How to deal with it? Some recipe

Learning rates a now adaptive

« adapt them to the loss landscape »

Can be made larger or small depending on the problem

How to deal with it? Some recipe

Learning rates are now adaptive

« adapt them to the loss landscape »

Can be made larger or small depending on the problem

Use Mini-batches for speed and stability

« Instead of computing the gradient on the whole dataset **use a subset** »

Smoother convergence + allows for larger learning rate

How to deal with it? Some recipe

Learning rates a now adaptive

« adapt them to the loss landscape »

Can be made larger or small depending on the problem

Use Mini-batches for speed and stability

« Instead of computing the gradient on the whole dataset **use a subset** »

Smoother convergence + allows for larger learning rate

Change the neural network architecture

Add more data

How to deal with it? Some recipe

Learning rates are now **adaptive**

« adapt them to the loss landscape »

Can be made larger or small **depending on the problem**

Use **Mini-batches** for speed and stability

« Instead of computing the gradient on the whole dataset **use a subset** »

Smoother convergence + allows for larger learning rate

Change the neural network architecture

Add more data

Those will always work!

Let's sum up so far

Let's sum up so far

1. The basic bloc to build a Neural Network

We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

2. The role of activation functions

They complexity the neural network and allow to fit more complex data

They are added at the end of each layer and there is a large variety of them

3. The cost function

It allows to measure how « far » is the neural network from the reality

Let's sum up so far

1. The basic bloc to build a Neural Network

We have seen the forward propagation, i.e. from an input how to use the Neural Network

This operation is call forward propagation, this operation is used at inference time

2. The role of activation functions

They complexity the neural network and allow to fit more complex data

They are added at the end of each layer and there is a large variety of them

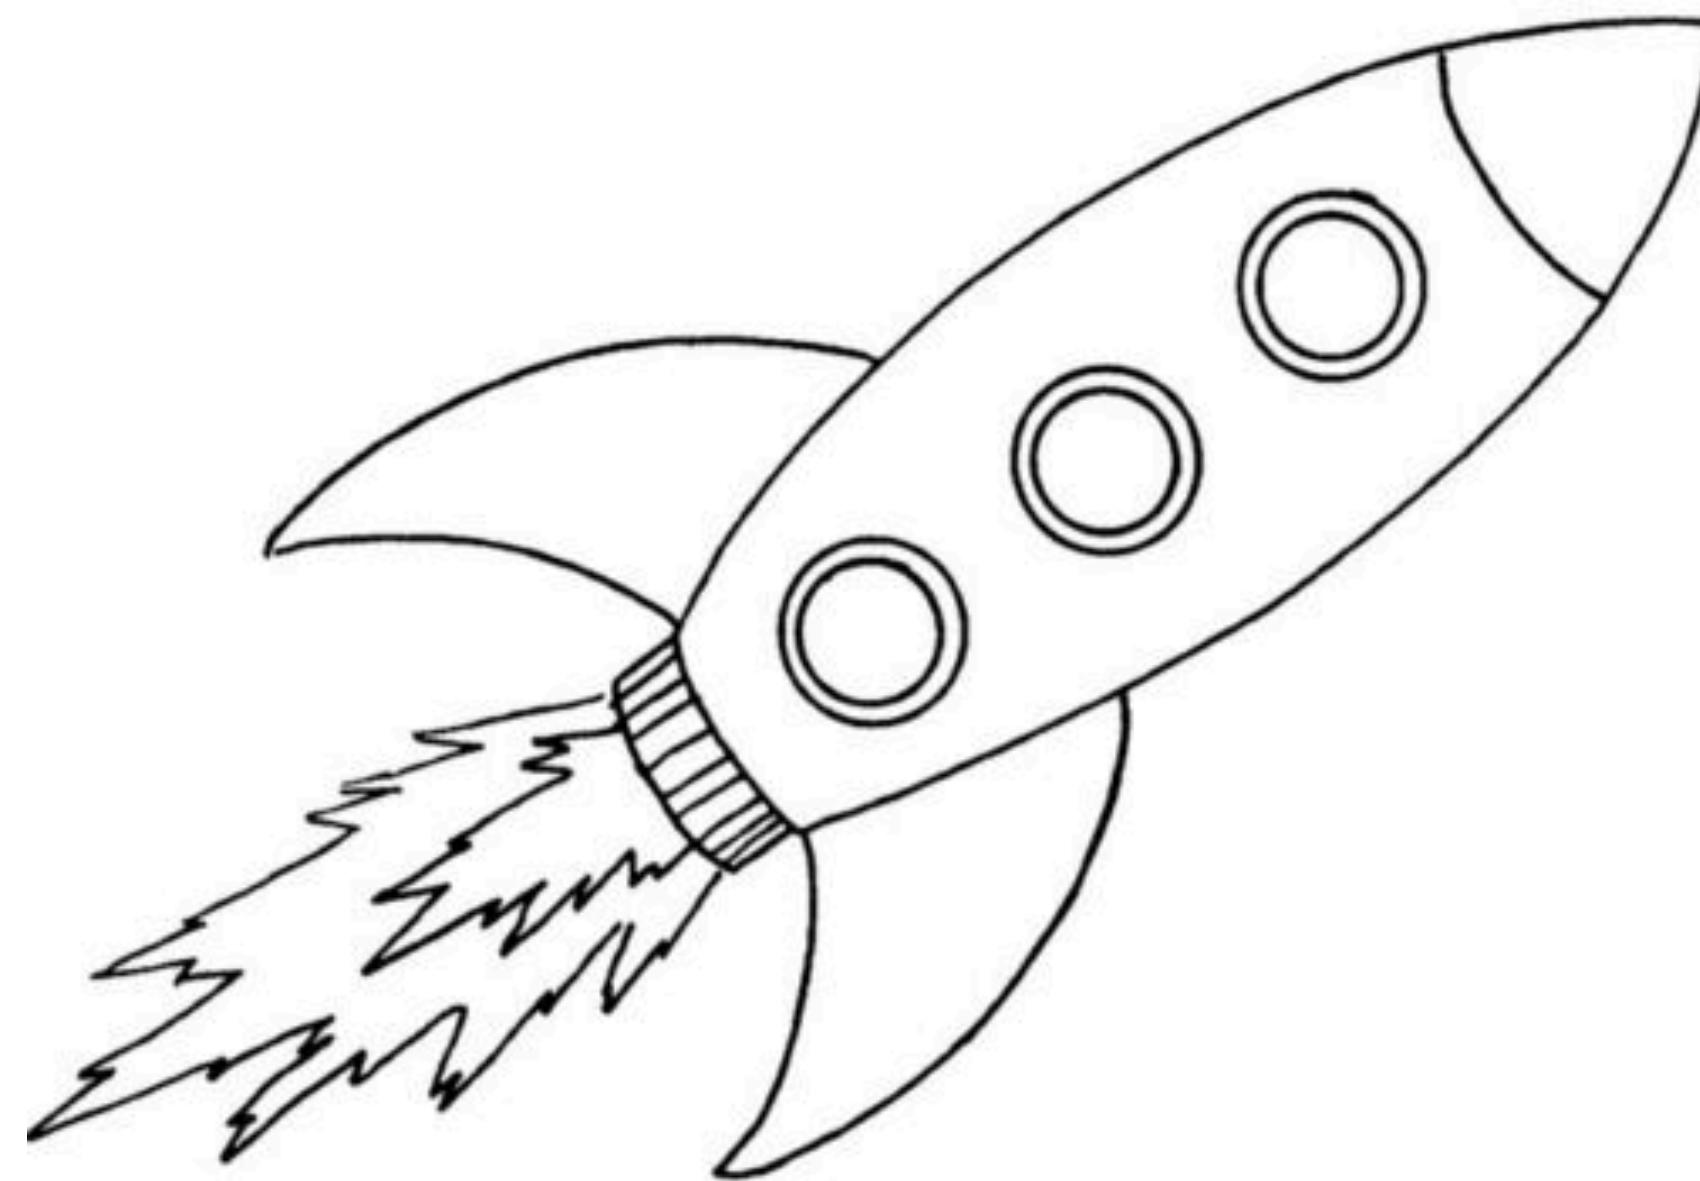
3. The cost function

It allows to measure how « far » is the neural network from the reality

4. The learning phase

The Gradient descent algorithm that allows to update the weights of the networks

From theory to practice



Splitting Data Into Training/Validation/Test Set

Splitting Data Into Training/Validation/Test Set

We have access to a dataset

Splitting Data Into Training/Validation/Test Set

We have access to a dataset

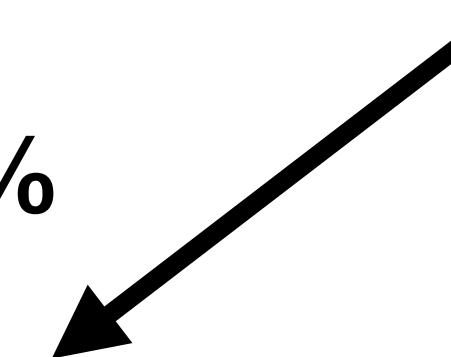
$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

Splitting Data Into Training/Validation/Test Set

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$

80%



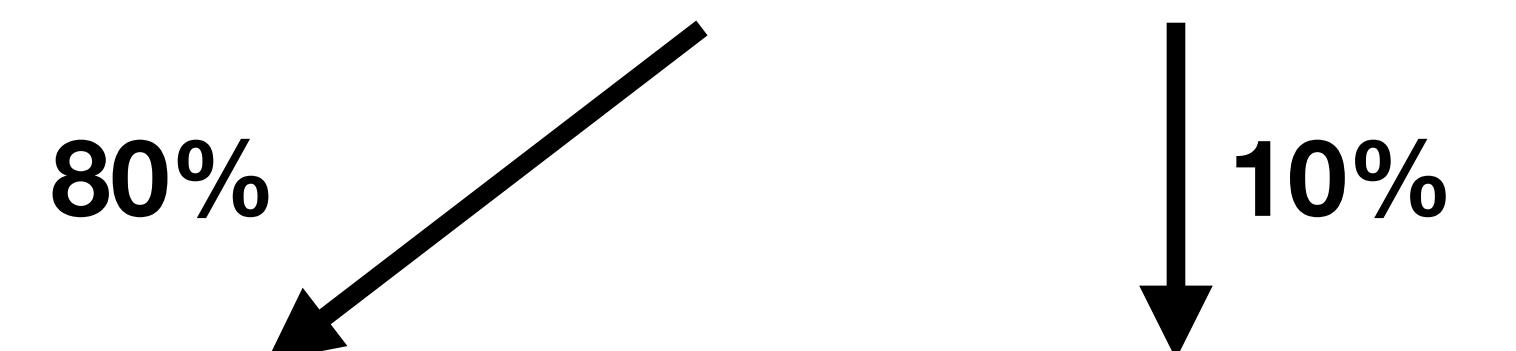
Training Set

Use for the learning
algorithm

Splitting Data Into Training/Validation/Test Set

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$



Training Set

Use for the learning
algorithm

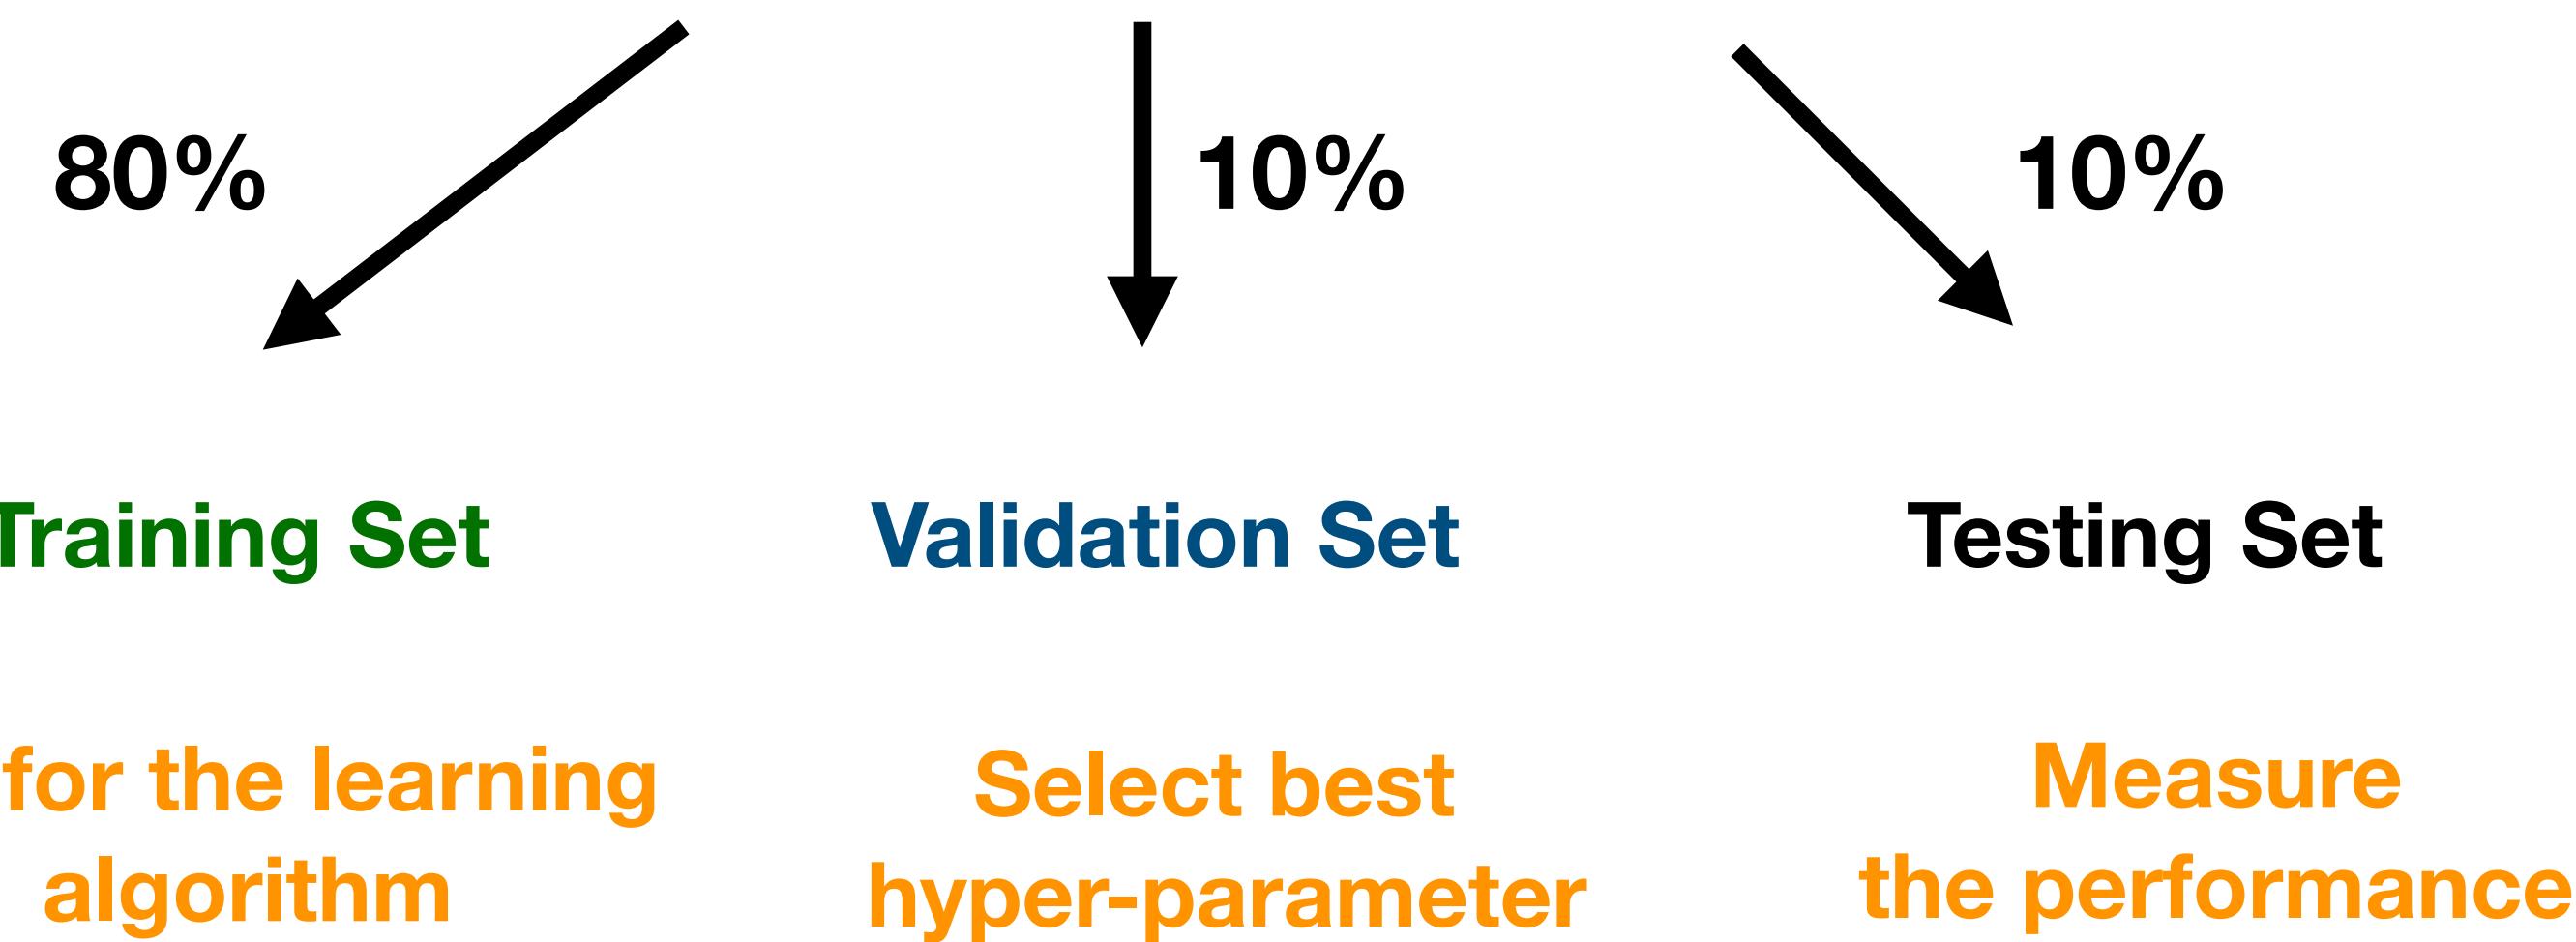
Validation Set

Select best
hyper-parameter

Splitting Data Into Training/Validation/Test Set

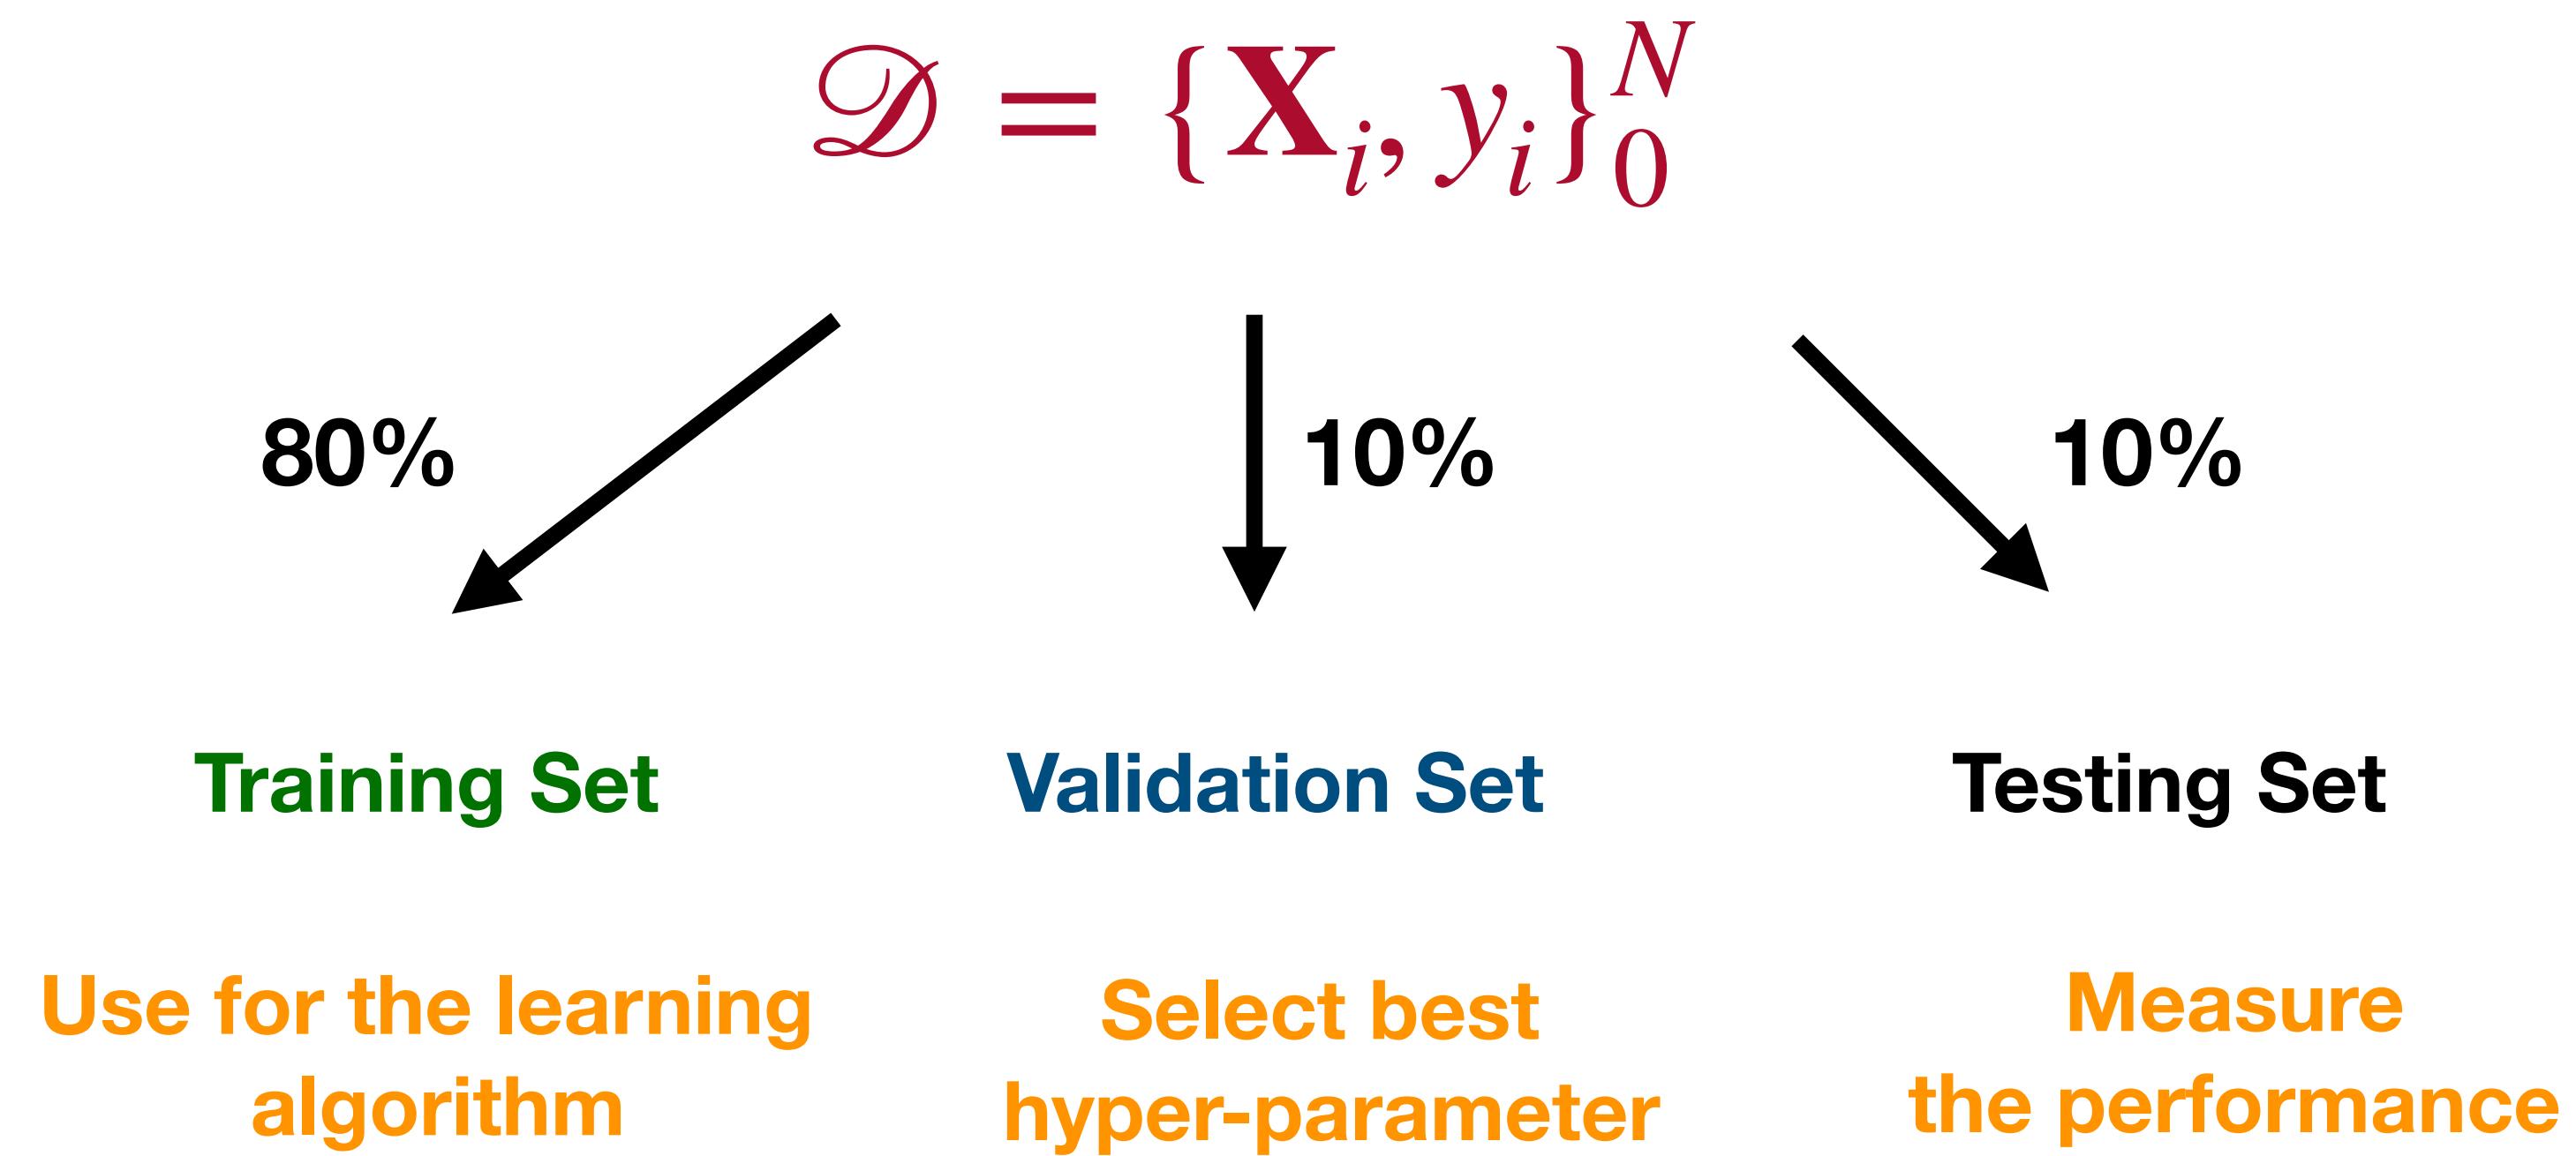
We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$



Splitting Data Into Training/Validation/Test Set

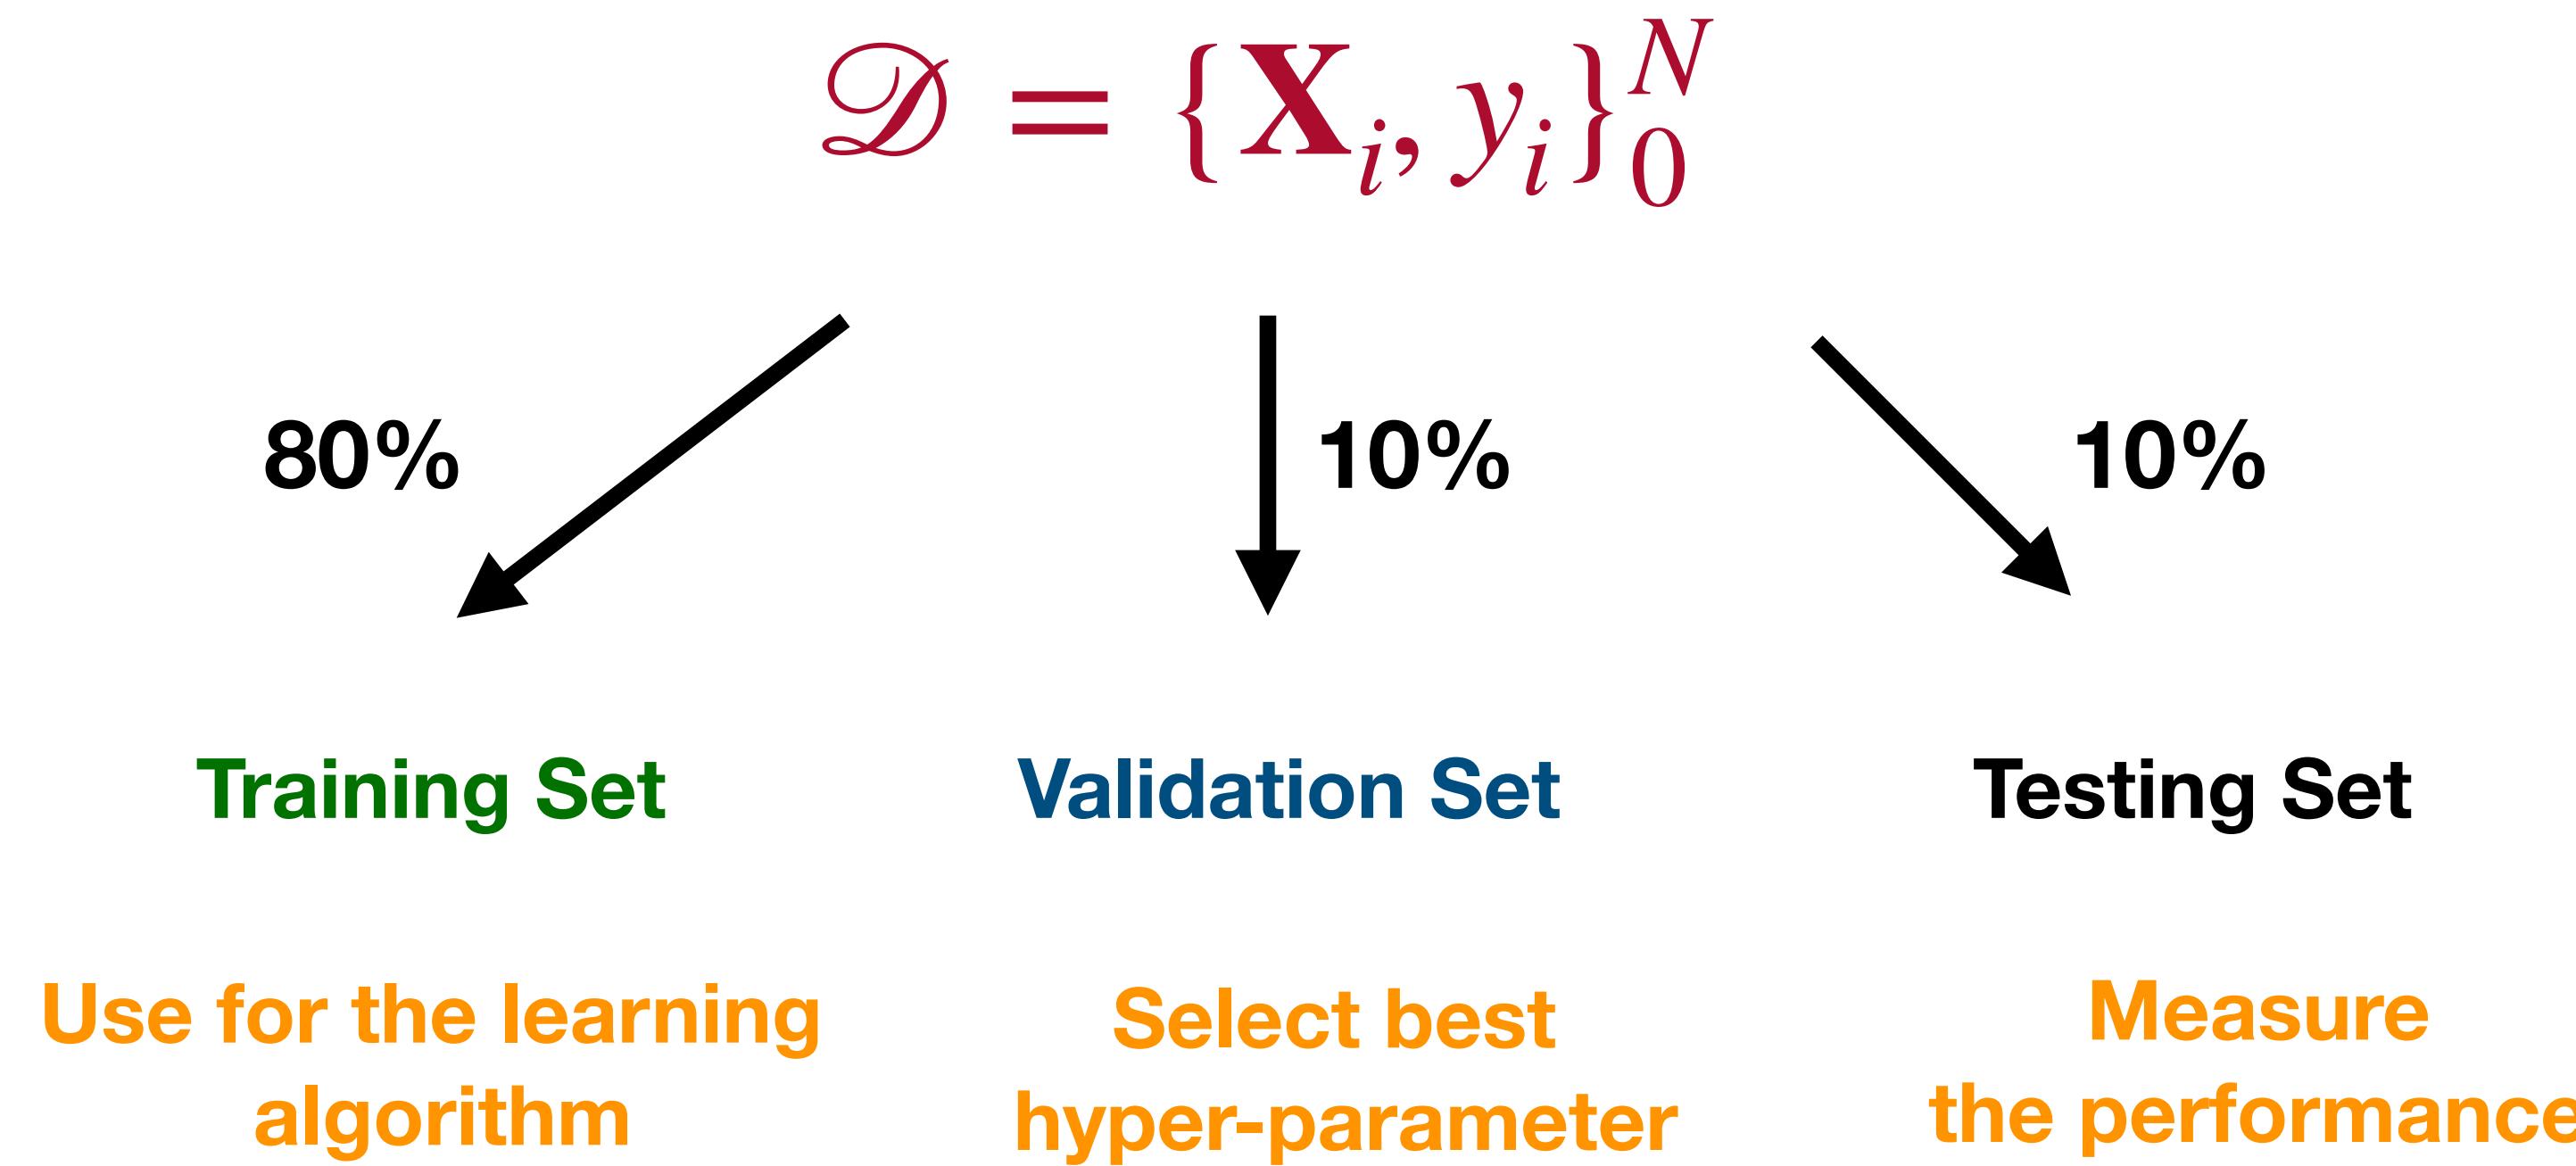
We have access to a dataset



Best practices:

Splitting Data Into Training/Validation/Test Set

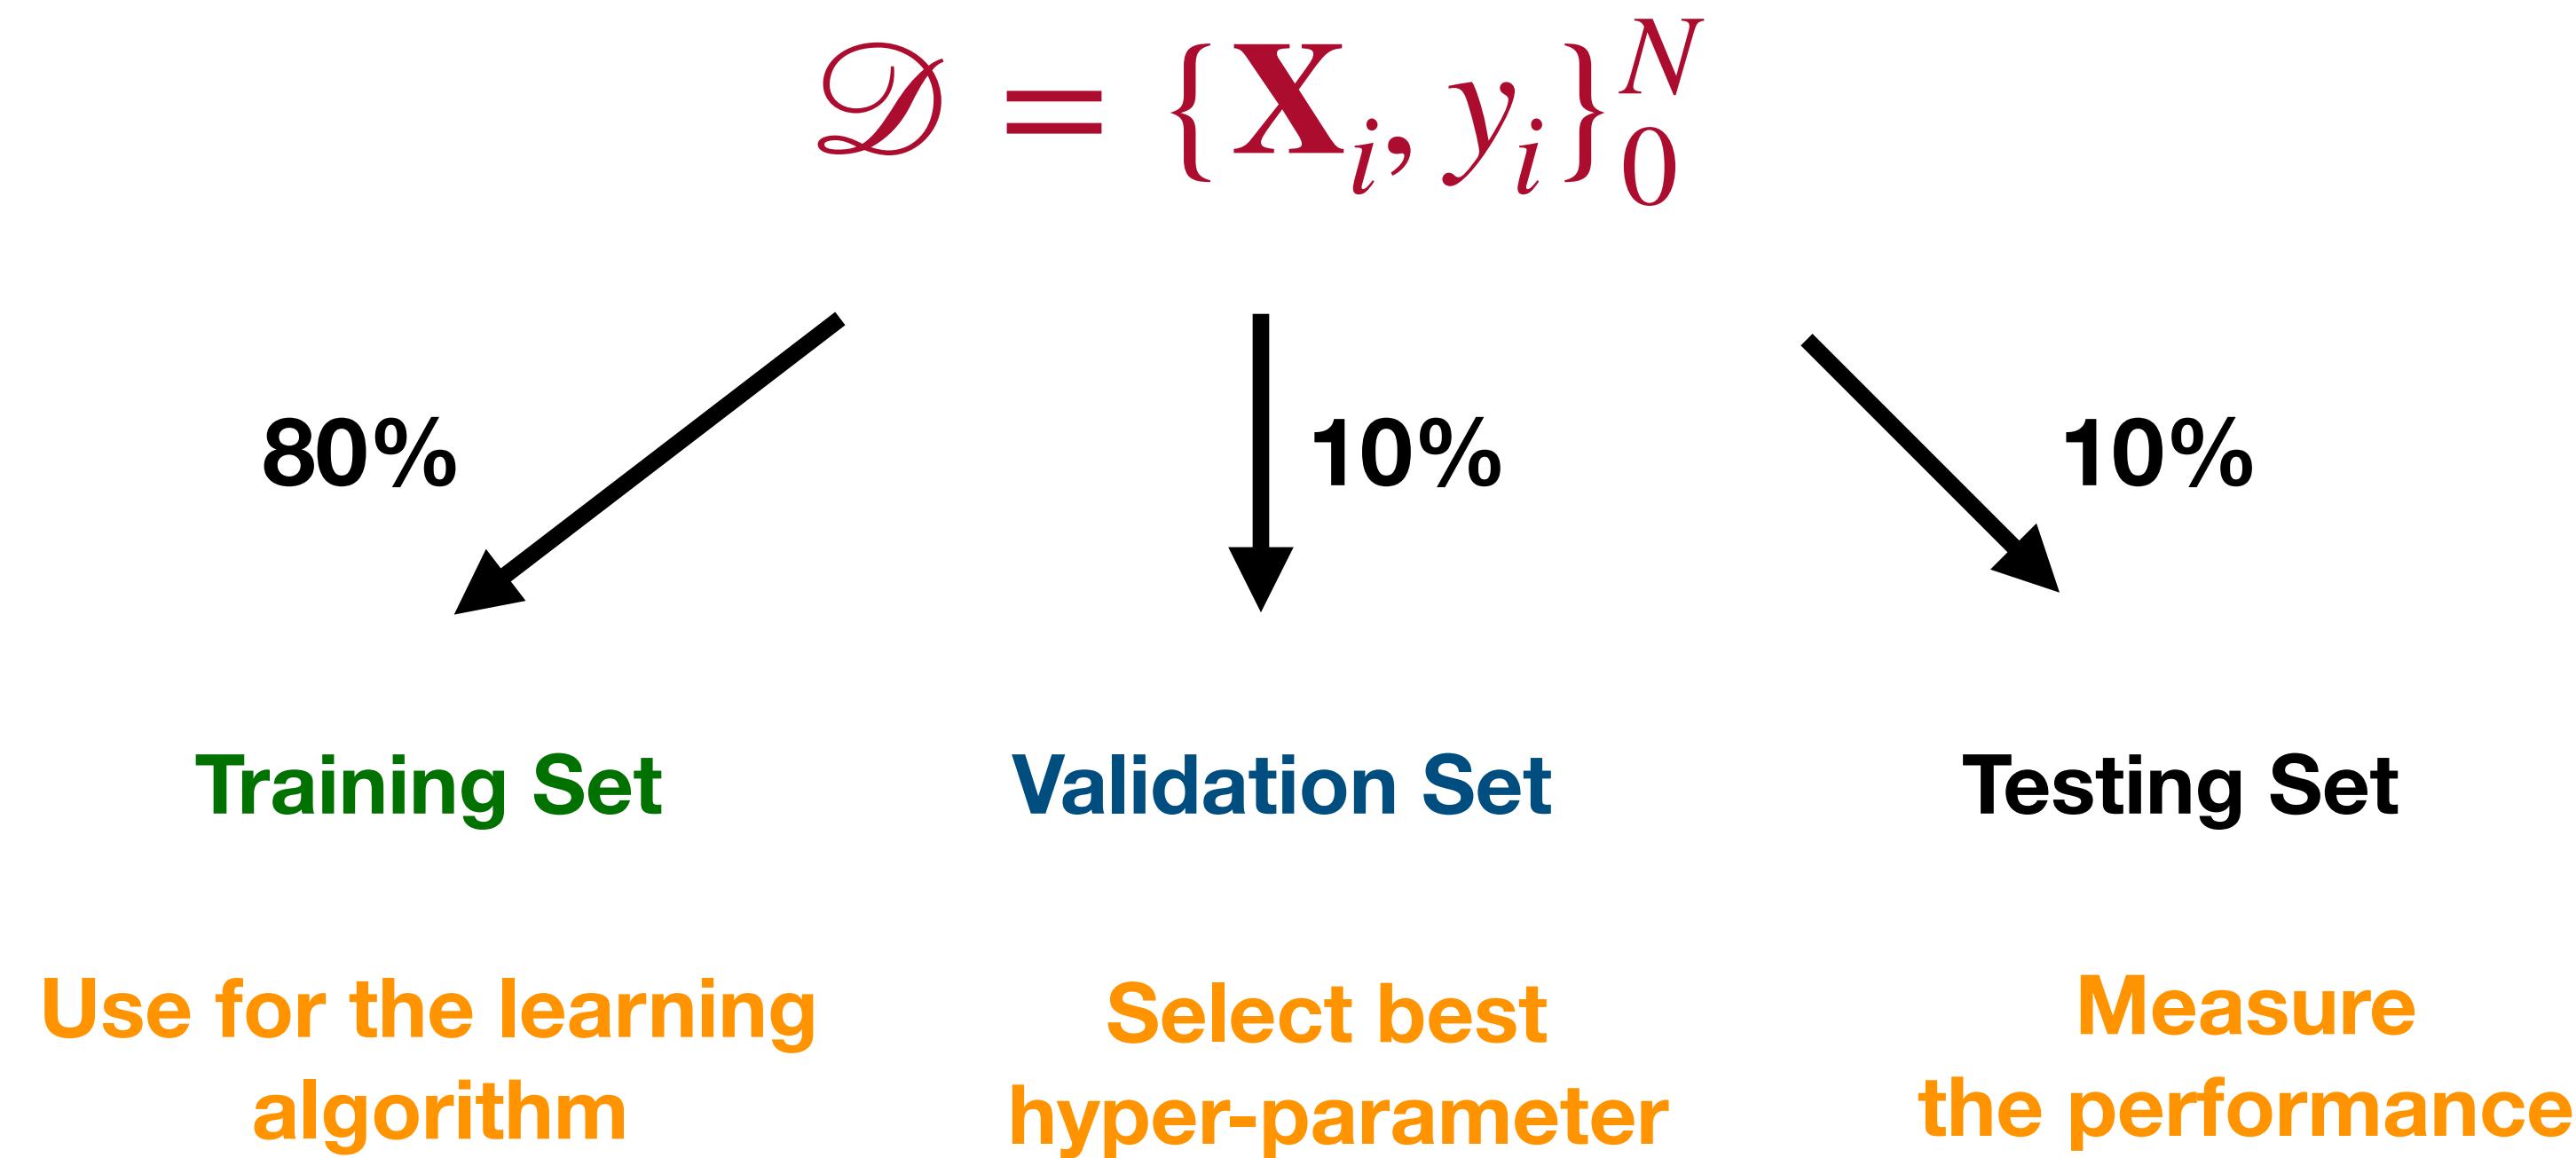
We have access to a dataset



Best practices: 1. ensure that your results are independent of the split

Splitting Data Into Training/Validation/Test Set

We have access to a dataset



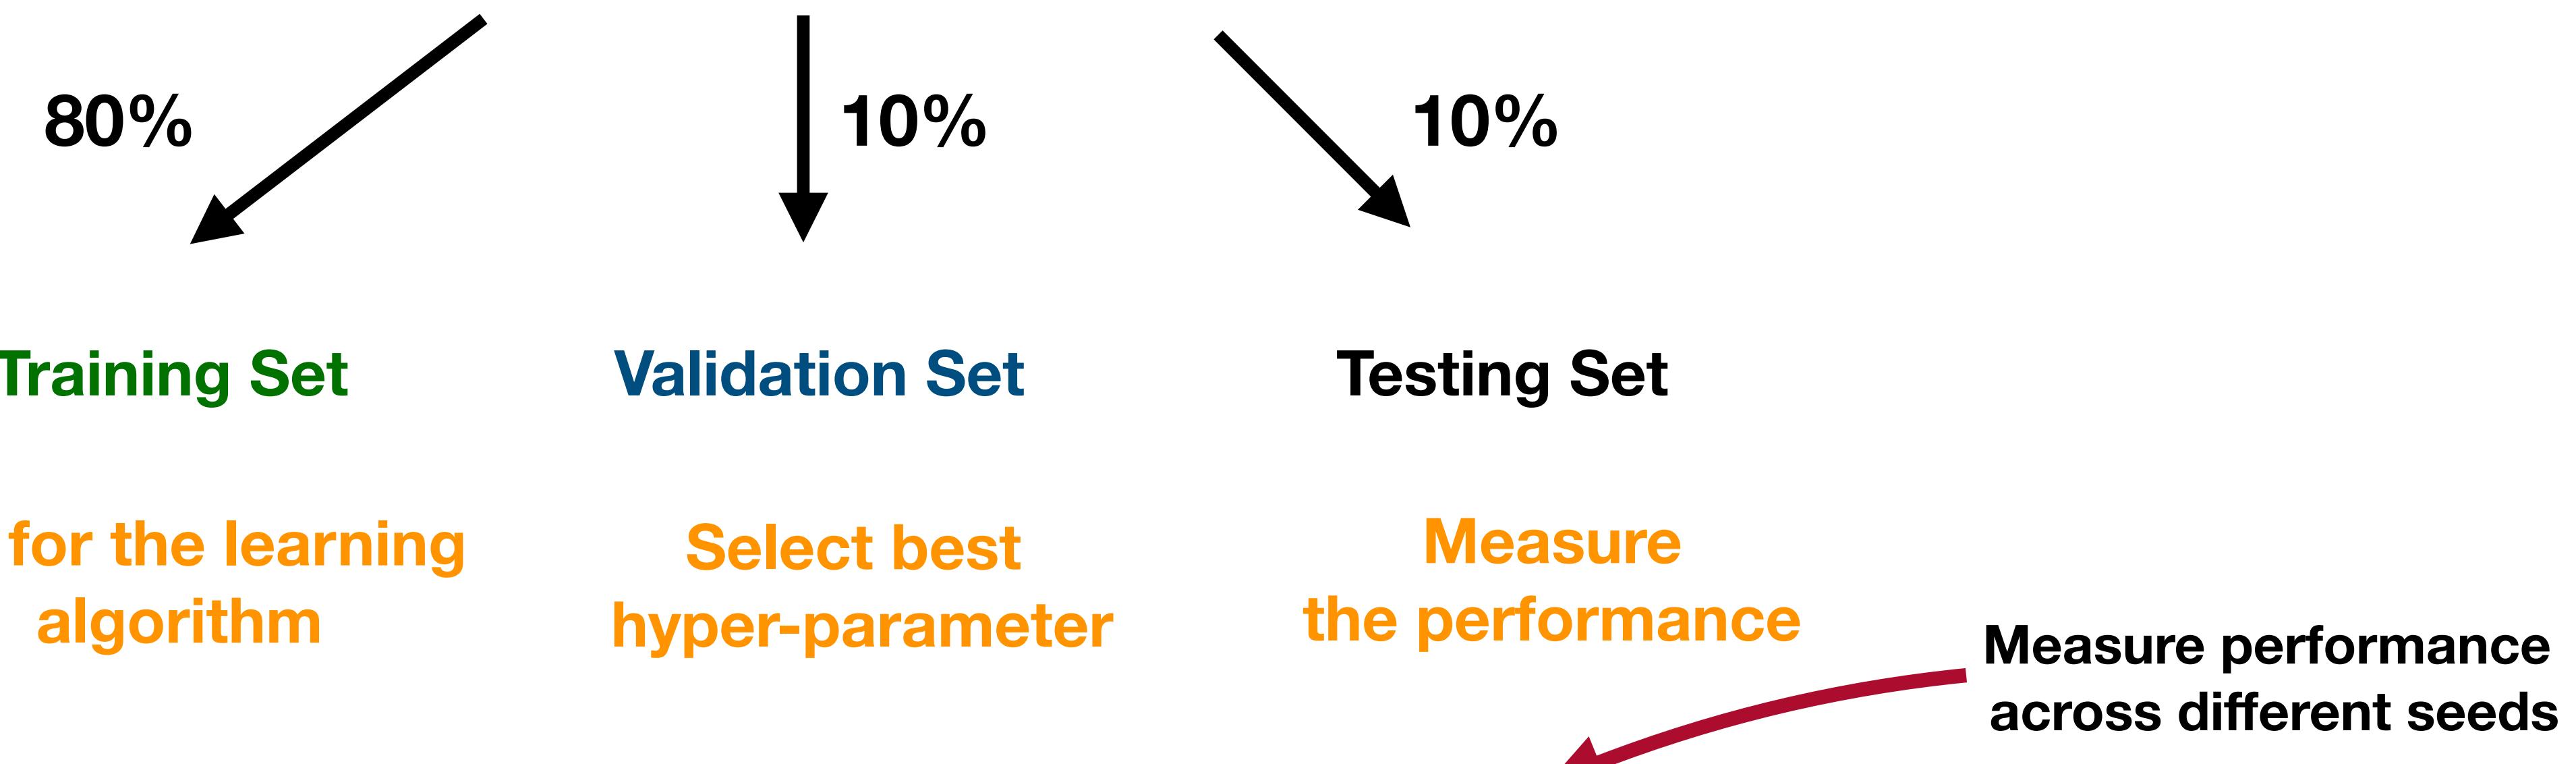
Best practices: 1. ensure that your results are independent of the split

2. make sure that the proportion of the labels in the split are uniform

Splitting Data Into Training/Validation/Test Set

We have access to a dataset

$$\mathcal{D} = \{\mathbf{X}_i, y_i\}_0^N$$



Best practices:

1. ensure that your results are independent of the split
2. make sure that the proportion of the labels in the split are uniform

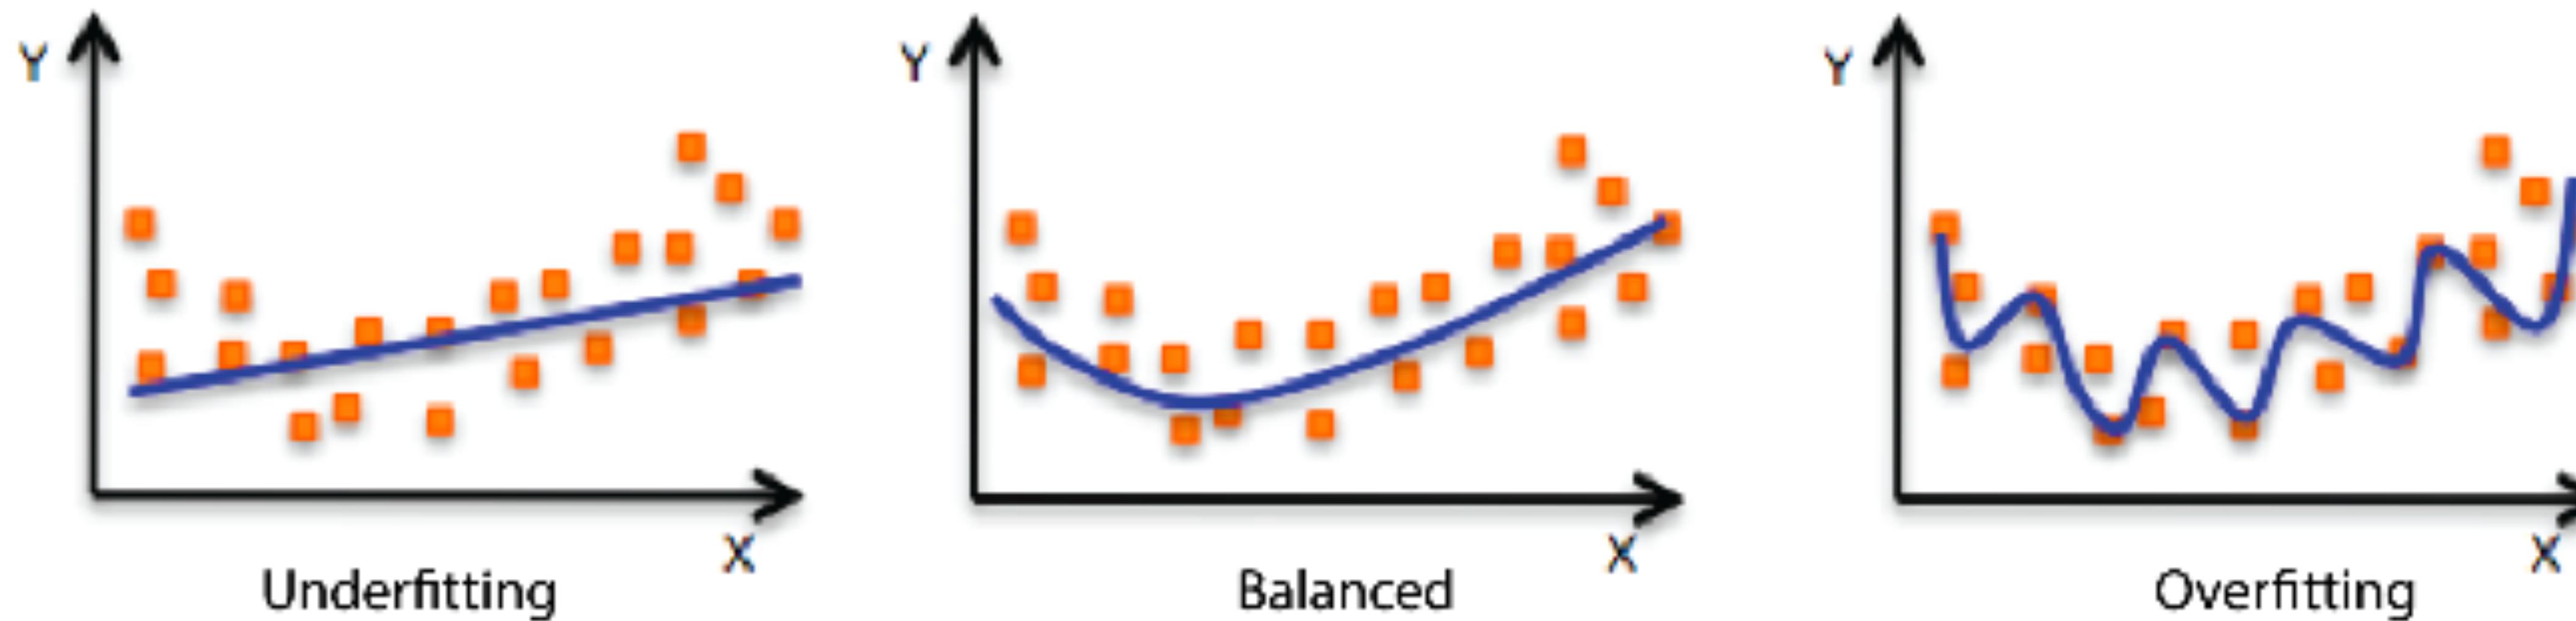
Overfitting vs Underfitting

Overfitting vs Underfitting

A gentle **reminder** from basic Machine Learning Classes

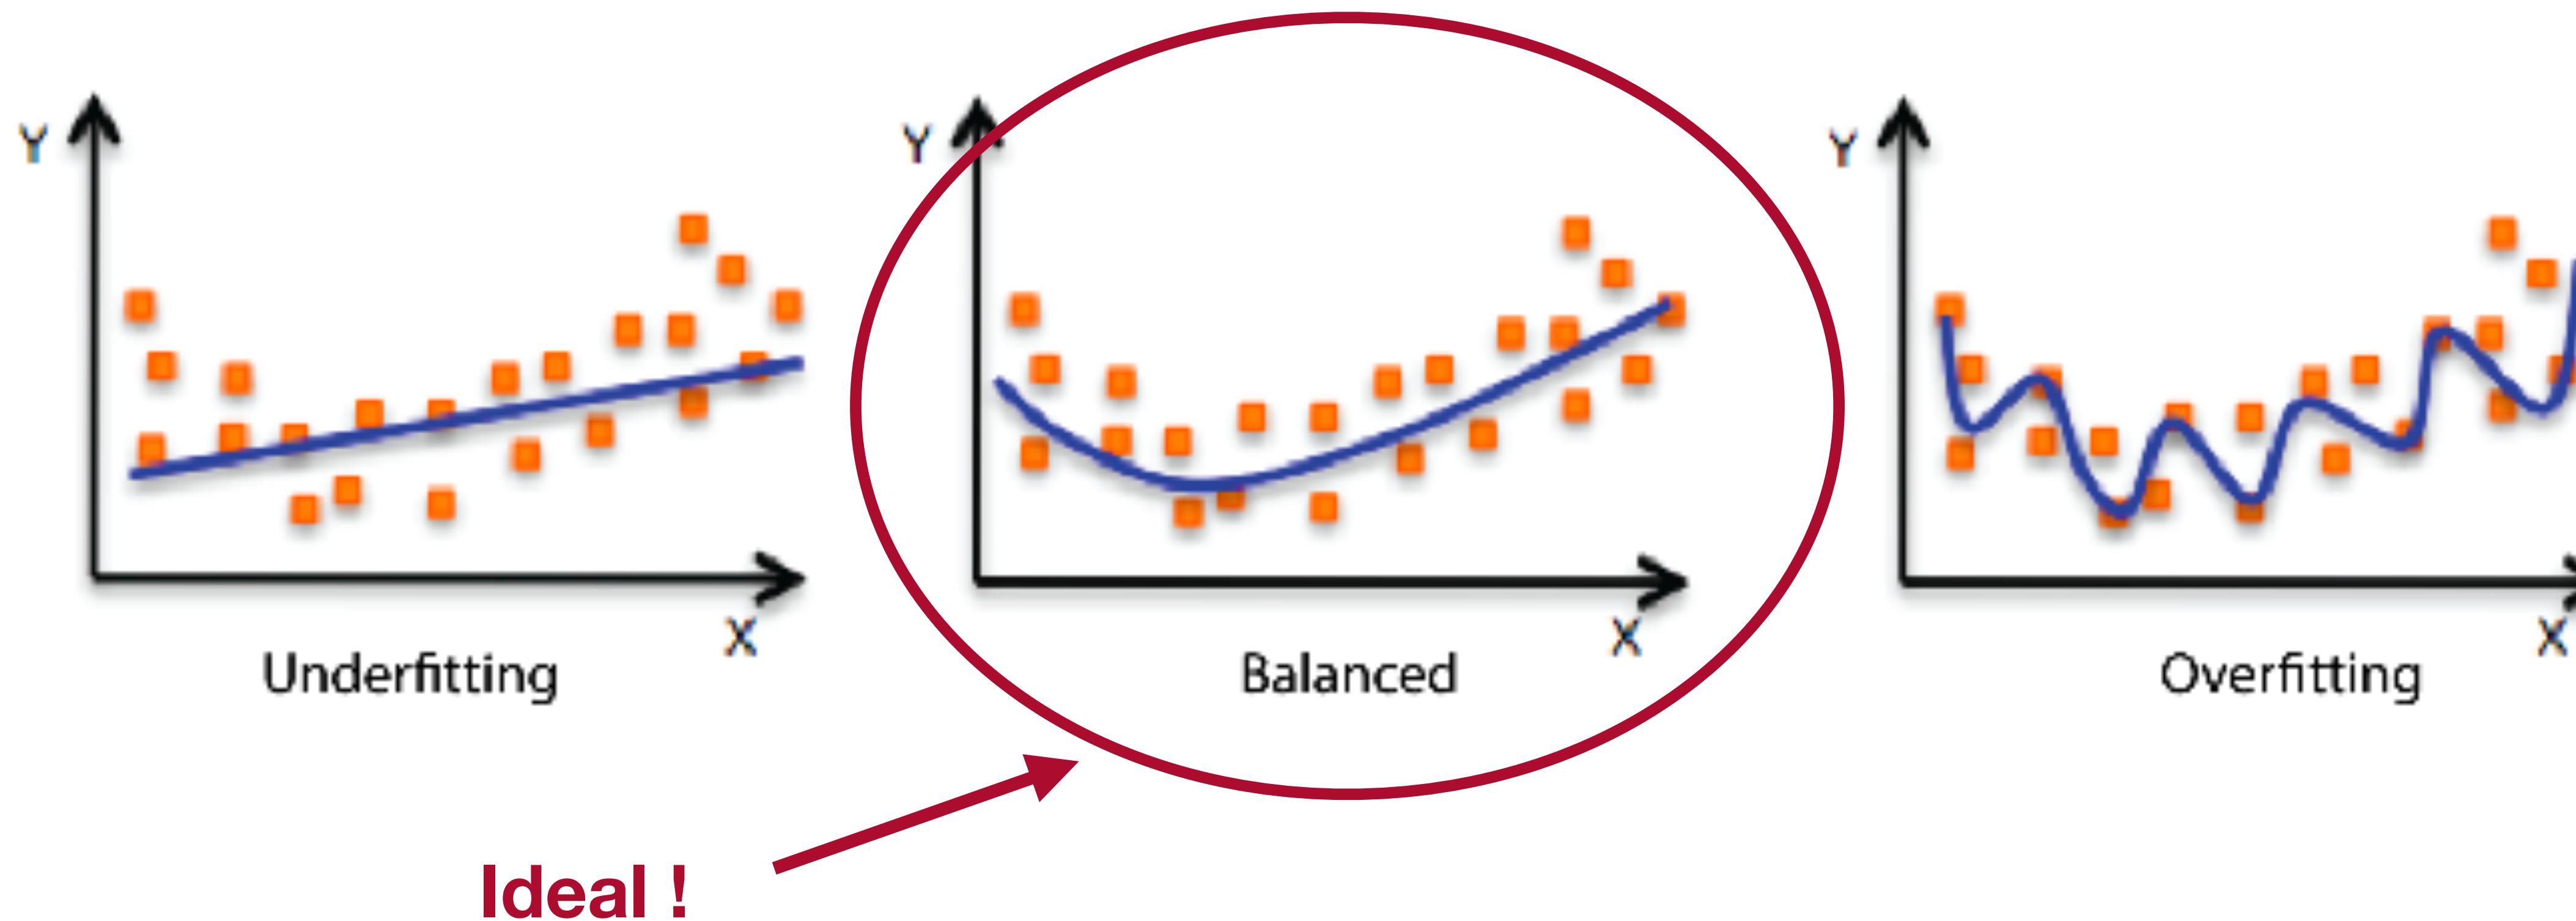
Overfitting vs Underfitting

A gentle **reminder** from basic Machine Learning Classes



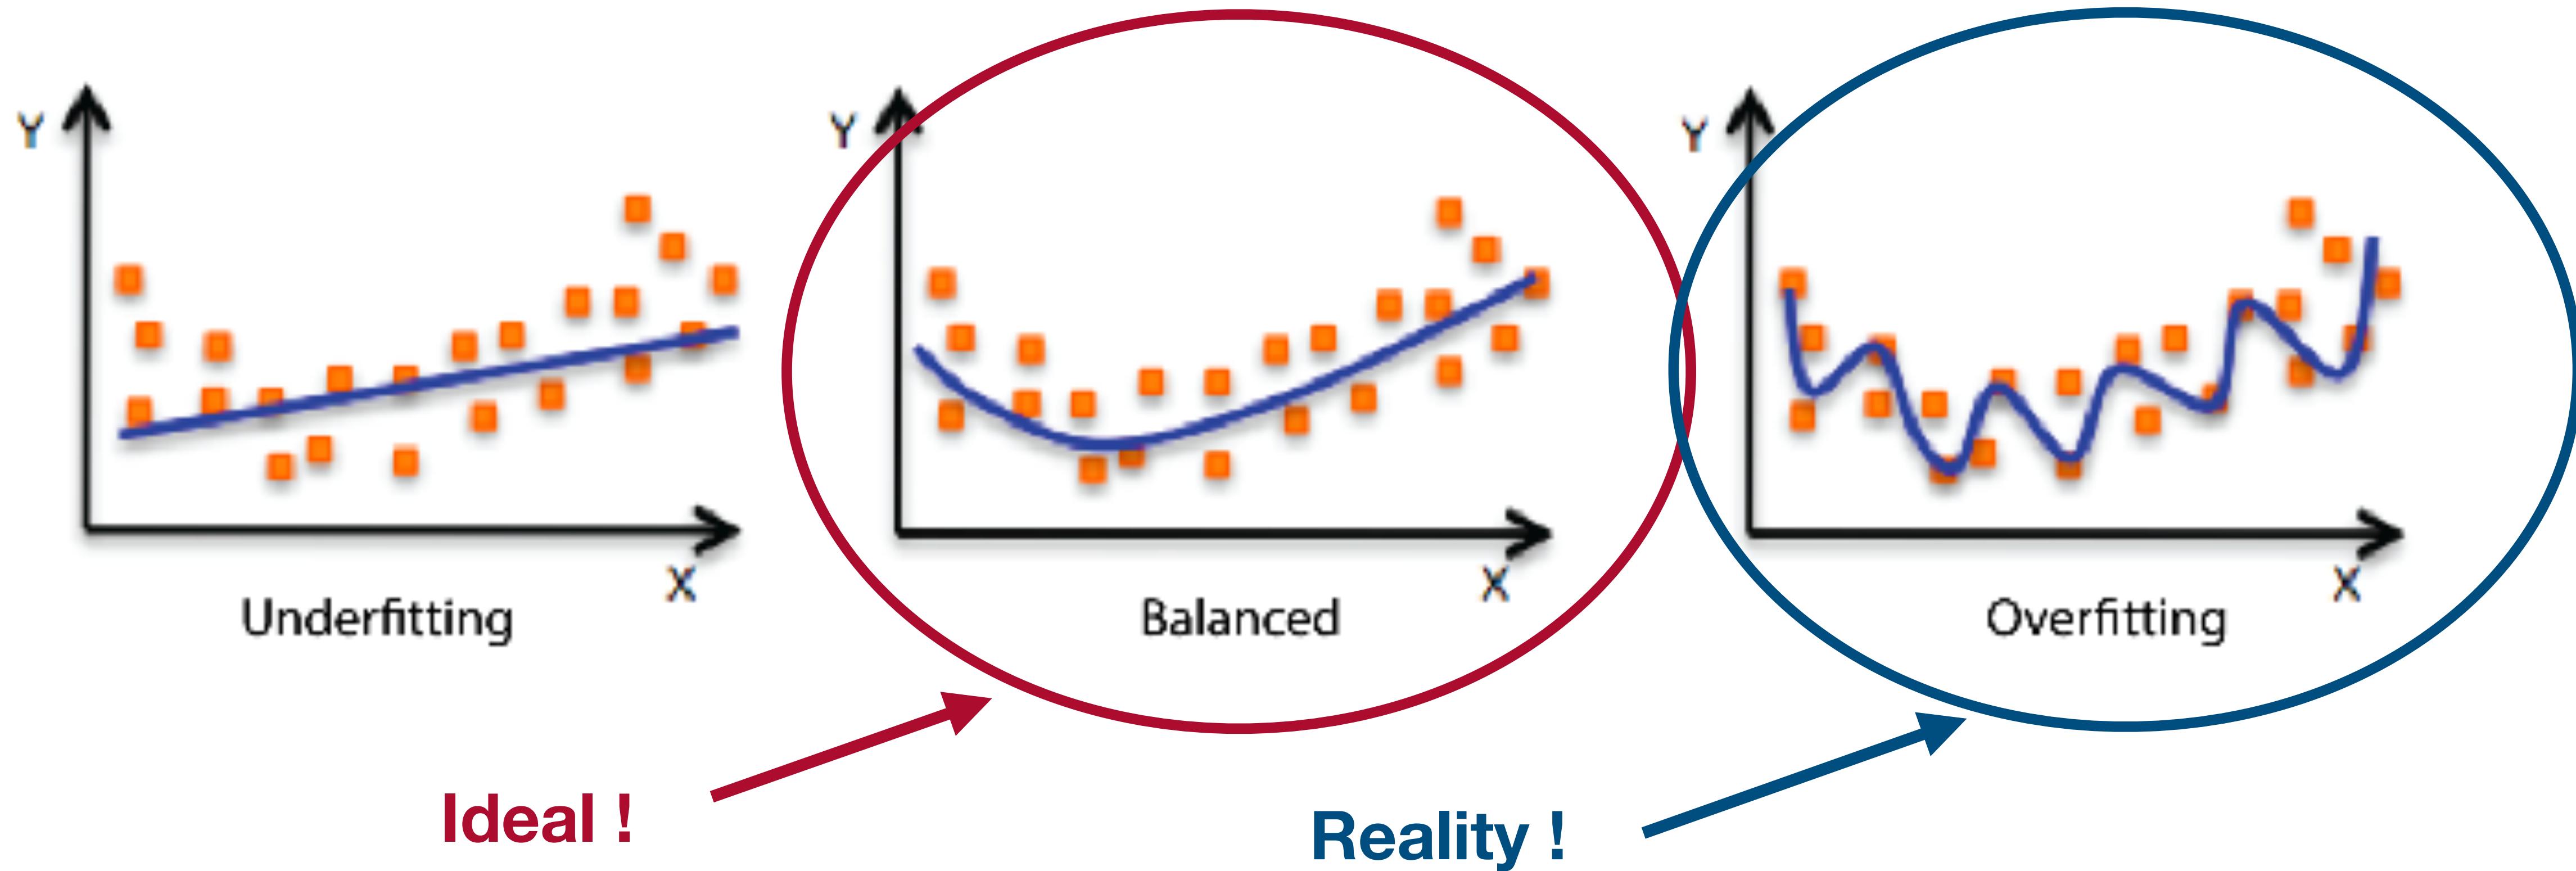
Overfitting vs Underfitting

A gentle **reminder** from basic Machine Learning Classes



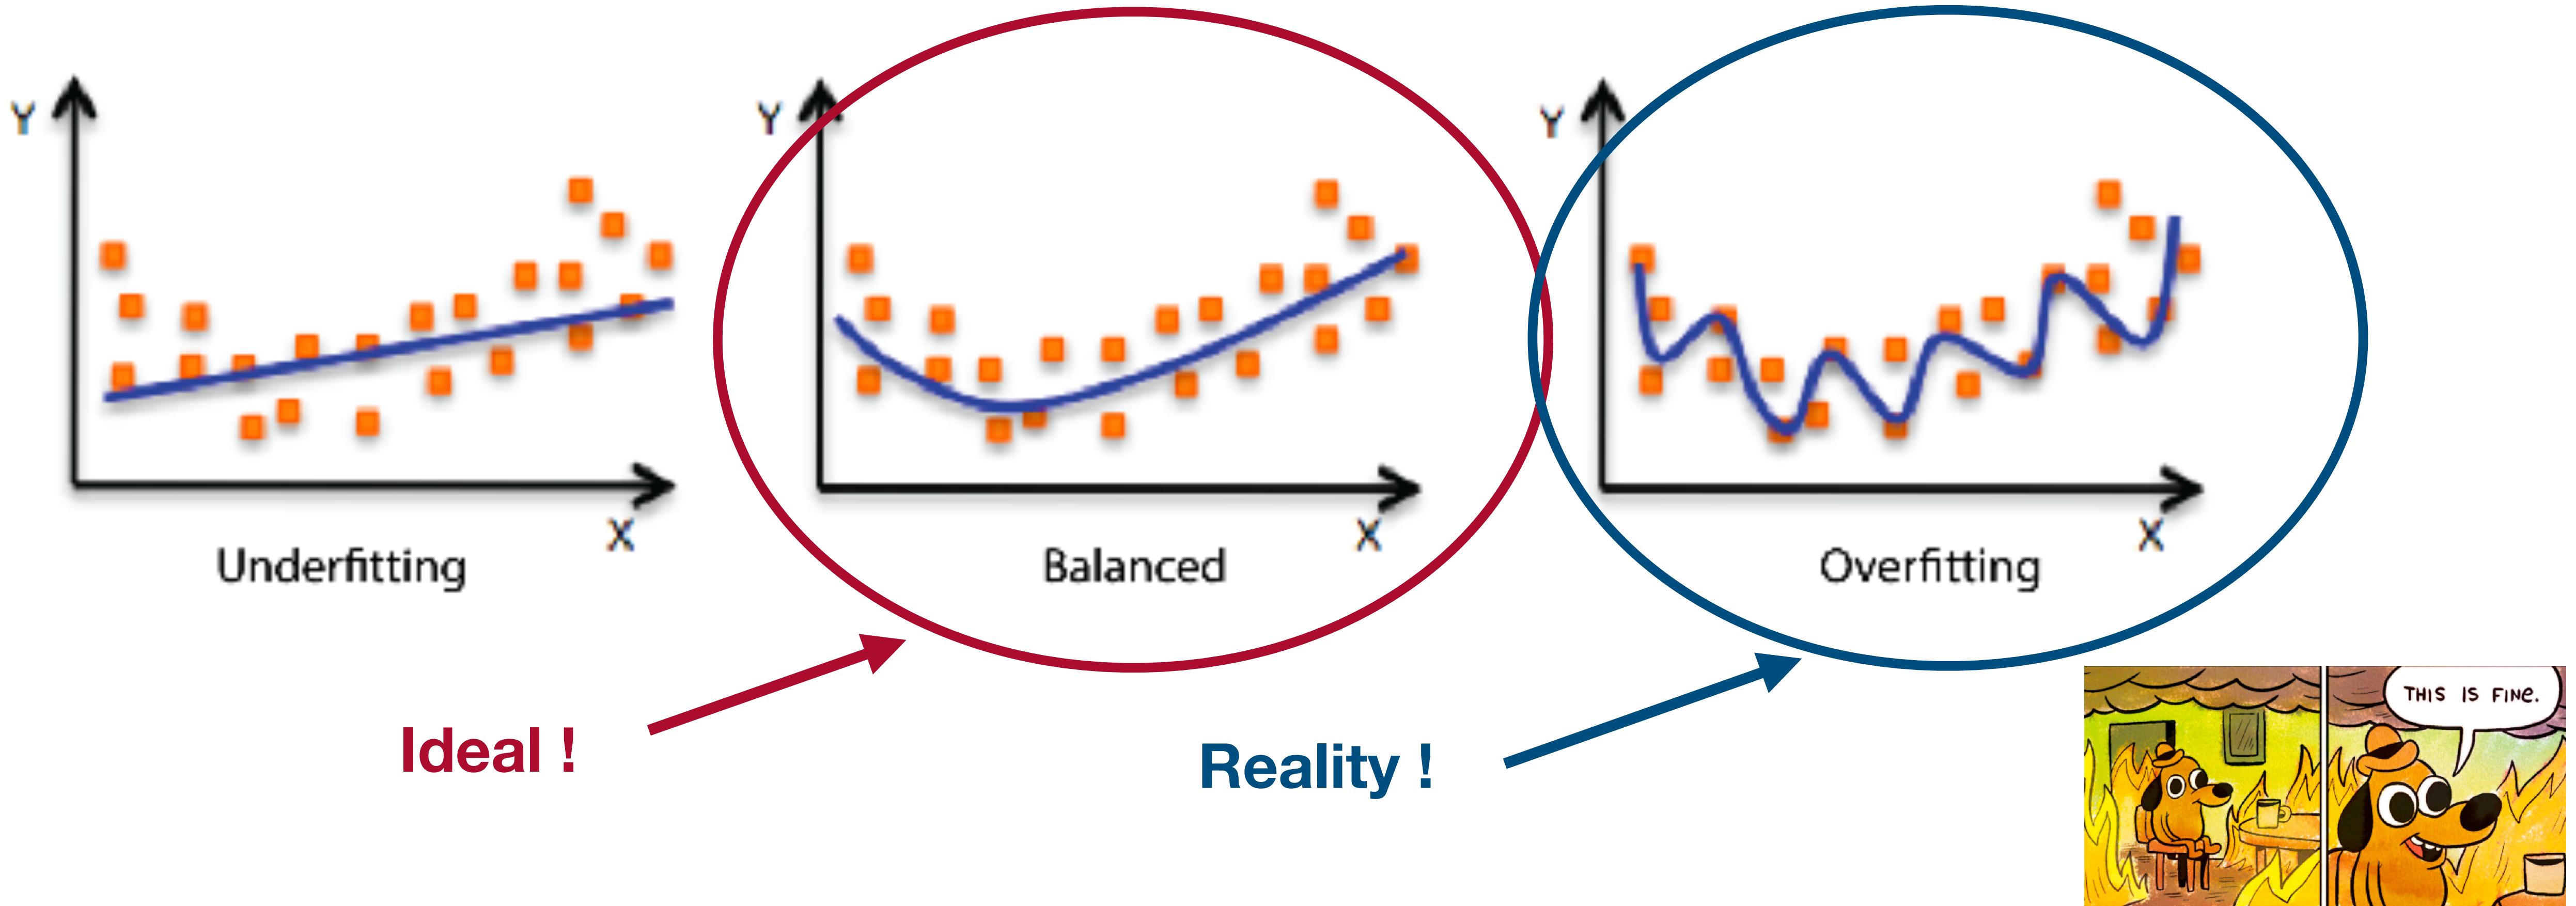
Overfitting vs Underfitting

A gentle **reminder** from basic Machine Learning Classes



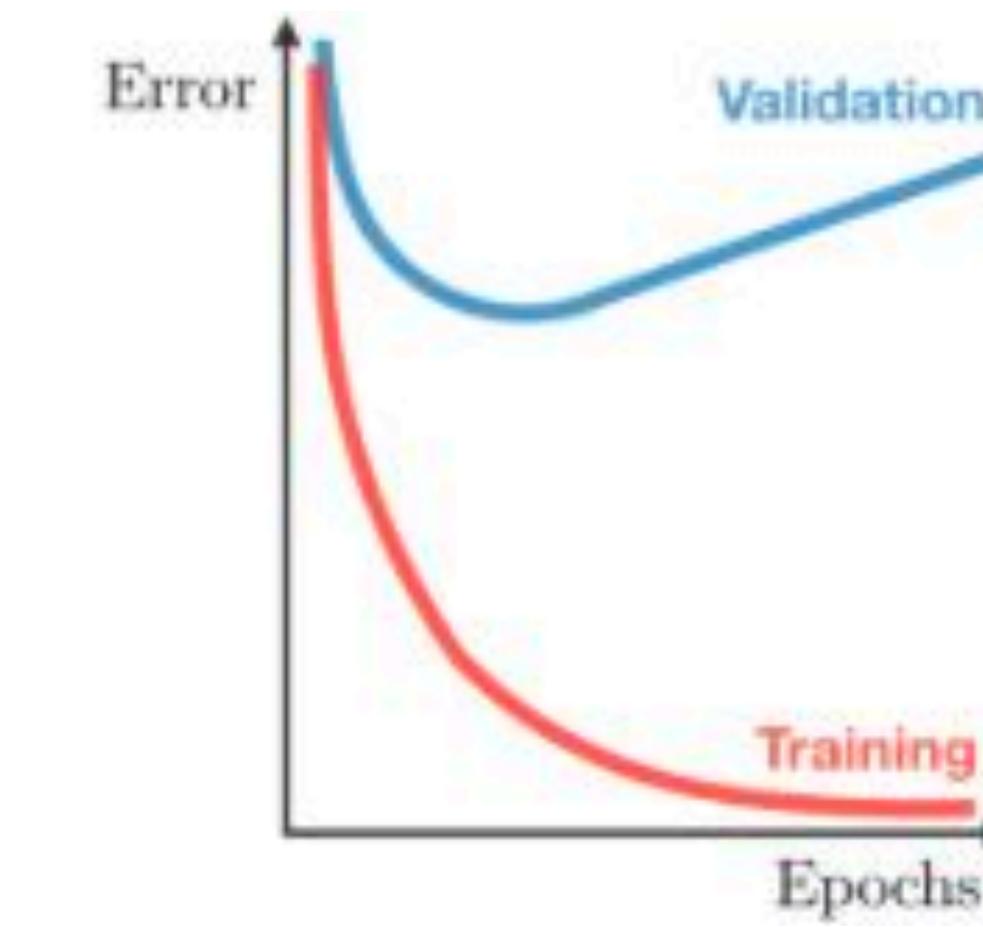
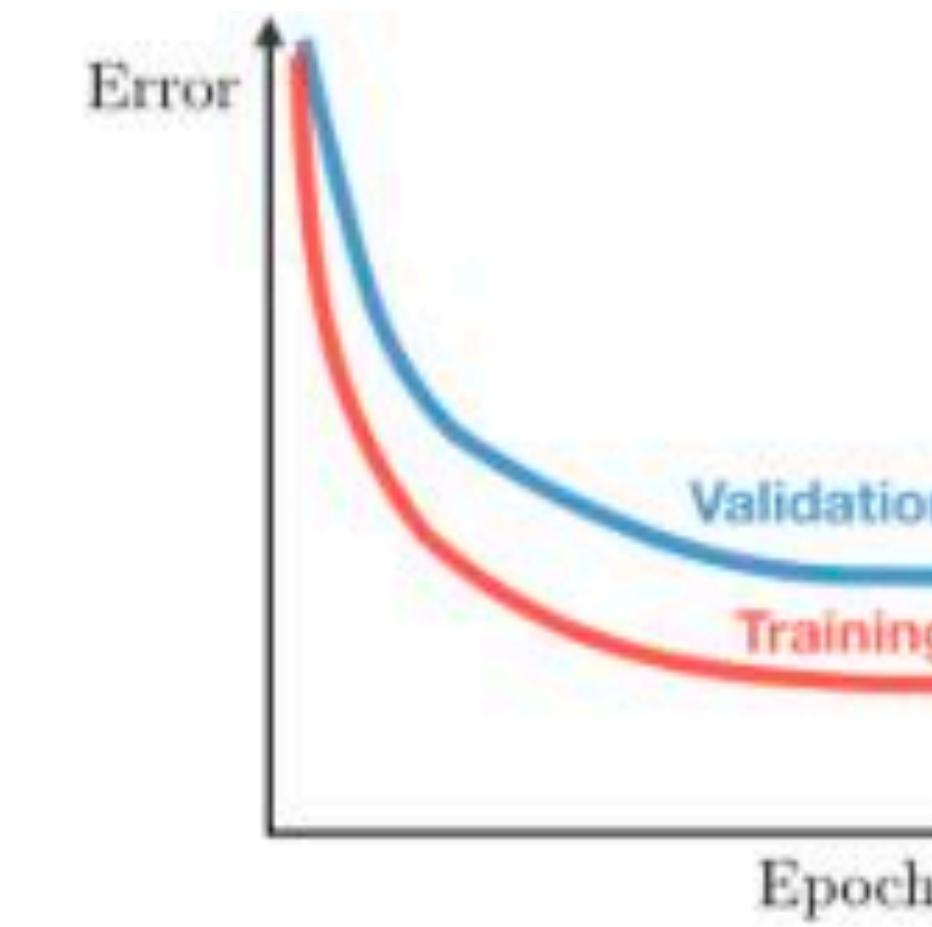
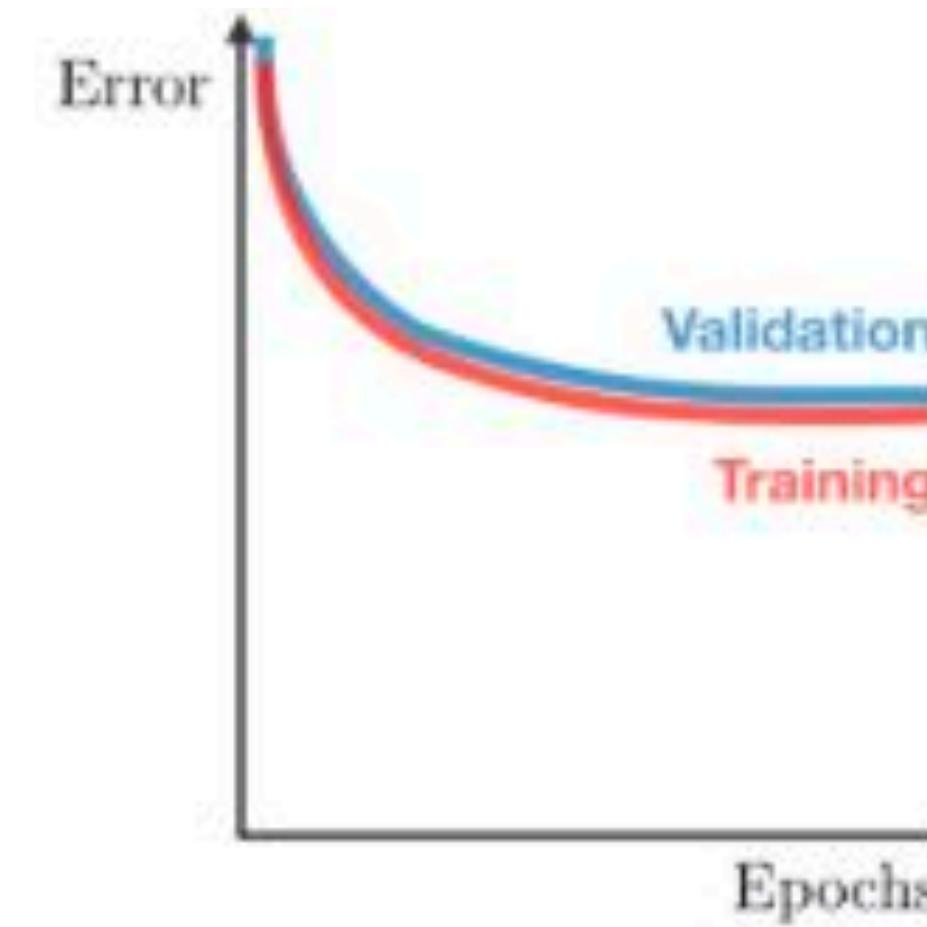
Overfitting vs Underfitting

A gentle **reminder** from basic Machine Learning Classes



Quiz: who is who ?

Which situation consist of overfitting/underfitting/ just right ?



In practice, when the validation loss starts to increase you are overfitting

When to stop training?

when the validation loss has converged or starts to increase

Fighting Overfitting with regularization

Fighting Overfitting with regularization





Regularization: the set of technics to prevent or fight overfitting



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 1: Dropout



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 1: Dropout

During training randomly set some activation to 0

Fighting Overfitting with regularization



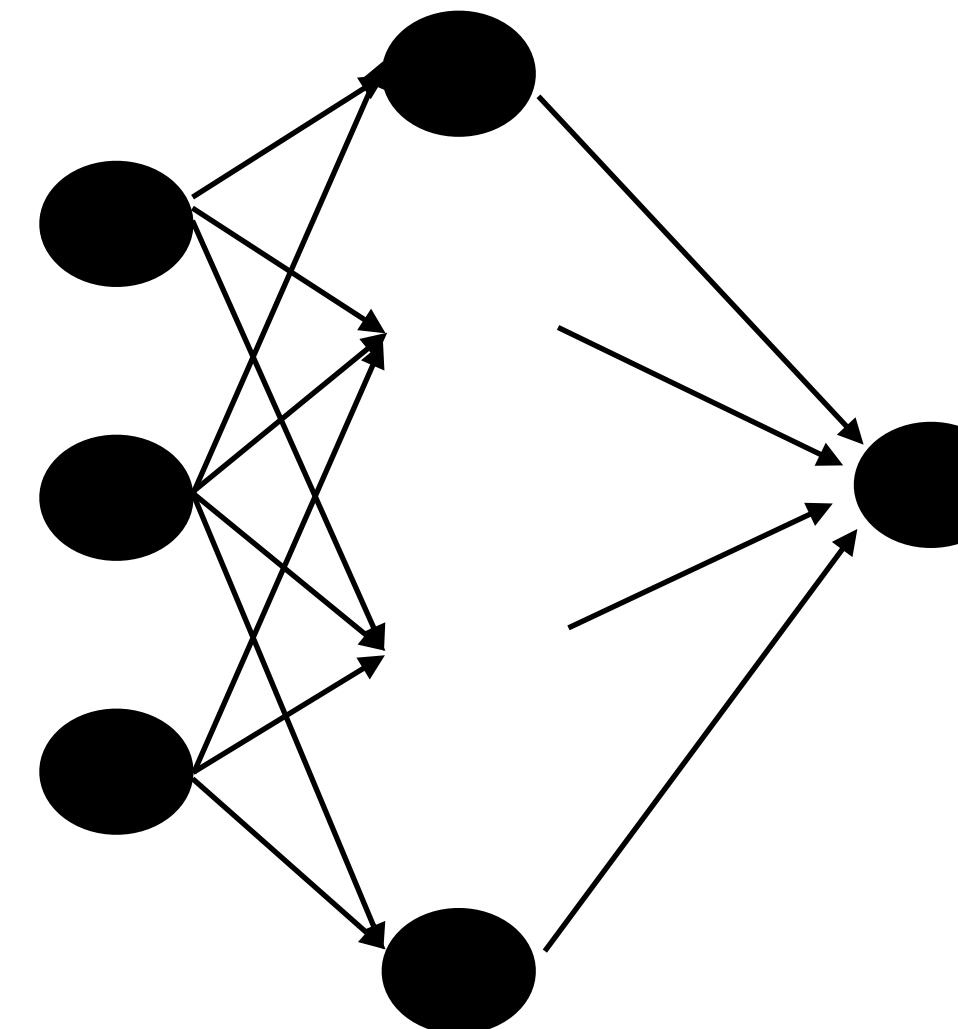
Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 1: Dropout

During training randomly set some activation to 0

Iter 1



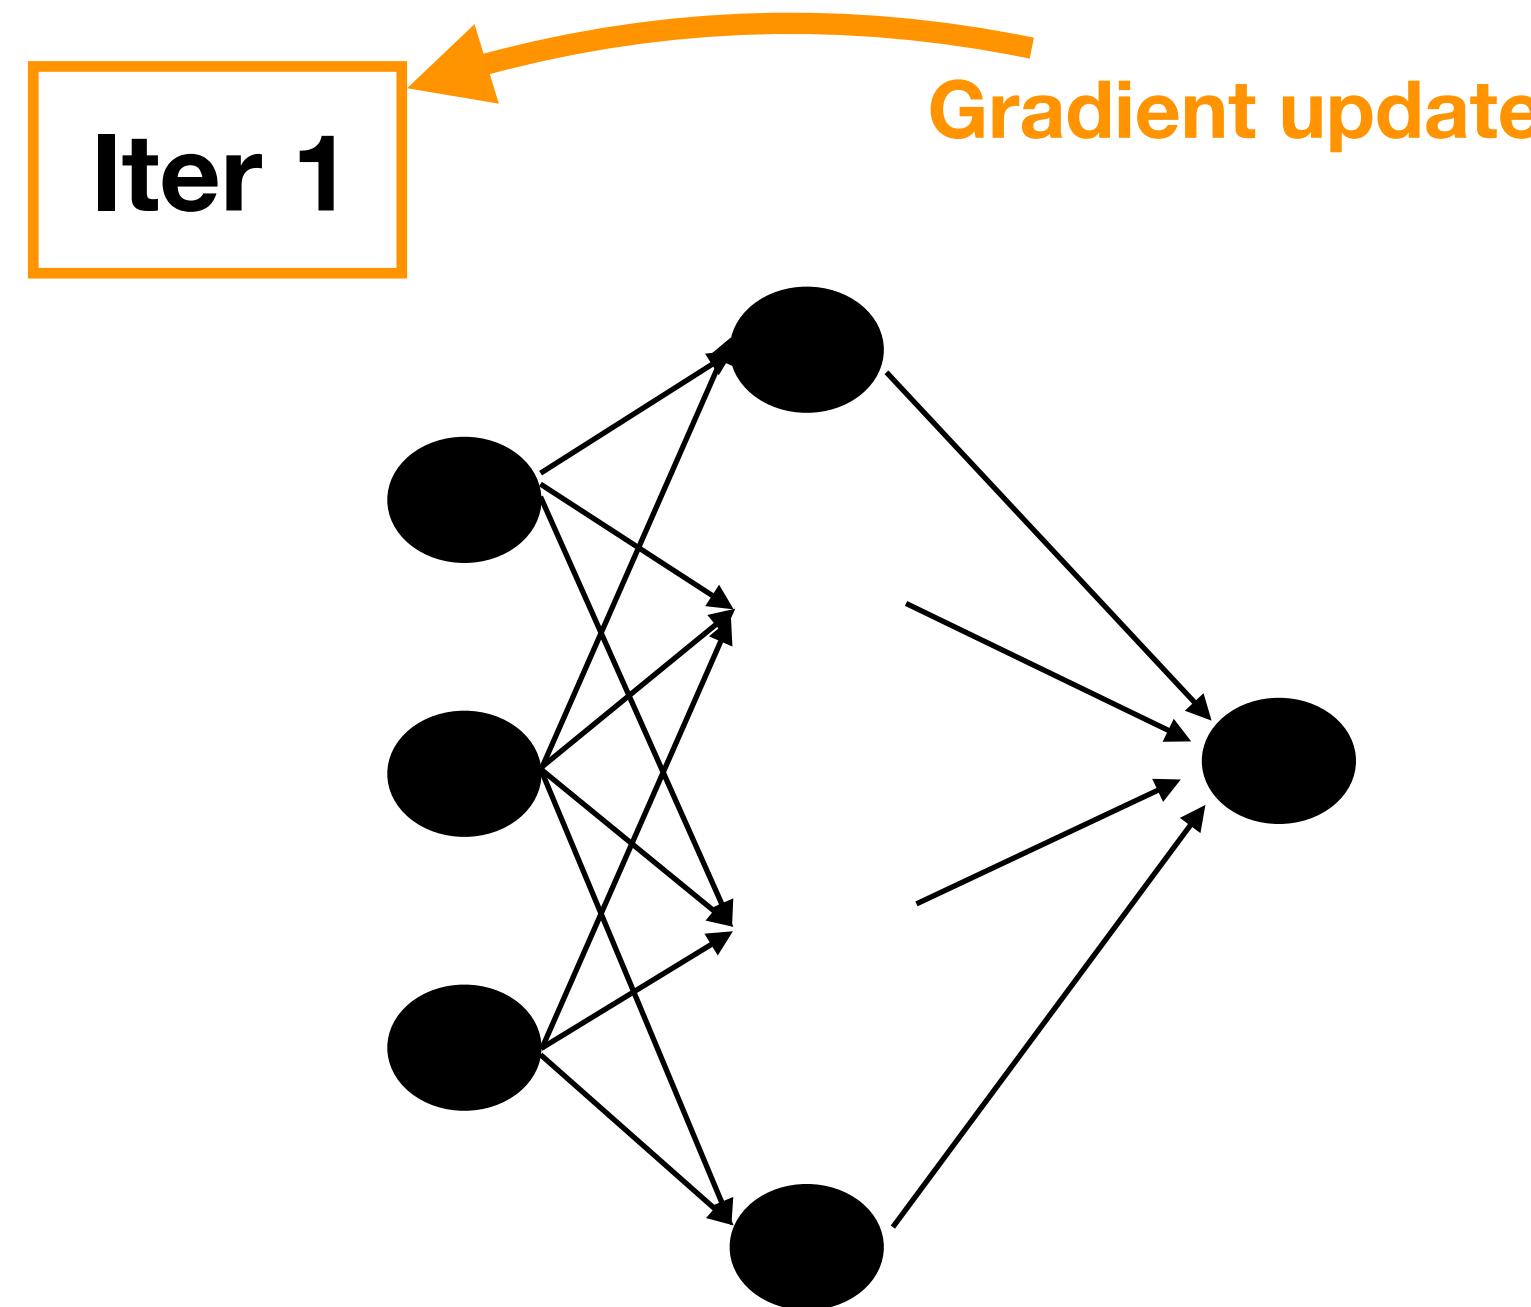
Fighting Overfitting with regularization



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 1: Dropout



During training randomly set some activation to 0

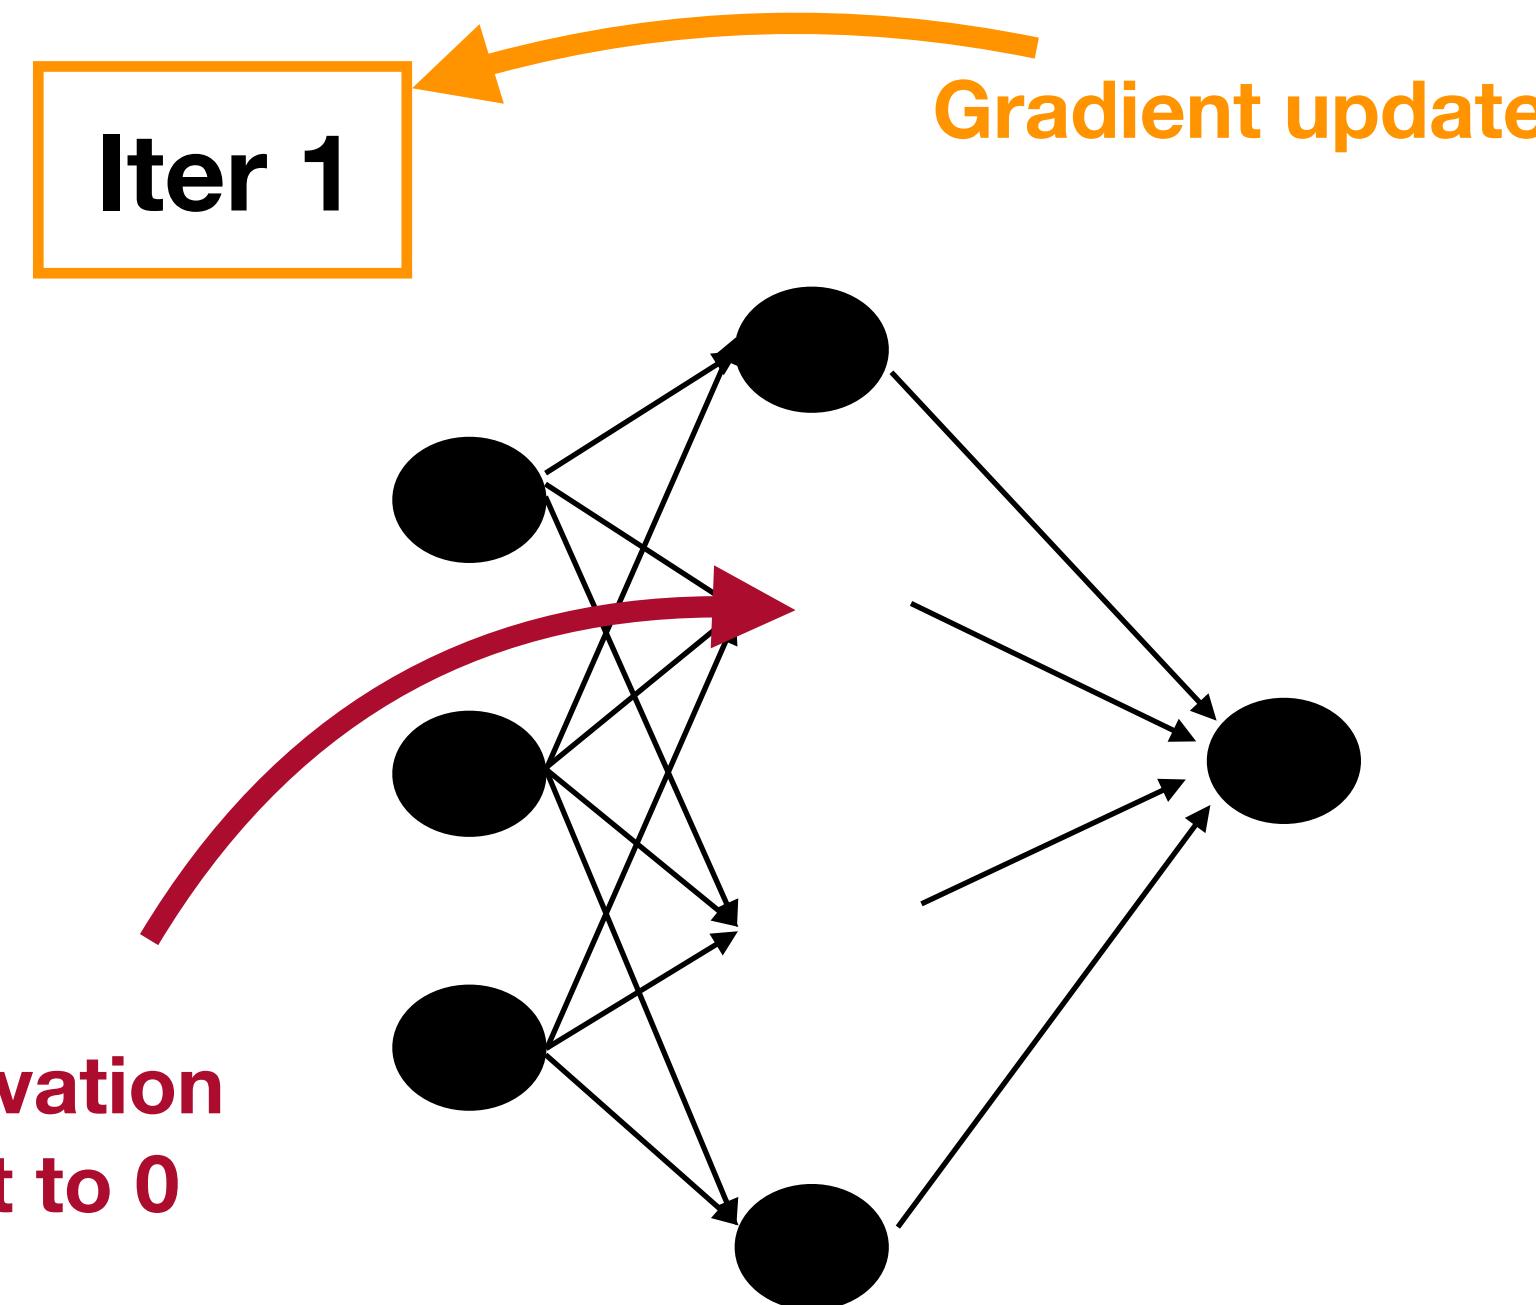
Fighting Overfitting with regularization



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 1: Dropout



During training randomly set some activation to 0

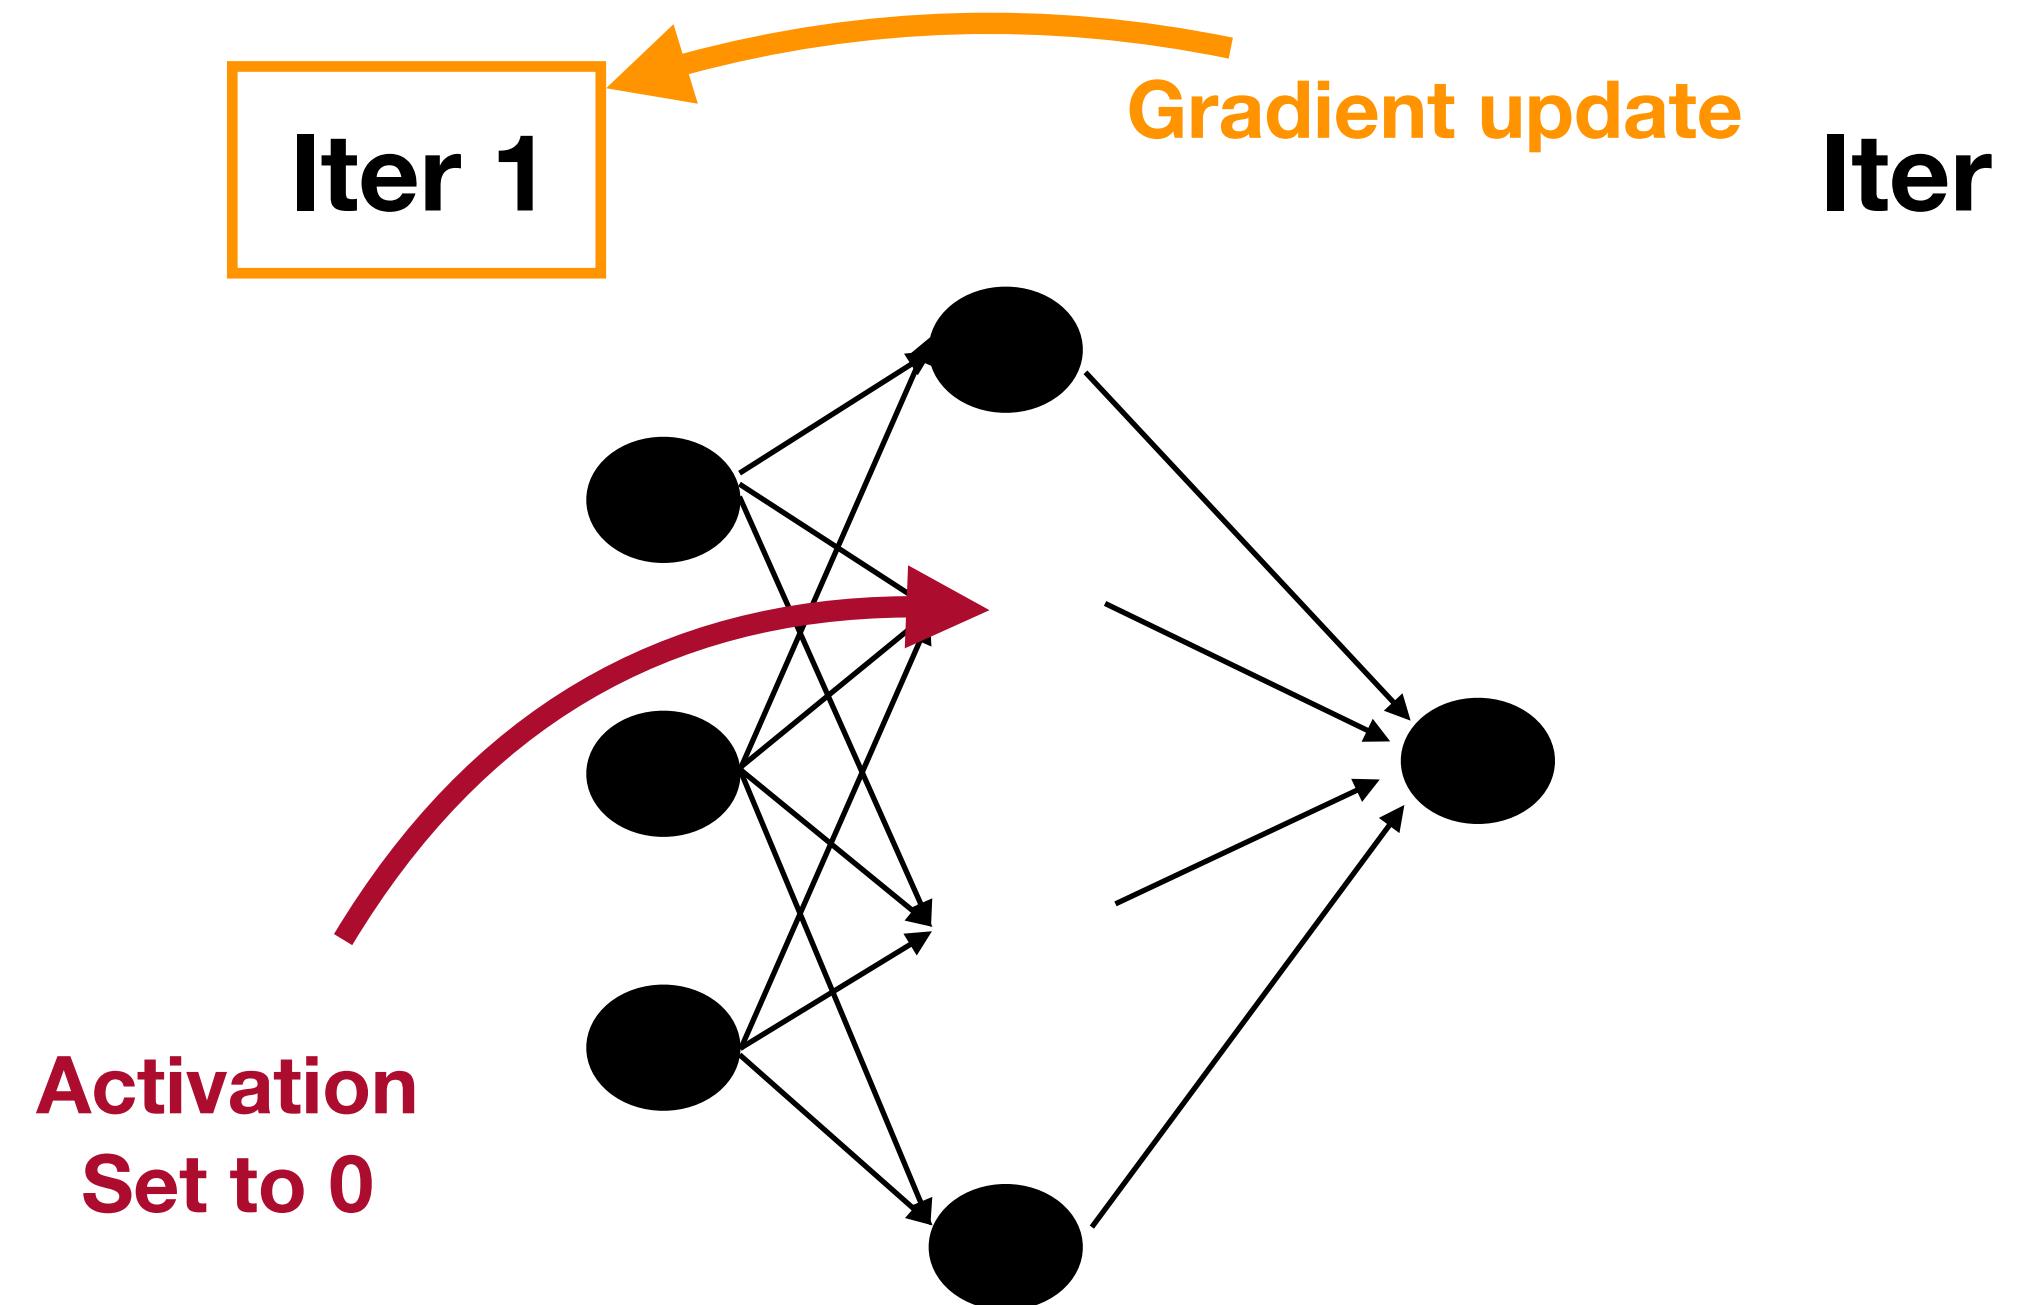
Fighting Overfitting with regularization



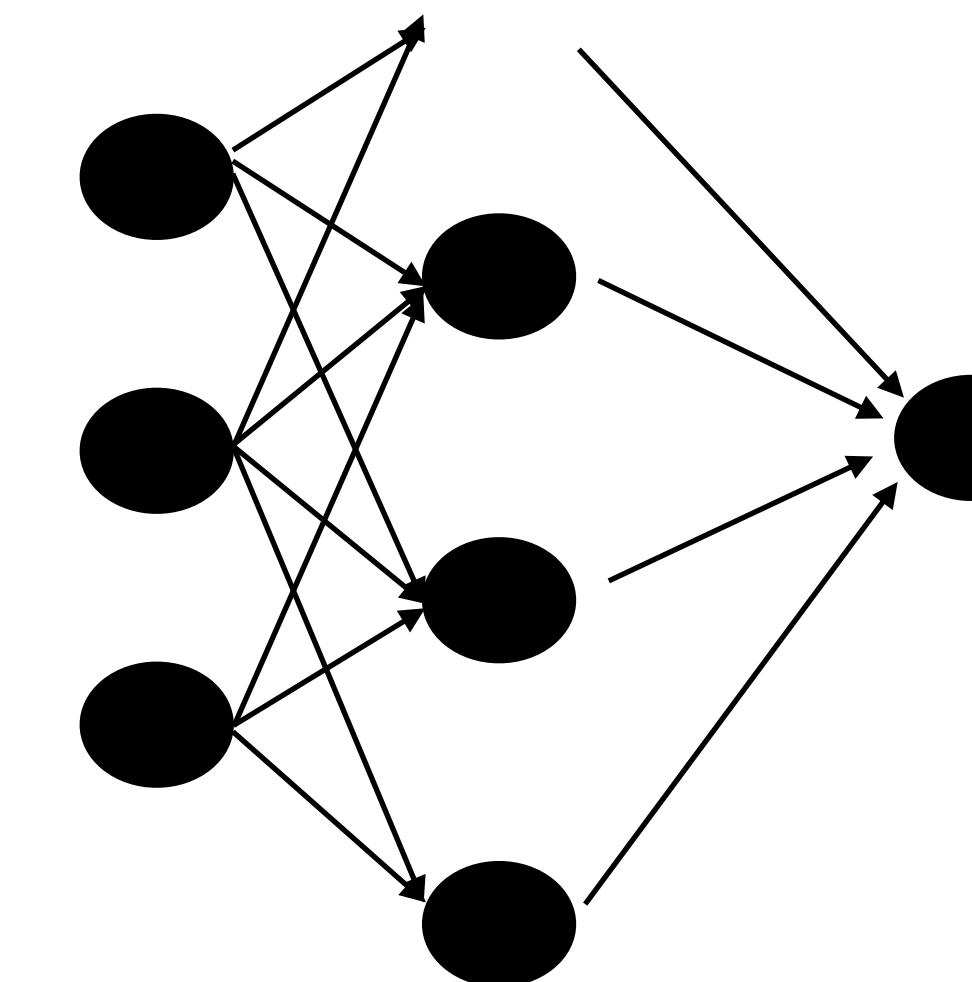
Regularization: the set of techniques to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 1: Dropout



During training randomly set some activation to 0



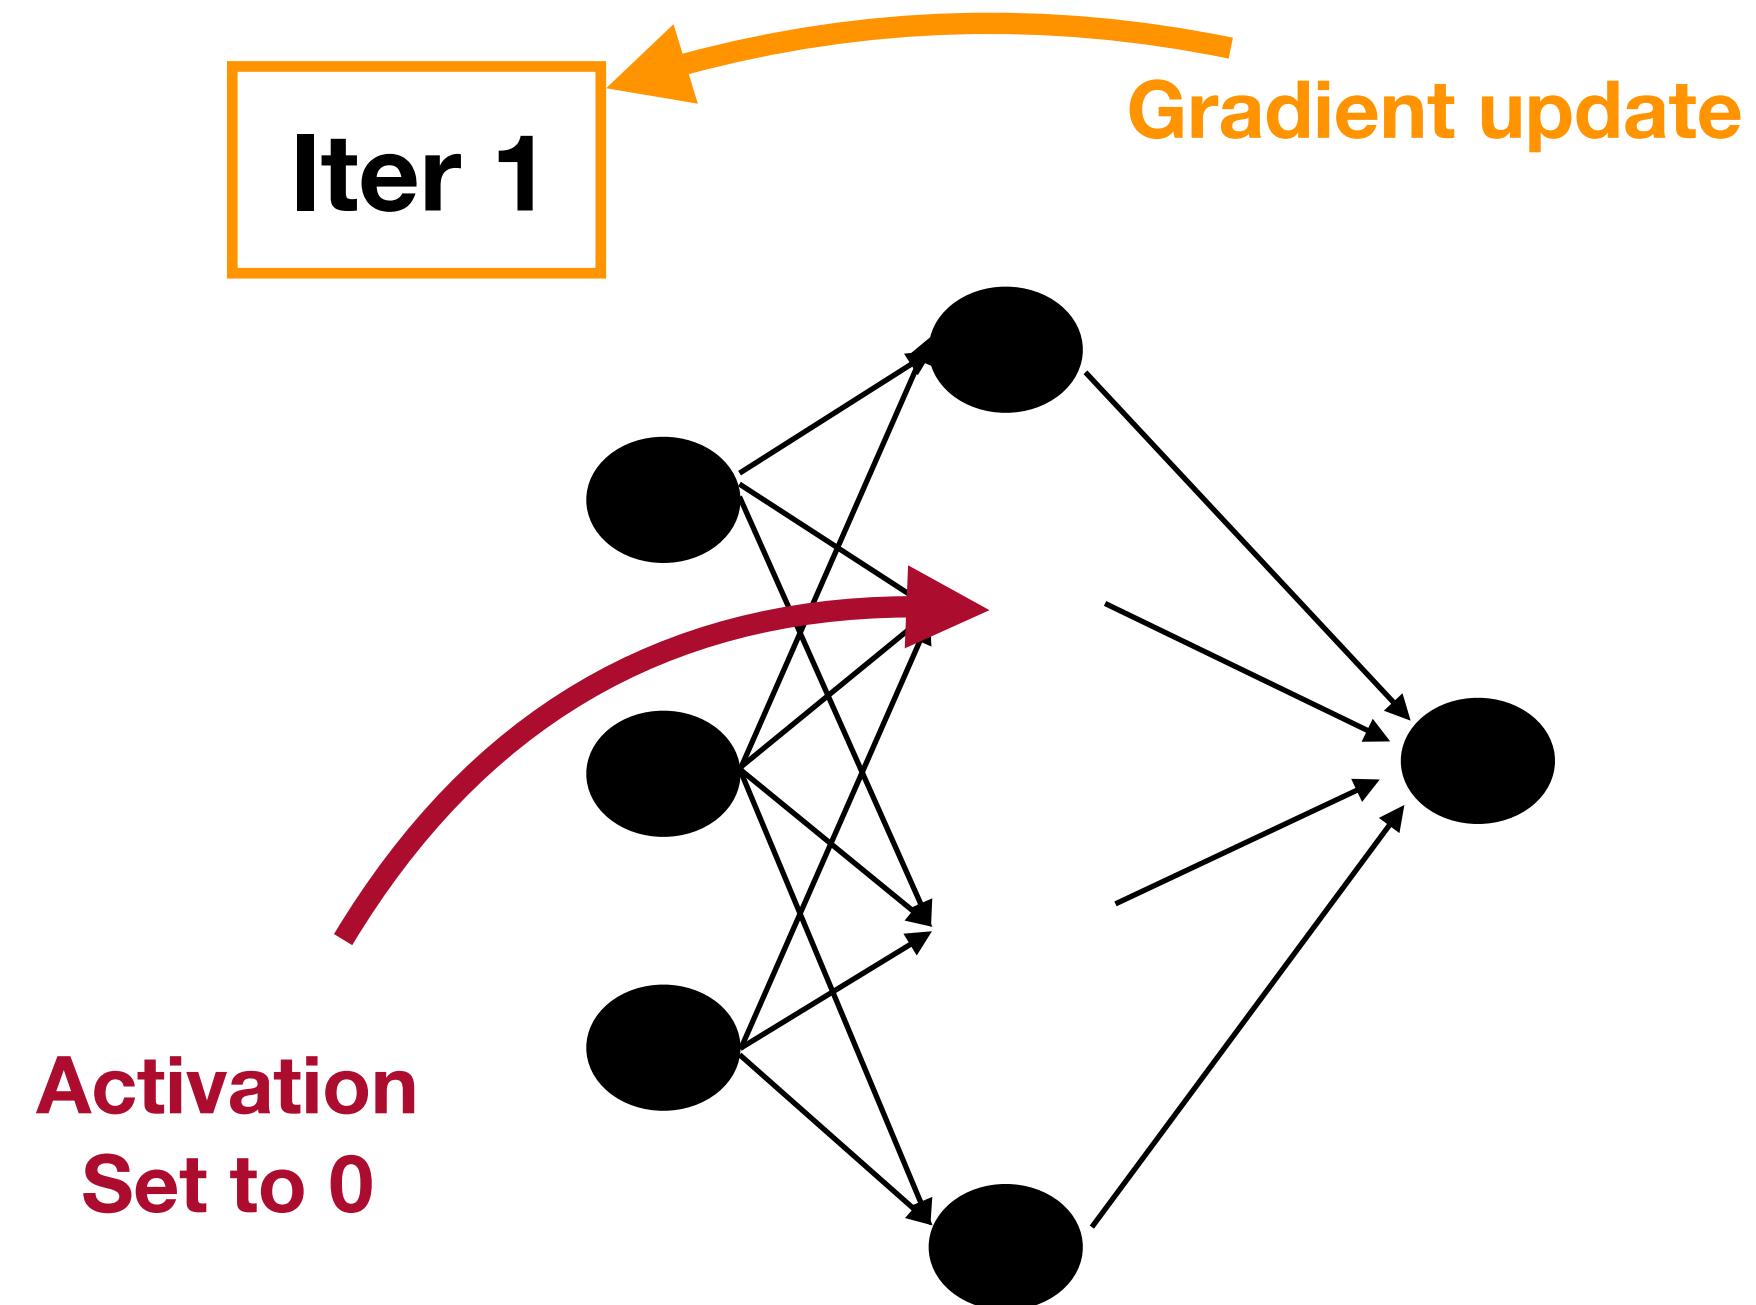
Fighting Overfitting with regularization



Regularization: the set of techniques to prevent or fight overfitting

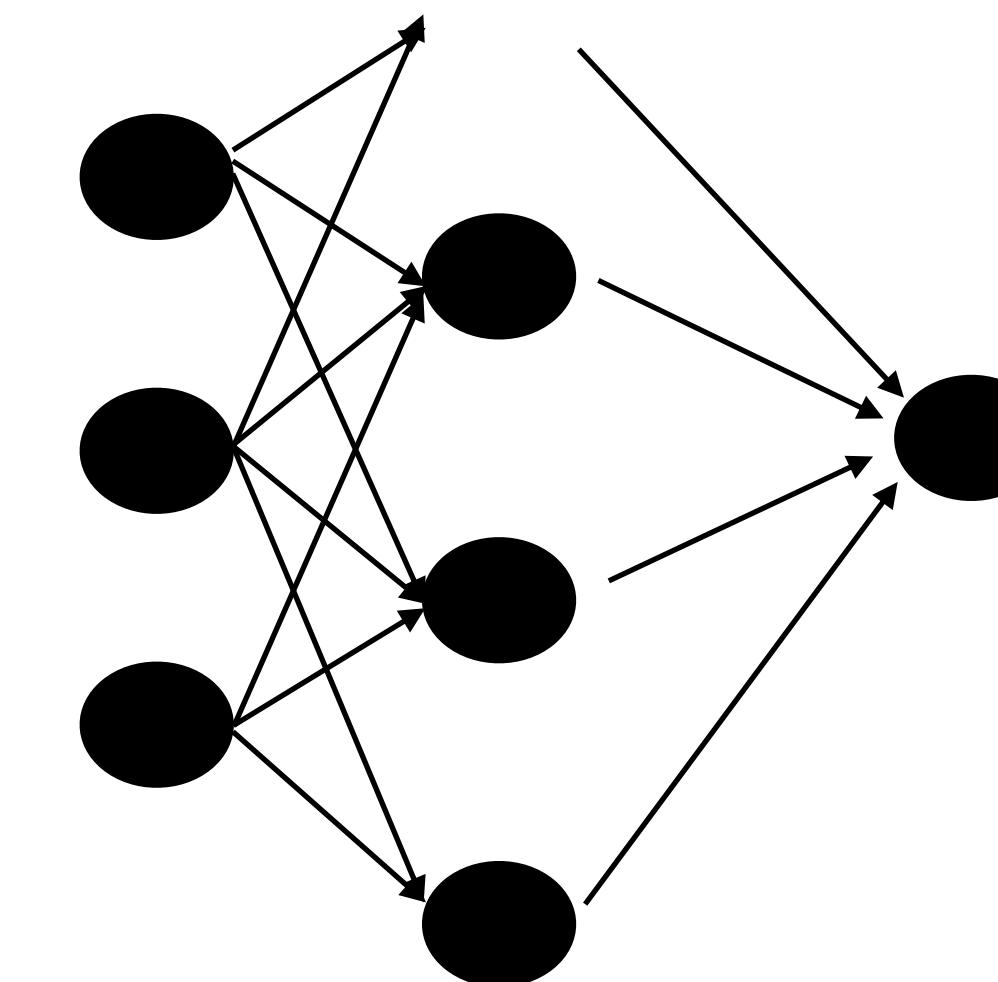
It improves the **generalization capability** of your model (i.e. better results)

Technique 1: Dropout

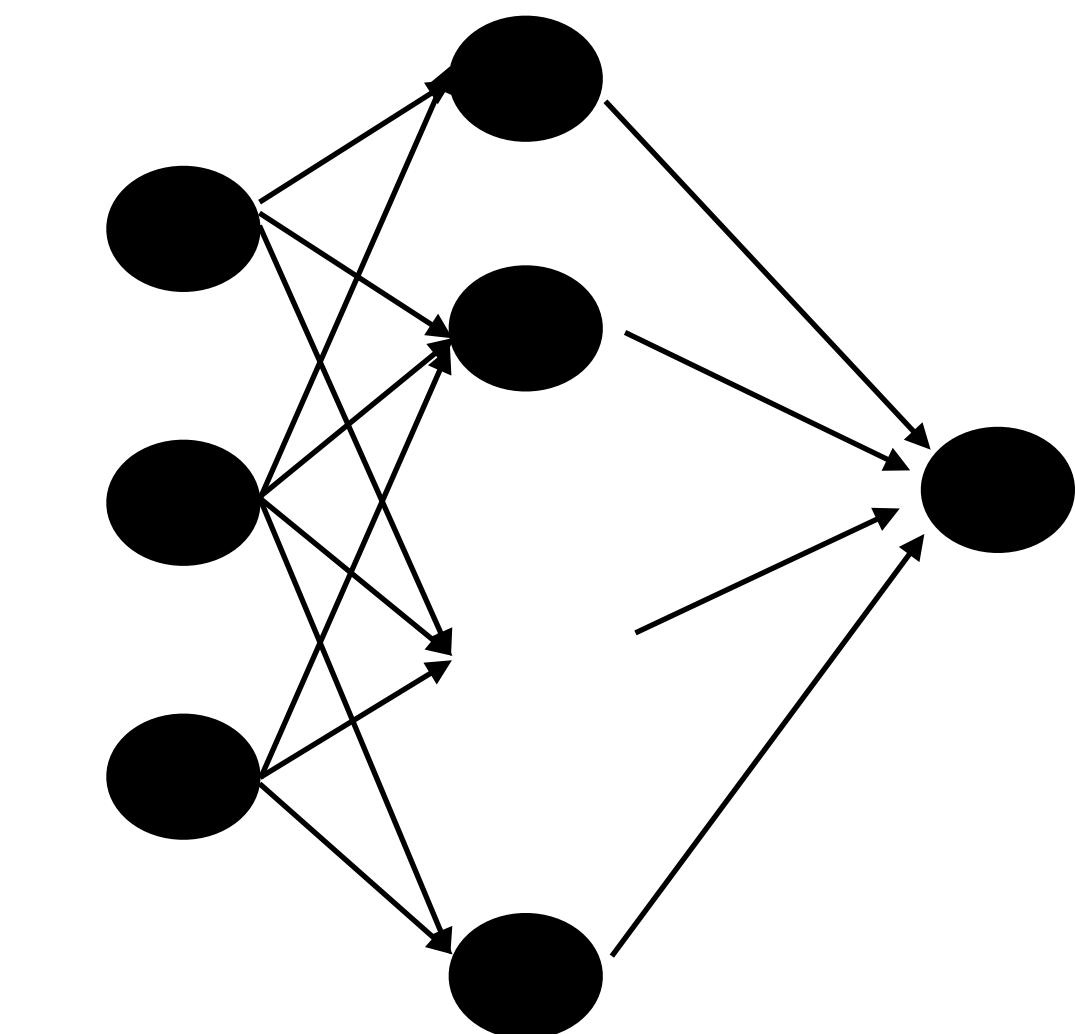


During training randomly set some activation to 0

Iter 2



Iter 3



Fighting Overfitting with regularization



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 2: Early Stopping



Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 2: Early Stopping

During training randomly set some activation to 0

Fighting Overfitting with regularization

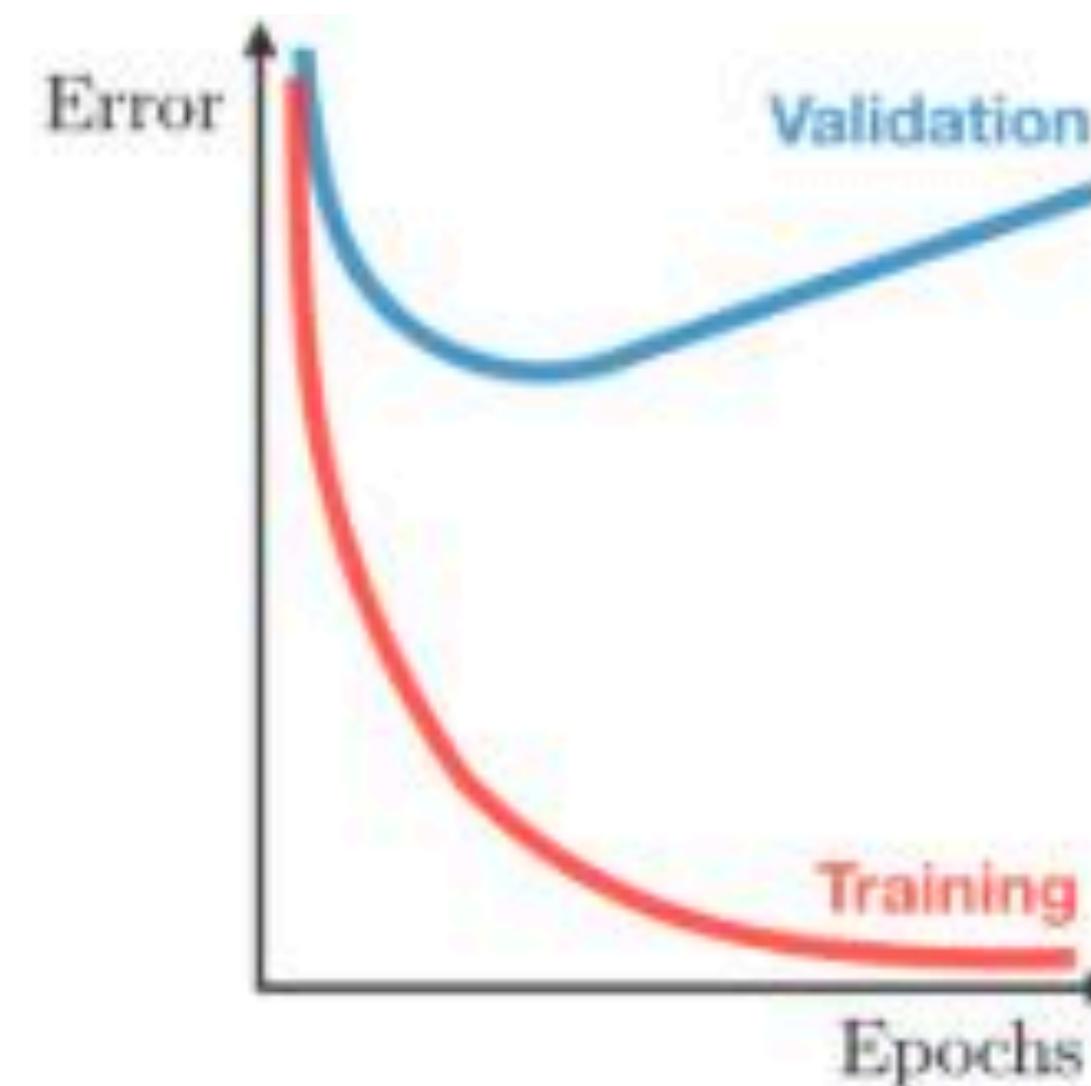


Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 2: Early Stopping

During training randomly set some activation to 0



Fighting Overfitting with regularization

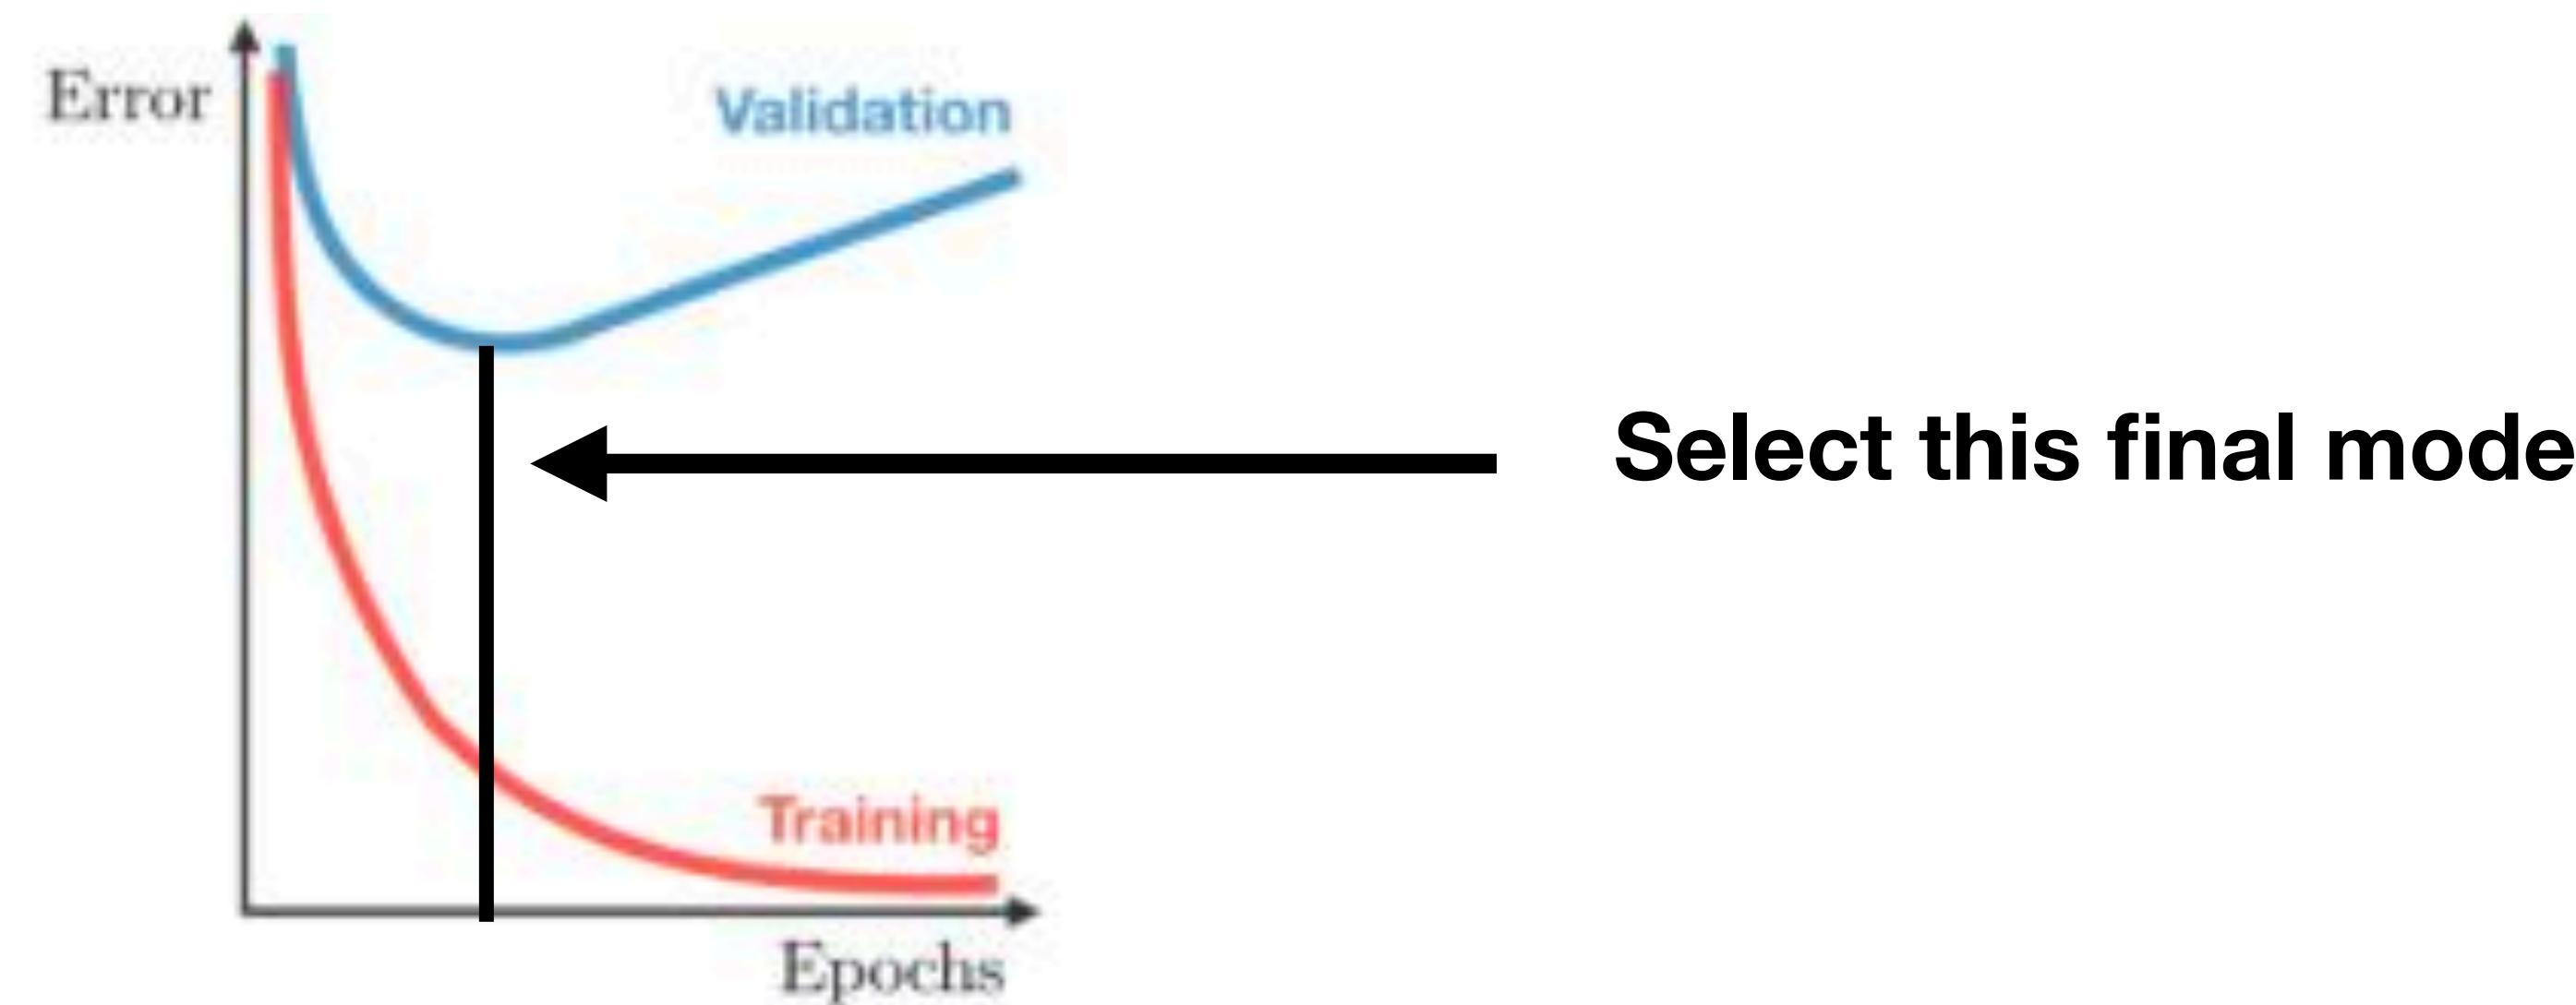


Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 2: Early Stopping

During training randomly set some activation to 0



Fighting Overfitting with regularization

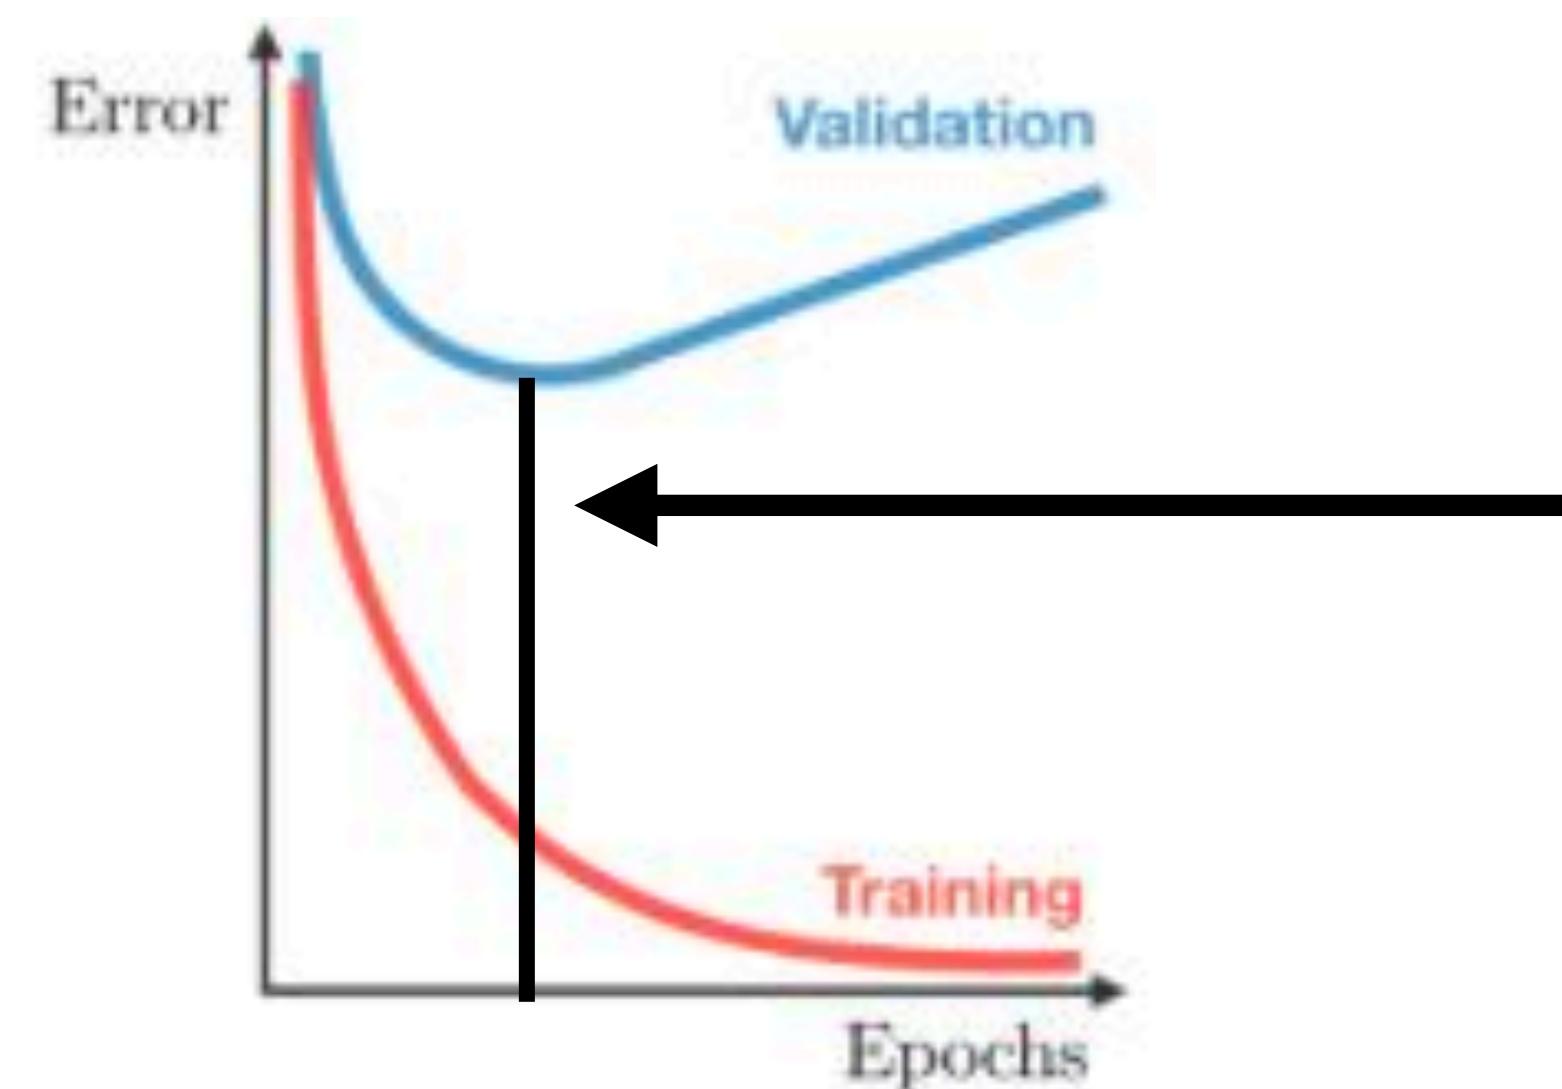


Regularization: the set of technics to prevent or fight overfitting

It improves the generalization capability of your model (i.e. better results)

Technique 2: Early Stopping

During training randomly set some activation to 0



Select this final model

This is the reason you need a validation set!

Let's sum up!



Procedure to use deep learning

Procedure to use deep learning

1. Get your **data clean**

Procedure to use deep learning

1. Get your **data clean**

2. **Split your data into train/val/test split**

Procedure to use deep learning

- 1. Get your **data clean****
- 2. **Split** your data into train/val/test split**
- 3. Define your **architecture/loss function/optimization algorithm****
- 4. Train and tune your algorithm by monitoring training/validation loss**

Procedure to use deep learning

- 1. Get your **data clean****
- 2. **Split** your data into train/val/test split**
- 3. Define your **architecture/loss function/optimization algorithm****
- 4. Train and tune your algorithm by monitoring training/validation loss**
- 5. Test your model on test set**

Procedure to use deep learning

- 1. Get your **data clean****
- 2. **Split** your data into train/val/test split**
- 3. Define your **architecture/loss function/optimization algorithm****
- 4. Train and tune your algorithm by monitoring training/validation loss**
- 5. Test your model on test set**
- 6. Deploy your model in production**

Procedure to use deep learning

1. Get your **data clean**

2. **Split** your data into train/val/test split

Training Loop

3. Define your **architecture/loss function/optimization algorithm**

4. Train and tune your algorithm by monitoring training/validation loss

5. Test your model on test set

6. Deploy your model in production

Procedure to use deep learning

1. Get your **data clean**

80% of the job is done here!

2. Split your data into train/val/test split

Training Loop

3. Define your **architecture/loss function/optimization algorithm**

4. Train and tune your algorithm by monitoring training/validation loss

5. Test your model on test set

6. Deploy your model in production

Data

Deep Learning Architecture

Putting Everything Together: Training

Data



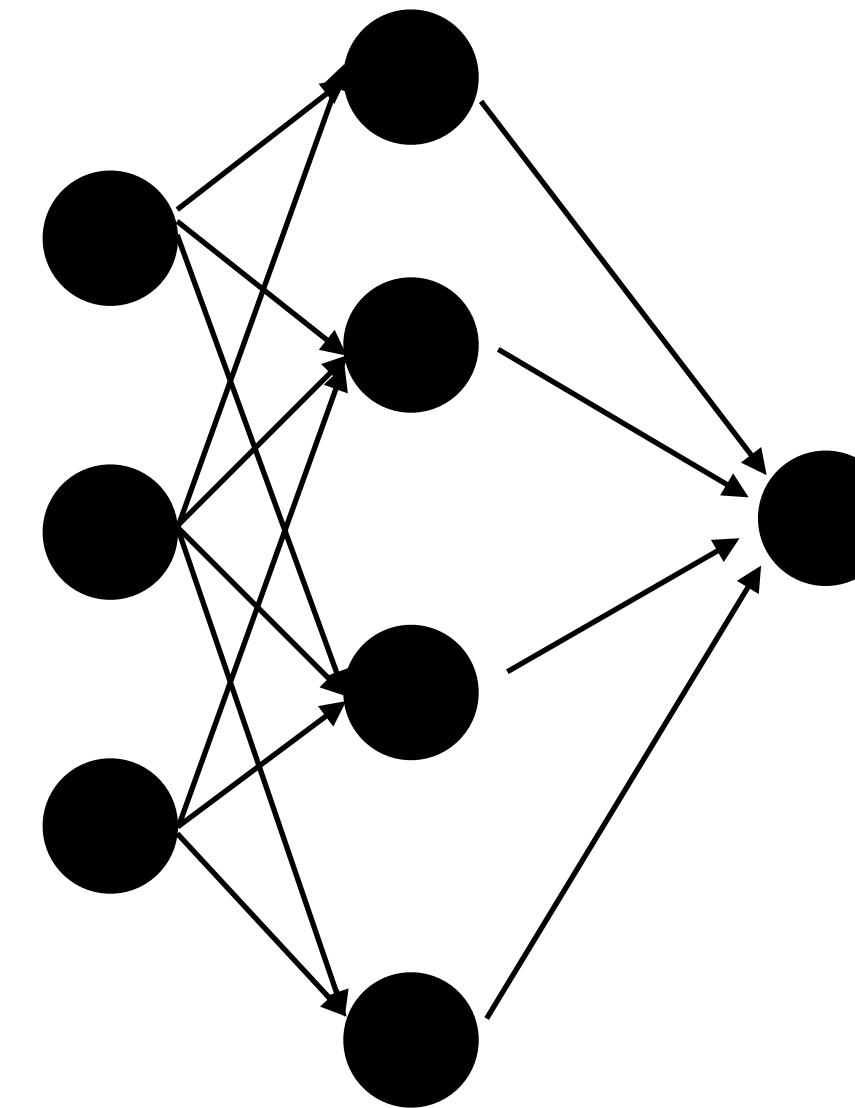
Deep Learning Architecture

Putting Everything Together: Training

Data



Deep Learning Architecture

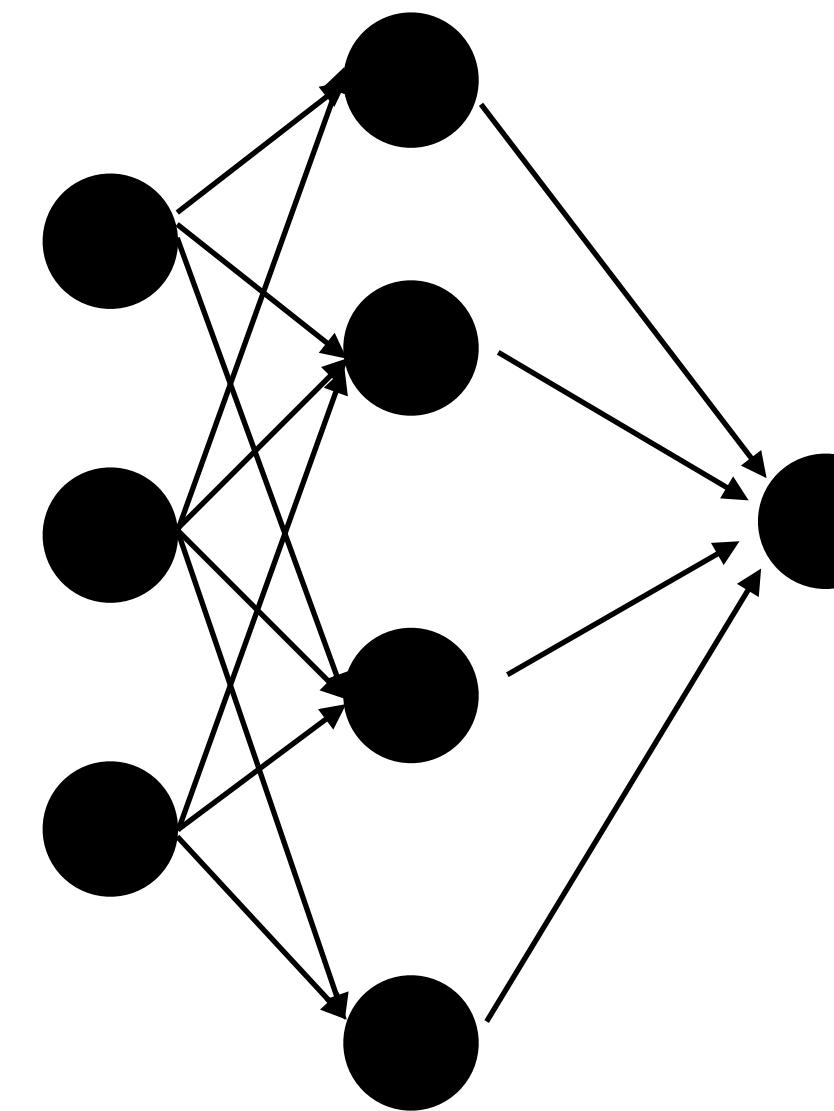


Putting Everything Together: Training

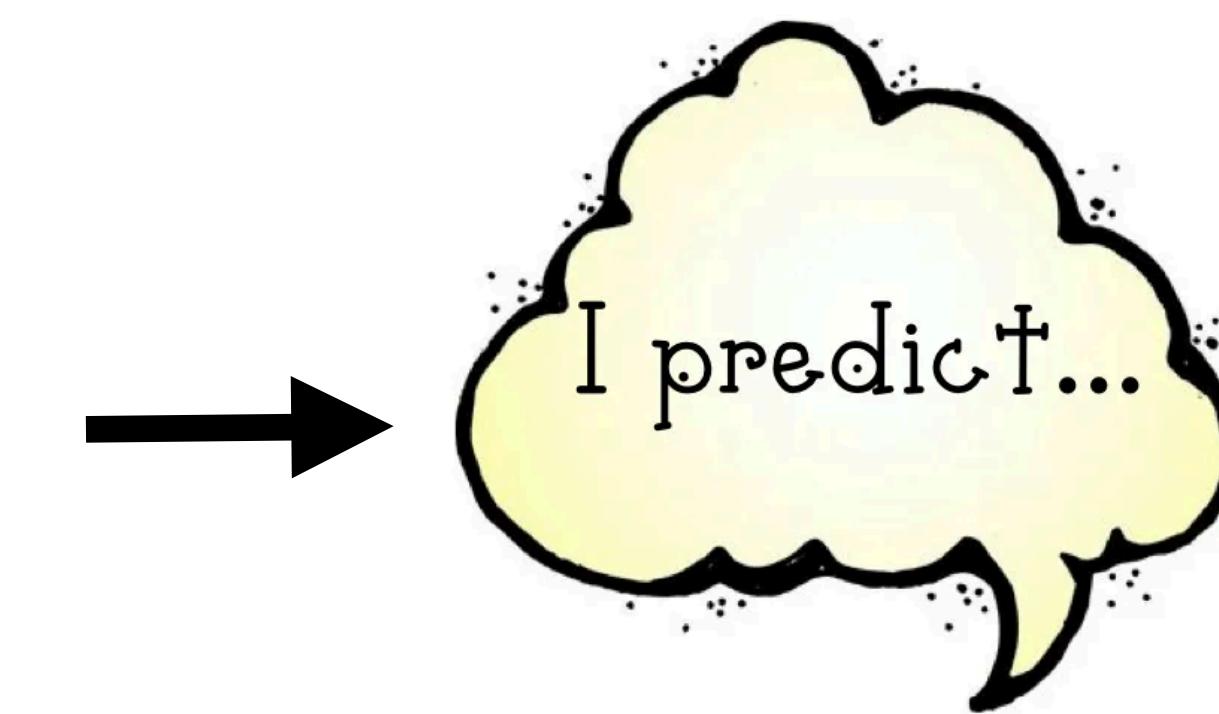
Data



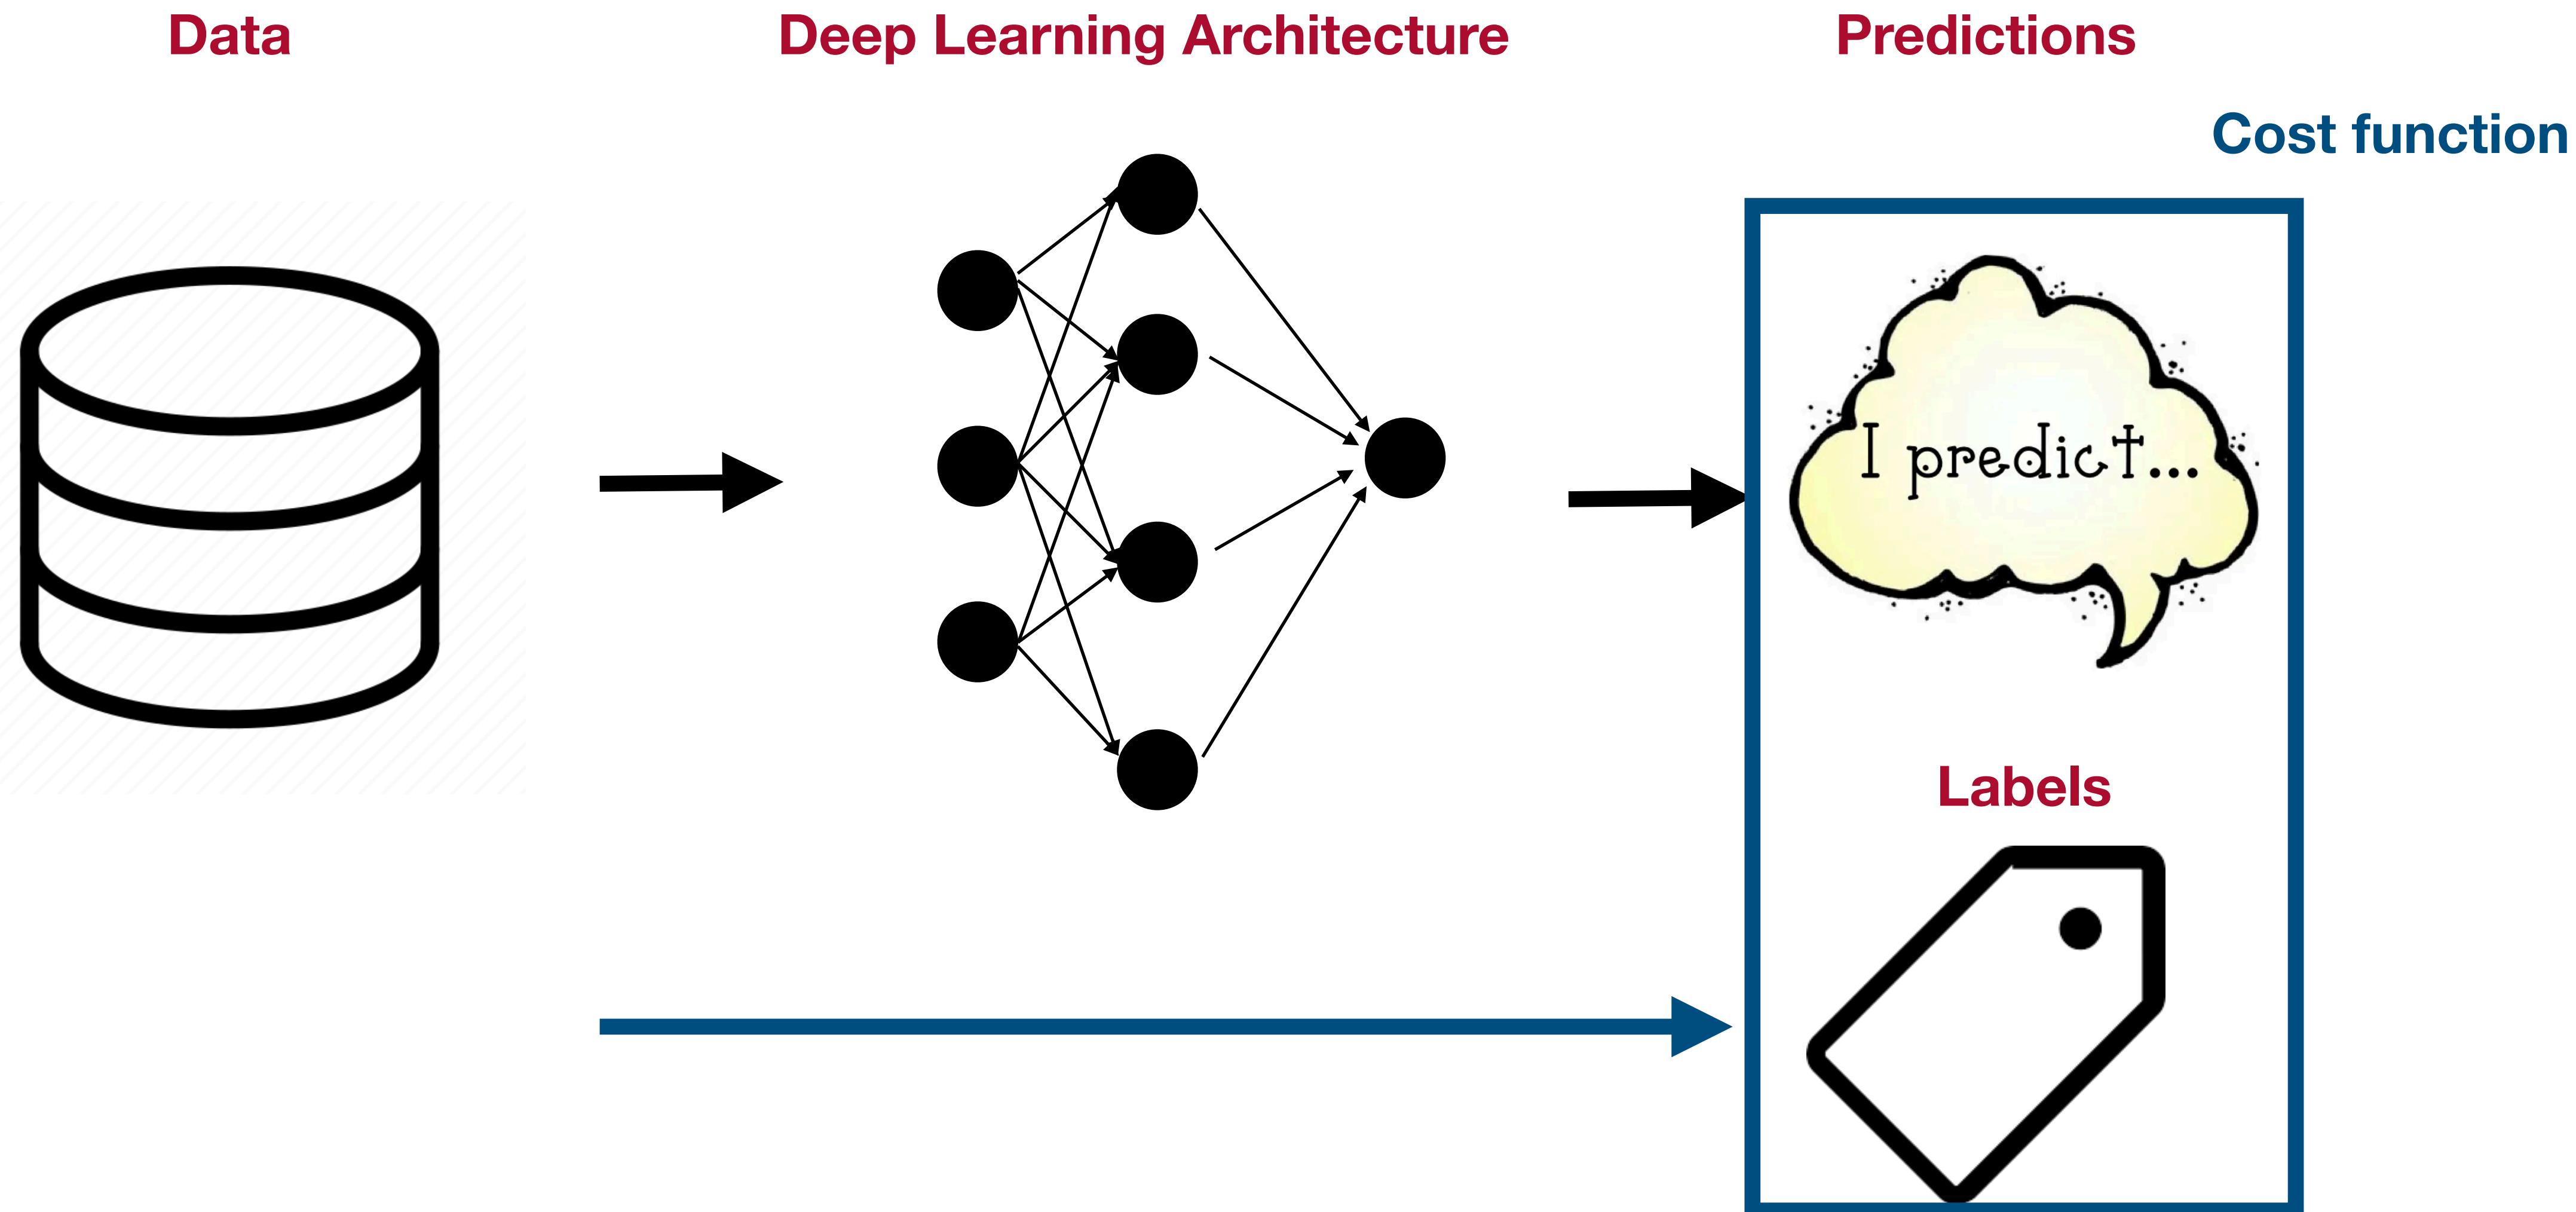
Deep Learning Architecture



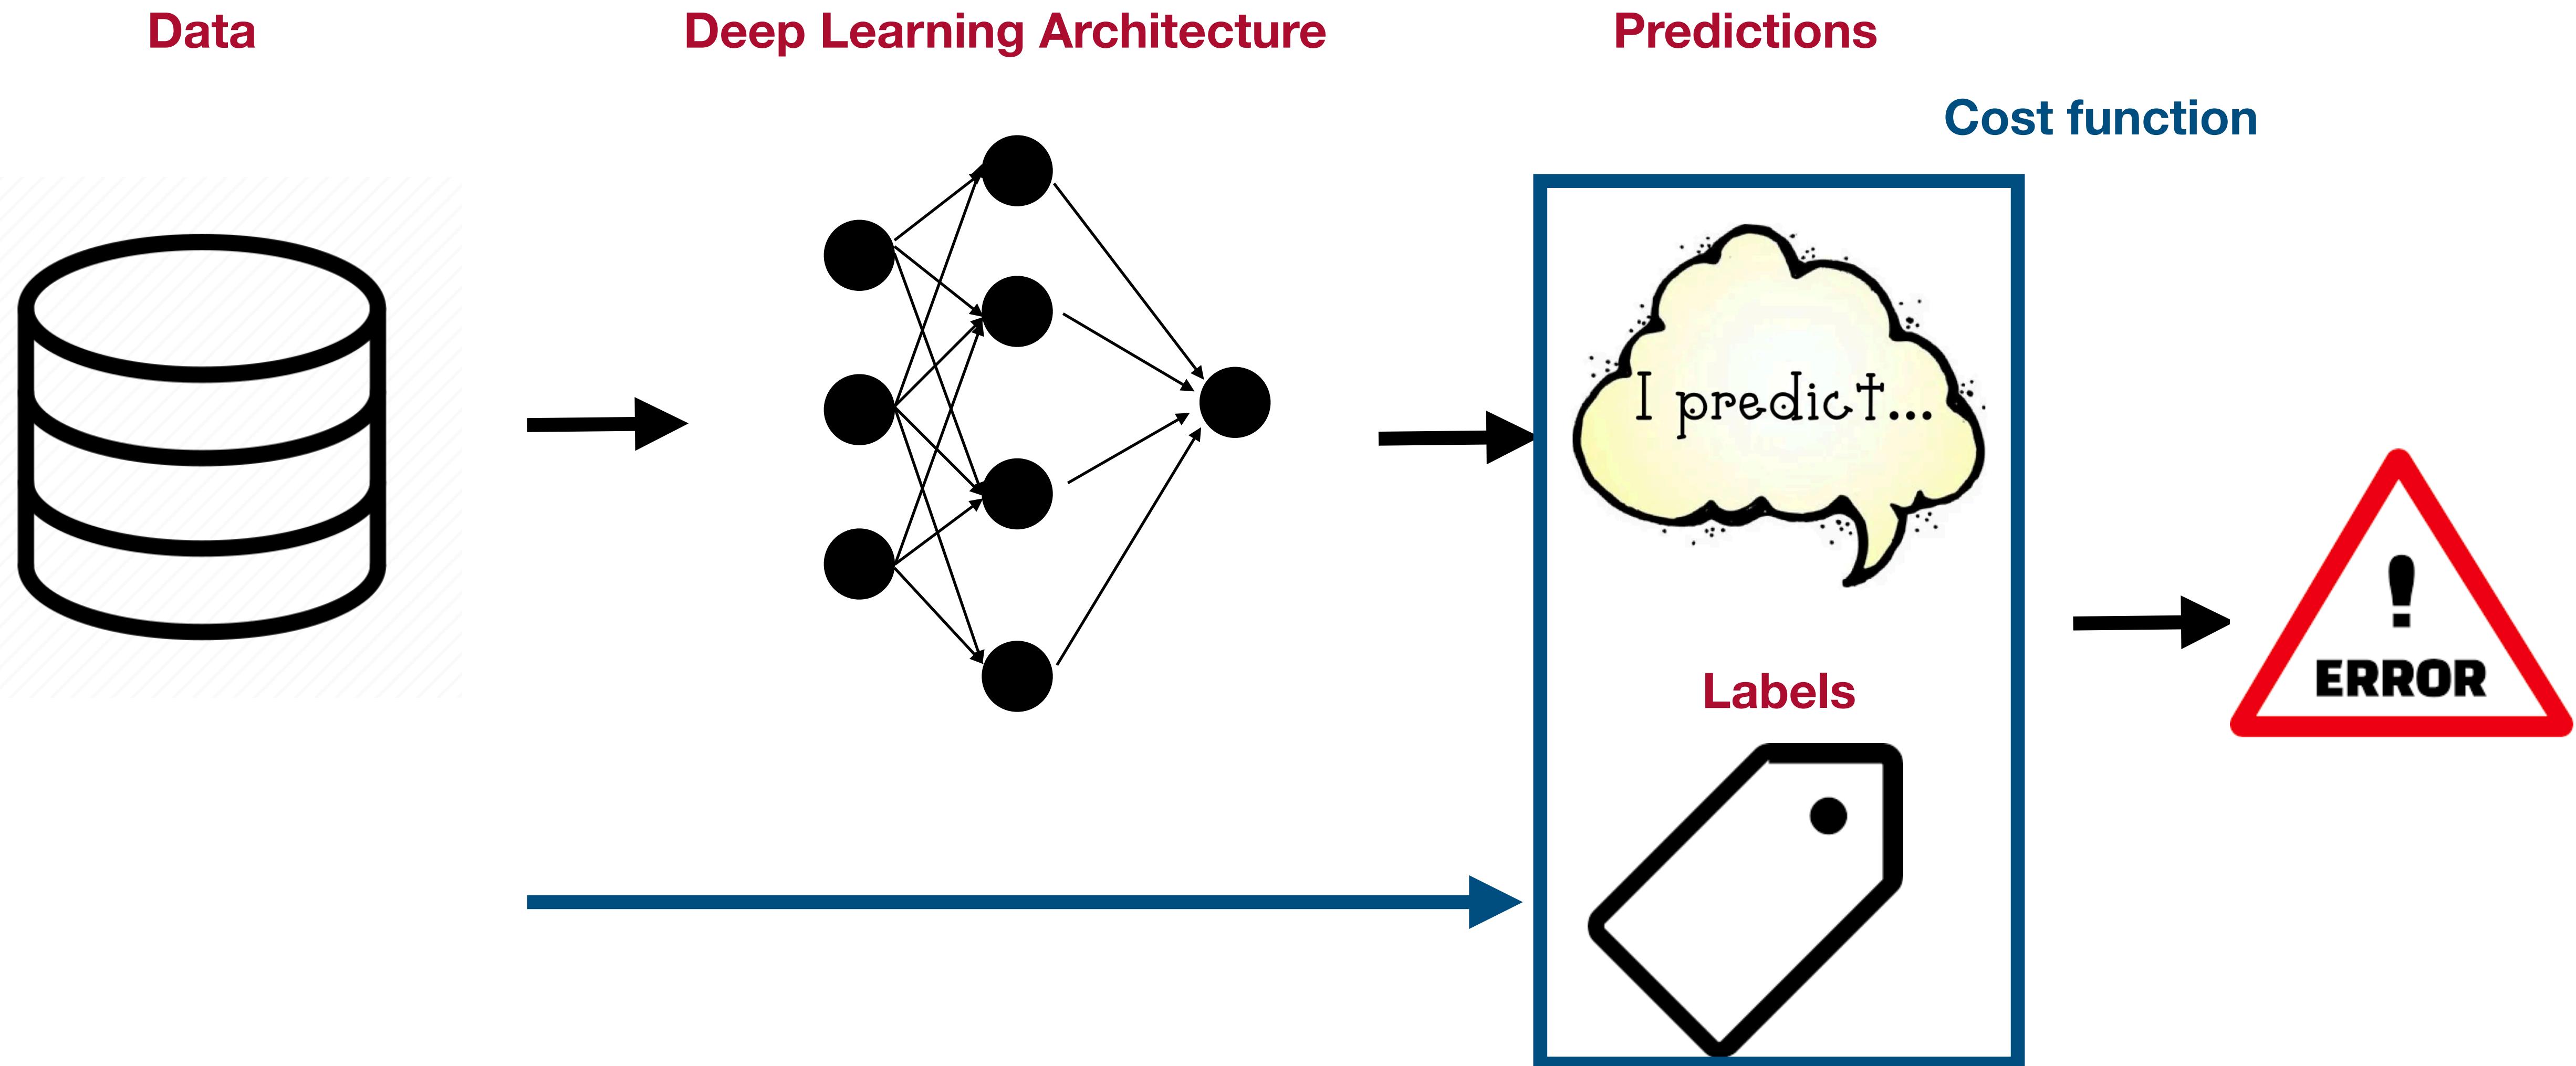
Predictions



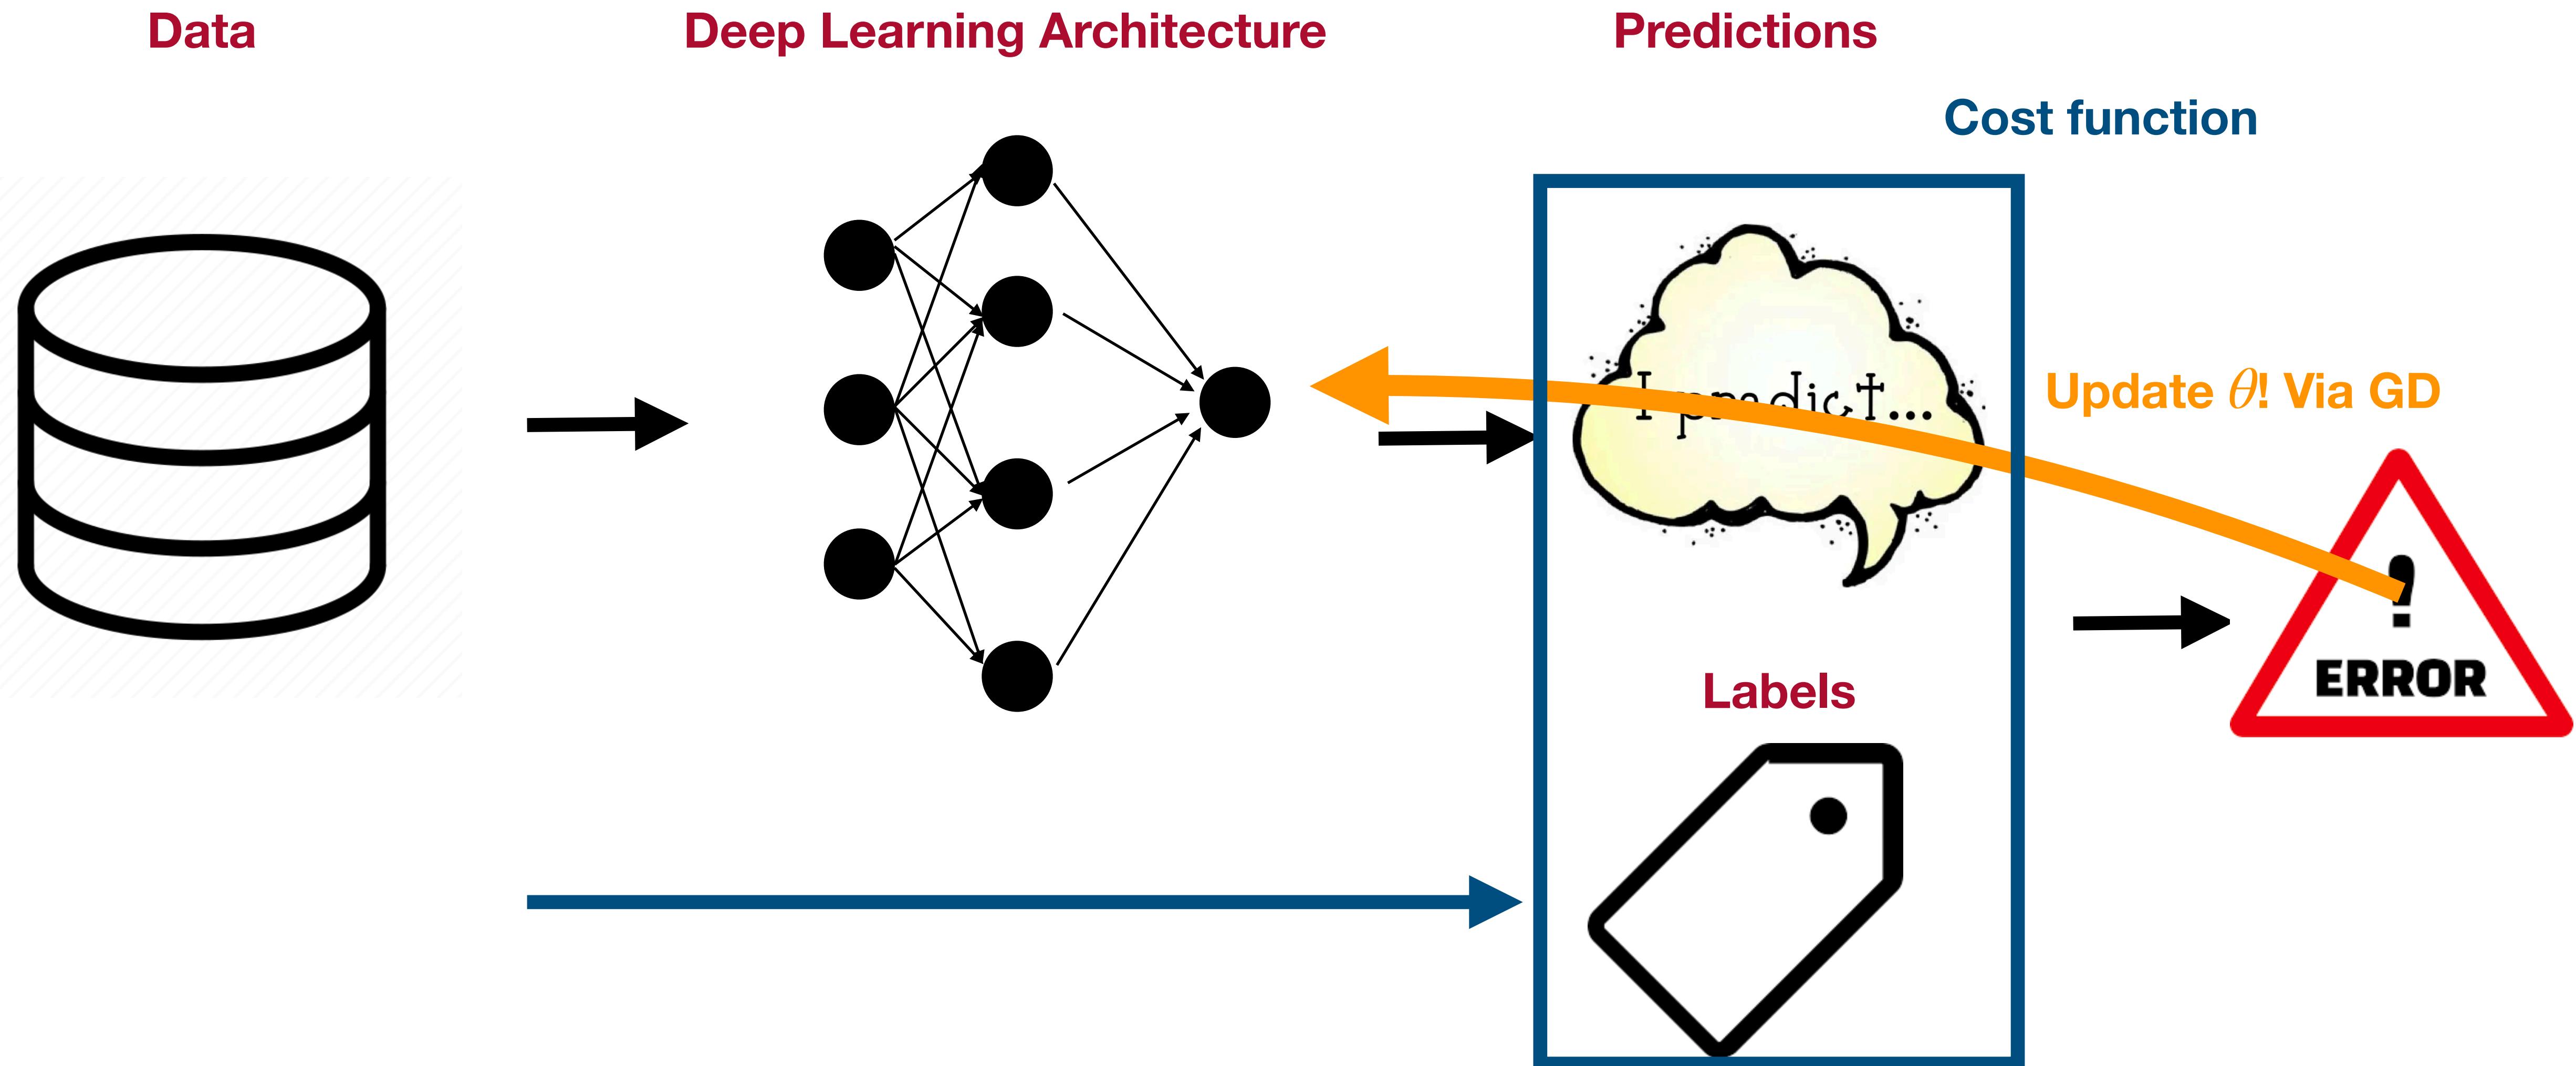
Putting Everything Together: Training



Putting Everything Together: Training



Putting Everything Together: Training



Putting Everything Together: Inference (deployment mode)

Data

Deep Learning Architecture

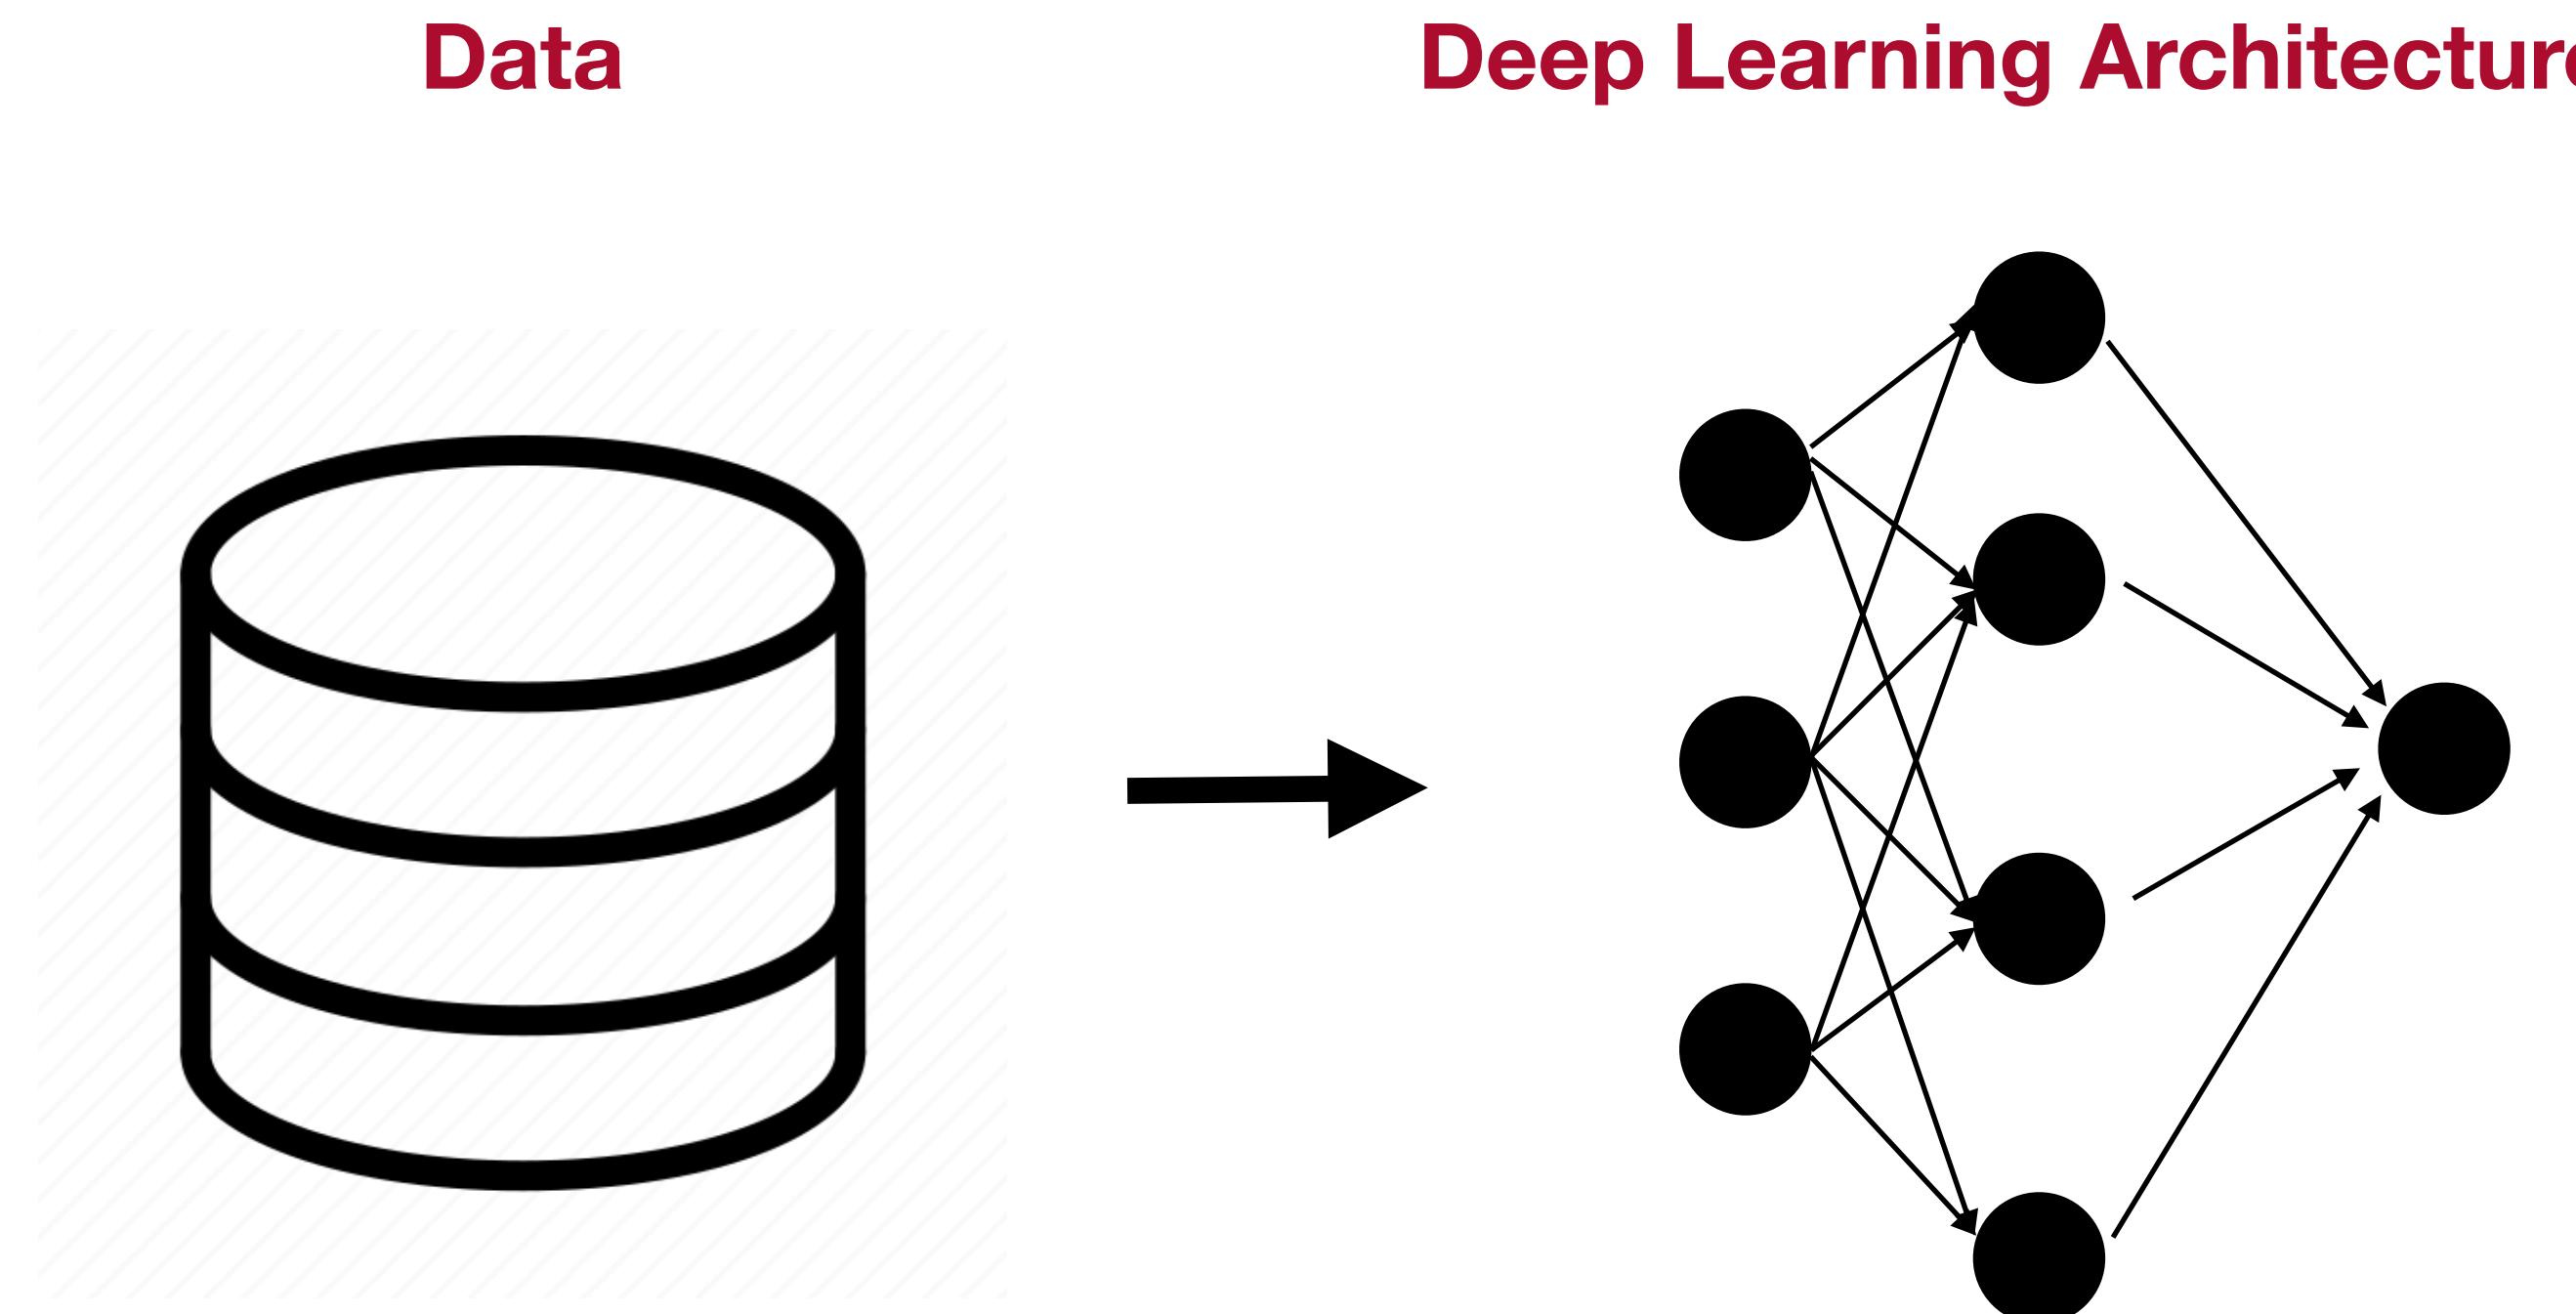
Putting Everything Together: Inference (deployment mode)

Data

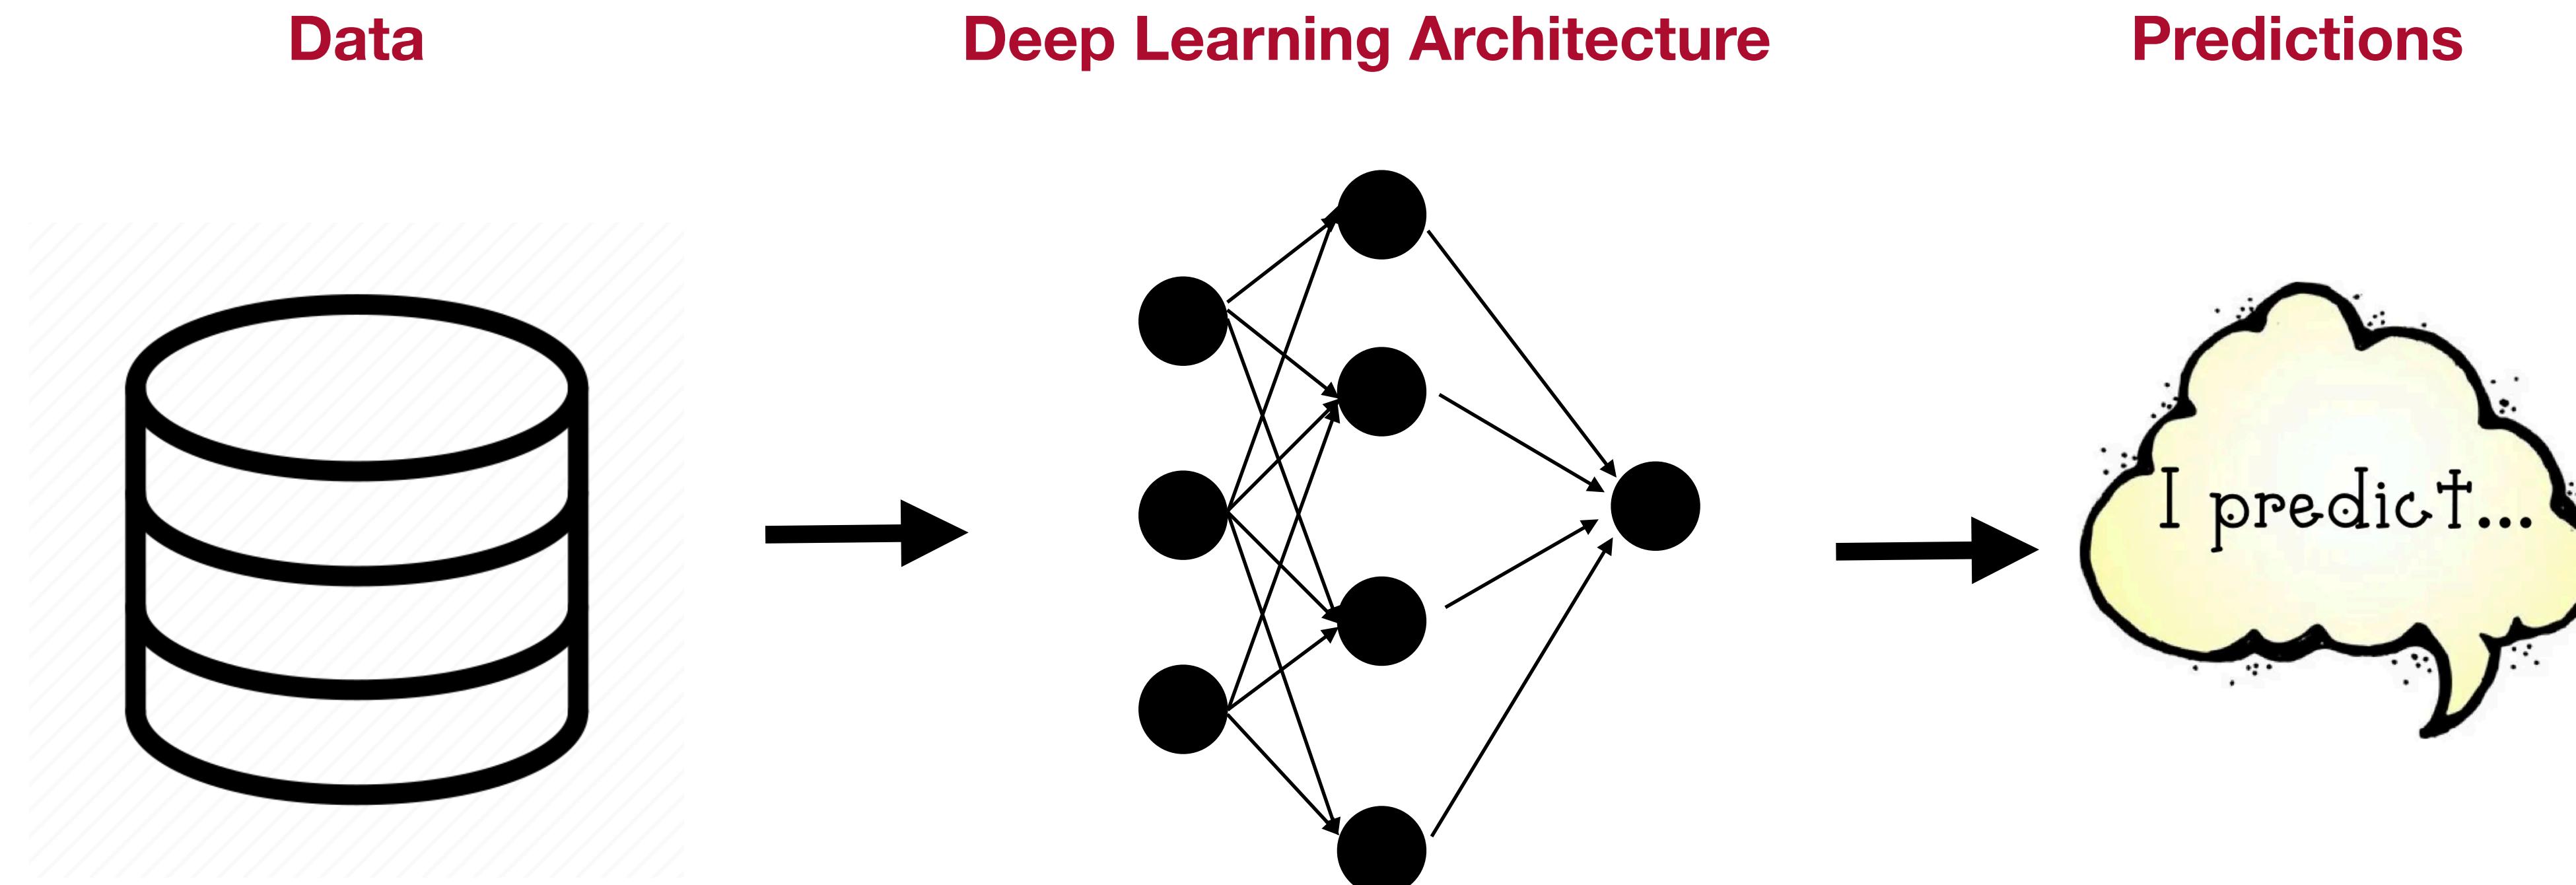


Deep Learning Architecture

Putting Everything Together: Inference (deployment mode)



Putting Everything Together: Inference (deployment mode)



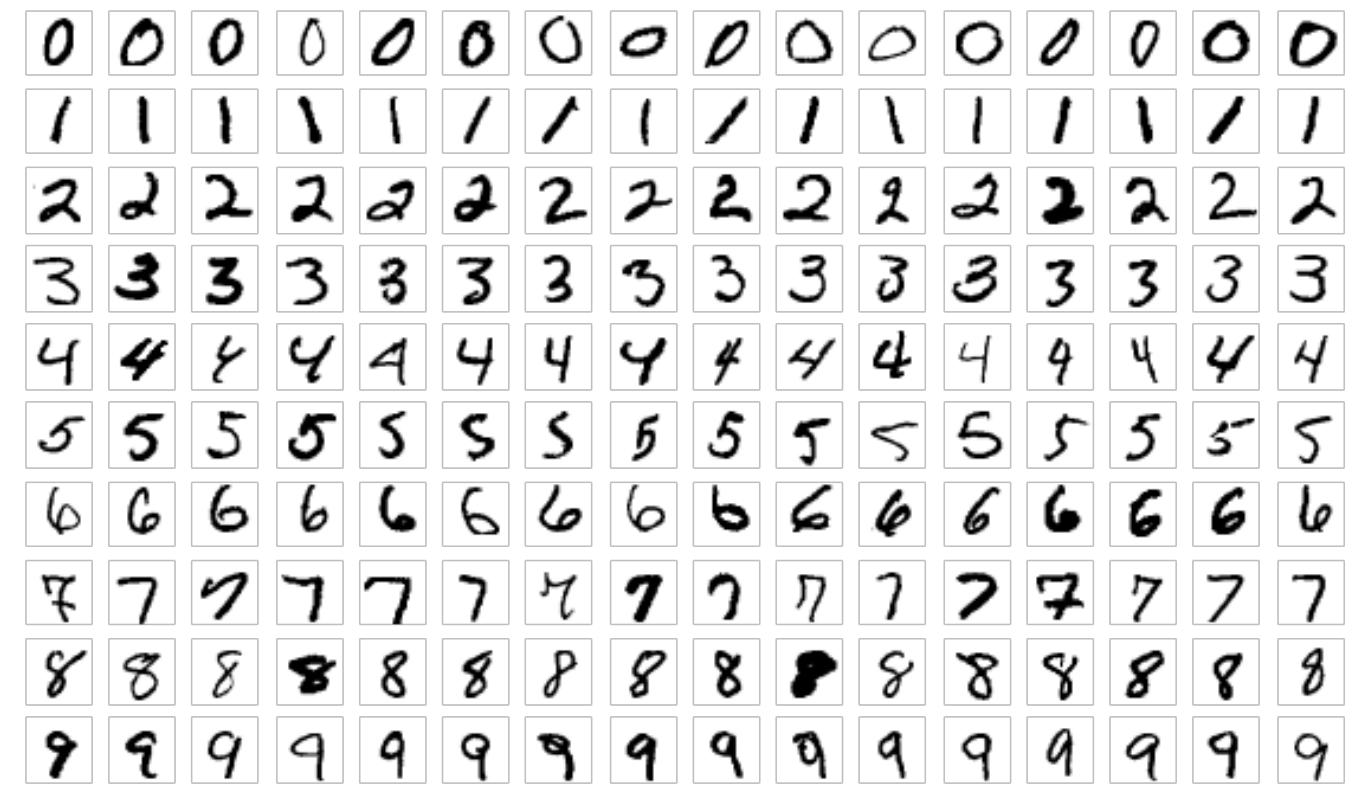
Let's code now



What will we code?

What will we code?

MNIST

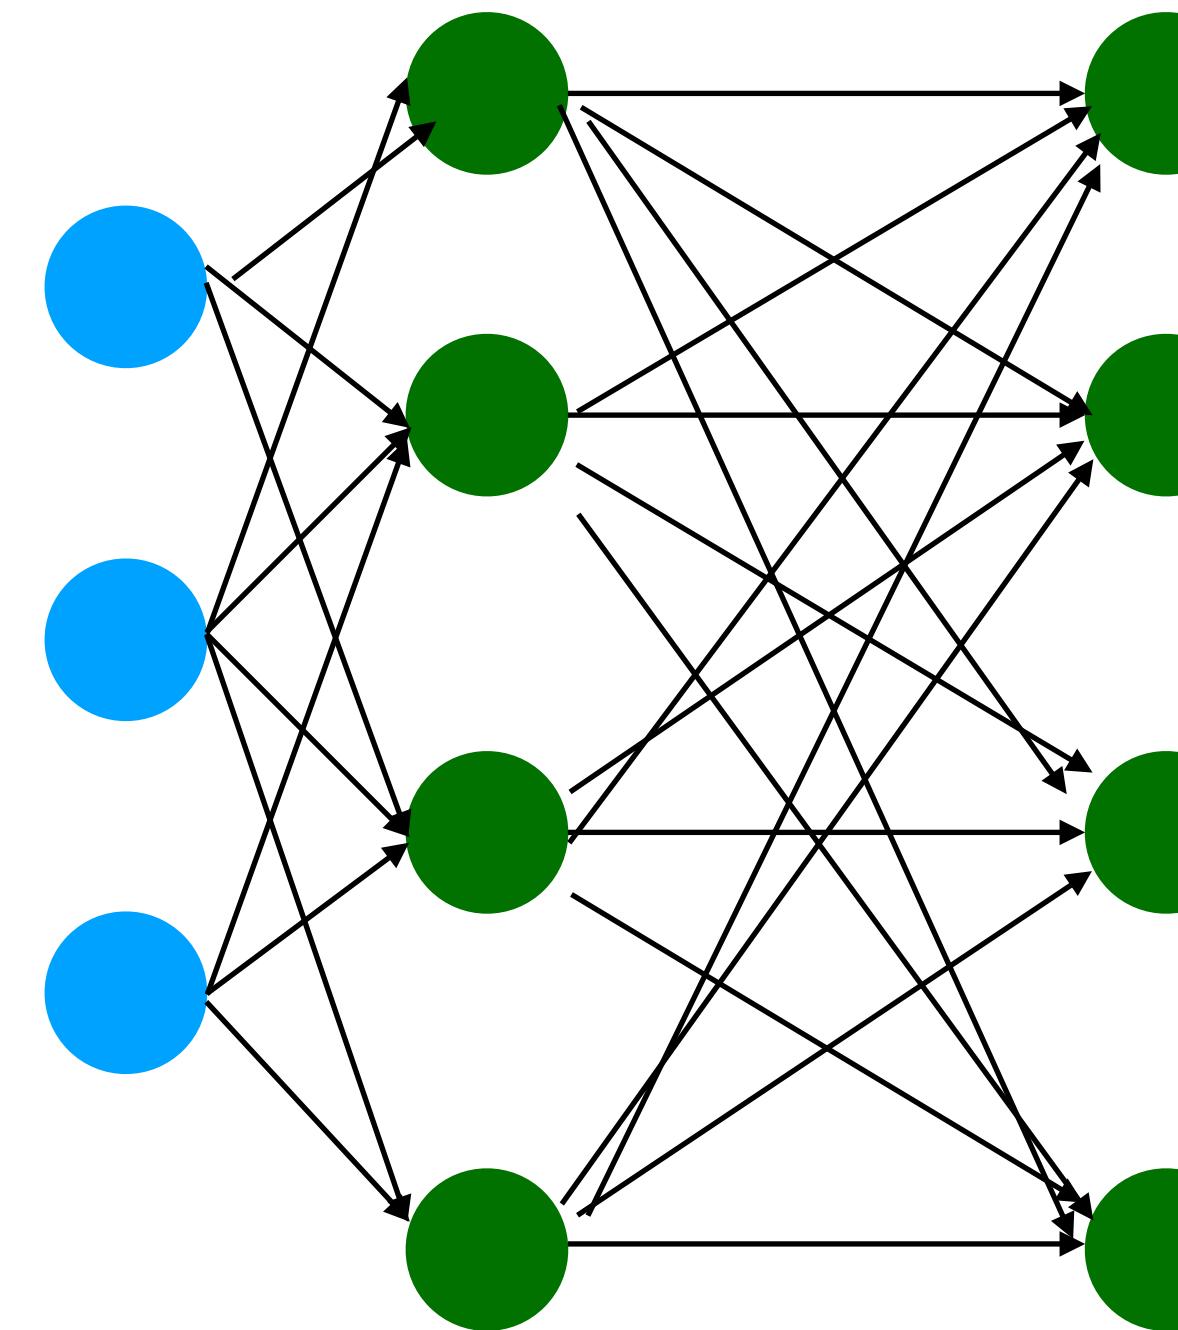


What will we code?

MNIST

A grid of 100 boxes arranged in 10 rows and 10 columns. Each box contains a handwritten digit from 0 to 9. The digits are written in a cursive or script-like style. The distribution of digits is as follows: Row 1: 0, 0, 0, 0, 0, 0, 0, 0, 0, 0; Row 2: 1, 1, 1, 1, 1, 1, 1, 1, 1, 1; Row 3: 2, 2, 2, 2, 2, 2, 2, 2, 2, 2; Row 4: 3, 3, 3, 3, 3, 3, 3, 3, 3, 3; Row 5: 4, 4, 4, 4, 4, 4, 4, 4, 4, 4; Row 6: 5, 5, 5, 5, 5, 5, 5, 5, 5, 5; Row 7: 6, 6, 6, 6, 6, 6, 6, 6, 6, 6; Row 8: 7, 7, 7, 7, 7, 7, 7, 7, 7, 7; Row 9: 8, 8, 8, 8, 8, 8, 8, 8, 8, 8; Row 10: 9, 9, 9, 9, 9, 9, 9, 9, 9, 9.

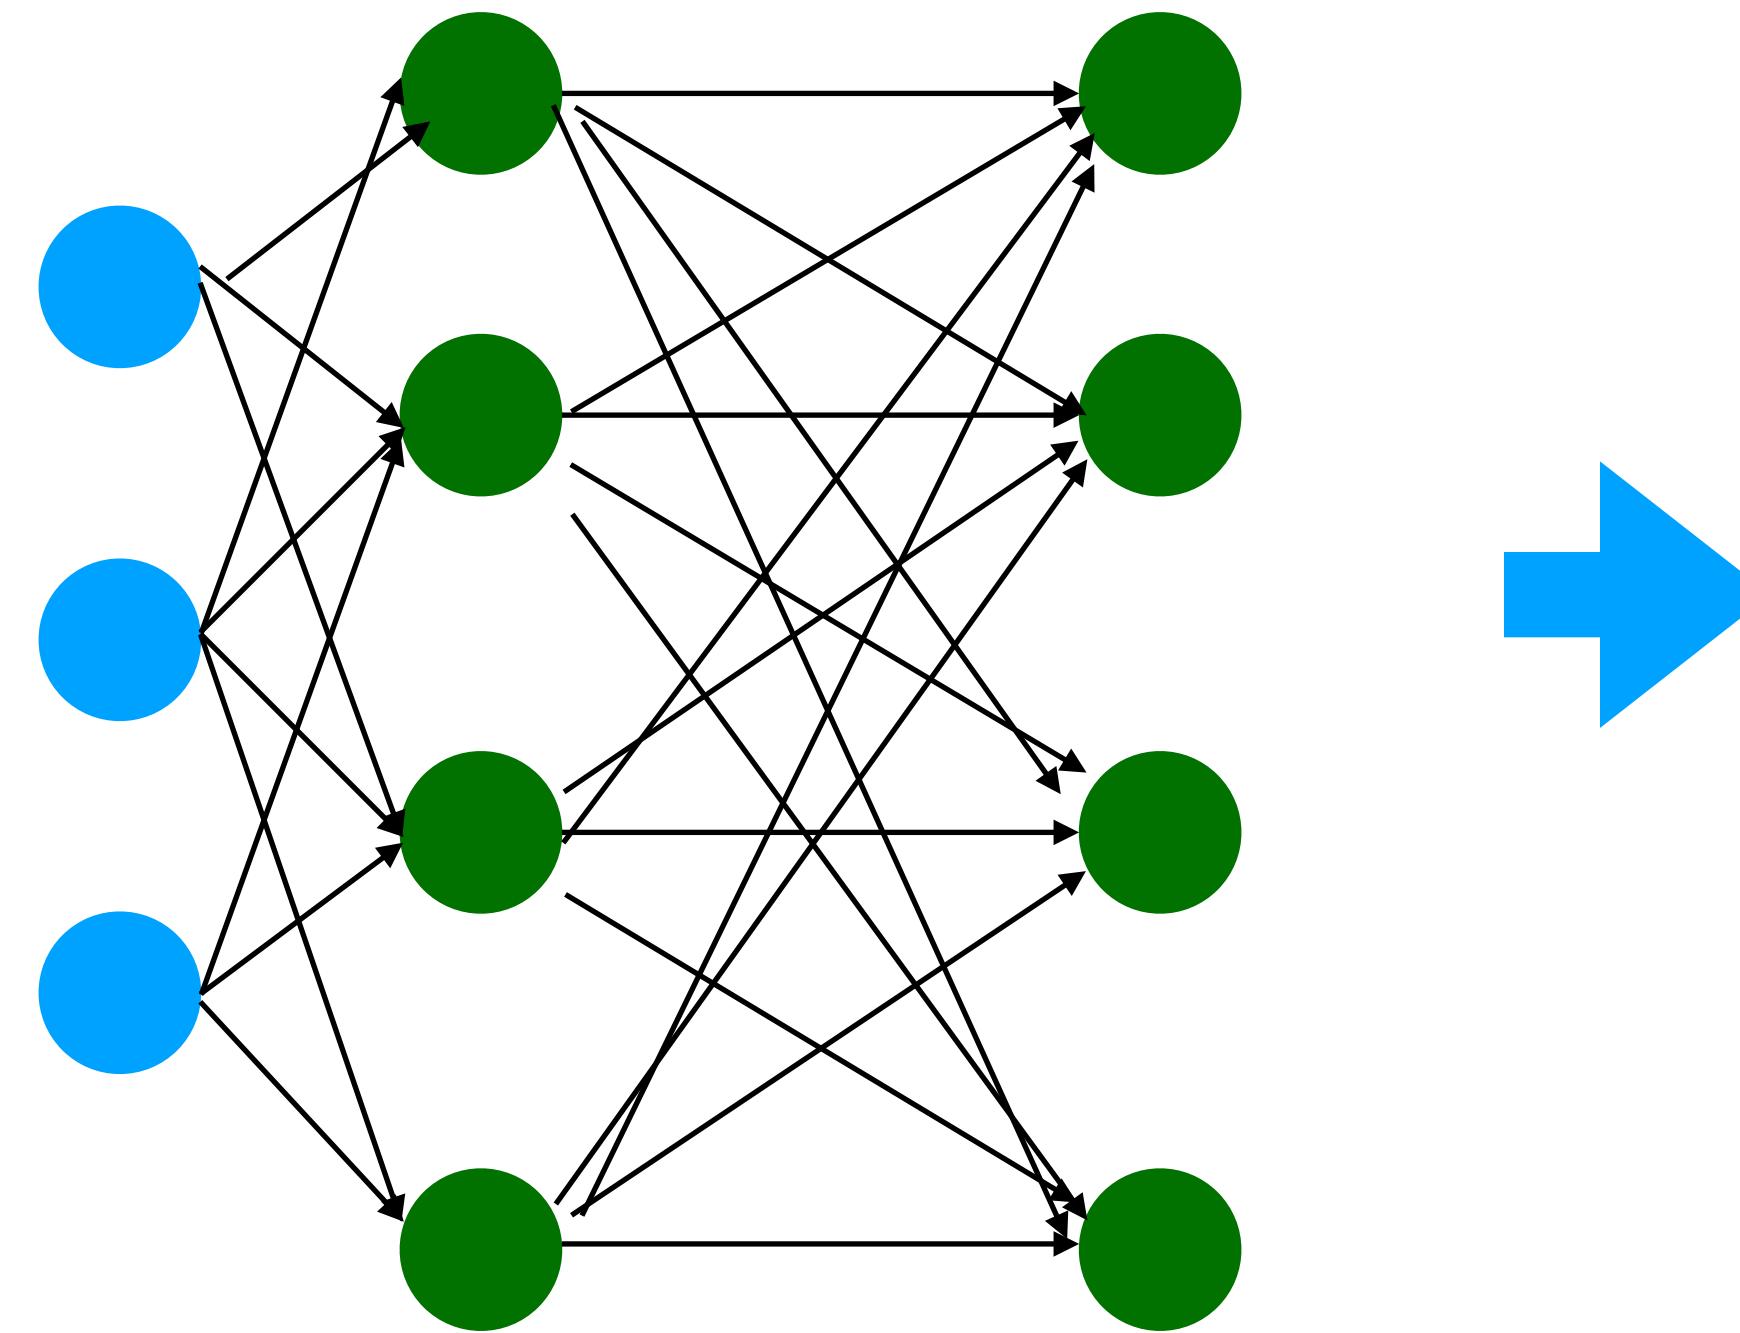
Neural Network 1 - hidden layer



What will we code?

MNIST

Neural Network 1 - hidden layer

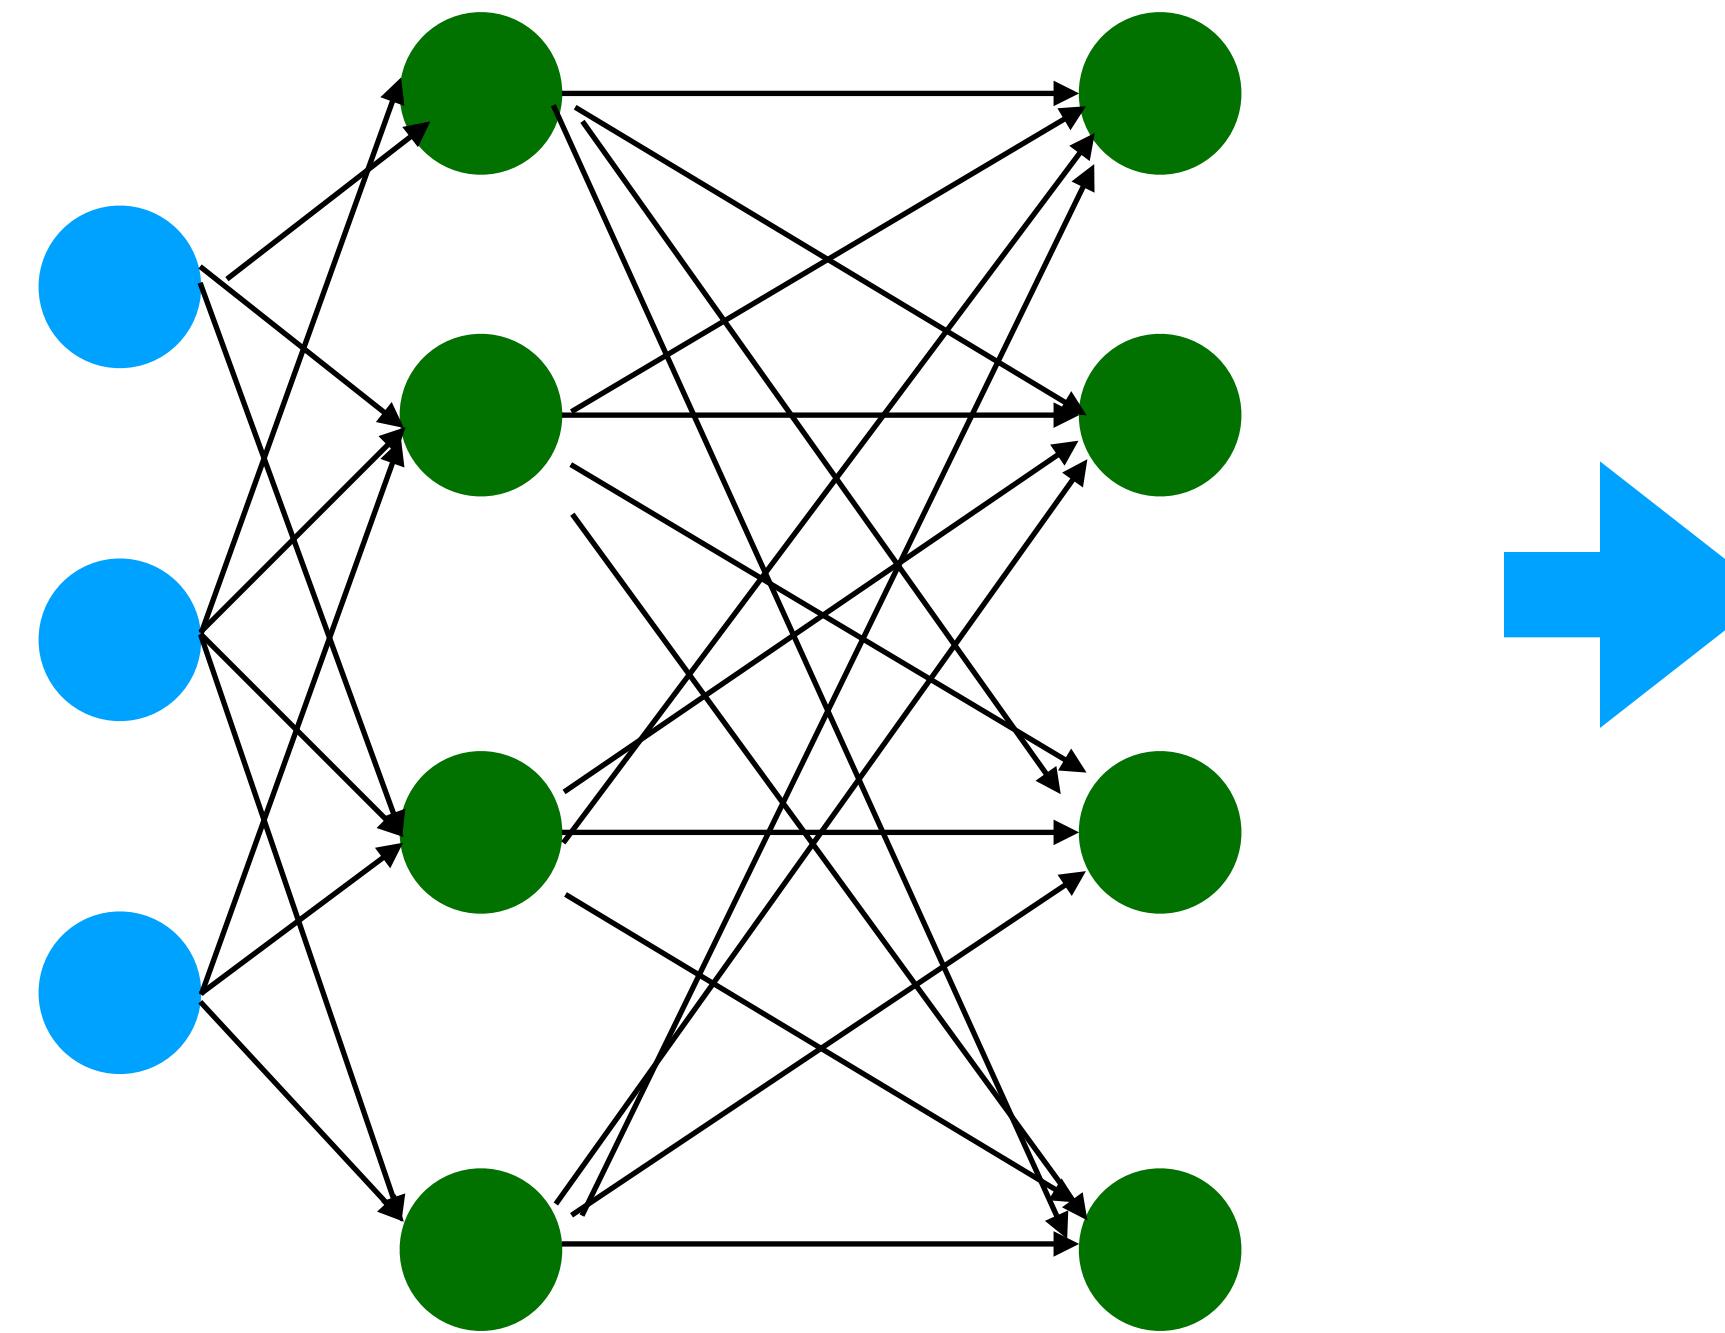


Automatically classify handwritten digits

What will we code?

MNIST

Neural Network 1 - hidden layer



Automatically classify handwritten digits

I want that you code, train and evaluate your first neural network !